

# **Elements of Machine Learning & Data Science**

Winter semester 2025/26

## **Lecture 3 – Decision Trees**

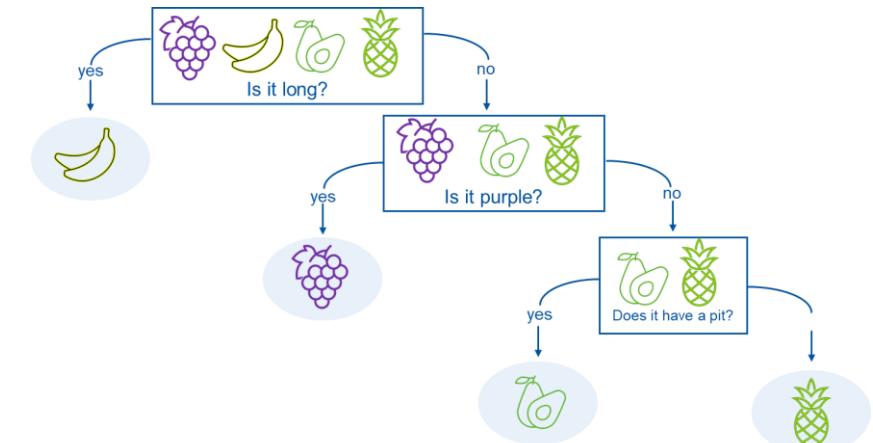
27.10.2025

Prof. Bastian Leibe

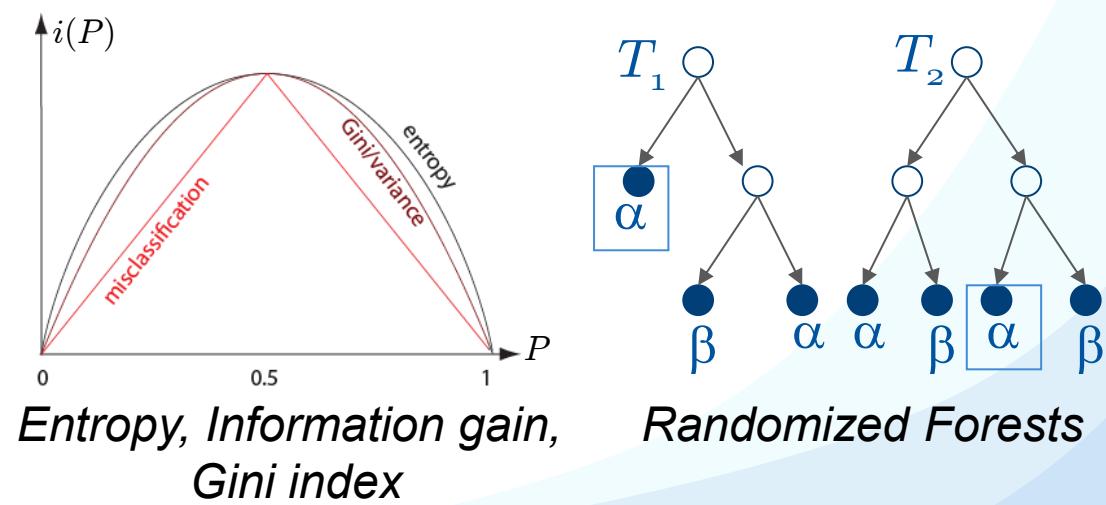
slides by Prof. Wil van der Aalst

# Overview of the Lecture Topics

1. Introduction to Data Science
- 2. Decision Trees**
3. Clustering
4. Frequent Itemsets
5. Association Rules
6. Time Series

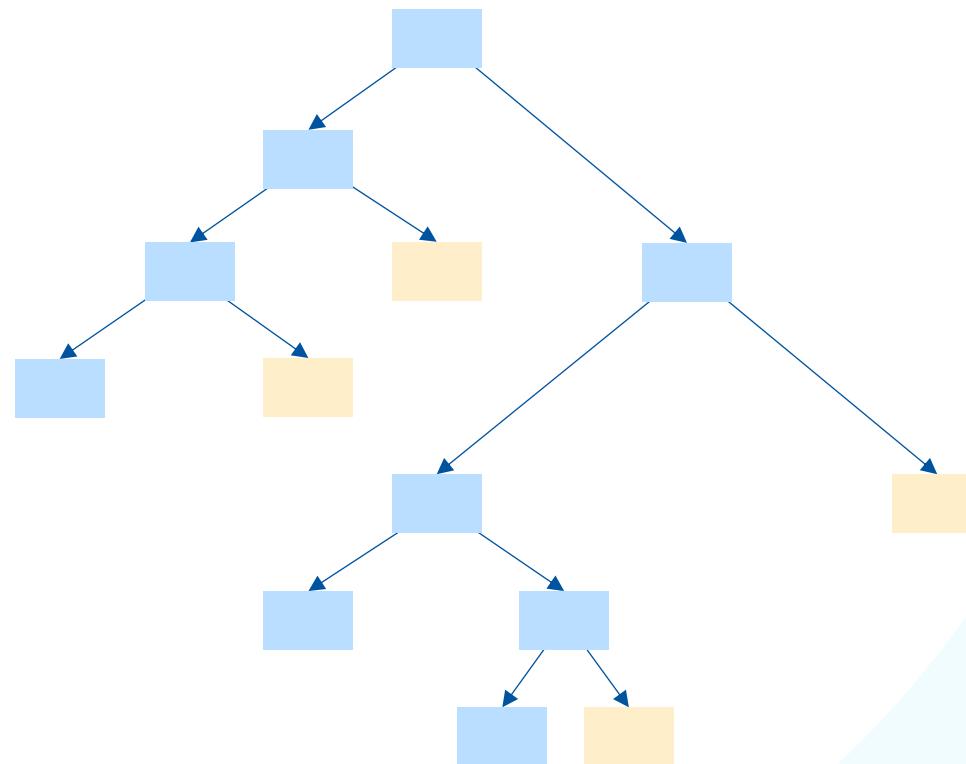


*Constructing Decision Trees*

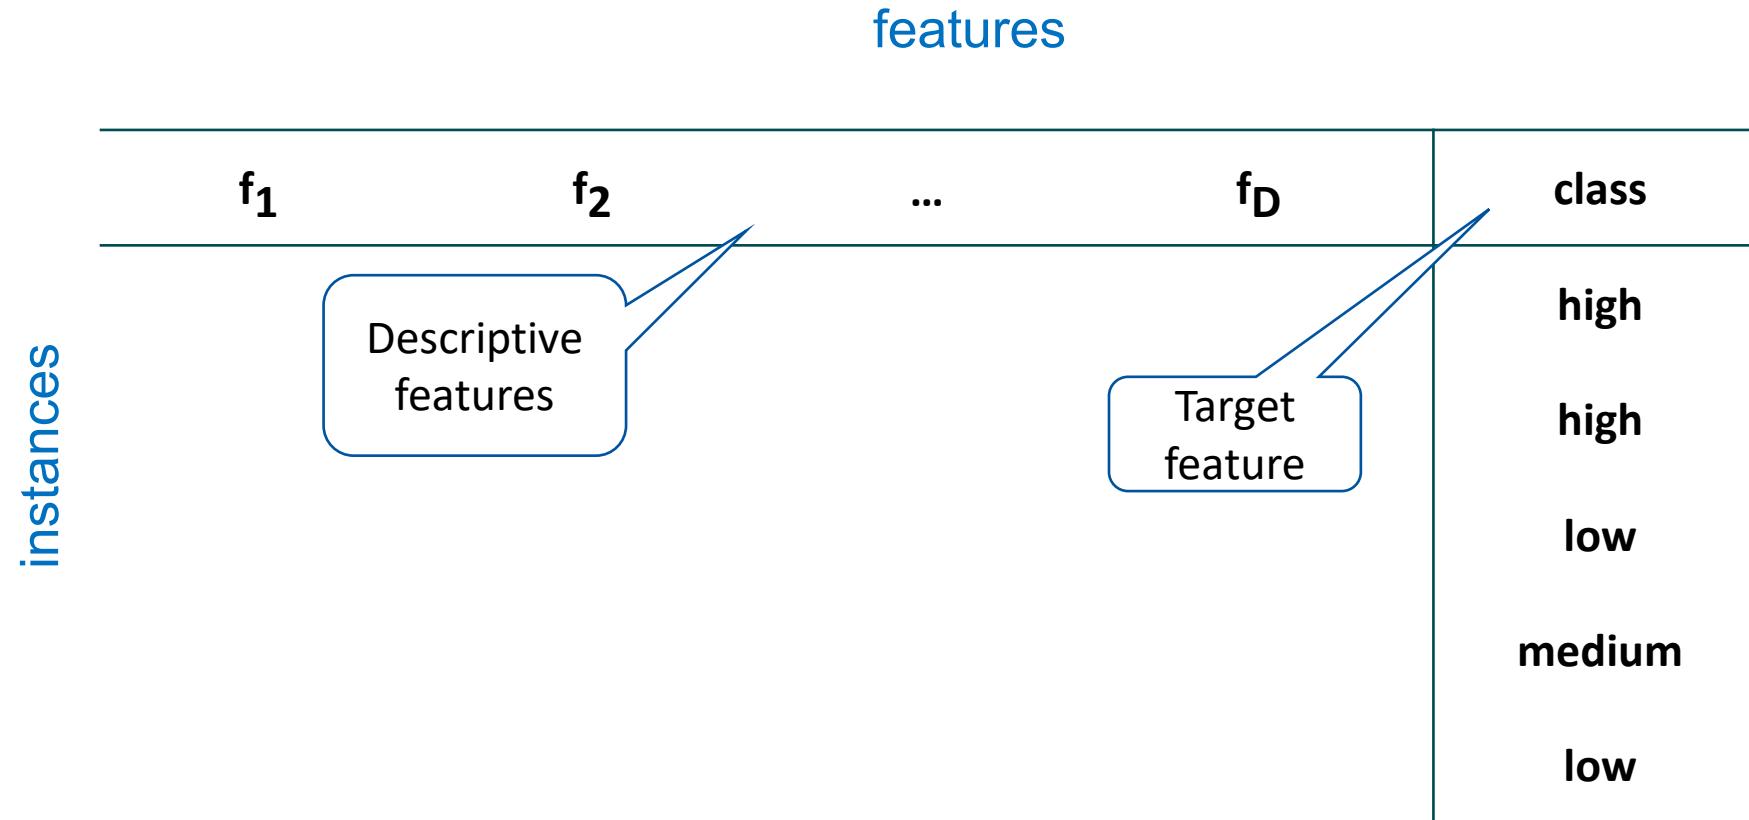


# Decision Trees

1. **Introduction to Decision Trees**
2. Entropy and Information Gain
3. ID3 Algorithm
4. Quantifying Information Gain
5. Pruning
6. Continuous Data
7. Ensembles

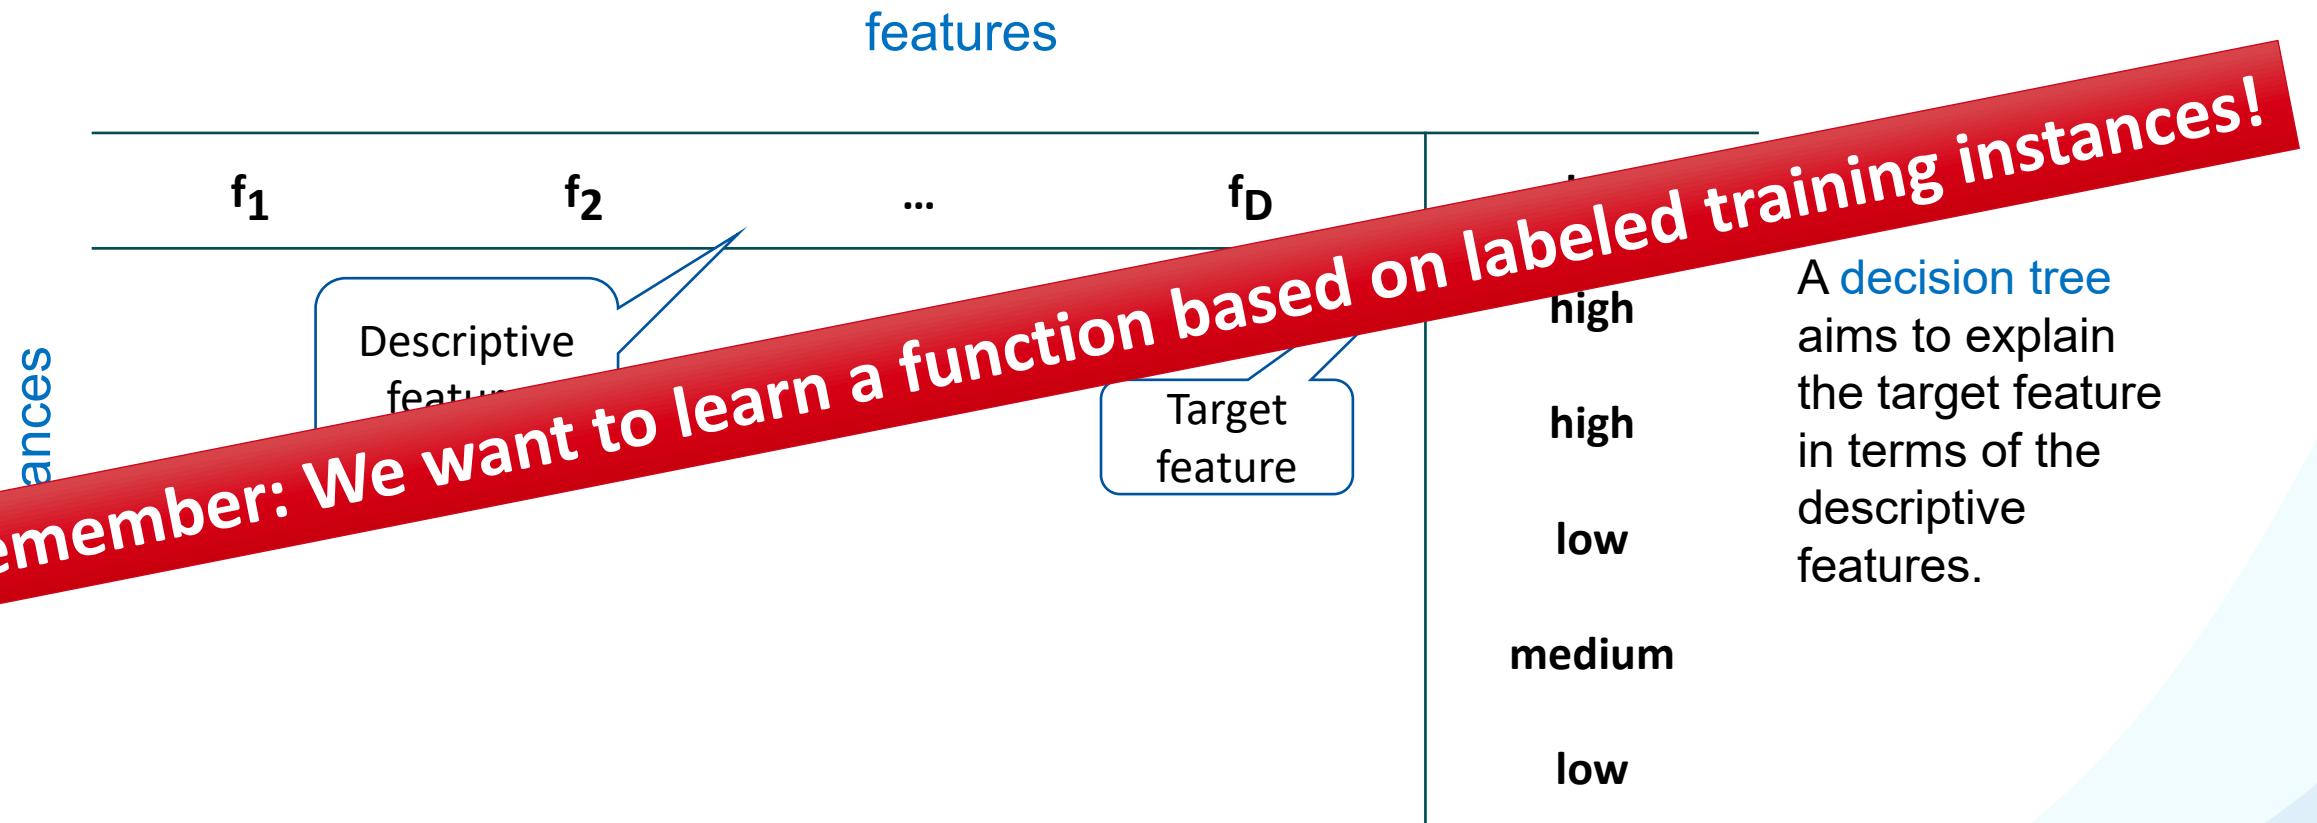


# Intuition and Interpretation

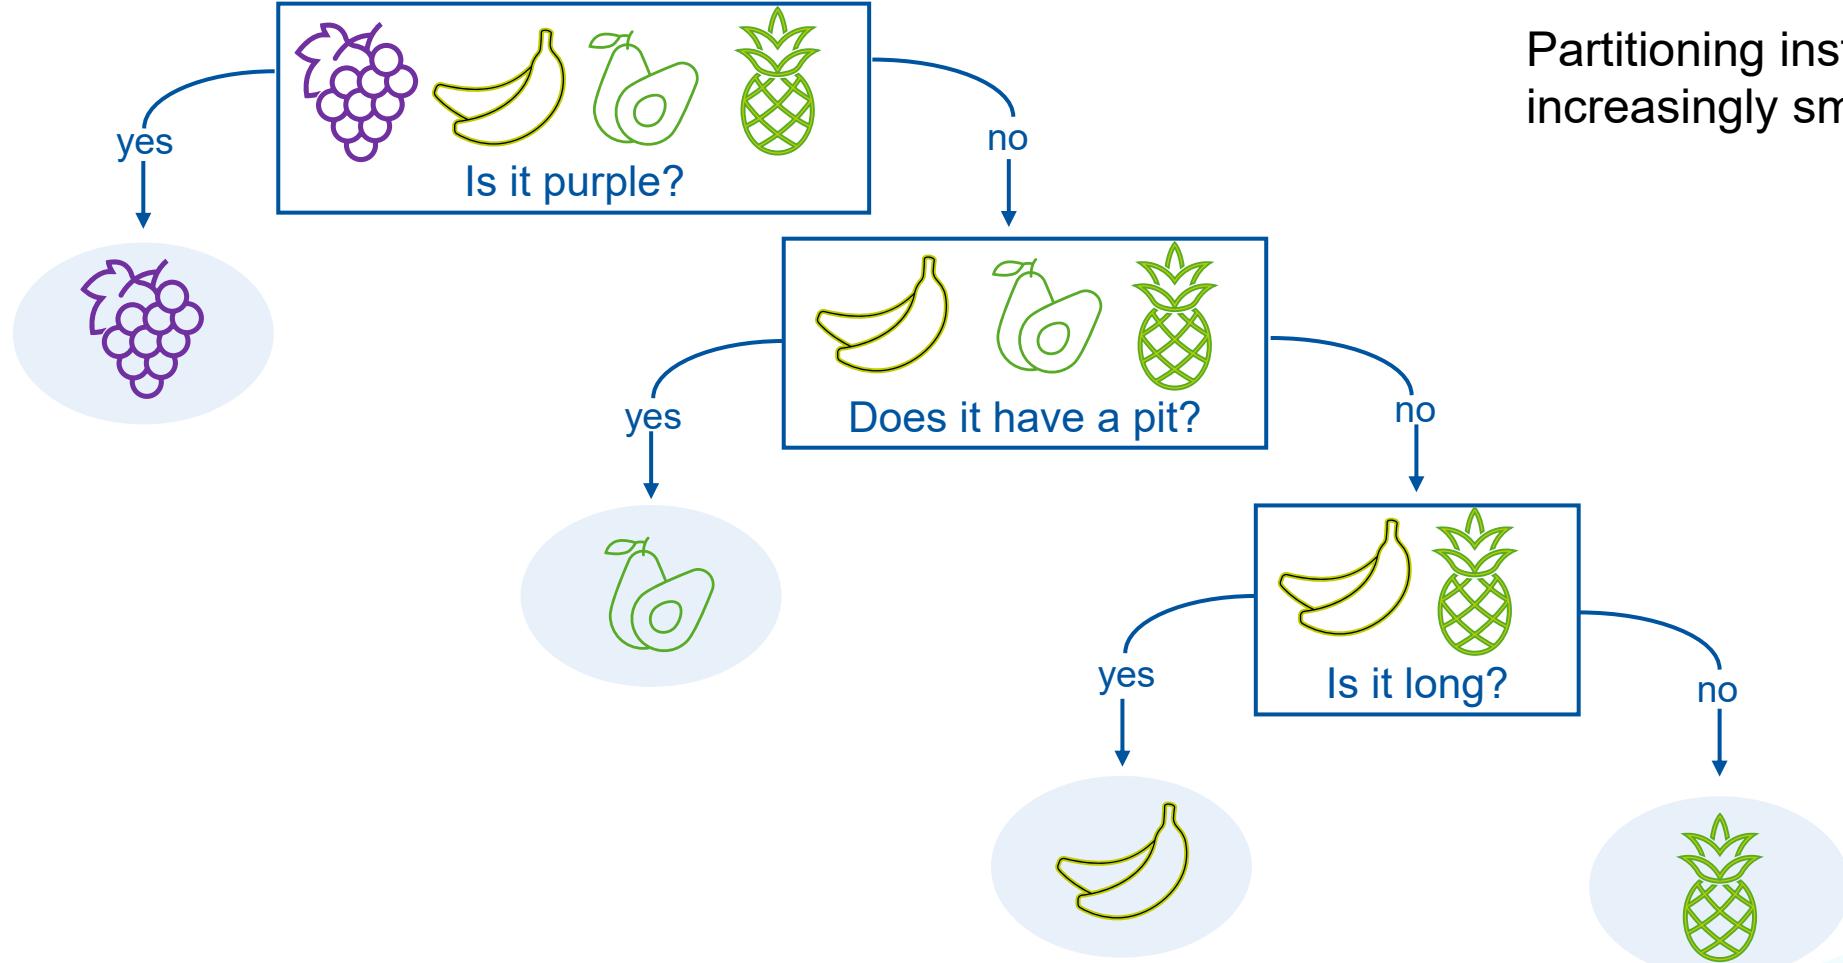


A decision tree aims to explain the target feature in terms of the descriptive features.

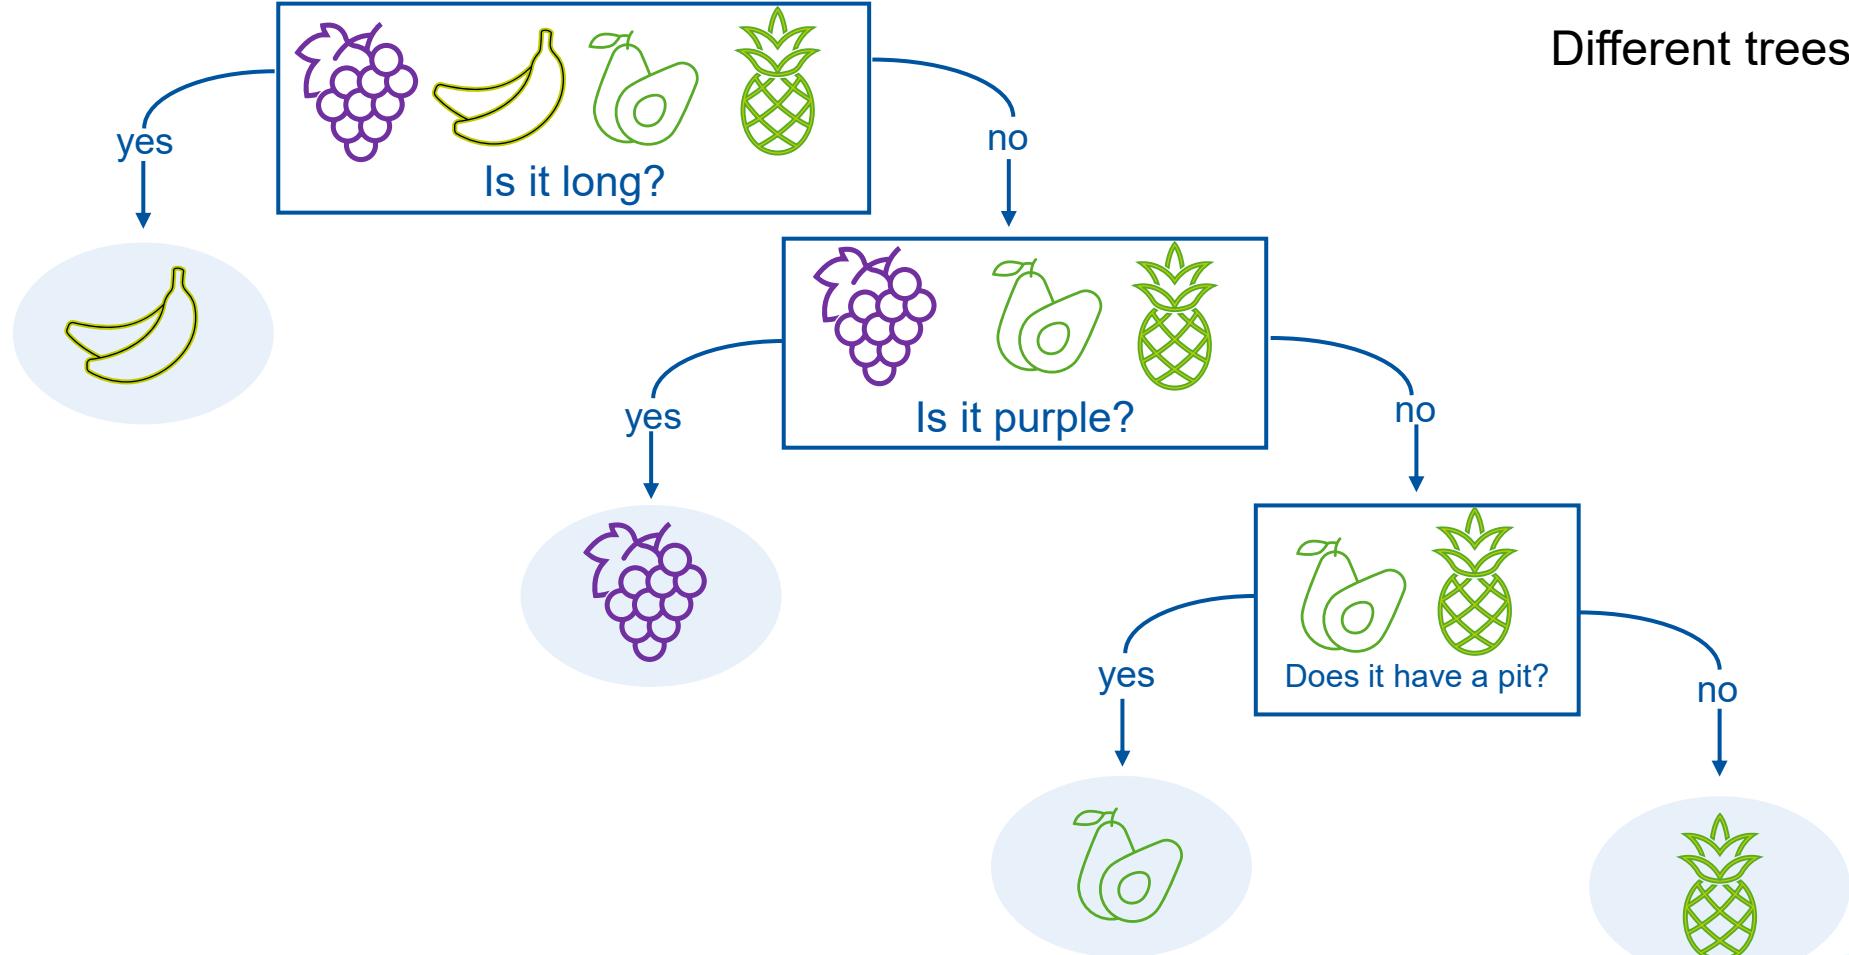
# Intuition and Interpretation



## Fruity Example



## Fruity Example



Different trees are possible

## Example 2

Rain	Wind	Temperature (°C)	Play tennis
Yes	Yes	15	No
No	No	34	Yes
Yes	No	23	Yes
Yes	Yes	20	Yes
No	Yes	28	No
...	...	...	...

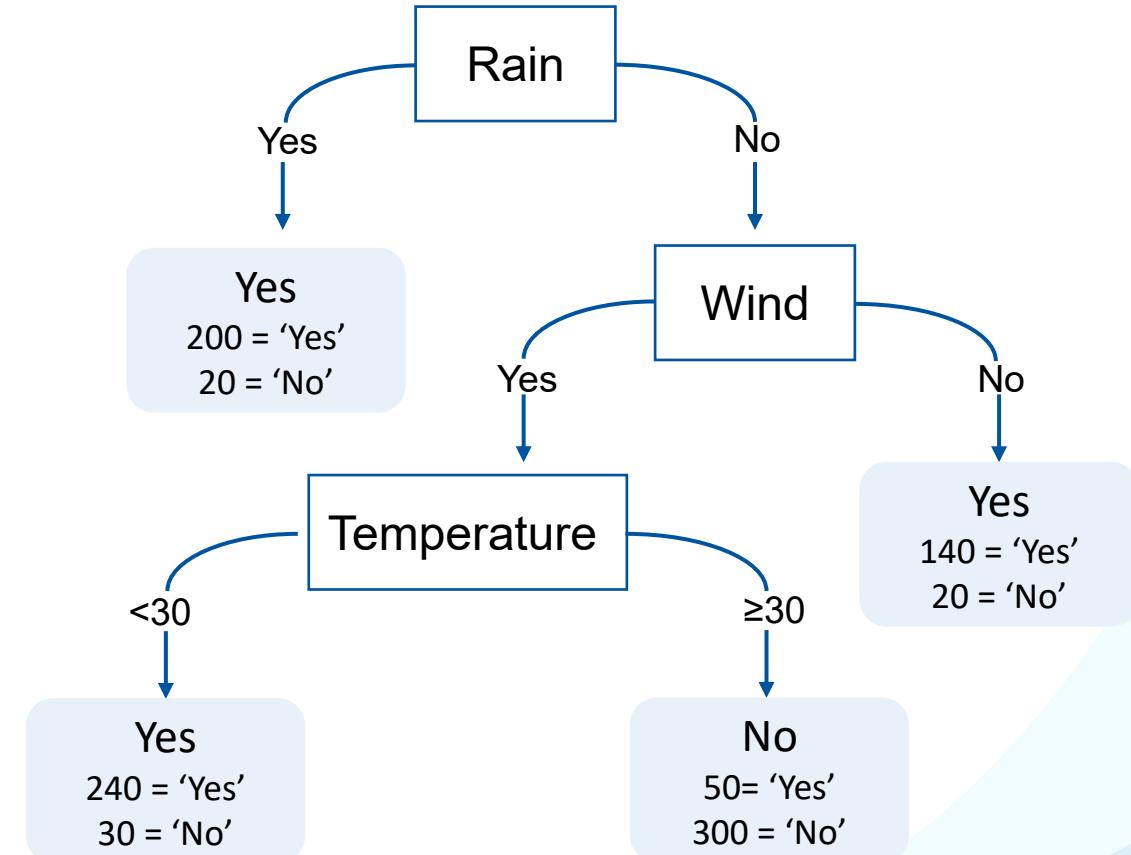
**Descriptive features**

**Target feature**

## Example 2

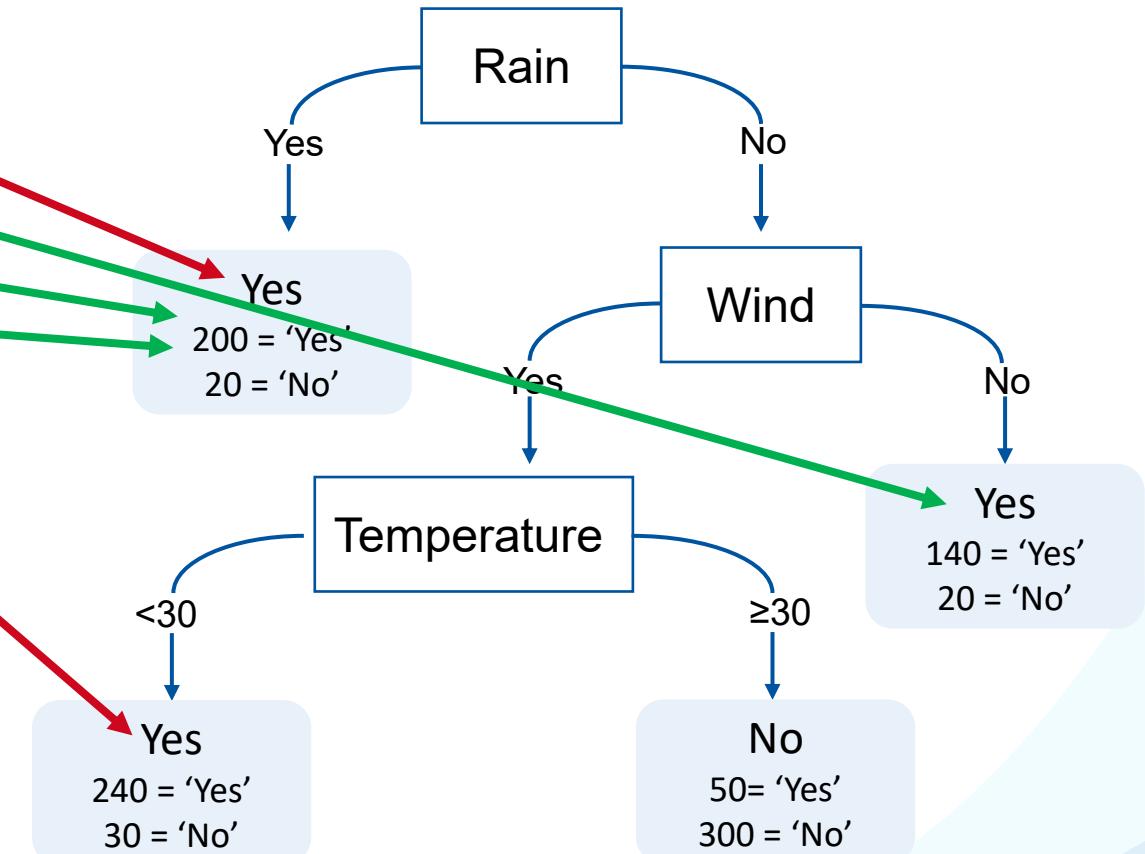
Rain	Wind	Temperature (°C)	Play tennis
Yes	Yes	15	No
No	No	34	Yes
Yes	No	23	Yes
Yes	Yes	20	Yes
No	Yes	28	No
...	...	...	...

1000 Instances



## Example 2

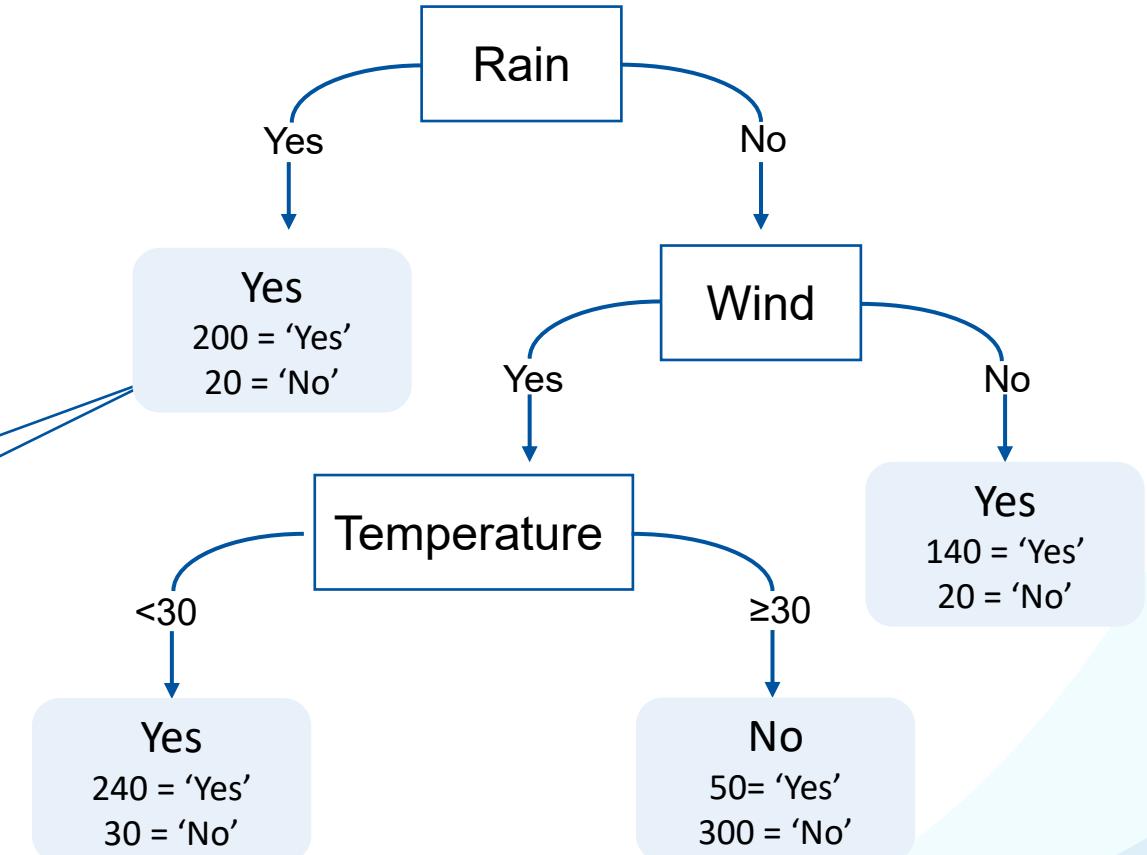
Rain	Wind	Temperature (°C)	Play tennis
Yes	Yes	15	No
No	No	34	Yes
Yes	No	23	Yes
Yes	Yes	20	Yes
No	Yes	28	No
...	...	...	...



## Example 2

Rain	Wind	Temperature (°C)	Play tennis
Yes	Yes	15	No
No	No	34	Yes
Yes	No	23	Yes
Yes	Yes	20	Yes
No	Yes	28	No
...	...	...	...

220 cases with *Rain = Yes* are classified as 'Yes' (Play tennis), but 20 are classified incorrectly



# Decision Tree Construction

## Tree Structure

- Three types of nodes: **root node**, **interior nodes** and **leaf nodes**
  - **Root node** refers to all instances
  - **Non-leaf nodes** **partition** the set of instances **based on a descriptive feature**
  - **Leaf nodes** have a label (target feature value)  
(usually based on the label of the majority of instances in this node)

# Decision Tree Construction

## Tree Structure

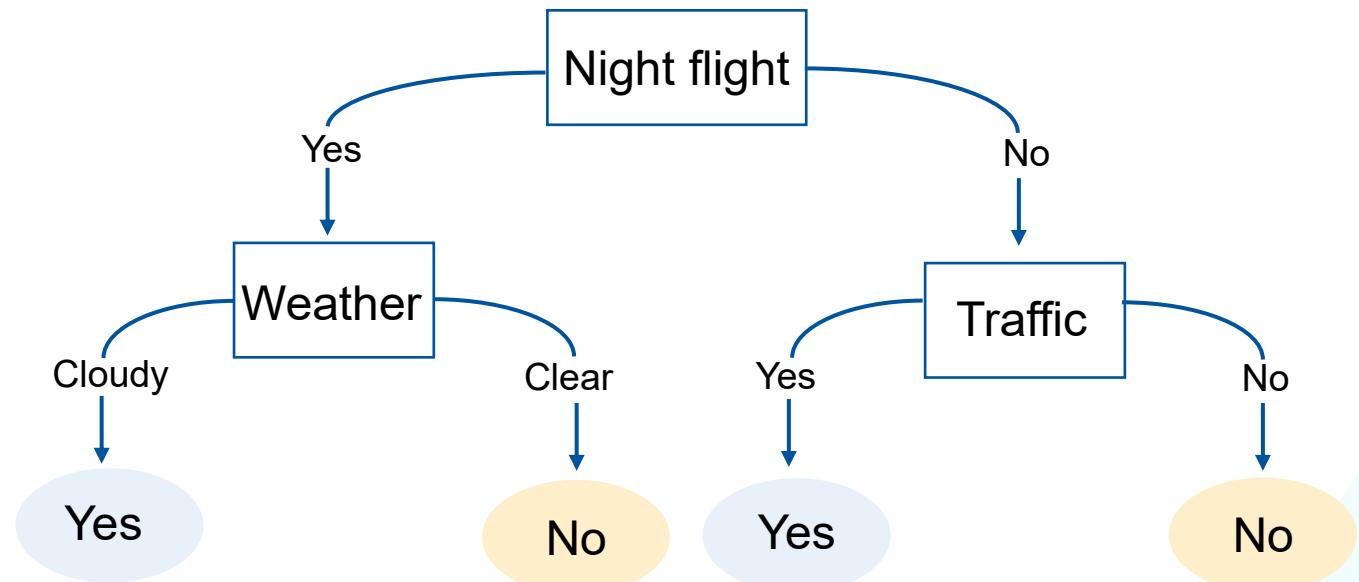
- Three types of nodes: **root node**, **interior nodes** and **leaf nodes**
  - **Root node** refers to all instances
  - **Non-leaf nodes** **partition** the set of instances **based on a descriptive feature**
  - **Leaf nodes** have a **label** (target feature value)  
(usually based on the label of the majority of instances in this node)

## There are two goals (often conflicting)

- The **tree is small and simple**
- The **leaves are homogeneous** in terms of the target feature

## Comparing Decision Trees (1/2)

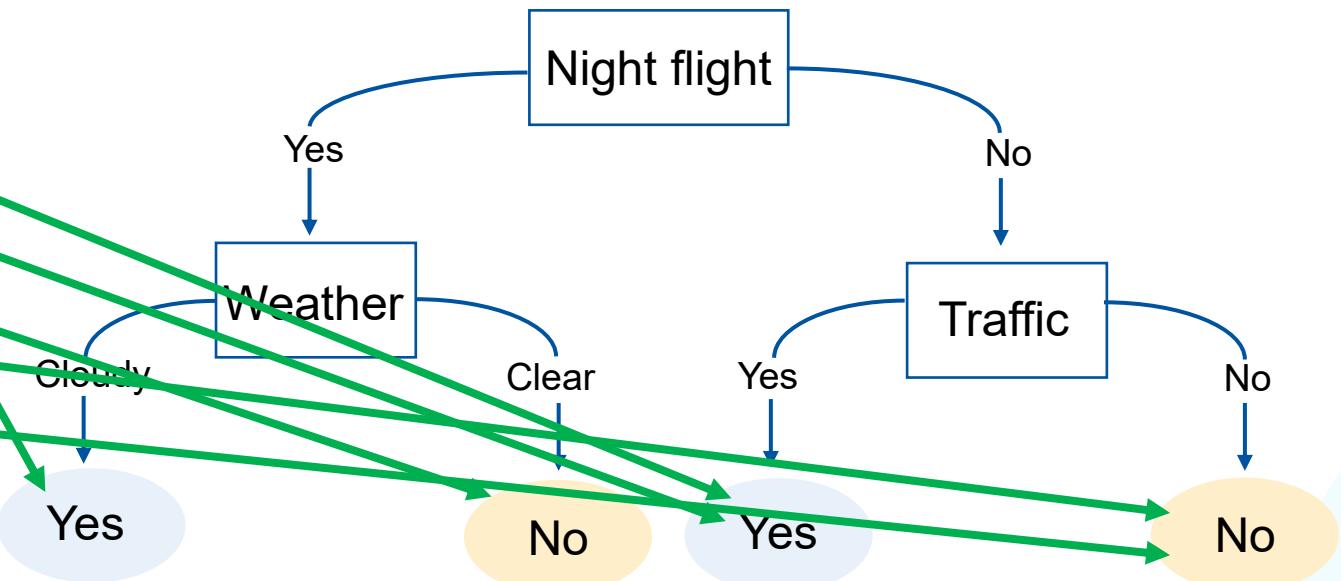
Weather	Traffic	Night flight	Flight delayed
Cloudy	No	Yes	Yes
Cloudy	Yes	No	Yes
Cloudy	Yes	No	Yes
Clear	Yes	Yes	No
Clear	No	No	No
Clear	No	No	No



# Comparing Decision Trees (1/2)

Weather	Traffic	Night flight	Flight delayed
Cloudy	No	Yes	Yes
Cloudy	Yes	No	Yes
Cloudy	Yes	No	Yes
Clear	Yes	Yes	No
Clear	No	No	No
Clear	No	No	No

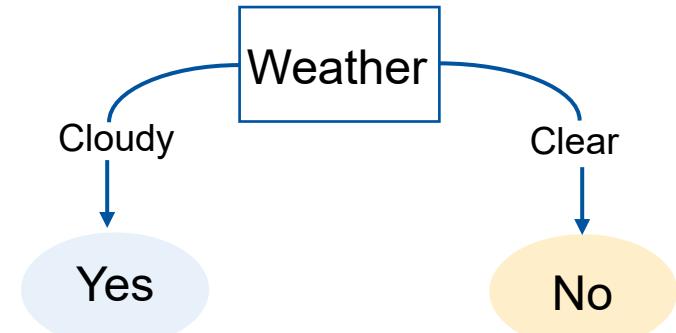
All instances correctly classified



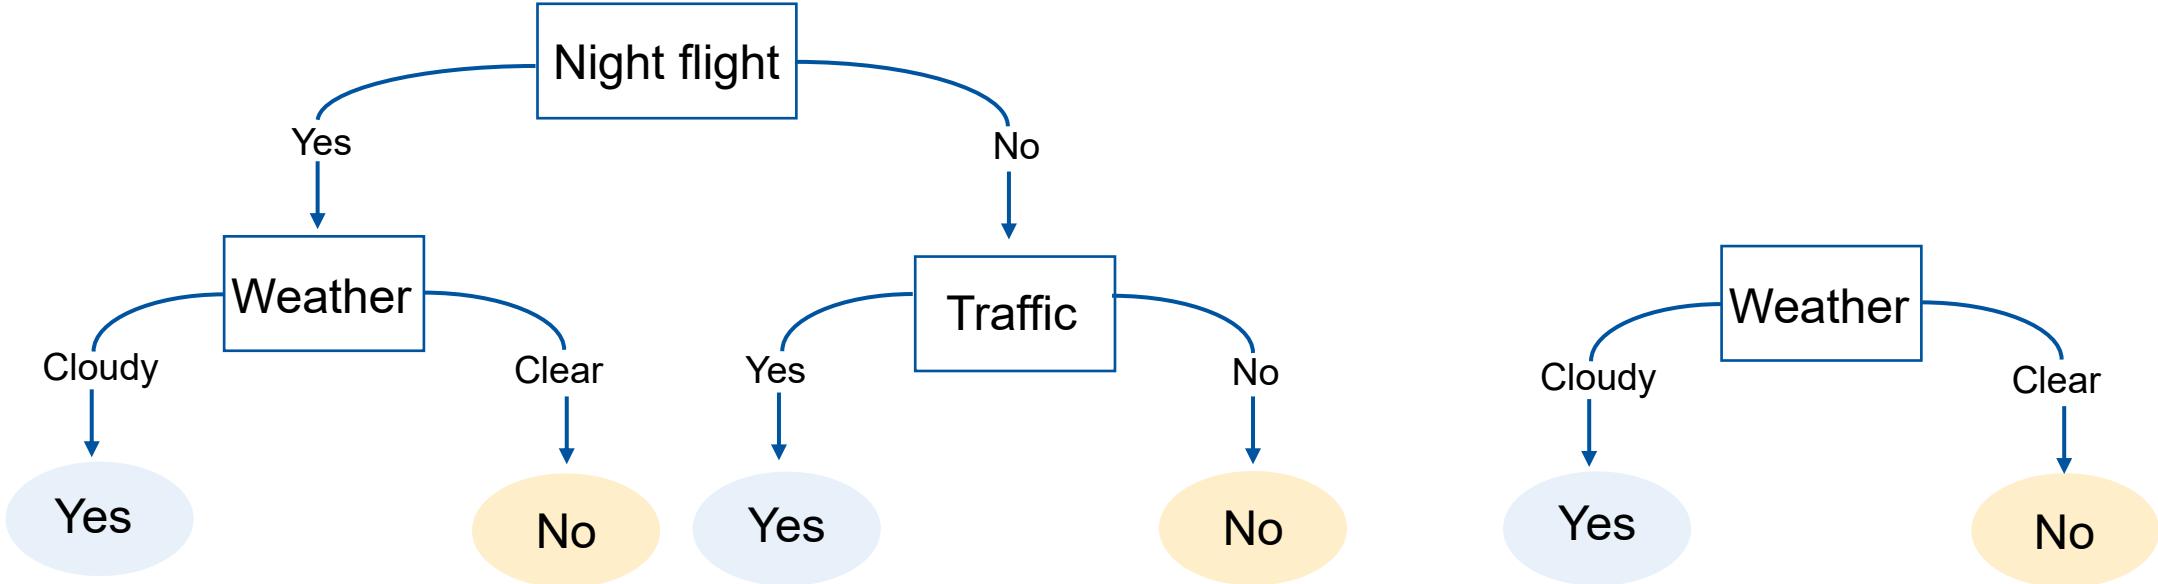
## Comparing Decision Trees (2/2)

Weather	Traffic	Night flight	Flight delayed
Cloudy	No	Yes	Yes
Cloudy	Yes	No	Yes
Cloudy	Yes	No	Yes
Clear	Yes	Yes	No
Clear	No	No	No
Clear	No	No	No

All instances correctly classified



## Comparing Decision Trees



Both trees correctly classify all observed instances, but the ‘simpler’ one seems ‘better’.

### Key concepts:

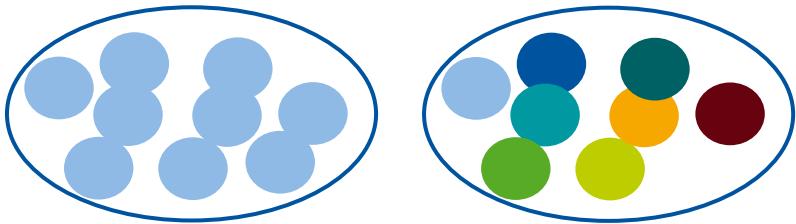
- avoid overfitting
- apply Occam's razor
- prefer shallow trees

## Characteristics Decision Trees

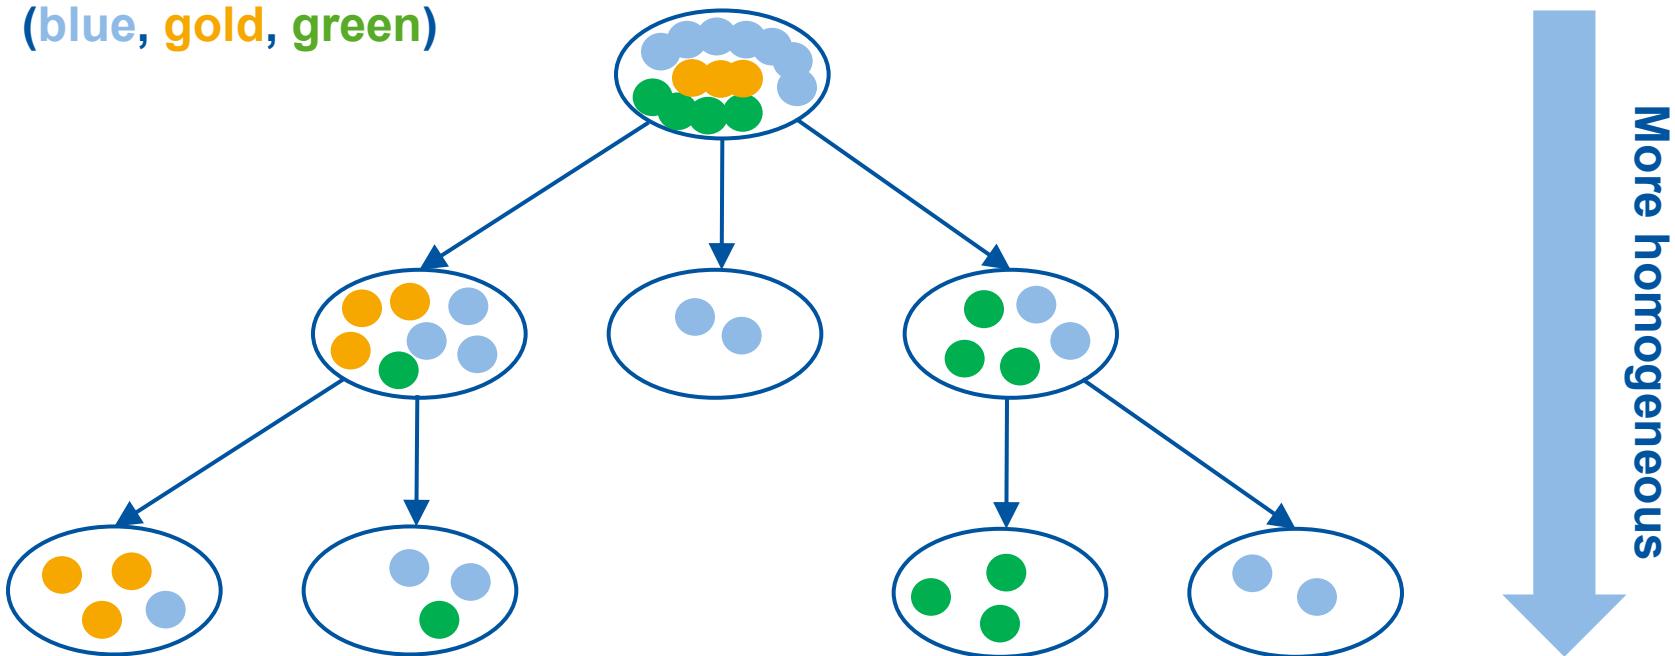
- A very simple model!
- In some cases, preferable to more complex and modern models (such as neural networks):
  - Fewer data points/attributes (managing [overfitting](#) is easier)
  - Well-suited for tabular data, where some attributes may be missing from table entries.
  - In domains where [explainability and transparency](#) are required
  - The choices of a tree are very easy to explain and show!
- There are [extensions of decision trees](#) that aim to combine simplicity and transparency with the ability to handle more complex data

# Decision Trees

1. Introduction to Decision Trees
2. **Entropy and Information Gain**
3. ID3 Algorithm
4. Quantifying Information Gain
5. Pruning
6. Continuous Data
7. Ensembles

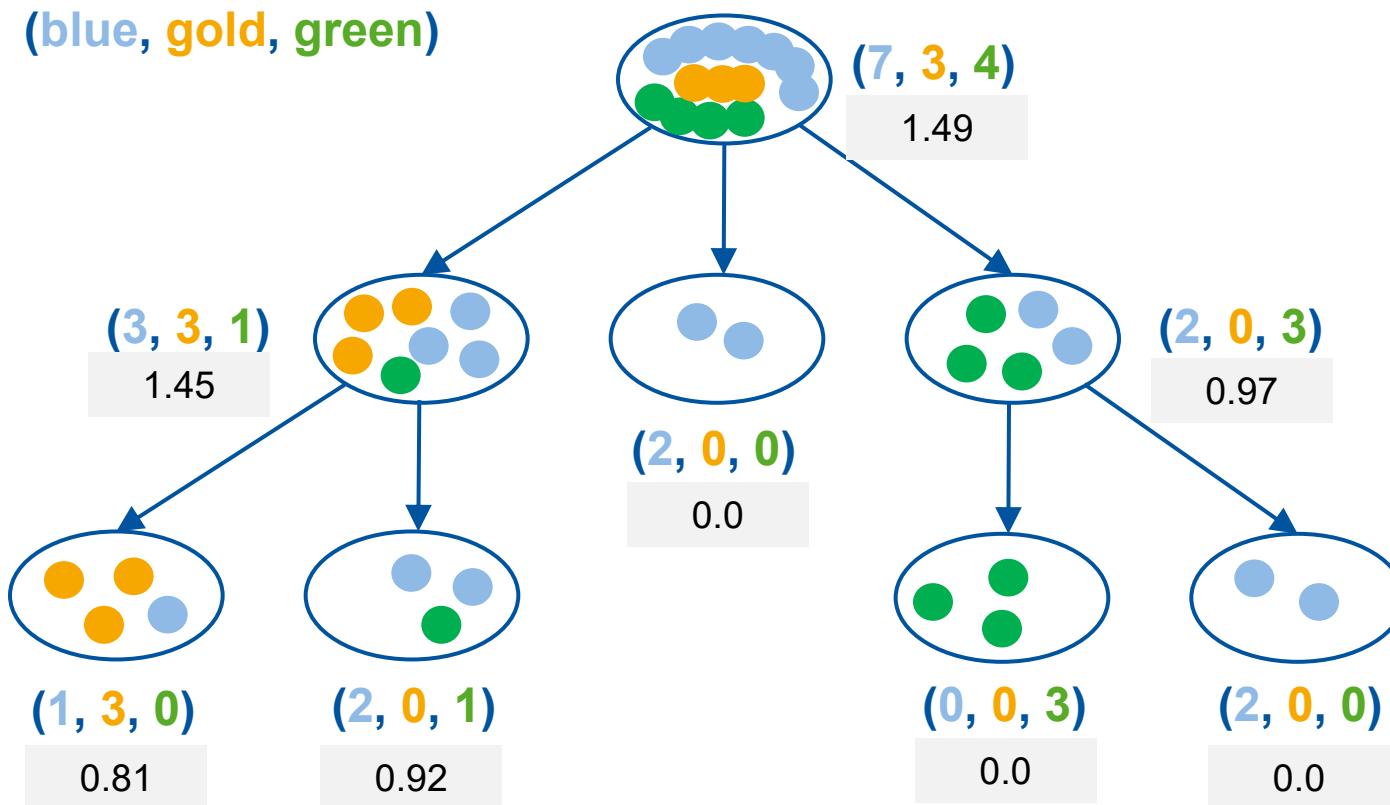


# Information Gain



Information gain = improvement in knowledge  
(predictability of target label in nodes)

# Entropy - Intuition



## Idea

- Measure of impurity
- Uncertainty when guessing
- Incompressibility

Worst case entropy for 3 values:  $\approx 1.58$

# Entropy - Formula

$$H(t) = - \sum_{k=1}^K (P(t = k) \cdot \log_s(P(t = k)))$$



$$H(\text{color}) = -\left(\frac{7}{14} \cdot \log_2\left(\frac{7}{14}\right) + \frac{3}{14} \cdot \log_2\left(\frac{3}{14}\right) + \frac{4}{14} \cdot \log_2\left(\frac{4}{14}\right)\right) \approx 1.49$$



$t$ : examined target feature ( $\text{color}$  in the example)



$K$ : number of possible values of the target feature ( $K = |\{\text{blue, gold, green}\}| = 3$  in the example)

$P(t = k) \in [0, 1]$ : probability that a random value in  $t$  equals the  $k$ th value in the set of possible values

$s$ : logarithm base (we use  $s = 2$  by convention)

## Entropy - Example

$$H(t) = - \sum_{k=1}^K (P(t = k) \cdot \log_s(P(t = k)))$$



$$H(\text{color}) = -\left(\frac{2}{5} \cdot \log_2\left(\frac{2}{5}\right) + \frac{0}{5} \cdot \log_2\left(\frac{0}{5}\right) + \frac{3}{5} \cdot \log_2\left(\frac{3}{5}\right)\right) \approx 0.97$$

## Entropy - Example

$$H(t) = - \sum_{k=1}^K (P(t = k) \cdot \log_s(P(t = k)))$$



$$H(\text{color}) = -\left(\frac{2}{5} \cdot \log_2\left(\frac{2}{5}\right) + \underbrace{\frac{0}{5} \cdot \log_2\left(\frac{0}{5}\right)}_{\text{Interpreted as 0}} + \frac{3}{5} \cdot \log_2\left(\frac{3}{5}\right)\right) \approx 0.97$$



$$H(\text{color}) = -\left(\frac{0}{3} \cdot \log_2\left(\frac{0}{3}\right) + \frac{0}{3} \cdot \log_2\left(\frac{0}{3}\right) + \frac{3}{3} \cdot \log_2\left(\frac{3}{3}\right)\right) = 0$$

## Questions

Suppose that we have  $K$  possible values (colors) and  $N$  instances (balls).



$$H(t) = - \sum_{k=1}^K (P(t = k) \cdot \log_s(P(t = k)))$$

What distribution of the  $N$  instances over the  $K$  possible values yields the lowest entropy?

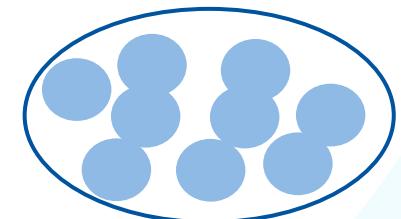
# Questions

Suppose that we have  $K$  possible values (colors) and  $N$  instances (balls).



$$H(t) = - \sum_{k=1}^K (P(t = k) \cdot \log_s(P(t = k)))$$

What distribution of the  $N$  instances over the  $K$  possible values yields the lowest entropy?



$$H(\text{color}) = -(1 \cdot \log_2(1)) = 0$$

→ all instances have the same value

## Questions

Suppose that we have  $K$  possible values (colors) and  $N$  instances (balls).



$$H(t) = - \sum_{k=1}^K (P(t=k) \cdot \log_s(P(t=k)))$$

What distribution of the  $N$  instances over the  $K$  possible values yields the **highest entropy**?

# Questions

Suppose that we have  $K$  possible values (colors) and  $N$  instances (balls).



$$H(t) = - \sum_{k=1}^K (P(t=k) \cdot \log_s(P(t=k)))$$

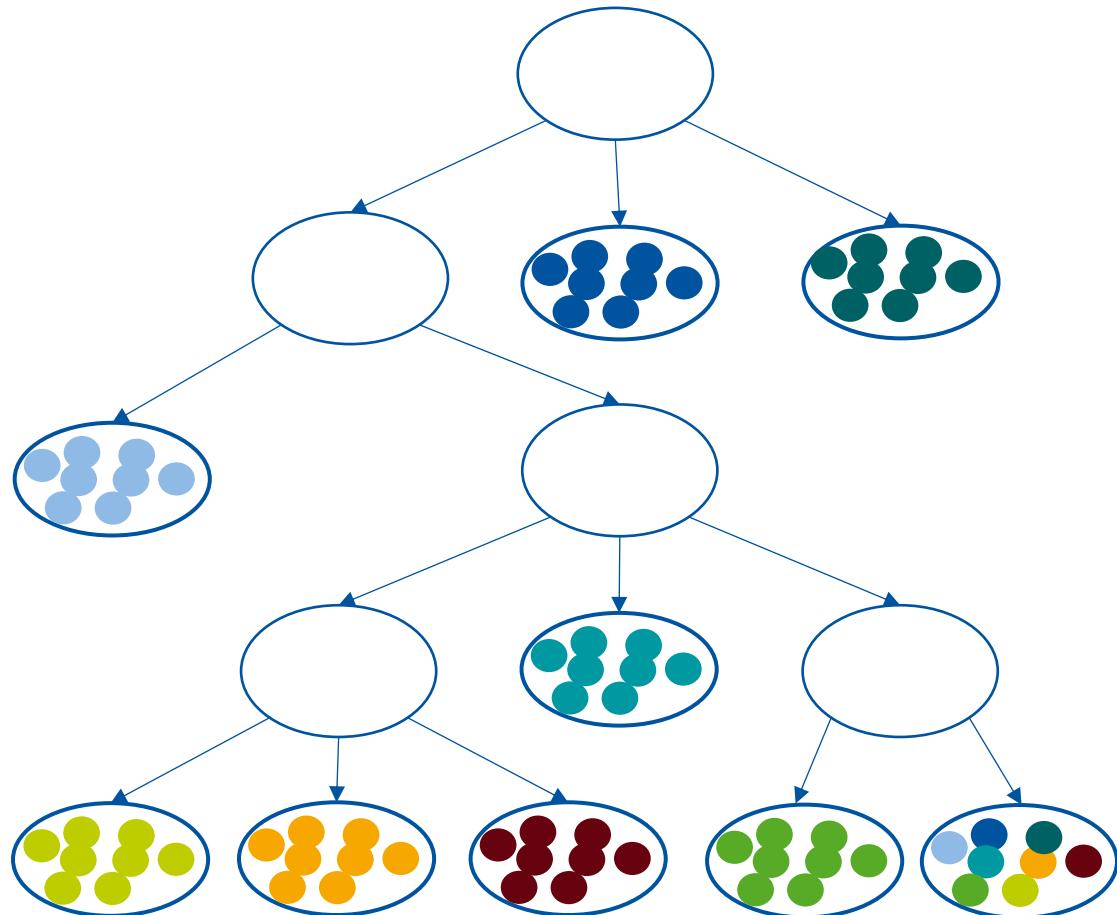
What distribution of the  $N$  instances over the  $K$  possible values yields the highest entropy?

→ Even distribution over all possible values

$$\begin{aligned} H(\text{color}) &= - \sum_{k=1}^K \left( \frac{1}{K} \cdot \log_2 \left( \frac{1}{K} \right) \right) \\ &= - \left( K \cdot \frac{1}{K} \cdot \log_2 \left( \frac{1}{K} \right) \right) \\ &= - \log_2 \left( \frac{1}{K} \right) \\ &= \log_2(K) \end{aligned}$$



## Overall Entropy

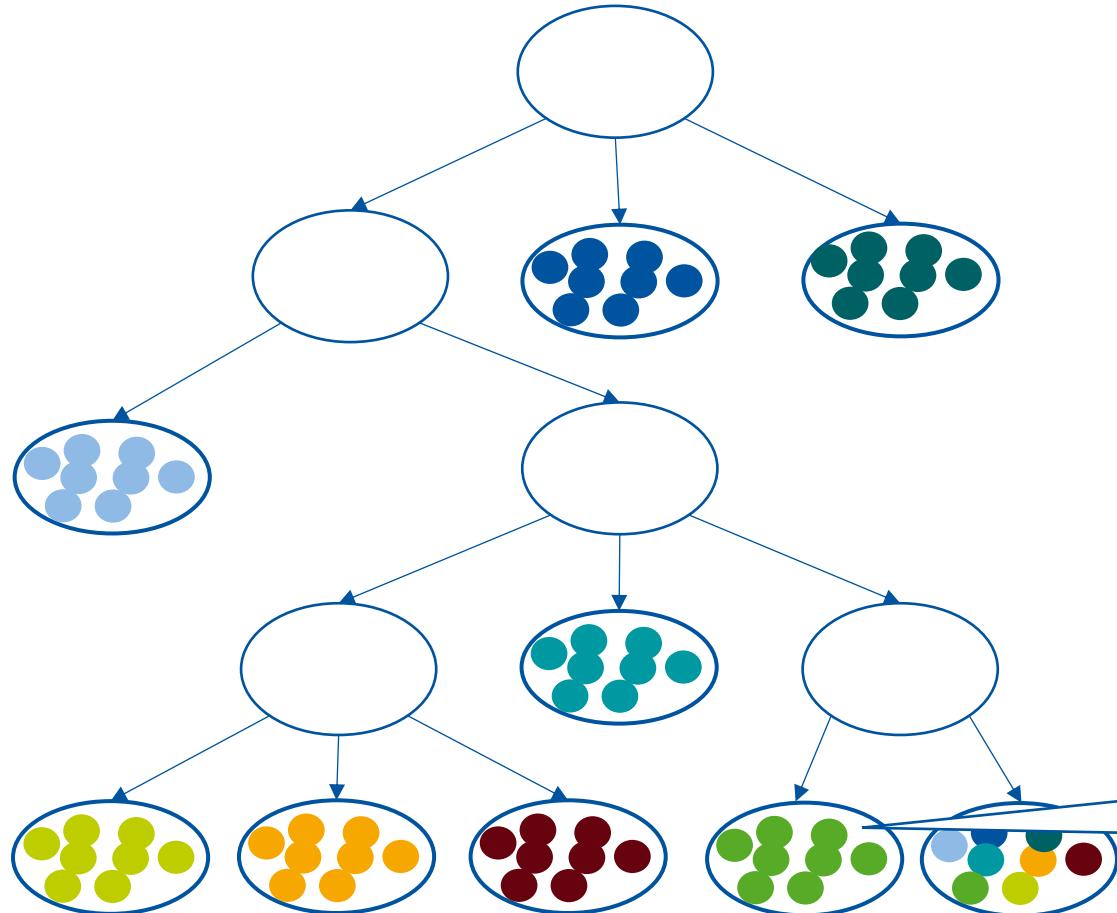


Overall entropy  $H_W$  is the weighted average of the individual entropies:

$$H_W(t) = \sum_{node \in nodes} \left( \frac{|node|}{N} \cdot H^{node}(t) \right)$$

Example:  $N = 72, K = 8$

# Overall Entropy



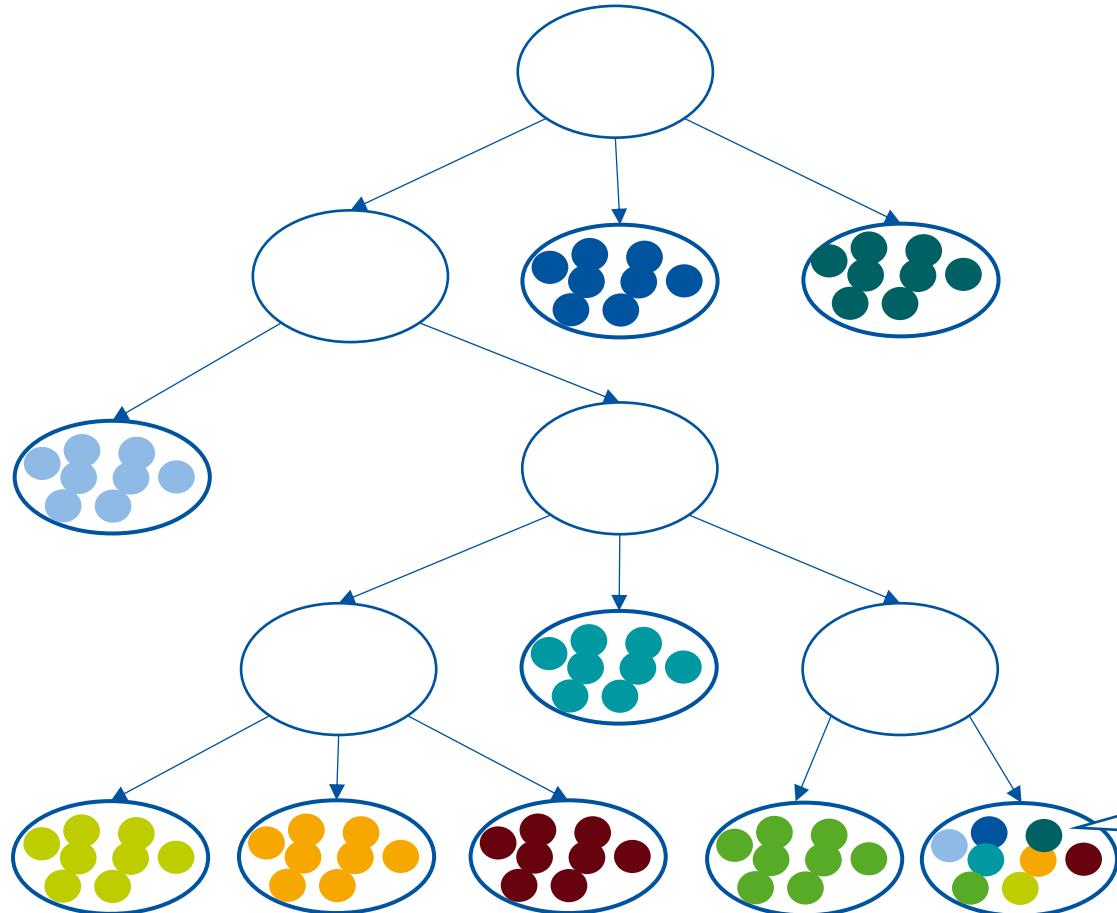
Overall entropy  $H_W$  is the weighted average of the individual entropies:

$$H_W(t) = \sum_{node \in nodes} \left( \frac{|node|}{N} \cdot H^{node}(t) \right)$$

Example:  $N = 72, K = 8$

8 homogeneously colored balls:  
 $H^{node}(\text{color}) = -\left(\frac{8}{8} \cdot \log_2\left(\frac{8}{8}\right)\right) = 0$

# Overall Entropy



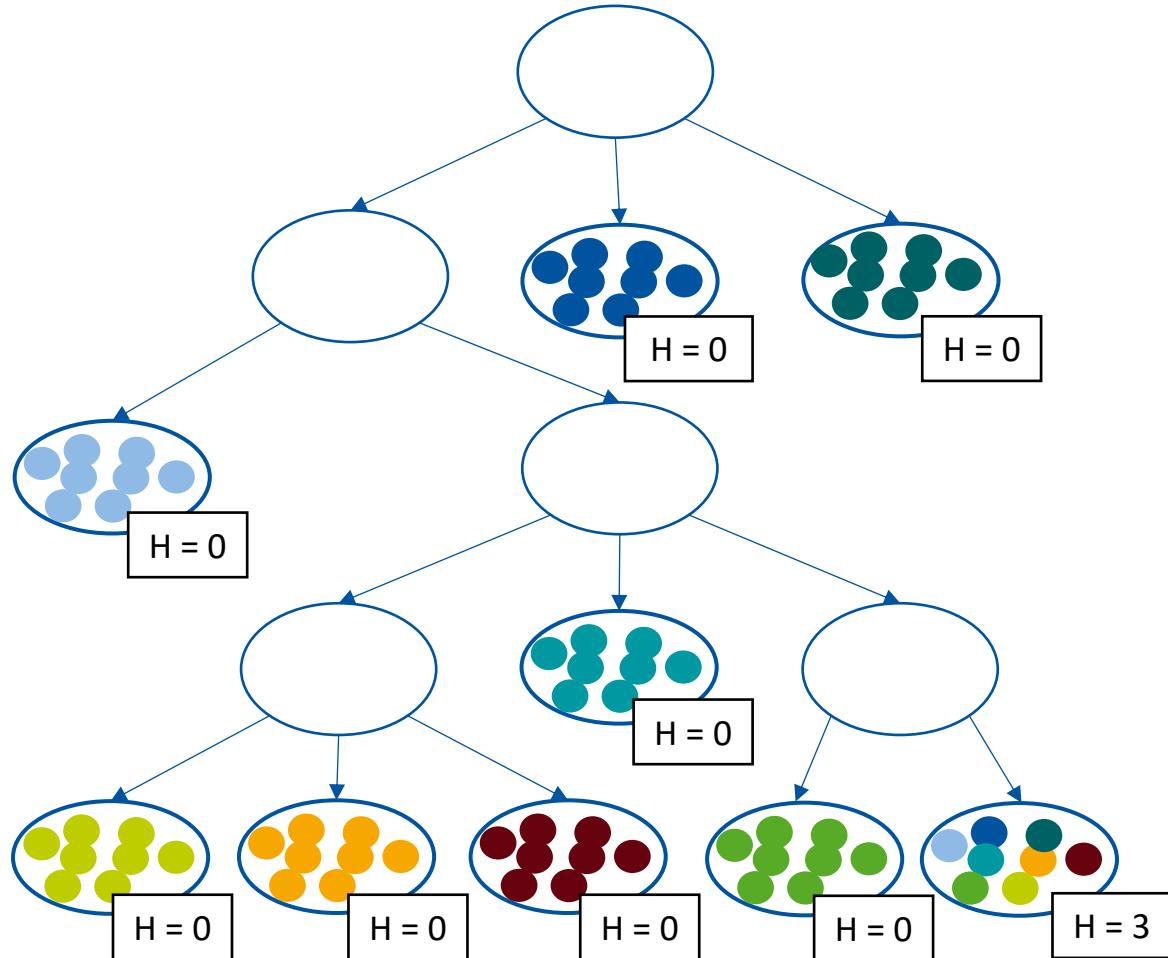
Overall entropy  $H_W$  is the weighted average of the individual entropies:

$$H_W(t) = \sum_{node \in nodes} \left( \frac{|node|}{N} \cdot H^{node}(t) \right)$$

Example:  $N = 72, K = 8$

Even distribution of 8 colors over 8 balls:  
$$H^{node}(color) = - \sum_{k=1}^8 \frac{1}{8} \cdot \log_2\left(\frac{1}{8}\right) = \log_2(8) = 3$$

# Overall Entropy

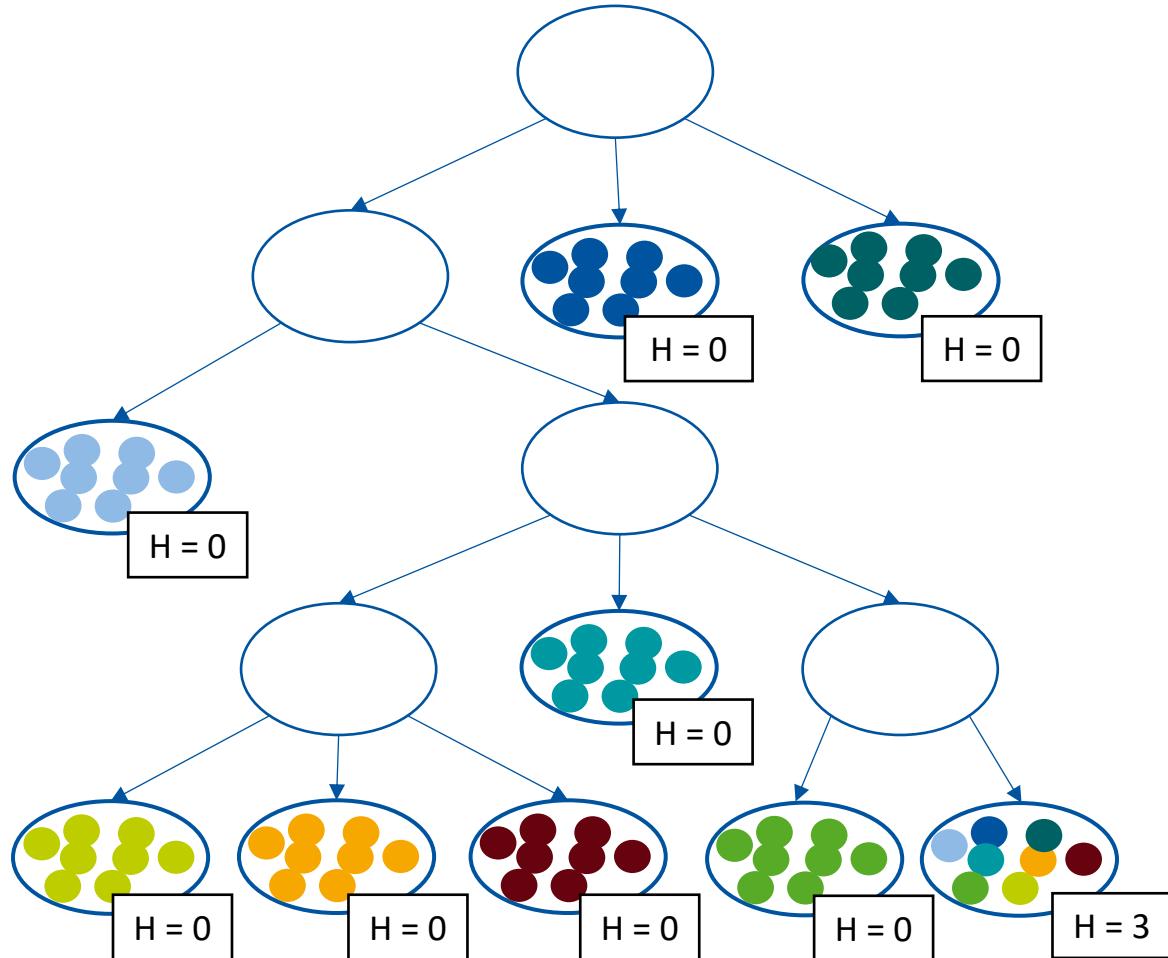


Overall entropy  $H_W$  is the weighted average of the individual entropies:

$$H_W(t) = \sum_{node \in nodes} \left( \frac{|node|}{N} \cdot H^{node}(t) \right)$$

Example:  $N = 72, K = 8$

# Overall Entropy



**Overall entropy**  $H_w$  is the weighted average of the individual entropies:

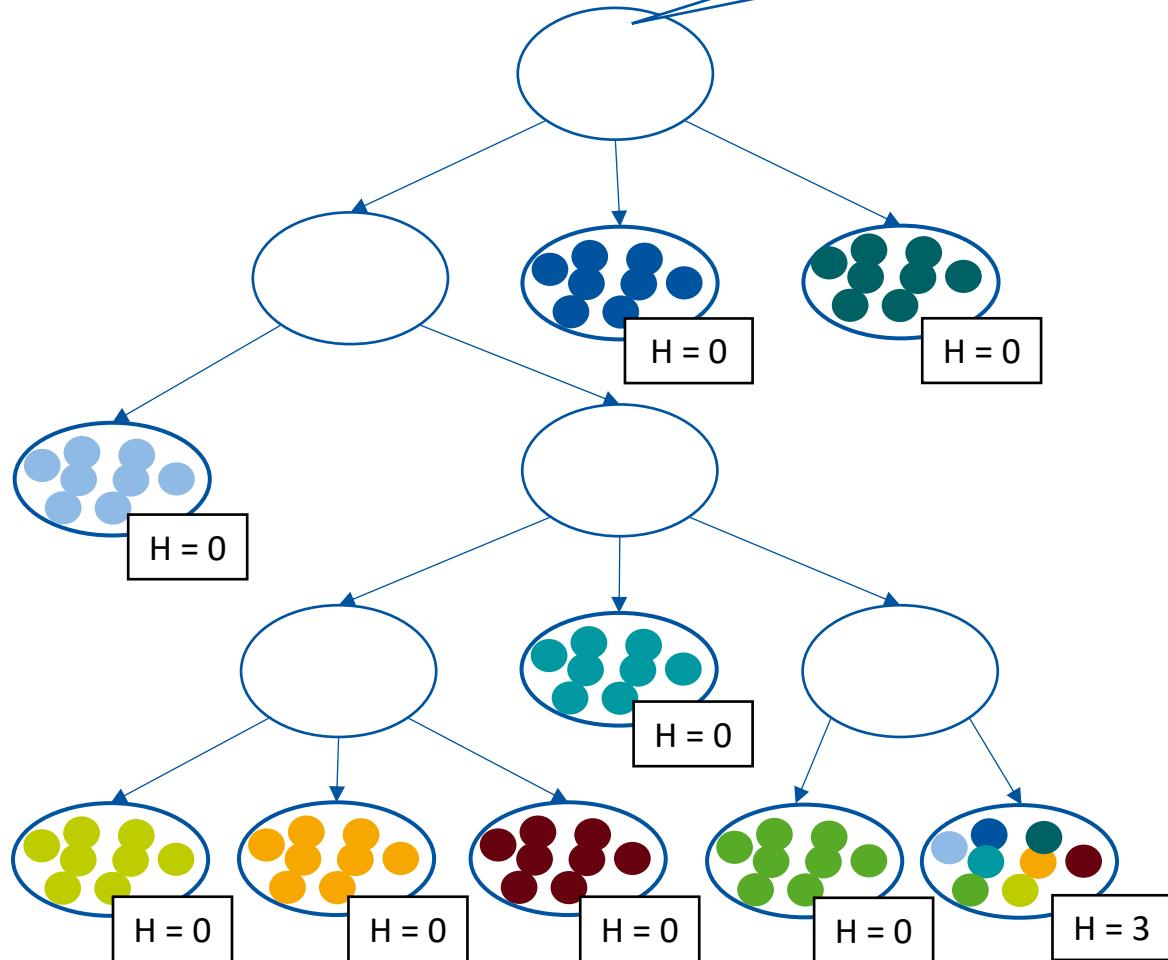
$$H_W(t) = \sum_{node \in nodes} \left( \frac{|node|}{N} \cdot H^{node}(t) \right)$$

**Example:**  $N = 72, K = 8$

$$H_W(\text{color}) = \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 \\ + \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 \\ + \frac{8}{72} \cdot 3 = \frac{24}{72} \approx 0.33$$

# Overall Entropy

Even distribution of 8 colors over 72 balls:



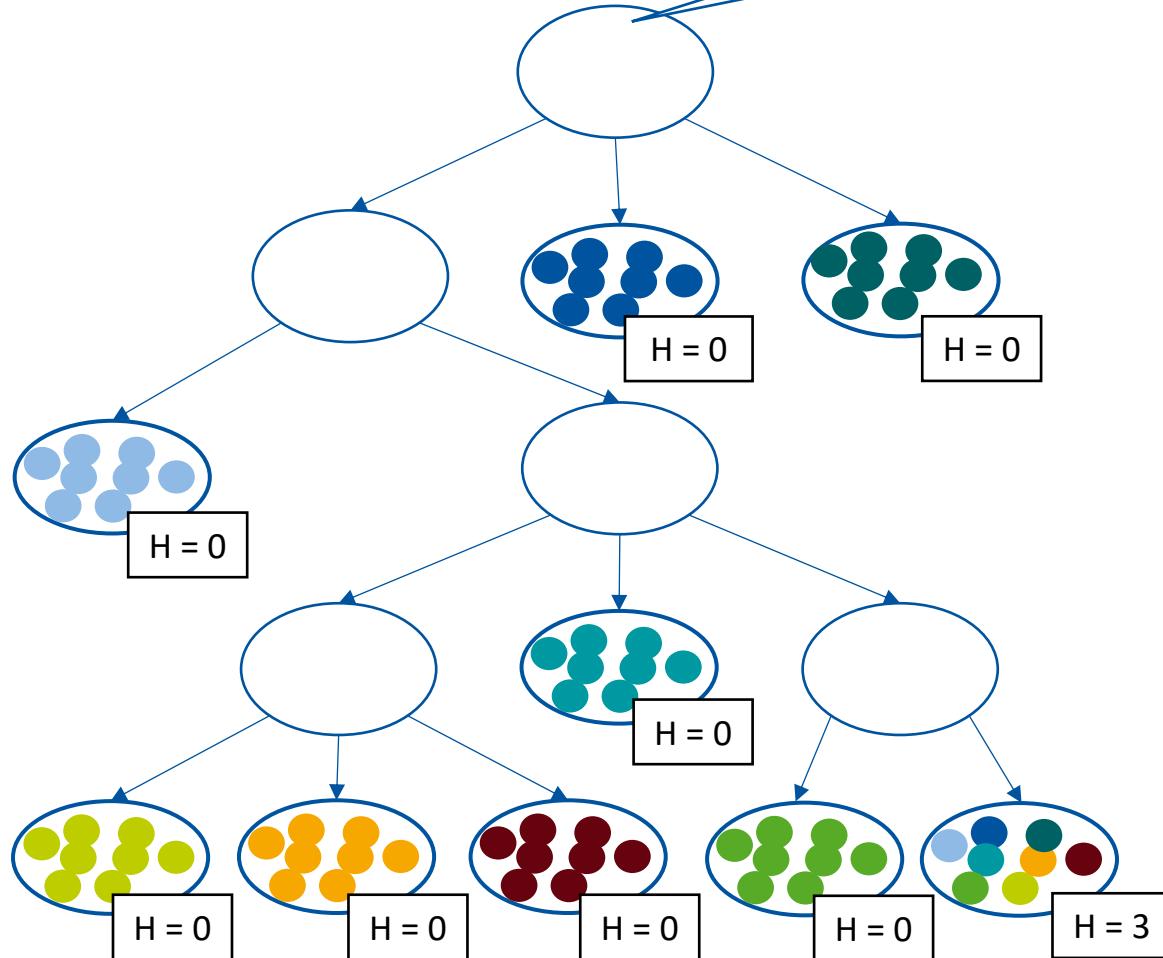
**Overall entropy**  $H_w$  is the weighted average of the individual entropies:

$$H_W(t) = \sum_{node \in nodes} \left( \frac{|node|}{N} \cdot H^{node}(t) \right)$$

**Example:**  $N = 72, K = 8$

$$H_W(\text{color}) = \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 \\ + \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 \\ + \frac{8}{72} \cdot 3 = \frac{24}{72} \approx 0.33$$

# Overall Entropy



Even distribution of 8 colors over 72 balls:

$$H_W(\text{color}) = \frac{72}{72} \cdot \left( -\sum_{k=1}^8 \left( \frac{9}{72} \cdot \log_2 \left( \frac{9}{72} \right) \right) \right) = \log_2(8) = 3$$

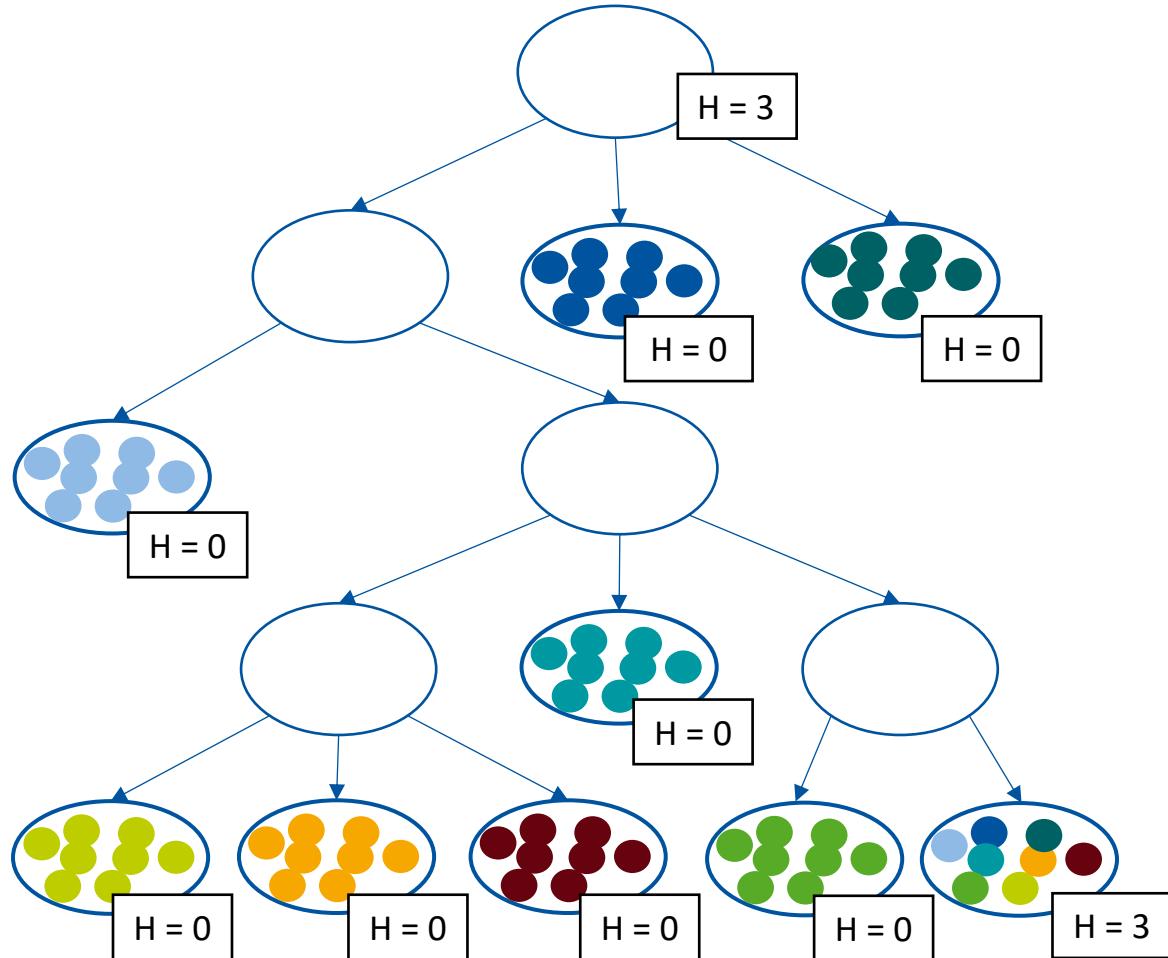
Overall entropy  $H_W$  is the weighted average of the individual entropies:

$$H_W(t) = \sum_{node \in nodes} \left( \frac{|node|}{N} \cdot H^{node}(t) \right)$$

Example:  $N = 72, K = 8$

$$\begin{aligned}
 H_W(\text{color}) &= \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 \\
 &\quad + \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 + \frac{8}{72} \cdot 0 \\
 &\quad + \frac{8}{72} \cdot 3 = \frac{24}{72} \approx 0.33
 \end{aligned}$$

# Information Gain



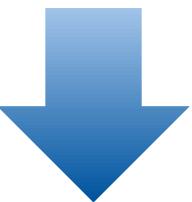
$$H_W(\text{color}) = 3$$



information loss  
 $\approx 2.66$



information gain  
 $\approx 2.66$



$$H_W(\text{color}) \approx 0.33$$

# Information Gain – Another Flight Example

$$H(\text{delayed}) = 1$$

Entropy before splitting

Weather	Traffic	Night flight	Flight delayed
Cloudy	No	No	Yes
Cloudy	Yes	No	Yes
Cloudy	Yes	No	Yes
Clear	Yes	Yes	No
Clear	No	Yes	No
Clear	No	No	No



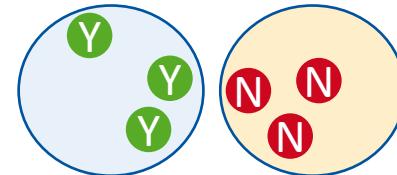
# Information Gain – Another Flight Example

$$H(\text{delayed}) = 1$$

Weather	Traffic	Night flight	Flight delayed
Cloudy	No	No	Yes
Cloudy	Yes	No	Yes
Cloudy	Yes	No	Yes
Clear	Yes	Yes	No
Clear	No	Yes	No
Clear	No	No	No



Weather	Flight delayed
Cloudy	Yes
Cloudy	Yes
Cloudy	Yes
Clear	No
Clear	No
Clear	No



# Information Gain – Another Flight Example

$$H^{cloudy}(\text{delayed}) = 0$$

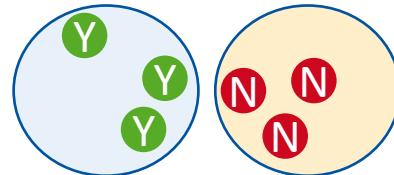
$$H^{clear}(\text{delayed}) = 0$$

Single leaf entropies after splitting based on *weather*

$$H(\text{delayed}) = 1$$

Weather	Traffic	Night flight	Flight delayed
Cloudy	No	No	Yes
Cloudy	Yes	No	Yes
Cloudy	Yes	No	Yes
Clear	Yes	Yes	No
Clear	No	Yes	No
Clear	No	No	No

Weather	Flight delayed
Cloudy	Yes
Cloudy	Yes
Cloudy	Yes
Clear	No
Clear	No
Clear	No



# Information Gain – Another Flight Example

$$H(\text{delayed}) = 1$$

Weather	Traffic	Night flight	Flight delayed
Cloudy	No	No	Yes
Cloudy	Yes	No	Yes
Cloudy	Yes	No	Yes
Clear	Yes	Yes	No
Clear	No	Yes	No
Clear	No	No	No



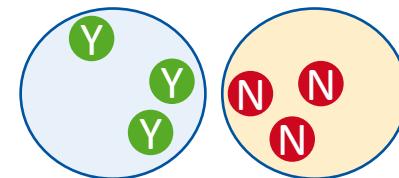
$$H^{\text{cloudy}}(\text{delayed}) = 0$$

$$H^{\text{clear}}(\text{delayed}) = 0$$

$$H_W^{\text{weather}}(\text{delayed}) = 0$$

Weighted entropy after splitting based on *weather*

Weather	Flight delayed
Cloudy	Yes
Cloudy	Yes
Cloudy	Yes
Clear	No
Clear	No
Clear	No



# Information Gain – Another Flight Example

$$H(\text{delayed}) = 1$$

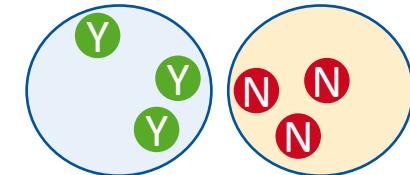
Weather	Traffic	Night flight	Flight delayed
Cloudy	No	No	Yes
Cloudy	Yes	No	Yes
Cloudy	Yes	No	Yes
Clear	Yes	Yes	No
Clear	No	Yes	No
Clear	No	No	No



$$H_{\text{W}}^{\text{weather}}(\text{delayed}) = 0$$

$$H_{\text{W}}^{\text{clear}}(\text{delayed}) = 0$$

Weather	Flight delayed
Cloudy	Yes
Cloudy	Yes
Cloudy	Yes
Clear	No
Clear	No
Clear	No

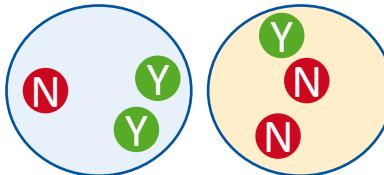


$$H_{\text{W}}^{\text{traffic\_yes}}(\text{delayed}) = 0.92$$

$$H_{\text{W}}^{\text{traffic\_no}}(\text{delayed}) = 0.92$$

$$H_{\text{W}}^{\text{traffic}}(\text{delayed}) = 0.92$$

Traffic	Flight delayed
No	Yes
Yes	Yes
Yes	Yes
Yes	No
No	No
No	No



# Information Gain – Another Flight Example

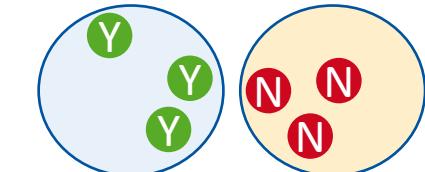
$$H(\text{delayed}) = 1$$

Weather	Traffic	Night flight	Flight delayed
Cloudy	No	No	Yes
Cloudy	Yes	No	Yes
Cloudy	Yes	No	Yes
Clear	Yes	Yes	No
Clear	No	Yes	No
Clear	No	No	No



$$H_W^{\text{weather}}(\text{delayed}) = 0$$

Weather	Flight delayed
Cloudy	Yes
Cloudy	Yes
Cloudy	Yes
Clear	No
Clear	No
Clear	No



$$H^{\text{cloudy}}(\text{delayed}) = 0$$

$$H^{\text{clear}}(\text{delayed}) = 0$$

$$H^{\text{traffic-yes}}(\text{delayed}) = 0.92$$

$$H^{\text{traffic-no}}(\text{delayed}) = 0.92$$

$$H_W^{\text{traffic}}(\text{delayed}) = 0.92$$

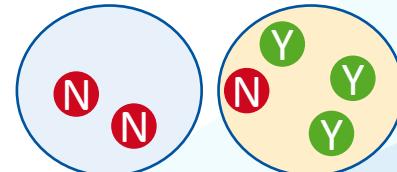
Traffic	Flight delayed
No	Yes
Yes	Yes
Yes	Yes
Yes	No
No	No
No	No

$$H^{\text{night-yes}}(\text{delayed}) = 0$$

$$H^{\text{night-no}}(\text{delayed}) \approx 0.81$$

$$H_W^{\text{night-flight}}(\text{delayed}) \approx 0.54$$

Night flight	Flight delayed
No	Yes
No	Yes
No	Yes
Yes	No
Yes	No
No	No



# Information Gain – Another Flight Example

$H(\text{delayed}) = 1$			
Weather	Traffic	Night flight	Flight delayed
Cloudy	No	No	Yes
...	...	...	...

$$H^{\text{cloudy}}(\text{delayed}) = 0$$

$$H^{\text{clear}}(\text{delayed}) = 0$$

$$H_W^{\text{weather}}(\text{delayed}) = 0$$

Weather	Flight delayed
Cloudy	Yes
...	...

$$H_W^{\text{traffic-yes}}(\text{delayed}) = 0.92$$

$$H_W^{\text{traffic-no}}(\text{delayed}) = 0.92$$

$$H_W^{\text{traffic}}(\text{delayed}) = 0.92$$

Traffic	Flight delayed
No	Yes
...	...

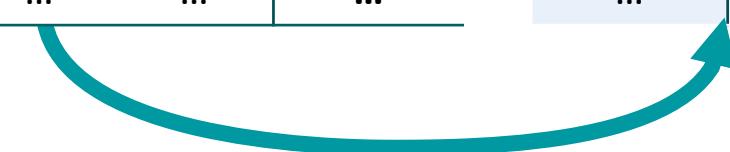
$$H_W^{\text{night-yes}}(\text{delayed}) = 0$$

$$H_W^{\text{night-no}}(\text{delayed}) \approx 0.81$$

$$H_W^{\text{night-flight}}(\text{delayed}) \approx 0.54$$

Night flight	Flight delayed
No	No
...	...

$$IG(\text{weather}) = H(\text{delayed}) - H_W^{\text{weather}}(\text{delayed}) = 1 - 0 = 1$$



# Information Gain – Another Flight Example

$$H(\text{delayed}) = 1$$

Weather	Traffic	Night flight	Flight delayed
Cloudy	No	Yes	Yes
...	...	...	...

$$\begin{aligned} H^{\text{cloudy}}(\text{delayed}) &= 0 \\ H^{\text{clear}}(\text{delayed}) &= 0 \end{aligned}$$

$$H_W^{\text{weather}}(\text{delayed}) = 0$$

Weather	Flight delayed
Cloudy	Yes
...	...

$$\begin{aligned} H_W^{\text{traffic-yes}}(\text{delayed}) &= 0.92 \\ H_W^{\text{traffic-no}}(\text{delayed}) &= 0.92 \end{aligned}$$

$$H_W^{\text{traffic}}(\text{delayed}) = 0.92$$

Traffic	Flight delayed
No	Yes
...	...

$$\begin{aligned} H^{\text{night-yes}}(\text{delayed}) &= 0 \\ H^{\text{night-no}}(\text{delayed}) &\approx 0.81 \end{aligned}$$

$$H_W^{\text{night-flight}}(\text{delayed}) \approx 0.54$$

Night flight	Flight delayed
Yes	No
...	...

$$IG(\text{weather}) = H(\text{delayed}) - H_W^{\text{weather}}(\text{delayed}) = 1 - 0 = 1$$

$$IG(\text{traffic}) = H(\text{delayed}) - H_W^{\text{traffic}}(\text{delayed}) = 1 - 0.92 = 0.08$$

# Information Gain – Another Flight Example

$H(\text{delayed}) = 1$			
Weather	Traffic	Night flight	Flight delayed
Cloudy	No	Yes	Yes
...	...	...	...

$$H^{\text{cloudy}}(\text{delayed}) = 0$$

$$H^{\text{clear}}(\text{delayed}) = 0$$

$$H_W^{\text{weather}}(\text{delayed}) = 0$$

Weather	Flight delayed
Cloudy	Yes
...	...

$$H_W^{\text{traffic-yes}}(\text{delayed}) = 0.92$$

$$H_W^{\text{traffic-no}}(\text{delayed}) = 0.92$$

$$H_W^{\text{traffic}}(\text{delayed}) = 0.92$$

Traffic	Flight delayed
No	Yes
...	...

$$H_W^{\text{night-yes}}(\text{delayed}) = 0$$

$$H_W^{\text{night-no}}(\text{delayed}) \approx 0.81$$

$$H_W^{\text{night-flight}}(\text{delayed}) \approx 0.54$$

Night flight	Flight delayed
Yes	No
...	...

$$IG(\text{weather}) = H(\text{delayed}) - H_W^{\text{weather}}(\text{delayed}) = 1 - 0 = 1$$

$$IG(\text{traffic}) = H(\text{delayed}) - H_W^{\text{traffic}}(\text{delayed}) = 1 - 0.92 = 0.08$$

$$IG(\text{night\_flight}) = H(\text{delayed}) - H_W^{\text{night\_flight}}(\text{delayed}) = 1 - 0.54 = 0.46$$

# Information Gain – Another Flight Example

$H(\text{delayed}) = 1$			
Weather	Traffic	Night flight	Flight delayed
Cloudy	No	Yes	Yes
...	...	...	...

$$H^{\text{cloudy}}(\text{delayed}) = 0$$

$$H^{\text{clear}}(\text{delayed}) = 0$$

$$H_W^{\text{weather}}(\text{delayed}) = 0$$

Weather	Flight delayed
Cloudy	Yes
...	...

$$H^{\text{traffic-yes}}(\text{delayed}) = 0.92$$

$$H^{\text{traffic-no}}(\text{delayed}) = 0.92$$

$$H_W^{\text{traffic}}(\text{delayed}) = 0.92$$

Traffic	Flight delayed
No	Yes
...	...

$$H^{\text{night-yes}}(\text{delayed}) = 0$$

$$H^{\text{night-no}}(\text{delayed}) \approx 0.81$$

$$H_W^{\text{night-flight}}(\text{delayed}) \approx 0.54$$

Night flight	Flight delayed
Yes	No
...	...

$$IG(\text{weather}) = H(\text{delayed}) - H_W^{\text{weather}}(\text{delayed}) = 1 - 0 = 1$$

**best**

$$IG(\text{traffic}) = H(\text{delayed}) - H_W^{\text{traffic}}(\text{delayed}) = 1 - 0.92 = 0.08$$

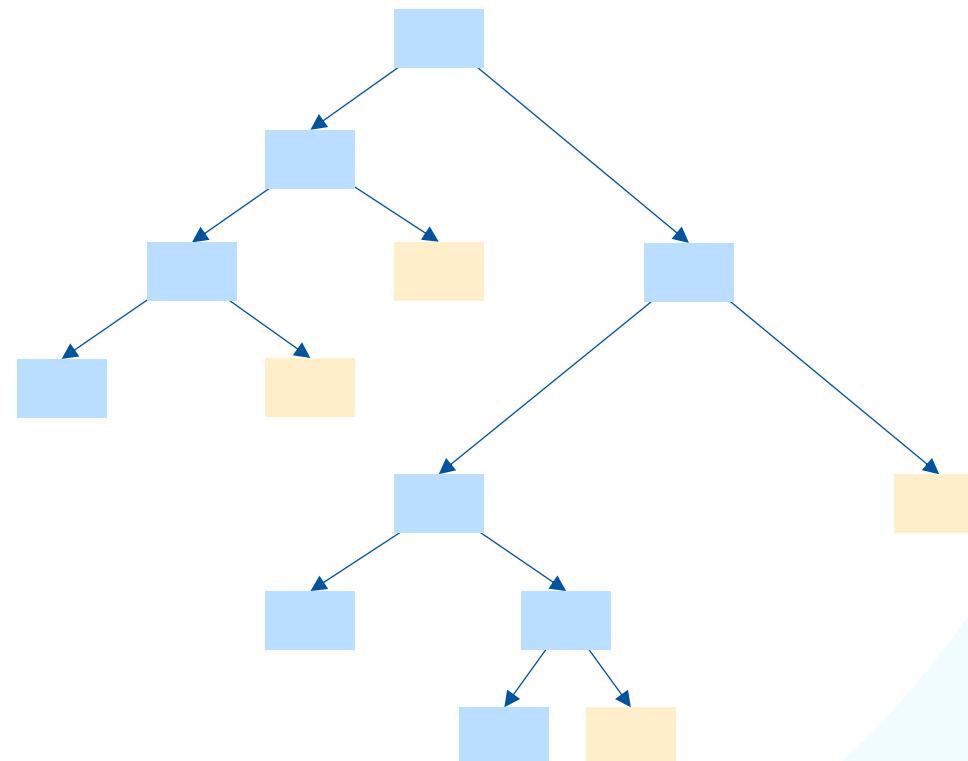
**worst**

$$IG(\text{night\_flight}) = H(\text{delayed}) - H_W^{\text{night\_flight}}(\text{delayed}) = 1 - 0.54 = 0.46$$

**intermediate**

# Decision Trees

1. Introduction to Decision Trees
2. Entropy and Information Gain
3. **ID3 Algorithm**
4. Quantifying Information Gain
5. Pruning
6. Continuous Data
7. Ensembles



# ID3 (Iterative Dichotomiser 3) - Key Idea

## Approach

1. For each candidate feature: calculate the resulting entropy splitting the dataset  $\mathcal{X}$  using the selected feature.
2. Split the set  $\mathcal{X}$  into subsets using the feature for which the resulting entropy (after splitting) is minimal (equivalently, information gain is maximal)
3. Create new decision tree leaf nodes based on that feature
4. Recurse on these subsets using remaining features (until stopping criteria are reached)

# When to Stop?

## Three stopping criteria

- When all of the instances have the same classification (**label = consensus value**)
- When there are no features left (**label = majority value**)
- When the dataset is empty (**label = majority parent**)

# ID3 Algorithm

**ID3 algorithm:**

1. **if** all the instances in  $X$  have the same classification
  - (a) **return** a decision tree with one leaf node with consensus value as a label
2. **else if** there are no features left
  - (a) **return** a decision tree with one leaf node with majority value as a label
3. **else if** the dataset is empty
  - (a) **return** a decision tree with one leaf node with majority parent value as a label

three  
stopping  
criteria

4. **else**

- (a) pick a feature that maximizes information gain
- (b) once a feature is picked along a path from the root, it cannot be used again
- (c) create subproblems based on the selected feature

recursively  
constructing  
the tree

## Example

$H(\text{Customer})$

$$= -\left(\frac{2}{7} \cdot \log_2\left(\frac{2}{7}\right) + \frac{3}{7} \cdot \log_2\left(\frac{3}{7}\right) + \frac{2}{7} \cdot \log_2\left(\frac{2}{7}\right)\right)$$

$$= 1.5567$$

ID	Insurance	Education	Employment	Customer
1	Yes	Bachelor	Employed	Basic
2	Yes	High school	Unemployed	Premium
3	Yes	Bachelor	Self-employed	Premium
4	No	Bachelor	Self-employed	Basic
5	No	Master	Employed	Economy
6	Yes	Bachelor	Retired	Economy
7	Yes	High school	Employed	Premium

## Example

$$H(\text{Customer}) = 1.5567$$

ID	Insurance	Education	Employment	Customer
1	Yes	Bachelor	Employed	Basic
2	Yes	High school	Unemployed	Premium
3	Yes	Bachelor	Self-employed	Premium
4	No	Bachelor	Self-employed	Basic
5	No	Master	Employed	Economy
6	Yes	Bachelor	Retired	Economy
7	Yes	High school	Employed	Premium

Split by feature	Possible Values	Instance Distribution	Entropy	Overall Entropy	Information Gain	
Insurance	No	4, 5	1	1.265	$1.5567 - 1.265 = 0.2917$	
	Yes	1, 2, 3, 6, 7	1.3710			
Education	High school	2, 7	0	0.8571	$1.5567 - 0.8571 = 0.6996$	
	Master	5	0			
Employment	Bachelor	1, 3, 4, 6	1.5	0.9650	$1.5567 - 0.9650 = 0.5917$	
	Employed	1, 5, 7	1.5850			
	Unemployed	2	0			
	Self-employed	3, 4	1			
	Retired	6	0			

## Example

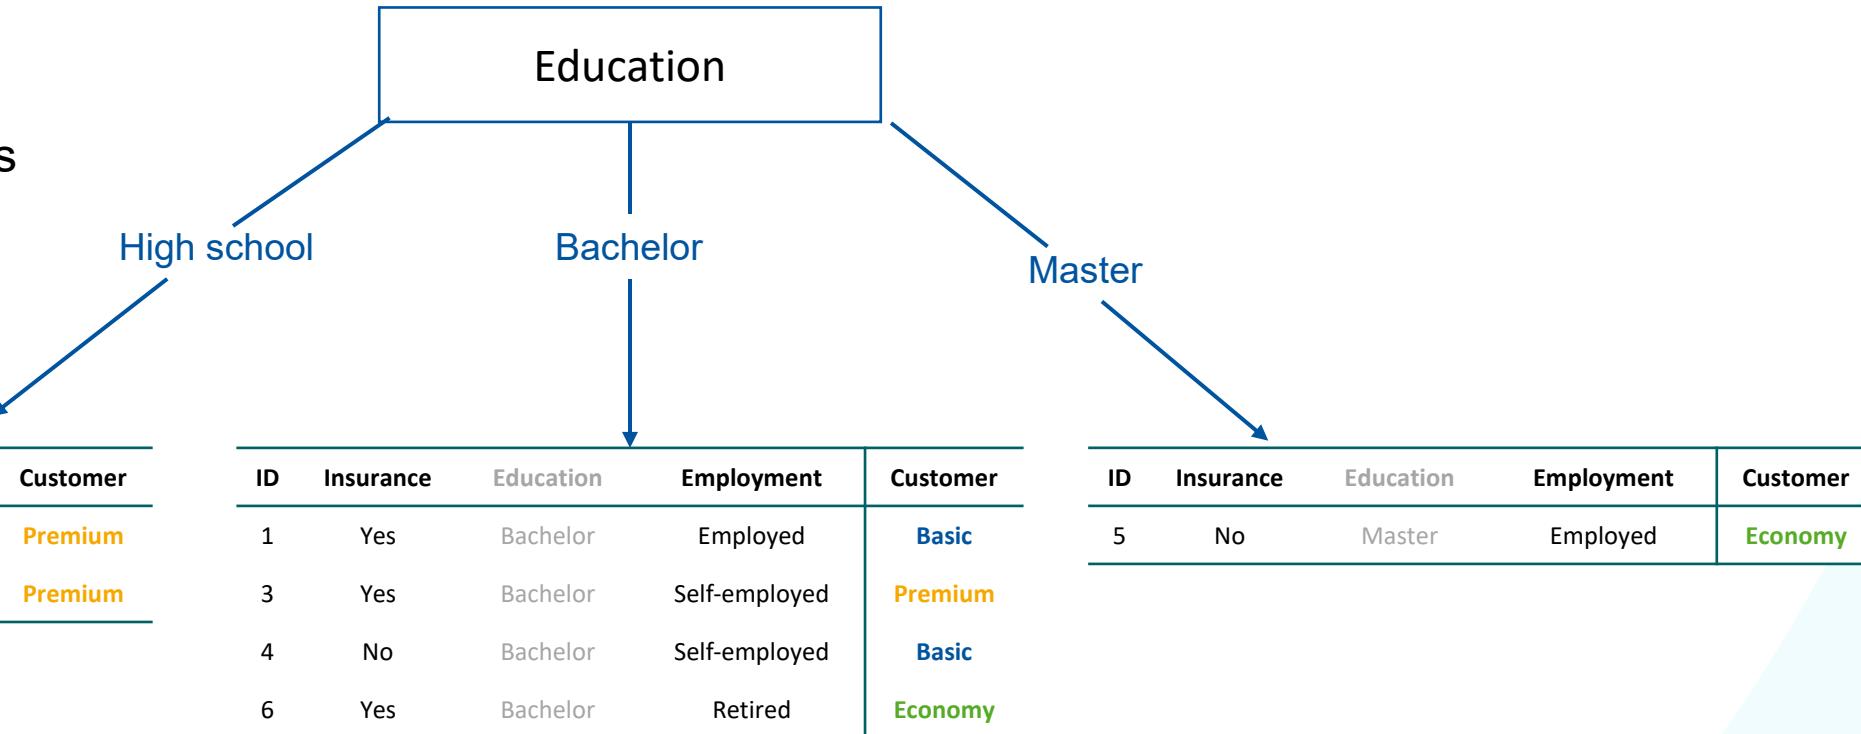
$$H(\text{Customer}) = 1.5567$$

ID	Insurance	Education	Employment	Customer
1	Yes	Bachelor	Employed	Basic
2	Yes	High school	Unemployed	Premium
3	Yes	Bachelor	Self-employed	Premium
4	No	Bachelor	Self-employed	Basic
5	No	Master	Employed	Economy
6	Yes	Bachelor	Retired	Economy
7	Yes	High school	Employed	Premium

Split by feature	Possible Values	Instances	Entropy	Overall Entropy	Information Gain	
Insurance	No	4, 5	1	1.265	$1.5567 - 1.265 = 0.2917$	
	Yes	1, 2, 3, 6, 7	1.3710			
Education	High school	2, 7	0	0.8571	$1.5567 - 0.8571 = 0.6996$	
	Master	5	0			
Employment	Bachelor	1, 3, 4, 6	1.5	0.9650	$1.5567 - 0.9650 = 0.5917$	
	Employed	1, 5, 7	1.5850			
	Unemployed	2	0			
	Self-employed	3, 4	1			
	Retired	6	0			

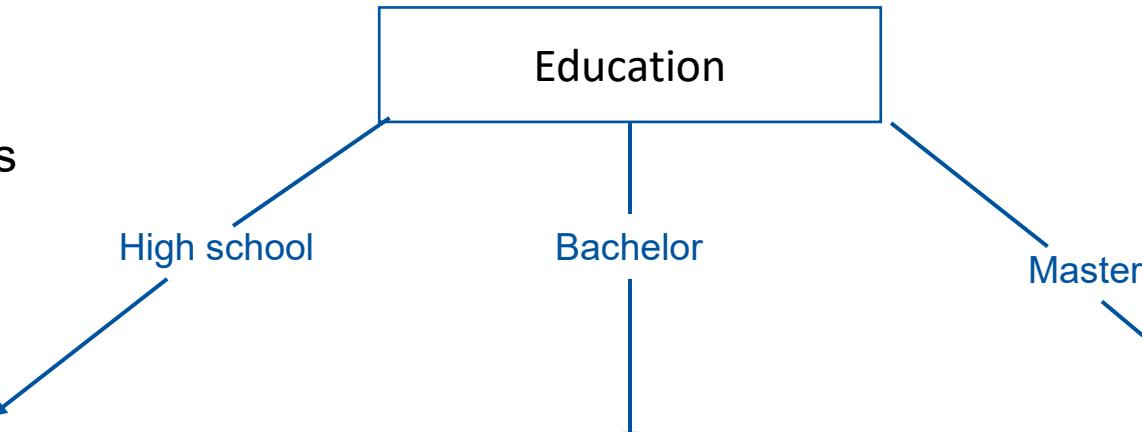
# Example

Recursion until  
stopping criteria holds



# Example

Recursion until  
stopping criteria holds



ID	Insurance	Education	Employment	Customer
2	Yes	High school	Unemployed	Premium
7	Yes	High school	Employed	Premium

**consensus**

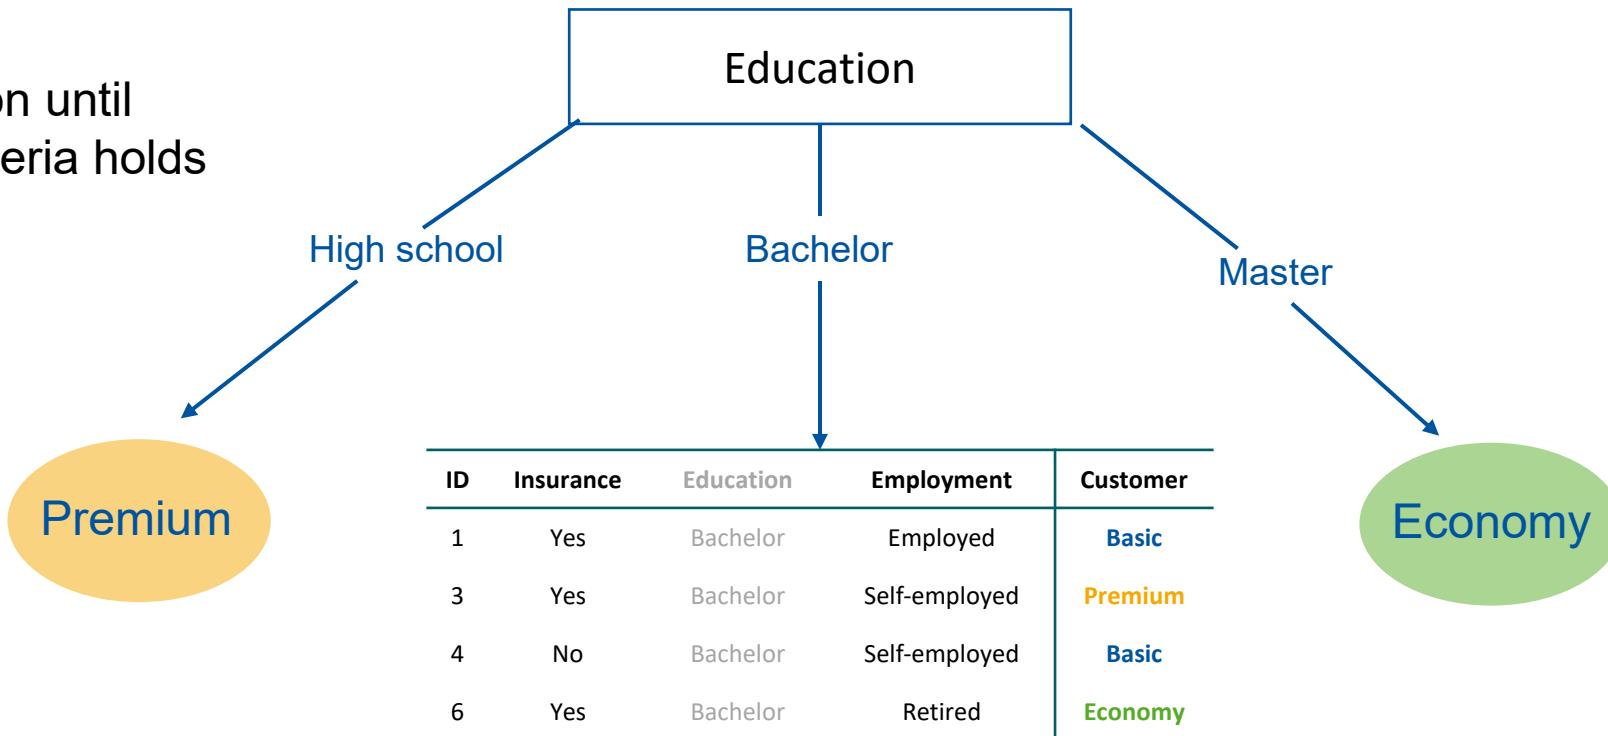
ID	Insurance	Education	Employment	Customer
1	Yes	Bachelor	Employed	Basic
3	Yes	Bachelor	Self-employed	Premium
4	No	Bachelor	Self-employed	Basic
6	Yes	Bachelor	Retired	Economy

ID	Insurance	Education	Employment	Customer
5	No	Master	Employed	Economy

**consensus**

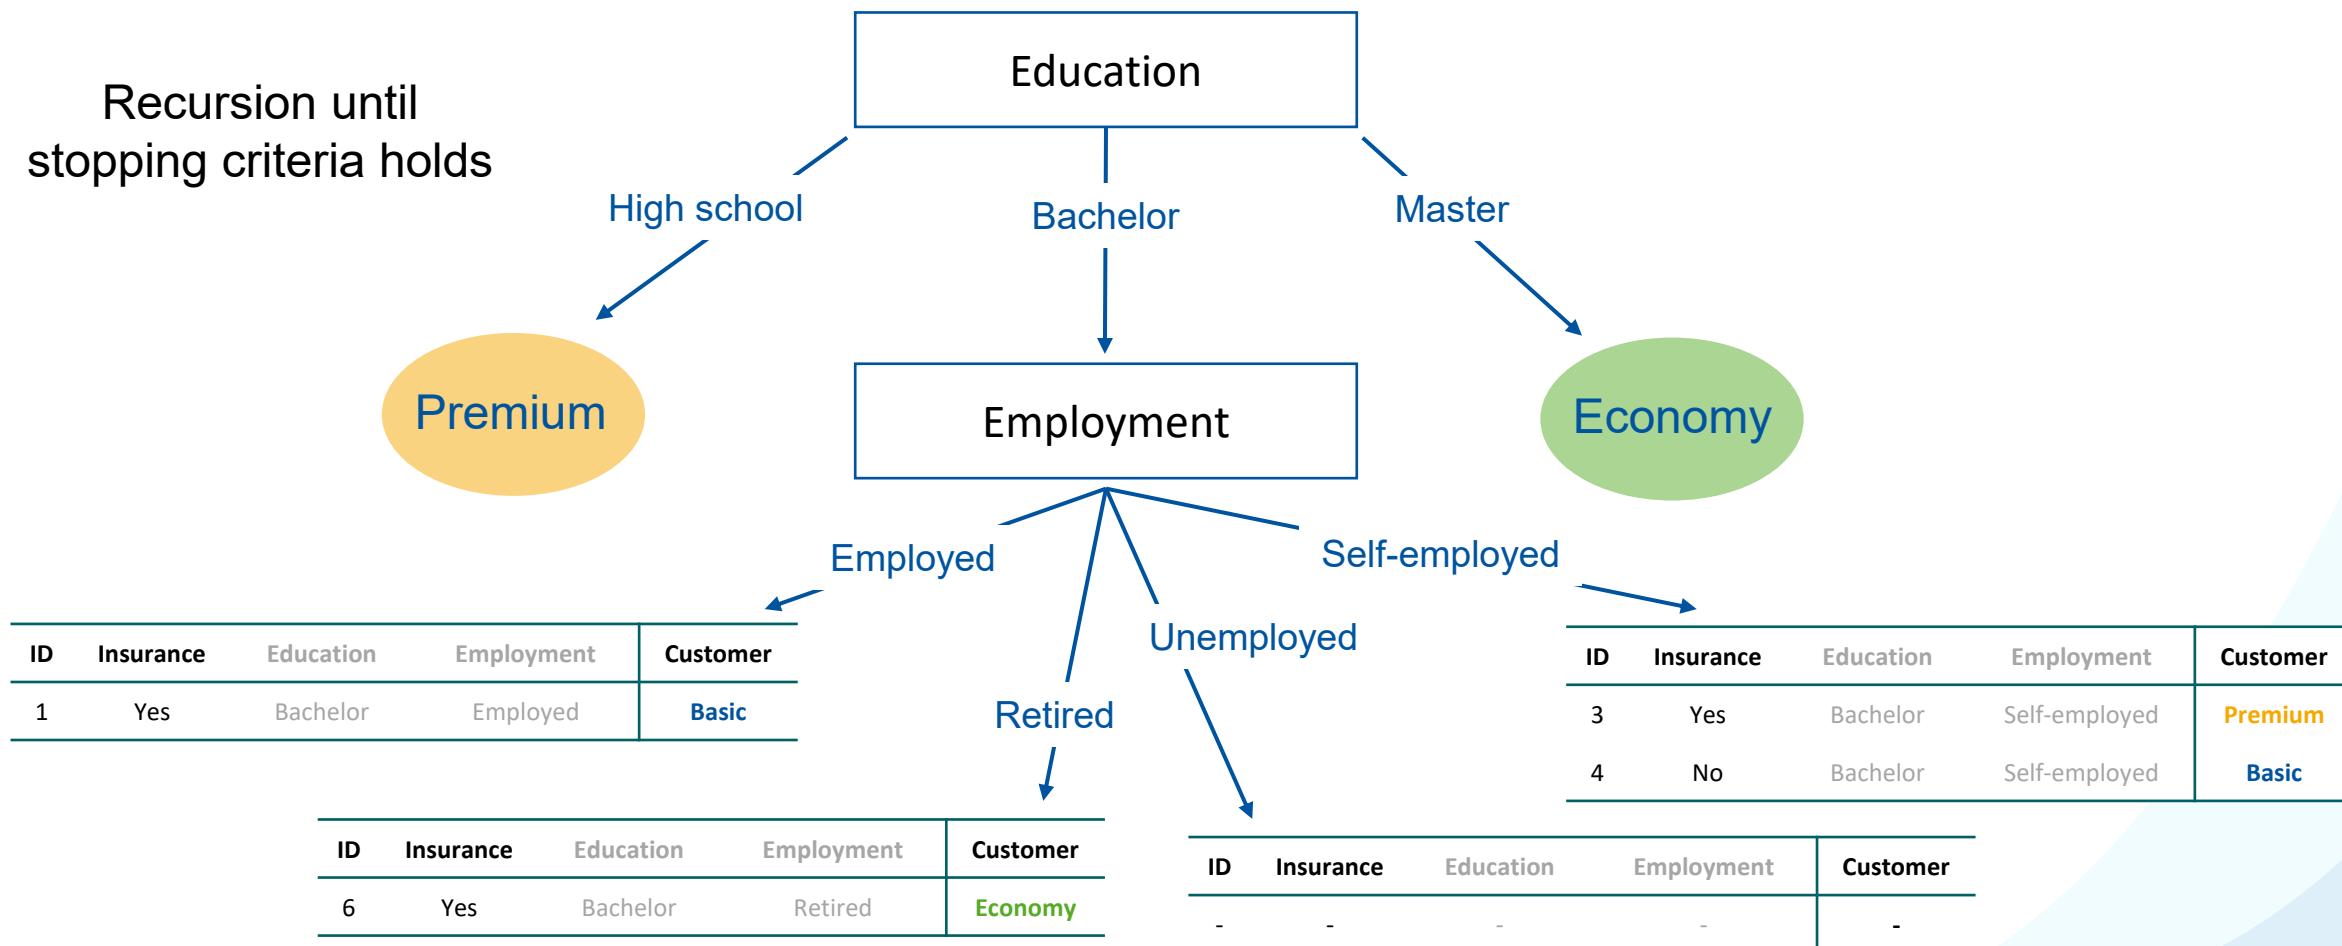
## Example

Recursion until  
stopping criteria holds



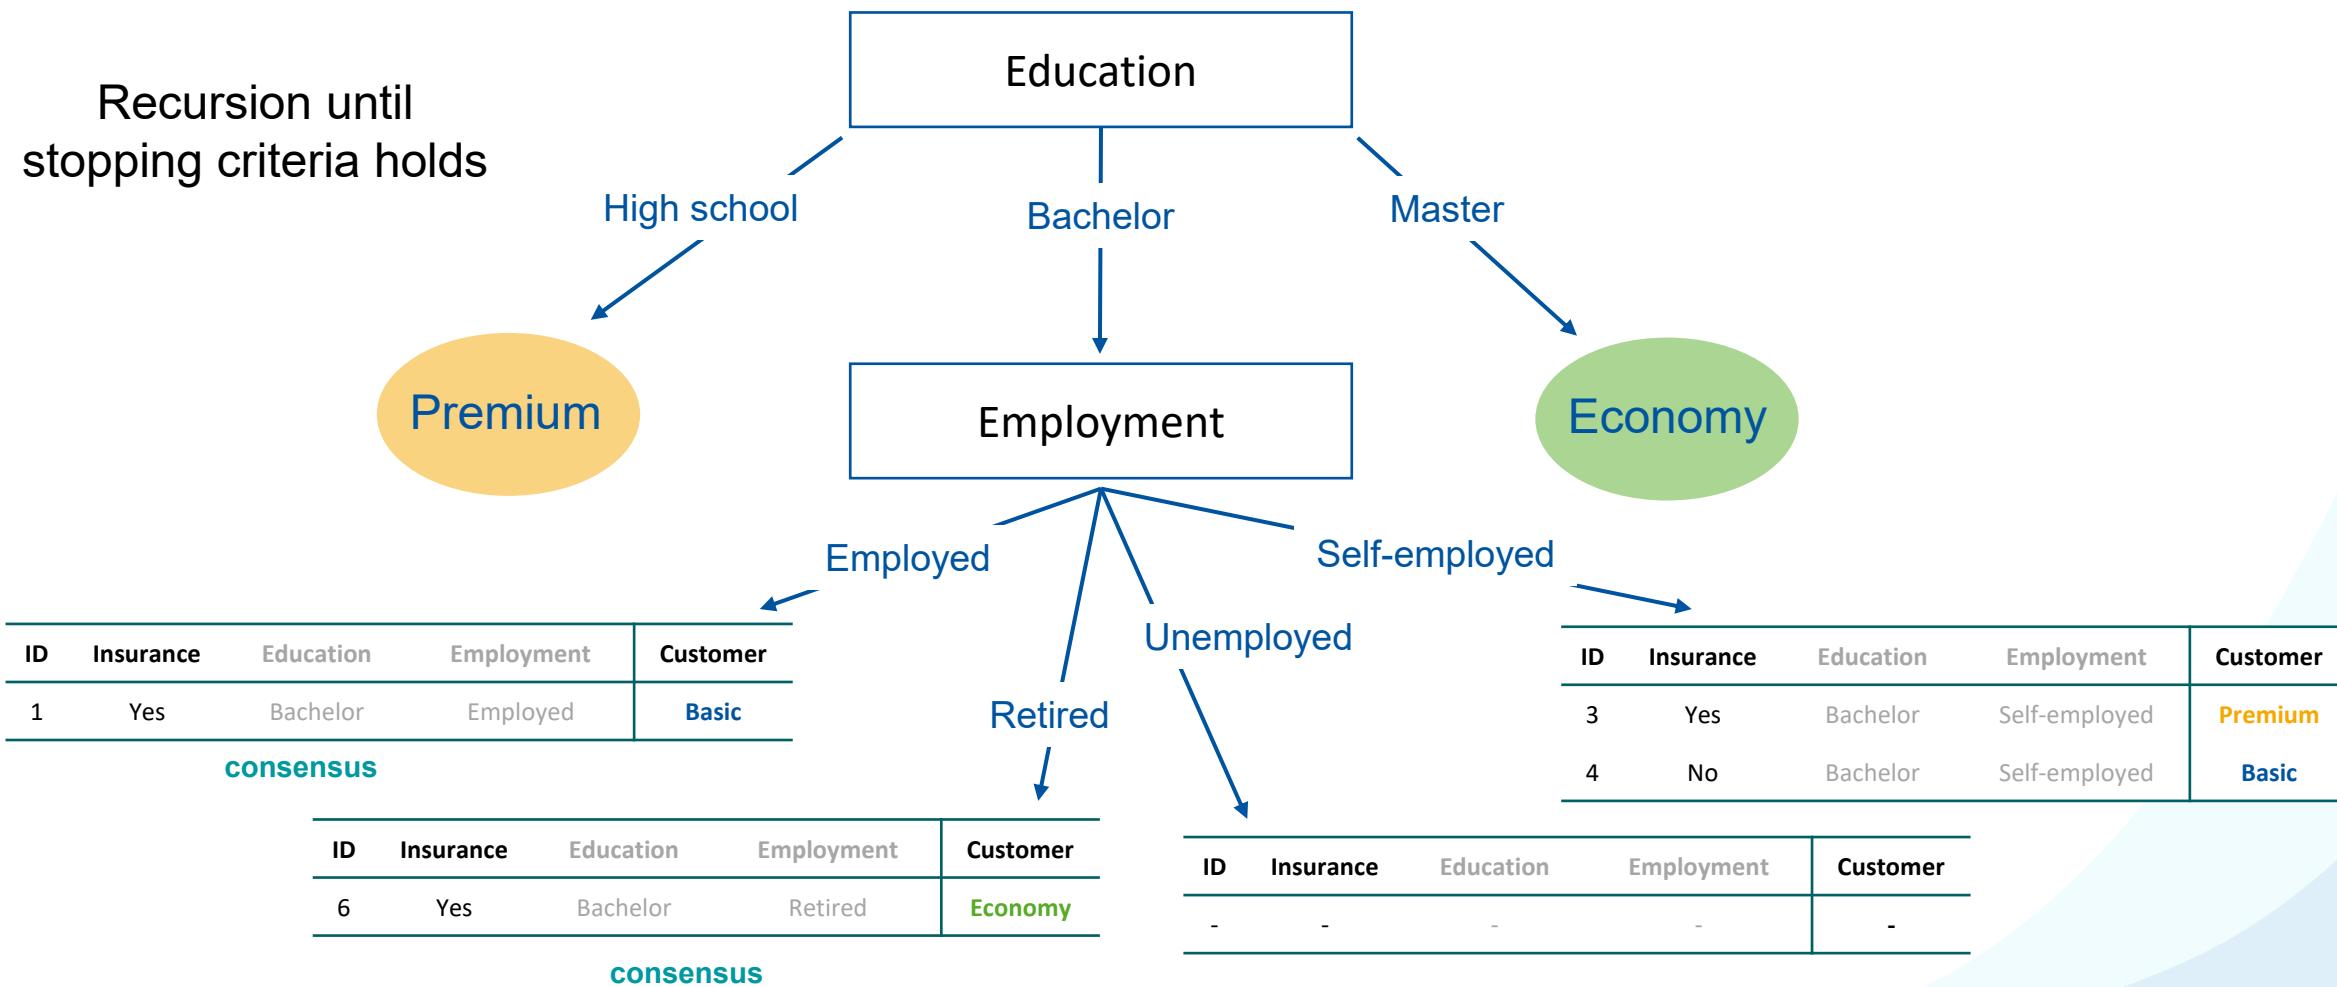
# Example

Recursion until  
stopping criteria holds



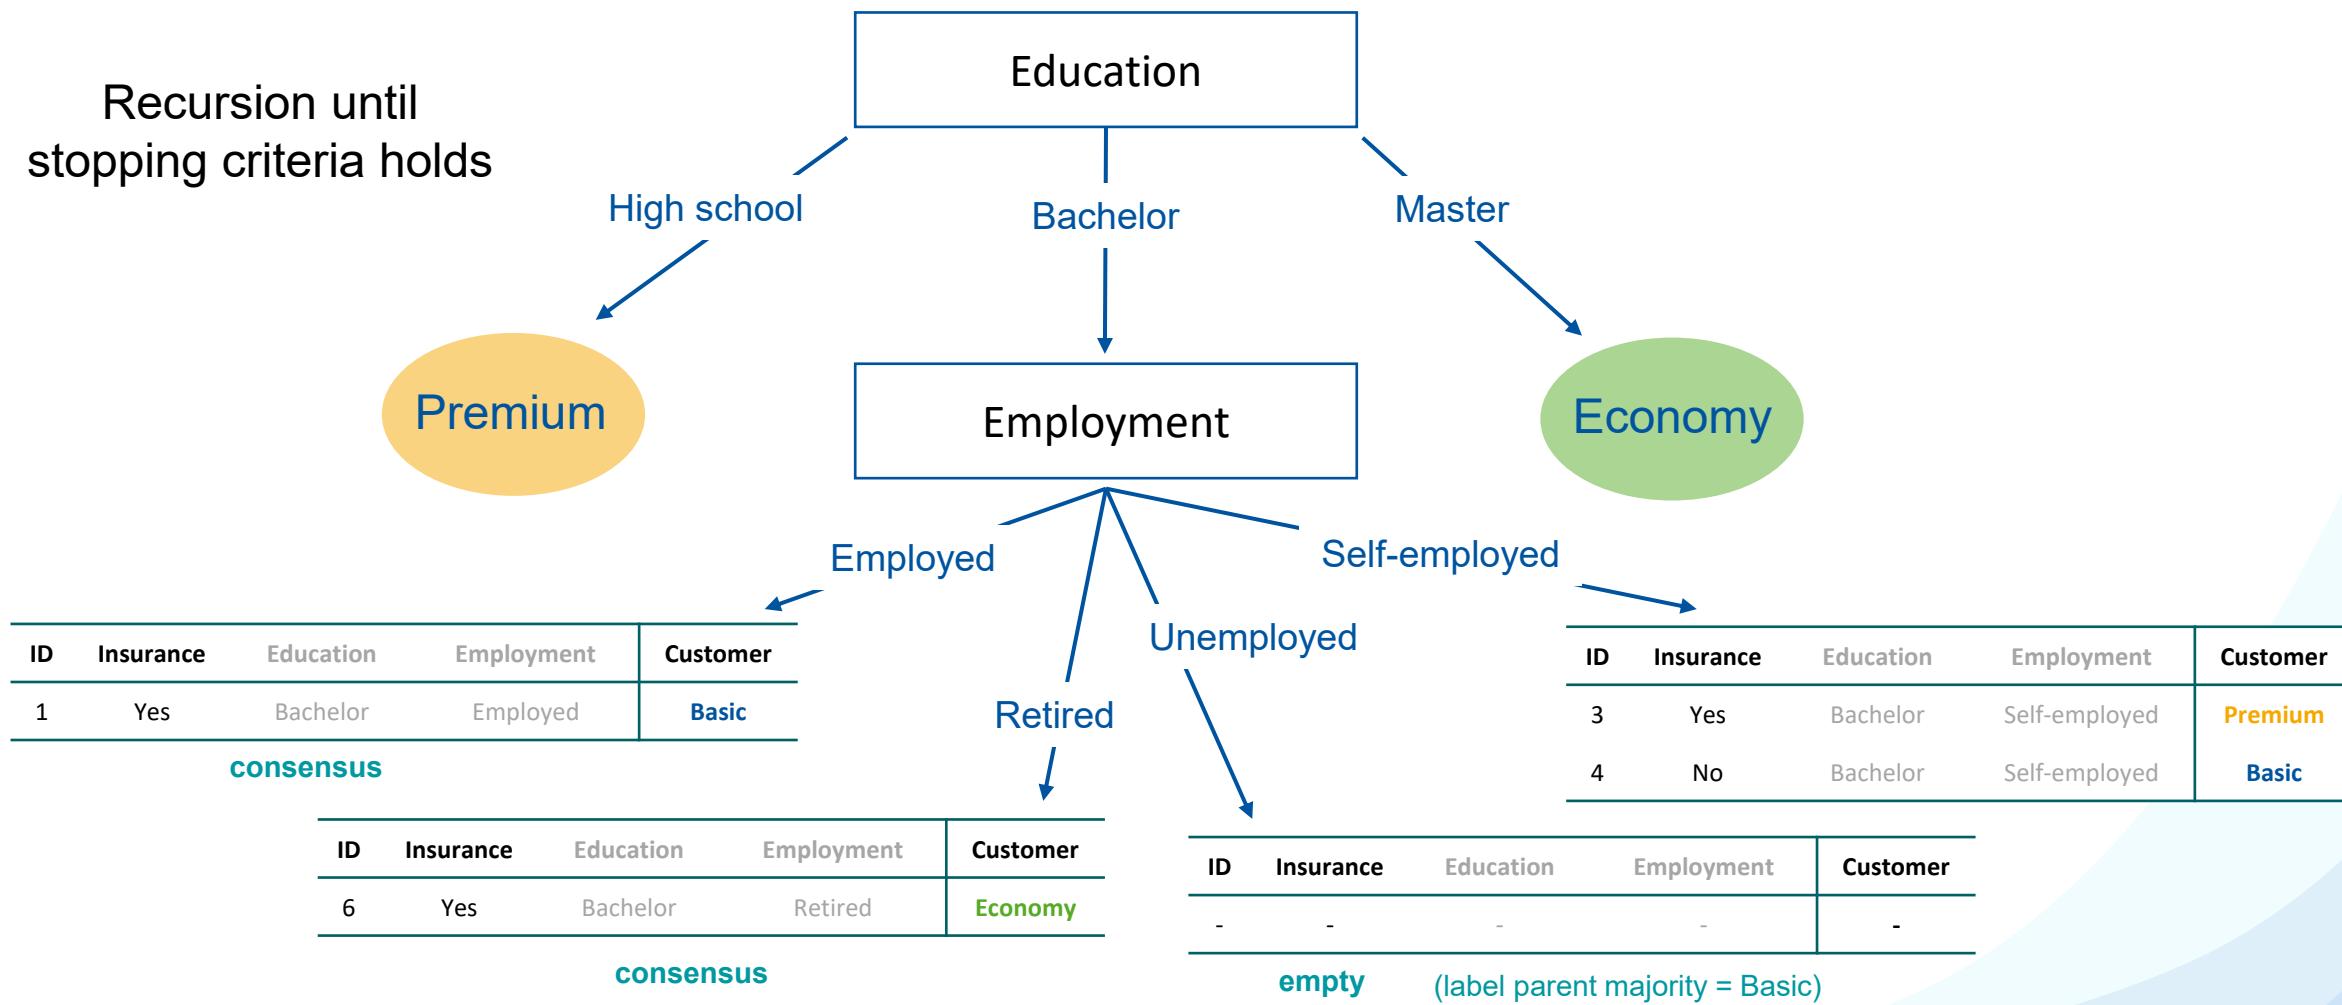
# Example

Recursion until  
stopping criteria holds



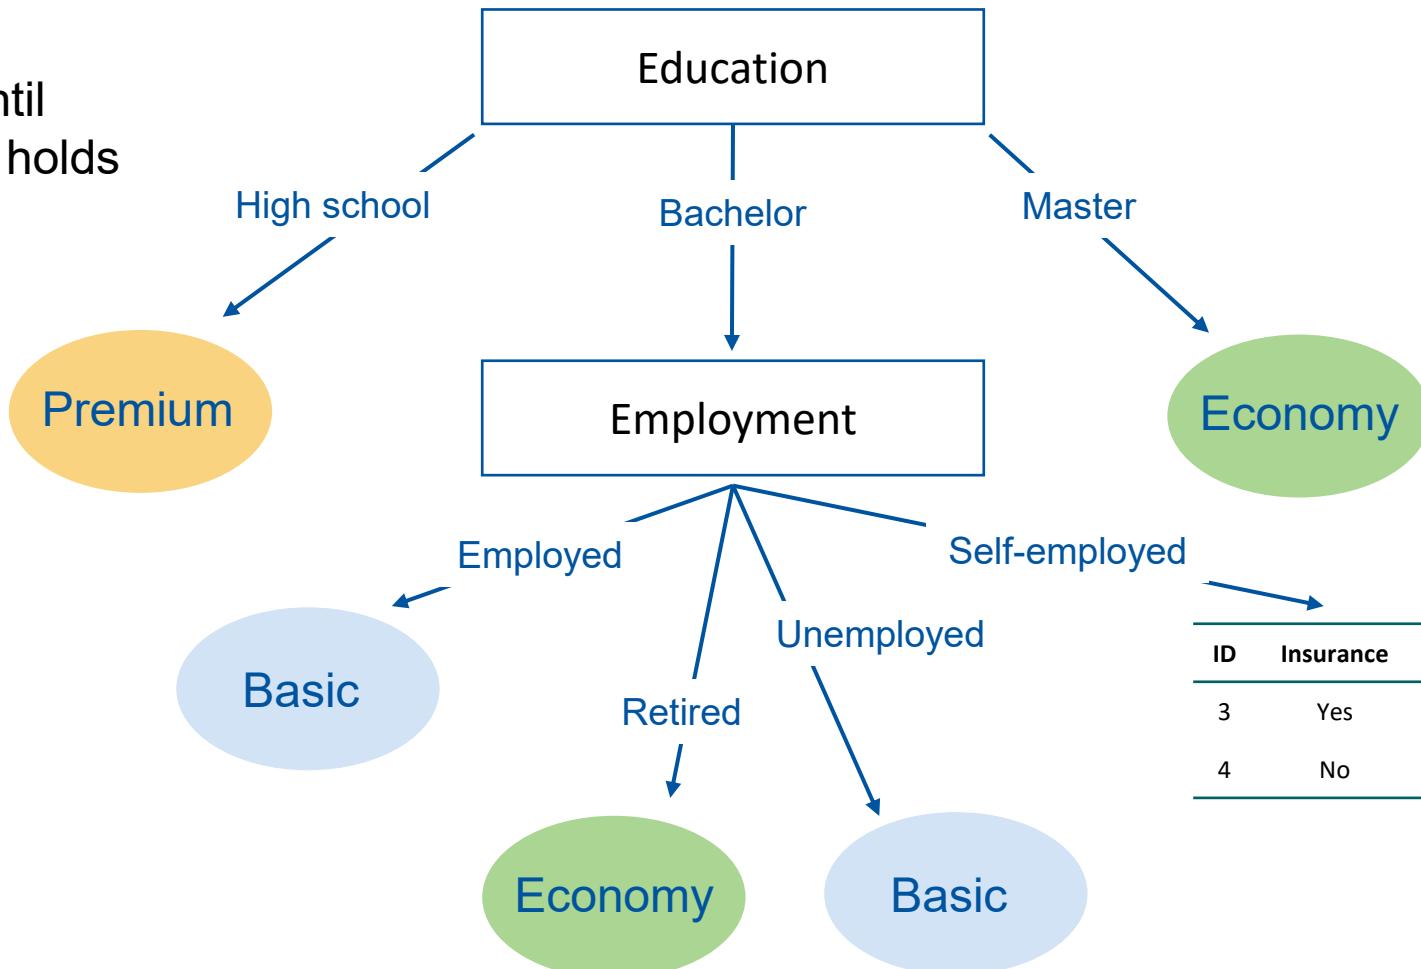
# Example

Recursion until  
stopping criteria holds



## Example

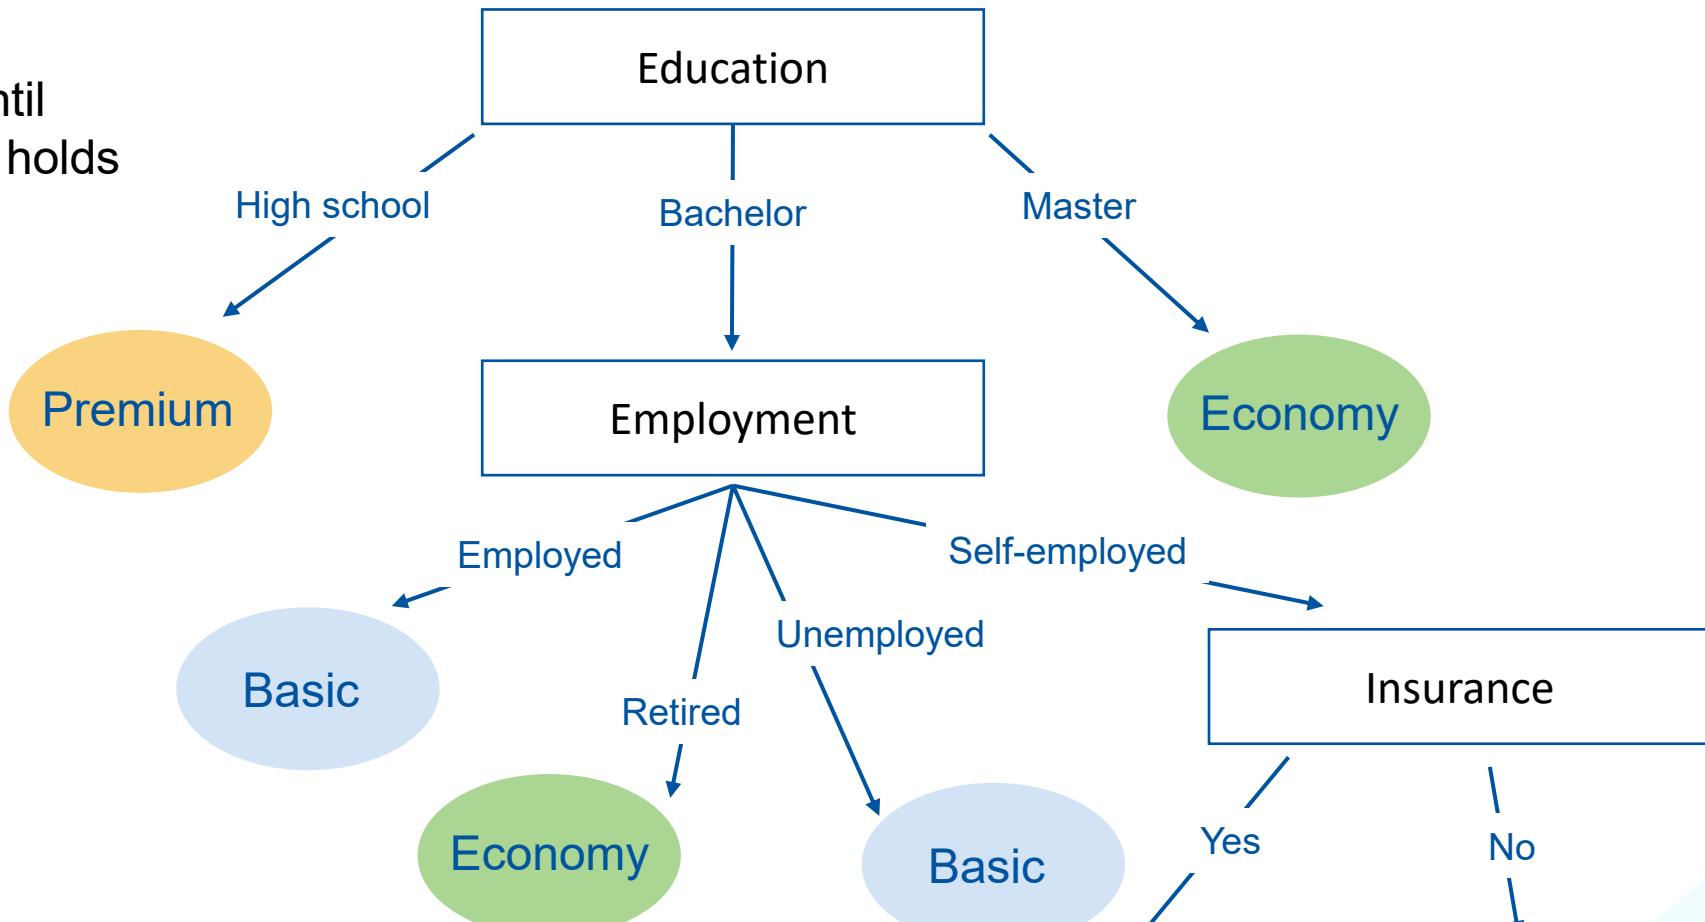
Recursion until  
stopping criteria holds



ID	Insurance	Education	Employment	Customer
3	Yes	Bachelor	Self-employed	Premium
4	No	Bachelor	Self-employed	Basic

# Example

Recursion until  
stopping criteria holds



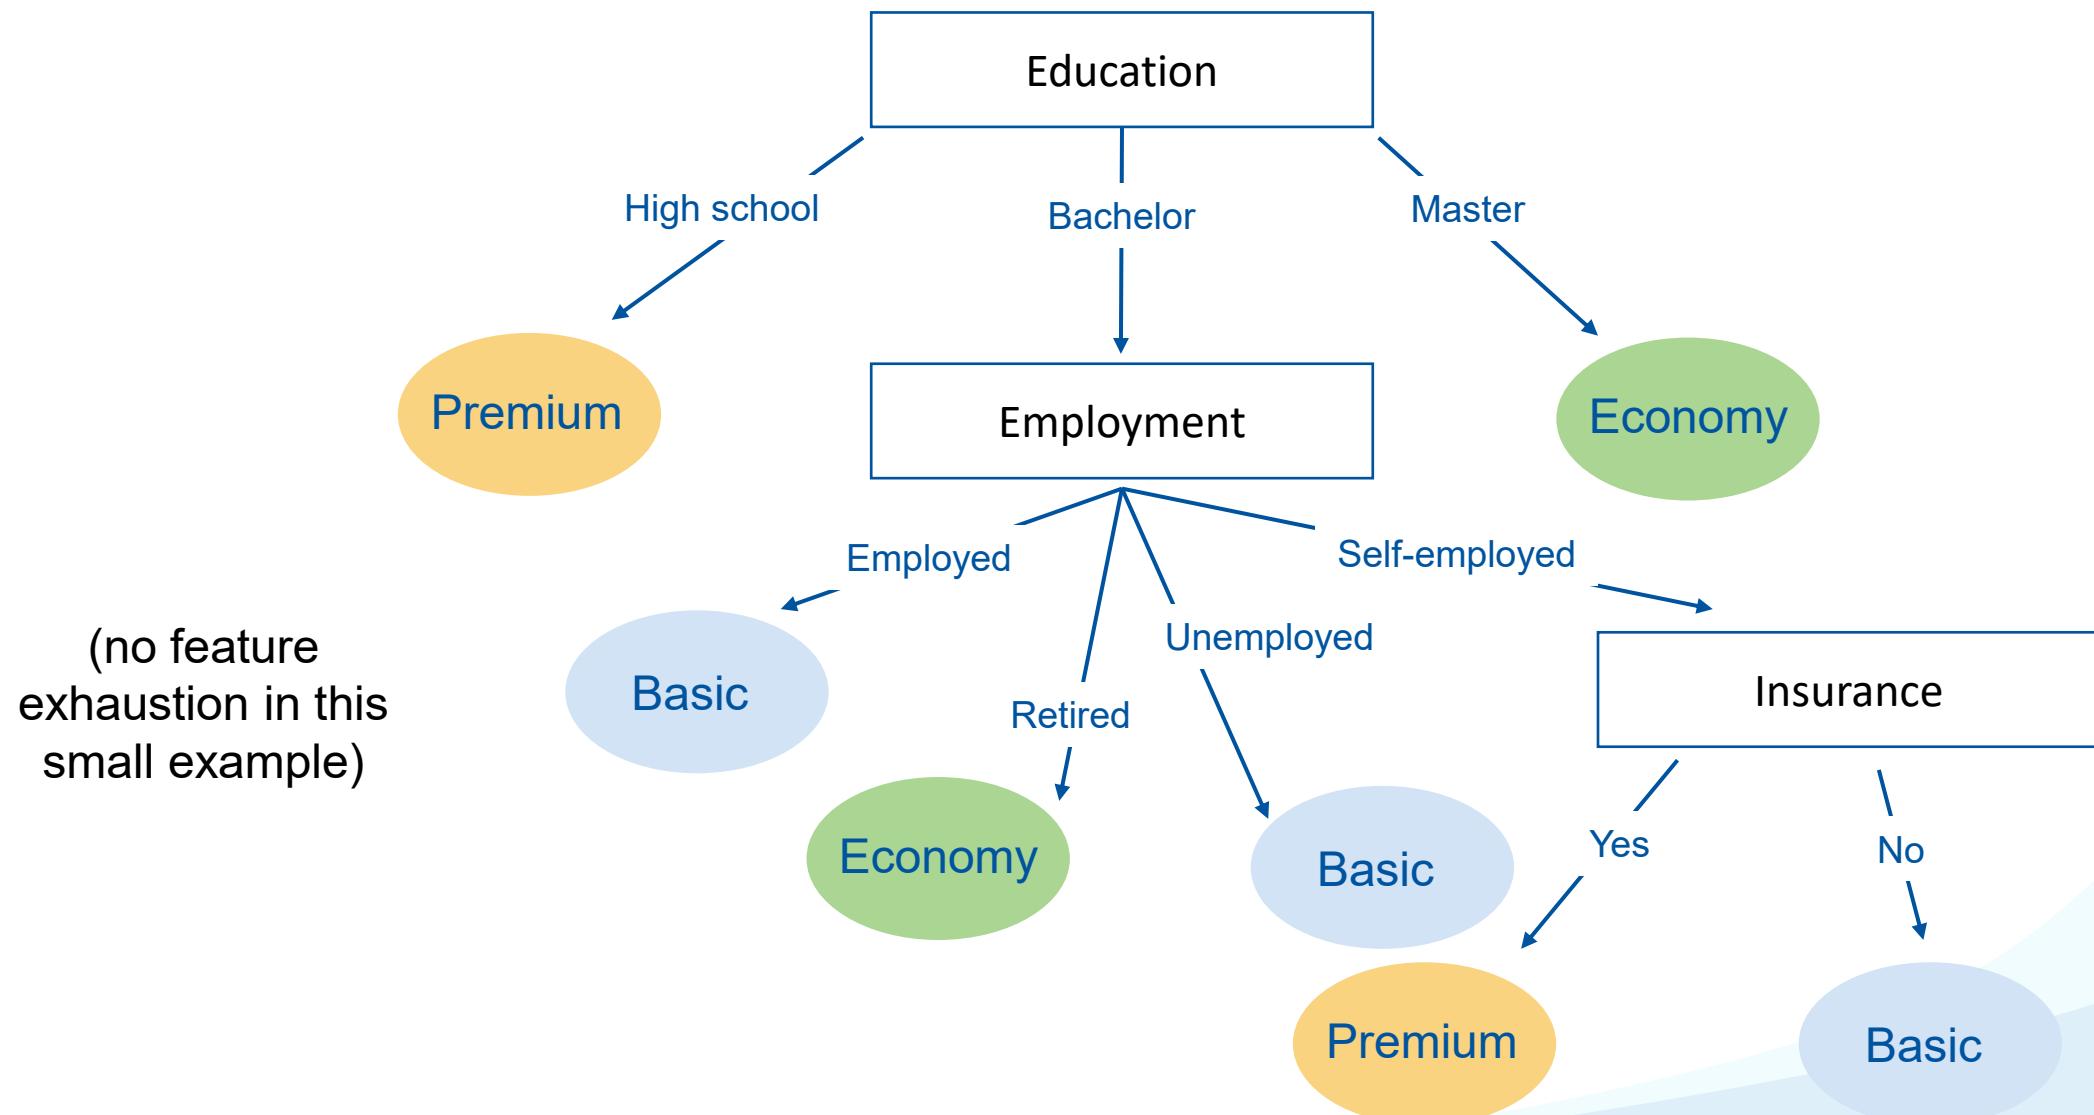
ID	Insurance	Education	Employment	Customer
3	Yes	Bachelor	Self-employed	Premium

consensus

ID	Insurance	Education	Employment	Customer
4	No	Bachelor	Self-employed	Basic

consensus

## Example



# Decision Trees

1. Introduction to Decision Trees
2. Entropy and Information Gain
3. ID3 Algorithm
4. **Quantifying Information Gain**
5. Pruning
6. Continuous Data
7. Ensembles

$$H(t) = - \sum_{k=1}^K (P(t = k) \cdot \log_s(P(t = k)))$$

$$H_W^d(t) = \sum_{node \in nodes(d)} \left( \frac{|node|}{N} \cdot H_{node}(t) \right)$$

$$IG(d) = H(t) - H_W^d(t)$$

## Alternative Information Gain Notions

- Information gain aims to measure the improvement in purity / predictability / compressibility
- Example approaches:
  - Entropy-based information gain (IG)
  - Information gain ratio (GR)
  - Gini index (Gini)
  - Chi-square ( $\chi^2$ )

# Alternative Information Gain Notions

- Information gain aims to measure the improvement in purity / predictability / compressibility
- Example approaches:
  - **Entropy-based information gain (IG)**
  - Information gain ratio (GR)
  - Gini index (Gini)
  - Chi-square ( $\chi^2$ )

Seen before:

Entropy of target feature  $t$  before splitting

$$H(t) = - \sum_{k=1}^K (P(t = k) \cdot \log_s(P(t = k)))$$

$$H_W^d(t) = \sum_{node \in nodes(d)} \left( \frac{|node|}{N} \cdot H_{node}(t) \right)$$

Weighted entropy of target feature  $t$  after splitting based on  $d$

$$IG(d) = H(t) - H_W^d(t)$$

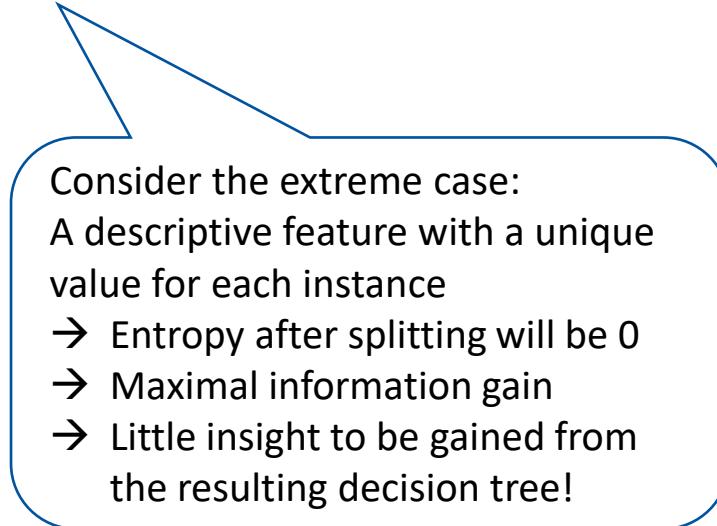
# Alternative Information Gain Notions

- Information gain aims to measure the improvement in purity / predictability / compressibility
- Example approaches:
  - Entropy-based information gain (IG)
  - **Information gain ratio (GR)**
  - **Gini index (Gini)**
  - **Chi-square ( $\chi^2$ )**

not covered in this lecture

## Information Gain Ratio

- Entropy-based information gain favors features with many different values (split in many subsets decreases entropy)
- **Information gain ratio** addresses this issue



Consider the extreme case:  
A descriptive feature with a unique  
value for each instance  
→ Entropy after splitting will be 0  
→ Maximal information gain  
→ Little insight to be gained from  
the resulting decision tree!

# Information Gain Ratio

- Entropy-based information gain favors features with many different values (split in many subsets decreases entropy)
- Information gain ratio addresses this issue:

Information gain when splitting based on descriptive feature  $d$

$$GR(d) = \frac{IG(d)}{H(d)}$$

Entropy of descriptive feature  $d$

→ we can think of it as making an absolute value relative

# Information Gain Ratio

- Entropy-based information gain favors features with many different values (split in many subsets decreases entropy)
- Information gain ratio addresses this issue:

$$GR(d) = \frac{IG(d)}{H(d)} = \frac{H(t) - H_W^d(t)}{-\sum_{k=1}^K (P(d=k) \cdot \log_2(P(d=k)))}$$

Information gain when splitting based on descriptive feature  $d$

Entropy of target feature  $t$

Overall entropy of target feature  $t$  after splitting based on  $d$

Entropy of descriptive feature  $d$

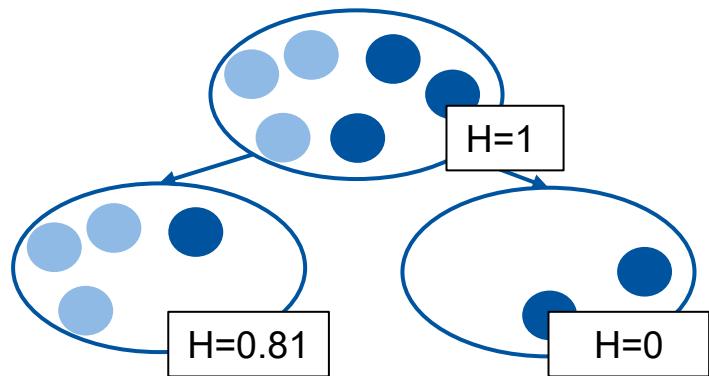
$d$  can take  $K$  possible values

Probability of  $d$  taking the  $k$ th possible value

→ we can think of it as making an absolute value relative

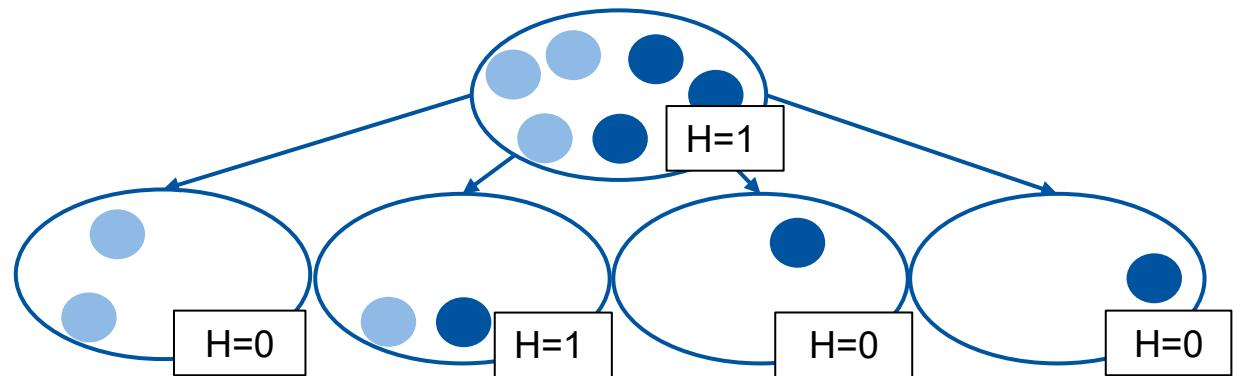
# Information Gain Ratio - Example

split based on feature  $d$



$$H_W^d(\text{color}) = \frac{4}{6} \cdot 0.81 + \frac{2}{6} \cdot 0 = 0.54$$
$$IG(d) = 0.46$$

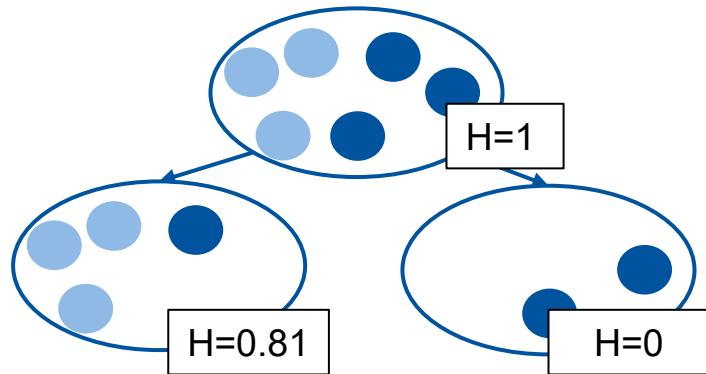
split based on feature  $d'$



$$H_W^{d'}(\text{color}) = \frac{2}{6} \cdot 0 + \frac{2}{6} \cdot 1 \frac{1}{6} \cdot 0 + \frac{1}{6} \cdot 0 = 0.33$$
$$IG(d') = 0.67$$

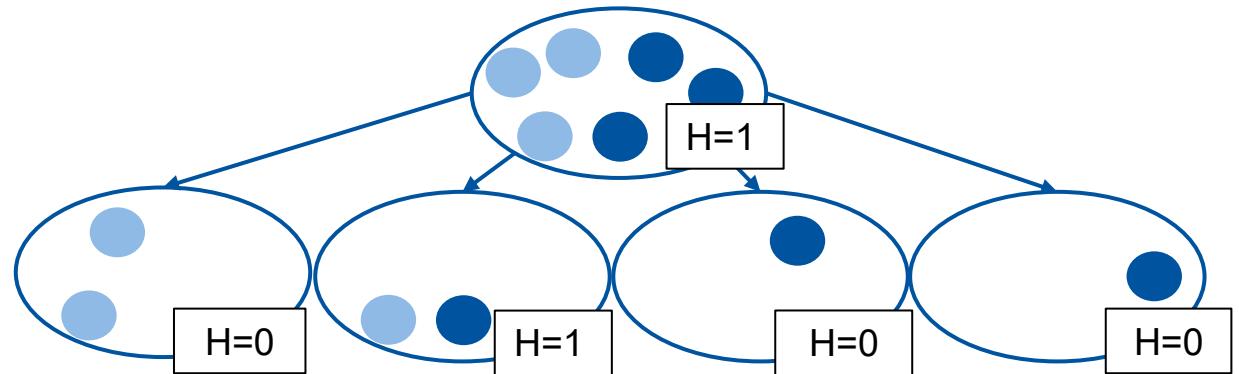
# Information Gain Ratio - Example

split based on feature  $d$



$$IG(d) = 0.46$$

split based on feature  $d'$

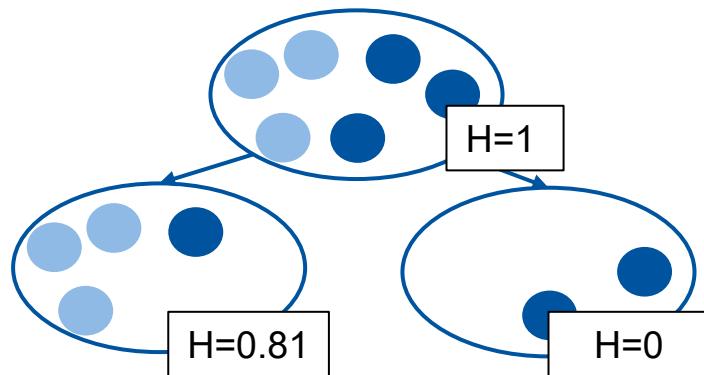


$$IG(d') = 0.67$$

$$GR(d) = \frac{IG(d)}{H(d)} = \frac{H(t) - H_W^d(t)}{-\sum_{k=1}^K (P(d=k) \cdot \log_2(P(d=k)))}$$

# Information Gain Ratio - Example

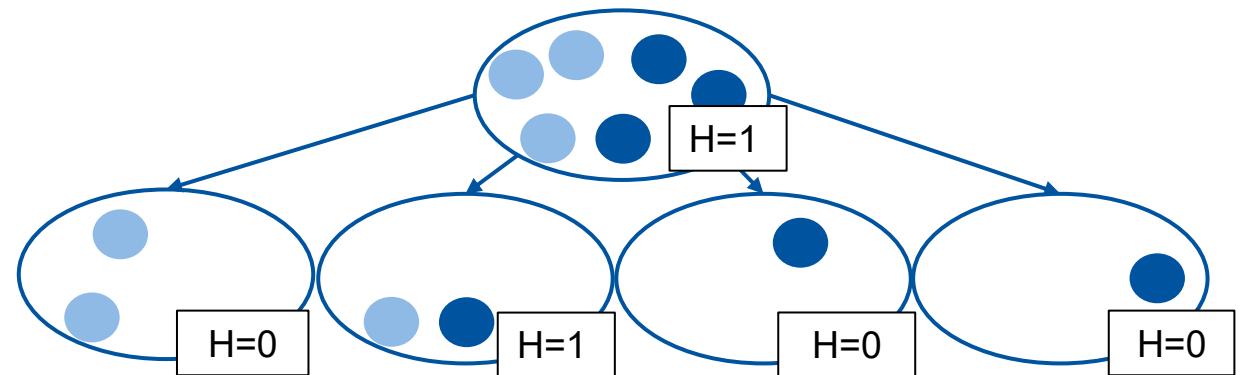
split based on feature  $d$



$$IG(d) = 0.46$$

$$\begin{aligned} GR(d) &= \frac{0.46}{-(\frac{4}{6} \cdot \log_2(\frac{4}{6}) + \frac{2}{6} \cdot \log_2(\frac{2}{6}))} \\ &= \frac{0.46}{0.92} = 0.5 \end{aligned}$$

split based on feature  $d'$



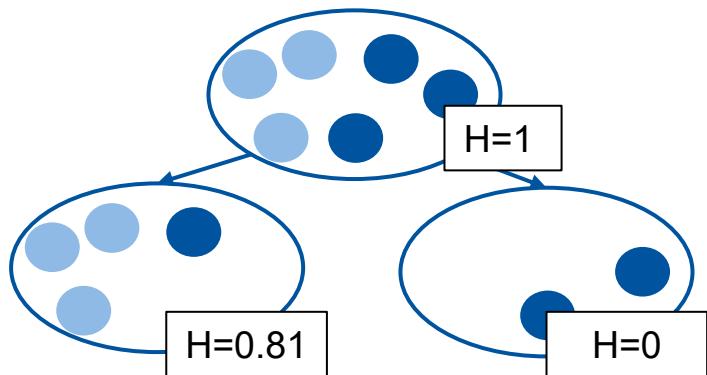
$$IG(d') = 0.67$$

*Feature  $d$  splits the 6 instances into one partition of size 4 and one partition of size 2*

$$GR(d) = \frac{IG(d)}{H(d)} = \frac{H(t) - H_W^d(t)}{-\sum_{k=1}^K (P(d=k) \cdot \log_2(P(d=k)))}$$

# Information Gain Ratio - Example

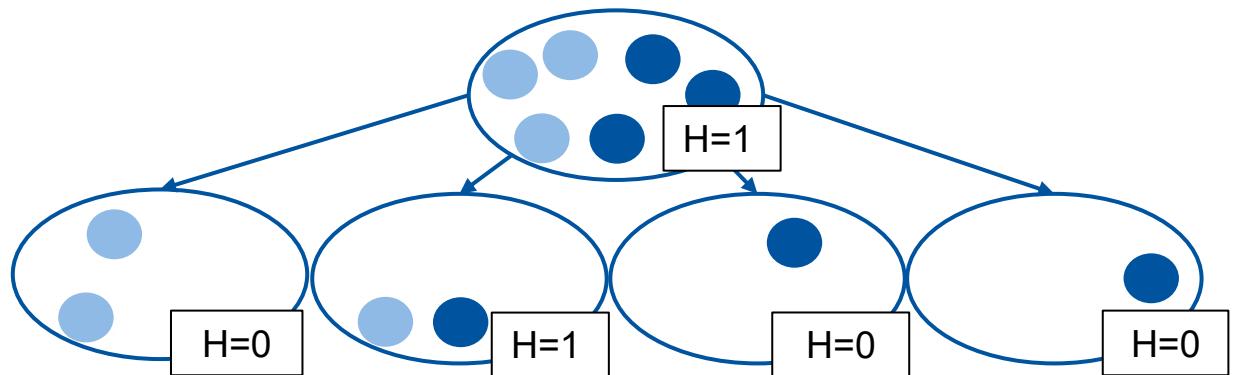
split based on feature  $d$



$$IG(d) = 0.46$$

$$\begin{aligned} GR(d) &= \frac{0.46}{-(\frac{4}{6} \cdot \log_2(\frac{4}{6}) + \frac{2}{6} \cdot \log_2(\frac{2}{6}))} \\ &= \frac{0.46}{0.92} = 0.5 \end{aligned}$$

split based on feature  $d'$



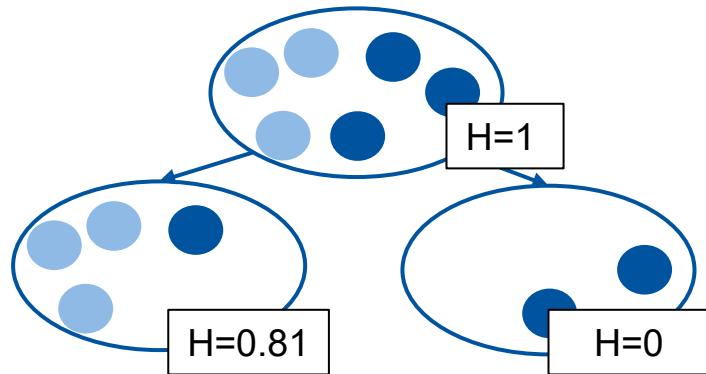
$$IG(d') = 0.67$$

$$\begin{aligned} GR(d') &= \frac{0.67}{-(\frac{2}{6} \cdot \log_2(\frac{2}{6}) + \frac{2}{6} \cdot \log_2(\frac{2}{6}) + \frac{1}{6} \cdot \log_2(\frac{1}{6}) + \frac{1}{6} \cdot \log_2(\frac{1}{6}))} \\ &= \frac{0.67}{1.92} = 0.35 \end{aligned}$$

$$GR(d) = \frac{IG(d)}{H(d)} = \frac{H(t) - H_W^d(t)}{-\sum_{k=1}^K (P(d=k) \cdot \log_2(P(d=k)))}$$

# Information Gain Ratio - Example

split based on feature  $d$



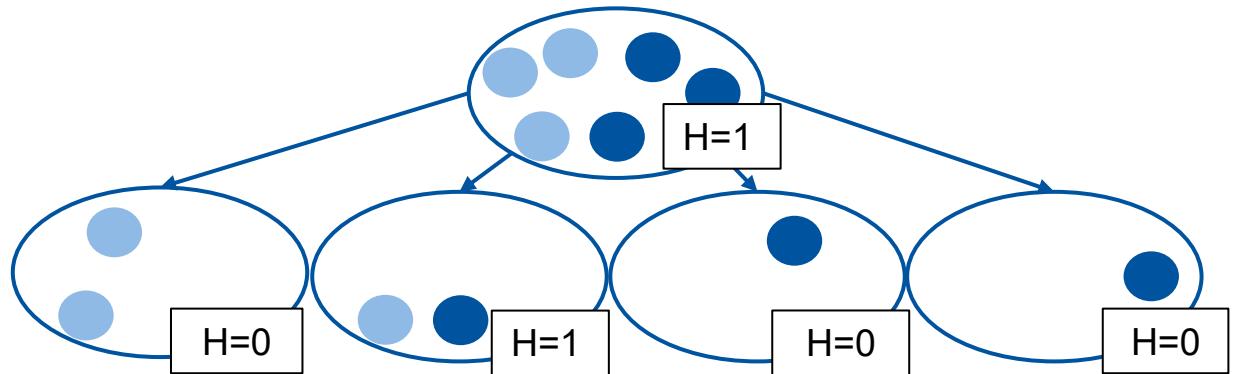
$$IG(d) = 0.46$$



$$\begin{aligned} GR(d) &= \frac{0.46}{-(\frac{4}{6} \cdot \log_2(\frac{4}{6}) + \frac{2}{6} \cdot \log_2(\frac{2}{6}))} \\ &= \frac{0.46}{0.92} = 0.5 \end{aligned}$$



split based on feature  $d'$



$$IG(d') = 0.67$$

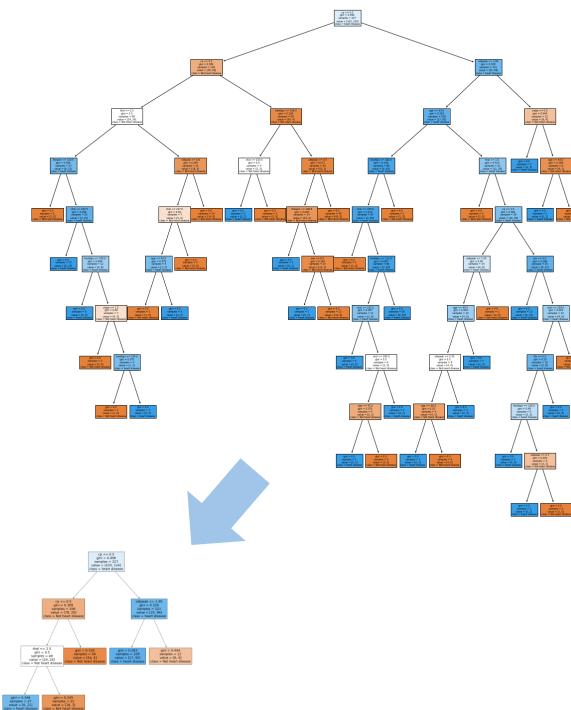


$$\begin{aligned} GR(d') &= \frac{0.67}{-(\frac{2}{6} \cdot \log_2(\frac{2}{6}) + \frac{2}{6} \cdot \log_2(\frac{2}{6}) + \frac{1}{6} \cdot \log_2(\frac{1}{6}) + \frac{1}{6} \cdot \log_2(\frac{1}{6}))} \\ &= \frac{0.67}{1.92} = 0.35 \end{aligned}$$

$$GR(d) = \frac{IG(d)}{H(d)} = \frac{H(t) - H_W^d(t)}{-\sum_{k=1}^K (P(d=k) \cdot \log_2(P(d=k)))}$$

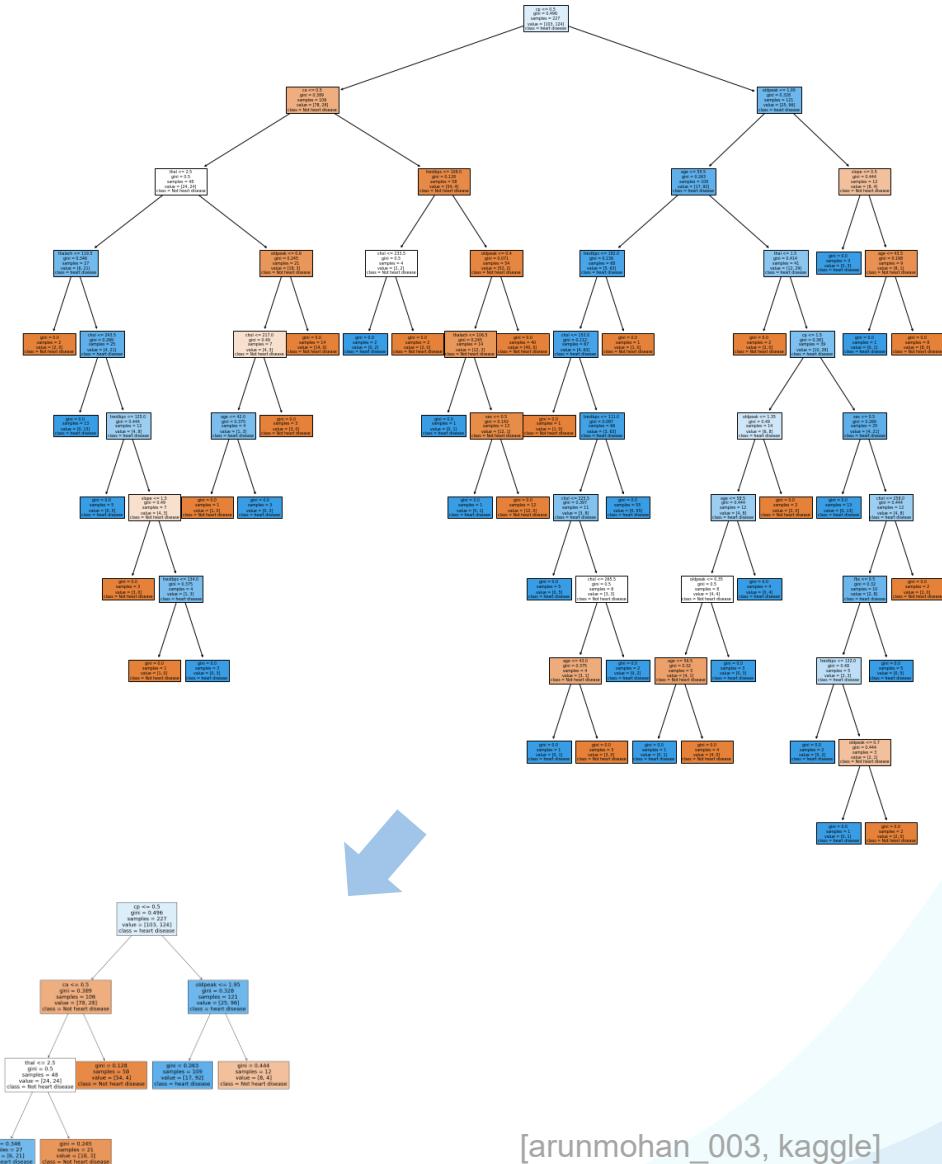
# Decision Trees

1. Introduction to Decision Trees
2. Entropy and Information Gain
3. ID3 Algorithm
4. Quantifying Information Gain
5. **Pruning**
6. Continuous Data
7. Ensembles



# Pruning Decision Trees

- Possible problems:
  - Decision tree is **overfitting** the data
  - Decision tree is too complex or too deep
- Two solution directions:
  - Pre-pruning** (early stopping/forward)
  - Post-pruning** (reduced error/backward)
- To generalize and **avoid overfitting**



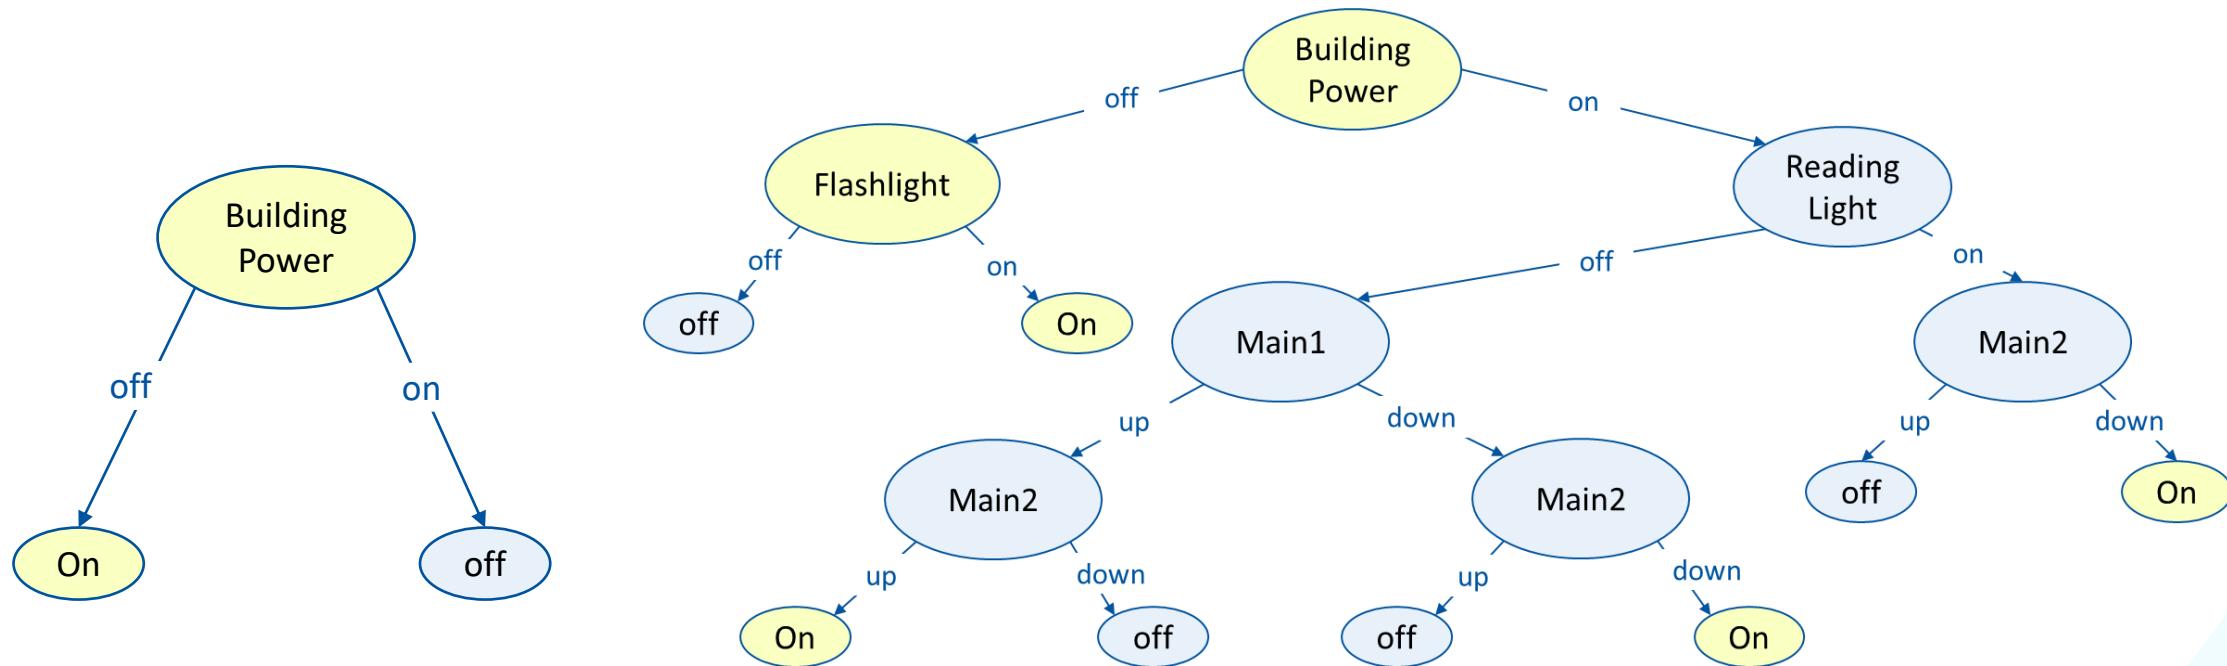
[arunmohan\_003, kaggle]

## Pre-pruning

- Idea: Stop creating subtrees and use **majority vote** to determine the label
- Many possible **stopping criteria**:
  - maximum tree depth
  - lower bound for number of instances before split
  - Lower bound for number of instances after split
  - lower bound for information gain
  - ...
- May create trees that are **not consistent** with respect to the data

# Pre-pruning – Enlightening Example

## Building the tree with pre-pruning

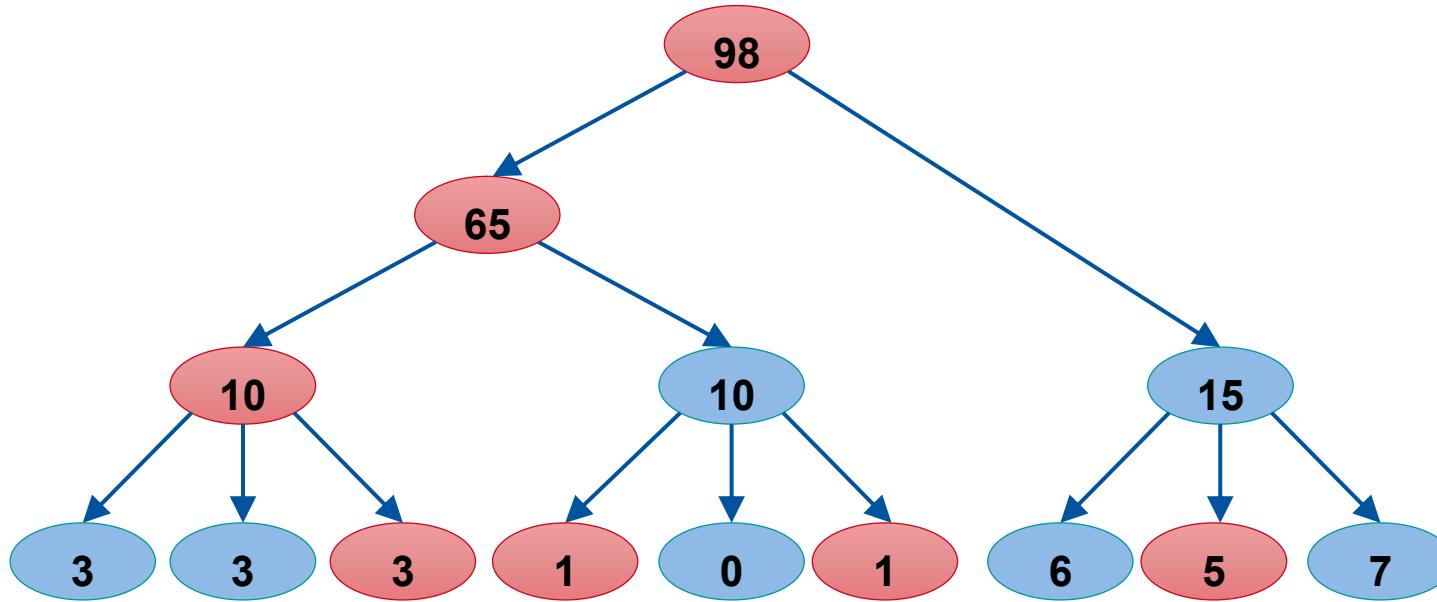


Pre-pruning is efficient, but we may miss interesting dependencies at lower levels of the tree

## Post-pruning

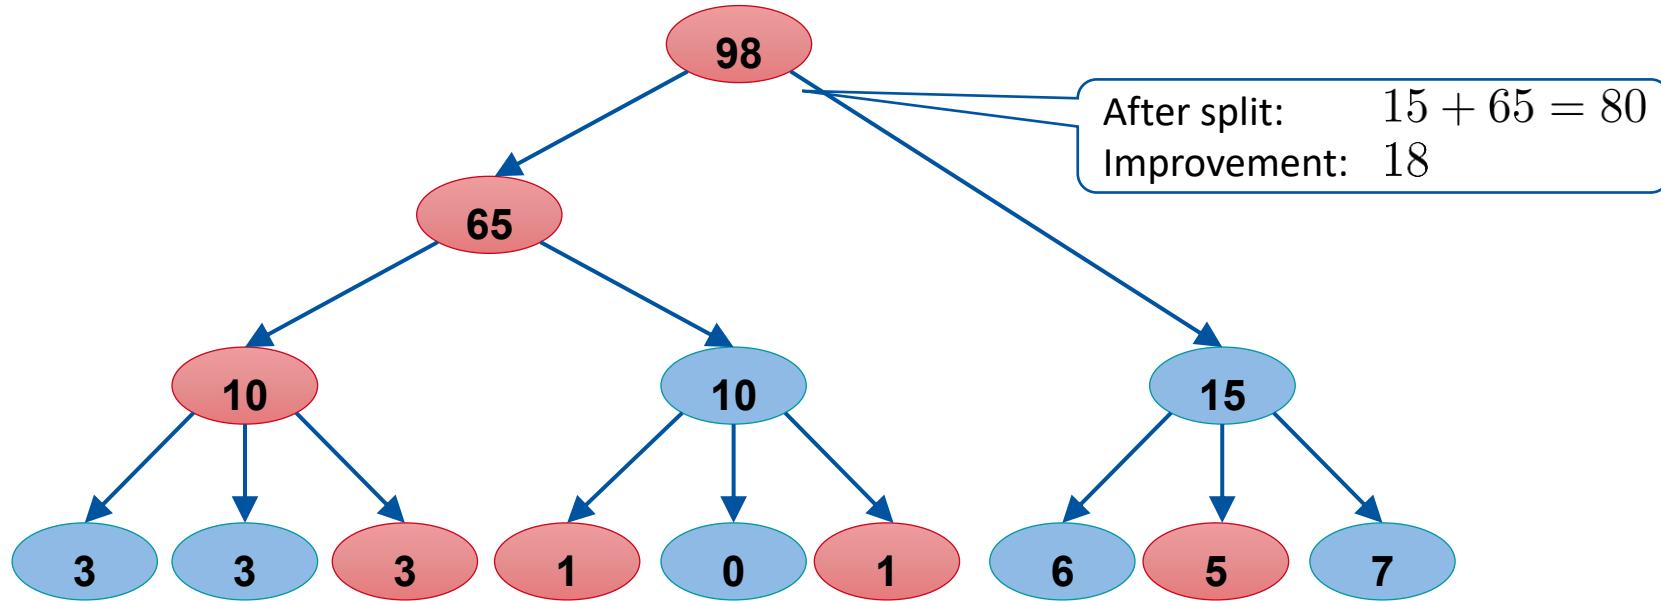
- First, build the **whole** decision tree; then **cut off branches** that do not add (much)
- Common approach is to **split the data** into a training set and a validation/test set
- Measure the **performance of splits** based on a validation/test set

# Post-pruning



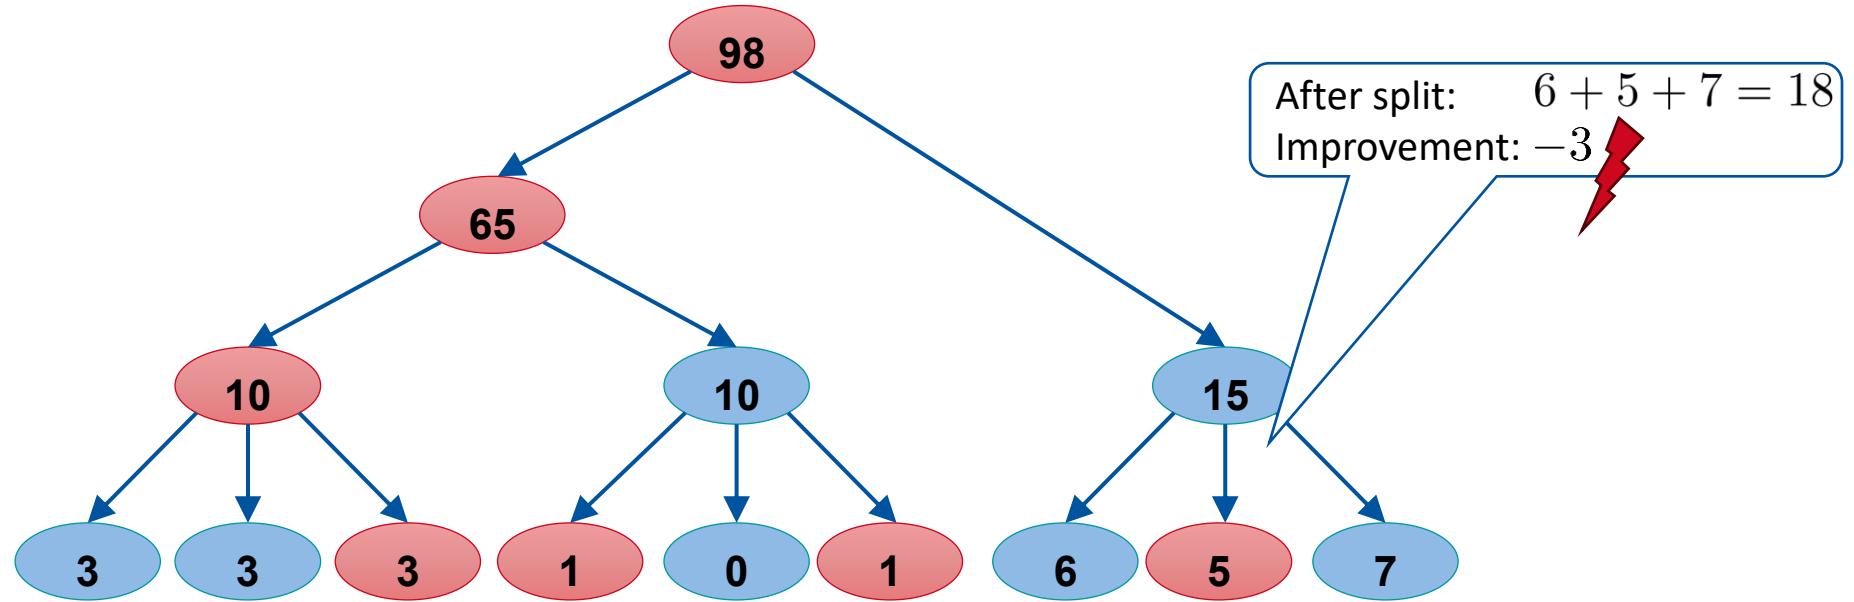
- Decision tree learned on a **training set**
- Numbers indicate misclassifications based on a **validation set**

# Post-pruning



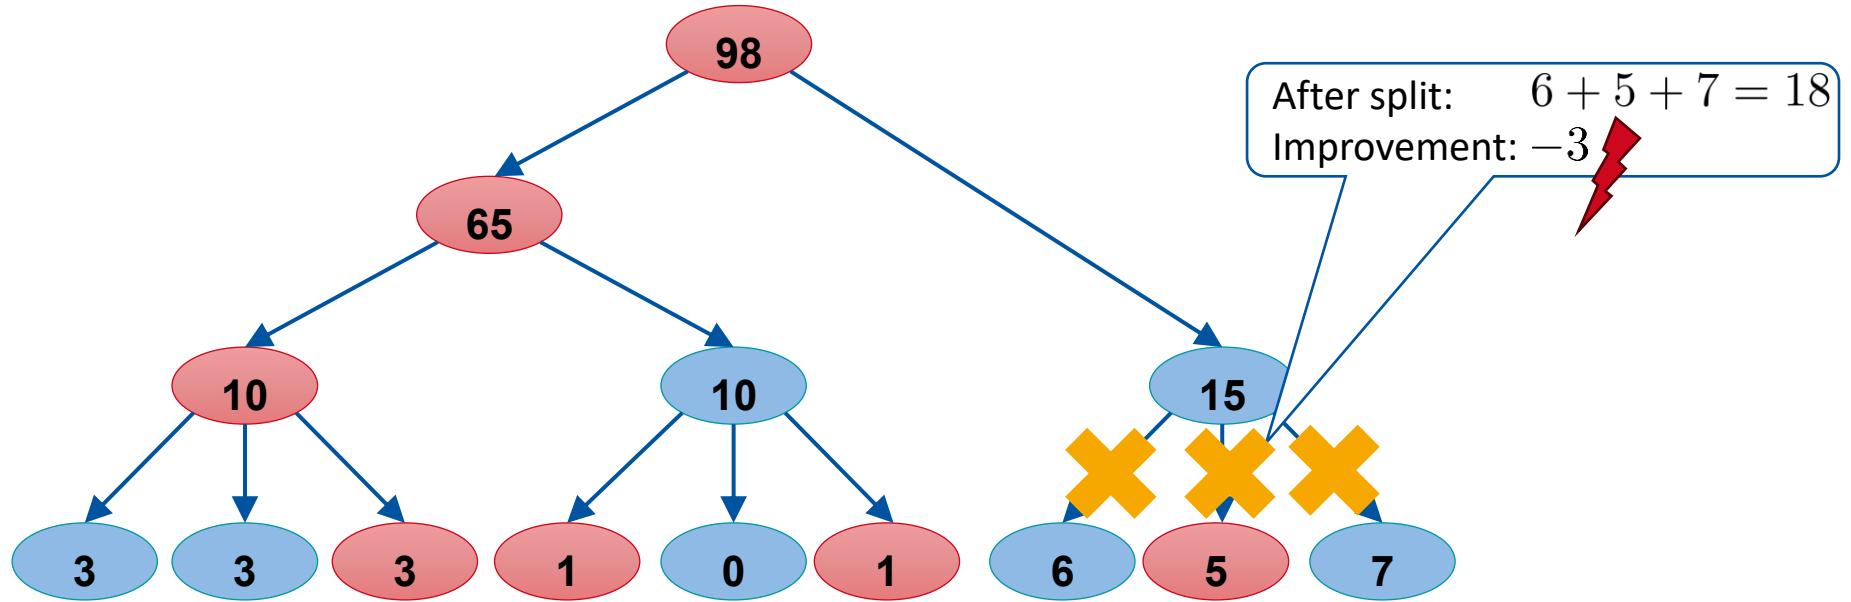
- Decision tree learned on a **training set**
- Numbers indicate misclassifications based on a **validation set**

# Post-pruning



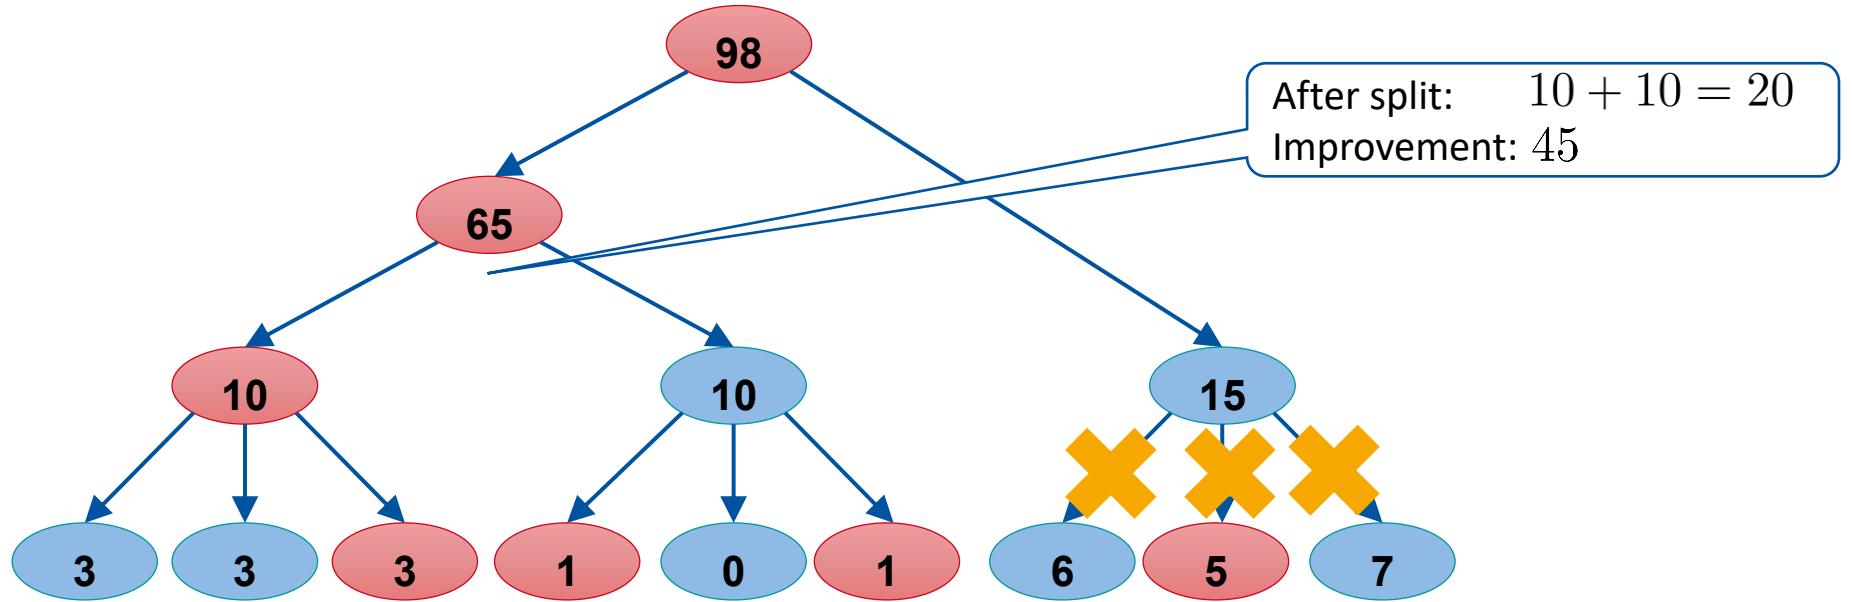
- Decision tree learned on a **training set**
- Numbers indicate misclassifications based on a **validation set**

# Post-pruning



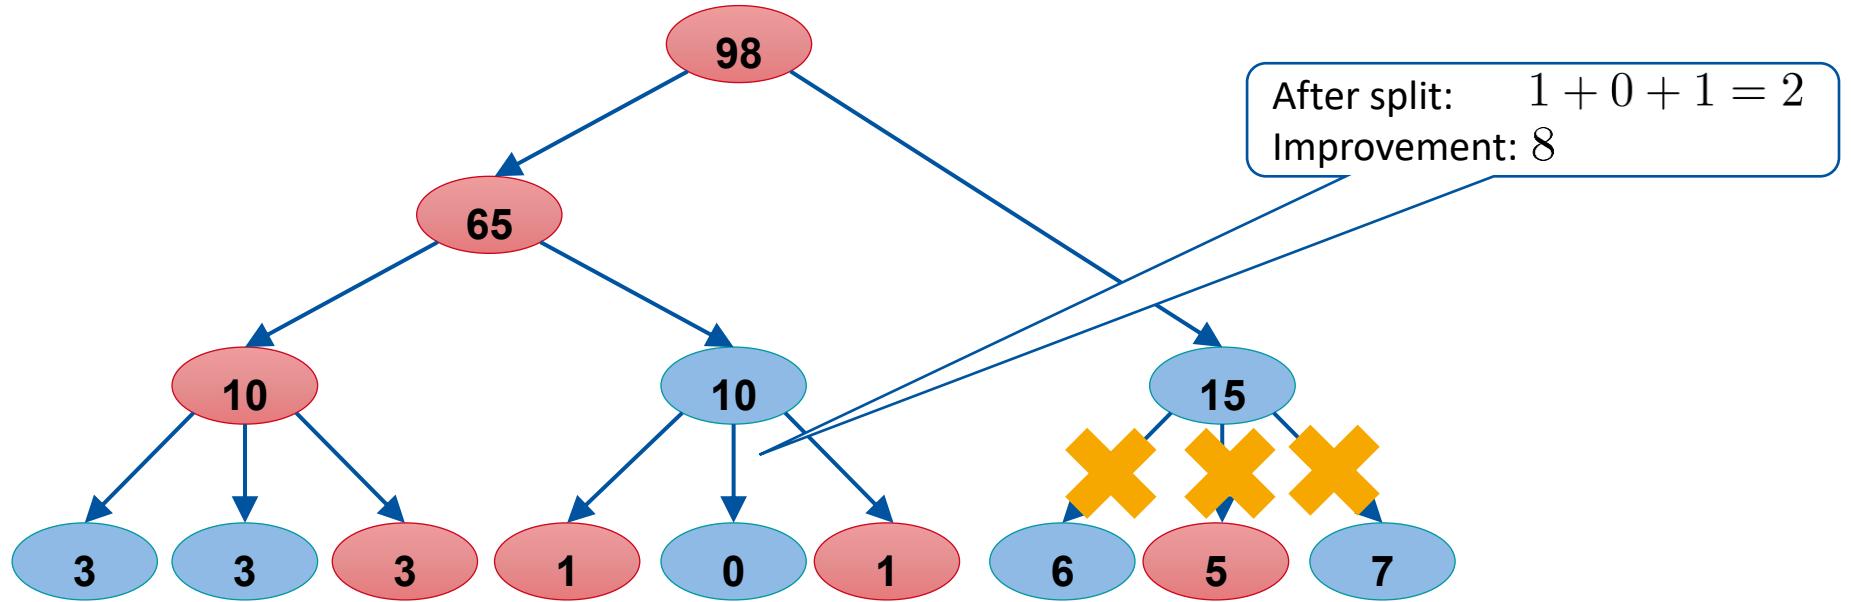
- Decision tree learned on a **training set**
- Numbers indicate misclassifications based on a **validation set**

# Post-pruning



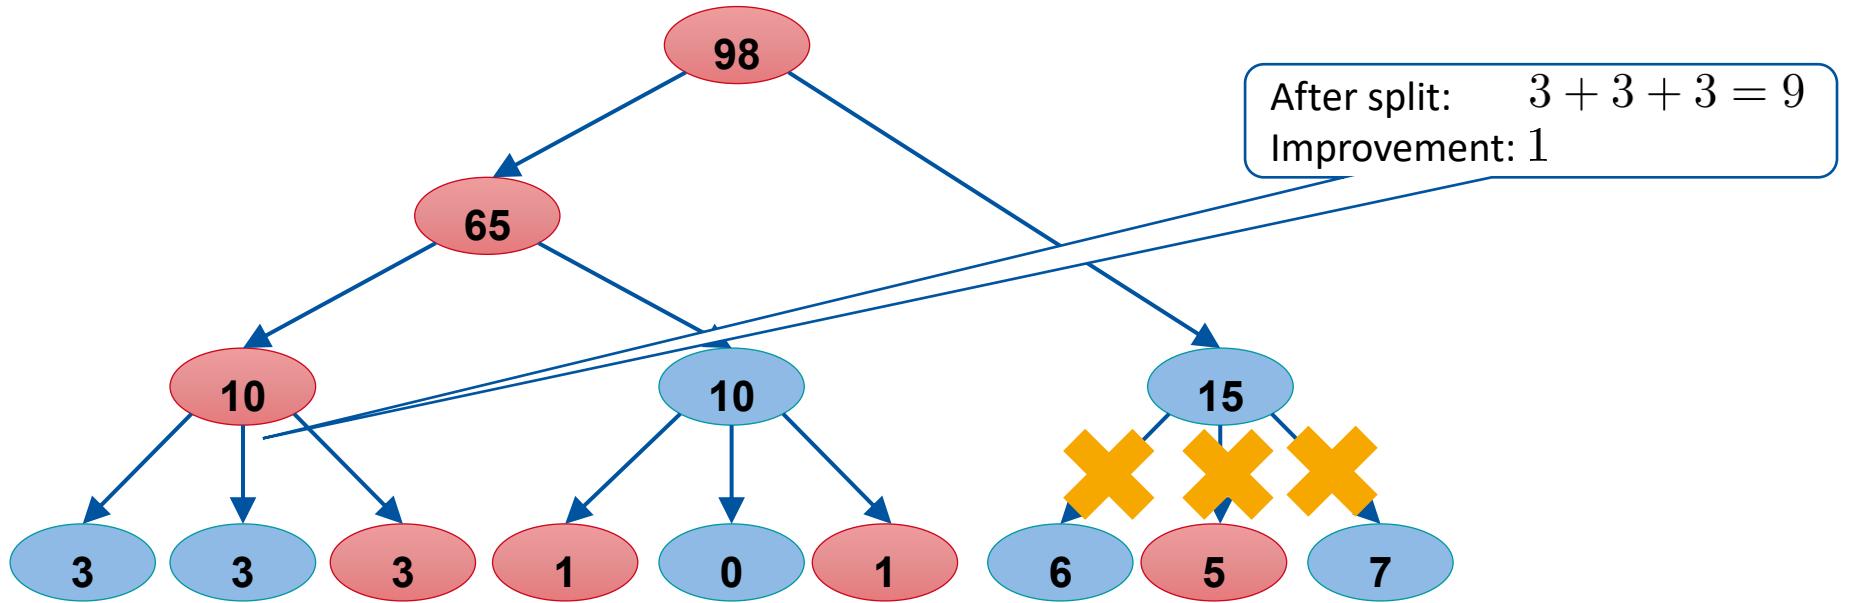
- Decision tree learned on a **training set**
- Numbers indicate misclassifications based on a **validation set**

# Post-pruning



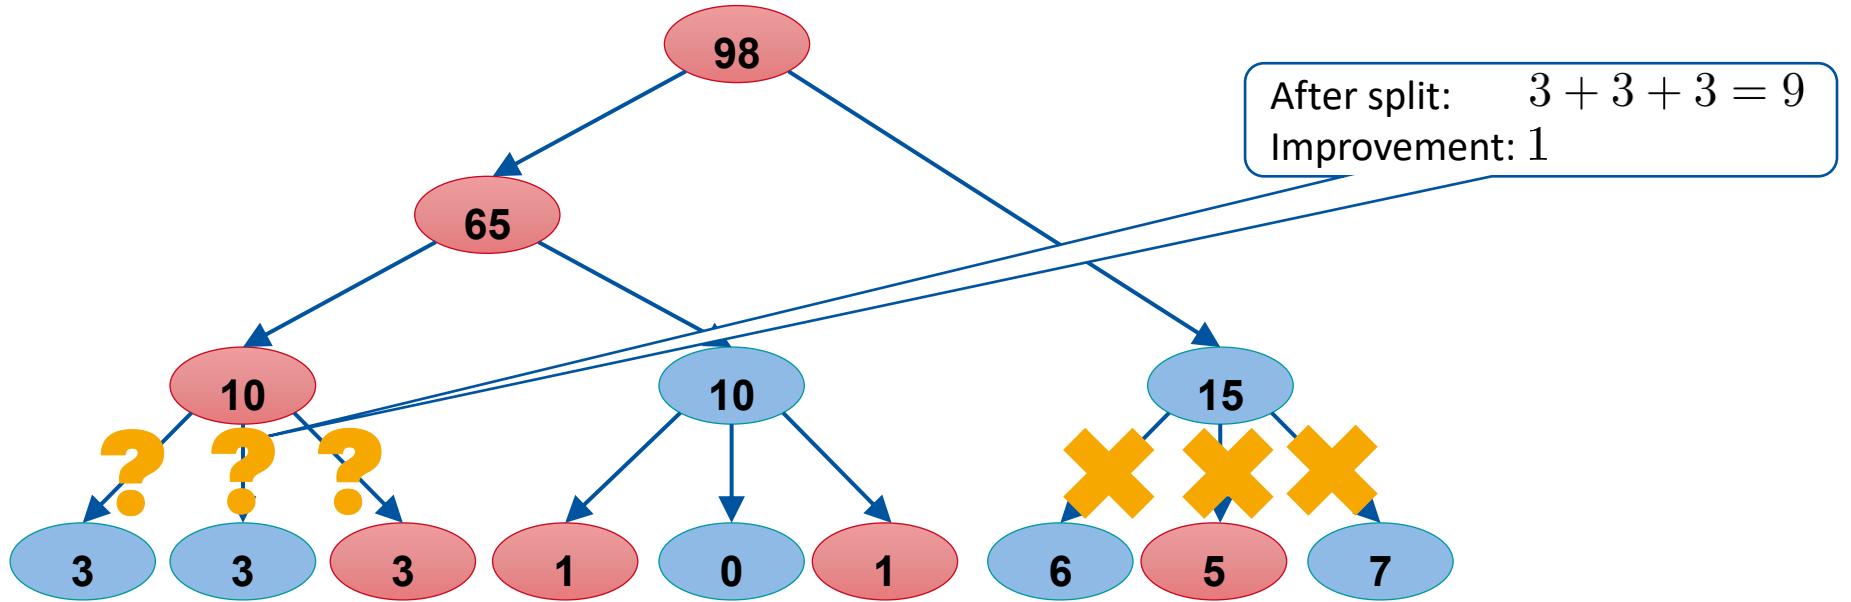
- Decision tree learned on a **training set**
- Numbers indicate misclassifications based on a **validation set**

# Post-pruning



- Decision tree learned on a **training set**
- Numbers indicate misclassifications based on a **validation set**

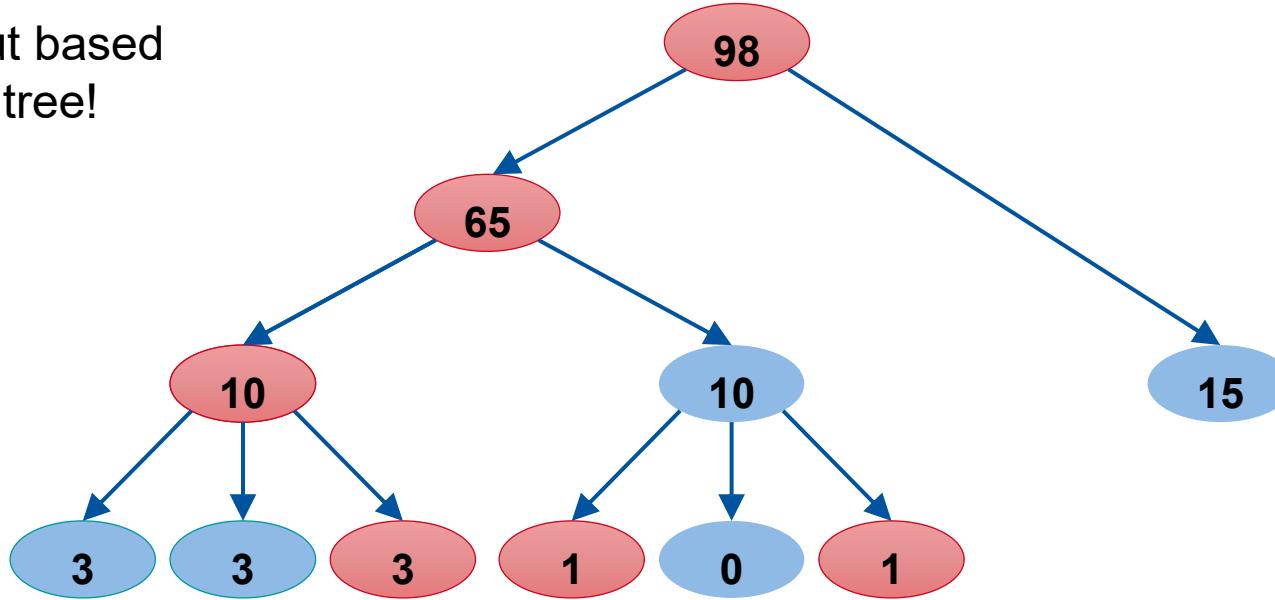
# Post-pruning



- Decision tree learned on a **training set**
- Numbers indicate misclassifications based on a **validation set**

## Post-pruning

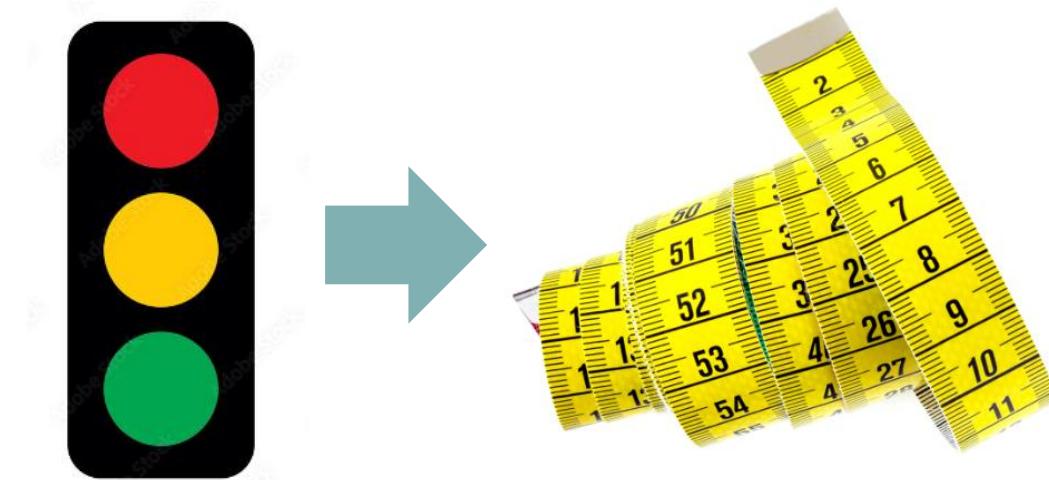
Less efficient, but based on the complete tree!



- Many variants
  - Bottom up instead of top-down
  - Quantification of Performance (e.g., normalize misclassifications)
  - Various pruning thresholds (in this lecture, subtrees are pruned when misclassification increases)
  - ...

# Decision Trees

1. Introduction to Decision Trees
2. Entropy and Information Gain
3. ID3 Algorithm
4. Quantifying Information Gain
5. Pruning
6. **Continuous Data**
7. Ensembles



# Dealing with Continuous Variables

- So far we assumed features were **categorical**
- We can use **binning** to make continuous features categorical

The diagram illustrates a dataset structure. On the left, a vertical column labeled "instances" lists five rows. To the right of this column is a horizontal row labeled "features". This row contains five columns:  $f_1$ ,  $f_2$ , ...,  $f_D$ , and "class". The "class" column contains numerical values: 5043, 4598, 3248, 5466, and 7682. A callout bubble points to the "class" column with the text "continuous target feature". Another callout bubble points to the  $f_2$  column with the text "continuous descriptive features".

instances	features				class
	$f_1$	$f_2$	...	$f_D$	
	high	88		59.99	5043
	high	76		50.00	4598
	low	32		39.50	3248
	low	89		49.99	5466
	high	21		59.99	7682

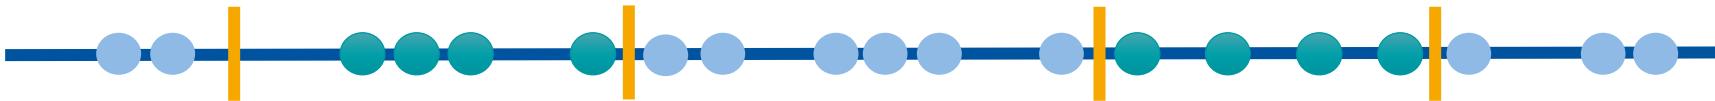
## Continuous Descriptive Features

- Challenge: determine suitable boundaries (infinite number of thresholds is possible)



# Continuous Descriptive Features

- Challenge: determine suitable boundaries (infinite number of thresholds is possible)
- Idea:
  - sort instances based on the continuous descriptive feature
  - look for changes in target feature labels
- Change points are candidate thresholds
- Select the threshold with the highest information gain



## Continuous Descriptive Features - Example

ID	Insurance	Income	Employment	Customer
1	Yes	3500	Employed	Basic
2	Yes	0	Unemployed	Premium
3	Yes	1000	Self-employed	Premium
4	No	2000	Self-employed	Basic
5	No	5000	Employed	Economy
6	Yes	5100	Retired	Economy
7	Yes	3000	Employed	Premium

sort 

## Continuous Descriptive Features - Example

ID	Insurance	Income	Employment	Customer
2	Yes	0	Unemployed	Premium
3	Yes	1000	Self-employed	Premium
4	No	2000	Self-employed	Basic
7	Yes	3000	Employed	Premium
1	Yes	3500	Employed	Basic
5	No	5000	Employed	Economy
6	Yes	5100	Retired	Economy

sort 

## Continuous Descriptive Features - Example

ID	Insurance	Income	Employment	Customer
2	Yes	0	Unemployed	Premium
3	Yes	1000	Self-employed	Premium
4	No	2000	Self-employed	Basic
7	Yes	3000	Employed	Premium
1	Yes	3500	Employed	Basic
5	No	5000	Employed	Economy
6	Yes	5100	Retired	Economy

Change in target feature:  
candidate threshold

# Continuous Descriptive Features - Example

ID	Insurance	Income	Employment	Customer
2	Yes	0	Unemployed	Premium
3	Yes	1000	Self-employed	Premium
4	No	2000	Self-employed	Basic
7	Yes	3000	Employed	Premium
1	Yes	3500	Employed	Basic
5	No	5000	Employed	Economy
6	Yes	5100	Retired	Economy

Change in target feature:  
candidate threshold

Thresholds: middle values of continuous feature in between changed target features

# Continuous Descriptive Features - Example

ID	Insurance	Income	Employment	Customer
2	Yes	0	Unemployed	Premium
3	Yes	1000	Self-employed	Premium
4	No	2000	Self-employed	Basic
7	Yes	3000	Employed	Premium
1	Yes	3500	Employed	Basic
5	No	5000	Employed	Economy
6	Yes	5100	Retired	Economy

Change in target feature:  
candidate threshold

Thresholds: middle values of continuous feature in between changed target features

# Continuous Descriptive Features - Example

ID	Insurance	Income	Employment	Customer
2	Yes	0	Unemployed	Premium
3	Yes	1500	Self-employed	Premium
4	No	2500	Self-employed	Basic
7	Yes	3250	Employed	Premium
1	Yes	4250	Employed	Basic
5	No	5000	Employed	Economy
6	Yes	5100	Retired	Economy

Four candidate thresholds

Thresholds: middle values of continuous feature in between changed target features

## Continuous Descriptive Features - Example

Threshold	Instances
$\geq 1500$	2, 3 1, 4, 5, 6, 7
$\geq 2500$	2, 3, 4 1, 5, 6, 7
$\geq 3250$	2, 3, 4, 7 1, 5, 6
$\geq 4250$	1, 2, 3, 4, 7 5, 6

# Continuous Descriptive Features - Example

Threshold	Instances	Partition Entropy	Overall Entropy	Information Gain
$\geq 1500$	2, 3	0	1.0871	0.1981
	1, 4, 5, 6, 7	1.5219		
$\geq 2500$	2, 3, 4	0.9183	1.2507	0.306
	1, 5, 6, 7	1.5		
$\geq 3250$	2, 3, 4, 7	0.8113	0.8572	0.6995
	1, 5, 6	0.9183		
$\geq 4250$	1, 2, 3, 4, 7	0.9710	0.6935	0.8631
	5, 6	0		

Compute  
as usual

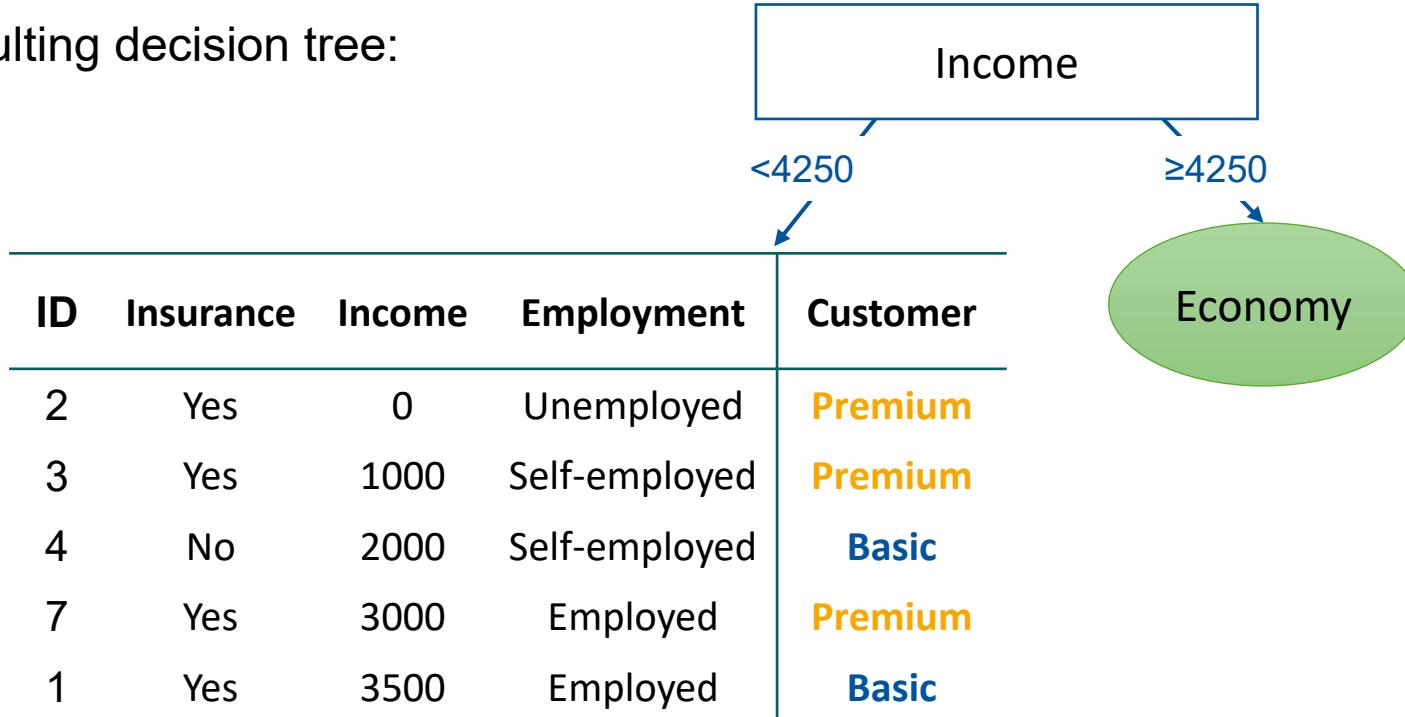
# Continuous Descriptive Features - Example

Threshold	Instances	Partition Entropy	Overall Entropy	Information Gain
$\geq 1500$	2, 3	0	1.0871	0.1981
	1, 4, 5, 6, 7	1.5219		
$\geq 2500$	2, 3, 4	0.9183	1.2507	0.306
	1, 5, 6, 7	1.5		
$\geq 3250$	2, 3, 4, 7	0.8113	0.8572	0.6995
	1, 5, 6	0.9183		
$\geq 4250$	1, 2, 3, 4, 7	0.9710	0.6935	0.8631
	5, 6	0		

best

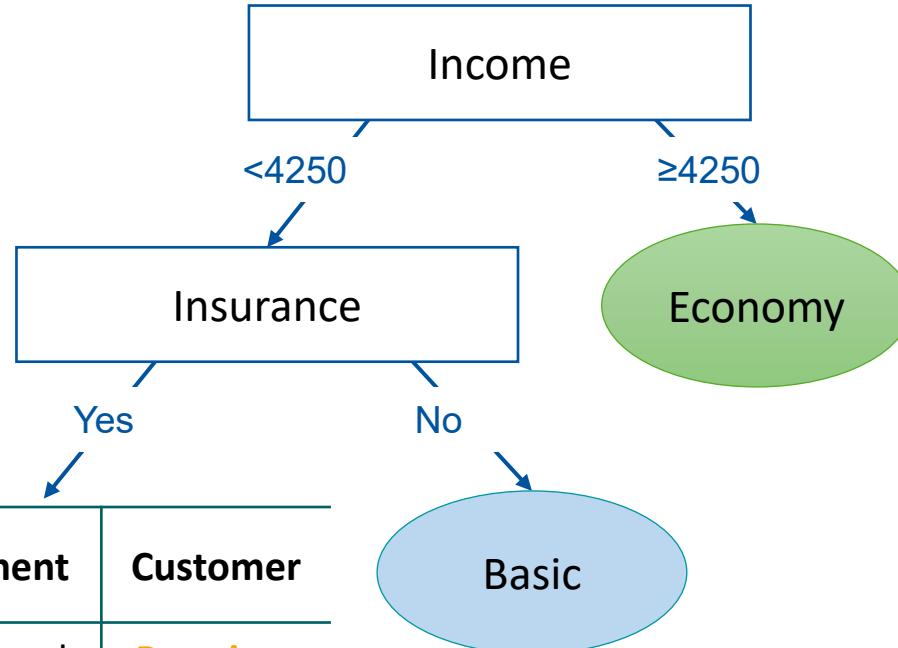
## Continuous Descriptive Features - Example

Resulting decision tree:



# Continuous Descriptive Features - Example

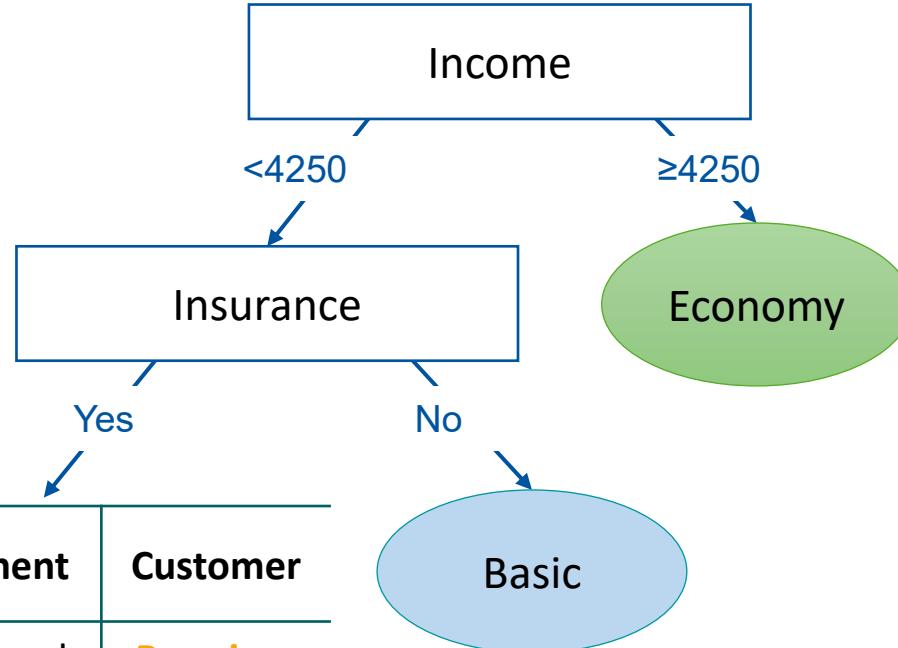
Resulting decision tree:



ID	Insurance	Income	Employment	Customer
2	Yes	0	Unemployed	Premium
3	Yes	1000	Self-employed	Premium
7	Yes	3000	Employed	Premium
1	Yes	3500	Employed	Basic

# Continuous Descriptive Features - Example

Resulting decision tree:

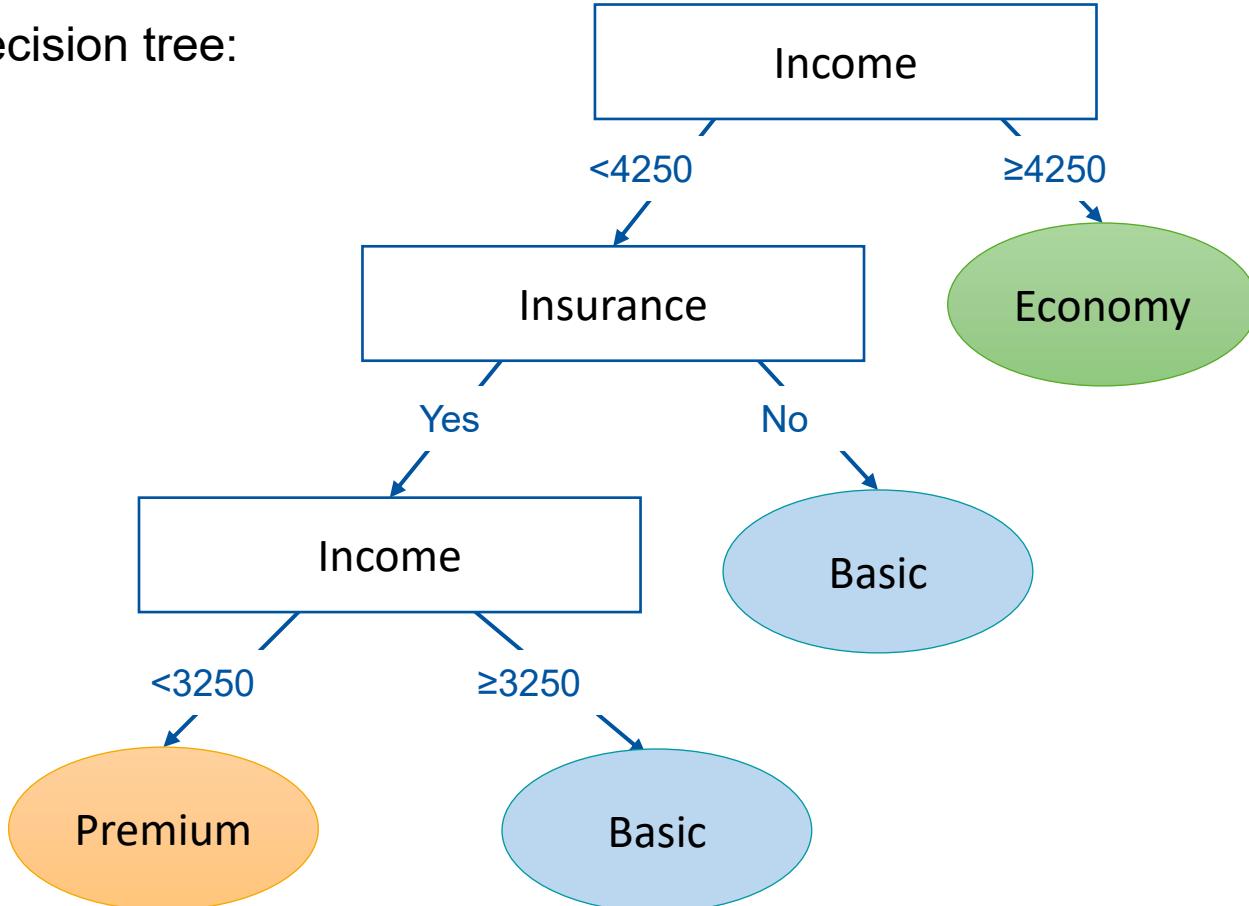


ID	Insurance	Income	Employment	Customer
2	Yes	0	Unemployed	Premium
3	Yes	1000	Self-employed	Premium
7	Yes	3000	Employed	Premium
1	Yes	3500	Employed	Basic

The same continuous feature can now be used multiple times!

# Continuous Descriptive Features - Example

Resulting decision tree:



The same continuous feature can now be used multiple times!

## Continuous Target Features

- Goal: find descriptive features that ‘nicely’ partition the target feature axis
- Impurity = Variance within a partition
- We cannot use the target feature itself
- We ‘color the dots’ based on a selected descriptive feature



# Continuous Target Features

- Goal: find descriptive features that ‘nicely’ partition the target feature axis
- Impurity = Variance within a partition
- We cannot use the target feature itself
- We ‘color the dots’ based on a selected descriptive feature

Intuition: instances described as green are predicted to have a value in this range

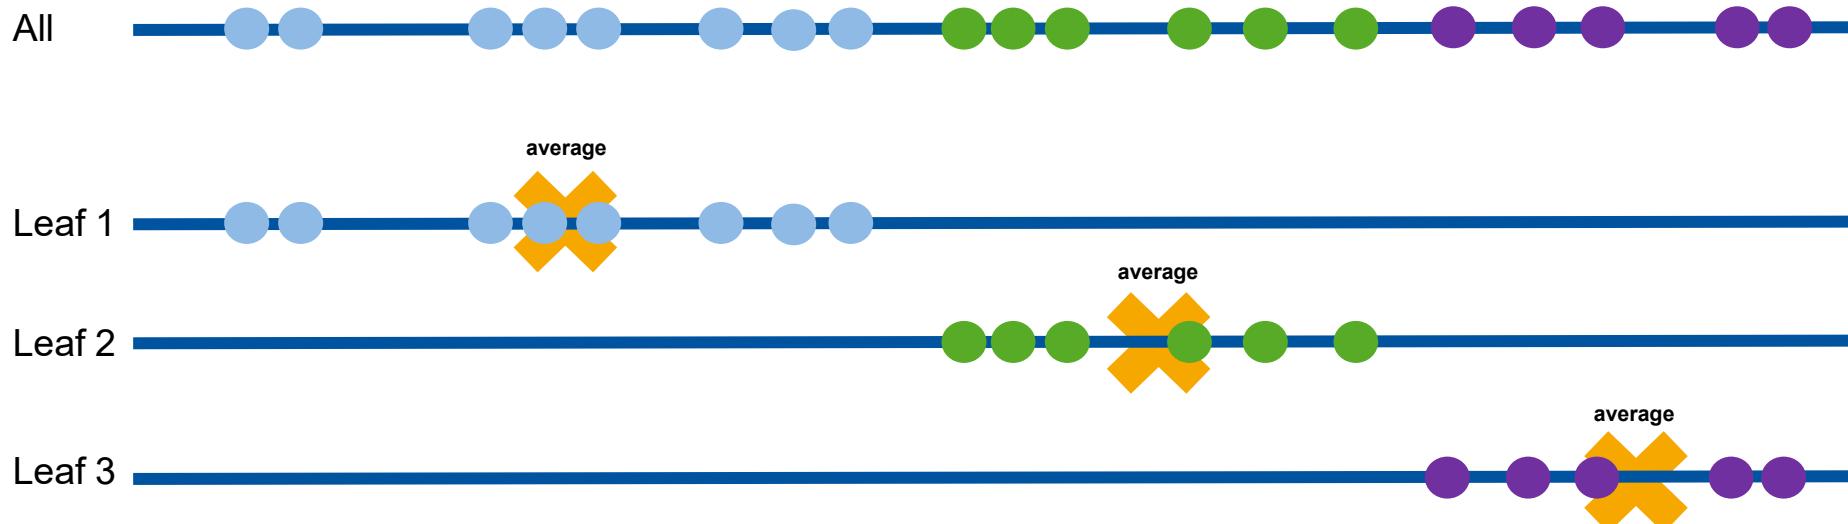


# Continuous Target Features

Good Classification



- Three leaves (purple, green, blue show mapping based on descriptive feature)
- Impurity as measure of quality: variance within a leaf of the decision tree

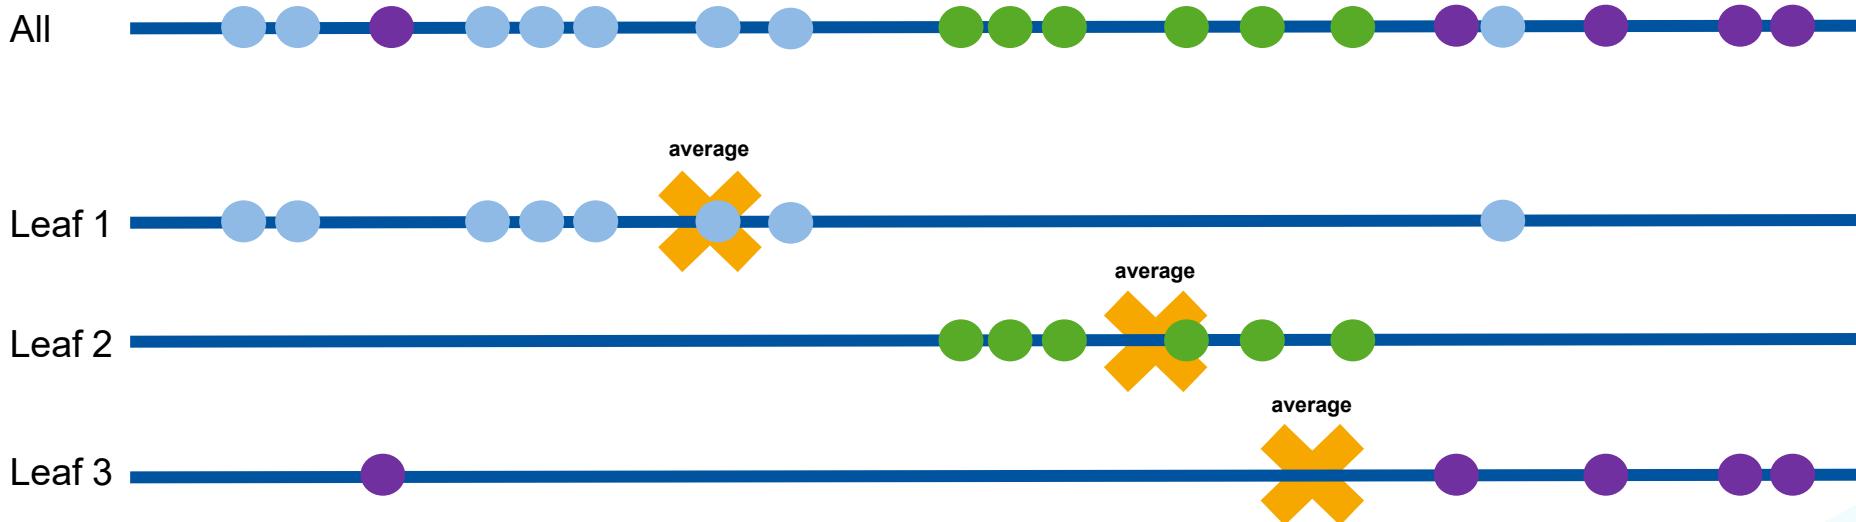


# Continuous Target Features



## Reasonable Classification

Variance within Leaf 1 and Leaf 3 increased with respect to the ‘good classification’

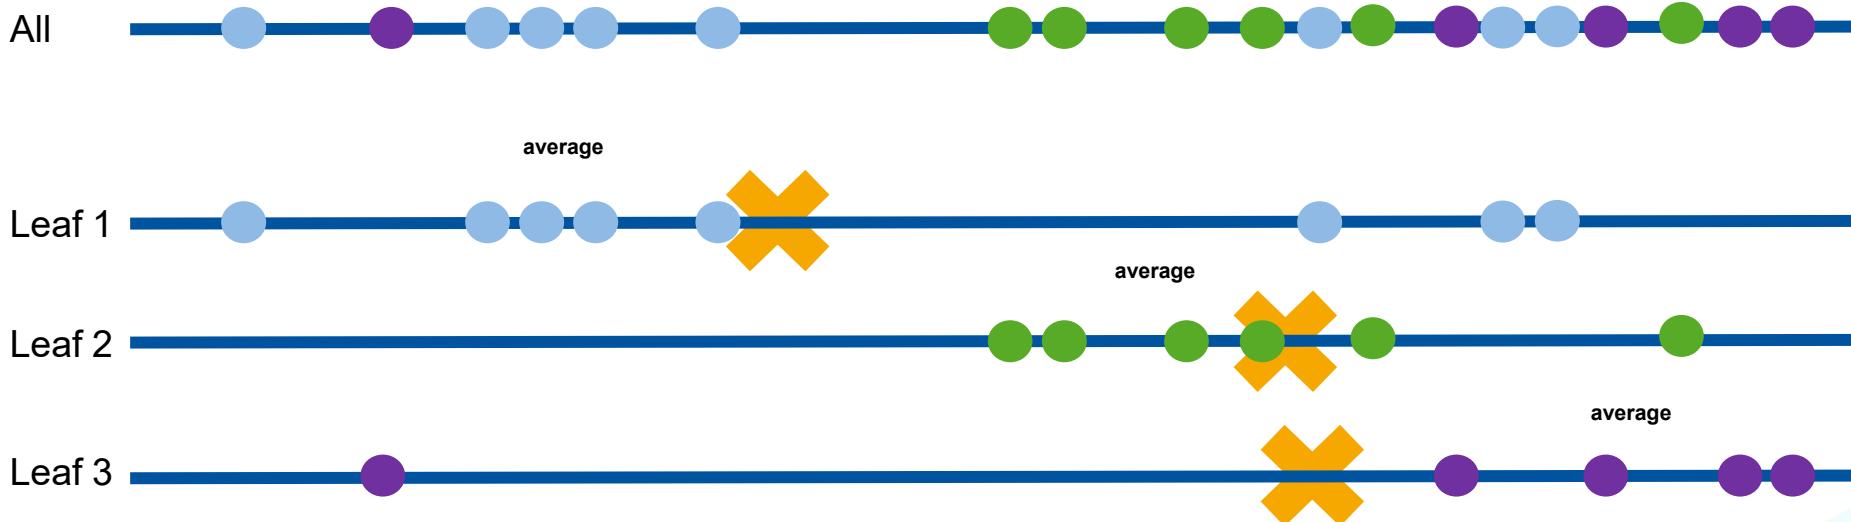


# Continuous Target Features



Poor Classification

Variance within all leaves is high compared to the 'good classification'

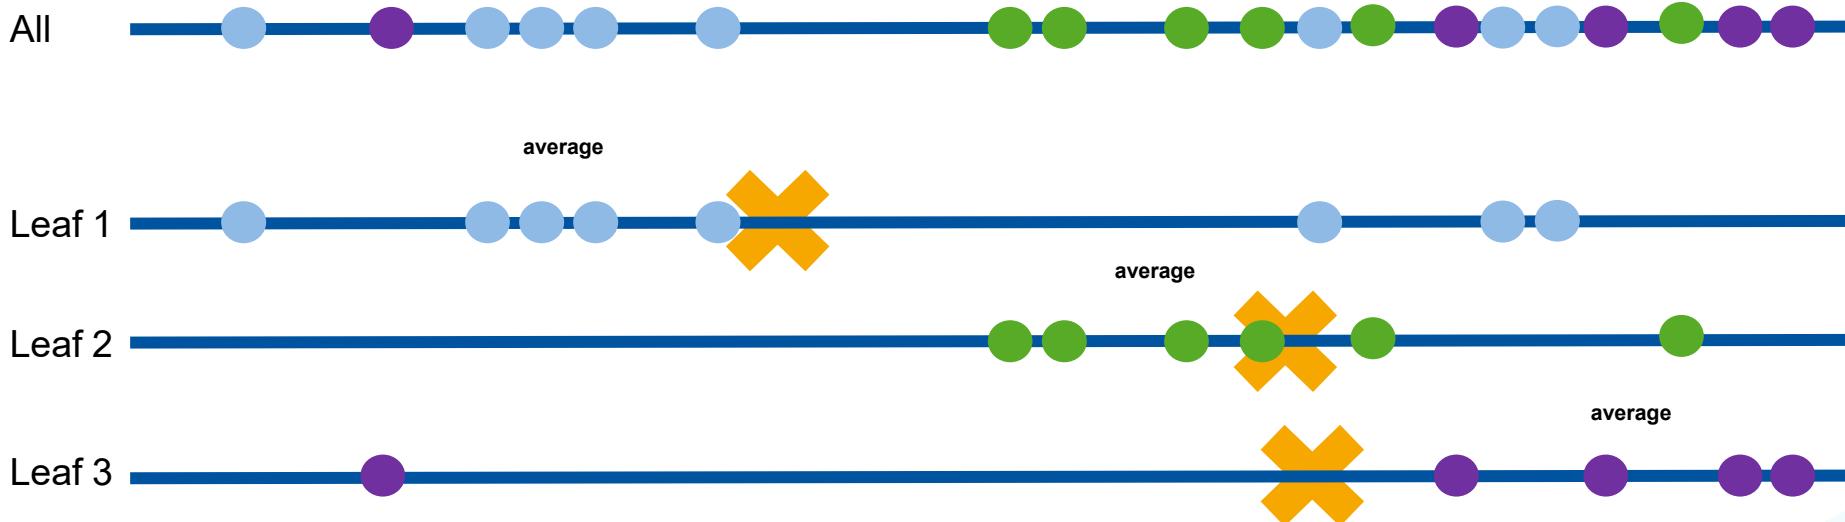


# Impurity

Variance in a Node/Leaf

Number of instances  
Target value of instance  $i$   
Mean of target values

$$Var(t) = \frac{\sum_{i=1}^N (t_i - \bar{t})^2}{N-1}$$



# Adapting the ID3 Algorithm

**ID3 algorithm:**

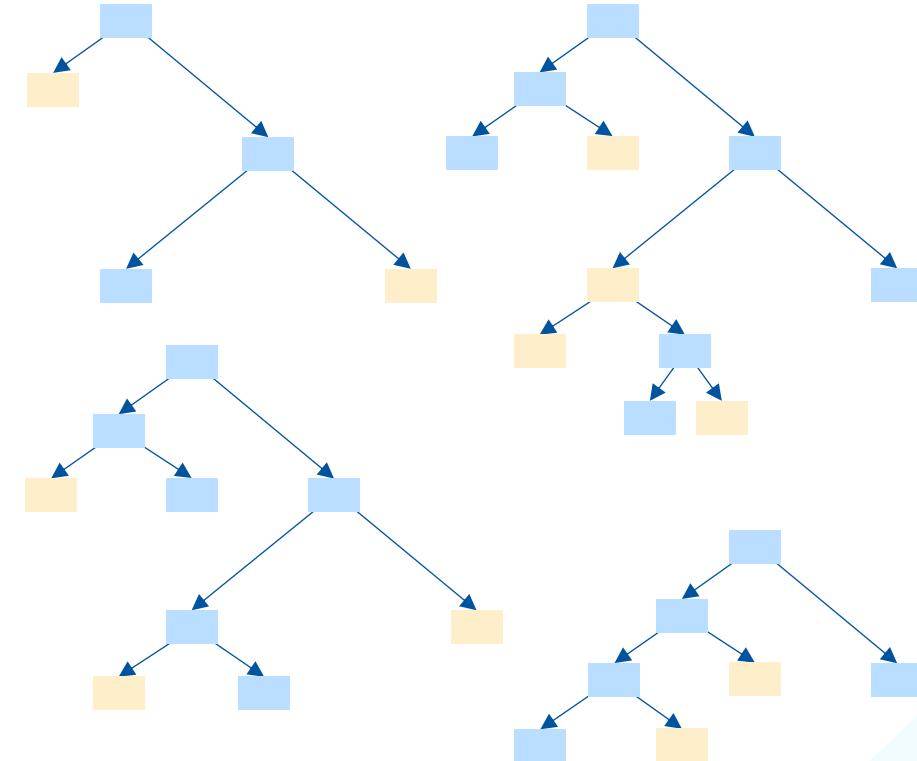
1. **if** all the instances in the dataset have the same classification
  - (a) **return** a decision tree with one leaf node with consensus value as a label
2. **else if** there are no features left
  - (a) **return** a decision tree with one leaf node with majority value as a label
3. **else if** the dataset is empty
  - (a) **return** a decision tree with one leaf node with majority parent value as a label
4. **else**
  - (a) pick a feature that lowers the weighted variance most within the subtrees
  - (b) once a feature is picked along a path from the root, it cannot be used again
  - (c) create subproblems based on the selected feature

Stopping criteria  
(as before)

Instead of  
maximizing  
information gain

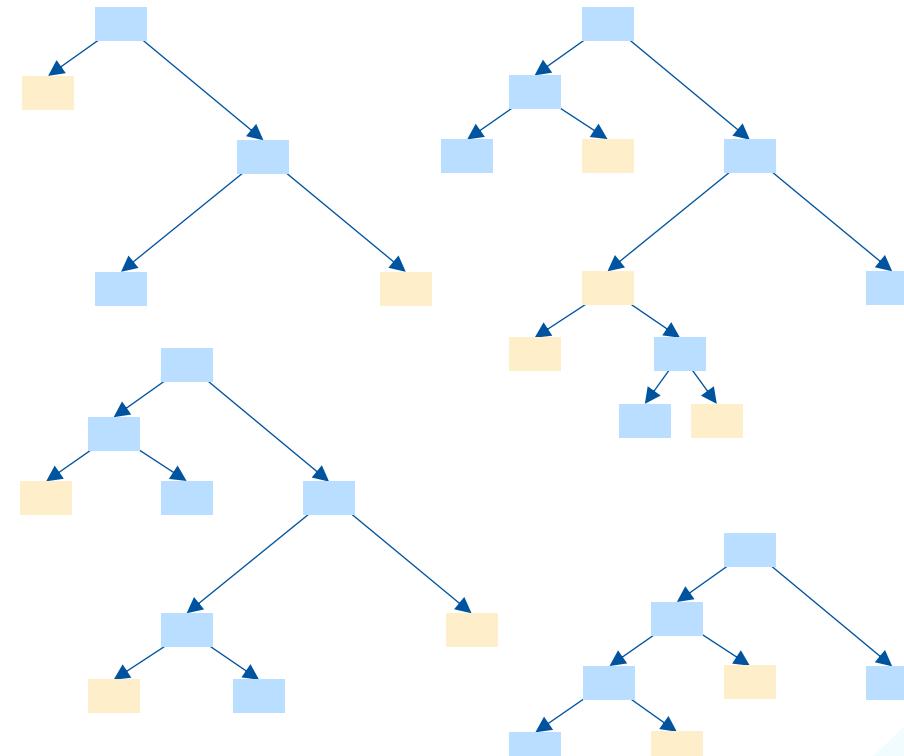
# Decision Trees

1. Introduction to Decision Trees
2. Entropy and Information Gain
3. ID3 Algorithm
4. Quantifying Information Gain
5. Pruning
6. Continuous Data
7. **Ensembles**



# Ensembles: Main Idea

- Rather than creating a single decision tree, we aim to create a **set of trees** (called a model **ensemble**)
- Models should complement each other
- Different models can "vote" on the label (votes may be weighted)
- Multiple trees may give different answers (select the most frequent value or the average)
- Many variations of the same idea...



# Ensembles: From one to multiple

- **Bagging** (sample instances): Trees focus on subsets of instances (i.e., rows)
- **Subspace sampling** (sample features): Trees focus on subsets of features (i.e., columns)
- **Random forest** (combine bagging and subspace sampling): Trees focus on subsets of instances and features
- **Boosting** (focus on errors): Create additional trees giving more weight to incorrectly classified instances.



## One glimpse into the toolbox...



- Not one specific ‘decision tree algorithm’
- Any variations are possible by combining ideas
- There is no best solution, it all depends on your data and goal

## Performance on unseen test data is what counts



- Avoid overfitting the data!
- Split data into training and test data
- Evaluation methods such as accuracy and confusion matrix will be discussed later

## Decision Trees - Conclusion

- Supervised learning aims to explain the target feature in terms of descriptive features
- Decision trees are easy to understand and interpret
- Focus on categorical variables but extensions to continuous data are possible
- Many variations based on the basic ID3 algorithm
  - Pruning
  - Ensembles
  - Information gain definitions
  - ...
- You have seen conceptual examples – implementations will differ in design choices, e.g., how to handle border cases, pick thresholds...