



Visual Computing
Institute

RWTH AACHEN
UNIVERSITY

Elements of Machine Learning & Data Science

Winter semester 2025/26

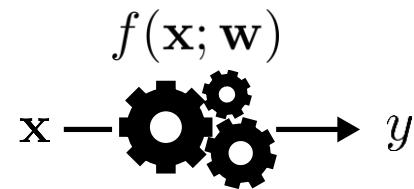
Lecture 9 – Bayes Decision Theory

24.11.2025

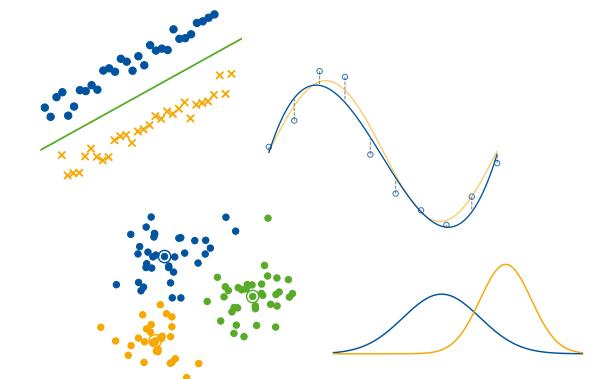
Prof. Bastian Leibe

Machine Learning Topics

- 8. **Introduction to ML**
- 9. Probability Density Estimation
- 10. Linear Discriminants
- 11. Linear Regression
- 12. Logistic Regression
- 13. Support Vector Machines
- 14. Neural Network Basics



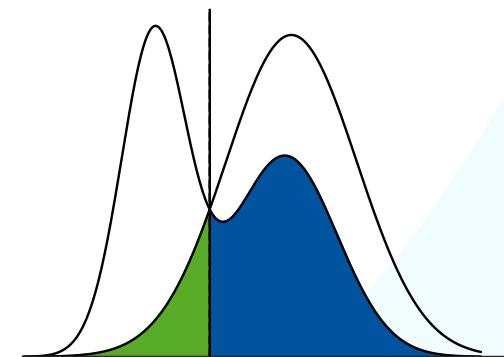
Machine Learning
Concepts



Forms of Machine Learning

$$p(\mathcal{C}|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C})p(\mathcal{C})}{p(\mathbf{x})}$$

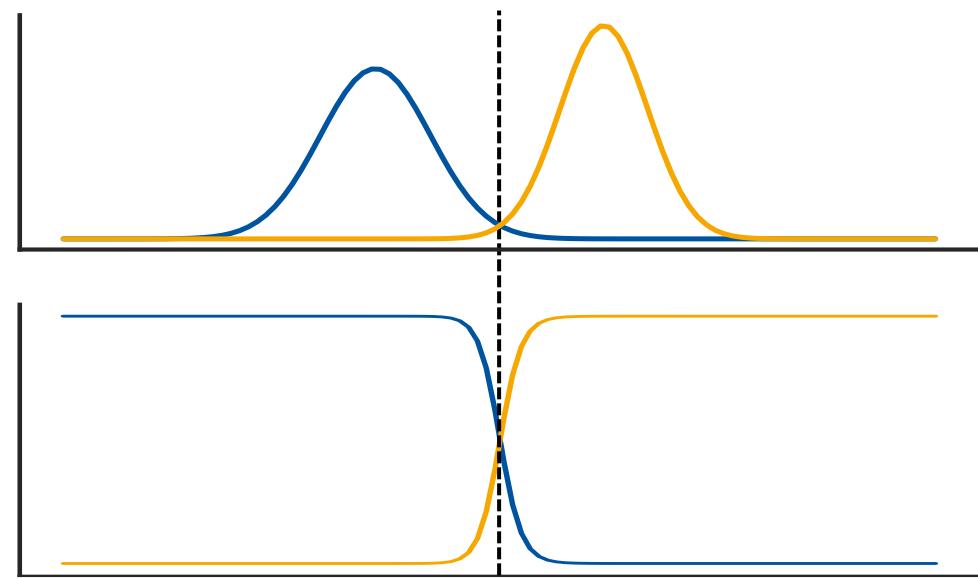
Bayes Decision Theory



Bayes Optimal
Classification

Introduction

1. Motivation
2. Forms of learning
3. Terms, Concepts, and Notation
4. **Bayes Decision Theory**

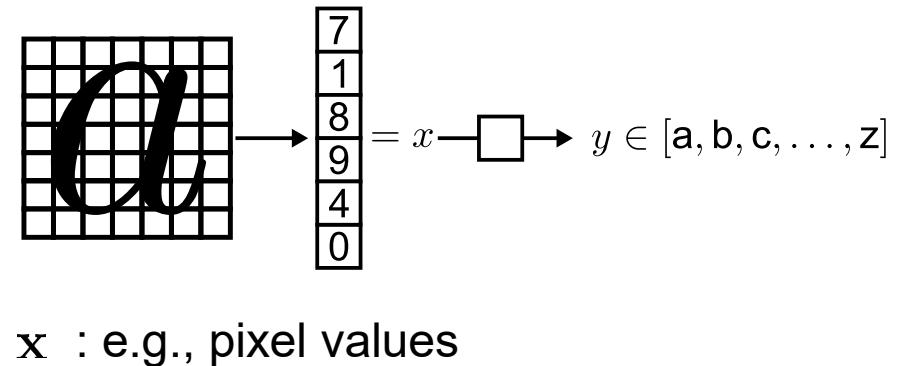


Bayes Decision Theory

- Goal: predict an output class \mathcal{C} from measurements x , by minimizing the probability of misclassification.
- *How can we make such decisions optimally?*
- Bayes Decision Theory gives us the tools for this
 - Based on Bayes' Theorem:

$$p(Y|X) = \frac{p(X|Y)p(Y)}{p(X)}$$

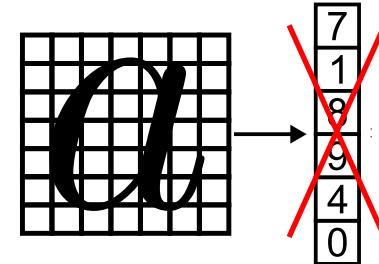
Example: handwritten character recognition



- In the following, we will introduce its basic concepts...

Core Concept: Priors

- What can we tell about the outcome of an experiment *before* making any measurements?
- The **a-priori probability** $p(\mathcal{C})$ captures the probability distribution over the different class outcomes
 - Based on previously observed data
 - i.e., independent of the actual measurement
- The prior probabilities over all possible class outcomes sum to one.



Example: in English text, the letter “e” makes up ~13% of all letters:

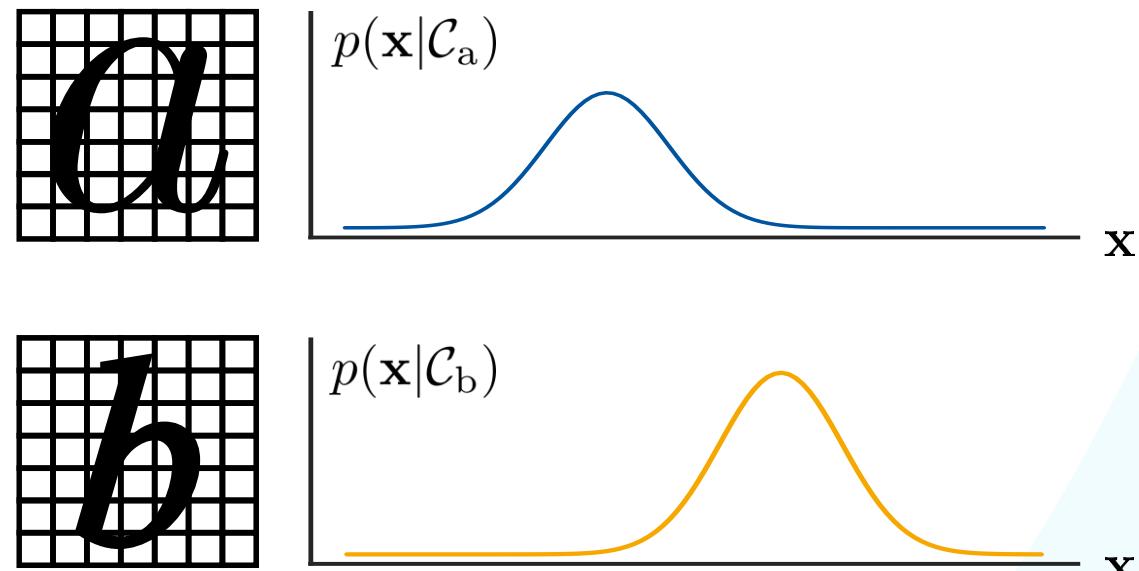
$$p(\mathcal{C}_e) = 0.13$$

And there are 26 letters in the English alphabet:

$$\sum_{\alpha \in \{a, \dots, z\}} p(\mathcal{C}_\alpha) = 1$$

Core Concept: Likelihood

- How *likely* is it that we *observe* a certain measurement \mathbf{x} *given* an example of class \mathcal{C} ?
- This is expressed by the **likelihood** $p(\mathbf{x}|\mathcal{C})$
 - It is called a *class-conditional distribution*, since it specifies the distribution of \mathbf{x} conditioned on the class \mathcal{C} .
 - We can estimate the likelihood from the distribution of measurements \mathbf{x} observed on the given training data.
- Here, \mathbf{x} measures certain properties of the input data.
 - E.g., the fraction of black pixels
 - We simply treat it as a vector $\mathbf{x} \in \mathbb{R}^D$.



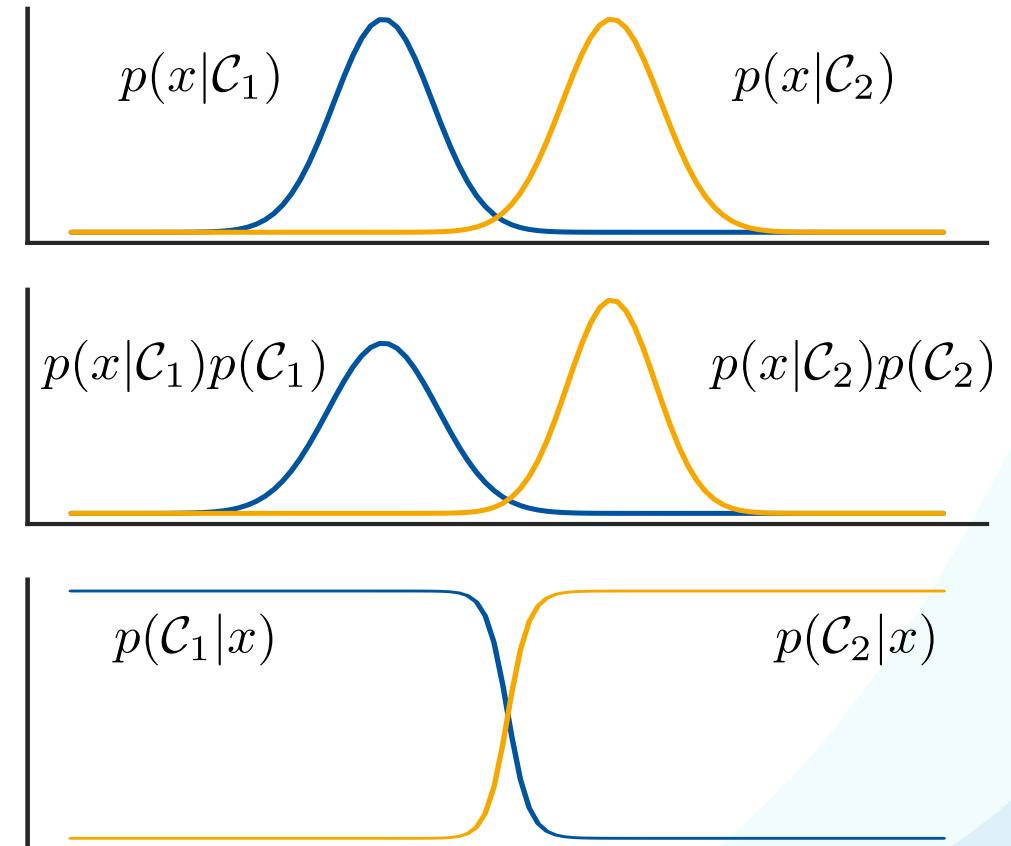
Core Concept: Posterior

- What is the probability for class \mathcal{C}_k if we made a measurement \mathbf{x} ?
- This **a-posteriori probability** $p(\mathcal{C}_k|\mathbf{x})$ can be computed via Bayes' Theorem after we observed \mathbf{x} :

$$p(\mathcal{C}_k|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_k)p(\mathcal{C}_k)}{p(\mathbf{x})} = \frac{p(\mathbf{x}|\mathcal{C}_k)p(\mathcal{C}_k)}{\sum_j p(\mathbf{x}|\mathcal{C}_j)p(\mathcal{C}_j)}$$

- *This is usually what we're interested in!*
- Interpretation

$$\text{posterior} = \frac{\text{likelihood} \cdot \text{prior}}{\text{normalization factor}}$$



Making Optimal Decisions

- Goal: minimize the probability of misclassification.

$$p(\text{mistake}) = p(x \in \mathcal{R}_1, \mathcal{C}_2) + p(x \in \mathcal{R}_2, \mathcal{C}_1)$$

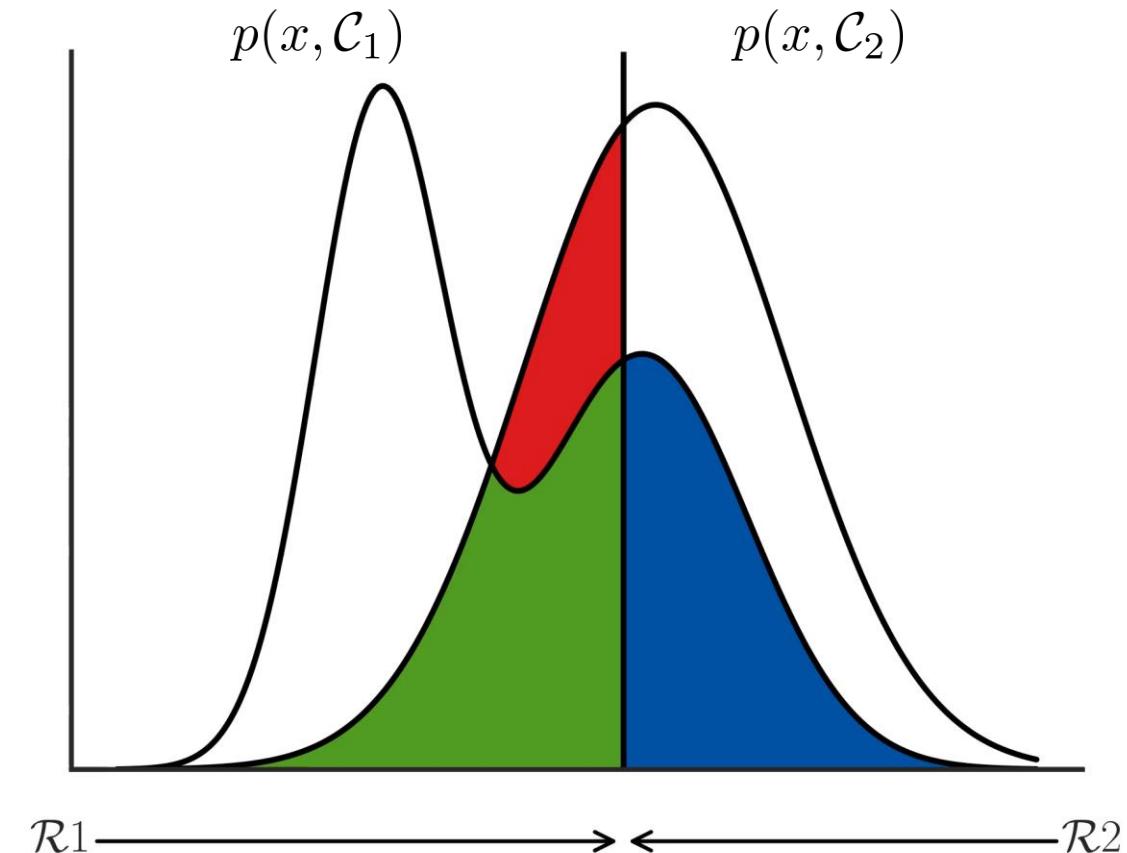
$$= \int_{\mathcal{R}_1} p(x, \mathcal{C}_2) dx + \int_{\mathcal{R}_2} p(x, \mathcal{C}_1) dx$$

$$= \int_{\mathcal{R}_1} p(\mathcal{C}_2|x)p(x) dx + \int_{\mathcal{R}_2} p(\mathcal{C}_1|x)p(x) dx$$

- Note:

+ = constant

We can only reduce



\mathcal{R}_1 and \mathcal{R}_2 are the **decision regions** after setting a decision threshold.

Making Optimal Decisions

- Goal: minimize the probability of misclassification.

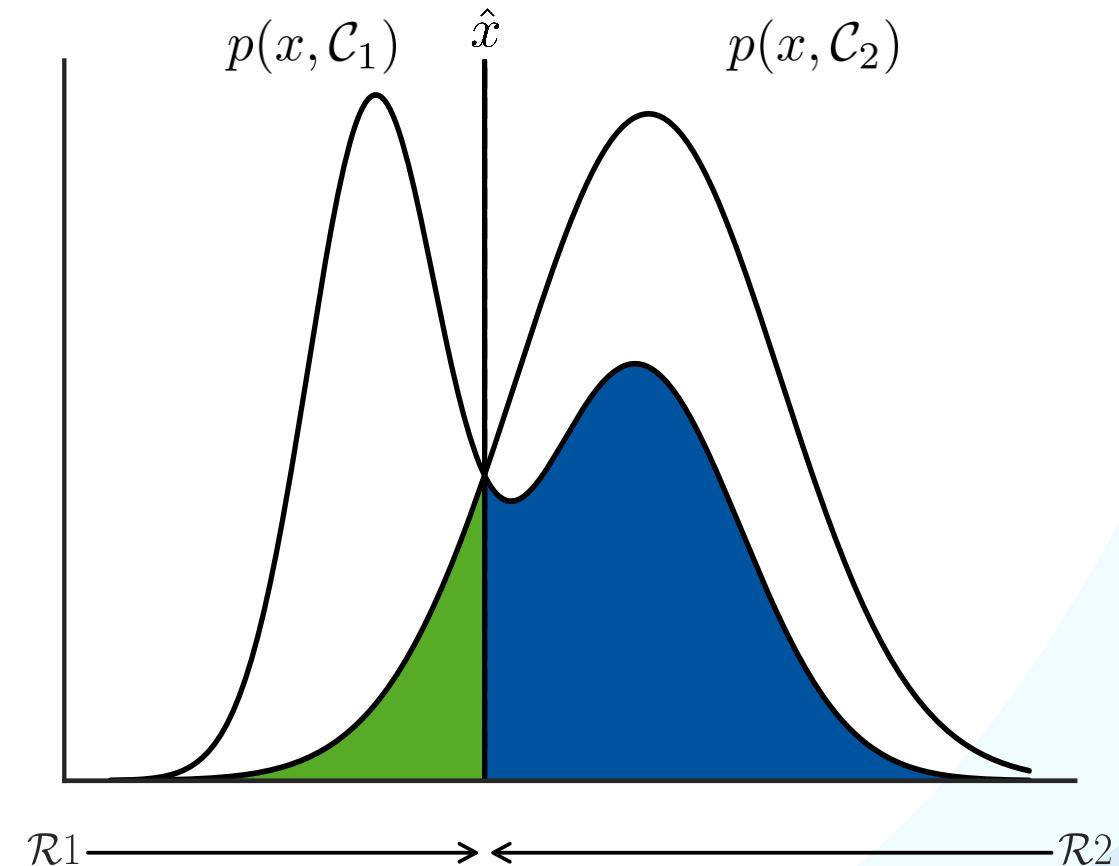
$$\begin{aligned}
 p(\text{mistake}) &= p(x \in \mathcal{R}_1, \mathcal{C}_2) + p(x \in \mathcal{R}_2, \mathcal{C}_1) \\
 &= \int_{\mathcal{R}_1} p(x, \mathcal{C}_2) dx + \int_{\mathcal{R}_2} p(x, \mathcal{C}_1) dx \\
 &= \int_{\mathcal{R}_1} p(\mathcal{C}_2|x)p(x) dx + \int_{\mathcal{R}_2} p(\mathcal{C}_1|x)p(x) dx
 \end{aligned}$$

- Note:

$$\text{blue} + \text{green} = \text{red}$$

We can only reduce

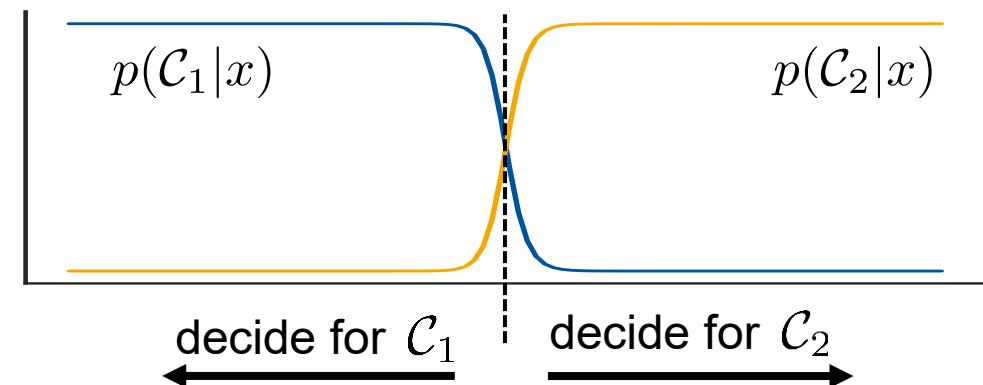
- Minimal error at the intersection \hat{x}*



\mathcal{R}_1 and \mathcal{R}_2 are the **decision regions** after setting a decision threshold.

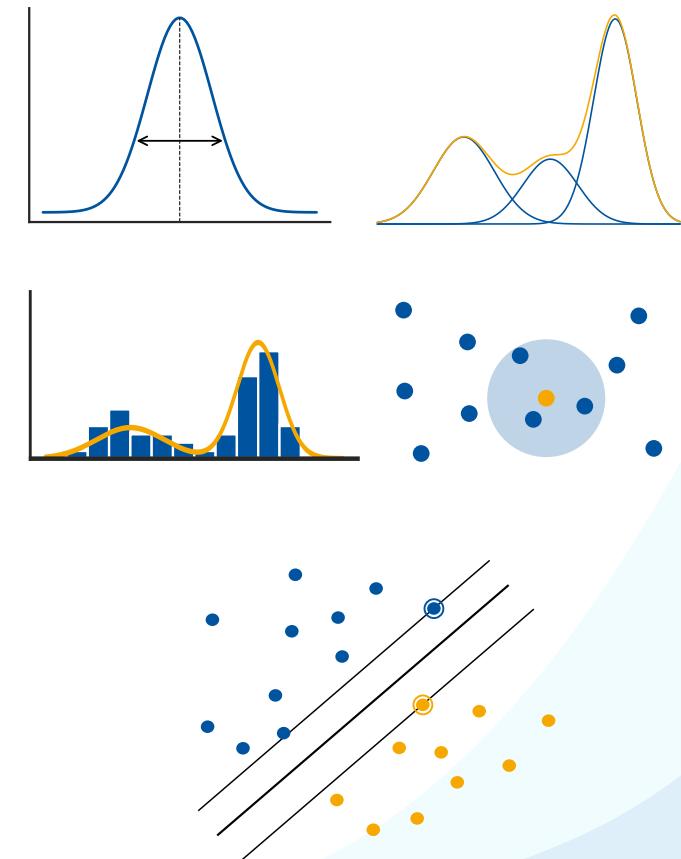
Making Optimal Decisions

- Our goal is to minimize the probability of a misclassification.
- The optimal decision rule is: decide for \mathcal{C}_1 iff $p(\mathcal{C}_1|\mathbf{x}) > p(\mathcal{C}_2|\mathbf{x})$
- Or for multiple classes: decide for \mathcal{C}_k iff $p(\mathcal{C}_k|\mathbf{x}) > p(\mathcal{C}_j|\mathbf{x}) \forall j \neq k$
- *Once we can estimate posterior probabilities, we can use this rule to build classifiers.*



Remainder of Today and Next Lectures...

- Ways how to estimate the probability densities $p(\mathbf{x}|\mathcal{C}_k)$
 - **Parametric methods**
 - Gaussian distribution
 - Mixtures of Gaussians
 - **Non-parametric methods**
 - Histograms
 - k-Nearest Neighbor
 - Kernel Density Estimation
- Ways to directly model the posteriors $p(\mathcal{C}_k|\mathbf{x})$
 - Linear discriminants
 - Logistic regression, SVMs, Neural Networks, ...



Machine Learning Topics

8. Introduction to ML

9. Probability Density Estimation

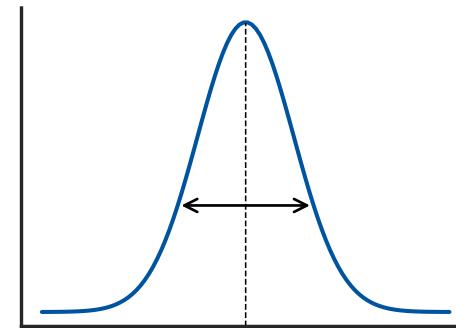
10. Linear Discriminants

11. Linear Regression

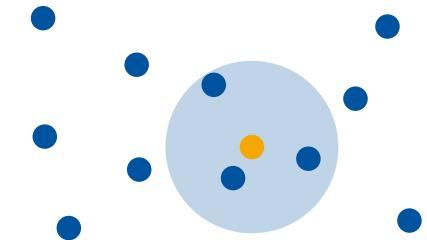
12. Logistic Regression

13. Support Vector Machines

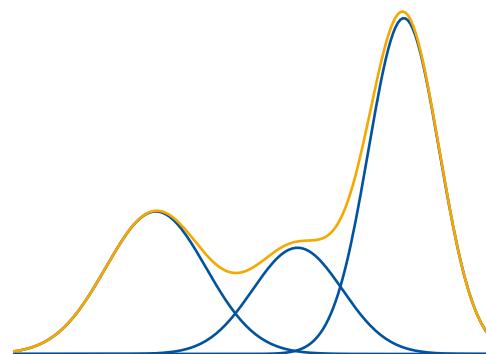
14. Neural Network Basics



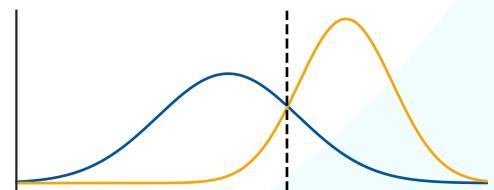
Parametric Methods
& ML-Algorithm



Nonparametric Methods



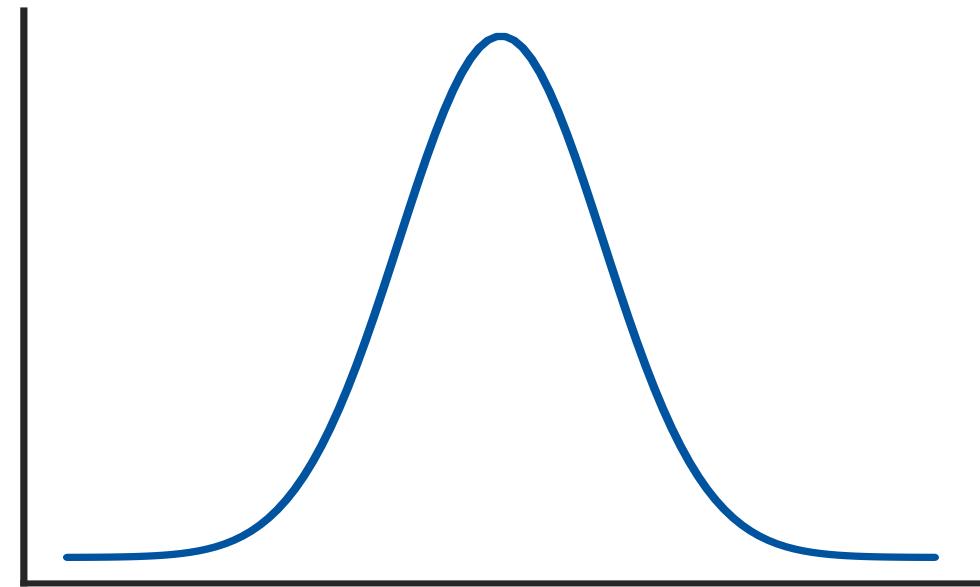
Mixtures of Gaussians
& EM-Algorithm



Bayes Classifiers

Probability Density Estimation

1. Probability Distributions
2. Parametric Methods
3. Nonparametric Methods
4. Mixture Models
5. Bayes Classifier
6. K-NN Classifier



Probability Distributions

- Up to now: Bayes optimal classification based on $p(\mathbf{x}|\mathcal{C}_k)$ and $p(\mathcal{C}_k)$.
- *How can we estimate (= learn) those probability densities?*
 - Supervised training case: data and class labels are known.
 - Estimate the probability density for each class \mathcal{C}_k separately.
$$p(\mathbf{x}|\mathcal{C}_k) \quad \text{given } \mathbf{x} \in \mathcal{C}_k$$
 - (For simplicity of notation, we will drop the class label \mathcal{C}_k in the following $\Rightarrow p(\mathbf{x})$).
- *First, we look at the Gaussian distribution in more detail...*

The Gaussian (or Normal) Distribution

- One-dimensional (univariate) case:

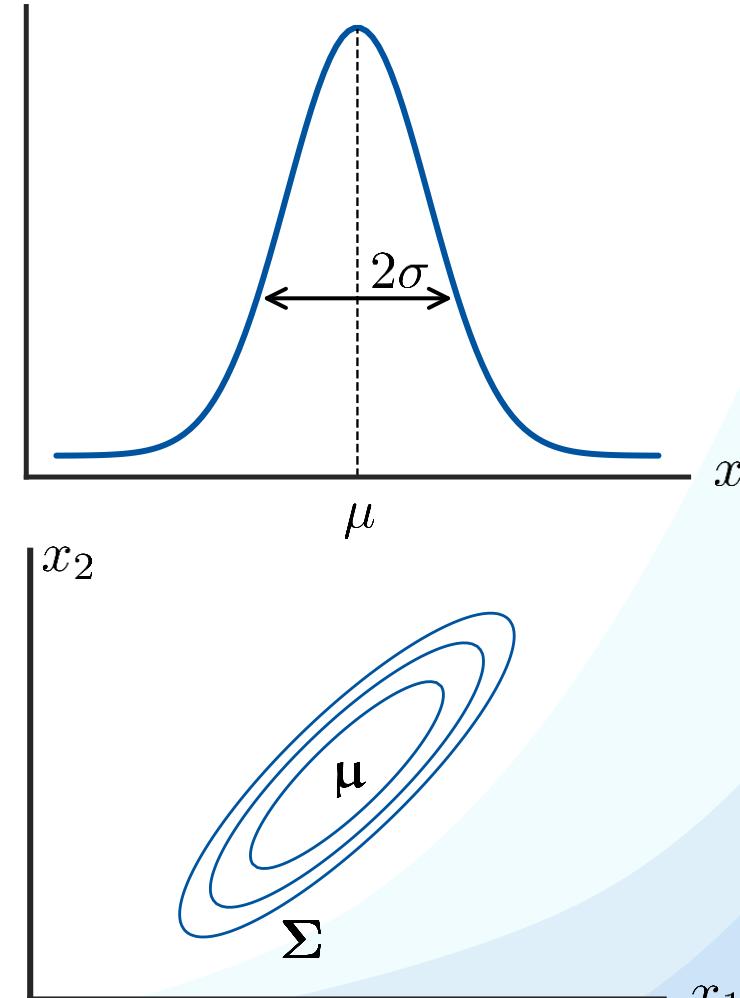
$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

Mean Variance

- Multi-dimensional (multivariate) case:

$$\mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{D/2}|\boldsymbol{\Sigma}|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\mathbf{x}-\boldsymbol{\mu})\right)$$

Mean vector Covariance matrix

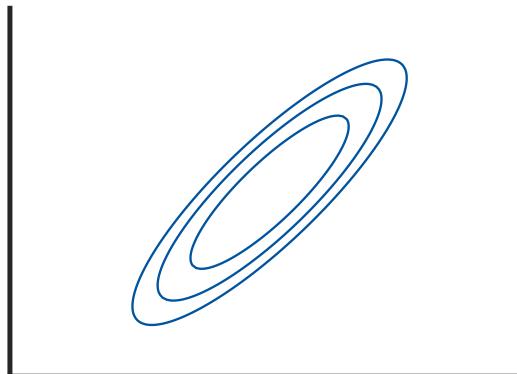


Gaussian Distribution: Shape

Full covariance matrix:

$$\Sigma = [\sigma_{ij}]$$

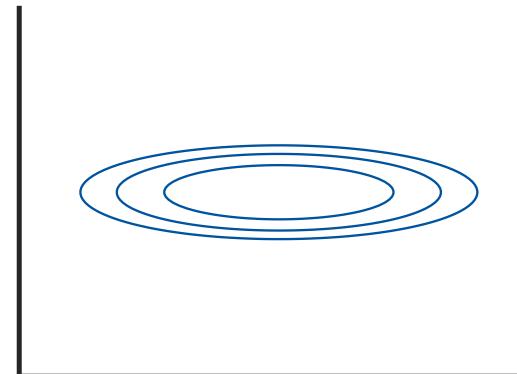
General ellipsoid shape



Diagonal covariance matrix:

$$\Sigma = \text{diag}\{\sigma_i\}$$

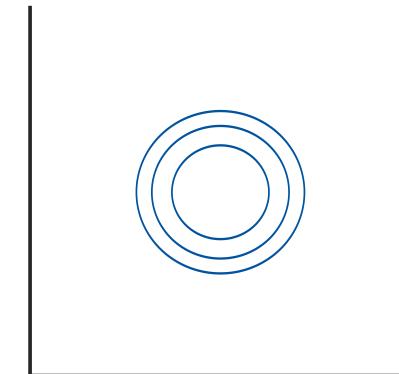
Axis-aligned ellipsoid



Uniform variance:

$$\Sigma = \sigma^2 \mathbf{I}$$

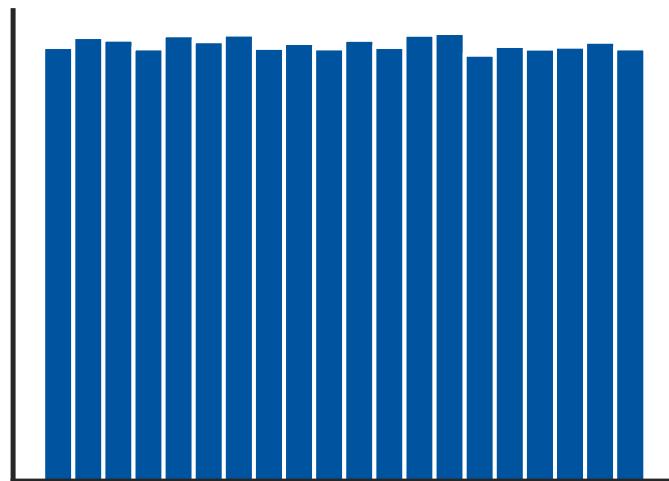
Hypersphere



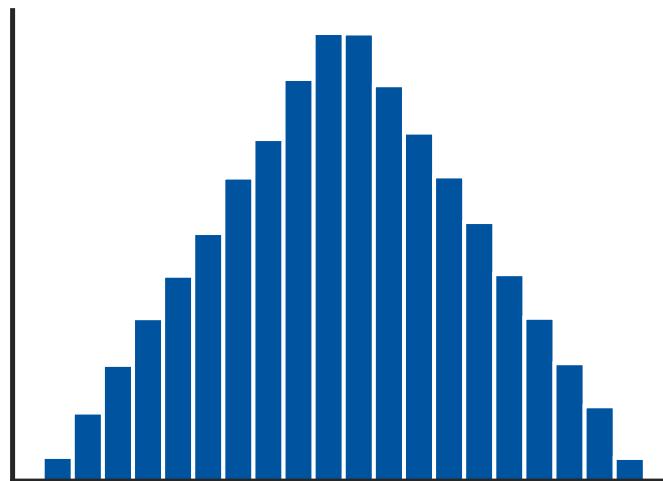
Gaussian Distribution: Motivation

- Central Limit Theorem
 - The distribution of a sum of N *i.i.d.* random variables becomes increasingly Gaussian as N grows.
 - In practice, the convergence to a Gaussian can be very rapid.
 - This makes the Gaussian interesting for many applications.
- Example: Sum over N uniform $[0,1]$ random variables.

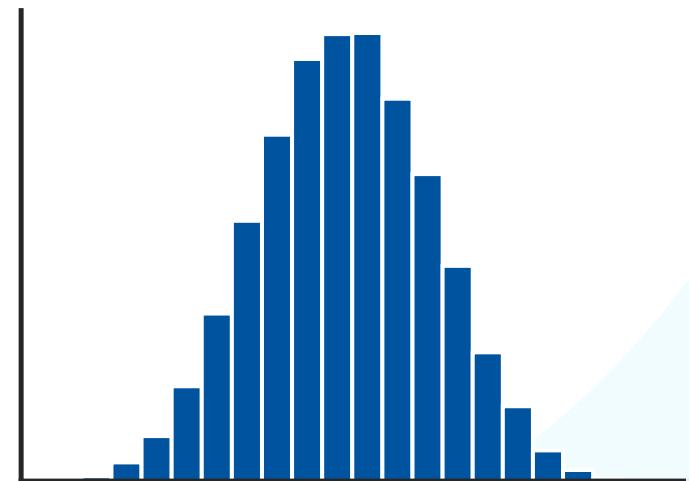
i.i.d. = independent and identically distributed



$N = 1$



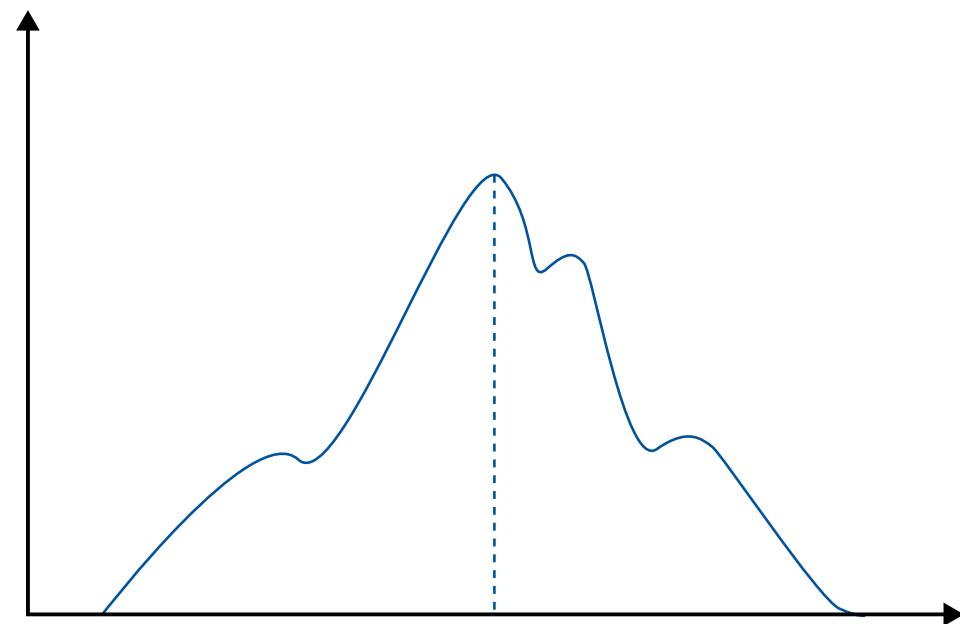
$N = 2$



$N = 5$

Probability Density Estimation

1. Probability Distributions
2. **Parametric Methods**
3. Nonparametric Methods
4. Mixture Models
5. Bayes Classifier
6. K-NN Classifier

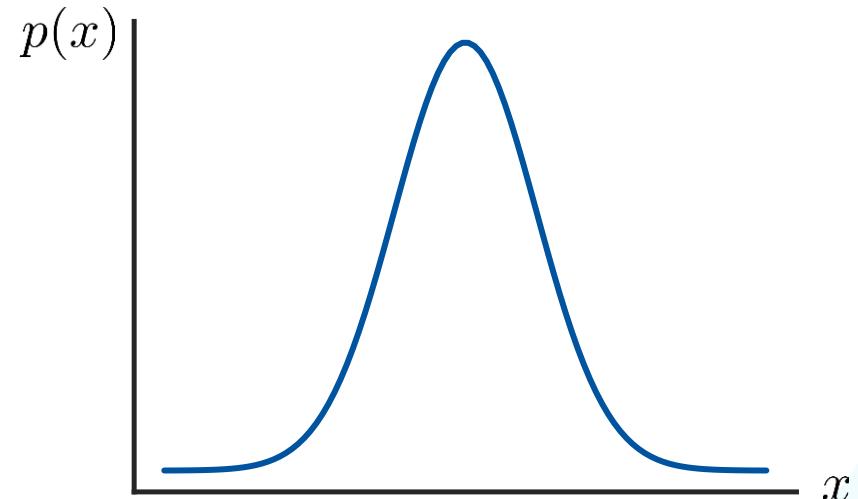


Parametric Methods

- In **parametric methods**, we assume that we know the parametric form of the underlying data distribution.
 - I.e., the equation of the pdf with parameters θ .

Example: $p(x) = \mathcal{N}(x|\mu, \sigma)$

$$\theta = (\mu, \sigma)$$



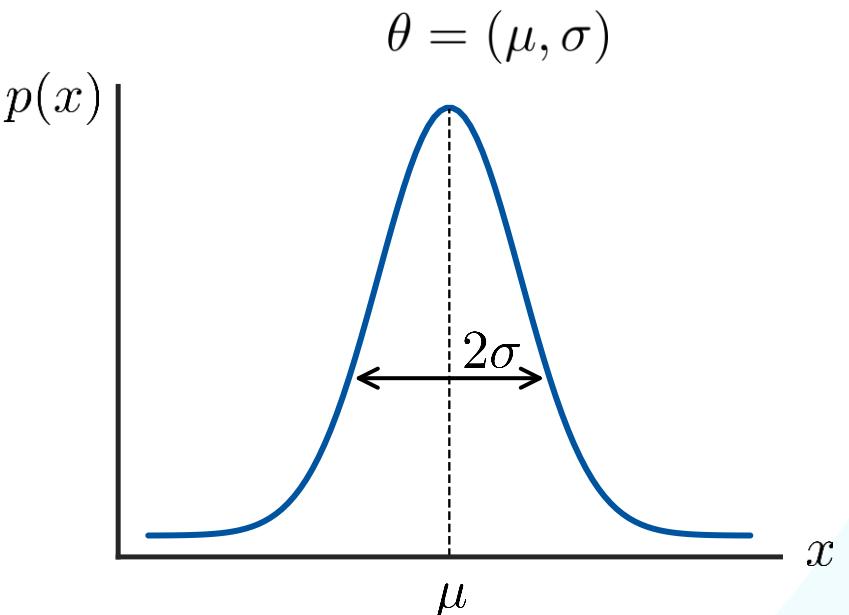
Parametric Methods

- In **parametric methods**, we assume that we know the parametric form of the underlying data distribution.
 - I.e., the equation of the pdf with parameters θ .
- Goal: Estimate θ from training data $\mathcal{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$.
- **Likelihood** of θ :

$$L(\theta) = p(\mathcal{X}|\theta)$$

Probability that the data \mathcal{X} was indeed generated by a distribution with parameters θ .

Example: $p(x) = \mathcal{N}(x|\mu, \sigma)$



Maximum Likelihood Approach

- Idea: Find optimal parameters by maximizing $L(\theta)$.
- Computation of the likelihood:

- Single data point (e.g., for Gaussian):

$$p(x_n|\theta) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x_n - \mu)^2}{2\sigma^2}\right)$$

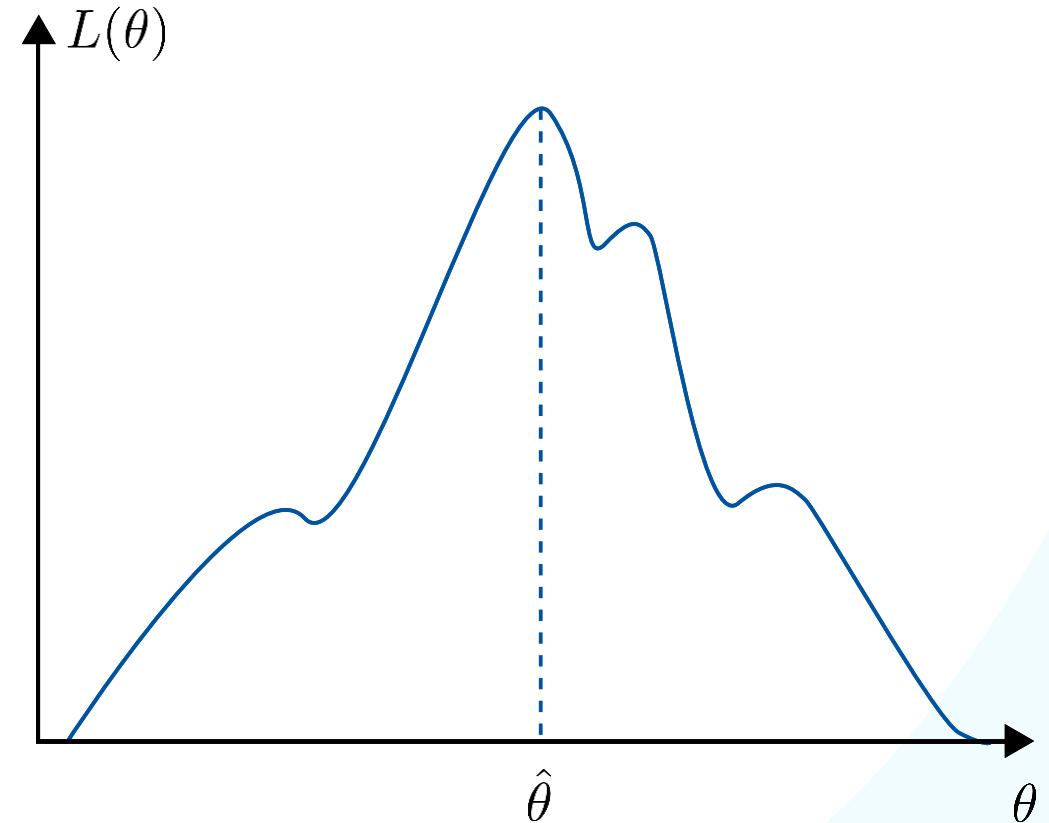
- Assumption: all data points are independent

$$L(\theta) = p(\mathcal{X}|\theta) = \prod_{n=1}^N p(x_n|\theta)$$

- Negative Log-Likelihood (“Energy”):

$$E(\theta) = -\ln L(\theta) = -\sum_{n=1}^N \ln p(x_n|\theta)$$

Maximizing the likelihood \Leftrightarrow minimizing the negative log-likelihood.



- Minimizing the negative log-likelihood:
 - Take the derivative and set it to zero.

$$\frac{\partial}{\partial \theta} E(\theta) = -\frac{\partial}{\partial \theta} \sum_{n=1}^N \ln p(x_n|\theta) = -\sum_{n=1}^N \frac{\frac{\partial}{\partial \theta} p(x_n|\theta)}{p(x_n|\theta)} \stackrel{!}{=} 0$$

- Log-likelihood for Normal distribution (1D case):

$$\begin{aligned}\frac{\partial}{\partial \mu} E(\mu, \sigma) &= -\sum_{n=1}^N \frac{\frac{\partial}{\partial \mu} p(x_n|\mu, \sigma)}{p(x_n|\mu, \sigma)} \\ &= -\sum_{n=1}^N -\frac{2(x_n - \mu)}{2\sigma^2} \frac{p(x_n|\mu, \sigma)}{p(x_n|\mu, \sigma)}\end{aligned}$$

$$p(x_n|\mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x_n - \mu)^2}{2\sigma^2}\right)$$

$$\frac{\partial}{\partial \mu} p(x_n|\mu, \sigma) = -\frac{2(x_n - \mu)}{2\sigma^2} p(x_n|\mu, \sigma)$$

$$=\frac{1}{\sigma^2} \sum_{n=1}^N (x_n - \mu) = \cancel{\frac{1}{\sigma^2}} \left(\sum_{n=1}^N x_n - N\mu \right) \stackrel{!}{=} 0 \iff \hat{\mu} = \frac{1}{N} \sum_{n=1}^N x_n$$

- By minimizing the negative log-likelihood, we found: Similarly, we can derive:

$$\hat{\mu} = \frac{1}{N} \sum_{n=1}^N x_n$$

sample mean

$$\hat{\sigma}^2 = \frac{1}{N} \sum_{n=1}^N (x_n - \hat{\mu})^2$$

sample variance

- $\hat{\theta} = (\hat{\mu}, \hat{\sigma})$ is the Maximum Likelihood estimate for the parameters of a Gaussian distribution.
 - This is a very important result.
 - Unfortunately, it is wrong...

- To be precise, the result is not wrong, but **biased**.
- Assume the samples x_1, x_2, \dots, x_N come from a true Gaussian distribution with mean μ and variance σ^2
 - It can be shown that the expected estimates are then

$$\mathbb{E}[\mu_{\text{ML}}] = \mu$$

$$\mathbb{E}[\sigma_{\text{ML}}^2] = \left(\frac{N-1}{N}\right)\sigma^2$$

\Rightarrow *The ML estimate will underestimate the true variance!*

- We can correct for this bias:

$$\hat{\sigma}^2 = \frac{N}{N-1}\sigma_{\text{ML}}^2 = \frac{1}{N-1} \sum_{n=1}^N (x_n - \hat{\mu})^2$$

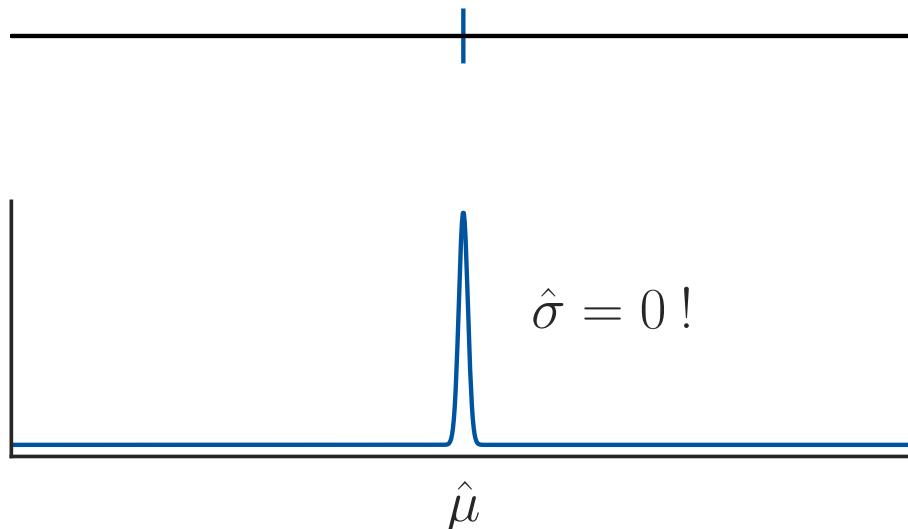
corrected, unbiased estimate

- Maximum Likelihood has several significant limitations.
 - It systematically underestimates the variance of the distribution!
 - E.g., consider the estimate for a single sample:

$$N = 1, \mathcal{X} = \{x_1\}$$

$$\hat{\mu} = \frac{1}{N} \sum_{n=1}^N x_n = x_1$$

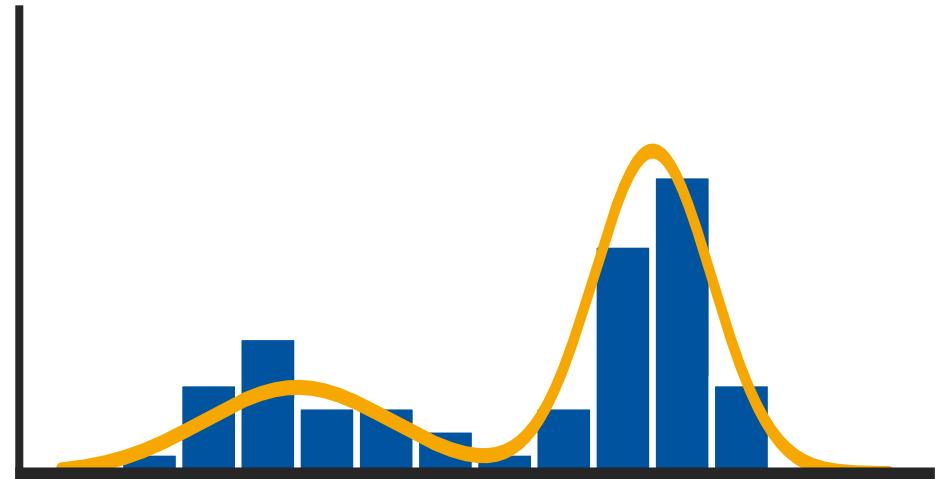
$$\hat{\sigma}^2 = \frac{1}{N} \sum_{n=1}^N (x_n - \hat{\mu})^2 = 0$$



- We say ML overfits to the observed data.
- *We will still often use Maximum Likelihood, but it is important to know about this effect.*

Probability Density Estimation

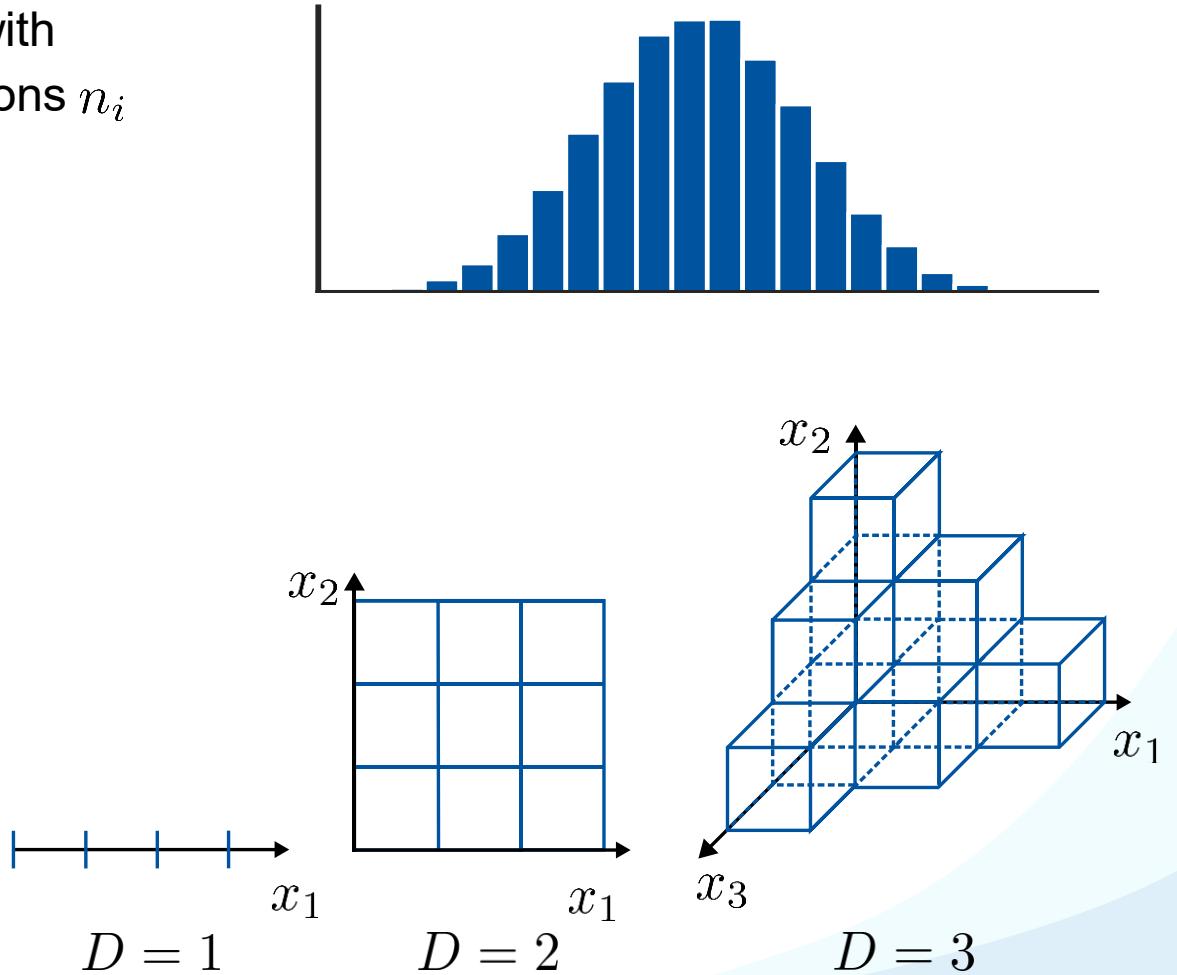
1. Probability Distributions
2. Parametric Methods
3. **Nonparametric Methods**
 - a) **Histograms**
 - b) Kernel Methods & k-Nearest Neighbors
4. Mixture Models
5. Bayes Classifier
6. K-NN Classifier



Histograms

- Partition the data space into N distinct bins with widths Δ_i and count the number of observations n_i in each bin.
- Then, $p_i = \frac{n_i}{N\Delta_i}$.
- Often the same width is used for all bins.
- This can be done, in principle, for any dimensionality D .

...but the required number of bins grows exponentially with D !



The bin width Δ acts as a smoothing factor.

Not smooth enough



About ok



Too smooth



Advantages

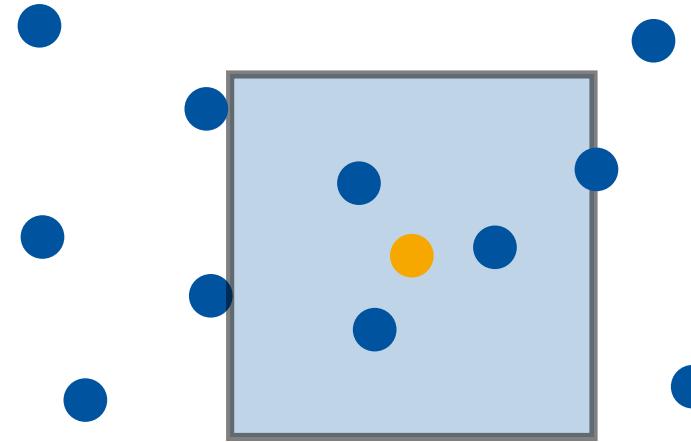
- Very general method. In the limit ($N \rightarrow \infty$), every probability density can be represented.
- No need to store the data points once histogram is computed.

Limitations

- Rather brute-force.
- Discontinuities at bin edges.
- Choosing right bin size is hard.
- Unsuitable for high-dimensional feature spaces.

Probability Density Estimation

1. Probability Distributions
2. Parametric Methods
3. **Nonparametric Methods**
 - a) Histograms
 - b) **Kernel Methods & k-Nearest Neighbors**
4. Mixture Models
5. Bayes Classifier
6. K-NN Classifier



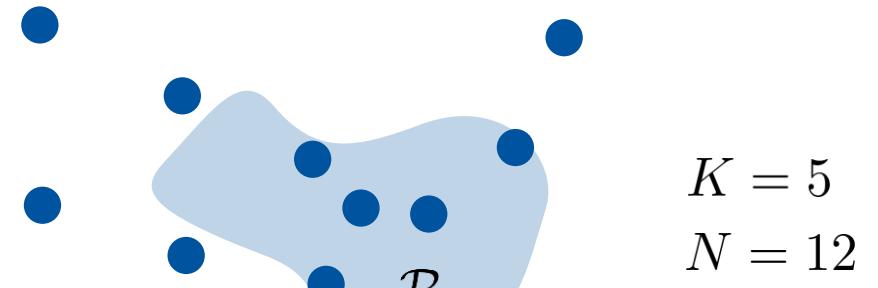
Kernel Methods and k-Nearest Neighbors

- Data point \mathbf{x} comes from pdf $p(\mathbf{x})$.
 - Probability that \mathbf{x} falls into small region \mathcal{R} :

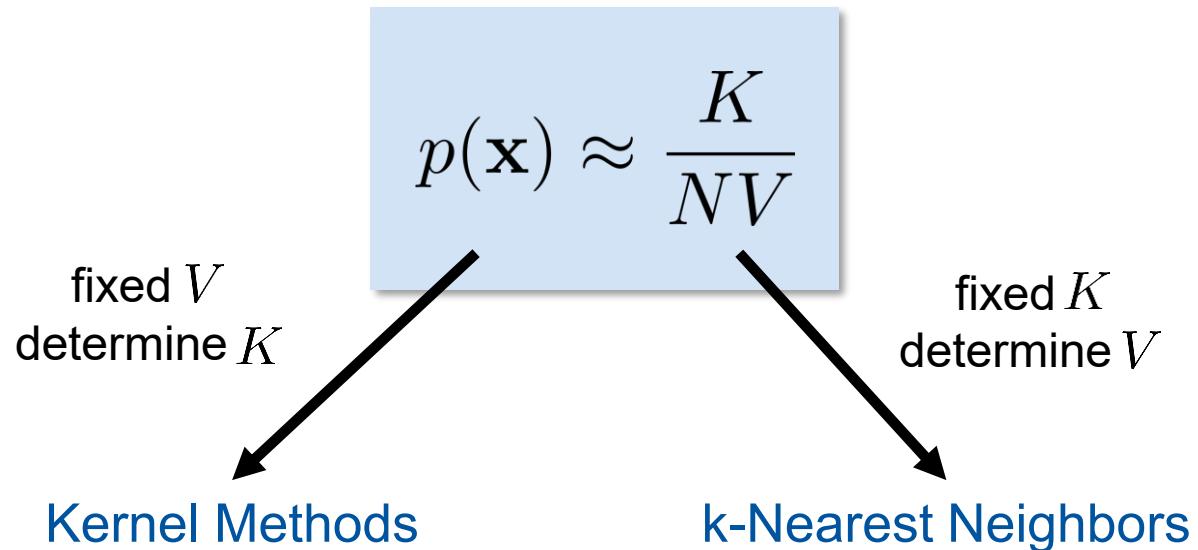
$$P = \int_{\mathcal{R}} p(y) dy \approx p(\mathbf{x})V$$

- Estimate $p(\mathbf{x})$ from samples
 - Let K be the number of samples that fall into \mathcal{R} .
 - If the number of samples N is sufficiently large, we can estimate P as:

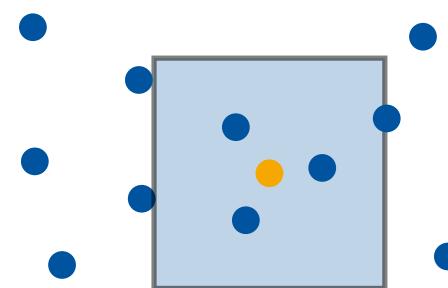
$$P = \frac{K}{N} \Rightarrow p(\mathbf{x}) \approx \frac{K}{NV}$$



For sufficiently small \mathcal{R} ,
 $p(\mathbf{x})$ is roughly constant.
 V : volume of \mathcal{R} .



Example: Determine
the number K of data
points inside a fixed
hypercube



Kernel Methods

- Hypercube of dimension D with edge length h :

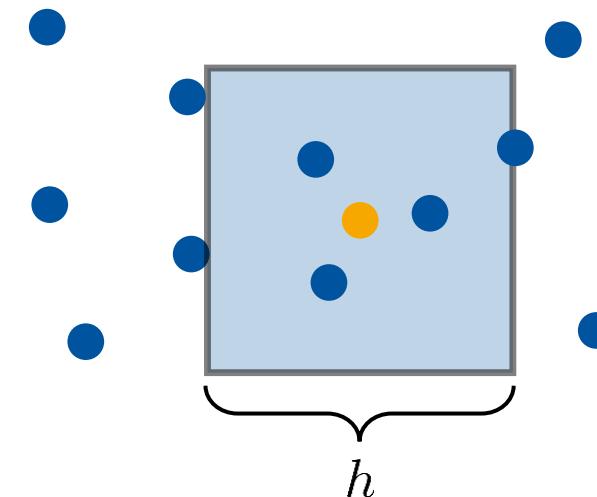
$$k(\mathbf{u}) = \begin{cases} 1, & \text{if } |u_i| \leq \frac{1}{2}h, \ i = 1, \dots, D \\ 0, & \text{otherwise} \end{cases}$$

$$K = \sum_{n=1}^N k(\mathbf{x} - \mathbf{x}_n) \quad V = \int k(\mathbf{u}) d\mathbf{u} = h^D$$

- Probability density estimate:

$$p(\mathbf{x}) \approx \frac{K}{NV} = \frac{1}{Nh^D} \sum_{n=1}^N k(\mathbf{x} - \mathbf{x}_n)$$

- This method is known as **Parzen Window** estimation.



- In general, we can use any kernel such that

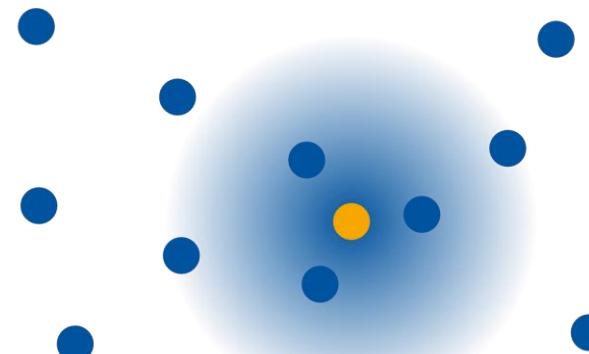
$$k(\mathbf{u}) \geq 0, \quad \int k(\mathbf{u}) d\mathbf{u} = 1$$

$$K = \sum_{n=1}^N k(\mathbf{x} - \mathbf{x}_n)$$

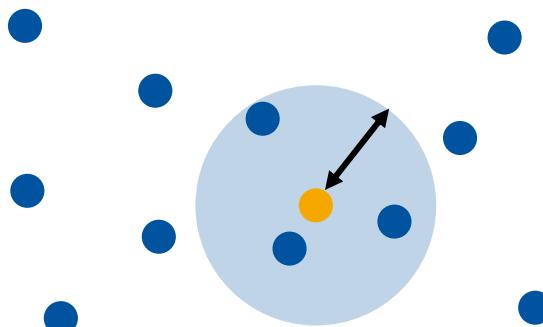
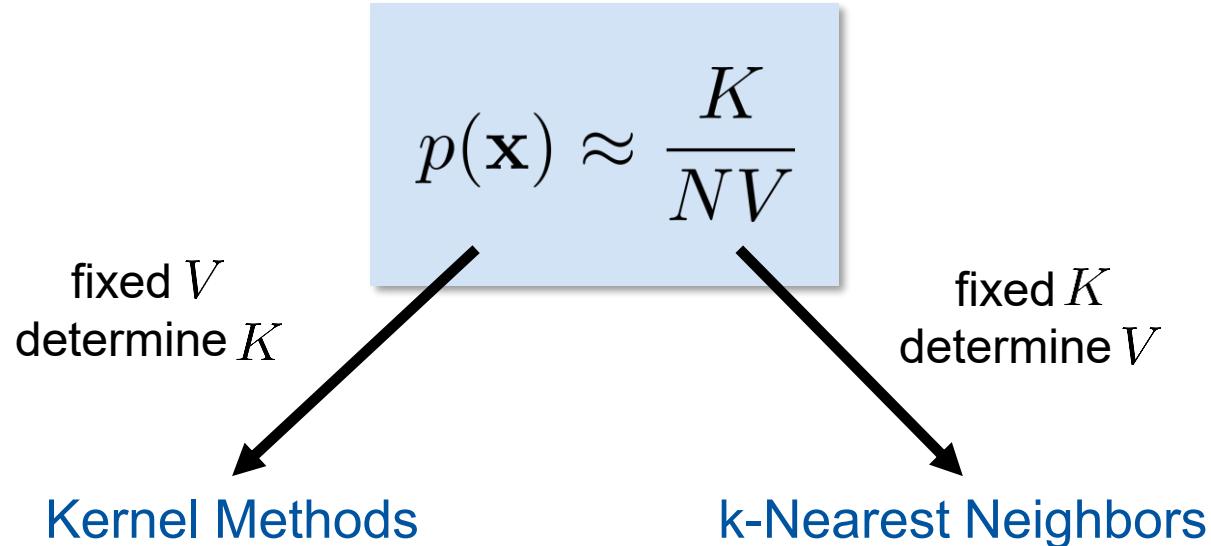
- Then, we get the probability density estimate

$$p(\mathbf{x}) \approx \frac{K}{NV} = \frac{1}{N} \sum_{n=1}^N k(\mathbf{x} - \mathbf{x}_n)$$

- This is known as [Kernel Density Estimation](#).



*E.g., a Gaussian kernel
for smoother boundaries.*



Increase the volume V
until the K next data
points are found.

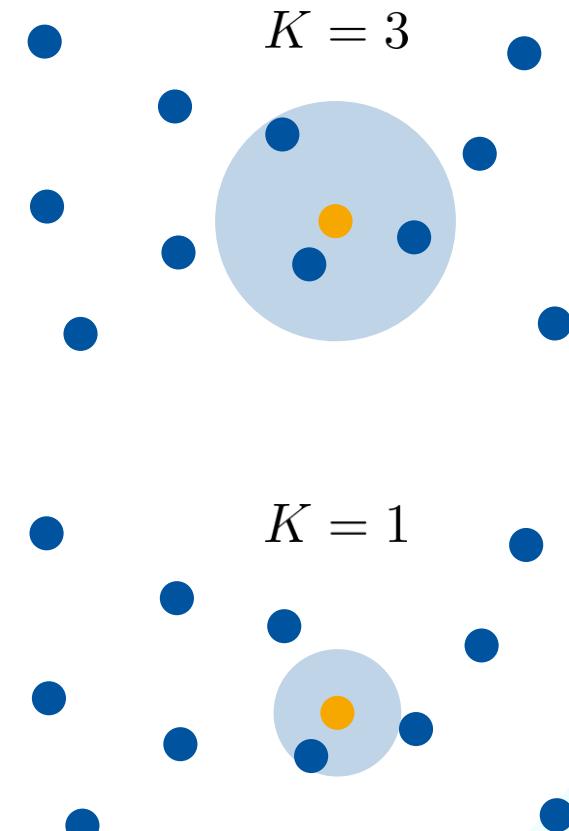
k-Nearest Neighbors

- Fix K , estimate V from the data.
- Consider a hypersphere centered on \mathbf{x} and let it grow to a volume V^* that includes K of the given N data points.
- Then

$$p(\mathbf{x}) \approx \frac{K}{NV^*}$$

- Side note:
 - Strictly speaking, the model produced by k-NN is not a true density model, because the integral over all space diverges.
 - E.g. consider $K = 1$ and a sample exactly on a data point.

$$V^* = 0 \quad \Rightarrow \quad p(\mathbf{x}) \approx \frac{K}{N \cdot 0}$$



Advantages

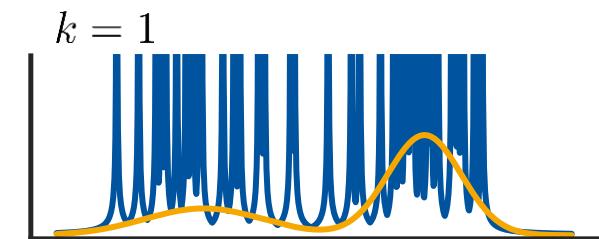
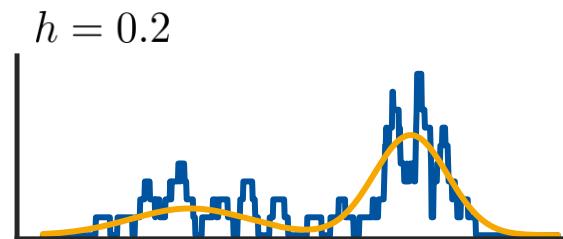
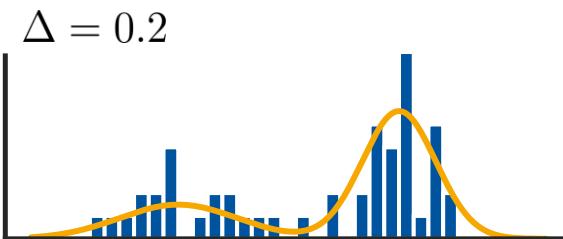
- Very general. In the limit ($N \rightarrow \infty$), every probability density can be represented.
- No computation during training phase
 - Just need to store training set

Limitations

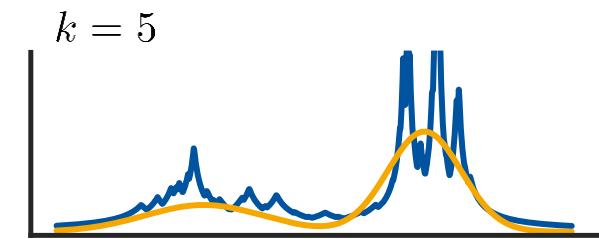
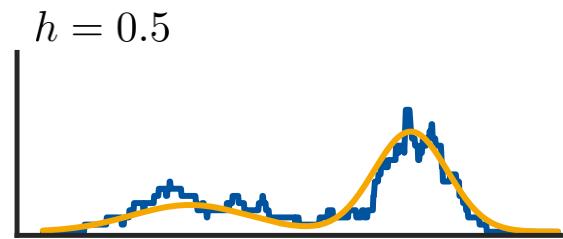
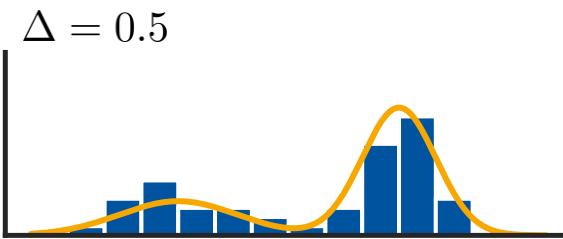
- Requires storing and computing with the entire dataset.
 - Computational costs linear in the number of data points.
 - Can be improved through efficient storage structures (at the cost of some computation during training).
- Choosing the kernel size/ K is a hyperparameter optimization problem.

Bias-Variance Tradeoff

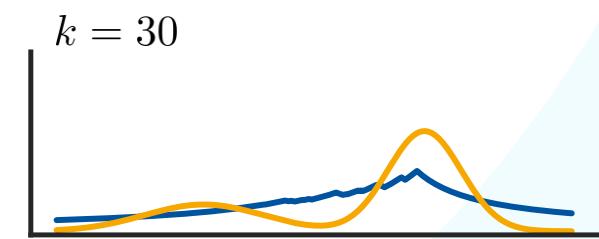
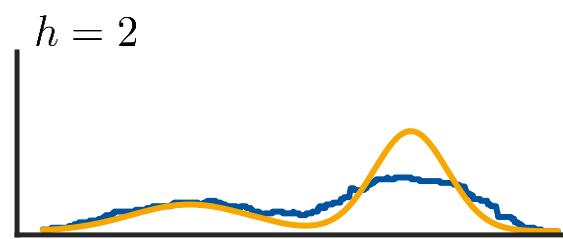
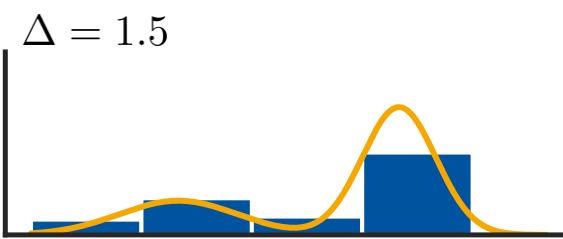
Not smooth enough
Too much variance



About ok



Too smooth
Too much bias



Histograms:
Bin width Δ

KDE:
Kernel size h

k-NN:
of neighbors k

References and Further Reading

- More information in Bishop's book
 - Gaussian distribution and ML: Ch. 1.2.4 and 2.3.1-2.3.4.
 - Nonparametric methods: Ch. 2.5.

Christopher M. Bishop
Pattern Recognition and Machine Learning
Springer, 2006

