



Visual Computing
Institute

RWTH AACHEN
UNIVERSITY

Elements of Machine Learning & Data Science

Winter semester 2025/26

Lecture 10 – Probability Density Estimation

25.11.2025

Prof. Bastian Leibe

Machine Learning Topics

8. Introduction to ML

9. Probability Density Estimation

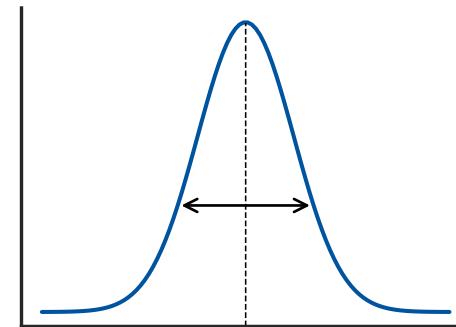
10. Linear Discriminants

11. Linear Regression

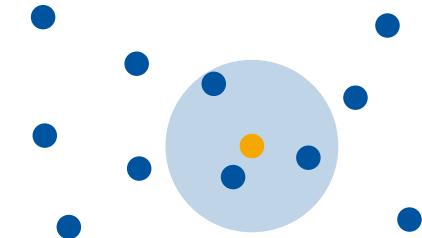
12. Logistic Regression

13. Support Vector Machines

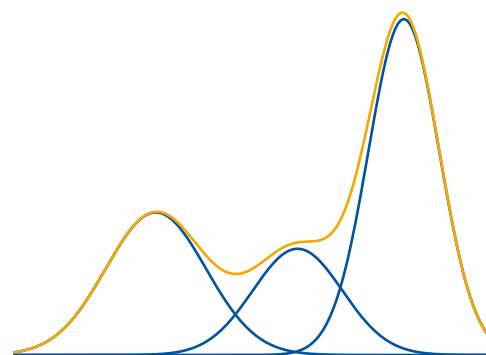
14. Neural Network Basics



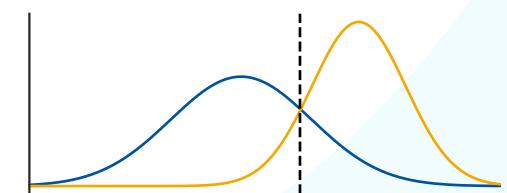
Parametric Methods
& ML-Algorithm



Nonparametric Methods



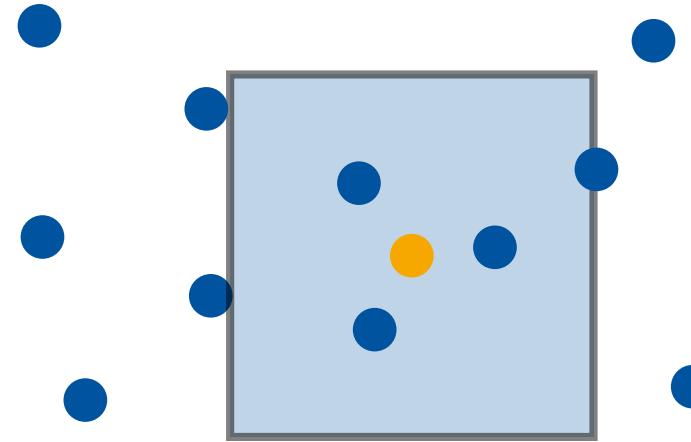
Mixtures of Gaussians
& EM-Algorithm



Bayes Classifiers

Probability Density Estimation

1. Probability Distributions
2. Parametric Methods
3. **Nonparametric Methods**
 - a) Histograms
 - b) **Kernel Methods & k-Nearest Neighbors**
4. Mixture Models
5. Bayes Classifier
6. K-NN Classifier



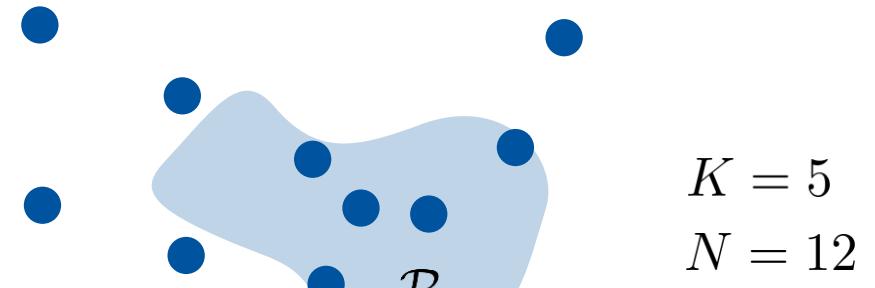
Kernel Methods and k-Nearest Neighbors

- Data point \mathbf{x} comes from pdf $p(\mathbf{x})$.
 - Probability that \mathbf{x} falls into small region \mathcal{R} :

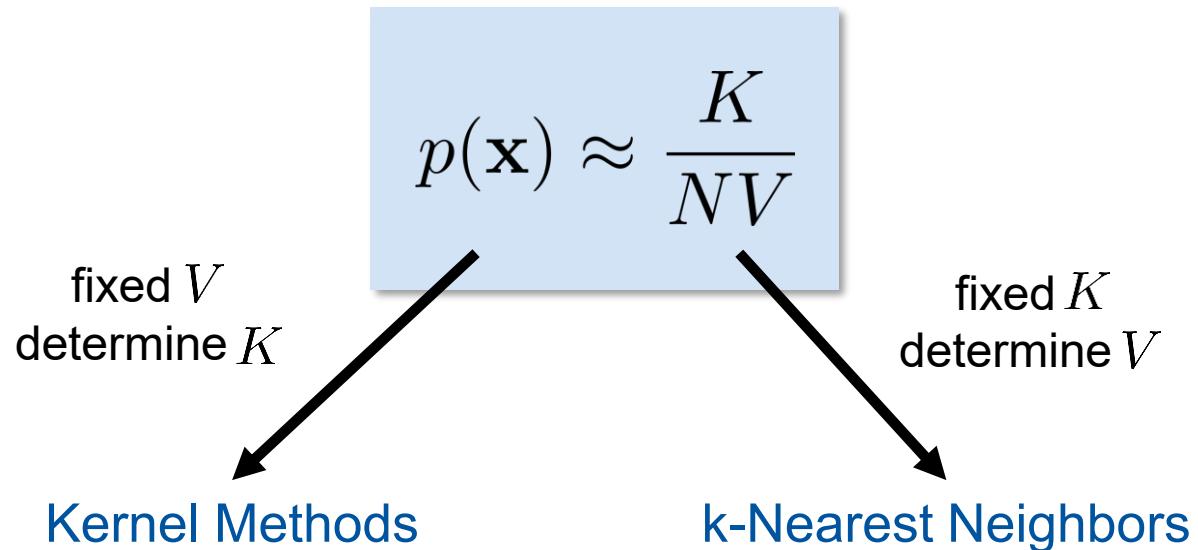
$$P = \int_{\mathcal{R}} p(y) dy \approx p(\mathbf{x})V$$

- Estimate $p(\mathbf{x})$ from samples
 - Let K be the number of samples that fall into \mathcal{R} .
 - If the number of samples N is sufficiently large, we can estimate P as:

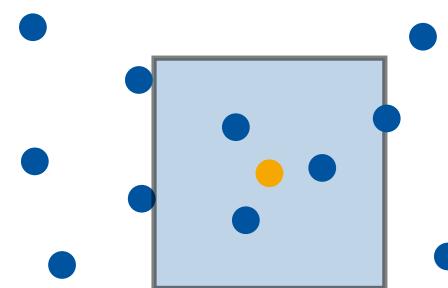
$$P = \frac{K}{N} \Rightarrow p(\mathbf{x}) \approx \frac{K}{NV}$$



For sufficiently small \mathcal{R} ,
 $p(\mathbf{x})$ is roughly constant.
 V : volume of \mathcal{R} .



Example: Determine
the number K of data
points inside a fixed
hypercube



Kernel Methods

- Hypercube of dimension D with edge length h :

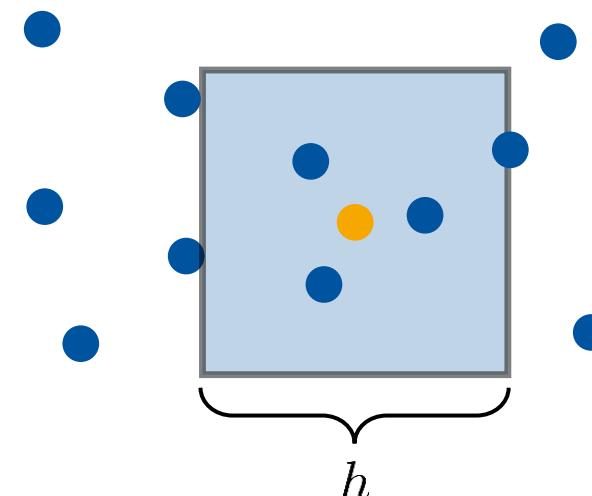
$$k(\mathbf{u}) = \begin{cases} 1, & \text{if } |u_i| \leq \frac{1}{2}h, \ i = 1, \dots, D \\ 0, & \text{otherwise} \end{cases}$$

$$K = \sum_{n=1}^N k(\mathbf{x} - \mathbf{x}_n) \quad V = \int k(\mathbf{u}) d\mathbf{u} = h^D$$

- Probability density estimate:

$$p(\mathbf{x}) \approx \frac{K}{NV} = \frac{1}{Nh^D} \sum_{n=1}^N k(\mathbf{x} - \mathbf{x}_n)$$

- This method is known as **Parzen Window** estimation.



- In general, we can use any kernel such that

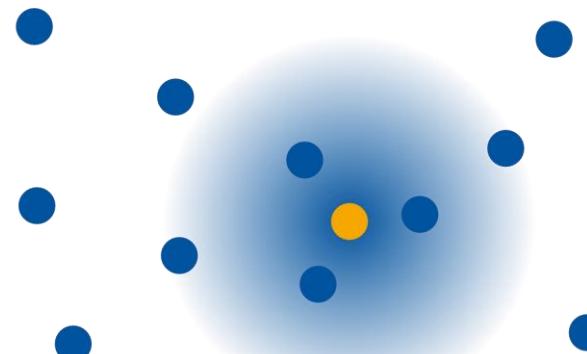
$$k(\mathbf{u}) \geq 0, \quad \int k(\mathbf{u}) d\mathbf{u} = 1$$

$$K = \sum_{n=1}^N k(\mathbf{x} - \mathbf{x}_n)$$

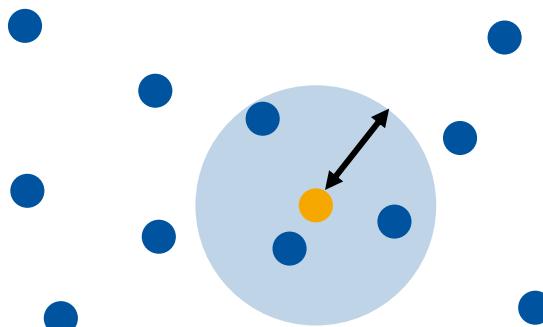
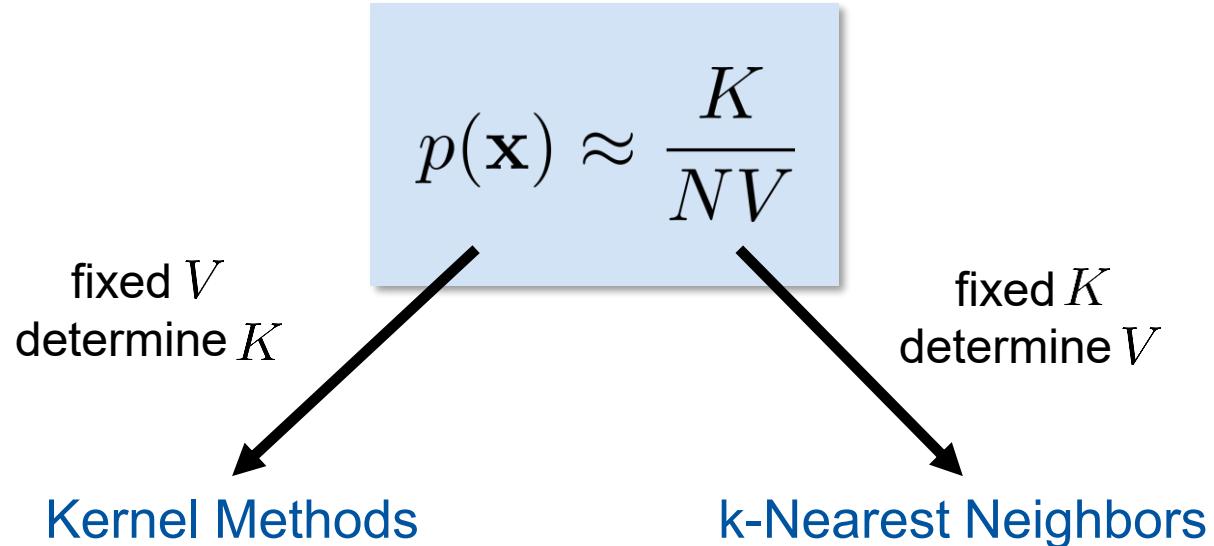
- Then, we get the probability density estimate

$$p(\mathbf{x}) \approx \frac{K}{NV} = \frac{1}{N} \sum_{n=1}^N k(\mathbf{x} - \mathbf{x}_n)$$

- This is known as [Kernel Density Estimation](#).



*E.g., a Gaussian kernel
for smoother boundaries.*



Increase the volume V
until the K next data
points are found.

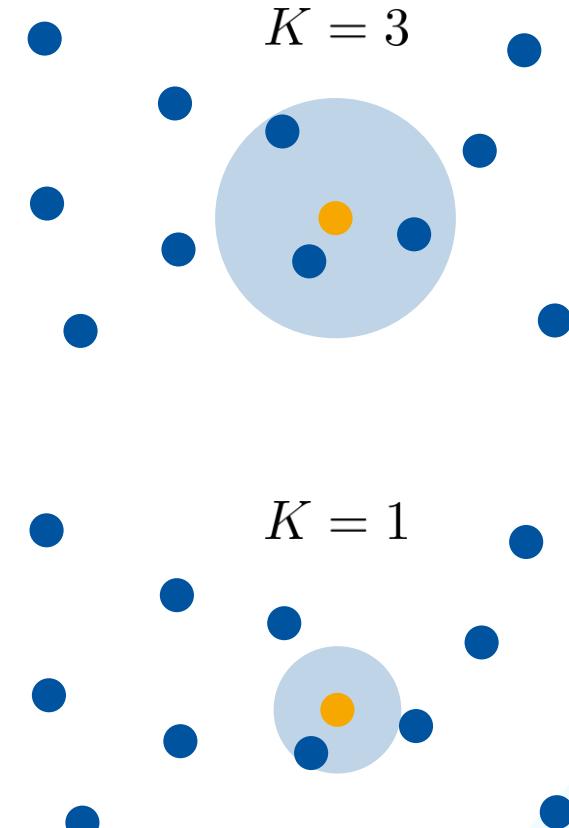
k-Nearest Neighbors

- Fix K , estimate V from the data.
- Consider a hypersphere centered on \mathbf{x} and let it grow to a volume V^* that includes K of the given N data points.
- Then

$$p(\mathbf{x}) \approx \frac{K}{NV^*}$$

- Side note:
 - Strictly speaking, the model produced by k-NN is not a true density model, because the integral over all space diverges.
 - E.g. consider $K = 1$ and a sample exactly on a data point.

$$V^* = 0 \quad \Rightarrow \quad p(\mathbf{x}) \approx \frac{K}{N \cdot 0}$$



Advantages

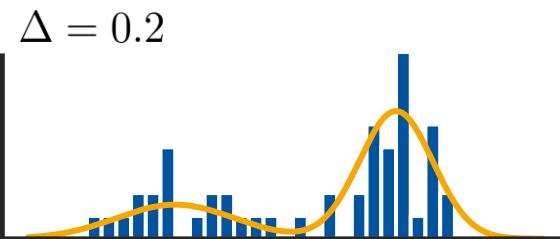
- Very general. In the limit ($N \rightarrow \infty$), every probability density can be represented.
- No computation during training phase
 - Just need to store training set

Limitations

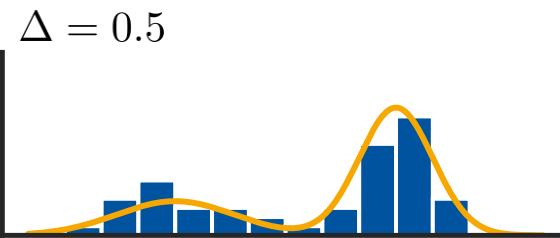
- Requires storing and computing with the entire dataset.
 - Computational costs linear in the number of data points.
 - Can be improved through efficient storage structures (at the cost of some computation during training).
- Choosing the kernel size/ K is a hyperparameter optimization problem.

Bias-Variance Tradeoff

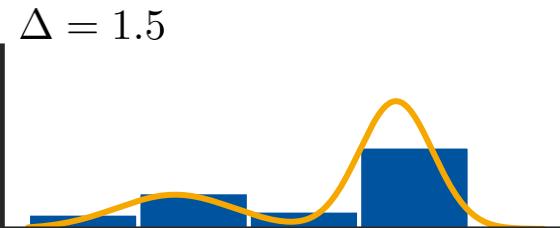
Not smooth enough
Too much variance



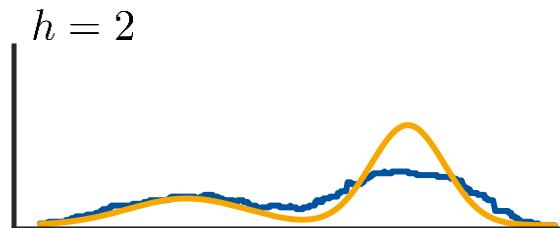
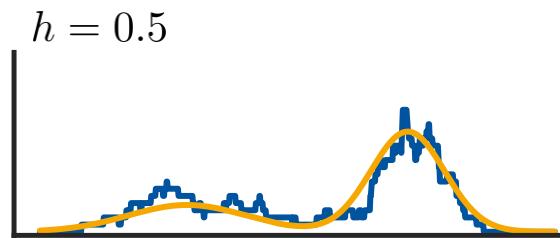
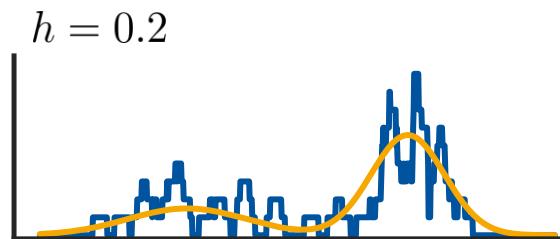
About ok



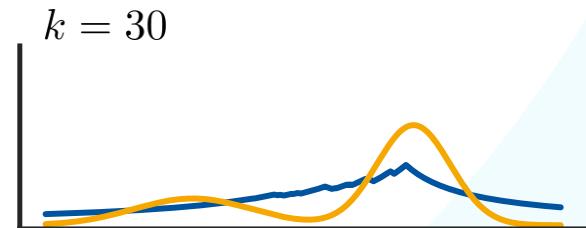
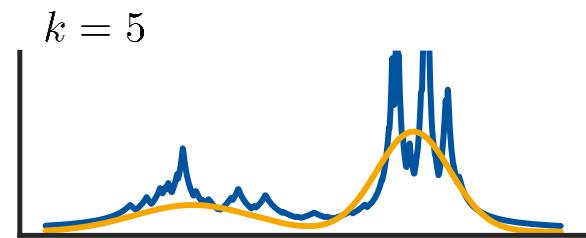
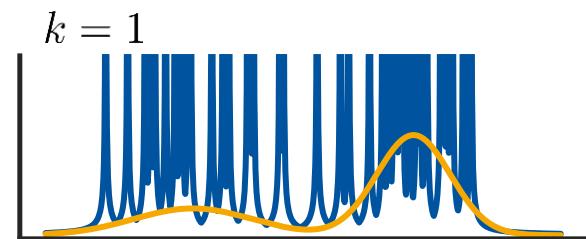
Too smooth
Too much bias



Histograms:
Bin width Δ



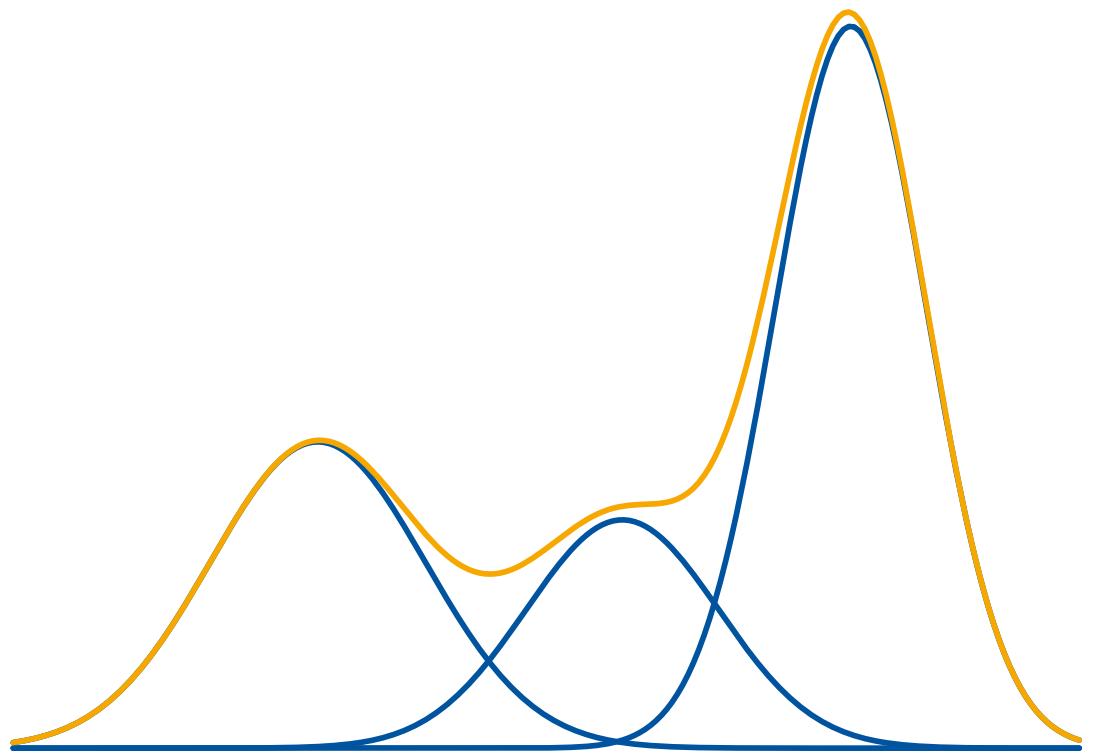
Parzen Window:
Kernel size h



k-NN:
of neighbors k

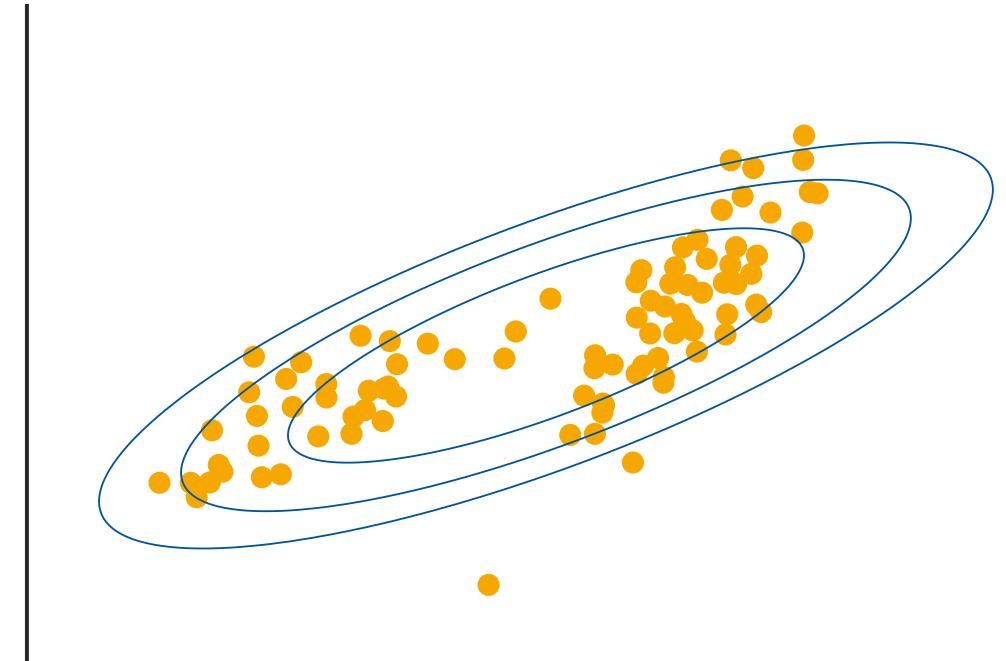
Probability Density Estimation

1. Probability Distributions
2. Parametric Methods
3. Nonparametric Methods
- 4. Mixture Models**
5. Bayes Classifier
6. K-NN Classifier



Mixture Models

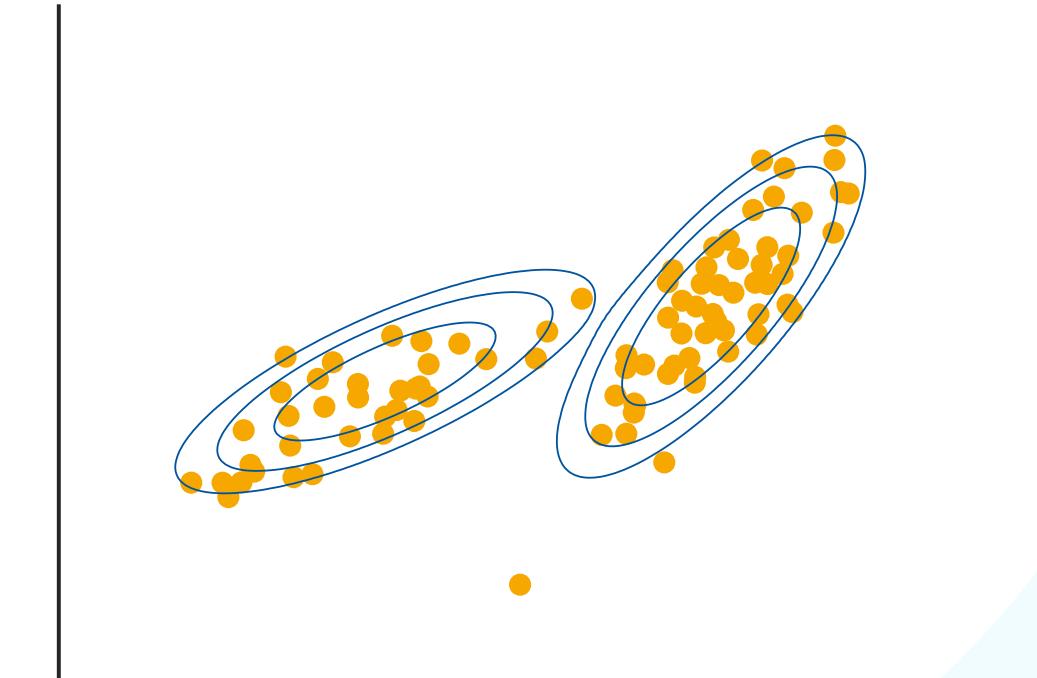
- Often, a single parametric representation is not enough.
- Struggle to fit multimodal data



Single Gaussian

Mixture Models

- Often, a single parametric representation is not enough.
- Struggle to fit multimodal data
- Mixture models combine multiple densities into a single distribution.
 - Improves modeling of multimodal data.



Mixture of two Gaussians

Mixture of Gaussians (MoG)

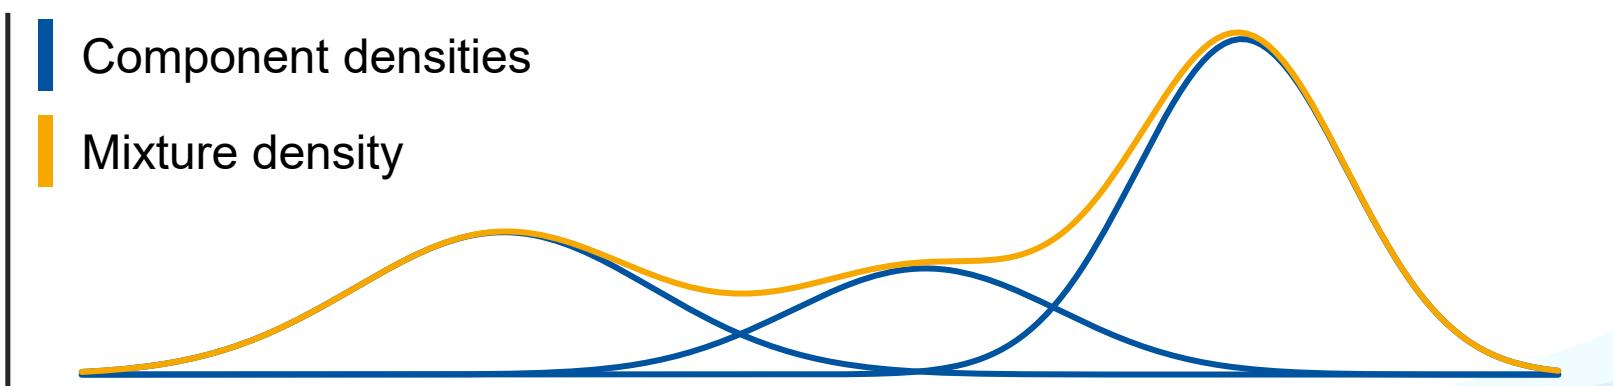
- This is the sum of M individual Normal distributions:

$$p(x|\theta) = \sum_{j=1}^M p(x|\theta_j)p(j)$$

Likelihood of measurement x given mixture component j

Prior of component j

- In the limit, every smooth distribution can be approximated this way (if M is large enough).



- For Gaussians, the complete mixture model is given as:

$$p(x|\theta) = \sum_{j=1}^M p(x|\theta_j)p(j)$$

$$p(x|\theta_j) = \mathcal{N}(x|\mu_j, \sigma_j^2) = \frac{1}{\sqrt{2\pi}\sigma_j} \exp\left(-\frac{(x - \mu_j)^2}{2\sigma_j^2}\right)$$

$$p(j) = \pi_j \text{ with } 0 \leq \pi_j \leq 1 \text{ and } \sum_{j=1}^M \pi_j = 1$$

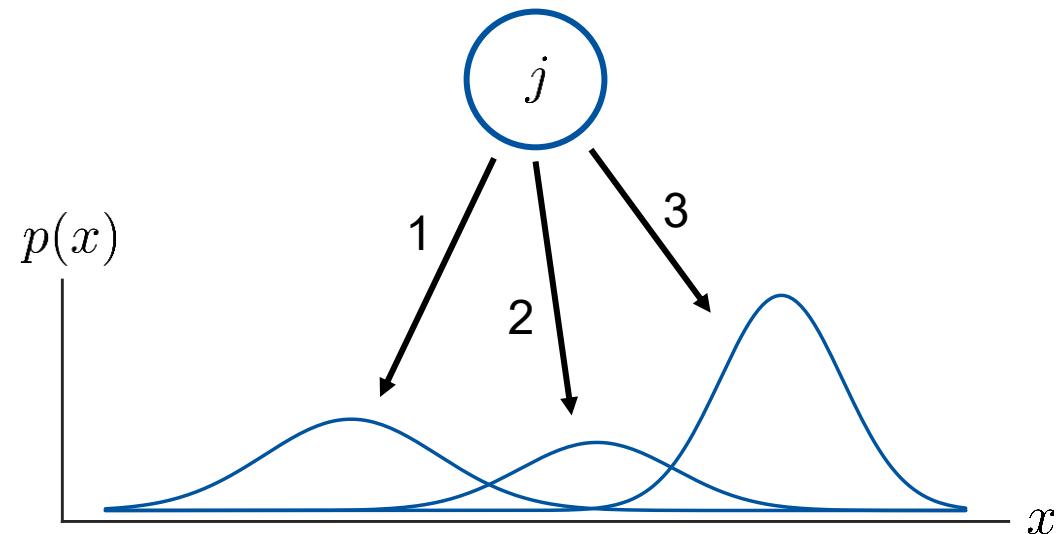
Note: this integrates to 1

$$\int p(x|\theta)dx = 1$$

Total parameters:

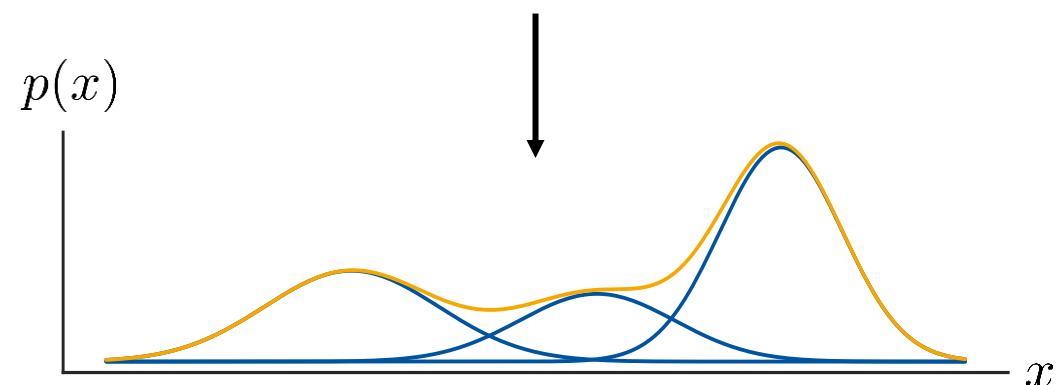
$$\theta = (\pi_1, \mu_1, \sigma_1, \dots, \pi_M, \mu_M, \sigma_M)$$

- MoGs are **generative models**: we can easily sample from them.



$$p(j) = \pi_j$$

Sample the component using its “weight”



$$p(x|\theta_j)$$

Sample from the mixture component

$$p(x|\theta) = \sum_{j=1}^M p(x|\theta_j)p(j)$$

The result is a sample from the mixture density

Learning a Mixture Model

- Apply Maximum Likelihood
 - Minimize $E = -\ln L(\theta) = -\sum_{n=1}^N \ln p(\mathbf{x}_n | \theta)$
 - Let's first look at μ_j :

$$\frac{\partial E}{\partial \mu_j} = 0$$

- We can already see that this will be difficult:

$$\ln p(\mathcal{X} | \boldsymbol{\pi}, \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \sum_{n=1}^N \ln \left(\sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \right)$$

Steps for Maximum Likelihood:

1. Express Likelihood $L(\theta)$
2. Apply negative logarithm to get $E(\theta)$
3. Take derivative, set to zero
4. Solve for parameters

Learning a Mixture Model

- Apply Maximum Likelihood
 - Minimize $E = -\ln L(\theta) = -\sum_{n=1}^N \ln p(\mathbf{x}_n | \theta)$
 - Let's first look at μ_j :

$$\frac{\partial E}{\partial \mu_j} = 0$$

- We can already see that this will be difficult:

$$\ln p(\mathcal{X} | \boldsymbol{\pi}, \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \sum_{n=1}^N \ln \left(\sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \right)$$

Steps for Maximum Likelihood:

1. Express Likelihood $L(\theta)$
2. Apply negative logarithm to get $E(\theta)$
3. Take derivative, set to zero
4. Solve for parameters

This will cause problems!

$$\begin{aligned}
 \frac{\partial E}{\partial \mu_j} &= - \sum_{n=1}^N \frac{\frac{\partial}{\partial \mu_j} p(\mathbf{x}_n | \theta_j)}{\sum_{k=1}^K p(\mathbf{x}_n | \theta_k)} \\
 &= - \sum_{n=1}^N \left(\Sigma^{-1} (\mathbf{x}_n - \mu_j) \frac{p(\mathbf{x}_n | \theta_j)}{\sum_{k=1}^K p(\mathbf{x}_n | \theta_k)} \right) \\
 &= - \Sigma^{-1} \sum_{n=1}^N (\mathbf{x}_n - \mu_j) \frac{\pi_j \mathcal{N}(\mathbf{x}_n | \mu_j, \Sigma_j)}{\sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k)} \stackrel{!}{=} 0
 \end{aligned}$$

$$\begin{aligned}
 \frac{\partial}{\partial \mu_j} \mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k) &= \\
 \Sigma^{-1} (\mathbf{x}_n - \mu_j) \mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k)
 \end{aligned}$$

↓

$$= \gamma_j(\mathbf{x}_n)$$

“responsibility” of
component j for \mathbf{x}_n

We thus obtain $\mu_j = \frac{\sum_{n=1}^N \gamma_j(\mathbf{x}_n) \mathbf{x}_n}{\sum_{n=1}^N \gamma_j(\mathbf{x}_n)}$

$$\boldsymbol{\mu}_j = \frac{\sum_{n=1}^N \gamma_j(\mathbf{x}_n) \mathbf{x}_n}{\sum_{n=1}^N \gamma_j(\mathbf{x}_n)}$$

$$\gamma_j(\mathbf{x}_n) = \frac{\pi_j \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)}{\sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}$$

- There is no direct analytical solution!

$$\frac{\partial E}{\partial \boldsymbol{\mu}_j} = f(\pi_1, \boldsymbol{\mu}_1, \boldsymbol{\Sigma}_1, \dots, \pi_M, \boldsymbol{\mu}_M, \boldsymbol{\Sigma}_M)$$

- Complex gradient function (non-linear mutual dependencies)
- Optimization of one Gaussian depends on all other Gaussians!
- Standard solution: iterative optimization with EM algorithm

The EM Algorithm

- The Expectation-Maximization (EM) Algorithm alternates between two steps:
 - **E-Step:** softly assign samples to mixture components:

$$\gamma_j(\mathbf{x}_n) \leftarrow \frac{\pi_j \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)}{\sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)} \quad \forall j = 1, \dots, K, n = 1, \dots, N$$

- **M-Step:** re-estimate parameters of each component based on the soft assignments:

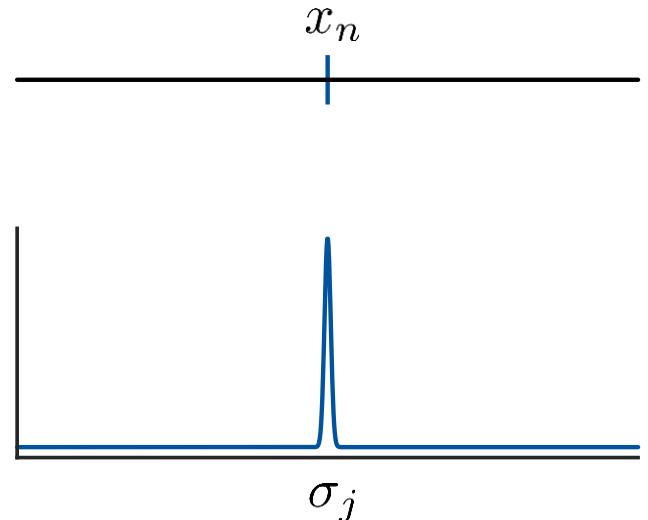
$$\begin{aligned} \hat{N}_j &\leftarrow \sum_{n=1}^N \gamma_j(\mathbf{x}_n) & \hat{\boldsymbol{\mu}}_j^{\text{new}} &\leftarrow \frac{1}{\hat{N}_j} \sum_{n=1}^N \gamma_j(\mathbf{x}_n) \mathbf{x}_n \\ \hat{\pi}_j^{\text{new}} &\leftarrow \frac{\hat{N}_j}{N} & \hat{\boldsymbol{\Sigma}}_j^{\text{new}} &\leftarrow \frac{1}{\hat{N}_j} \sum_{n=1}^N \gamma_j(\mathbf{x}_n) (\mathbf{x}_n - \hat{\boldsymbol{\mu}}_j^{\text{new}}) (\mathbf{x}_n - \hat{\boldsymbol{\mu}}_j^{\text{new}})^T \end{aligned}$$

Practical Advice

- When implementing EM, we need to take care to avoid singularities in the estimation!
 - Mixture components may collapse on single data points.
 - E.g. consider the case $\Sigma_k = \sigma_k^2 \mathbf{I}$ (this also holds in general)
 - Assume component j is exactly centered on data point \mathbf{x}_n . This data point will then contribute a term in the likelihood function

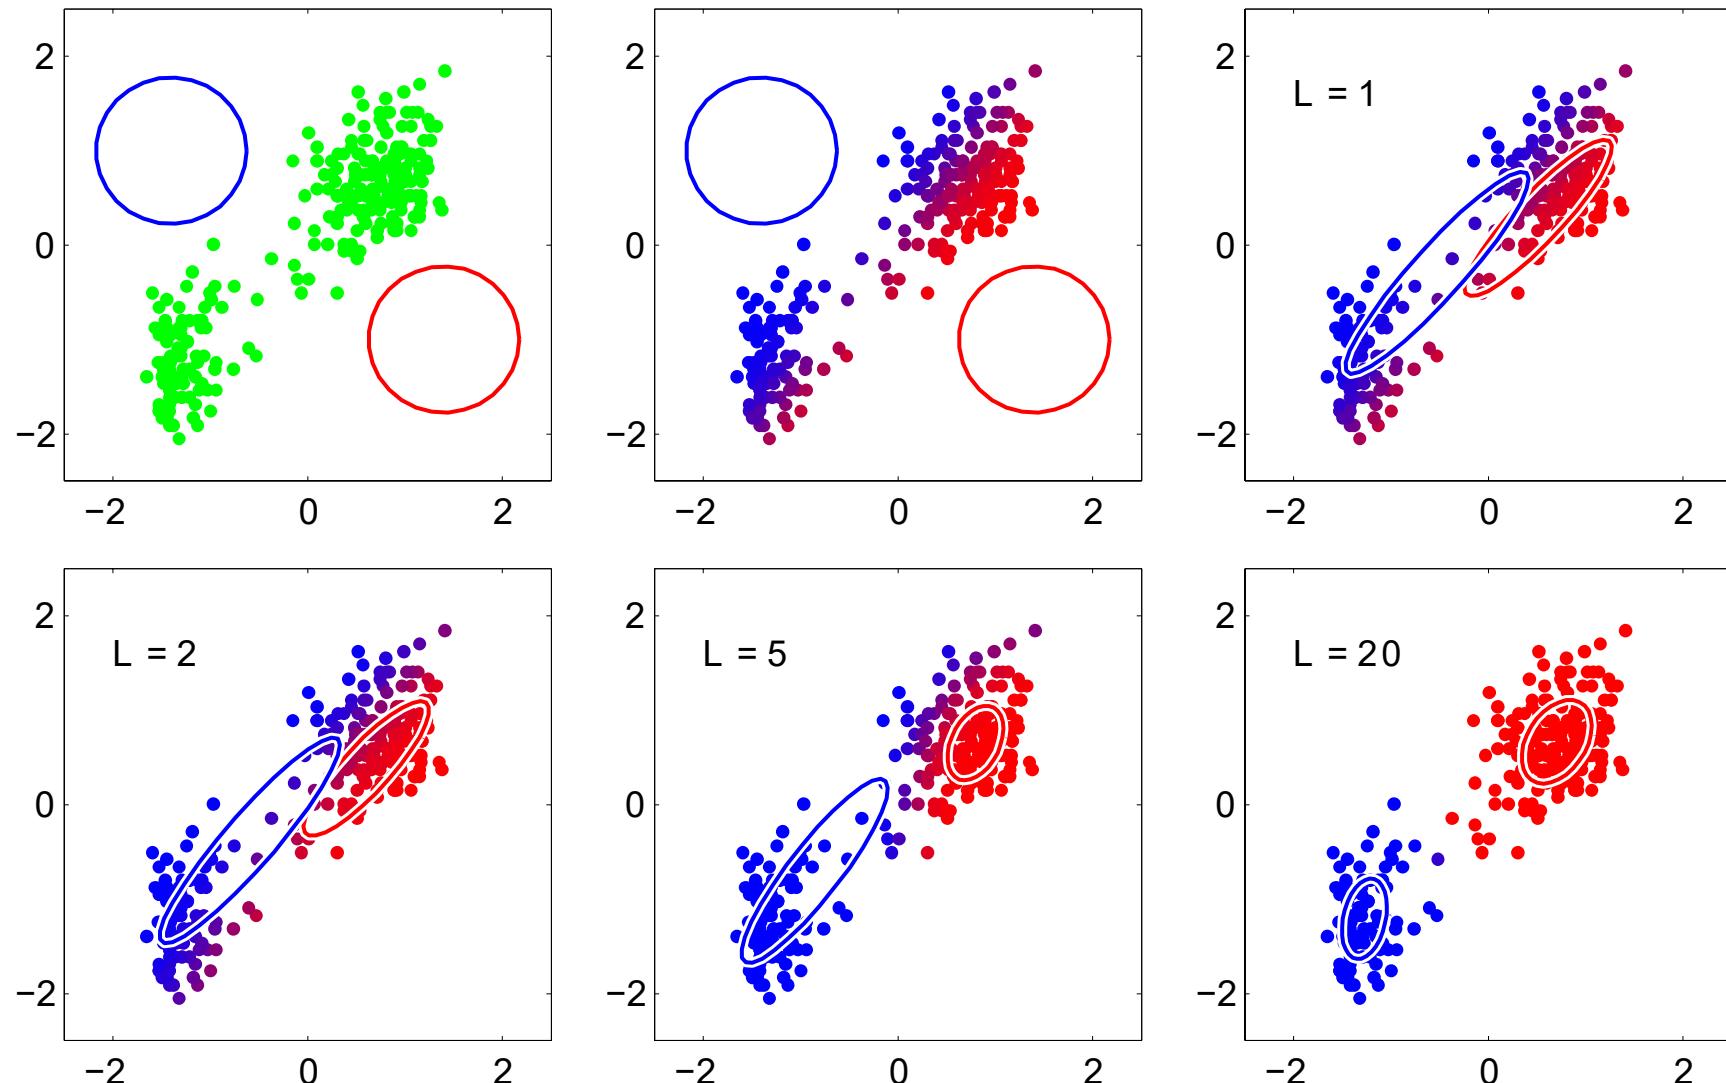
$$\mathcal{N}(\mathbf{x}_n | \mathbf{x}_n, \sigma_j^2 \mathbf{I}) = \frac{1}{\sqrt{2\pi}\sigma_j}$$

- For $\sigma_j \rightarrow 0$, this term goes to infinity!
- We need to introduce **regularization** to avoid this.
 - Enforce minimum width for the Gaussians



Instead of Σ^{-1} , use $(\Sigma + \sigma_{\min} \mathbf{I})^{-1}$.

EM Algorithm – An Example



Discussion: Mixture Models

Advantages

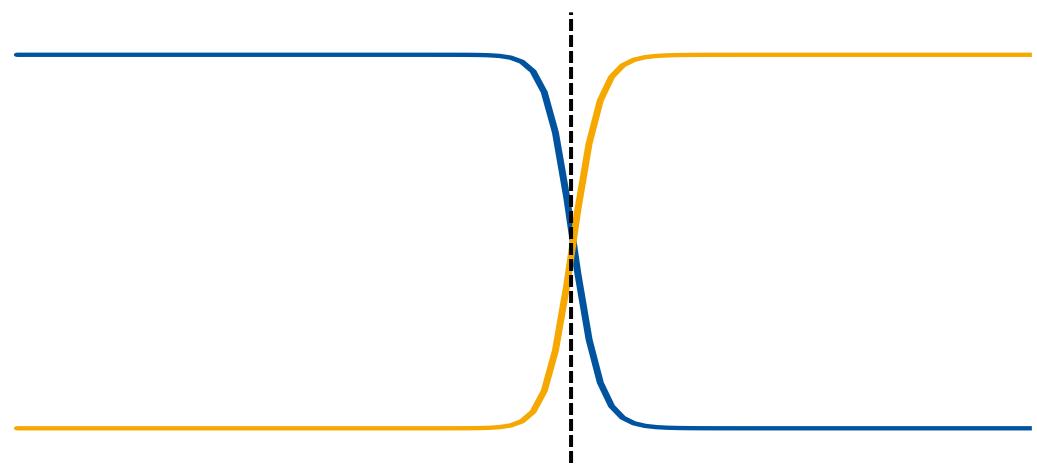
- Very general, can represent any continuous distribution.
- Once trained, is very fast to evaluate.

Limitations

- Need to apply regularization to avoid numerical instabilities.
- Choosing the right number of mixture components is hard.
- The EM algorithm is computationally expensive.
 - Especially for high-dim. problems.
 - Very sensitive to initialization.
 - *Practical Tip:* Run k-Means first and initialize clusters with k-Means result

Probability Density Estimation

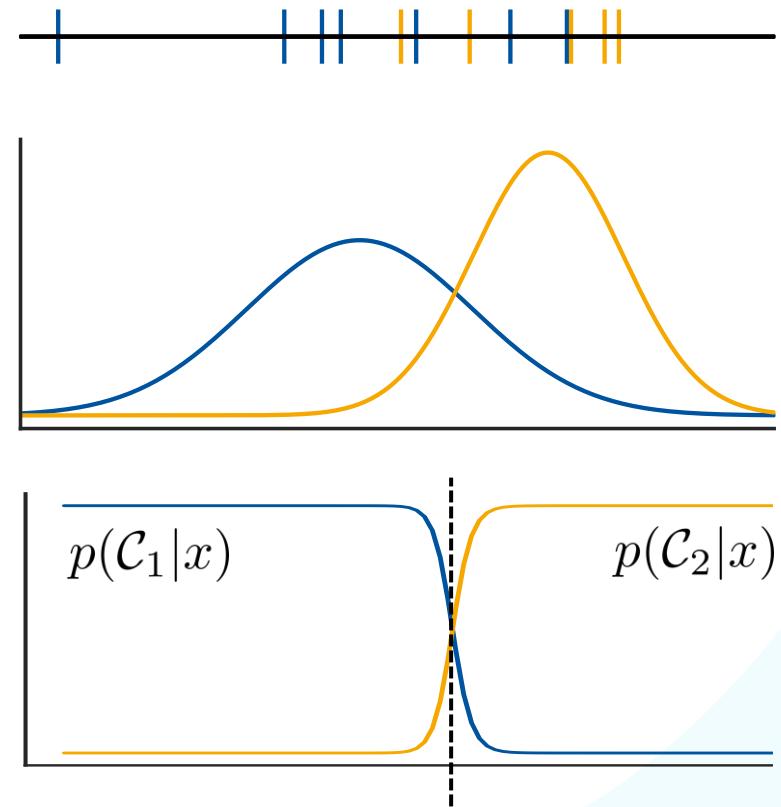
1. Probability Distributions
2. Parametric Methods
3. Nonparametric Methods
4. Mixture Models
5. **Bayes Classifier**
6. K-NN Classifier



Bayes Classifier

- We know how to estimate probability densities from data.
- We can now use [Bayes Decision Theory](#) to build a classifier:
 - Estimate likelihoods & priors from data.
 - Calculate posterior with Bayes' Theorem.
 - Decide for class with highest posterior probability:

$$p(\mathcal{C}_1|x) > p(\mathcal{C}_2|x)$$



Likelihood-Ratio Test

- Assume we want to classify an observation x into one of two classes $\mathcal{C}_1, \mathcal{C}_2$.

- Decide for \mathcal{C}_1 if

$$p(\mathcal{C}_1|x) > p(\mathcal{C}_2|x)$$

- This is equivalent to

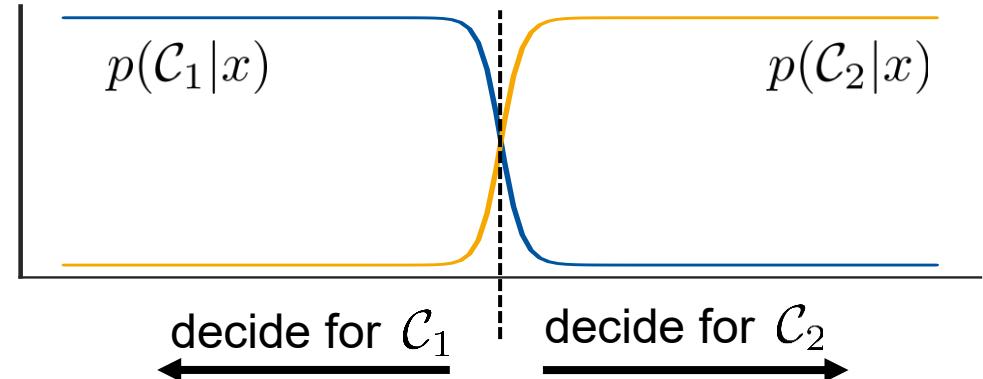
$$p(x|\mathcal{C}_1)p(\mathcal{C}_1) > p(x|\mathcal{C}_2)p(\mathcal{C}_2)$$

- Which again is equivalent to

$$\frac{p(x|\mathcal{C}_1)}{p(x|\mathcal{C}_2)} > \frac{p(\mathcal{C}_2)}{p(\mathcal{C}_1)}$$



Decision threshold θ



$$p(\mathcal{C}|x) = \frac{p(x|\mathcal{C})p(\mathcal{C})}{p(x)}$$

Decision Functions

- We can find a decision function based on probability densities.
 - Determine class-conditional densities $p(x|\mathcal{C}_k)$ for each class individually.
 - Separately infer the prior class probabilities $p(\mathcal{C}_k)$.
 - Then use Bayes' theorem and/or the likelihood-ratio test.
- Alternative: solve the inference problem of determining the posterior class probabilities directly.
 - Then use Bayes' decision theory to assign each new observation to its class.

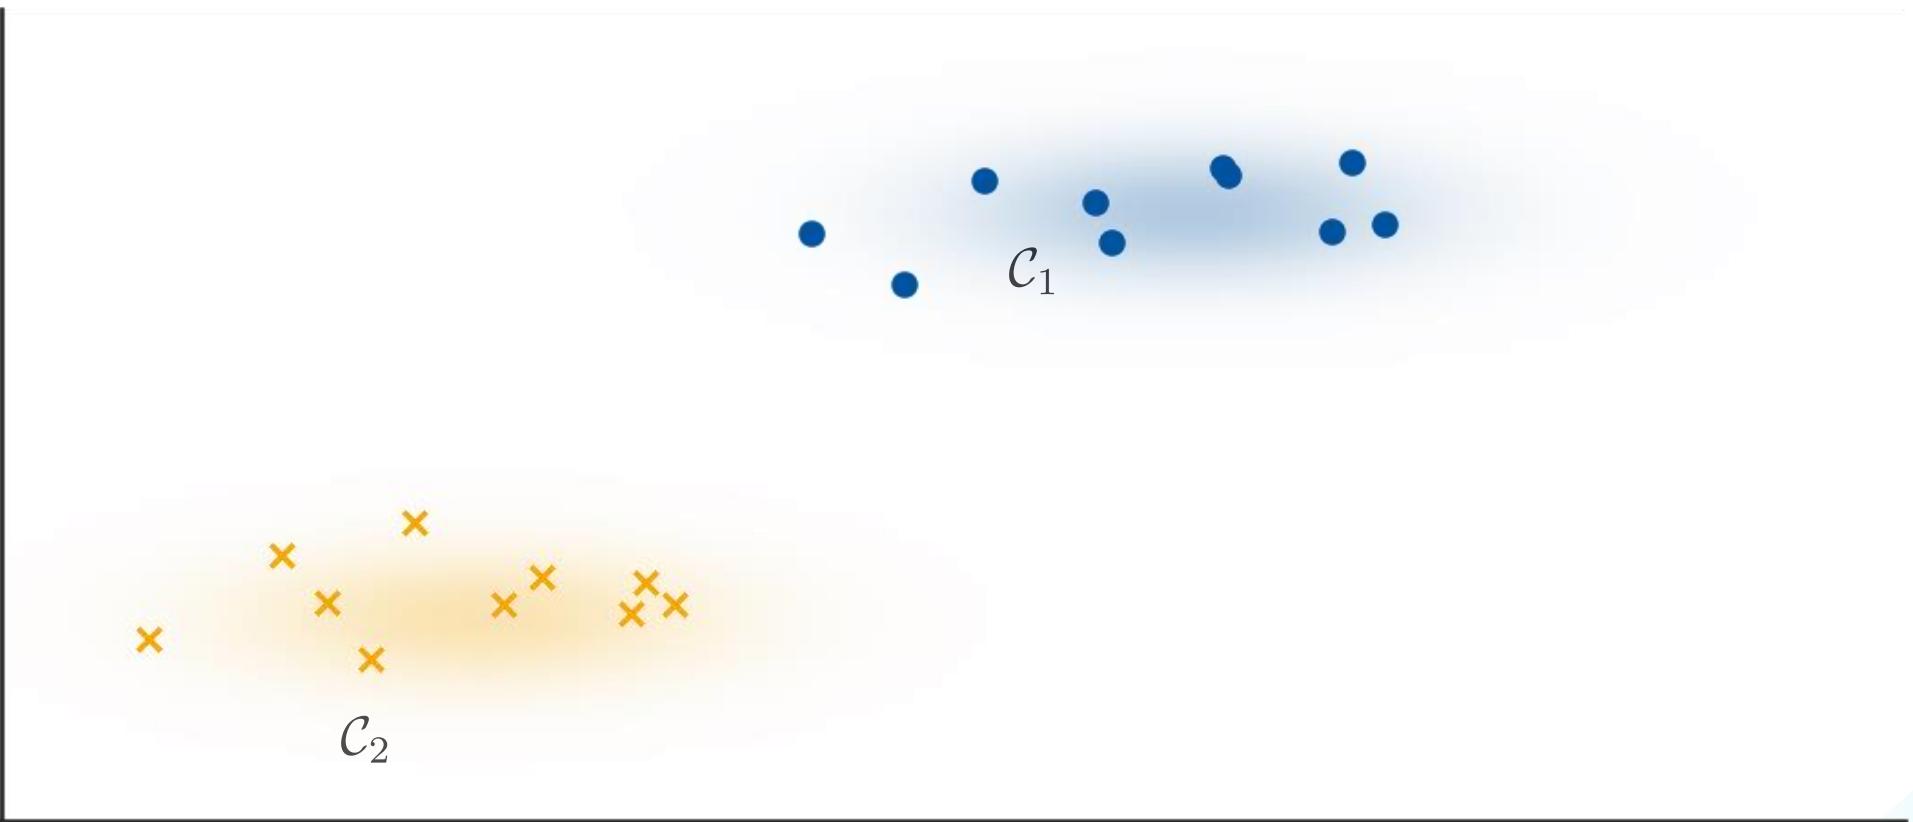
Generative methods:

$$y_k(x) \propto p(x|\mathcal{C}_k)p(\mathcal{C}_k)$$

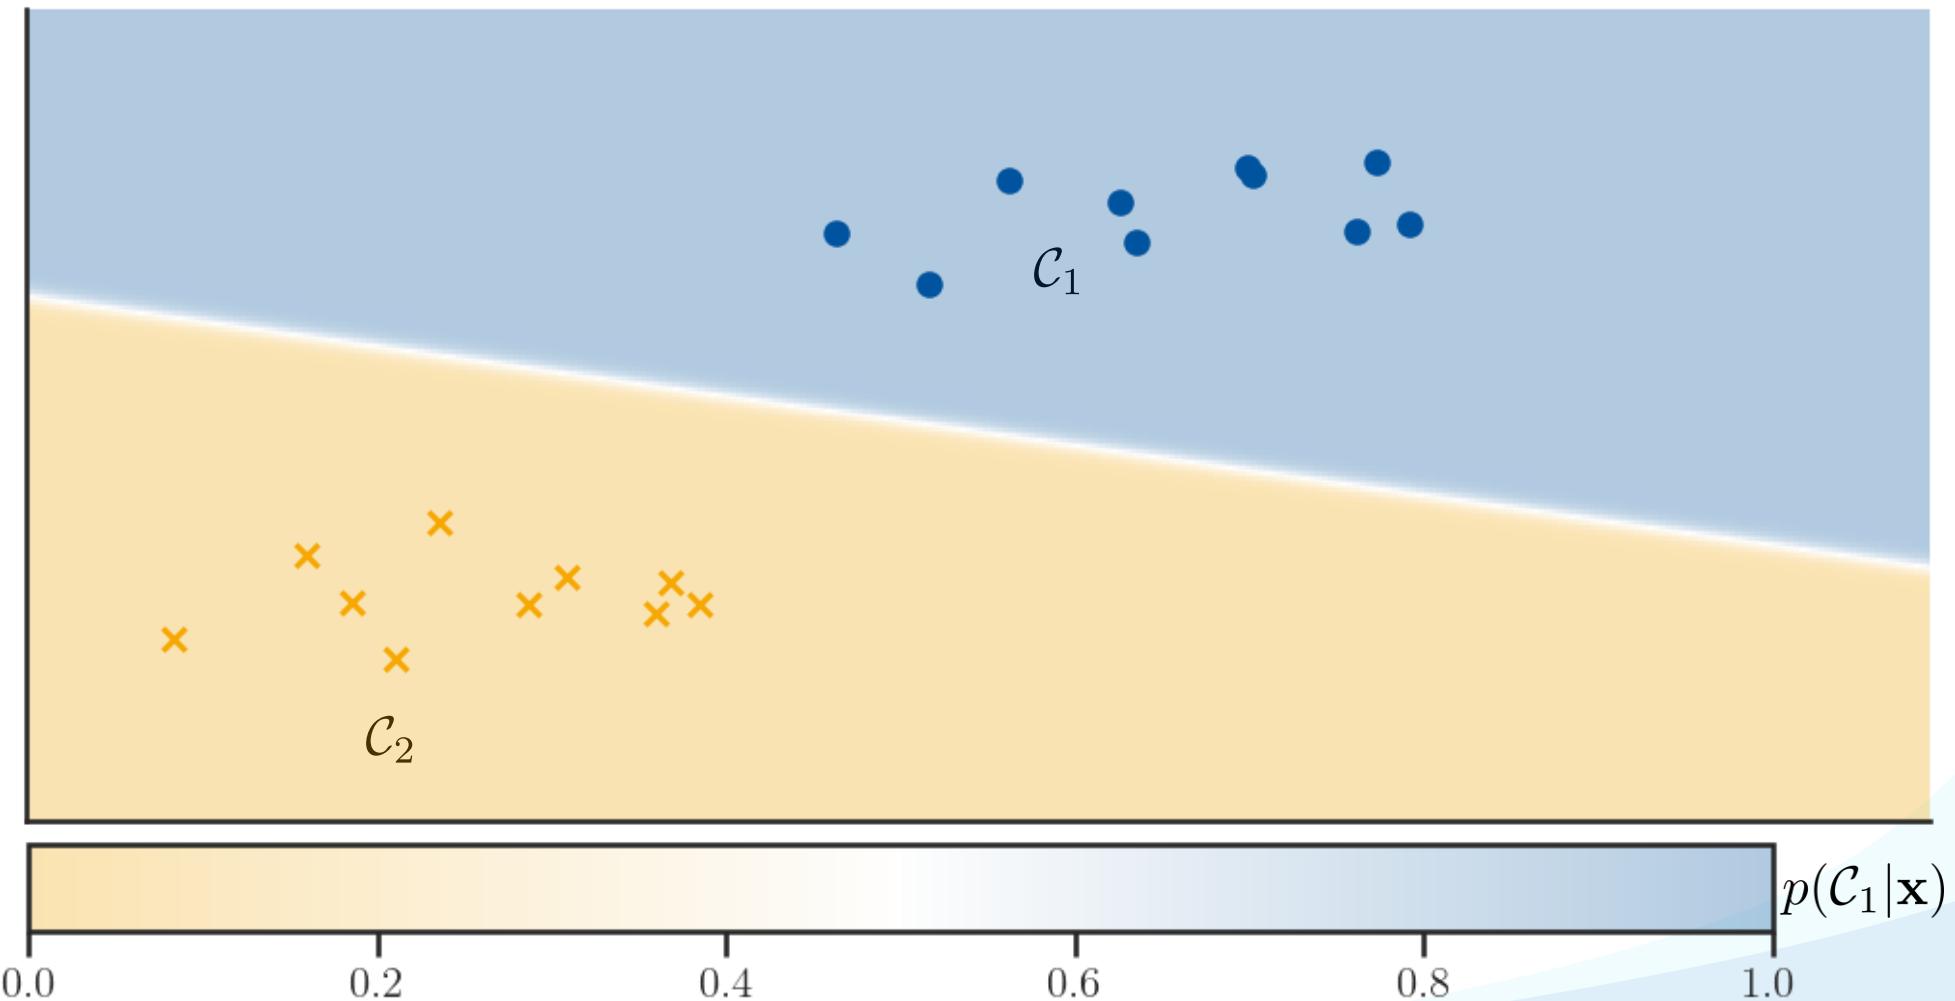
Discriminative methods:

$$y_k(x) = p(\mathcal{C}_k|x)$$

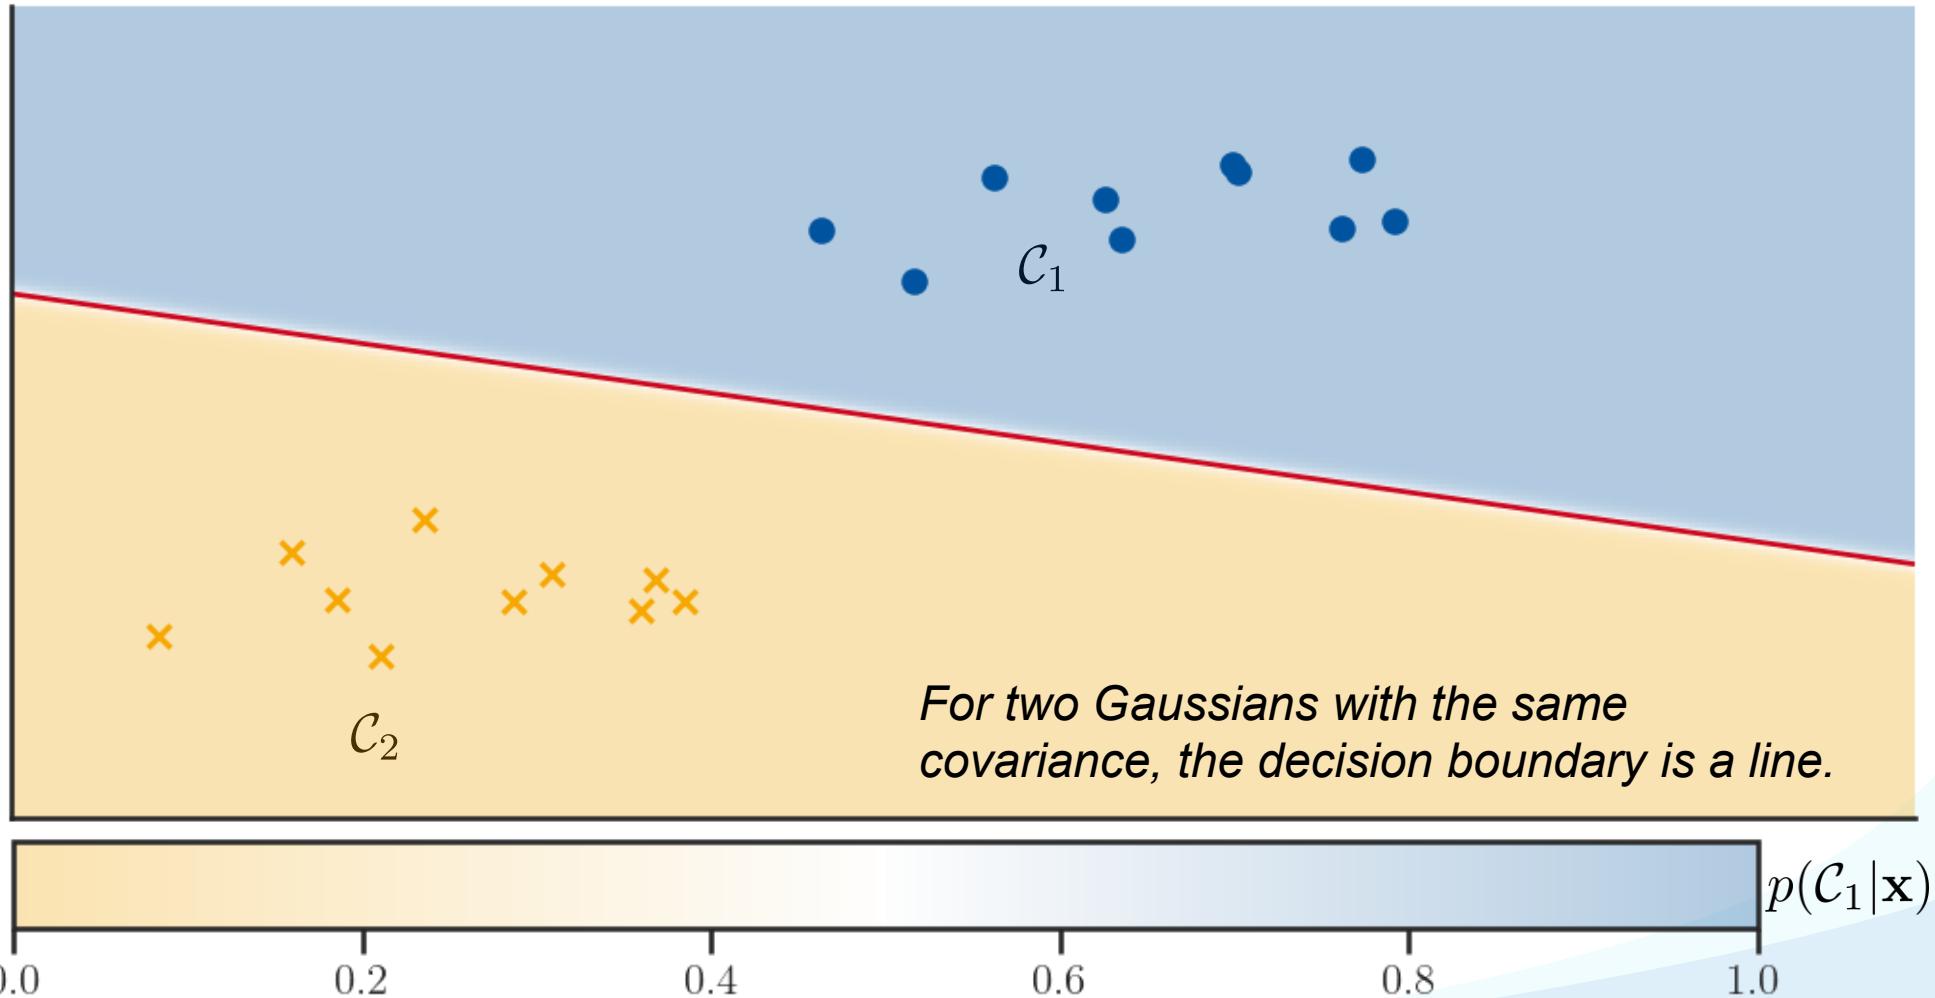
Example



Example

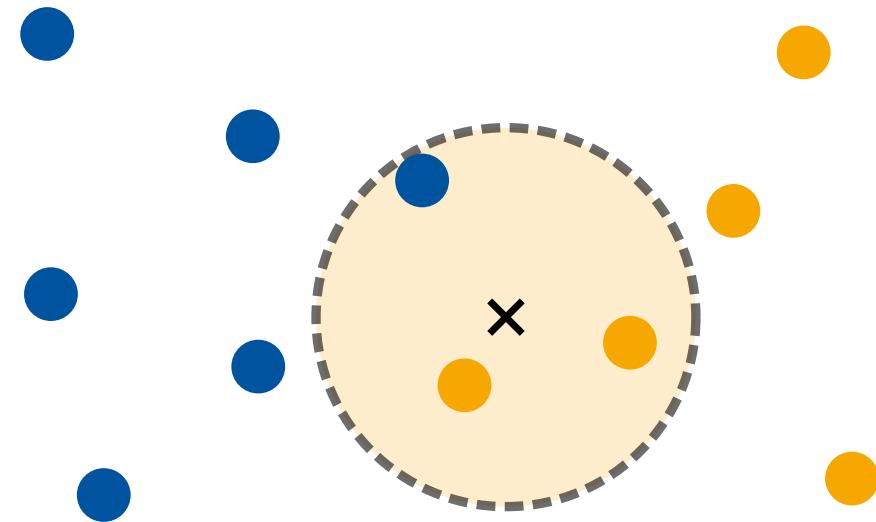


Example



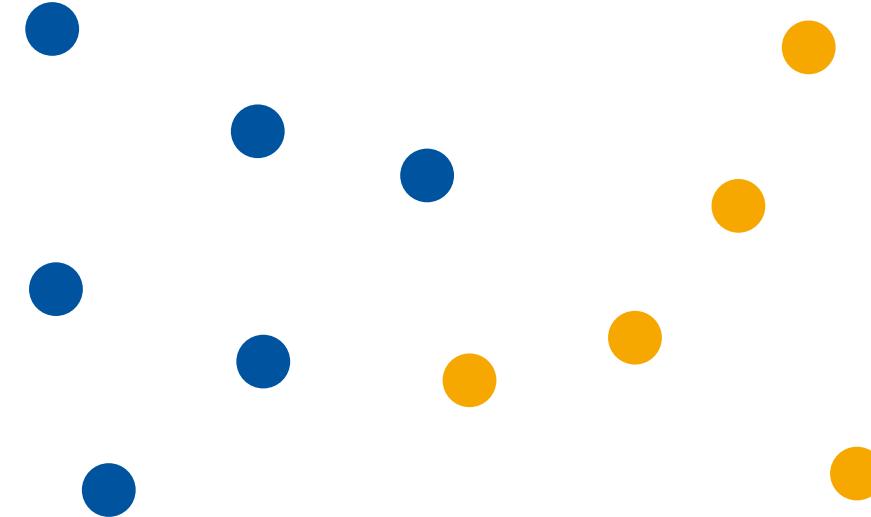
Probability Density Estimation

1. Probability Distributions
2. Parametric Methods
3. Nonparametric Methods
4. Mixture Models
5. Bayes Classifier
6. **K-NN Classifier**



K-NN Classifier

- Combine K-NN density estimation with Bayes Decision Theory: [K-NN Classifier](#)



K-NN Classifier

- Combine K-NN density estimation with Bayes Decision Theory: [K-NN Classifier](#)

- Determine the class-conditional densities

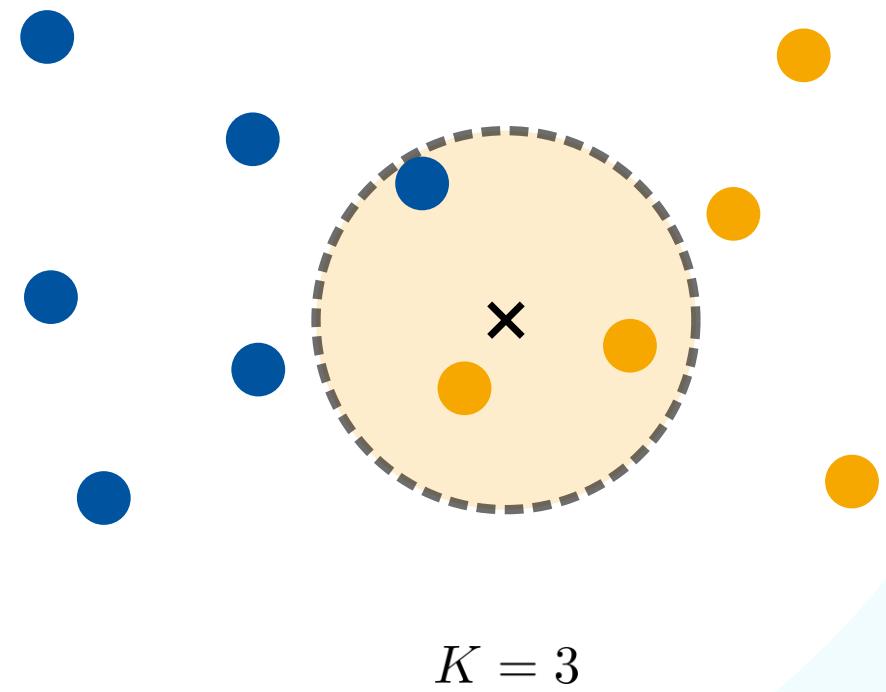
$$p(\mathbf{x}|\mathcal{C}_j) \approx \frac{K_j}{N_j V} \quad p(\mathbf{x}) \approx \frac{K}{NV}$$

- Determine the prior probabilities

$$p(\mathcal{C}_j) \approx \frac{N_j}{N}$$

- Use Bayes' theorem to compute the posterior

$$p(\mathcal{C}_j|\mathbf{x}) \approx p(\mathbf{x}|\mathcal{C}_j)p(\mathcal{C}_j) \frac{1}{p(\mathbf{x})}$$



K-NN Classifier

- Combine K-NN density estimation with Bayes Decision Theory: [K-NN Classifier](#)

- Determine the class-conditional densities

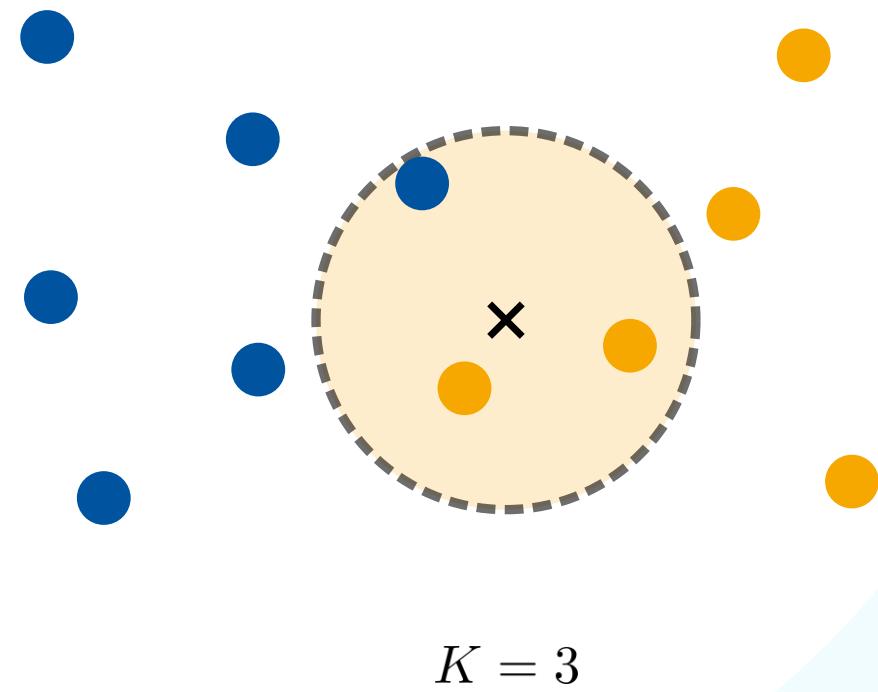
$$p(\mathbf{x}|\mathcal{C}_j) \approx \frac{K_j}{N_j V} \quad p(\mathbf{x}) \approx \frac{K}{NV}$$

- Determine the prior probabilities

$$p(\mathcal{C}_j) \approx \frac{N_j}{N}$$

- Use Bayes' theorem to compute the posterior

$$p(\mathcal{C}_j|\mathbf{x}) \approx \frac{K_j}{N_j V} p(\mathcal{C}_j) \frac{1}{p(\mathbf{x})}$$



K-NN Classifier

- Combine K-NN density estimation with Bayes Decision Theory: [K-NN Classifier](#)

- Determine the class-conditional densities

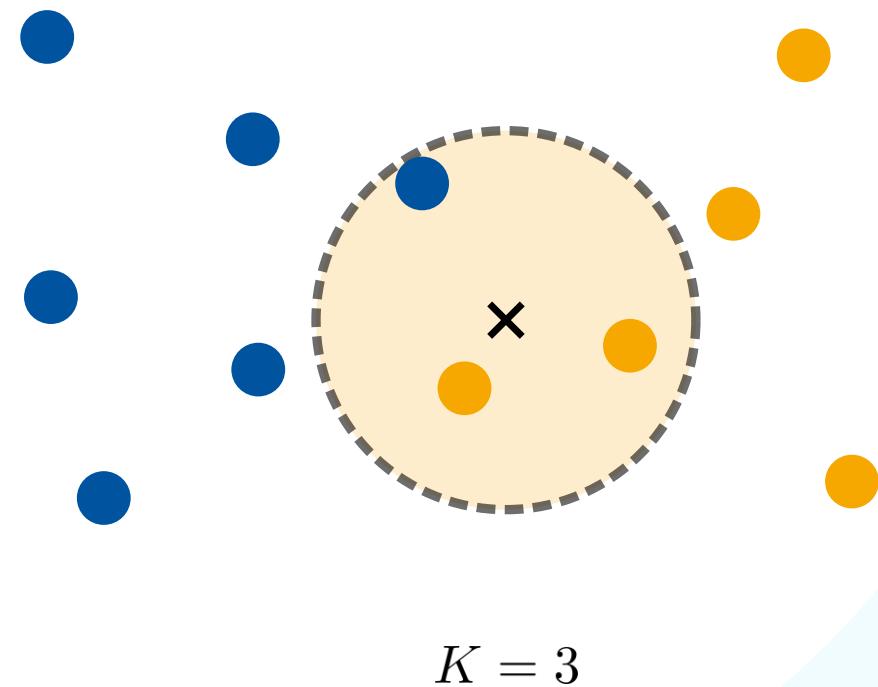
$$p(\mathbf{x}|\mathcal{C}_j) \approx \frac{K_j}{N_j V} \quad p(\mathbf{x}) \approx \frac{K}{NV}$$

- Determine the prior probabilities

$$p(\mathcal{C}_j) \approx \frac{N_j}{N}$$

- Use Bayes' theorem to compute the posterior

$$p(\mathcal{C}_j|\mathbf{x}) \approx \frac{K_j}{N_j V} \frac{N_j}{N} \frac{1}{p(\mathbf{x})}$$



K-NN Classifier

- Combine K-NN density estimation with Bayes Decision Theory: [K-NN Classifier](#)

- Determine the class-conditional densities

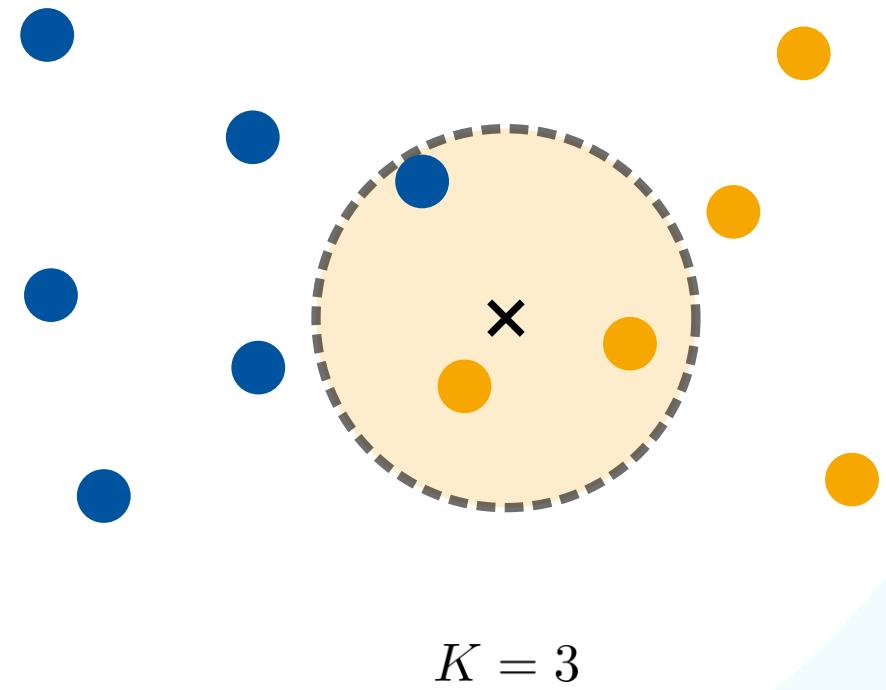
$$p(\mathbf{x}|\mathcal{C}_j) \approx \frac{K_j}{N_j V} \quad p(\mathbf{x}) \approx \frac{K}{NV}$$

- Determine the prior probabilities

$$p(\mathcal{C}_j) \approx \frac{N_j}{N}$$

- Use Bayes' theorem to compute the posterior

$$p(\mathcal{C}_j|\mathbf{x}) \approx \frac{K_j}{N_j V} \frac{N_j}{N} \frac{NV}{K}$$



K-NN Classifier

- Combine K-NN density estimation with Bayes Decision Theory: [K-NN Classifier](#)

- Determine the class-conditional densities

$$p(\mathbf{x}|\mathcal{C}_j) \approx \frac{K_j}{N_j V} \quad p(\mathbf{x}) \approx \frac{K}{NV}$$

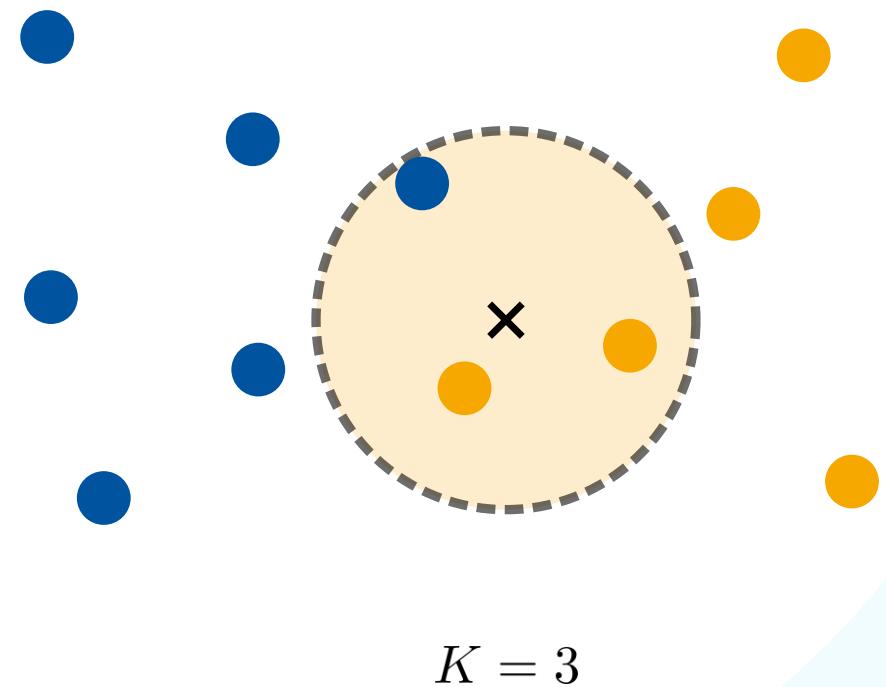
- Determine the prior probabilities

$$p(\mathcal{C}_j) \approx \frac{N_j}{N}$$

- Use Bayes' theorem to compute the posterior

$$p(\mathcal{C}_j|\mathbf{x}) \approx \frac{K_j}{N_j V} \frac{N_j}{N} \frac{NV}{K} = \frac{K_j}{K}$$

\Rightarrow Decide for the majority class among the neighbors.

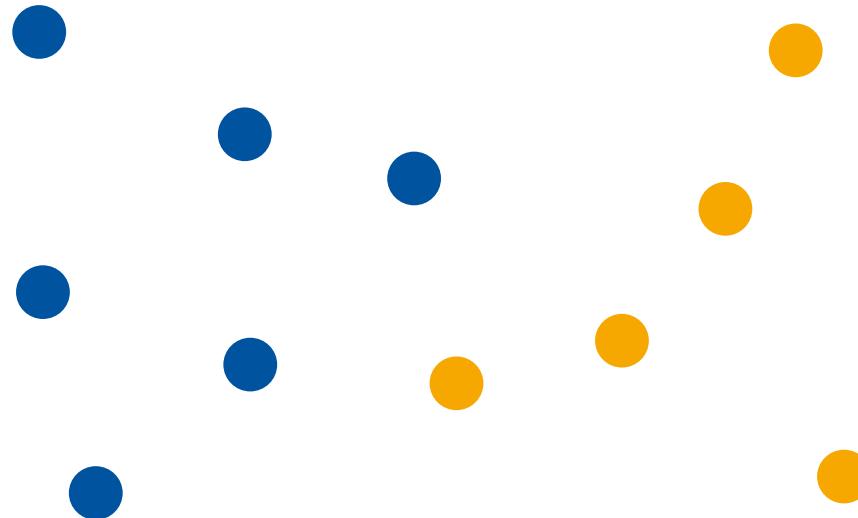


K-NN Classifier

- Combine K-NN density estimation with Bayes Decision Theory: [K-NN Classifier](#)
- Algorithm

Given a new sample x :

1. Find the K training samples with the smallest distance to x .
2. Assign the majority label of those samples to x .

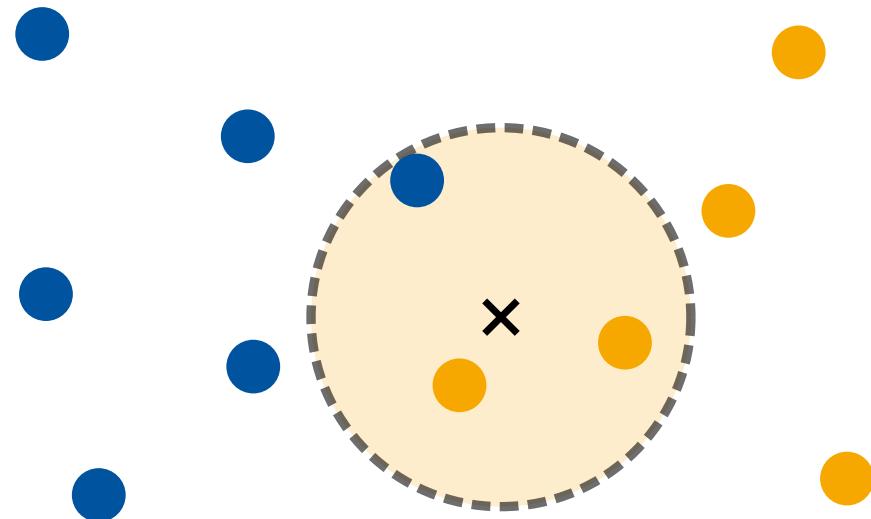


K-NN Classifier

- Combine K-NN density estimation with Bayes Decision Theory: [K-NN Classifier](#)
- Algorithm

Given a new sample \mathbf{x} :

1. Find the K training samples with the smallest distance to \mathbf{x} .
2. Assign the majority label of those samples to \mathbf{x} .



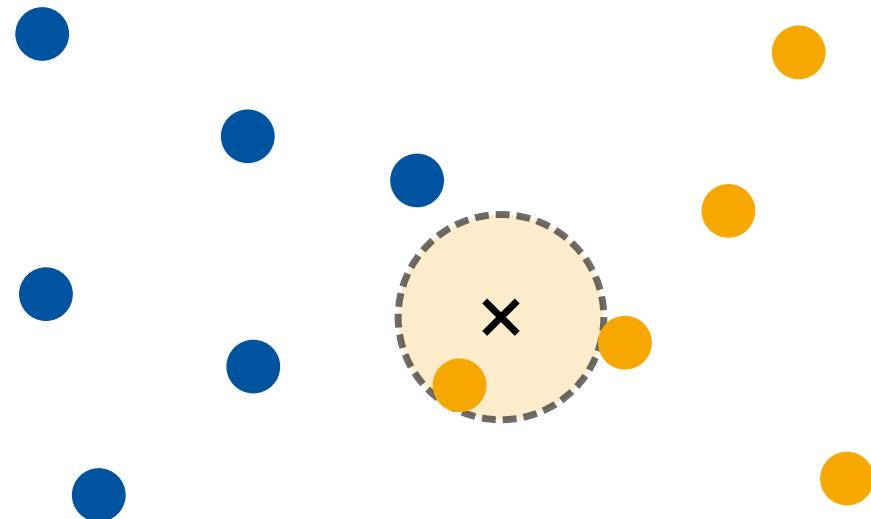
K-NN Classifier

- Combine K-NN density estimation with Bayes Decision Theory: [K-NN Classifier](#)
- Algorithm

Given a new sample \mathbf{x} :

1. Find the K training samples with the smallest distance to \mathbf{x} .
 2. Assign the majority label of those samples to \mathbf{x} .
- Special case: 1-NN Classifier.

Theoretical guarantee: Never worse than 2x the error of the optimal classifier!



Example

 $K = 1$  $K = 3$  $K = 15$

Discussion: K-NN Classifier

Advantages

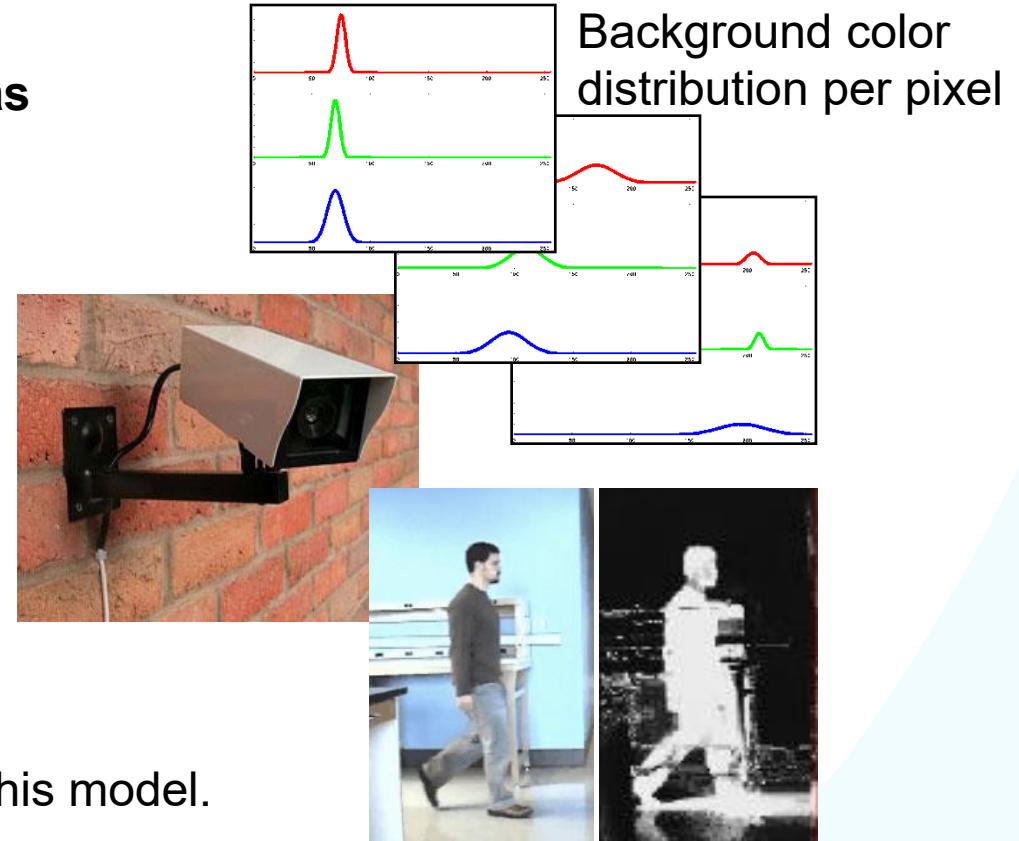
- Very simple, Bayes-optimal classifier.
- Needs no training.
- Can always be used as first estimate when working with a new dataset.
- Theoretical guarantees
 - Never worse than 2x optimal error

Limitations

- Requires storing the complete training set.
- Finding the k-nearest neighbors can become very expensive in high-dimensional spaces
- Theoretical optimality bound is often too loose to be of practical value.

Application Example: Background Models for Tracking

- **Example: Object tracking in static surveillance cameras**
 - Want to know if anybody enters a forbidden area
 - Challenge: many possible moving objects
- **Idea:** Train background color model for each pixel
 - Initialize with an empty scene.
 - Learn “common” appearance variation for each background pixel, e.g., by fitting a Gaussian distribution to the observed noise over several frames.
 - Evaluate the likelihood of observed pixel colors under this model.
⇒ *Anything that cannot be explained by the background model is labeled as foreground (=object).*



Application Example: Background Models for Tracking

- Problem: Outdoor scenes
 - Dynamic areas
 - Waving trees, rippling water, ...
- ⇒ *More flexible representation needed here!*
- Idea:
 - Use Kernel Density Estimation using the observed pixel values over a temporal window to model the “background” distribution for each pixel.
 - Again, evaluate the likelihood of the observed pixel color under this background model to detect “foreground” objects.

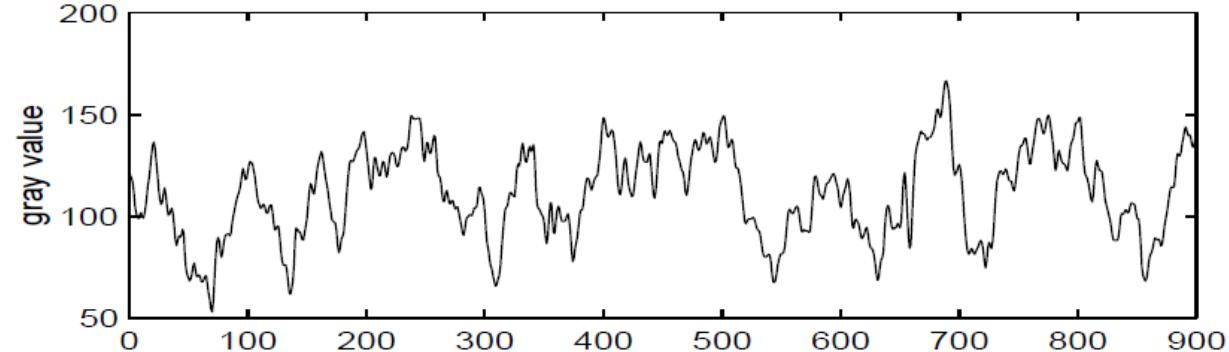


Image & Video source: A. Elgammal

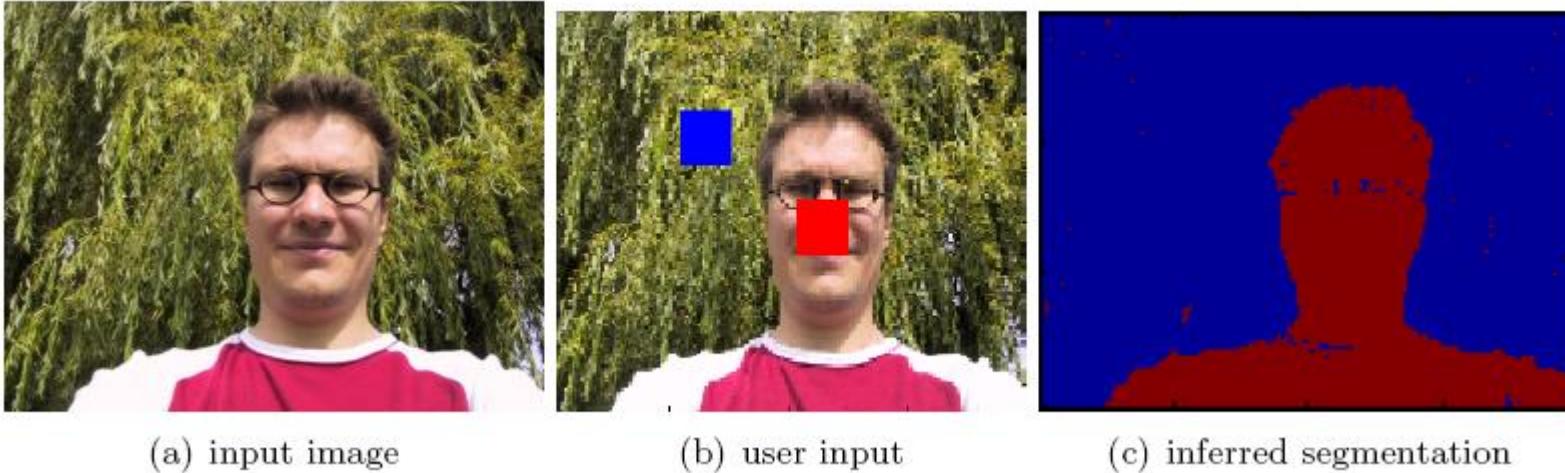
Application Example: Background Models for Tracking

- Results
 - Very robust foreground object detection in dynamic scenes
 - Automatic adaptation to varying weather conditions through temporal window



Video source: A. Elgammal

Application Example: Image Segmentation



- **Example: User assisted image segmentation**
 - User marks two regions for foreground and background.
 - Learn a MoG model for the color values in each region.
 - Use those models to classify all other pixels with a Bayes classifier (likelihood ratio test)
- ⇒ Simple, but effective segmentation procedure

References and Further Reading

- More information about EM and MoG estimation is available in Chapter 2.3.9 and the entire Chapter 9 of Bishop's book.

Christopher M. Bishop
Pattern Recognition and Machine Learning
Springer, 2006

