



Visual Computing  
Institute

RWTH AACHEN  
UNIVERSITY

# Elements of Machine Learning & Data Science

Winter semester 2025/26

## Lecture 11 – Linear Discriminants

01.12.2025

Prof. Bastian Leibe


# Announcements: Pre-recorded Videos

- **Pre-recorded videos available**
  - Extended explanations of key lecture topics
  - High production value
  - *Please use them as supplementary material*
- **For the next lecture**
  - Please watch the provided videos as preparation for the lecture
    - Motivation of Linear Regression
    - Least-Squares Regression
    - Regularization
    - Ridge Regression



# Machine Learning Topics


- 8. Introduction to ML
- 9. Probability Density Estimation
- 10. Linear Discriminants**
- 11. Linear Regression
- 12. Logistic Regression
- 13. Support Vector Machines
- 14. Neural Network Basics




Linear Discriminant Functions

$$E(\mathbf{w}) = \frac{1}{2} \sum_n (y(\mathbf{x}_n; \mathbf{w}) - t_n)^2$$

Least-Squares Classification




Activation Functions



Basis Functions

# Linear Discriminants

1. Motivation: Discriminant Functions
2. Linear Discriminant Functions
3. Least-Squares Classification
4. Generalized Linear Discriminants
5. Basis Functions



# Discriminant Functions


- Remember Bayes Decision Theory
  - Model the likelihood  $p(\mathbf{x}|\mathcal{C}_k)$  and the prior  $p(\mathcal{C}_k)$
  - From this, we can compute the posterior  $p(\mathcal{C}_k|\mathbf{x})$
  - Bayes optimal decision: Decide for class  $\mathcal{C}_1$  if

$$p(\mathcal{C}_1|\mathbf{x}) > p(\mathcal{C}_2|\mathbf{x}) \Leftrightarrow$$

$$\frac{p(\mathbf{x}|\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_2)} > \frac{p(\mathcal{C}_2)}{p(\mathcal{C}_1)}$$


*by comparing posteriors*

*by comparing likelihoods and priors*



# Discriminant Functions

- This chapter: Different approach
  - Directly represent the decision boundary with a **discriminant function**  $y(\mathbf{x})$
  - Without explicit modeling of probability densities!
    - *We will learn ways to define  $y(\mathbf{x})$  such that we still make a decision based on posteriors...*
    - *...but we don't have to. This framework gives us more flexibility.*



# Idea

- Formulate classification in terms of comparisons


- Bayes Decision Theory:

$$\begin{aligned} p(\mathcal{C}_1|\mathbf{x}) &> p(\mathcal{C}_2|\mathbf{x}) \\ \iff p(\mathcal{C}_1|\mathbf{x}) - p(\mathcal{C}_2|\mathbf{x}) &> 0 \\ \iff y(\mathbf{x}) &> 0 \end{aligned}$$

- More general:

- Define a **discriminant function**  $y(\mathbf{x})$
    - Classify  $\mathbf{x}$  as class  $\mathcal{C}_1$  if  $y(\mathbf{x}) > 0$

- Advantage: more flexibility



*E.g., we could now define*

$$y(\mathbf{x}) = p(\mathcal{C}_1|\mathbf{x}) - p(\mathcal{C}_2|\mathbf{x})$$

$$y(\mathbf{x}) = \ln \frac{p(\mathbf{x}|\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_2)} + \ln \frac{p(\mathcal{C}_1)}{p(\mathcal{C}_2)}$$

# Idea

- **Multi-class case**
  - Define **discriminant functions**
$$y_1(\mathbf{x}), \dots, y_K(\mathbf{x})$$
  - Classify  $\mathbf{x}$  as class  $\mathcal{C}_k$  if
$$y_k(\mathbf{x}) > y_j(\mathbf{x}) \quad \forall j \neq k$$
- Again, this is compatible with Bayes Decision Theory:

- E.g.,  $y_k(\mathbf{x}) = p(\mathcal{C}_k | \mathbf{x})$

$$y_k(\mathbf{x}) = p(\mathbf{x} | \mathcal{C}_k) p(\mathcal{C}_k)$$


$$y_k(\mathbf{x}) = \log p(\mathbf{x} | \mathcal{C}_k) + \log p(\mathcal{C}_k)$$

# Problem Formulation

- **General classification problem**
  - Goal: take a new input  $\mathbf{x}$  and assign it to one of  $K$  classes  $\mathcal{C}_k$
  - Given training set  $\mathcal{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$  with target values  $\mathcal{T} = \{\mathbf{t}_1, \dots, \mathbf{t}_N\}$   
 $\Rightarrow$  Learn a discriminant function  $y(\mathbf{x})$  to perform the classification
- **2-class problem**
  - **Binary target values**  $t_n \in \{-1, 1\}$  or  $t_n \in \{0, 1\}$
- **$K$ -class problem**
  - **1-of- $K$  coding scheme**, e.g.  $\mathbf{t}_n = (0, 1, 0, 0, 0)^T$

# Linear Discriminants

1. Motivation: Discriminant Functions
2. **Linear Discriminant Functions**
3. Least-Squares Classification
4. Generalized Linear Discriminants
5. Basis Functions



# Linear Discriminant Functions

- 2-class problem
  - $y(\mathbf{x}) > 0$ : Decide for class  $\mathcal{C}_1$ , else for  $\mathcal{C}_2$
- In the following, we focus on [linear discriminant functions](#):

$$y(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + w_0$$

weight vector


“bias”  
(= threshold)

- If a dataset can be perfectly classified by a linear discriminant, we call it [linearly separable](#).

# Intuition

$$y(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + w_0$$


- Graphical interpretation:
  - Normal equation of a hyperplane
  - Normal vector:  $\mathbf{w}$
  - Offset:  $\frac{-w_0}{\|\mathbf{w}\|}$



# Intuition

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

- Graphical interpretation:
  - Normal equation of a hyperplane
- The hyperplane is given by  $y(\mathbf{x}) = 0$ 
  - One side:  $y(\mathbf{x}) > 0$
  - Other side:  $y(\mathbf{x}) < 0$
- This hyperplane defines the decision boundary.



# Notation

- Let's look at the equation in detail

$$\begin{aligned}y(\mathbf{x}) &= \mathbf{w}^T \mathbf{x} + w_0 \\&= \sum_{i=1}^D w_i x_i + w_0\end{aligned}$$

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_D \end{bmatrix} \quad \mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_D \end{bmatrix}$$

- Alternative notation with extended vector


$$\begin{aligned}y(\mathbf{x}) &= \sum_{i=0}^D w_i x_i \quad \text{with } x_0 = 1 \\&= \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}\end{aligned}$$

$$\tilde{\mathbf{x}} = \begin{bmatrix} 1 \\ x_1 \\ x_2 \\ \vdots \\ x_D \end{bmatrix} \quad \tilde{\mathbf{w}} = \begin{bmatrix} w_0 \\ w_1 \\ w_2 \\ \vdots \\ w_D \end{bmatrix}$$

$D$  : Number of Dimensions

# Extension to Multiple Classes

- **K-class discriminant**
  - Combination of  $K$  linear functions
$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0}, \quad k = 1, \dots, K$$
- Interpretation
  - Decide for  $\mathcal{C}_k$  iff  $y_k(\mathbf{x}) > y_j(\mathbf{x}) \quad \forall j \neq k$
  - Resulting decision hyperplanes:
$$(\mathbf{w}_k - \mathbf{w}_j)^T \mathbf{x} + (w_{k0} - w_{j0}) = 0$$



# Extension to Multiple Classes

- **$K$ -class discriminant**

- Combination of  $K$  linear functions

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0}, \quad k = 1, \dots, K$$

- Interpretation


- Decide for  $\mathcal{C}_k$  iff  $y_k(\mathbf{x}) > y_j(\mathbf{x}) \quad \forall j \neq k$

- Resulting decision hyperplanes:

$$(\mathbf{w}_k - \mathbf{w}_j)^T \mathbf{x} + (w_{k0} - w_{j0}) = 0$$


- It can be shown that the decision regions of such a discriminant are always **singly connected** and **convex**.

⇒ Particularly suitable for problems with unimodal conditional densities  $p(\mathbf{x}|\mathbf{w}_i)$



# Linear Discriminants

1. Motivation: Discriminant Functions
2. Linear Discriminant Functions
3. **Least-Squares Classification**
4. Generalized Linear Discriminants
5. Basis Functions



# General Classification Problem

- $K$  Classes described by linear discriminant models:

$$y_k(\mathbf{x}) = \mathbf{w}_k^\top \mathbf{x} + w_{k0}, \quad k = 1, \dots, K$$

- Group them together using vector notation:

$$\mathbf{y}(\mathbf{x}) = \widetilde{\mathbf{W}}^\top \widetilde{\mathbf{x}} = \begin{bmatrix} y_1(\mathbf{x}) \\ y_2(\mathbf{x}) \\ \vdots \\ y_K(\mathbf{x}) \end{bmatrix} \quad \text{where}$$

$$\widetilde{\mathbf{W}} = [\widetilde{\mathbf{w}}_1, \dots, \widetilde{\mathbf{w}}_K] = \begin{bmatrix} w_{10} & \cdots & w_{K0} \\ \vdots & \ddots & \vdots \\ w_{1D} & \cdots & w_{KD} \end{bmatrix}$$
$$\widetilde{\mathbf{x}} = [1, x_1, \dots, x_D]^\top$$

- The output will be in **1-of- $K$**  notation.  
⇒ We can directly compare it to the target value

$$\mathbf{t} = [t_1, \dots, t_k]^\top$$

# Classification Problem for the Entire Dataset

- Write the classification output for the whole dataset:

$$\mathbf{Y}(\tilde{\mathbf{X}}) = \tilde{\mathbf{X}}\tilde{\mathbf{W}}$$

where

$$\tilde{\mathbf{W}} = [\tilde{\mathbf{w}}_1, \dots, \tilde{\mathbf{w}}_K] = \begin{bmatrix} w_{01} & \cdots & w_{0K} \\ \vdots & \ddots & \vdots \\ w_{D1} & \cdots & w_{DK} \end{bmatrix}$$

- Using the data matrix

$$\tilde{\mathbf{X}} = \begin{bmatrix} \tilde{\mathbf{x}}_1^T \\ \vdots \\ \tilde{\mathbf{x}}_N^T \end{bmatrix} = \begin{bmatrix} x_{10} & \cdots & x_{1D} \\ \vdots & \ddots & \vdots \\ x_{N0} & \cdots & x_{ND} \end{bmatrix}$$

- Similarly, we group the target vectors in a matrix

$$\mathbf{T} = \begin{bmatrix} \mathbf{t}_1^T \\ \vdots \\ \mathbf{t}_N^T \end{bmatrix} = \begin{bmatrix} t_{11} & \cdots & t_{1K} \\ \vdots & \ddots & \vdots \\ t_{N1} & \cdots & t_{NK} \end{bmatrix}$$

- We can now compare the classifier output with the target matrix:  $\tilde{\mathbf{X}}\tilde{\mathbf{W}} - \mathbf{T}$

# Defining the Classification Problem

- Comparing the classifier output with the target matrix:

$$\widetilde{\mathbf{X}} \widetilde{\mathbf{W}} - \mathbf{T}$$

- Goal: Choose  $\widetilde{\mathbf{W}}$  such that this difference becomes minimal
  - What does *minimal* mean here?
  - How strongly do we want to penalize deviations from the ideal target value?*
- Idea: define an **error function** that specifies the **loss** for each deviation

$$E(\widetilde{\mathbf{W}}) = \sum_{n=1}^N \sum_{k=1}^K L(y_k(\mathbf{x}_n; \mathbf{w}_k), t_{nk})$$

# Least-Squares Classification

- Simplest approach: minimize Sum-of-squares error

$$\begin{aligned} E(\mathbf{W}) &= \frac{1}{2} \sum_{n=1}^N \sum_{k=1}^K (y_k(\mathbf{x}_n; \mathbf{w}_k) - t_{nk})^2 \\ &= \frac{1}{2} \sum_{n=1}^N \sum_{k=1}^K (\mathbf{w}_k^\top \mathbf{x}_n - t_{nk})^2 \end{aligned}$$

*Simplified notation:  
Leaving out the  $\tilde{\mathbf{x}}$  ...*

- How do we minimize this function?
  - *Take the derivative and set it to zero...*

# Derivation

- Let's concentrate on the two-class case first:

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (\mathbf{w}^\top \mathbf{x}_n - t_n)^2 \quad t_n \in \{-1, 1\}$$

- Taking the derivative:

$$\begin{aligned} \frac{\partial E(\mathbf{w})}{\partial \mathbf{w}} &= \sum_{n=1}^N (\mathbf{w}^\top \mathbf{x}_n - t_n) \mathbf{x}_n \\ &= \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{t}) \end{aligned}$$

with  $\mathbf{X} = \begin{pmatrix} \mathbf{x}_1^\top \\ \vdots \\ \mathbf{x}_N^\top \end{pmatrix} \quad \mathbf{t} = (t_1, \dots, t_N)^\top$

Linear Algebra textbook:

$$\boxed{\frac{\partial \mathbf{a}^\top \mathbf{b}}{\partial \mathbf{a}} = \mathbf{b}}$$

- Setting the derivative to zero:

$$\frac{\partial E(\mathbf{w})}{\partial \mathbf{w}} = \mathbf{X}^T (\mathbf{X}\mathbf{w} - \mathbf{t}) \stackrel{!}{=} 0 \quad \text{"pseudo-inverse"}$$
$$\mathbf{X}^T \mathbf{X} \mathbf{w} = \mathbf{X}^T \mathbf{t}$$
$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{t}$$
$$= \mathbf{X}^\dagger \mathbf{t}$$

- We then obtain the discriminant function as


$$y(\mathbf{x}; \mathbf{w}) = \mathbf{w}^T \mathbf{x} = \mathbf{t}^T (\mathbf{X}^\dagger)^T \mathbf{x}$$

*Exact, closed-form solution!*


## Example: Two Classes




## Example: Two Classes



## Example: Two Classes




## Example: Two Classes



# Why Does This Happen?


$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

- Remember the interpretation of  $y(\mathbf{x})$ 
  - Normal equation of a hyperplane
- $y(\mathbf{x})$  measures the (signed) distance of the point  $\mathbf{x}$  from the hyperplane.
- However, we now compare it to a fixed target value of  $t_n \in \{-1, 1\}$  ...



# Linear Discriminants

1. Motivation: Discriminant Functions
2. Linear Discriminant Functions
3. Least-Squares Classification
4. **Generalized Linear Discriminants**
5. Basis Functions



# Generalized Linear Models

- So far: model classification by linear discriminant function

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

- Generalize this with an activation function  $g(\cdot)$ :


$$y(\mathbf{x}) = g(\mathbf{w}^T \mathbf{x} + w_0)$$

- Remarks
  - $g(\cdot)$  may be non-linear.
  - Decision surfaces correspond to

$$y(\mathbf{x}) = \text{const} \iff \mathbf{w}^T \mathbf{x} + w_0 = \text{const}$$

$\Rightarrow$  If  $g(\cdot)$  is monotonous (which is typically the case),  
the decision boundaries are still linear functions of  $\mathbf{x}$ .

Generalized Linear Model



# Activation Functions

- Recall least-squares classification:
  - Outliers have strong influence


$$E(\mathbf{w}) = \sum_{n=1}^N (y(\mathbf{x}_n; \mathbf{w}) - t_n)^2$$

- This is because the output  $y(\mathbf{x}; \mathbf{w})$  can grow arbitrarily large:

$$y(\mathbf{x}; \mathbf{w}) = \mathbf{w}^\top \mathbf{x} + w_0$$

- Choosing a suitable nonlinearity can limit those influences:

$$y(\mathbf{x}; \mathbf{w}) = g(\mathbf{w}^\top \mathbf{x} + w_0)$$




# The Logistic Sigmoid

- To limit the influence of outliers, we can use the **logistic sigmoid** function:


$$\sigma(x) = \frac{1}{1 + \exp(-x)}$$

- For 2-class problems, we scale it to the output range (-1,1) (known as **tangens hyperbolicus**):


$$\tanh(x) = 2\sigma(x) - 1$$




# Example




# Example




## Example



## Example



## Example



# Discussion: Activation Functions

## Advantages


- Nonlinearity gives more flexibility.
- Can be used to limit the effect of outliers.
- Choice of sigmoid actually has a nice probabilistic interpretation.

## Limitations

- Least-squares minimization in general no longer leads to a closed-form analytical solution.  
⇒ Need to apply iterative methods.

# Linear Discriminants

1. Motivation: Discriminant Functions
2. Linear Discriminant Functions
3. Least-Squares Classification
4. Generalized Linear Discriminants
5. **Basis Functions**




# Basis Functions

- So far: assumed linear separability
  - Very restrictive assumption, classical counterexample: XOR
  - We need non-linear decision boundaries...
- Solution: use non-linear **basis functions**  $\phi_j(\mathbf{x})$ :

$$y(\mathbf{x}) = \sum_{j=1}^M w_j \phi_j(\mathbf{x}) + w_0$$

- By choosing the right  $\phi$ , every continuous function can (in principle) be approximated with arbitrary accuracy



# Intuition


$$y_k(\mathbf{x}) = \sum_{j=0}^M w_{kj} \phi_j(\mathbf{x}) = \mathbf{w}_k^\top \phi(\mathbf{x})$$

This is still a **linear problem** in  $\phi(\mathbf{x})$ .


$\phi_j(\mathbf{x})$  are called **basis functions**.

But, depending on  $\phi(\cdot)$ , it may now be a nonlinear problem in  $\mathbf{x}$ .


Typically,  $\phi_0(\mathbf{x}) = 1$  so that  $w_0$  acts as a bias.



$$\phi : \mathbb{R}^D \mapsto \mathbb{R}^M$$




Usually,  $\phi$  maps into a higher-dimensional space.



Not linearly separable

$$\phi : \mathbb{R}^D \rightarrow \mathbb{R}^M$$

$M \gg D$



Linearly separable

## Example: Polynomial Basis Functions

- Polynomial basis functions map  $x$  to powers of  $x$ :


$$\phi(x) = (x^m, x^{m-1}, \dots, x, 1)^T$$


- When we optimize  $\mathbf{w}^T \phi(x)$  with polynomial basis functions, we implicitly optimize the coefficients of a polynomial in  $x$ :

$$\begin{aligned}y(x) &= \mathbf{w}^T \phi(x) \\&= w_m x^m + w_{m-1} x^{m-1} + \dots + w_1 x + w_0\end{aligned}$$


- As before, we decide for  $\mathcal{C}_1$  if  $y(x) > 0$ .

Let's use a third-degree polynomial:  $\phi(x) = (x^3, x^2, x, 1)^\top$






We can also extend this for higher dimensional inputs, e.g., 2D:



We can also extend this for higher dimensional inputs, e.g., 2D:



# Discussion: Basis Functions

## Advantages

- Basis functions allow us to address linearly non-separable problems
- The problem is still linear in  $\phi(\mathbf{x})$  (but may be nonlinear in  $\mathbf{x}$ ).
- We can think of  $\phi(\mathbf{x})$  as transforming the data into a feature space in which the problem is easier to solve.
- In general, it is easier to find a separating hyperplane in higher-dimensional spaces.


## Limitations

- The right choice of  $\phi(\mathbf{x})$  depends on the problem and is another hyperparameter to optimize.
- Flexibility is limited by the curse of dimensionality. Evaluating  $\mathbf{w}^T \phi(\mathbf{x})$  can be expensive in high-dimensional spaces.
- Choosing a higher-dimensional feature space  $\phi(\mathbf{x})$  increases the capacity of the classifier and may lead to overfitting.

## References and Further Reading

- More information about Linear Discriminants is available in Chapter 4.1 of Bishop's book.

Christopher M. Bishop  
Pattern Recognition and Machine Learning  
Springer, 2006



# Announcements: Pre-recorded Videos

- **Pre-recorded videos available**
  - Extended explanations of key lecture topics
  - High production value
  - *Please use them as supplementary material*
- **For the next lecture**
  - Please watch the provided videos as preparation for the lecture
    - Motivation of Linear Regression
    - Least-Squares Regression
    - Regularization
    - Ridge Regression

