



Visual Computing
Institute

RWTH AACHEN
UNIVERSITY

Elements of Machine Learning & Data Science

Winter semester 2025/26

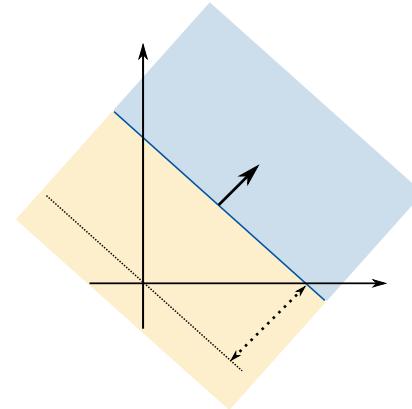
Lecture 12 – Linear Regression

02.12.2025

Prof. Bastian Leibe

Machine Learning Topics

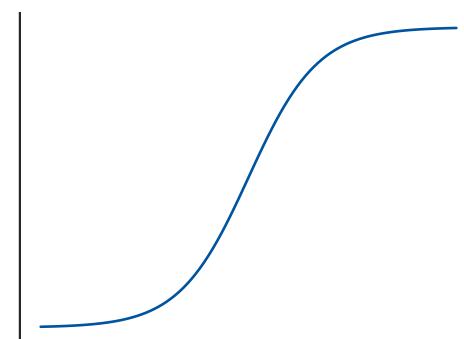
- 8. Introduction to ML
- 9. Probability Density Estimation
- 10. Linear Discriminants**
- 11. Linear Regression
- 12. Logistic Regression
- 13. Support Vector Machines
- 14. Neural Network Basics



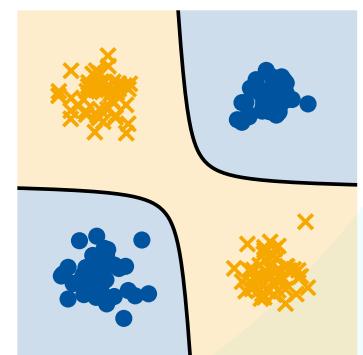
Linear Discriminant Functions

$$E(\mathbf{w}) = \frac{1}{2} \sum_n (y(\mathbf{x}_n; \mathbf{w}) - t_n)^2$$

Least-Squares Classification



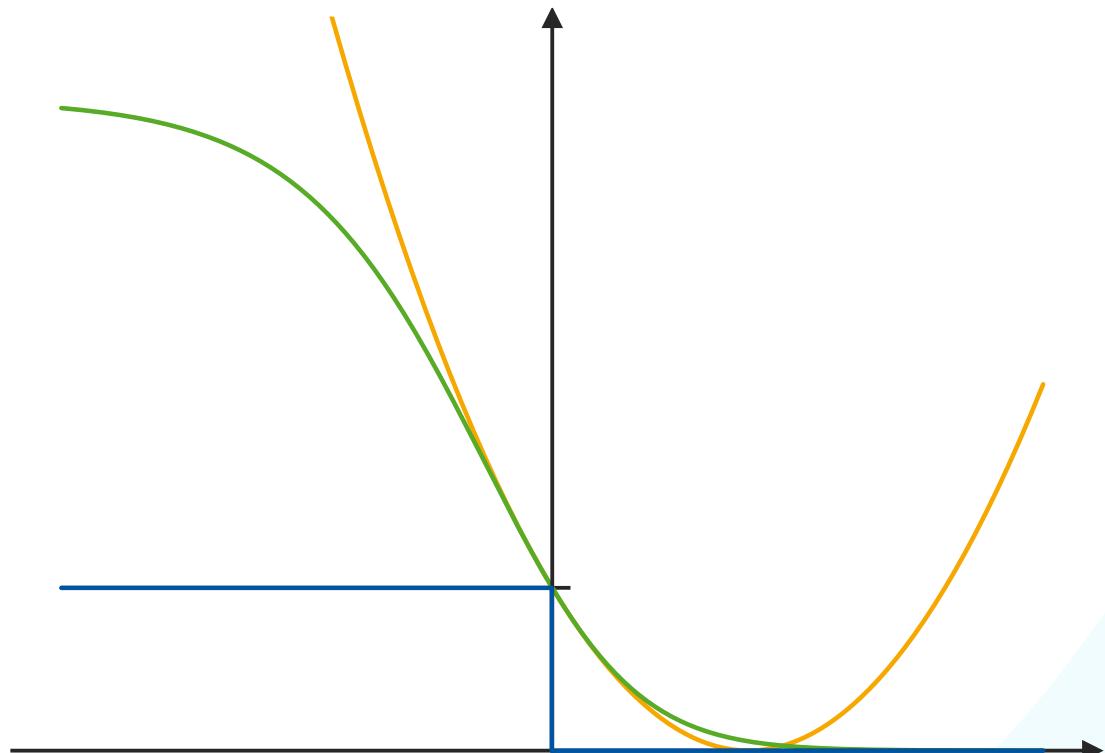
Activation Functions



Basis Functions

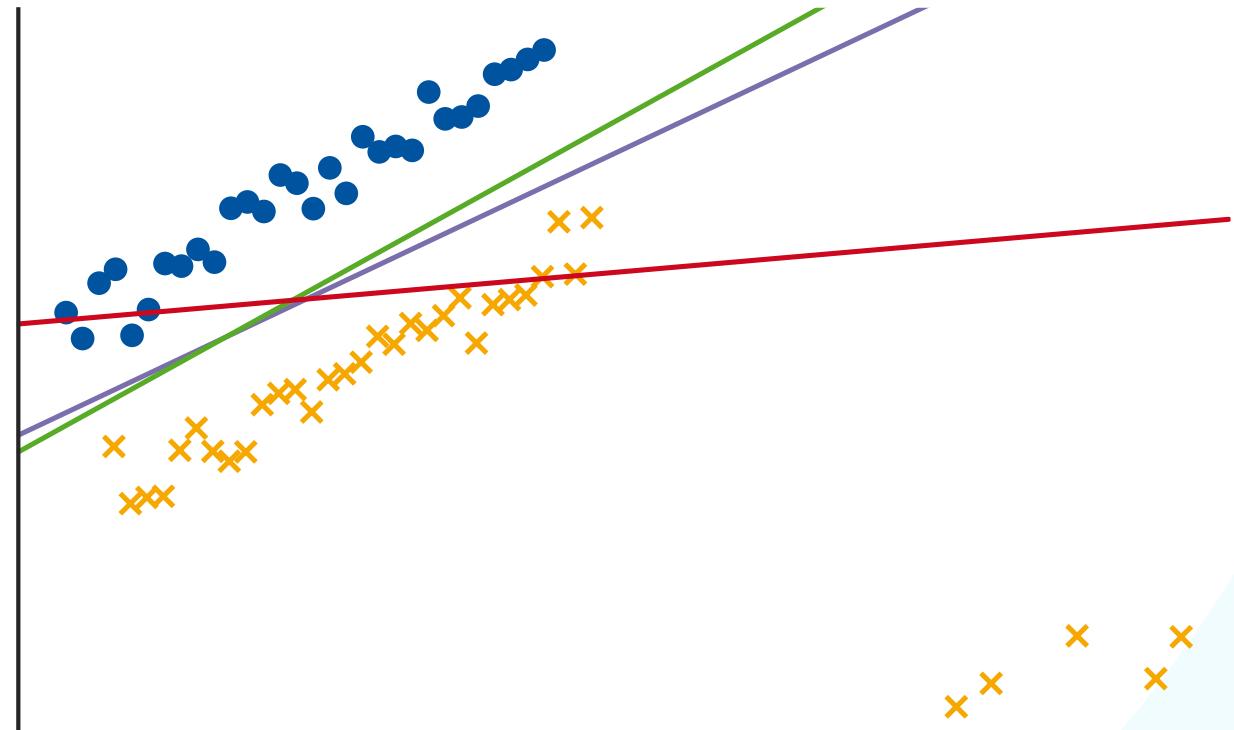
Linear Discriminants

1. Motivation: Discriminant Functions
2. Linear Discriminant Functions
3. Least-Squares Classification
4. Generalized Linear Discriminants
5. Basis Functions
6. **Error Function Analysis**

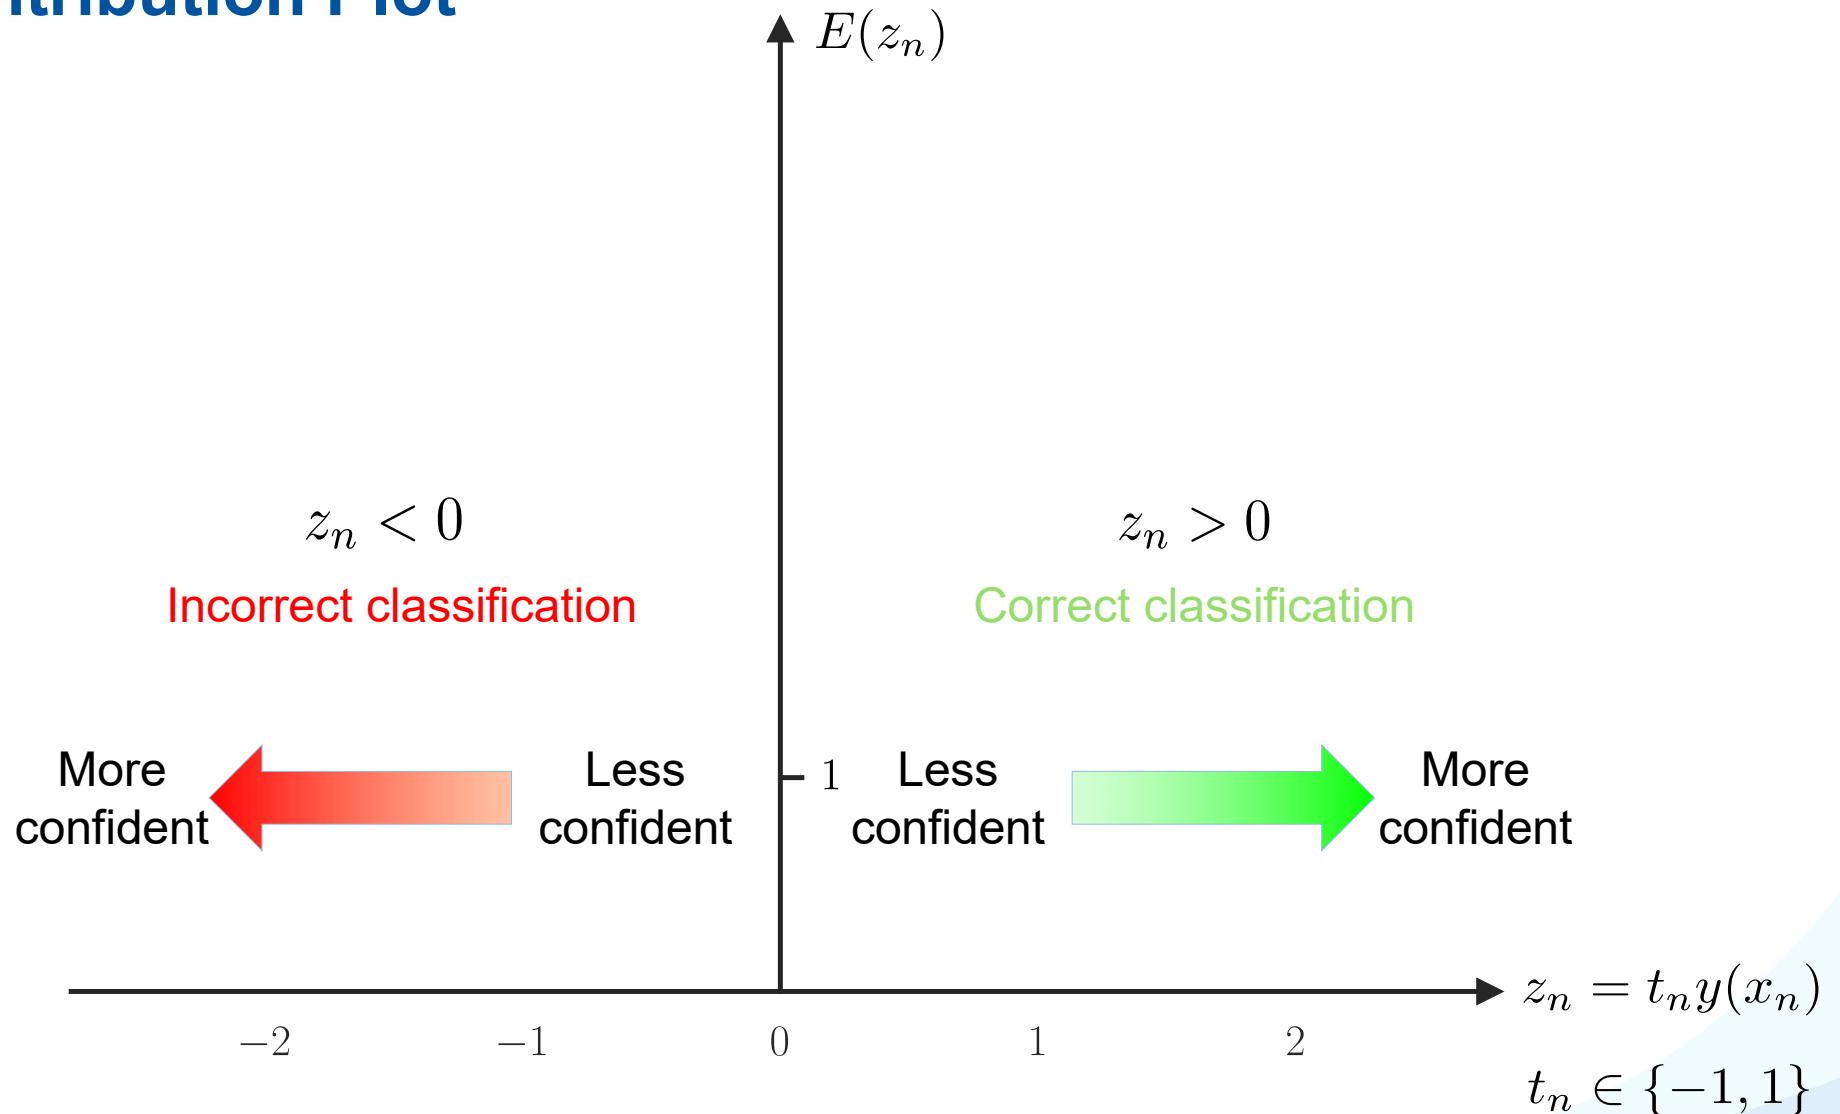


Error Function Analysis

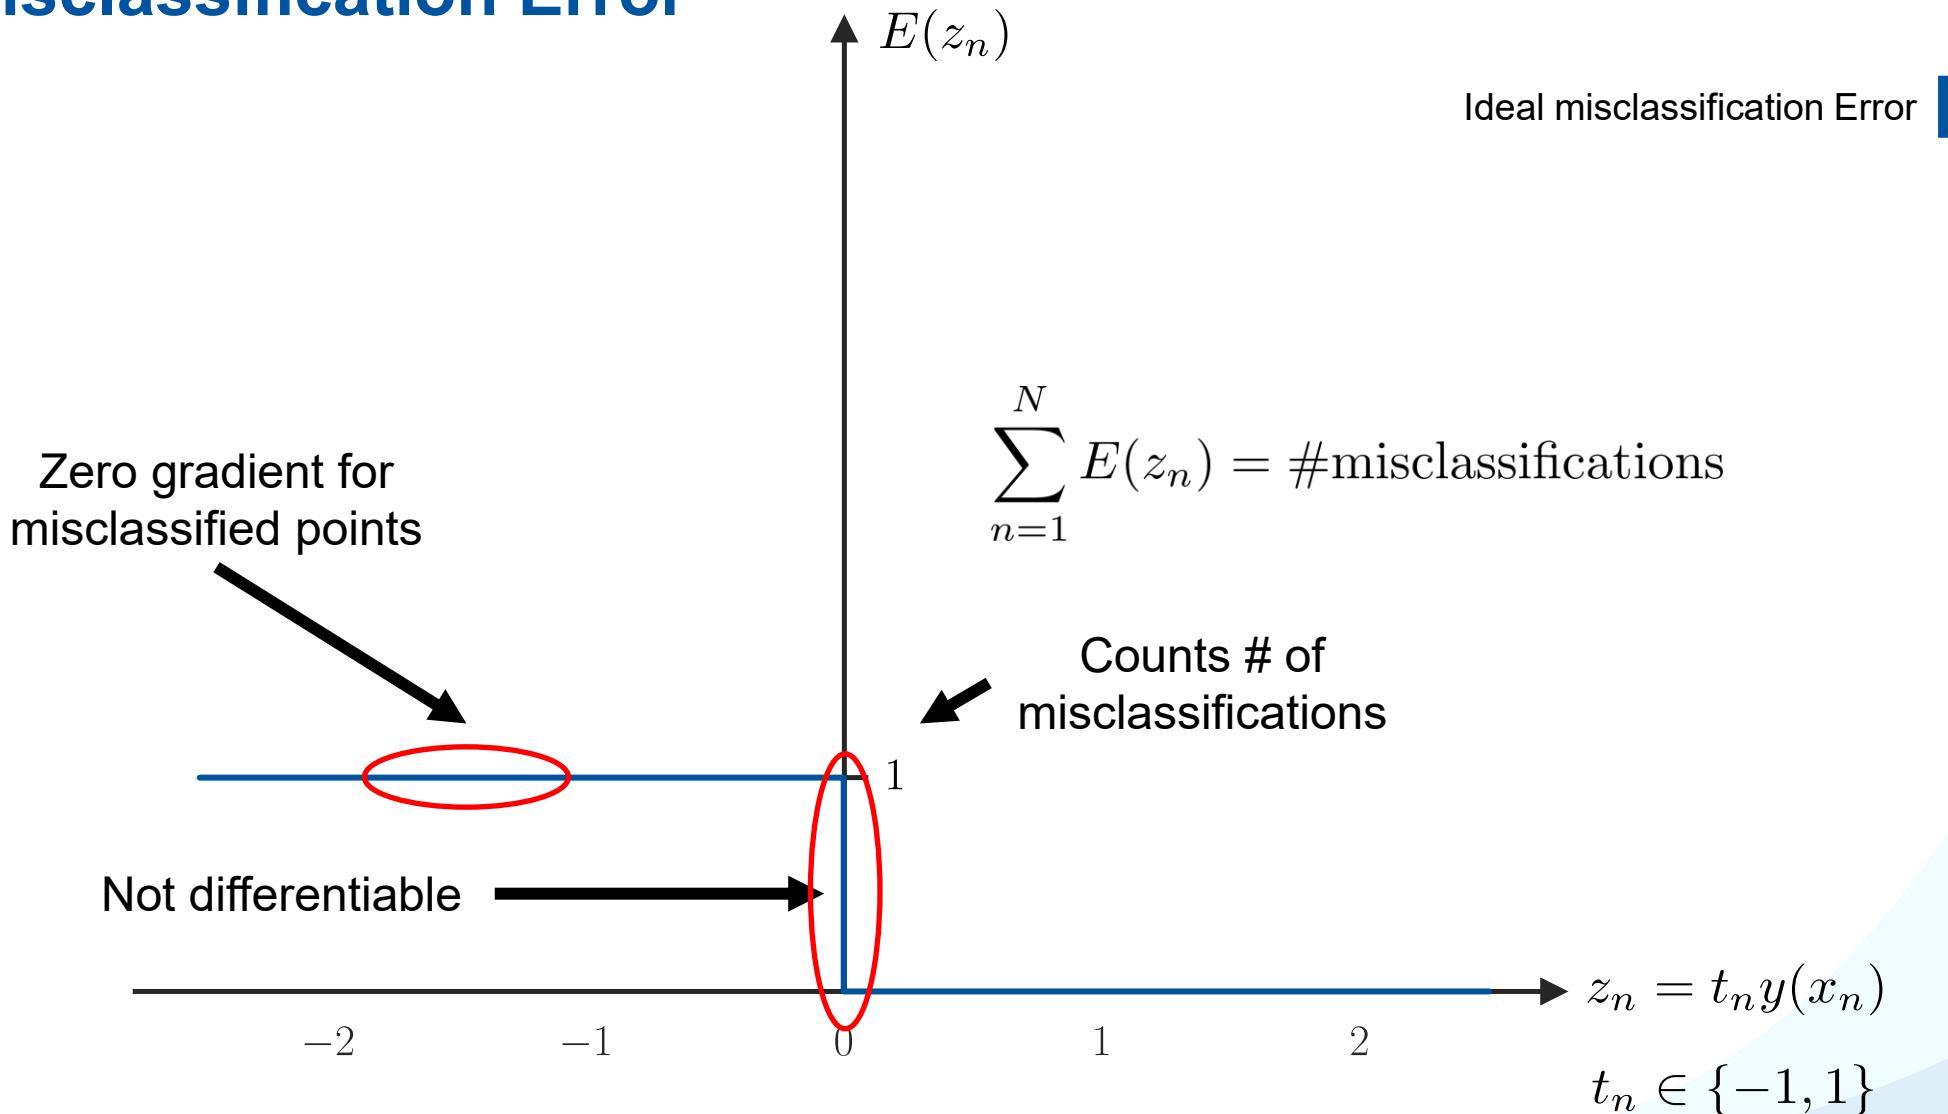
- We have seen how to learn **generalized linear discriminant** models by optimizing an error function.
 - We observed problems with **Least-Squares Classification** based on the squared error function.
 - In particular, sensitivity to outliers
 - Can we predict when such problems will occur?
- *Let's analyze the behavior of error functions in more detail...*



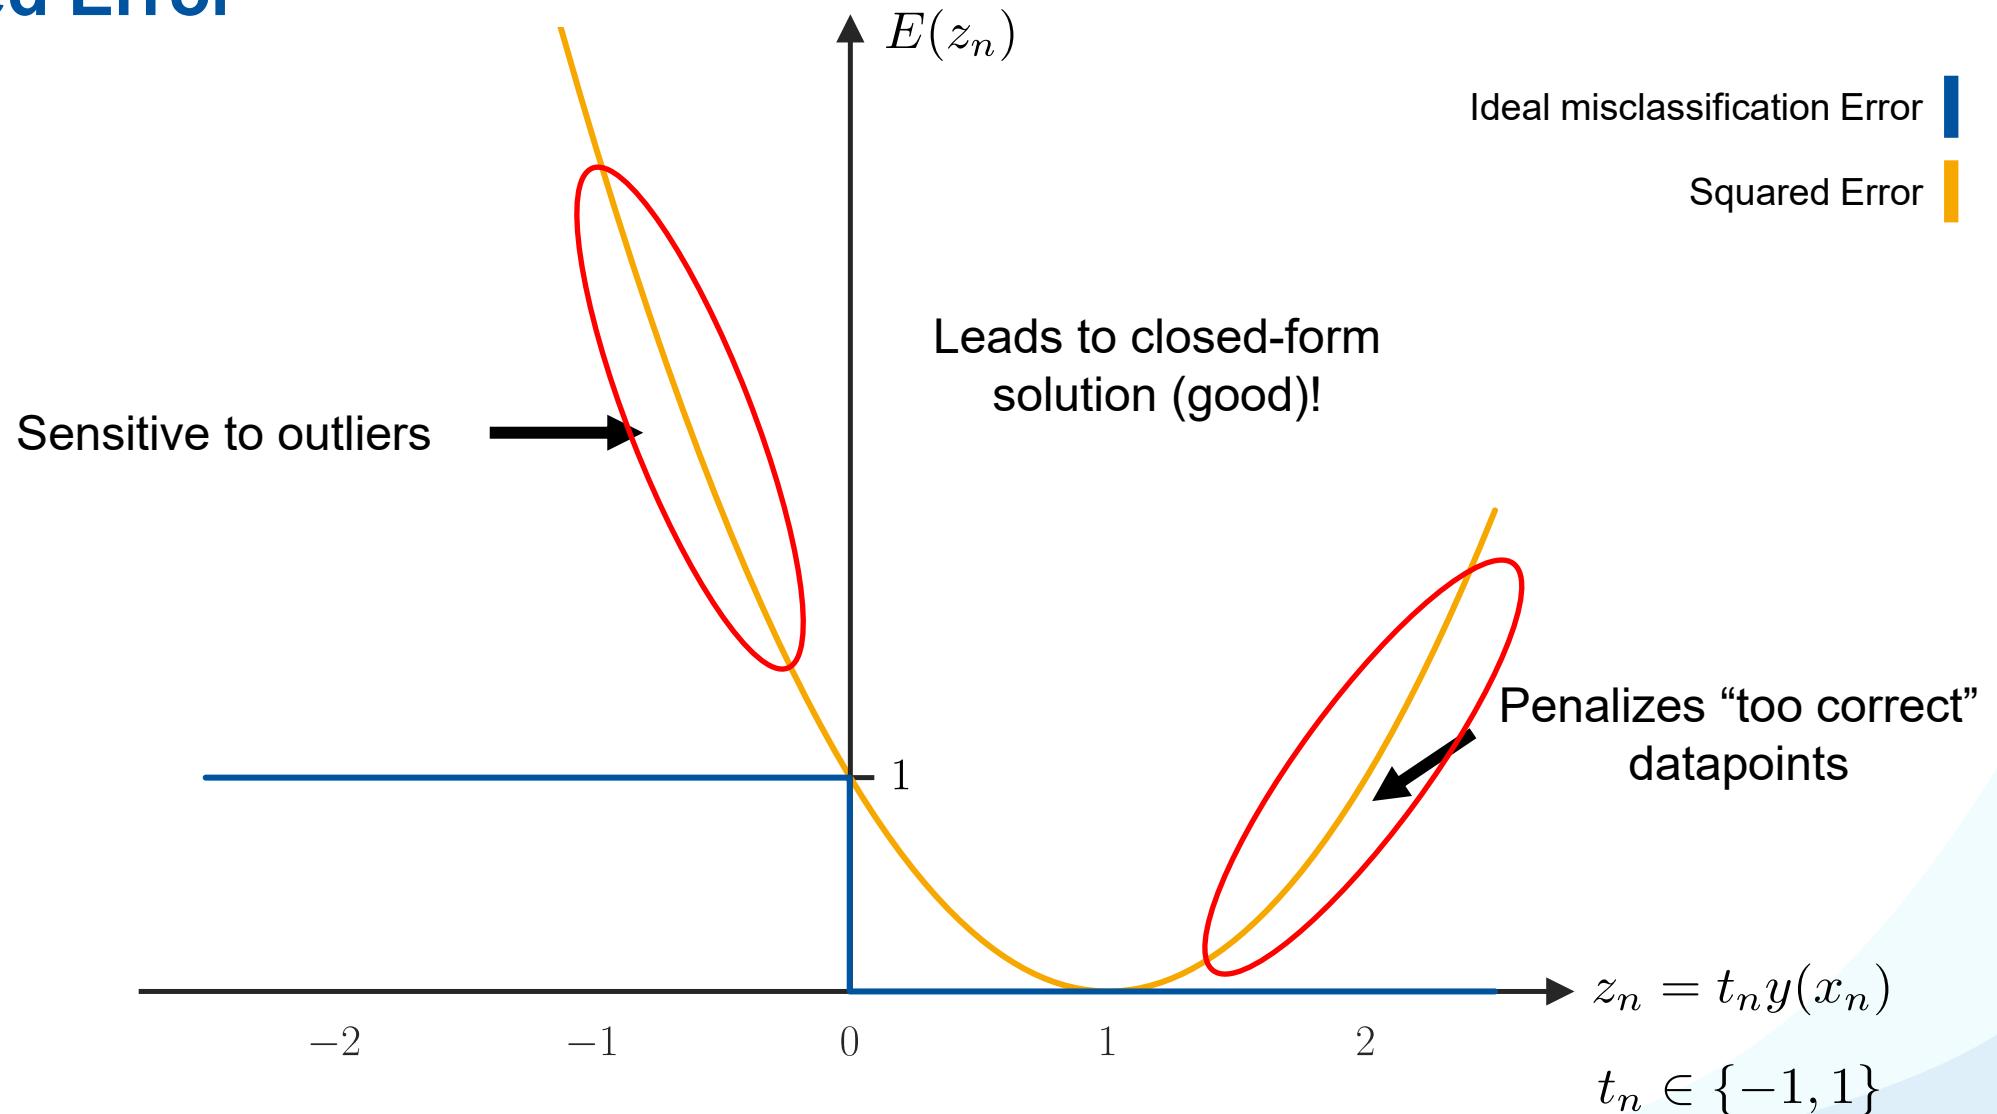
Error Contribution Plot



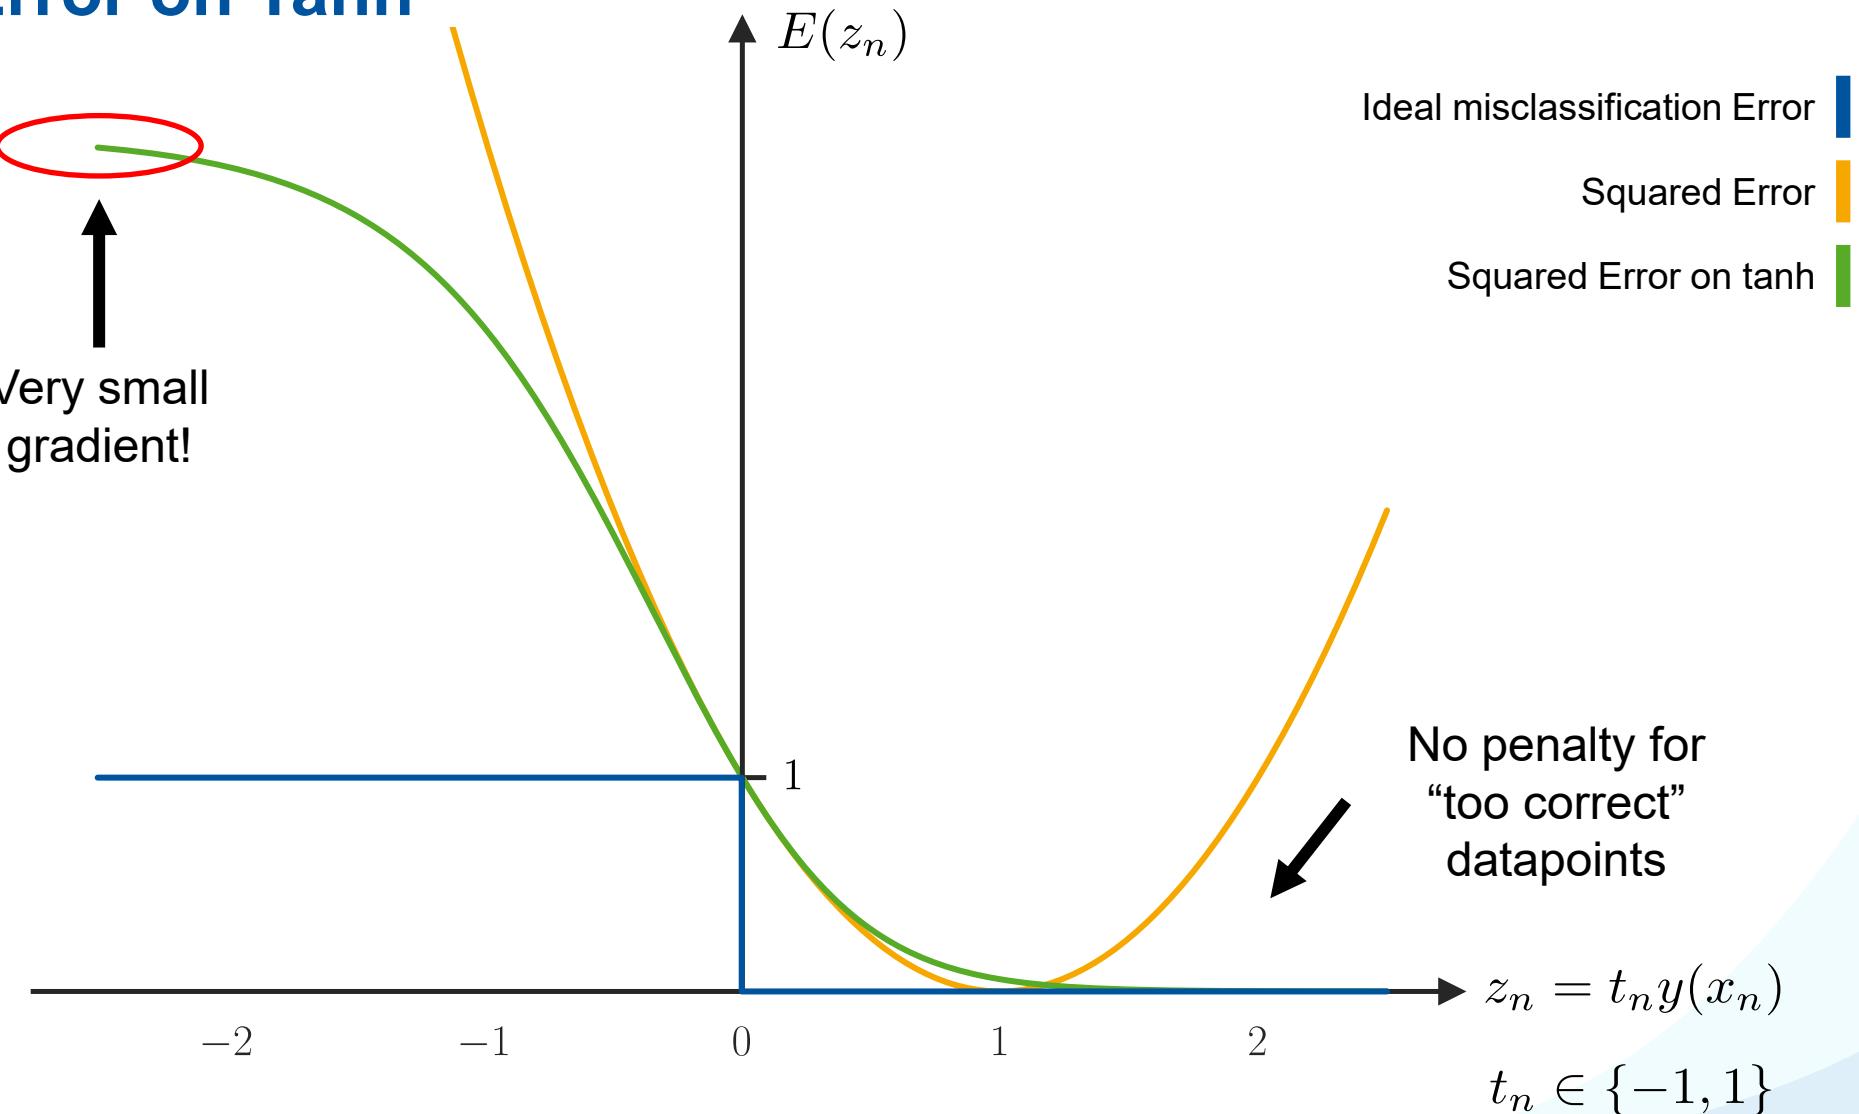
Ideal Misclassification Error



Squared Error

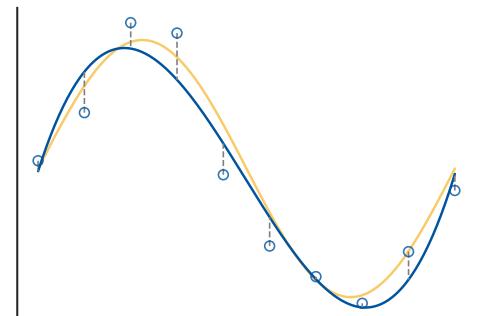


Squared Error on Tanh

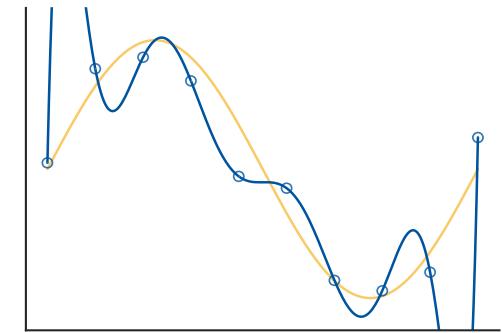


Machine Learning Topics

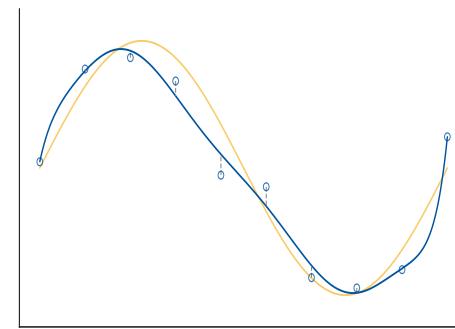
- 8. Introduction to ML
- 9. Probability Density Estimation
- 10. Linear Discriminants
- 11. Linear Regression**
- 12. Logistic Regression
- 13. Support Vector Machines
- 14. Neural Network Basics



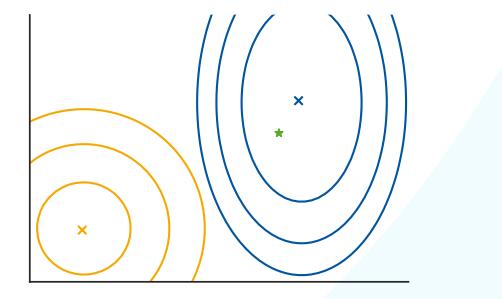
$y(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}) + w_0$
Linear Regression Functions



Overfitting



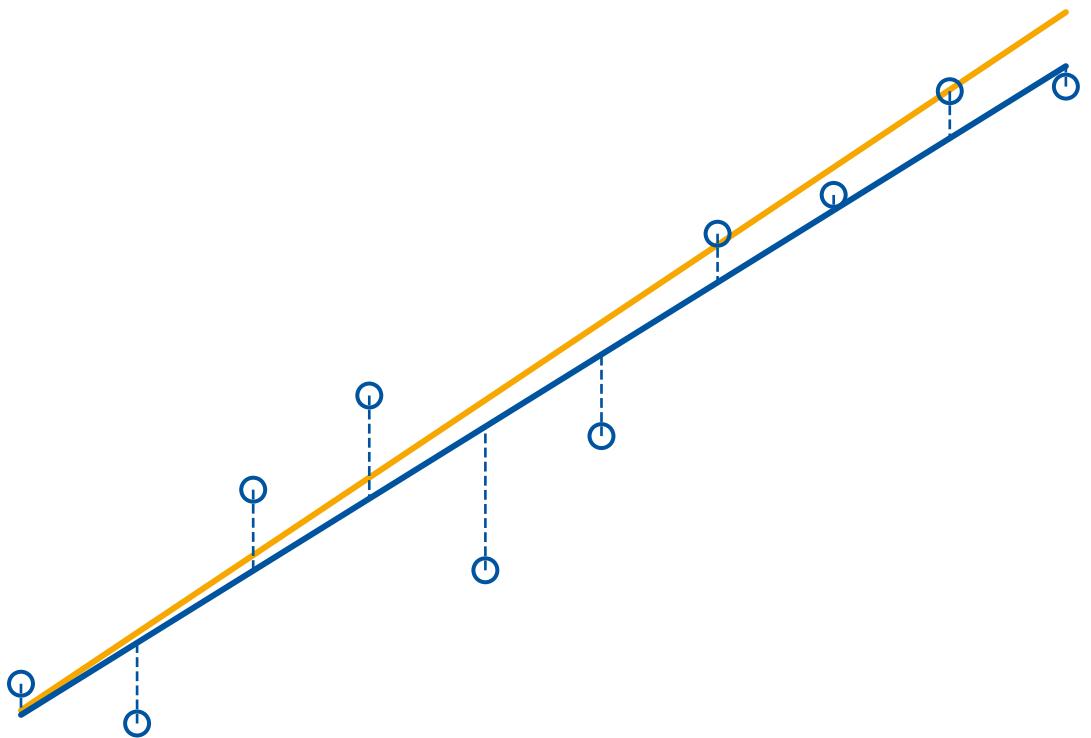
$E(\mathbf{w}) = L(\mathbf{w}) + \lambda \Omega(\mathbf{w})$
Regularization



$\mathbf{w} = (\Phi^T \Phi + \lambda \mathbf{I})^{-1} \Phi^T \mathbf{t}$
Ridge Regression

Linear Regression

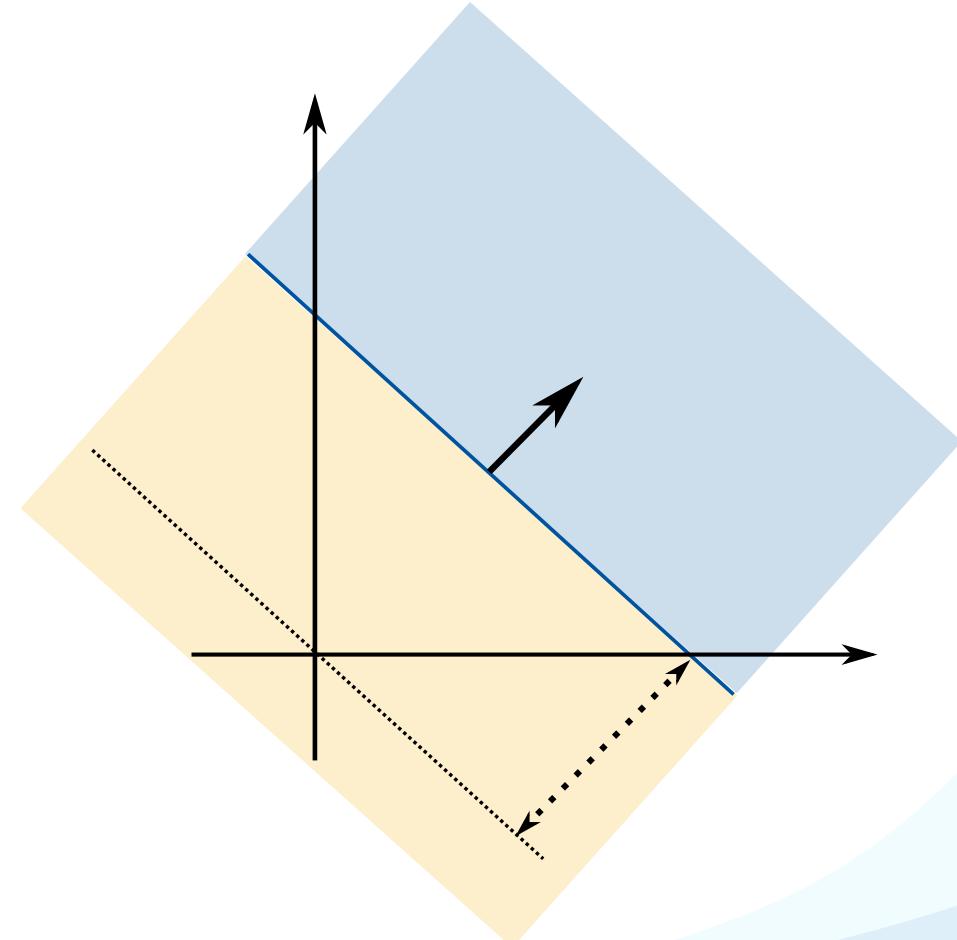
1. Motivation
2. Least-Squares Regression
3. Regularization
4. Ridge Regression
5. The Bias-Variance Tradeoff



Motivation: Linear Regression

- We have seen how to build classifiers with linear functions:

$$y(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}) + w_0$$

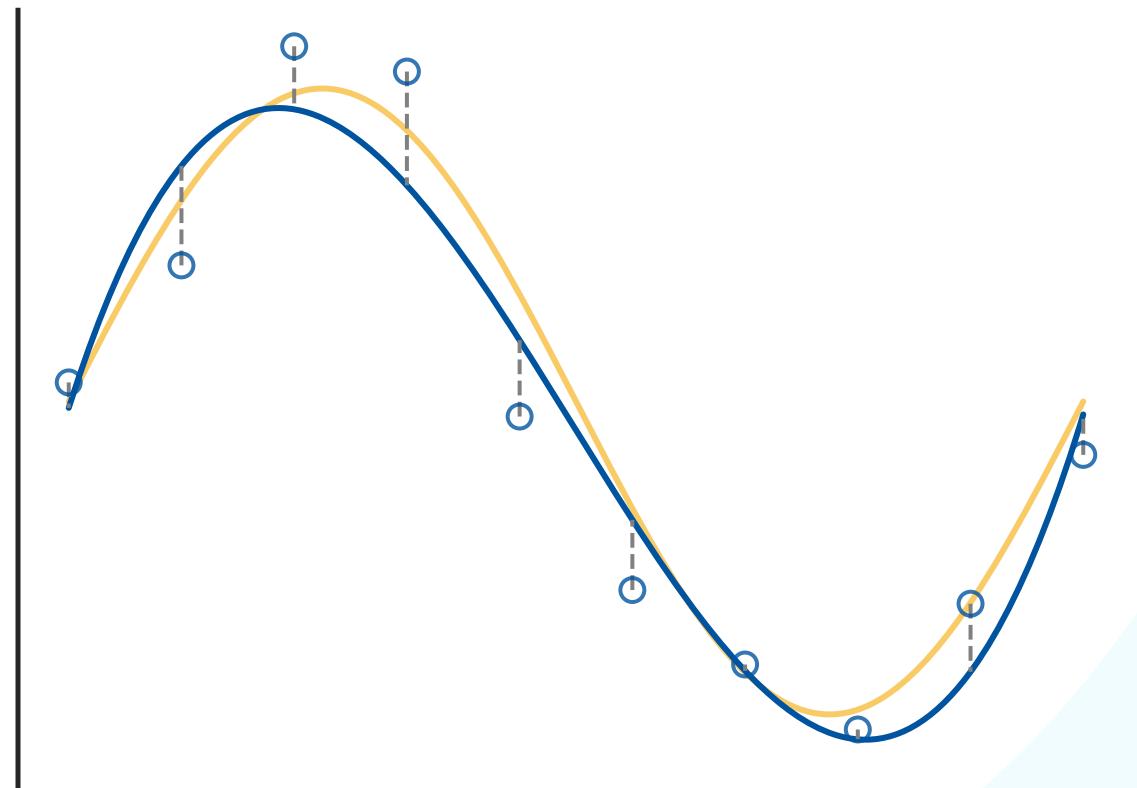


Motivation: Linear Regression

- We have seen how to build classifiers with linear functions:

$$y(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}) + w_0$$

- Now, we will use this model to estimate arbitrary functions using real-valued labels $t_n \in \mathbb{R}$.
- Key assumption: data is generated by some function $h(\mathbf{x})$ with Gaussian noise:
$$t_n = h(\mathbf{x}) + \epsilon$$
- This model is called [linear regression](#).



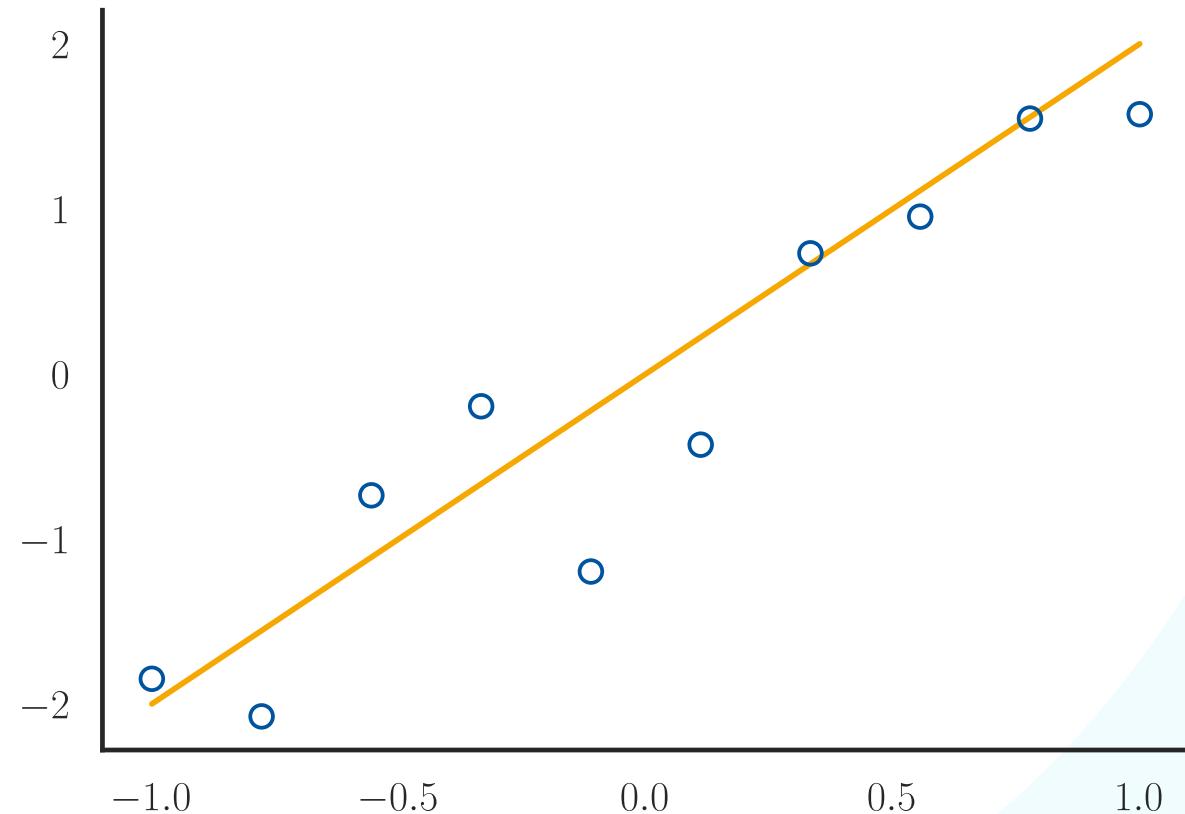
Example: Linear target functions

$$y(x) = w_1x + w_0$$

- We assume ground truth is a linear function:

$$f(x) = mx + b$$

- We try to find a line that minimizes the distance to the samples.



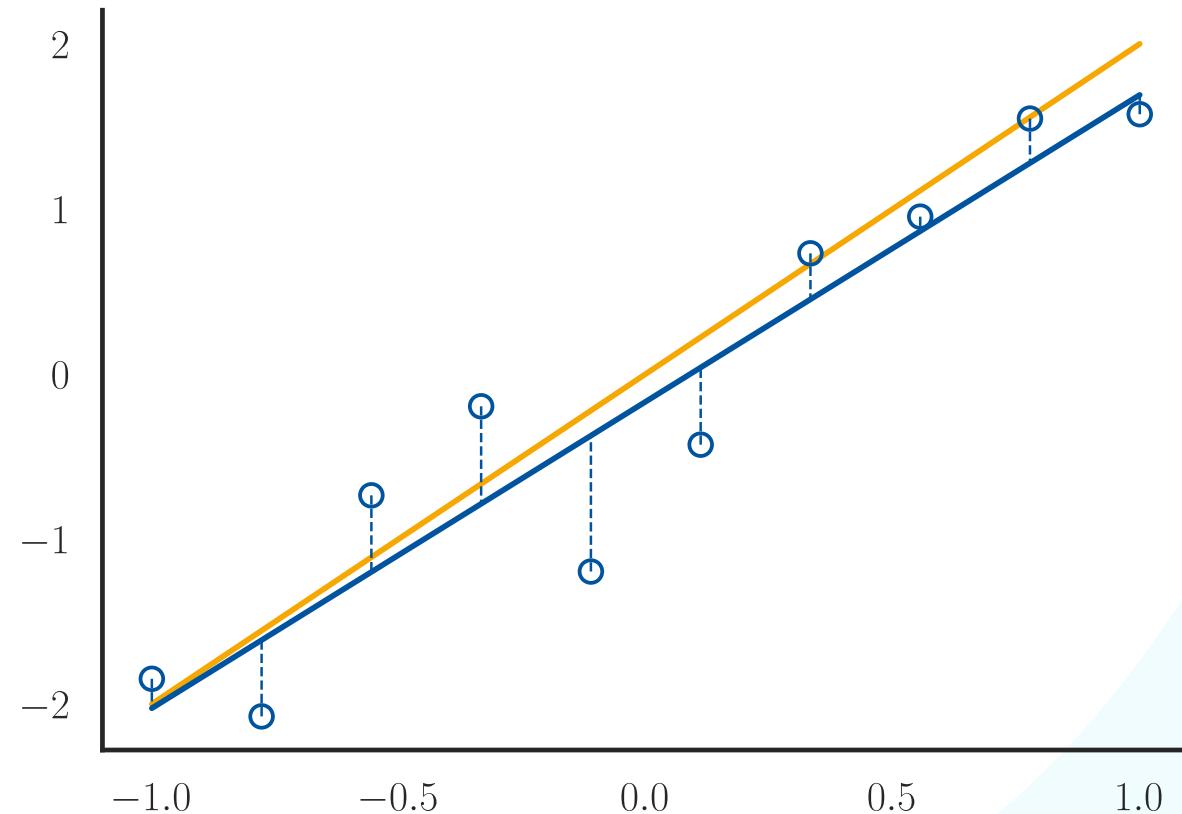
Example: Linear target functions

$$y(x) = w_1x + w_0$$

- We assume ground truth is a linear function:

$$f(x) = mx + b$$

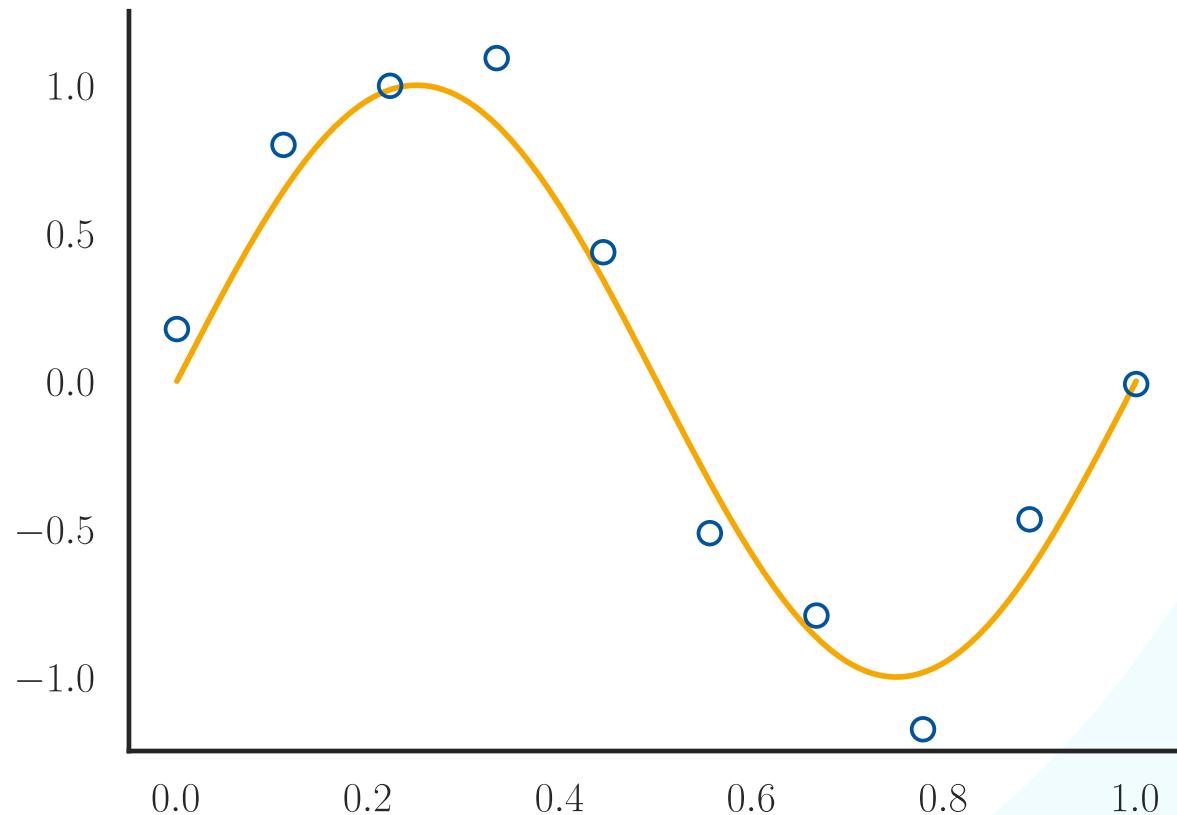
- We try to find a line that minimizes the distance to the samples.



Example: Non-linear target functions

$$y(x) = \mathbf{w}^T \phi(x) + w_0$$

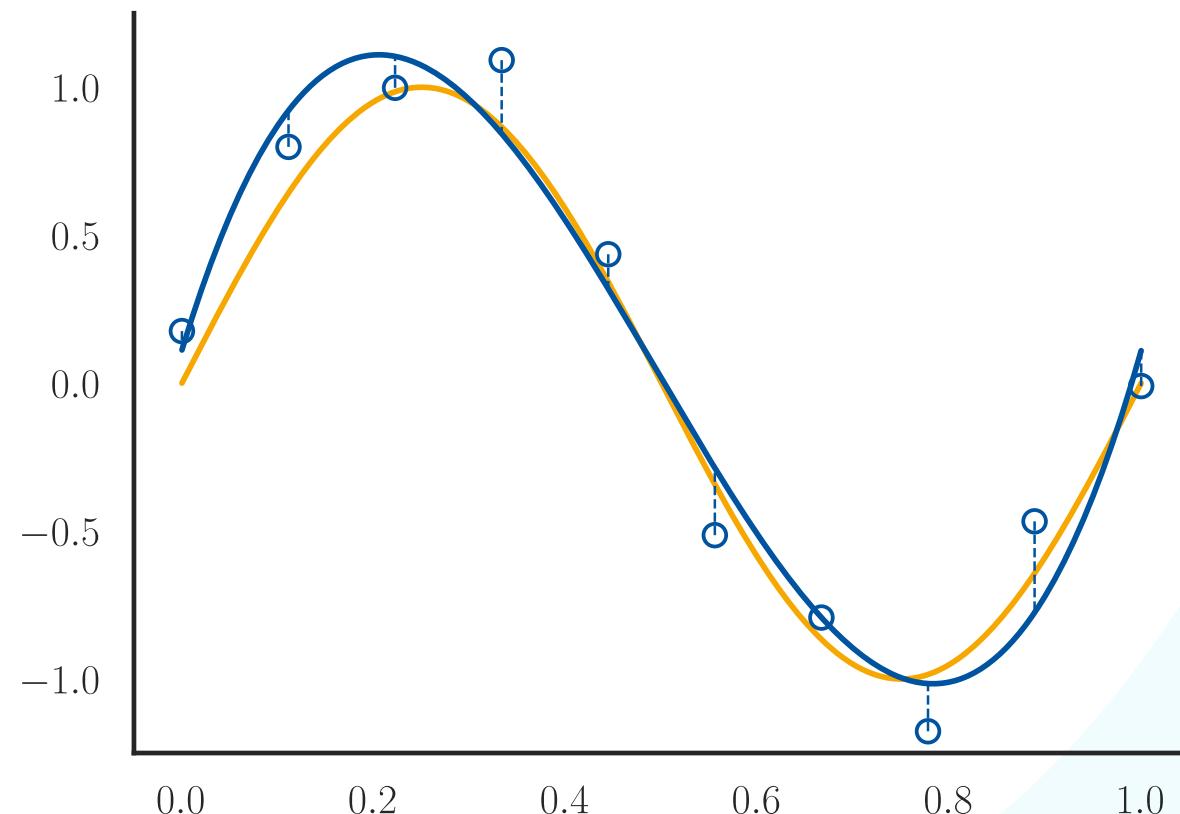
- We can use basis functions to fit arbitrary functions $f(\mathbf{x})$.
- For example, polynomial basis functions.



Example: Non-linear target functions

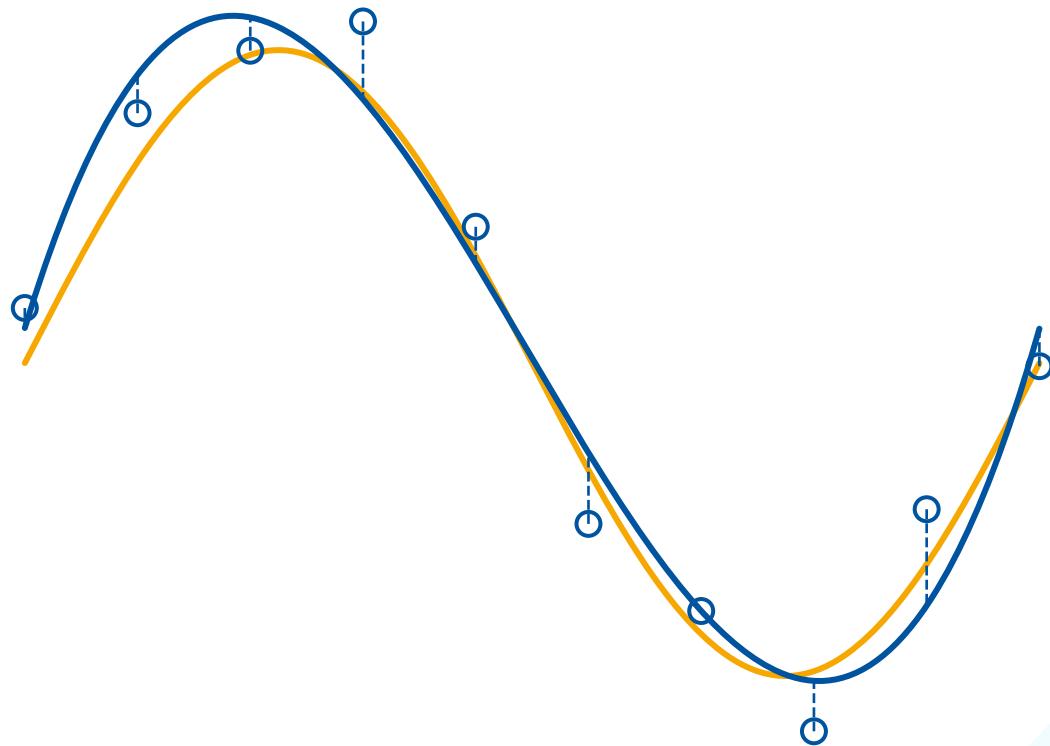
$$y(x) = \mathbf{w}^T \phi(x) + w_0$$

- We can use basis functions to fit arbitrary functions $f(\mathbf{x})$.
- For example, polynomial basis functions.
- Finding a good set of basis functions usually requires some insight into the data...



Linear Regression

1. Motivation
2. **Least-Squares Regression**
3. Regularization
4. Ridge Regression
5. The Bias-Variance Tradeoff



Least-Squares Regression

- We want to optimize the difference between our predictor $y(\mathbf{x}_n; \mathbf{w})$ and the targets t_n .
- The only difference is that our targets t_n are now continuous values.
- Again, use the familiar **squared error** objective:

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (y(\mathbf{x}_n; \mathbf{w}) - t_n)^2$$

- This has the same solution as for classification (normal equations).

$$y(\mathbf{x}_n; \mathbf{w}) = \mathbf{w}^\top \phi(\mathbf{x}_n)$$

$$\begin{aligned}\frac{\partial E(\mathbf{w})}{\partial \mathbf{w}} &= \sum_{n=1}^N (\mathbf{w}^\top \phi_n - t_n) \phi_n \\ &= \Phi^\top (\Phi \mathbf{w} - \mathbf{t}) \stackrel{!}{=} 0\end{aligned}$$

$$\Rightarrow \mathbf{w} = (\Phi^\top \Phi)^{-1} \Phi^\top \mathbf{t}$$

Example: Fitting a polynomial

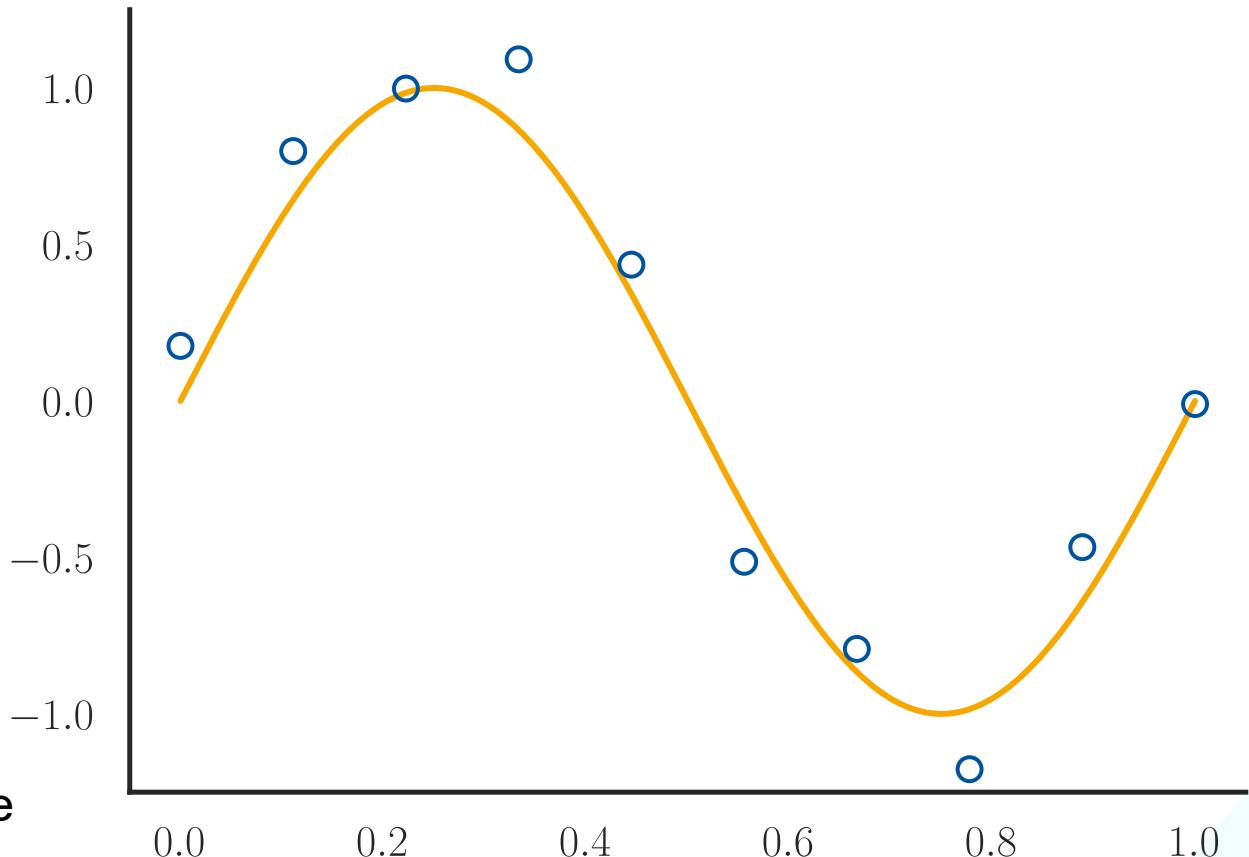
- This is clearly not a linear function.
- Let's use polynomial basis functions:

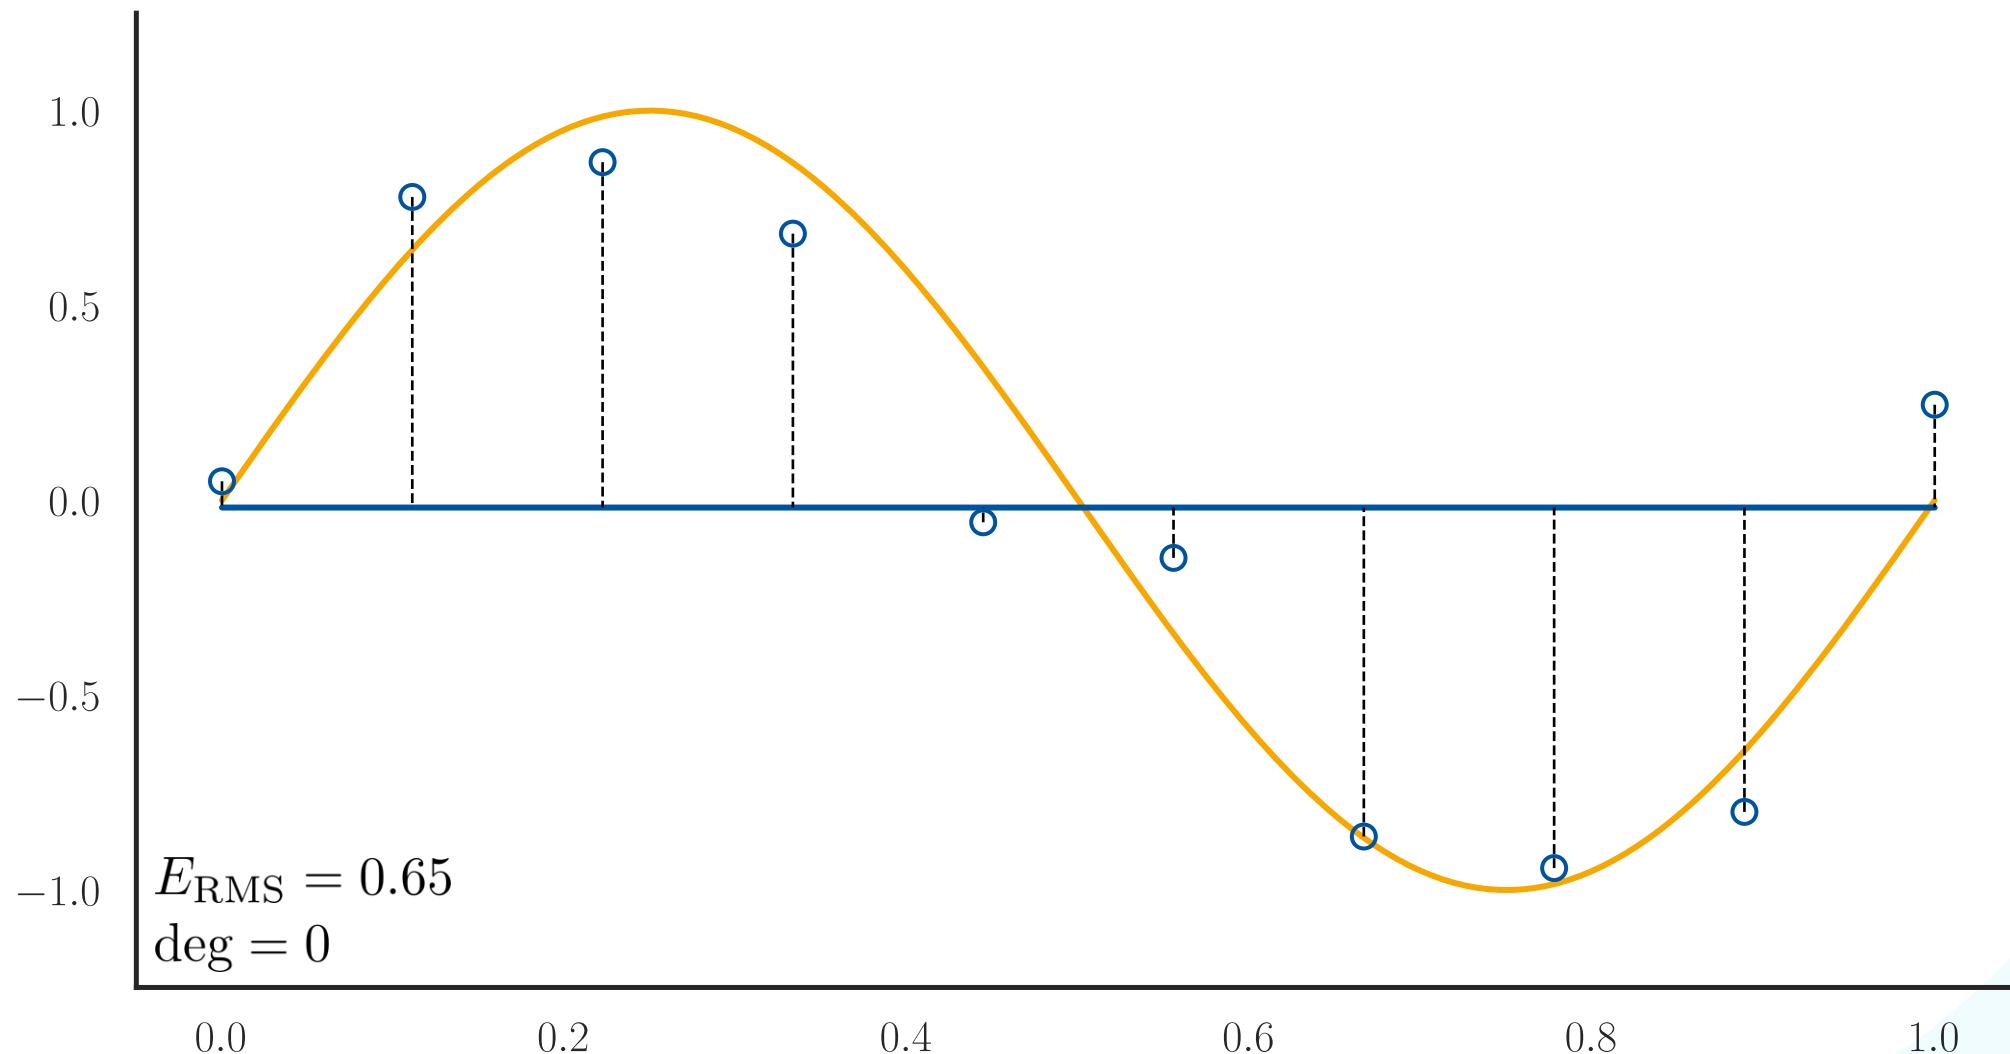
$$\phi_j(x) = (x^j)$$

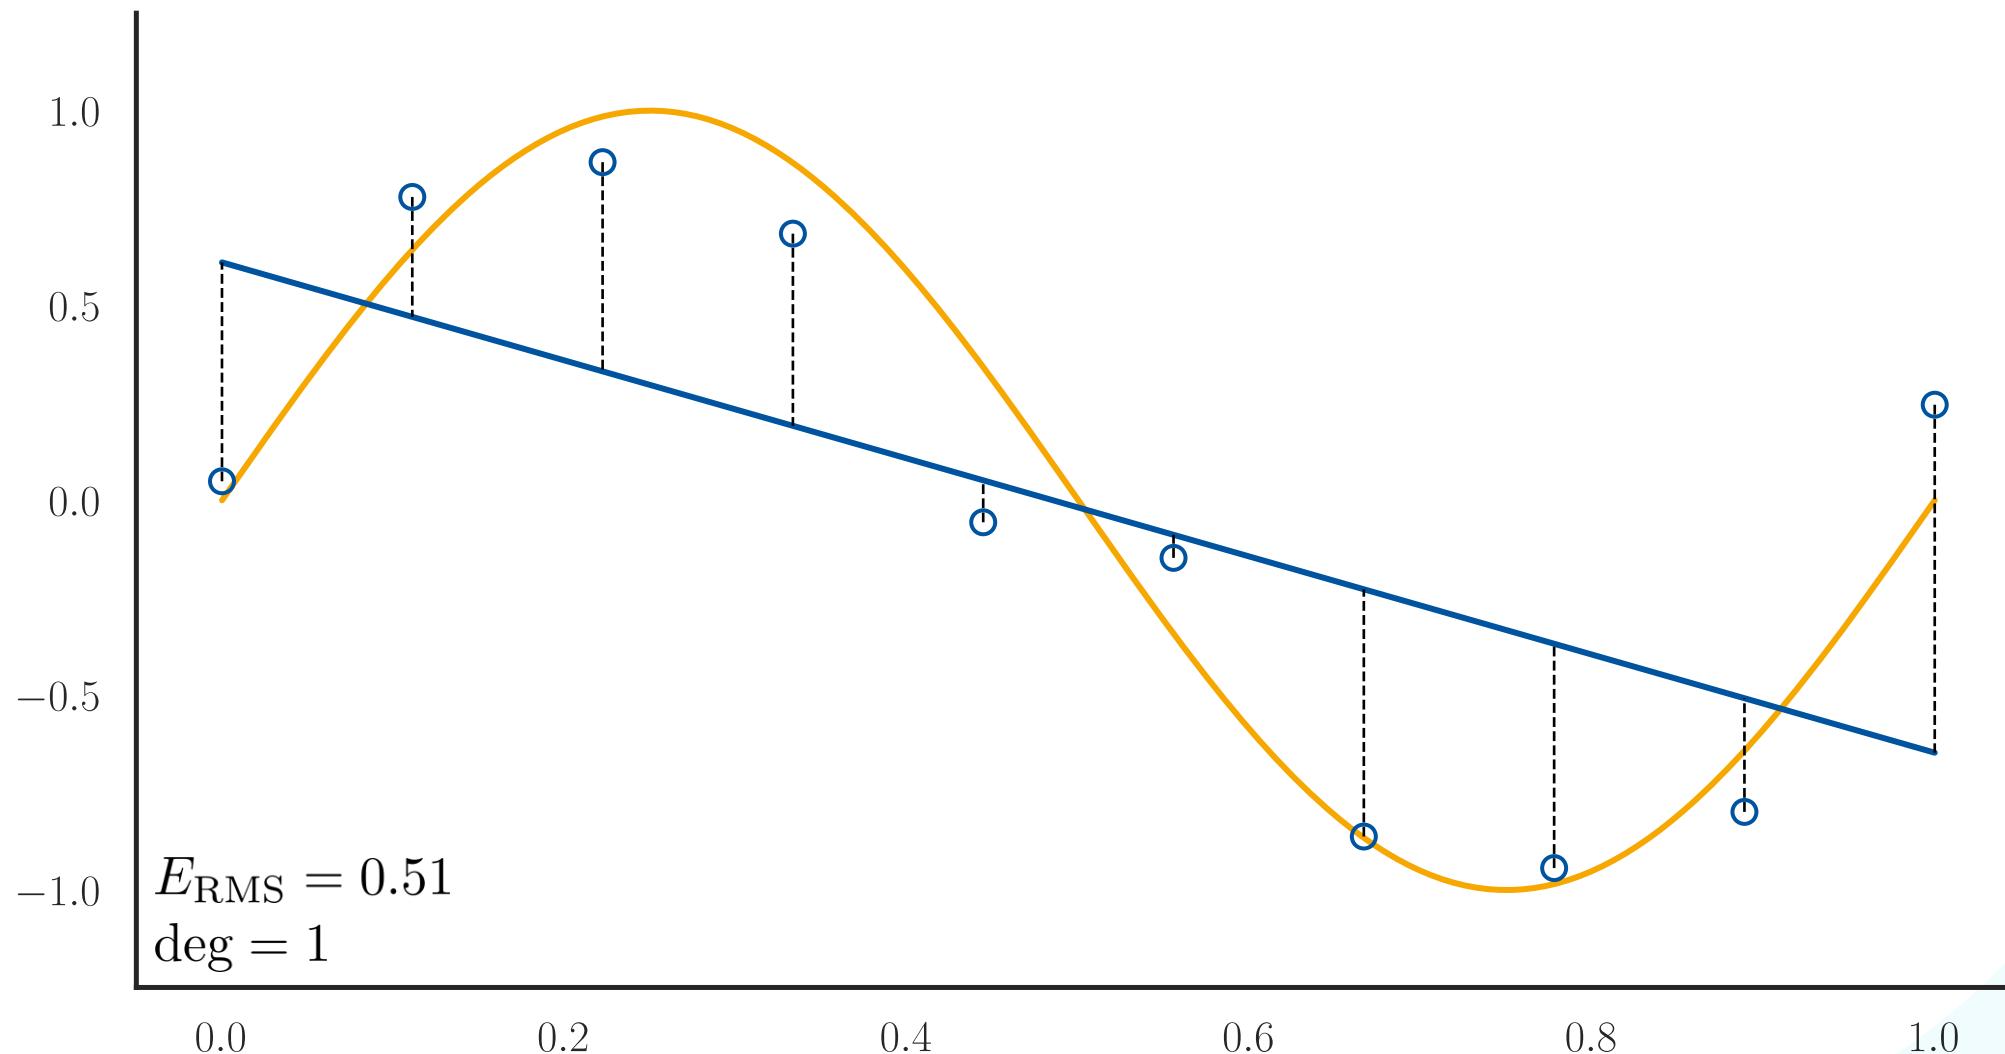
- Which degree should we use?
- Compare different models by their Root Mean Square Error:

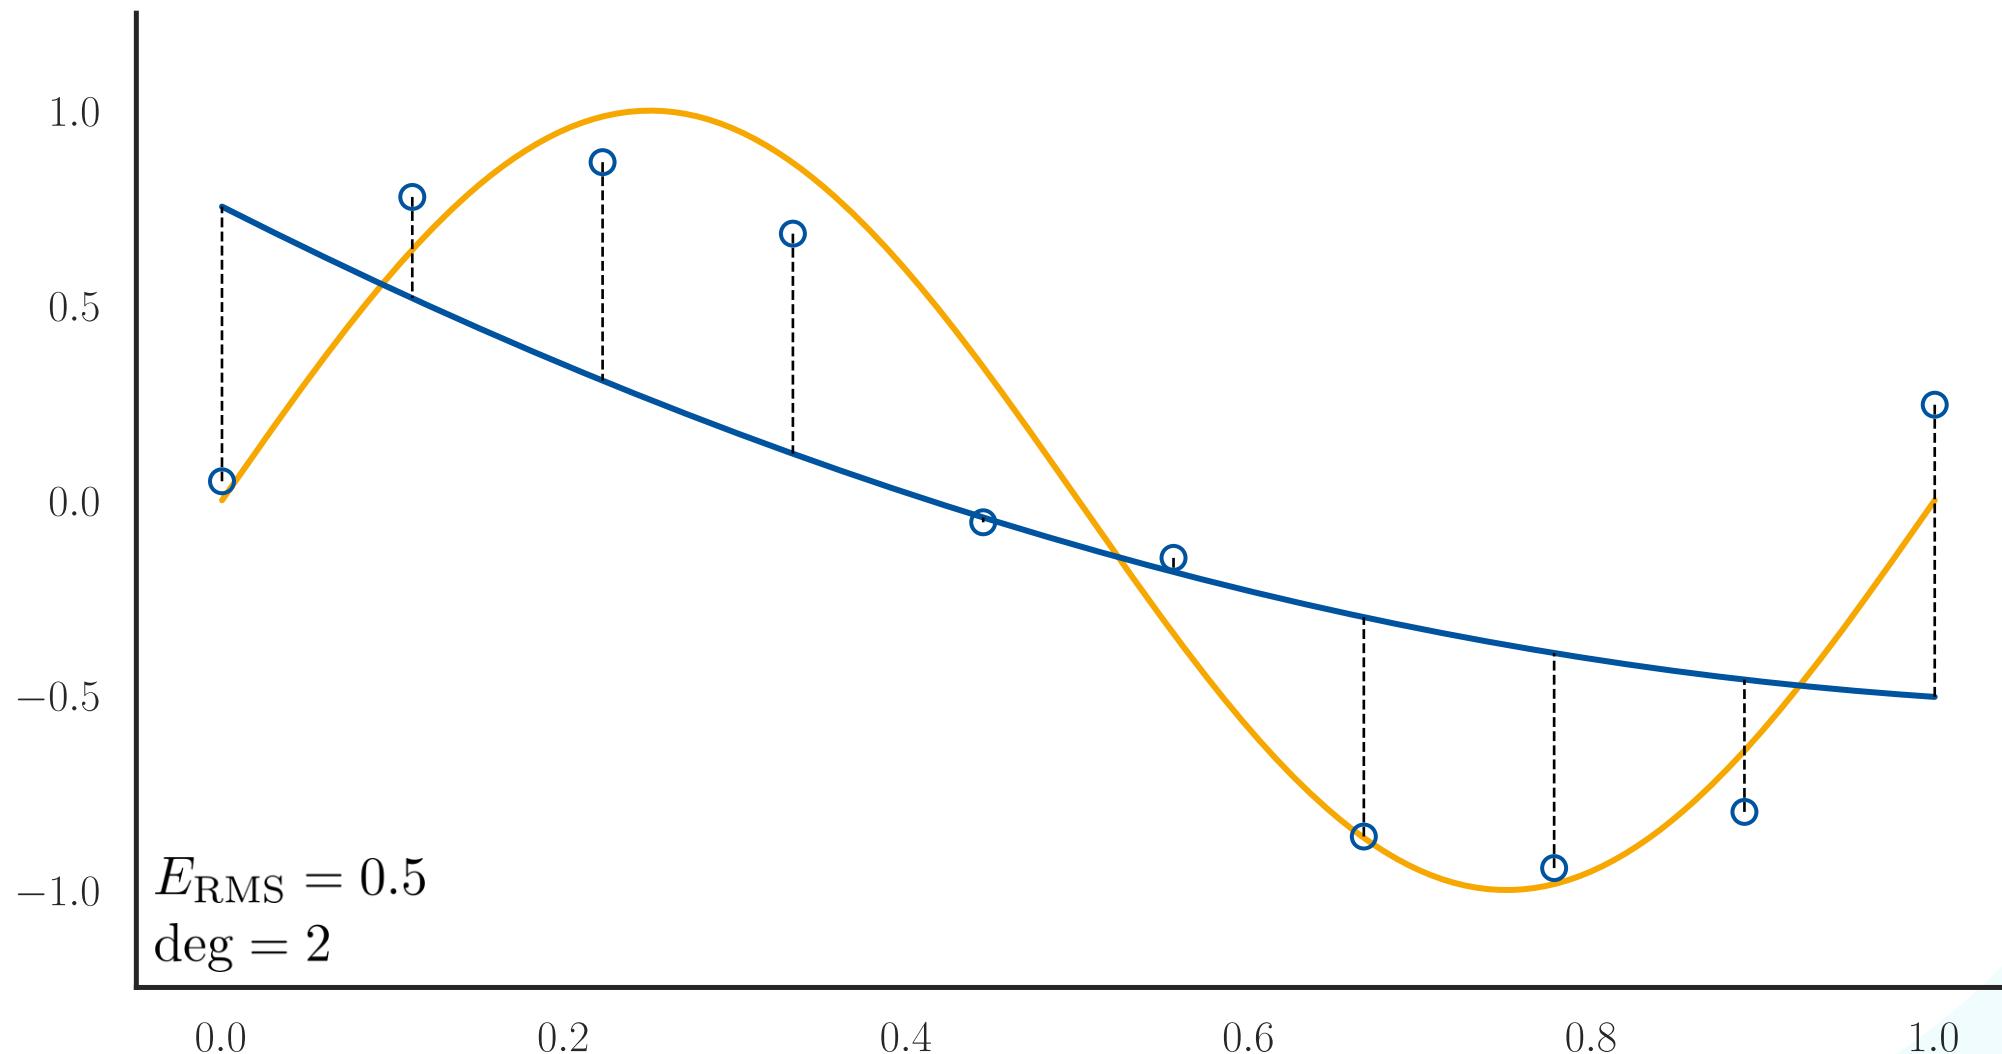
$$E_{\text{RMS}} = \sqrt{\frac{2E(\mathbf{w}^*)}{N}}$$

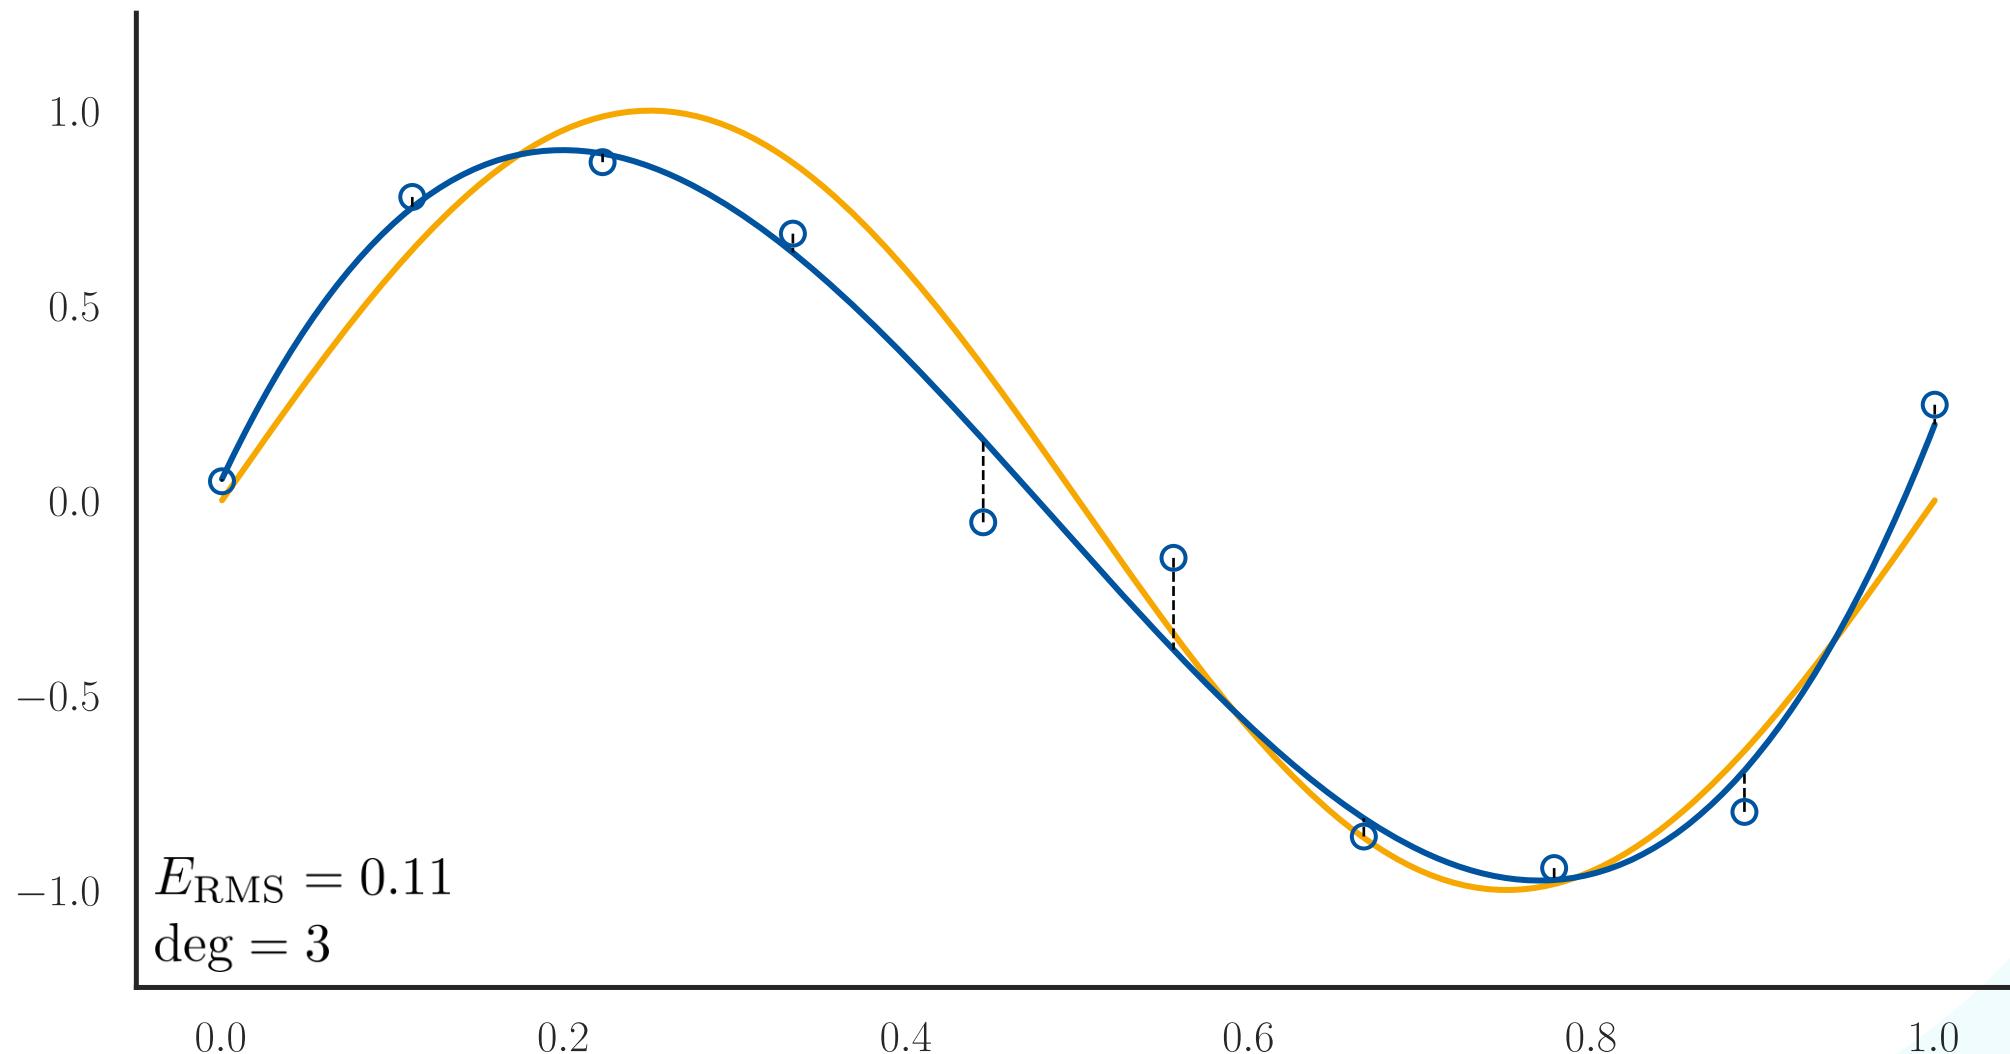
- RMS is independent of training set size and has the same scale as the data.

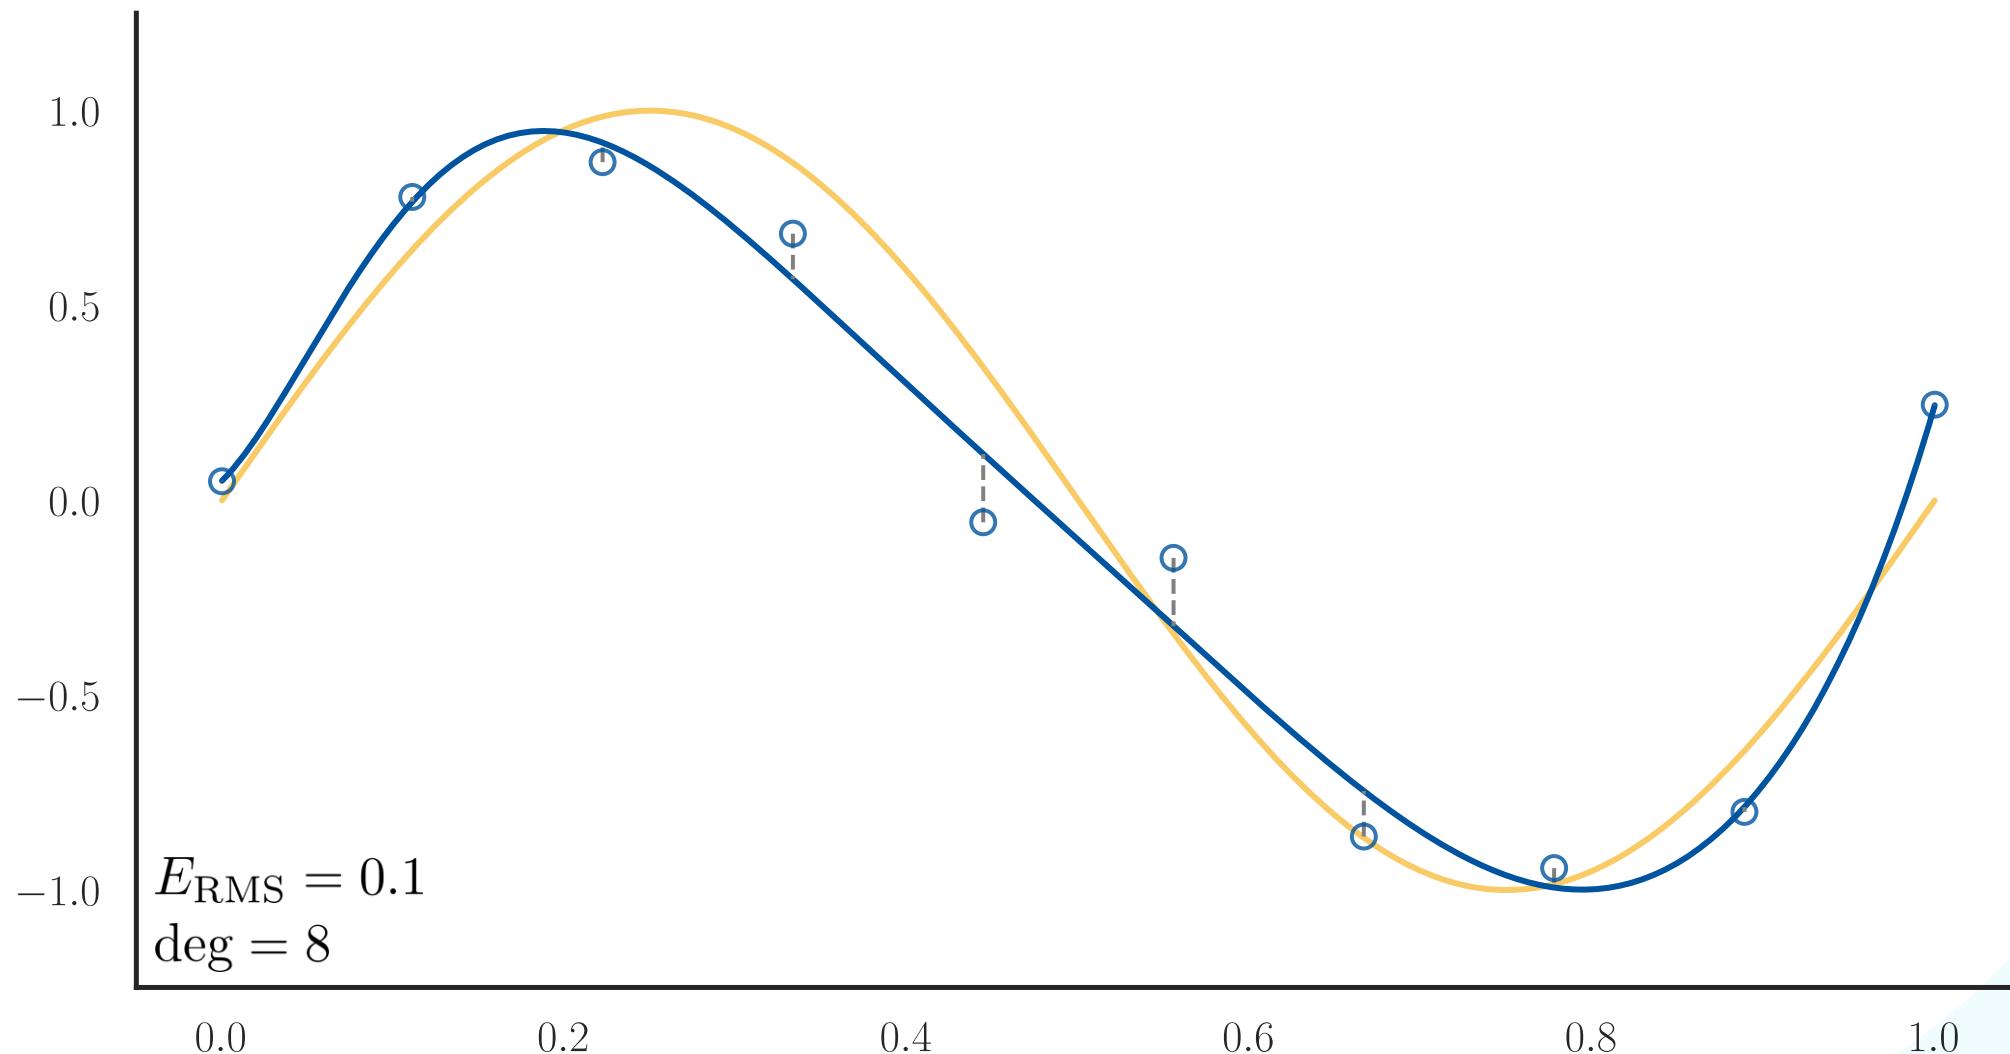


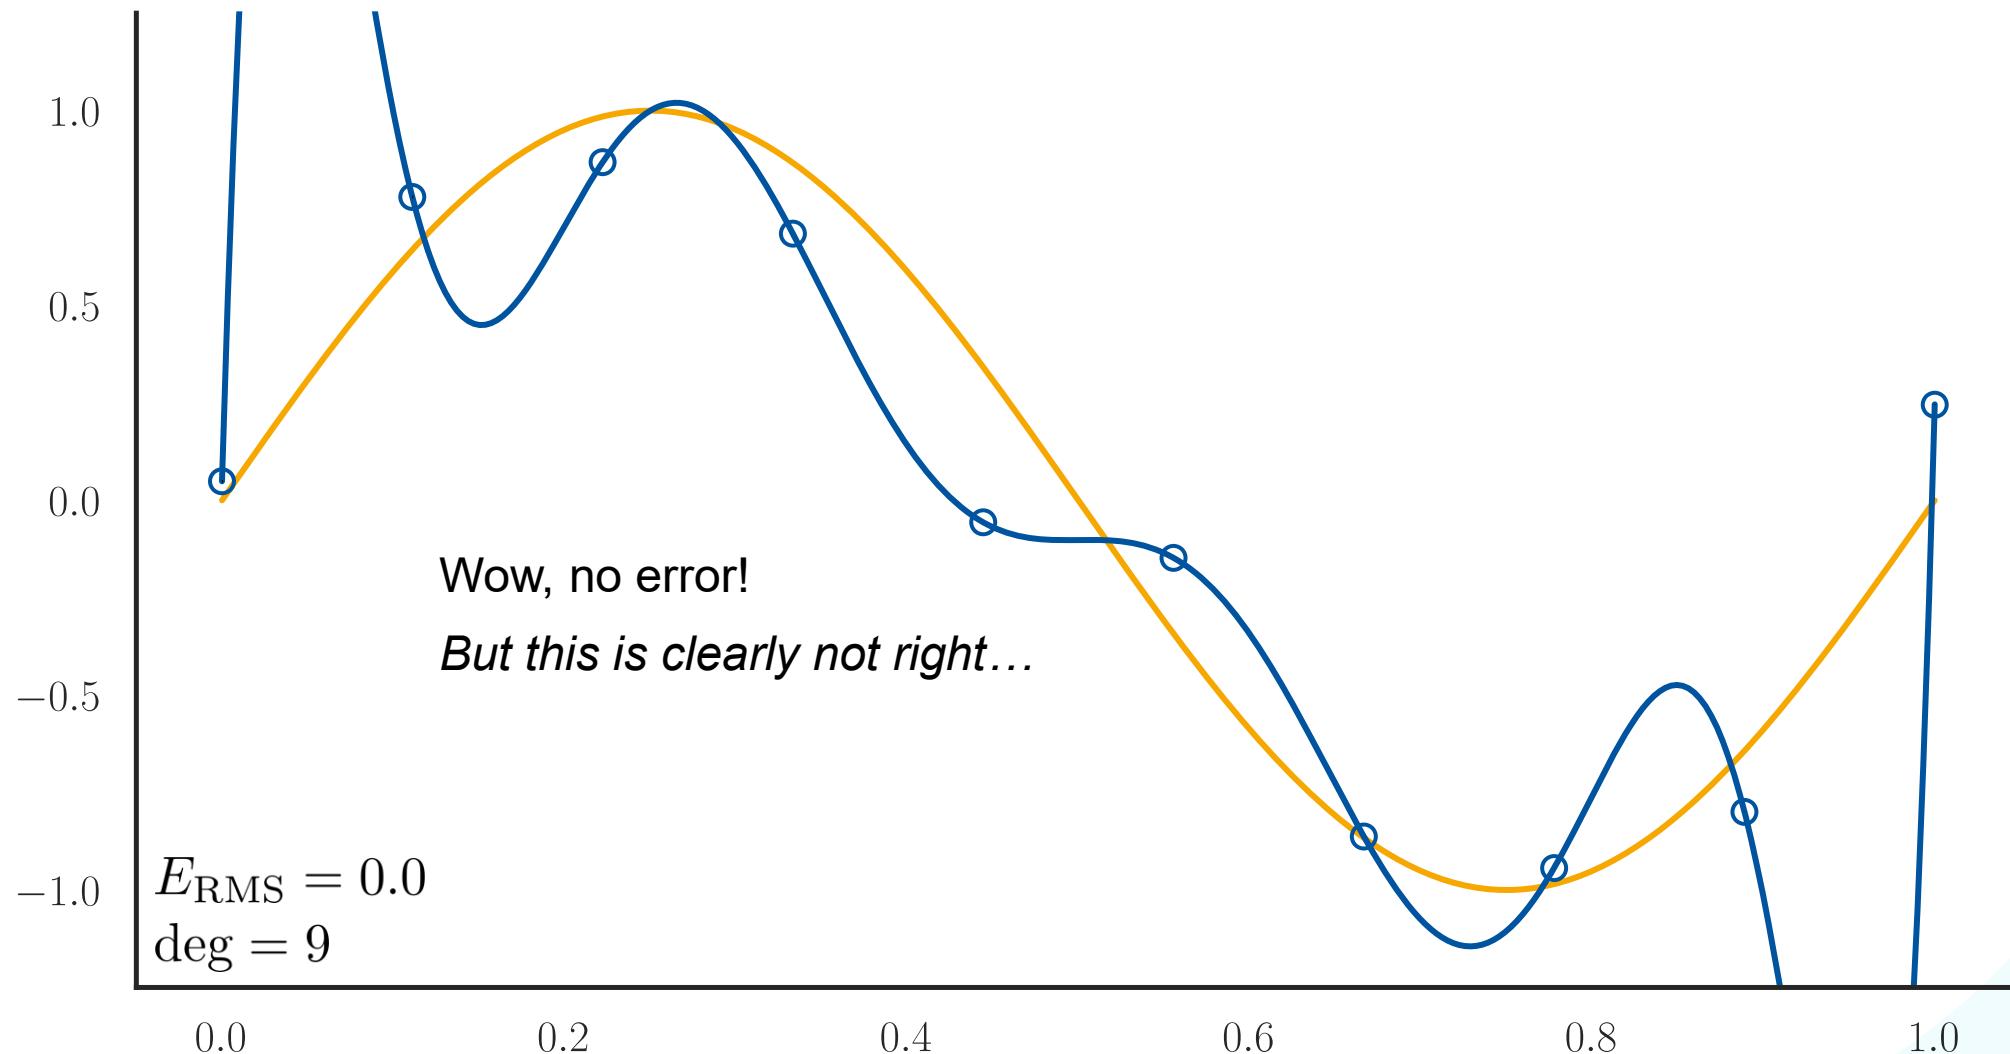


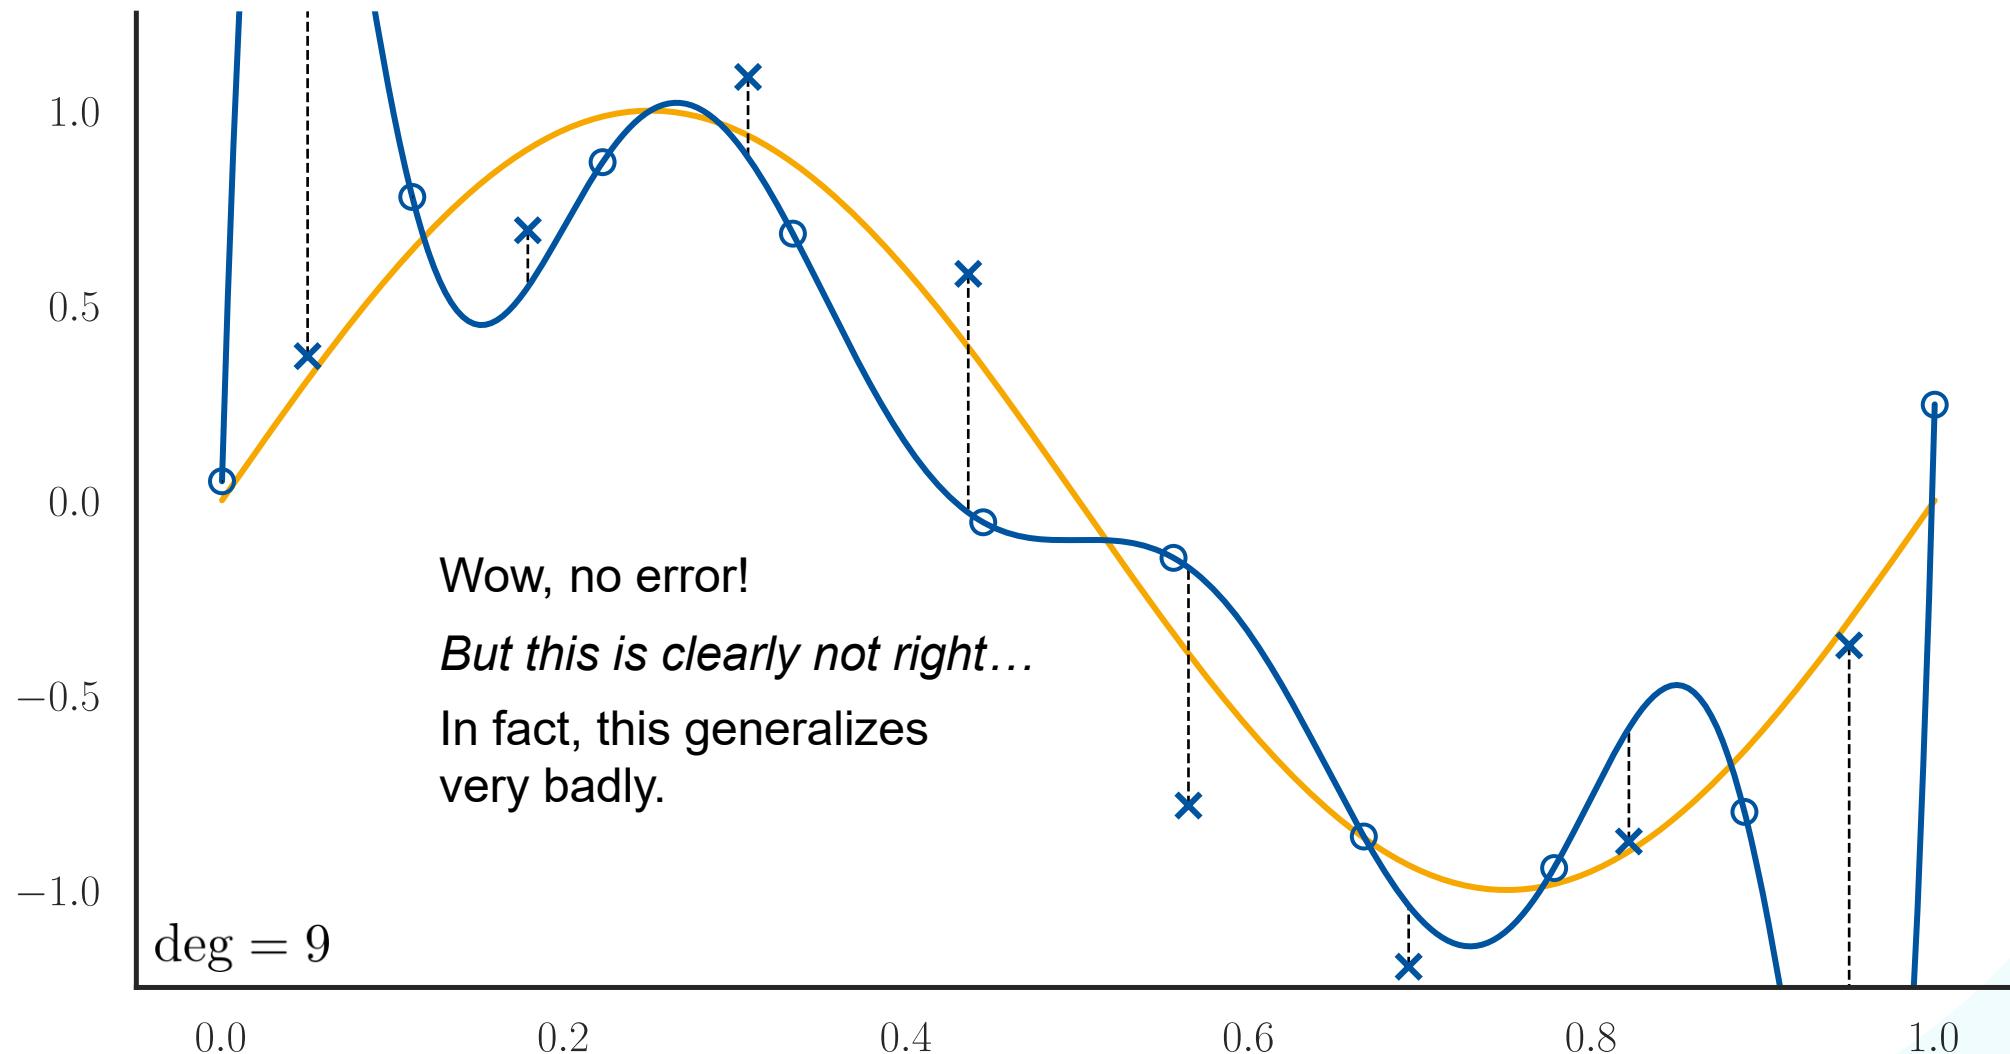






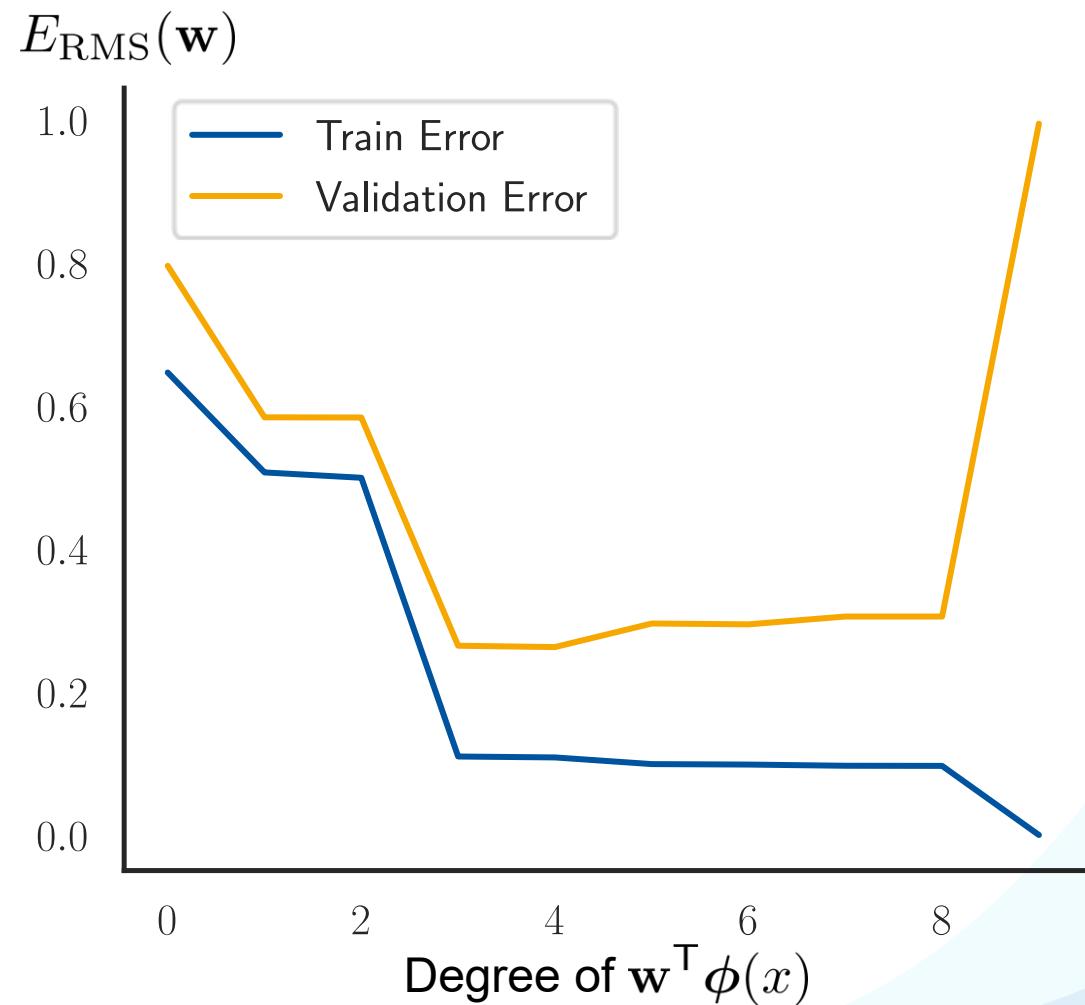


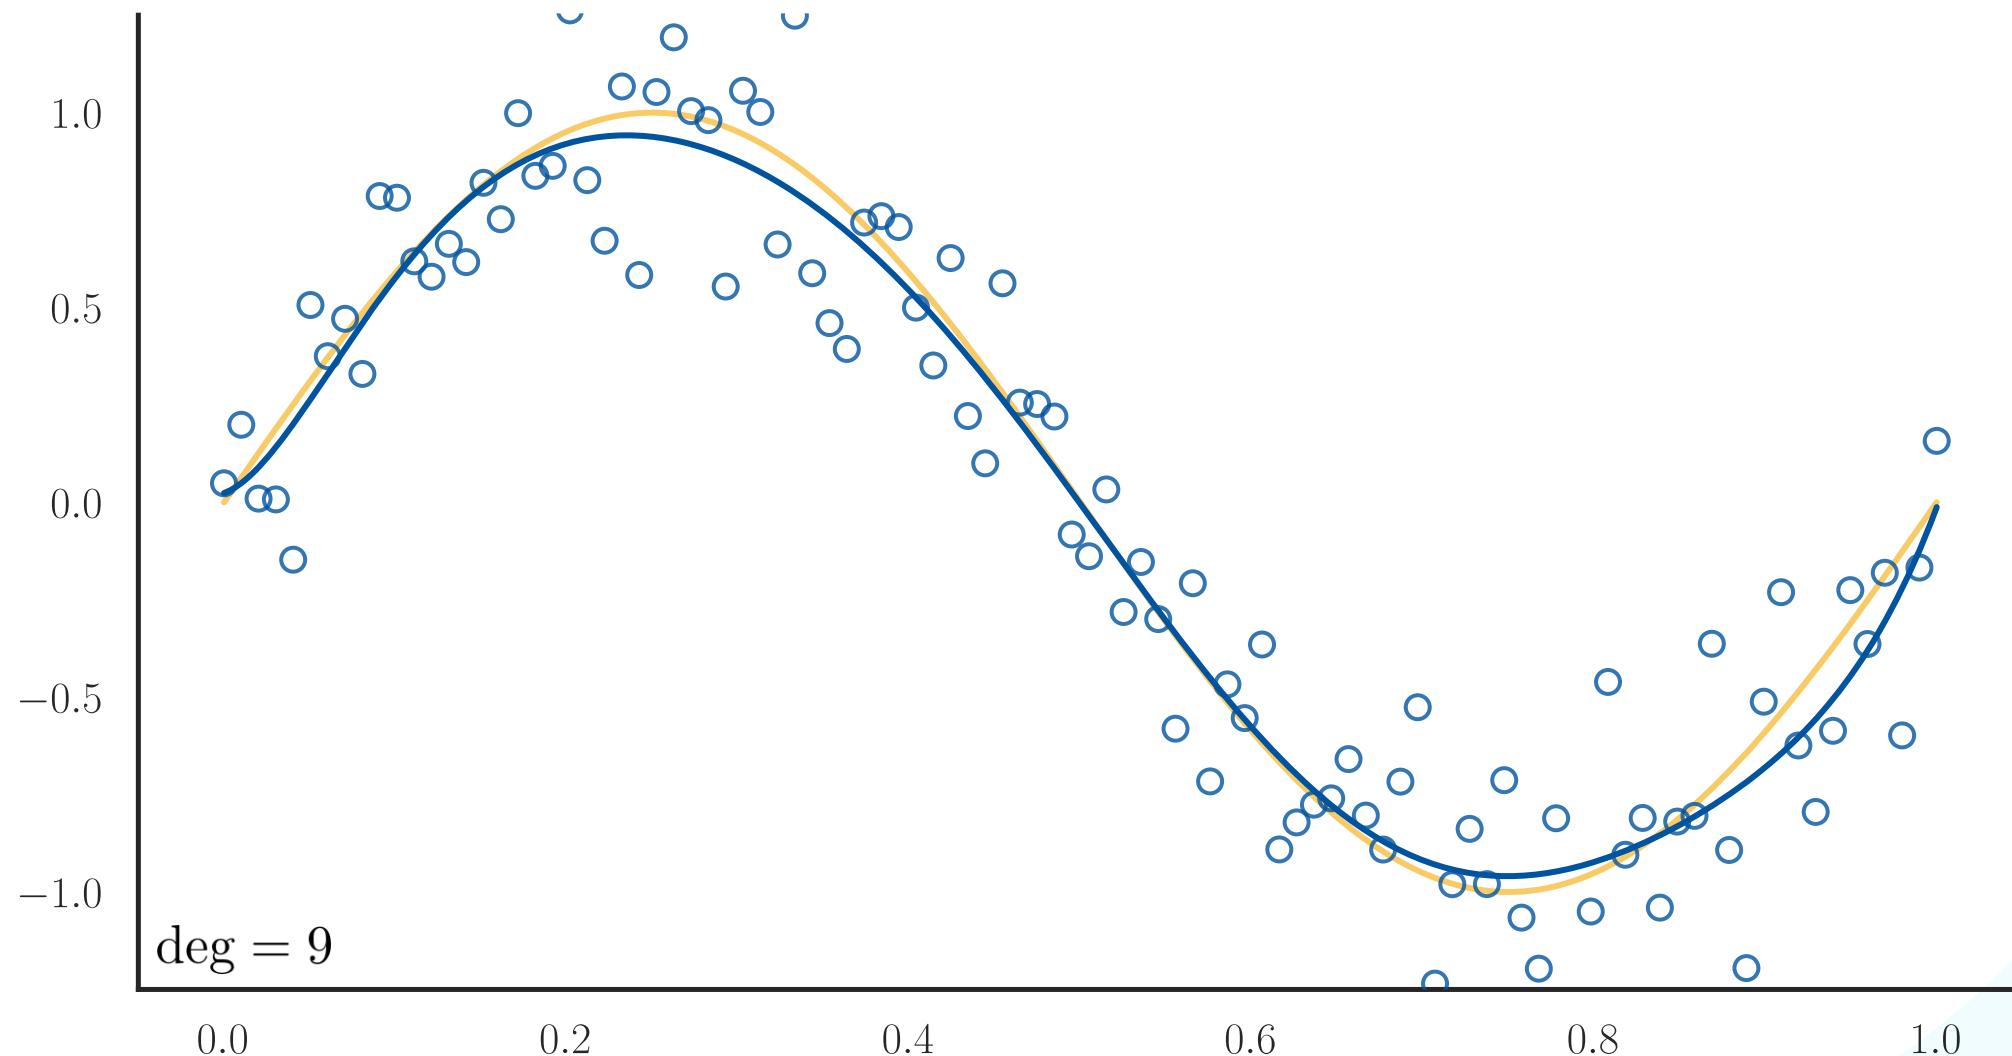




Overfitting

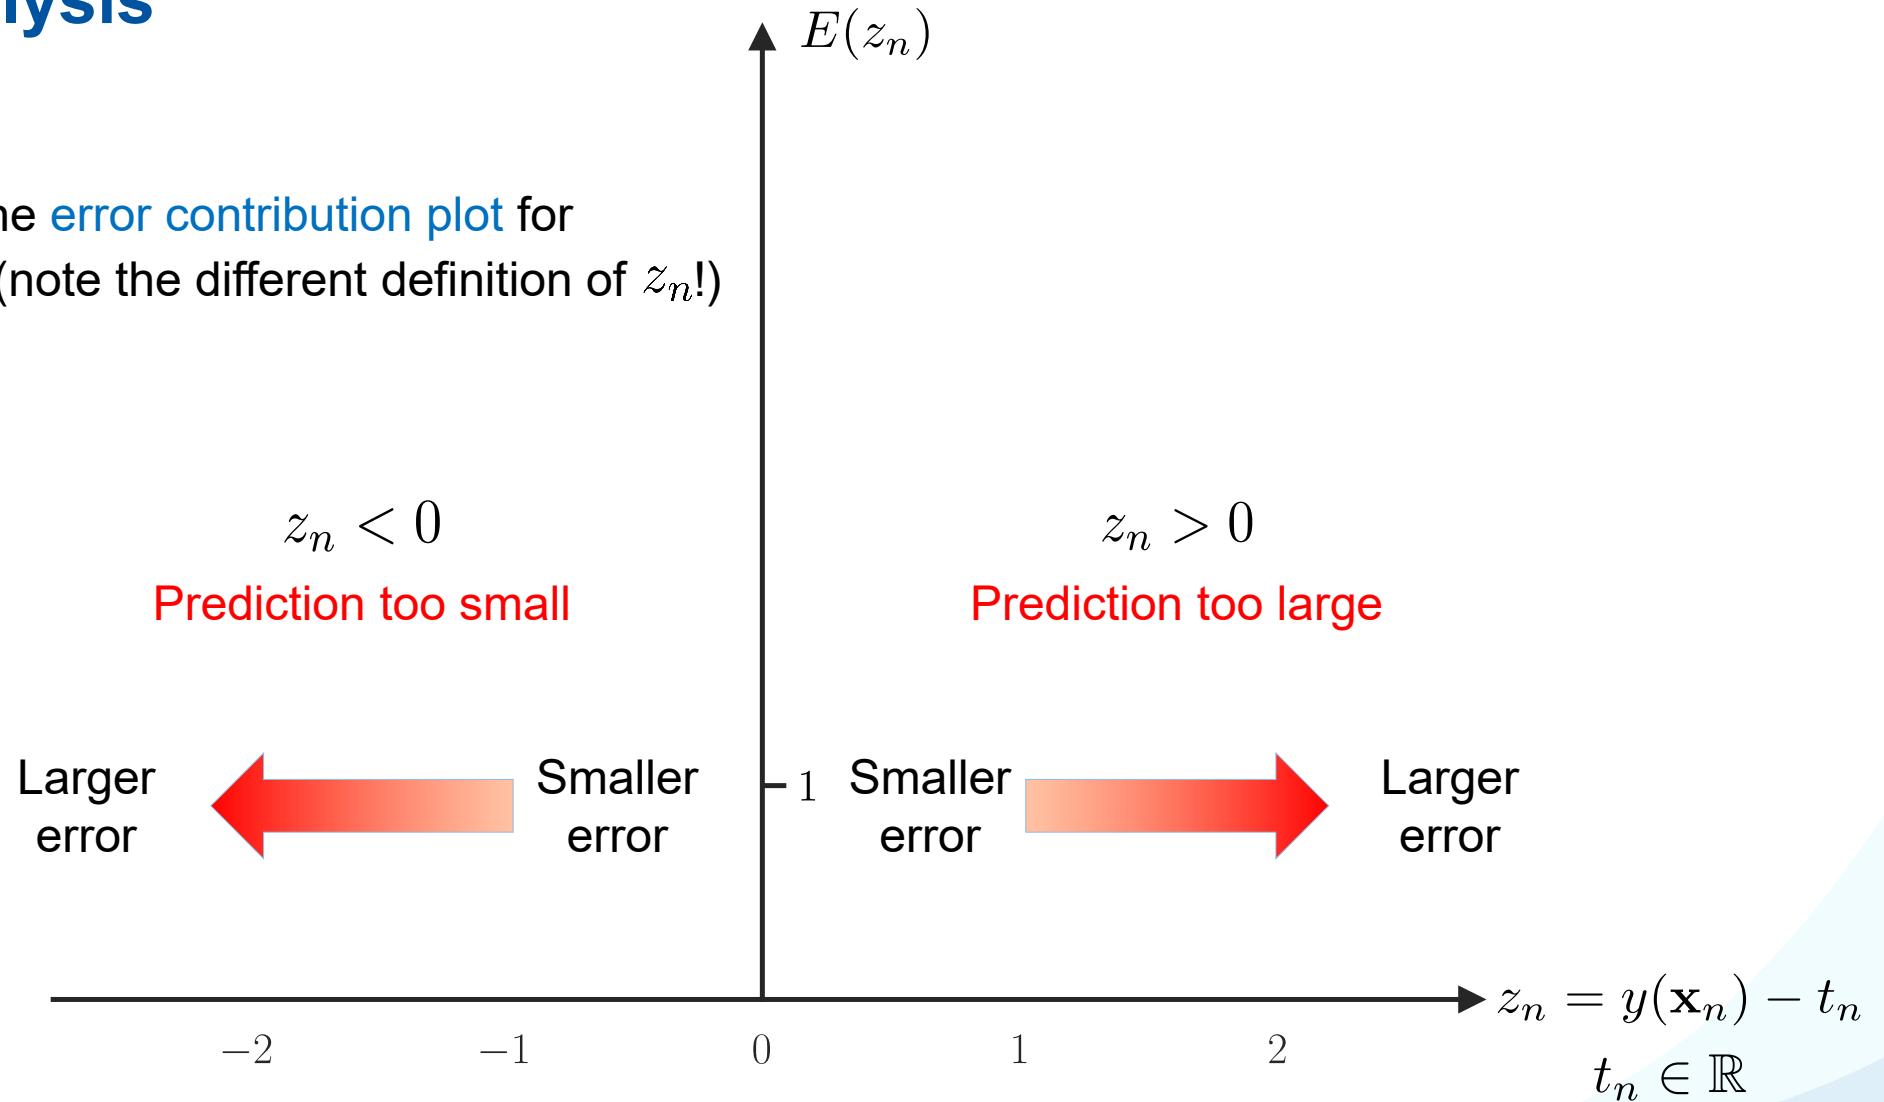
- We fit the dataset perfectly, but the resulting function is clearly not what we want.
- This phenomenon is called **overfitting**.
- Remember: we assume $t_n = h(\mathbf{x}_n) + \epsilon$.
- Our model is “too” powerful and models the noise instead of the underlying function!
- *What can we do to avoid overfitting?*
 - One solution: More data!



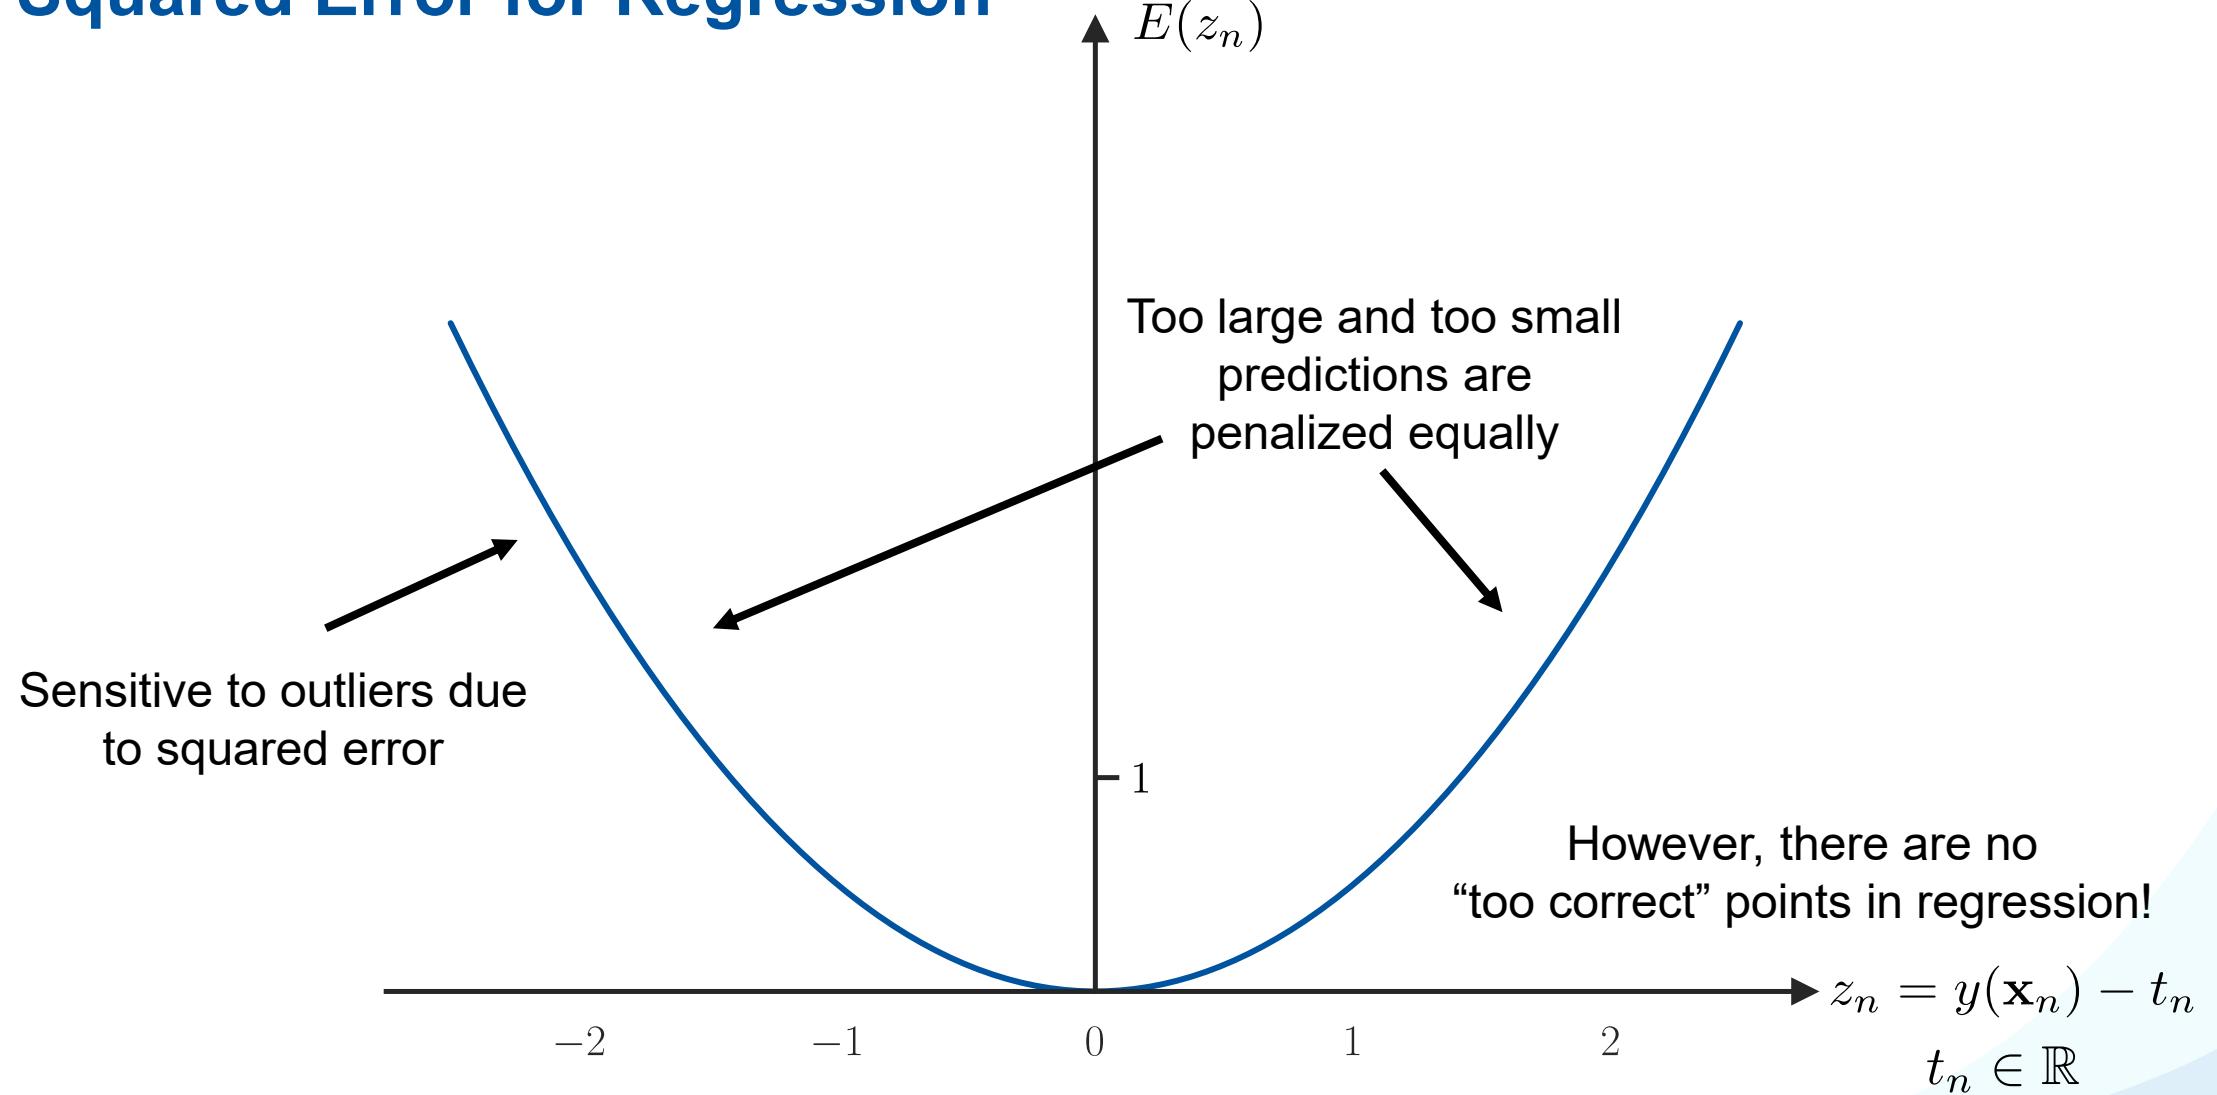


Error Analysis

- Variant of the error contribution plot for regression (note the different definition of z_n !)



Squared Error for Regression



Discussion: Least-Squares Regression

Advantages

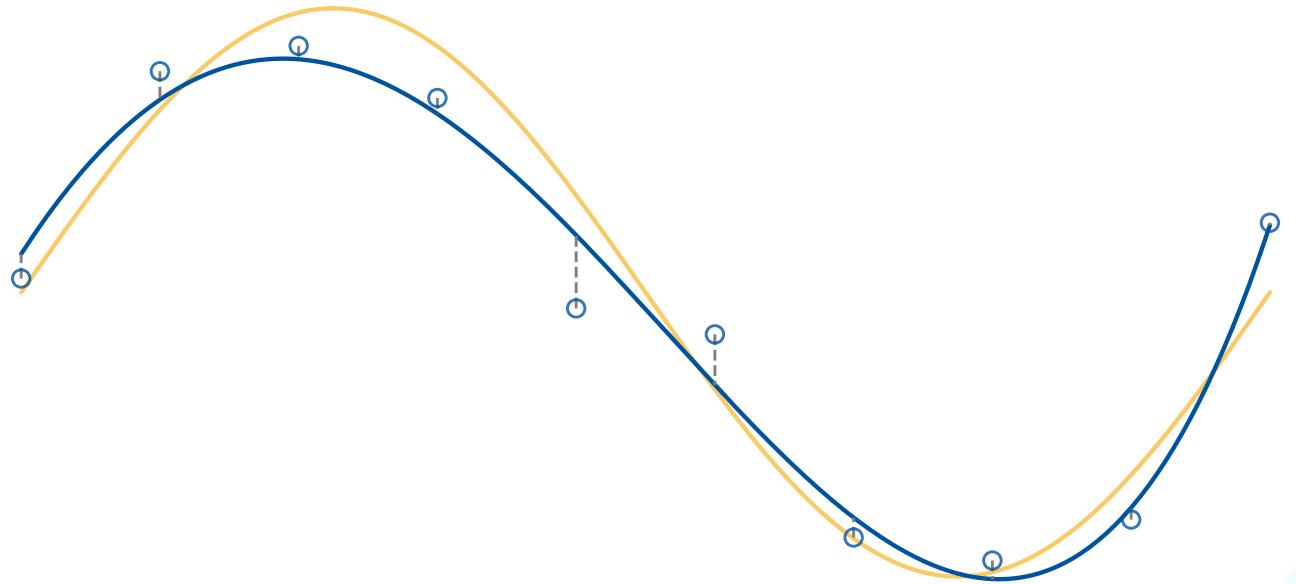
- Squared error leads to closed-form solution of the regression problem.
- We can use basis functions to fit non-linear functions while staying within the framework of linear regression.
- Polynomial basis functions with different degrees of the polynomial result in regression functions with different capacities to approximate the target function.
- We can compare their results using the RMS error.

Limitations

- The squared error for regression is not robust to outliers, but it does not exhibit the systematic problems of least-squares classification.
- Overfitting when the degree of the polynomial becomes too large.
- Overfitting is a function of the available amount of data (and is more likely to occur with small training sets)

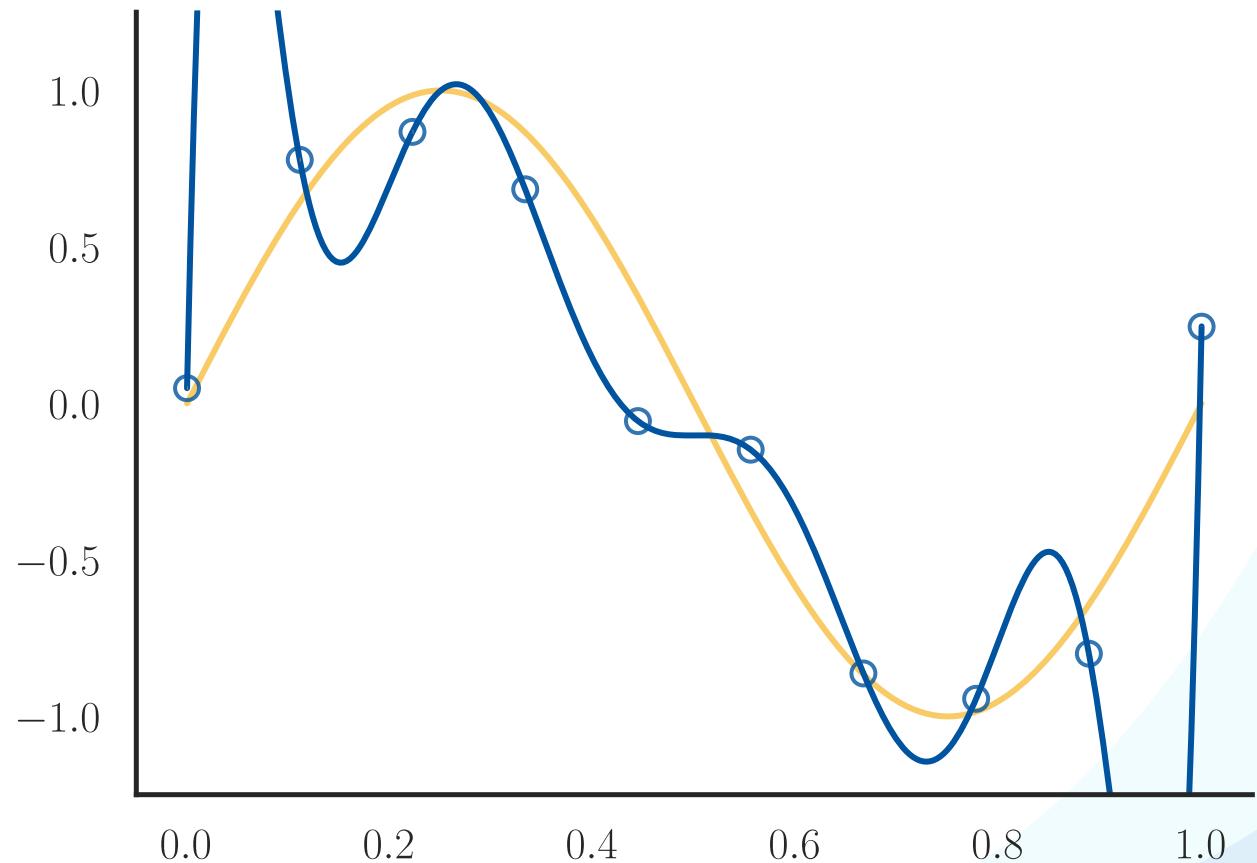
Linear Regression

1. Linear Regression
2. Least-Squares Regression
- 3. Regularization**
4. Ridge Regression
5. The Bias-Variance Tradeoff



Regularization

- With enough parameters, our model will overfit to the training set.



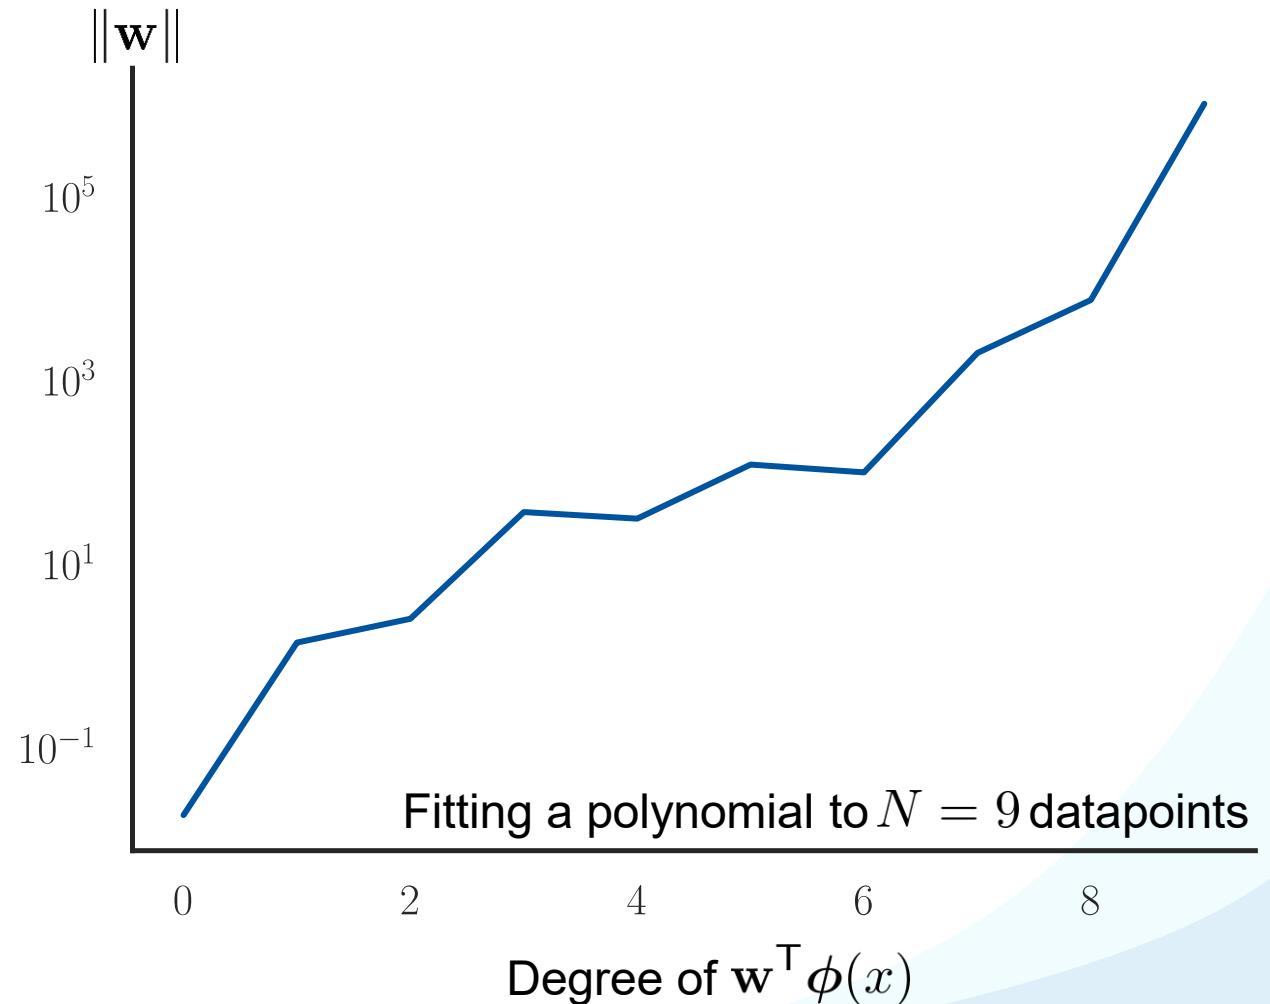
Regularization

- With enough parameters, our model will overfit to the training set.
- This leads to very large coefficient values w_i and thus to a large $\|\mathbf{w}\|$.
- Solution: penalize large parameters.

$$E(\mathbf{w}) = L(\mathbf{w}) + \lambda\Omega(\mathbf{w})$$

$$\Omega(\mathbf{w}) = \frac{1}{2}\|\mathbf{w}\|^2$$

- $L(\mathbf{w})$ is called the **loss** term. Here, we can use the familiar squared loss.
- $\Omega(\mathbf{w})$ is called the **regularizer**. Here, we use a squared regularizer.



Note: Excluding the Bias

- The bias w_0 is usually not regularized, since it does not change the functions' complexity.
- Therefore, we do not include it in $\Omega(\mathbf{w})$ here.
- We can fit the model without a bias by estimating \mathbf{w} on **centered** data:

$$t_n^c = t_n - \bar{t} \quad \mathbf{x}_n^c = \mathbf{x}_n - \bar{\mathbf{x}}$$

$$\bar{t} = \frac{1}{N} \sum_{n=1}^N t_n \quad \bar{\mathbf{x}} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

- And computing w_0 afterwards:

$$w_0 = \bar{t} - \mathbf{w}^\top \bar{\mathbf{x}}$$

$$y(\mathbf{x}; \mathbf{w}) = \mathbf{w}^\top \phi(\mathbf{x}) + w_0$$

$$L(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N ((\mathbf{w}^\top \phi(\mathbf{x}_n) + w_0) - t_n)^2$$

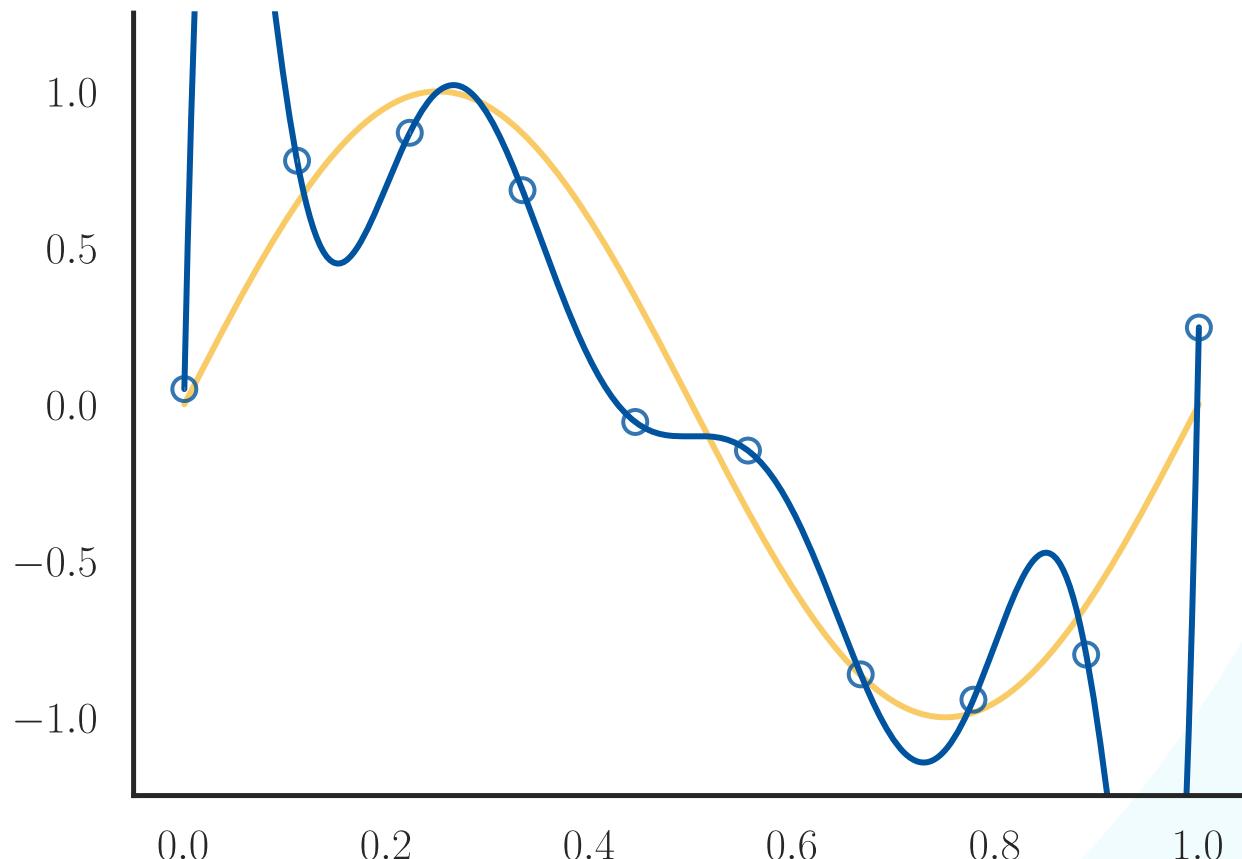
$$\Omega(\mathbf{w}) = \frac{1}{2} \|\mathbf{w}\|^2 = \frac{1}{2} \sum_{j=1}^M w_j^2$$

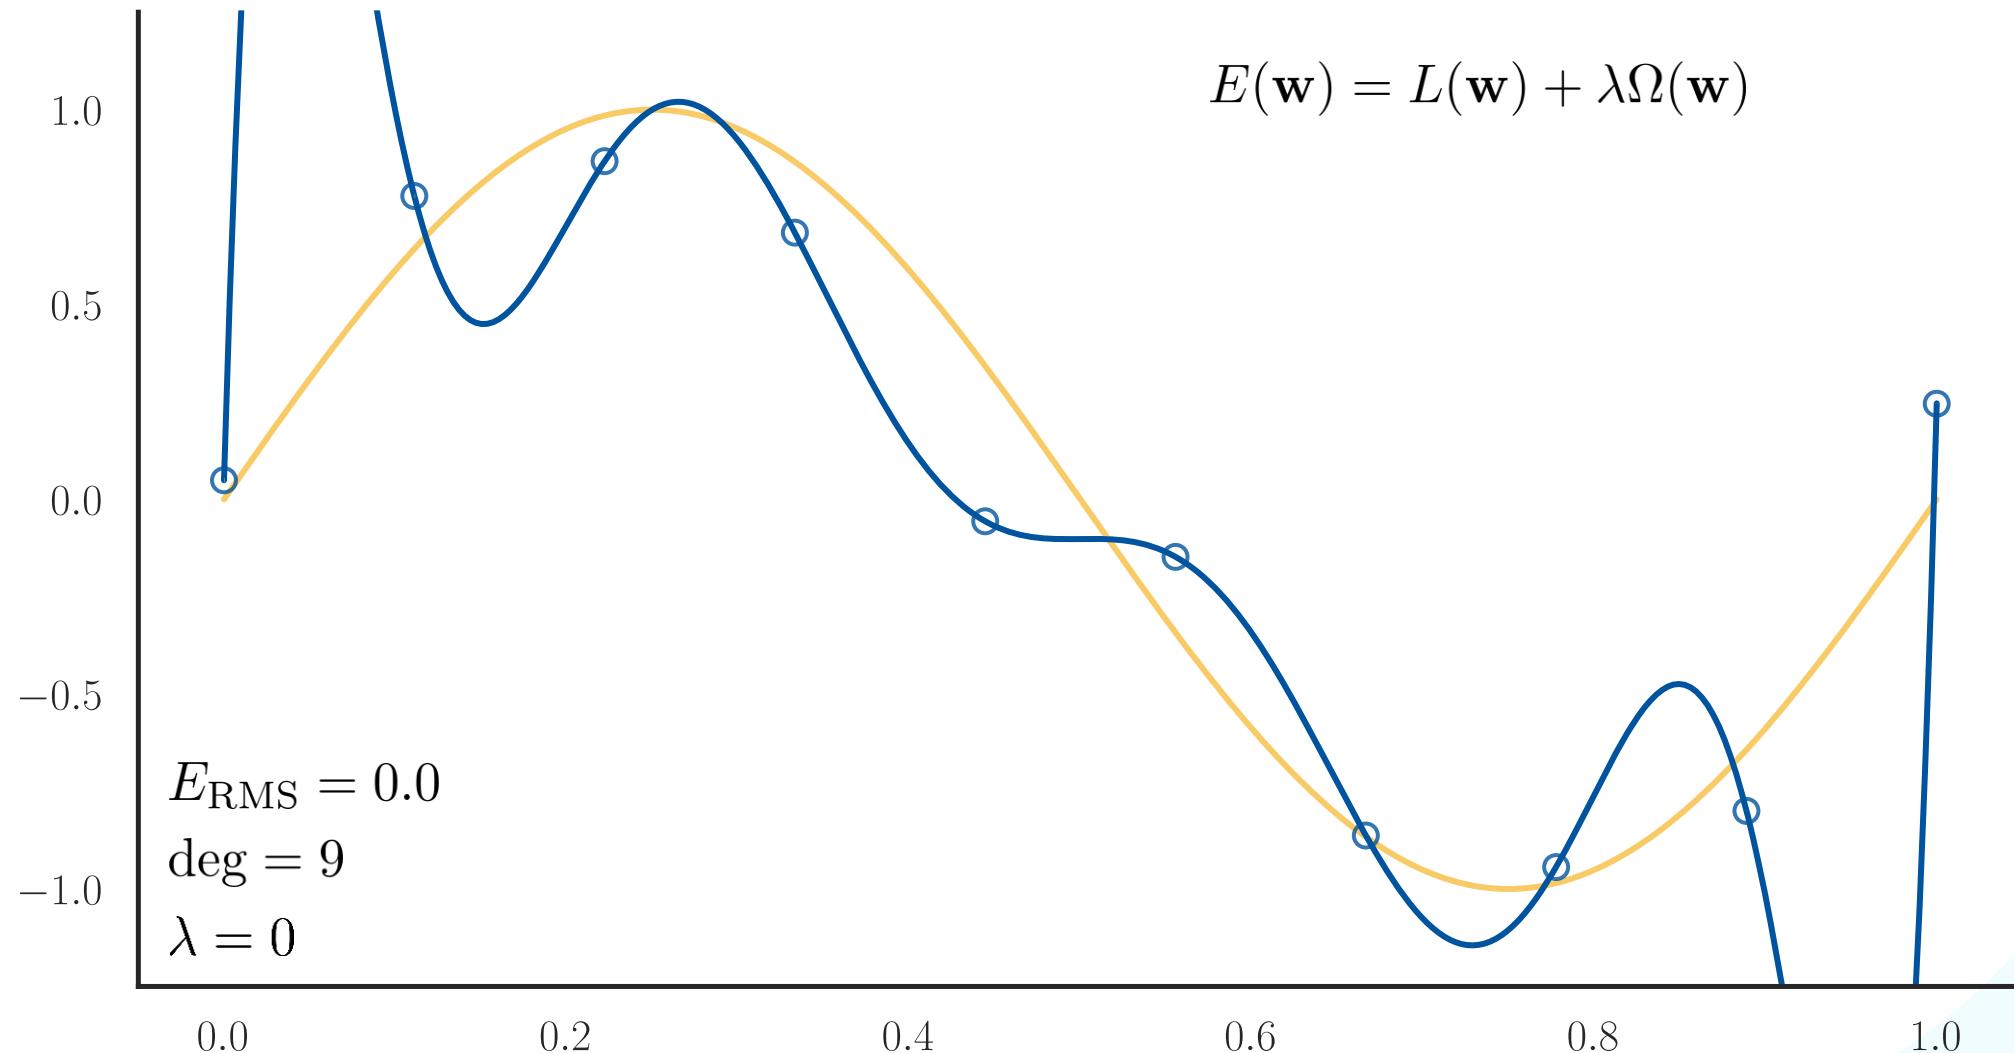
Example: Regularizing a Polynomial

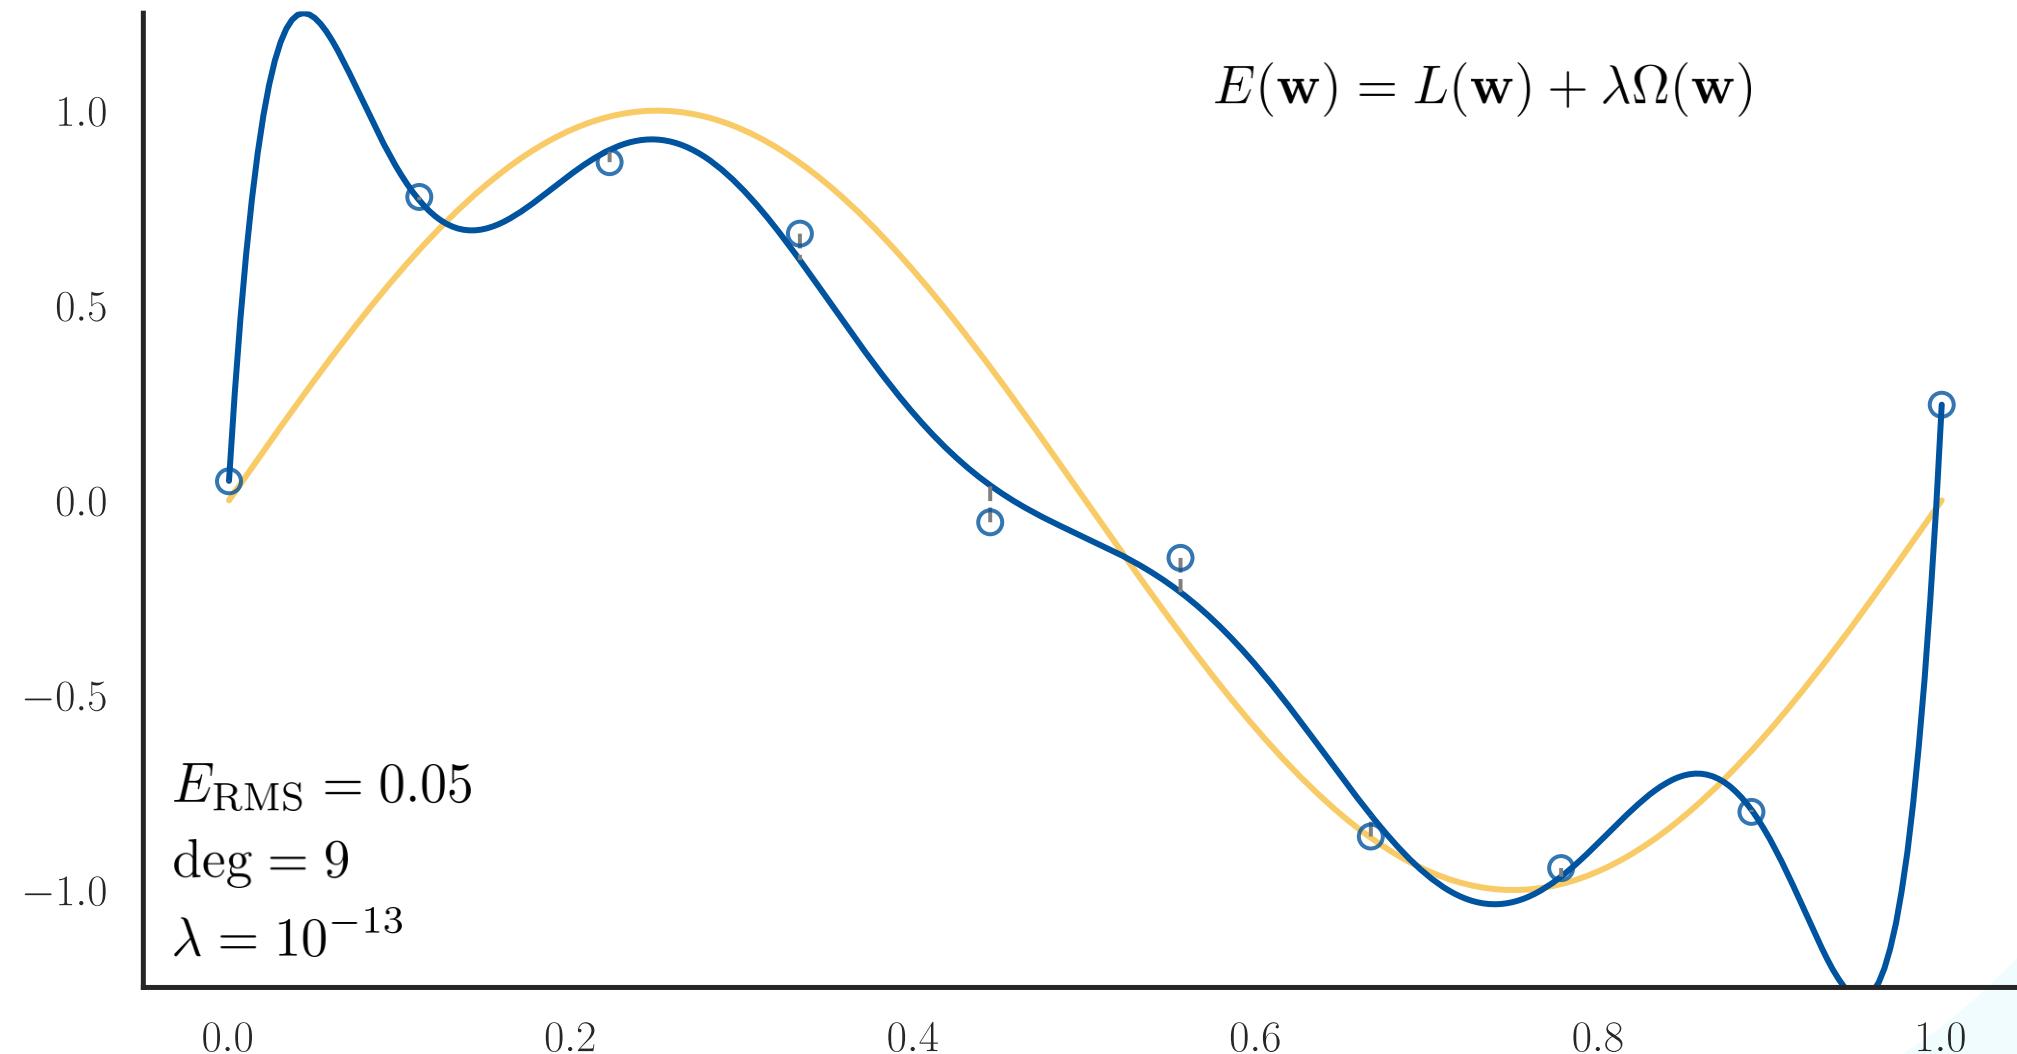
- Again, use polynomial basis functions:

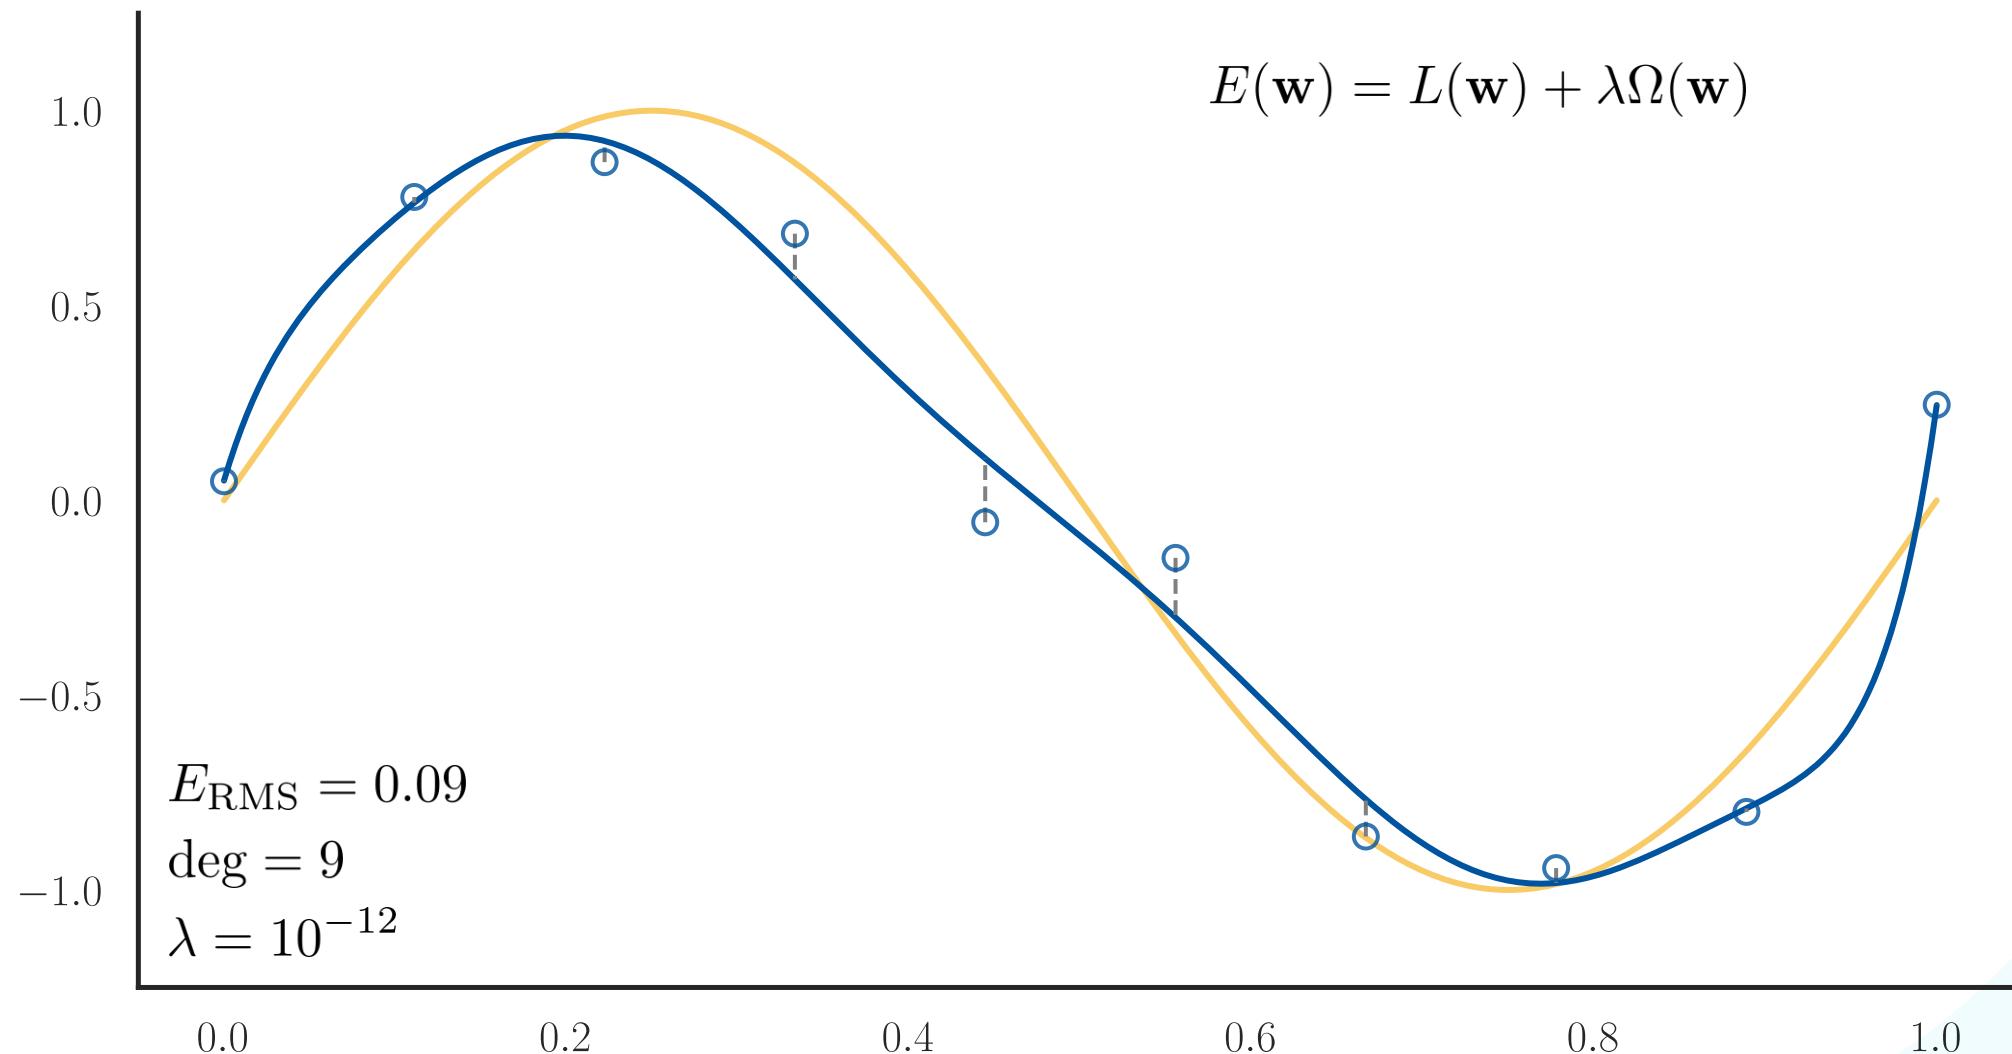
$$\phi_j(x) = (x^j)$$

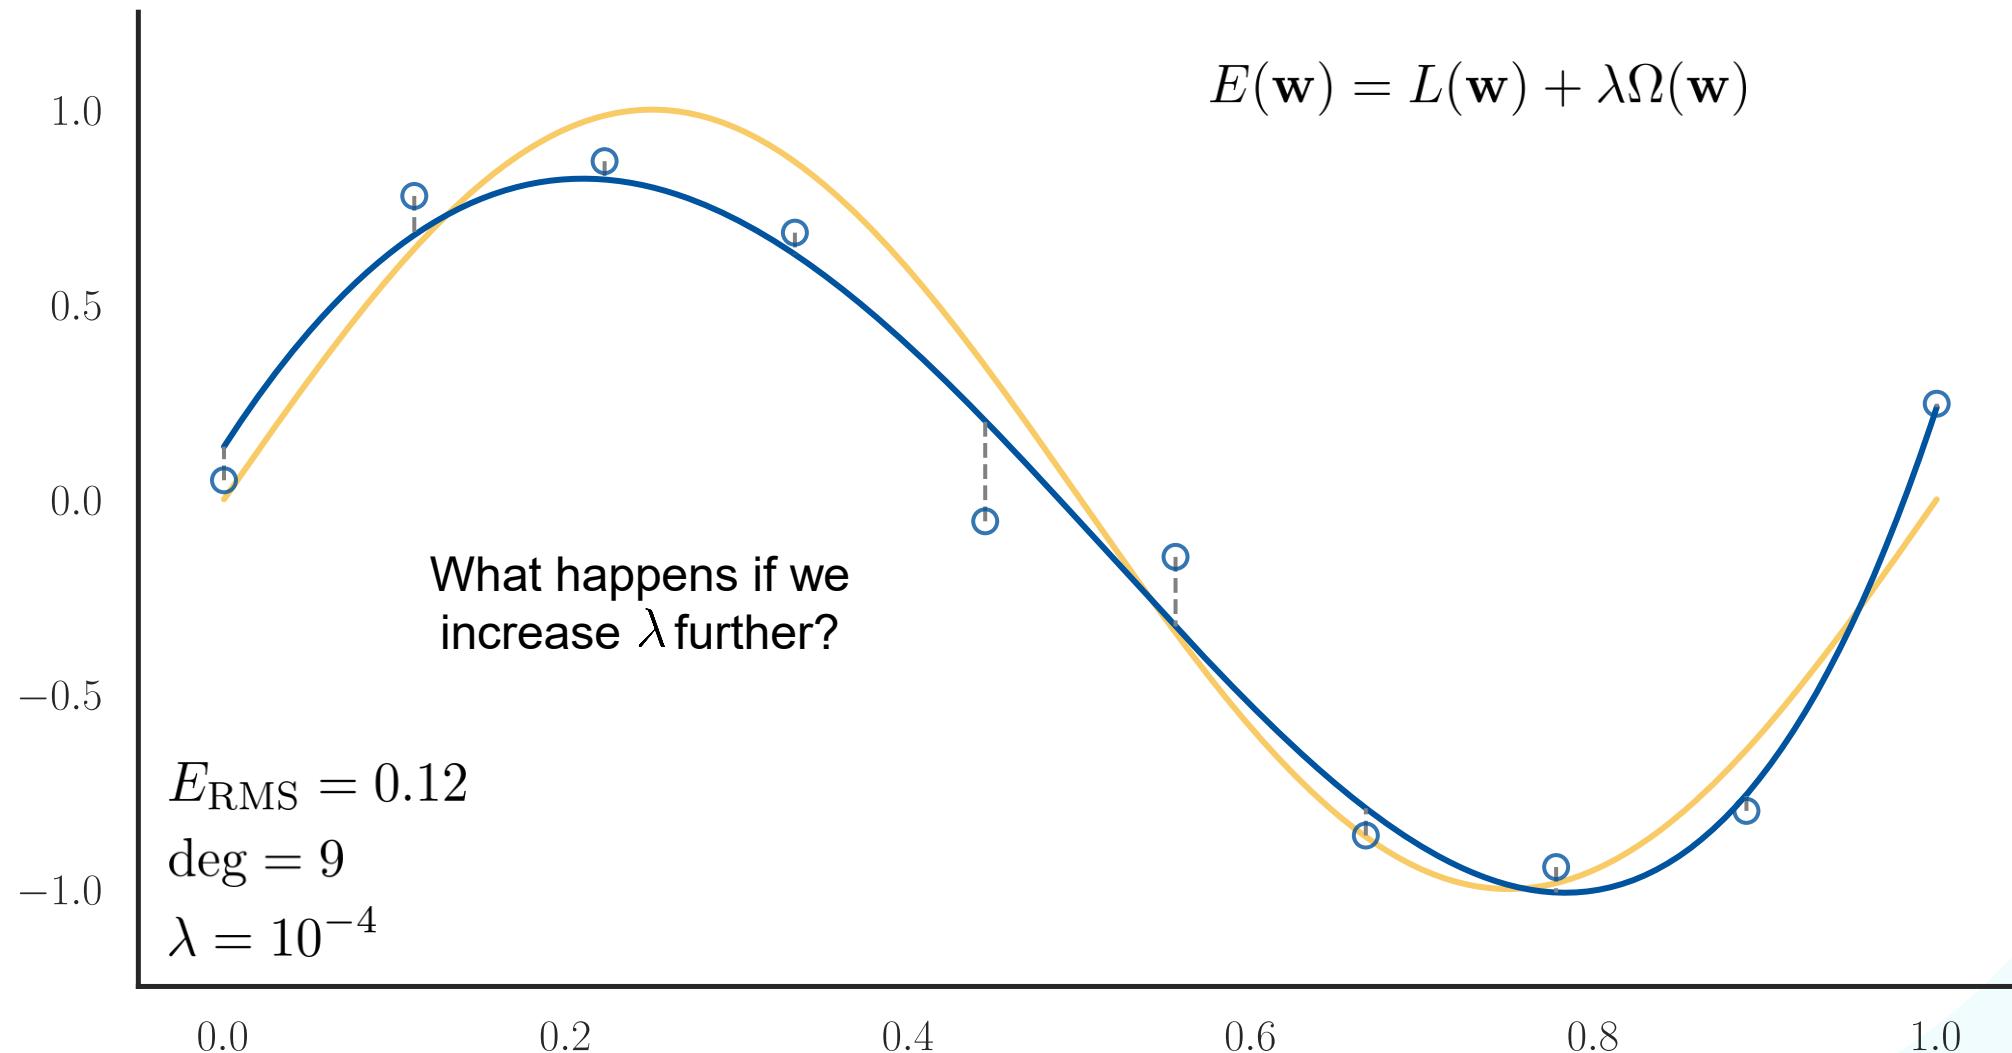
- Start off with an overfitting model.
- How much should we regularize?*

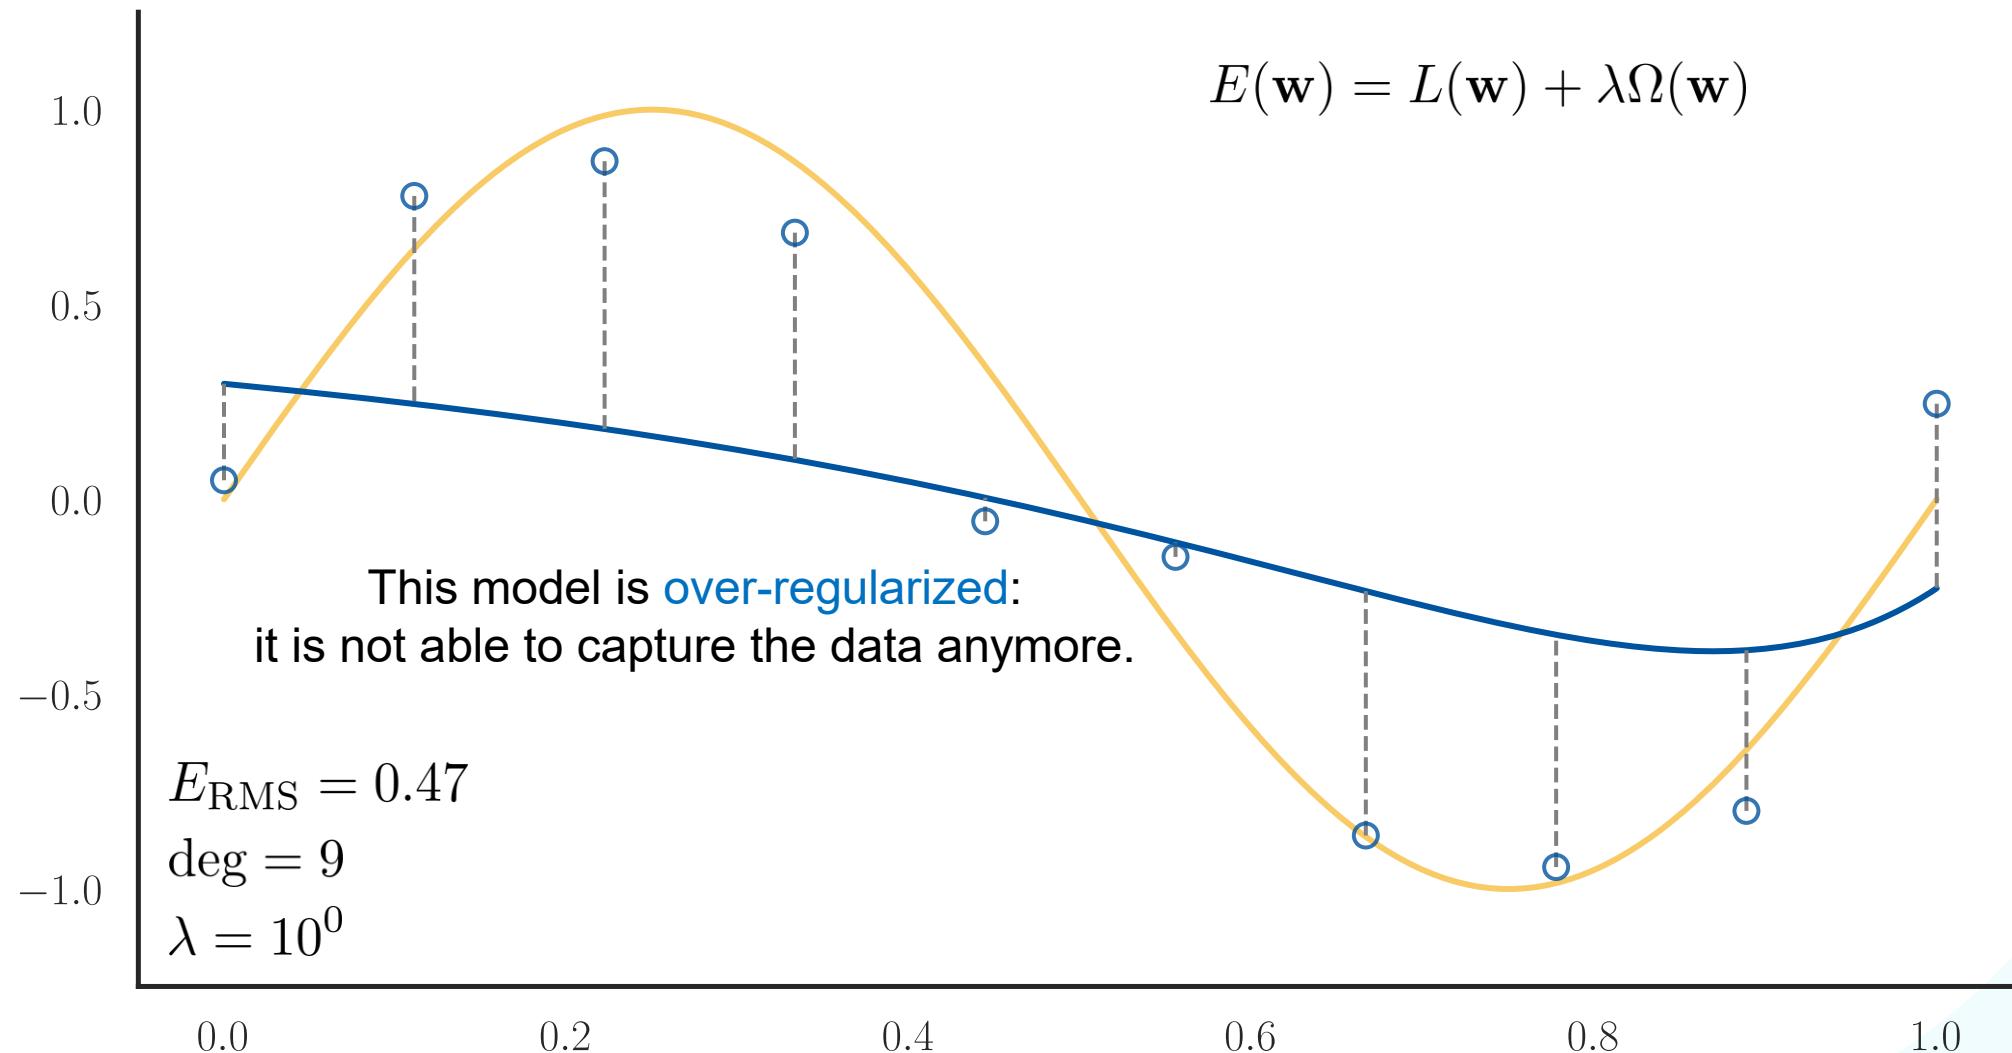






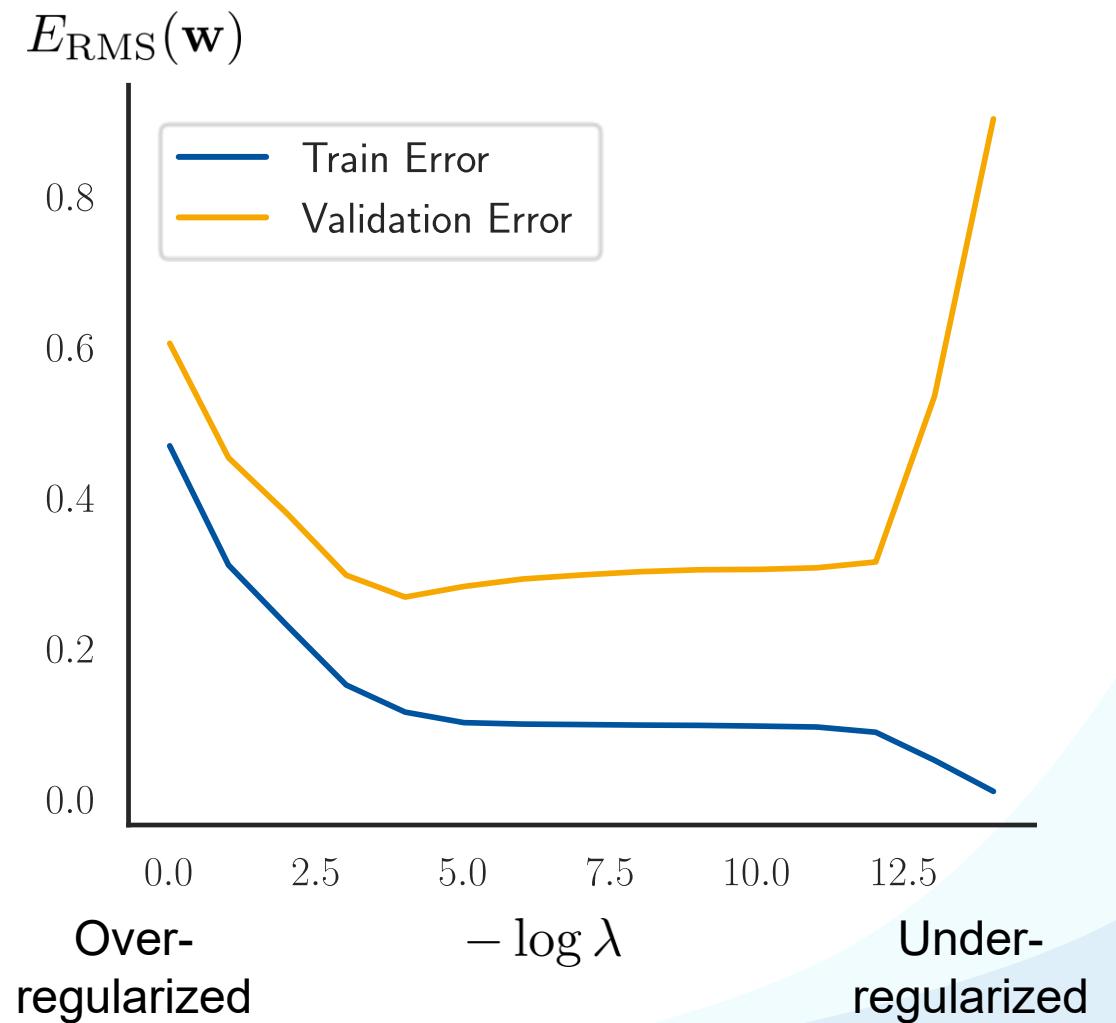






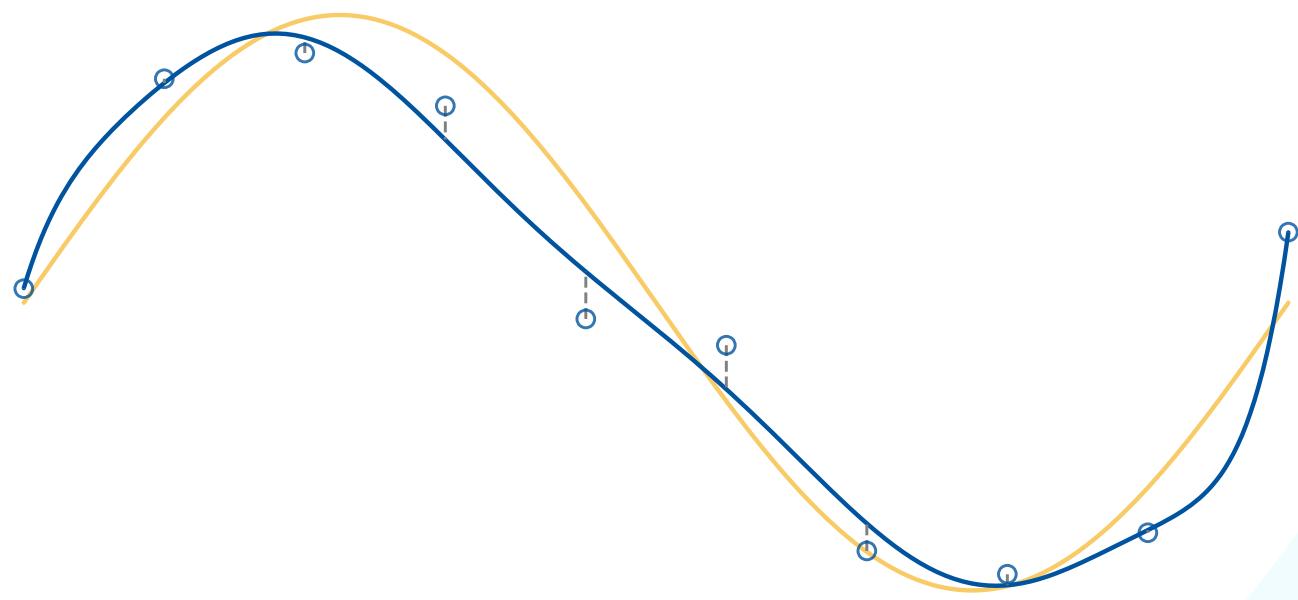
Choosing the right Regularization

- Regularization allows us to apply complex model on small datasets.
- However, we shifted the problem from selecting a suitable model to selecting a suitable regularization.
- The regularization factor λ becomes a [hyperparameter](#).



Linear Regression

1. Motivation
2. Least-Squares Regression
3. Regularization
4. **Ridge Regression**
5. The Bias-Variance Tradeoff



Ridge Regression

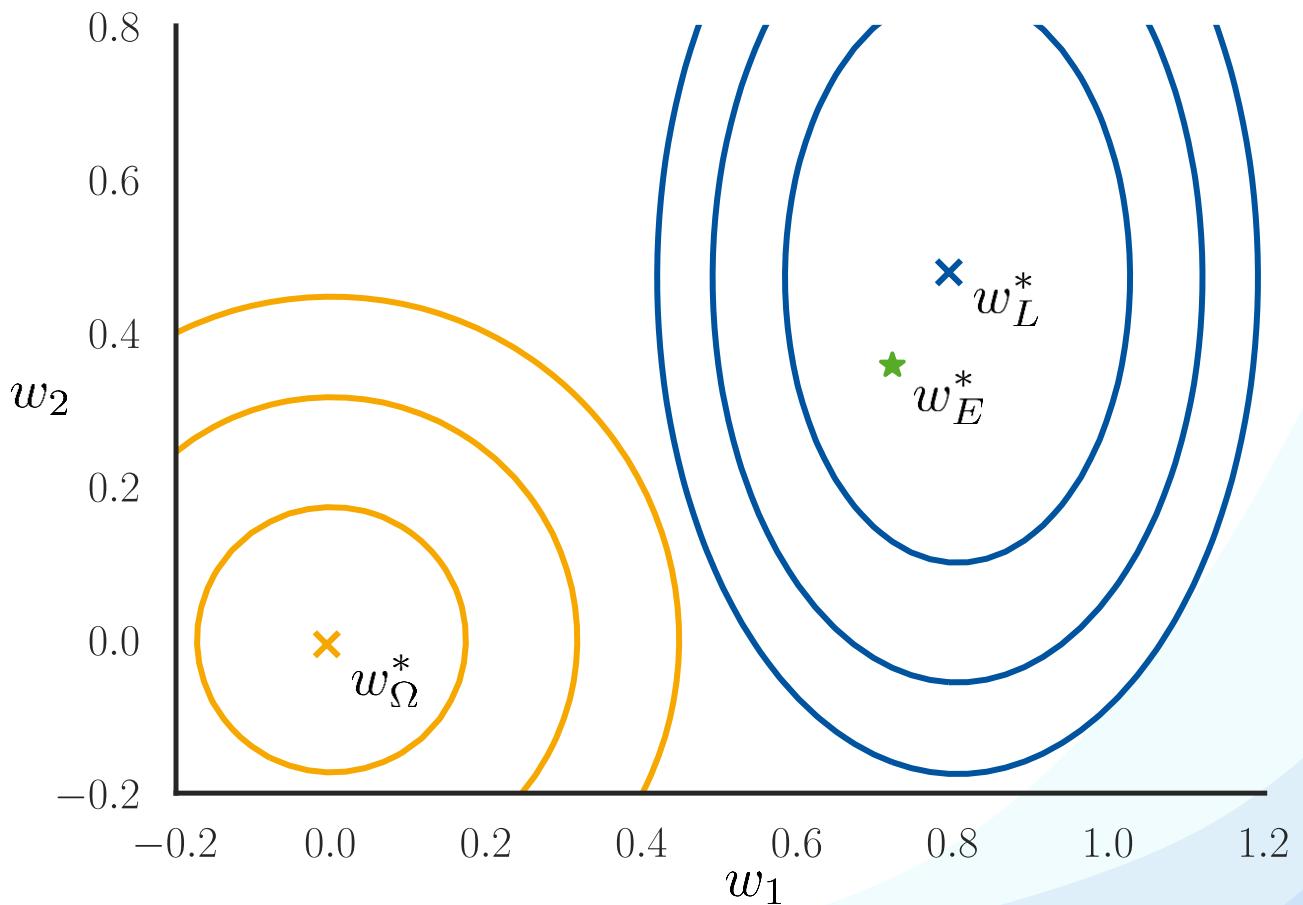
- We want to jointly minimize the squared error and the regularization term:

$$E(\mathbf{w}) = L(\mathbf{w}) + \lambda\Omega(\mathbf{w})$$

$$L(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (y(\mathbf{x}_n; \mathbf{w}) - t_n)^2$$

$$\Omega(\mathbf{w}) = \frac{1}{2} \|\mathbf{w}\|^2$$

- This model is called ridge regression.



Derivation

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (y(\mathbf{x}_n; \mathbf{w}) - t_n)^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

$$\begin{aligned}\nabla E(\mathbf{w}) &= \sum_{n=1}^N (\mathbf{w}^\top \boldsymbol{\phi}(\mathbf{x}_n) - t_n) \boldsymbol{\phi}(\mathbf{x}_n) + \lambda \mathbf{w} \stackrel{!}{=} 0 \\ \boldsymbol{\Phi}^\top (\boldsymbol{\Phi} \mathbf{w} - \mathbf{t}) + \lambda \mathbf{w} &= 0\end{aligned}$$

$$(\boldsymbol{\Phi}^\top \boldsymbol{\Phi} + \lambda \mathbf{I}) \mathbf{w} = \boldsymbol{\Phi}^\top \mathbf{t}$$

$$\mathbf{w} = (\boldsymbol{\Phi}^\top \boldsymbol{\Phi} + \lambda \mathbf{I})^{-1} \boldsymbol{\Phi}^\top \mathbf{t}$$



Effect of **regularization**: keeps the inverse well-conditioned.

References and Further Reading

- More information about [Linear Discriminants](#) is available in Chapter 4.1 of Bishop's book. For more information about [Linear Regression](#), read Chapter 3.1.

Christopher M. Bishop
Pattern Recognition and Machine Learning
Springer, 2006

