



Visual Computing  
Institute

RWTH AACHEN  
UNIVERSITY

# Elements of Machine Learning & Data Science

Winter semester 2025/26


## Lecture 13 – Logistic Regression

08.12.2025


Prof. Bastian Leibe

# Machine Learning Topics


- 8. Introduction to ML
- 9. Probability Density Estimation
- 10. Linear Discriminants
- 11. Linear Regression
- 12. Logistic Regression**
- 13. Support Vector Machines
- 14. Neural Network Basics




Logistic Regression  
Formulation



Parameter Efficiency




$E(\mathbf{w}) = -\ln p(\mathbf{t}|\mathbf{w})$   
Cross-Entropy Error



$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E(\mathbf{w})$   
Iterative Optimization

# Logistic Regression

- 1. Logistic Regression Formulation**
2. Motivation and Background
3. Iterative Optimization
4. First-Order Gradient Descent
5. Second-Order Gradient Descent
6. Error Function Analysis



# Motivation



- We have seen how to build probabilistic classifiers using Bayes' Theorem:

$$y_k(\mathbf{x}) = p(\mathcal{C}_k|\mathbf{x}) \propto p(\mathbf{x}|\mathcal{C}_k)p(\mathcal{C}_k)$$

- We have directly modeled the decision boundary with linear discriminants:

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

- In the following, we will combine those two ideas
  - We will model the posterior  $p(\mathcal{C}_k|\mathbf{x})$
  - But we will do that using a linear discriminant function  $y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$
- The resulting model will be called [logistic regression](#).



# Reminder: Probabilistic Classification

- Remember what we did in probabilistic classification
  - We modeled the likelihood of each class


$$p(\mathbf{x}|\mathcal{C}_k)$$

- We scaled the likelihoods with the priors

$$p(\mathbf{x}|\mathcal{C}_k)p(\mathcal{C}_k)$$

- We normalized to compute the posterior

$$p(\mathcal{C}_1|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)}$$



- Let's now start with the posterior and rewrite it



$$\begin{aligned}
 p(\mathcal{C}_1|\mathbf{x}) &= \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \\
 &= \frac{1}{1 + \frac{p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}} \\
 &= \frac{1}{1 + \exp(-a)} \quad =: \sigma(a)
 \end{aligned}$$

- This is the equation for the **logistic sigmoid** function  $\sigma(a)$

⇒ If we set

$$a = \ln \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)}$$

the logistic sigmoid expresses a posterior probability!



# Properties of the Logistic Sigmoid

- Definition:

$$\sigma(a) = \frac{1}{1 + \exp(-a)}$$

- Inverse (also known as **logit** function):


$$a = \ln\left(\frac{\sigma(a)}{1 - \sigma(a)}\right)$$

- Symmetry:

$$\sigma(-a) = 1 - \sigma(a)$$

- Derivative:

$$\frac{\partial \sigma(a)}{\partial a} = \sigma(a)(1 - \sigma(a))$$



# Logistic Regression

- We now define the **logistic regression** model
  - For the start, let us assume two classes  $\mathcal{C}_1, \mathcal{C}_2$ .
  - We model the class posteriors  $p(\mathcal{C}_k|\mathbf{x})$  as

$$p(\mathcal{C}_1|\mathbf{x}) = y(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x})$$



$$p(\mathcal{C}_2|\mathbf{x}) = 1 - p(\mathcal{C}_1|\mathbf{x})$$

**Logistic sigmoid**  
activation function:

$$\sigma(a) = \frac{1}{1 + \exp(-a)}$$

- I.e., we define a linear discriminant model

$$y(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x})$$

that is meant to represent the class posterior with  
the help of a logistic sigmoid activation function  $\sigma(a)$ .

- Our target labels are now  $t_n \in \{0, 1\}$ .

# Error Function

- Consider a data set  $\mathcal{D} = \{(\phi_1, t_1), \dots, (\phi_N, t_N)\}$ 
  - with data points  $\Phi = [\phi_1, \dots, \phi_N]^\top$ ,  $\phi_n = \phi(\mathbf{x}_n)$
  - And target labels  $\mathbf{t} = [t_1, \dots, t_N]^\top$ ,  $t_n \in \{0, 1\}$
- Maximum likelihood** approach
  - With  $y_n = p(\mathcal{C}_1 | \phi_n) = \sigma(\mathbf{w}^\top \phi_n)$
  - We model the probability of the target labels  $\mathbf{t}$  given our model parameters  $\mathbf{w}$  as

*Trick: use  $t_n$  as an indicator variable*

$$p(\mathbf{t} | \mathbf{w}) = \prod_{n=1}^N \begin{cases} p(\mathcal{C}_1 | \phi_n) & , t_n = 1 \\ p(\mathcal{C}_2 | \phi_n) & , t_n = 0 \end{cases} = \prod_{n=1}^N \begin{cases} y_n & , t_n = 1 \\ (1 - y_n) & , t_n = 0 \end{cases} = \prod_{n=1}^N y_n^{t_n} (1 - y_n)^{1-t_n}$$

- Maximum likelihood approach

$$p(\mathbf{t}|\mathbf{w}) = \prod_{n=1}^N y_n^{t_n} (1 - y_n)^{1-t_n}$$

- Define the error function as the negative log-likelihood

$$E(\mathbf{w}) = -\ln p(\mathbf{t}|\mathbf{w})$$

$$= -\sum_{n=1}^N (t_n \ln y_n + (1 - t_n) \ln(1 - y_n))$$

- This function is known as the binary cross-entropy error.

# Softmax Regression

- Multi-class extension of logistic regression
  - Generalization to  $K$  classes with target labels in 1-of- $K$  notation  $\mathbf{t}_n = [0, 1, \dots, 0]^T$
  - Again, we define a linear discriminant function that models the class posteriors

$$\mathbf{y}(\mathbf{x}; \mathbf{w}) = \begin{bmatrix} p(\mathcal{C}_1 | \mathbf{x}; \mathbf{w}) \\ p(\mathcal{C}_2 | \mathbf{x}; \mathbf{w}) \\ \vdots \\ p(\mathcal{C}_K | \mathbf{x}; \mathbf{w}) \end{bmatrix} = \frac{1}{\sum_{j=1}^K \exp(\mathbf{w}_j^\top \mathbf{x})} \begin{bmatrix} \exp(\mathbf{w}_1^\top \mathbf{x}) \\ \exp(\mathbf{w}_2^\top \mathbf{x}) \\ \vdots \\ \exp(\mathbf{w}_K^\top \mathbf{x}) \end{bmatrix}$$

- This makes use of the softmax function as a multi-class extension of the logistic sigmoid

$$\text{softmax}(\mathbf{a}) = \frac{\exp(a_k)}{\sum_{j=1}^K \exp(a_j)}$$

- We can write the **binary cross-entropy error** as

$$\begin{aligned} E(\mathbf{w}) &= - \sum_{n=1}^N (t_n \ln y_n + (1 - t_n) \ln(1 - y_n)) \\ &= - \sum_{n=1}^N \sum_{k=0}^1 (\mathbb{I}(t_n = k) \ln p(\mathcal{C}_k | \mathbf{x}_n; \mathbf{w})) \end{aligned}$$

**indicator function**

$$\mathbb{I}(\varphi) = \begin{cases} 1 & \text{if } \varphi \text{ is true} \\ 0 & \text{otherwise} \end{cases}$$


- Using one-hot labels  $\mathbf{t}_n$ , the generalization to  $K$  classes is:

$$E(\mathbf{w}) = - \sum_{n=1}^N \sum_{k=1}^K \left( \mathbb{I}(t_{kn} = 1) \ln \frac{\exp(\mathbf{w}_k^\top \mathbf{x})}{\sum_{j=1}^K \exp(\mathbf{w}_j^\top \mathbf{x})} \right)$$

- This function is known as the **multi-class cross-entropy error** or **softmax cross-entropy error**.

# Logistic Regression

1. Logistic Regression Formulation
2. **Motivation and Background**
3. Iterative Optimization
4. First-Order Gradient Descent
5. Second-Order Gradient Descent
6. Error Function Analysis




# Motivation: Why Logistic Regression?

- Logistic Regression uses models of the form

$$p(\mathcal{C}_1|\phi) = y(\phi) = \sigma(\mathbf{w}^T \phi)$$

$$p(\mathcal{C}_2|\phi) = 1 - p(\mathcal{C}_1|\phi)$$




- Interpretation

- We model the **class posteriors**  $p(\mathcal{C}_k|\phi)$ , as required to make Bayes optimal decisions.
  - We have seen previously that we can obtain  $p(\mathcal{C}_k|\phi) = p(\phi|\mathcal{C}_k)p(\mathcal{C}_k)$ .
  - However, here we model  $p(\mathcal{C}_k|\phi)$  as a **linear discriminant function**  $y(\phi) = \sigma(\mathbf{w}^T \phi)$  instead.

- *Why should we do this?*


- *What advantage does such a model have compared to direct modeling of the probabilities?*

## Example




Let's assume the  $p(\phi|\mathcal{C}_k)$  are modeled using Gaussians with equal covariances.

## Example



Let's assume the  $p(\phi|\mathcal{C}_k)$  are modeled using Gaussians with equal covariances.

## Example




Let's assume the  $p(\phi|\mathcal{C}_k)$  are modeled using Gaussians with equal covariances.

⇒ The decision boundary between them will be linear!

# Parameter Efficiency

- #Parameters needed for generative models:
    - Assuming an  $M$ -dimensional feature space
    - Prior  $p(\mathcal{C}_1)$  1
    - Means  $\mu_1, \mu_2$   $2M$
    - Covariances  $\Sigma$   $M(M + 1)/2$
- $$\Rightarrow \text{Total} \quad M(M + 5)/2 + 1$$
- 
- #Parameters needed for logistic regression:
    - Weights  $w$   $M$

$\Rightarrow$  For large  $M$ , logistic regression has clear advantages!



# Discussion: Logistic Regression

## Advantages


- Nice probabilistic interpretation, directly represents the posterior.
- Requires fewer parameters than modeling the likelihood + prior.
- Cross-Entropy error is convex: unique minimum exists.
- More robust than least-squares.

## Limitations

- No closed-form solution, requires iterative optimization approach.

# Logistic Regression


1. Logistic Regression Formulation
2. Motivation and Background
- 3. Iterative Optimization**
4. First-Order Gradient Descent
5. Second-Order Gradient Descent
6. Error Function Analysis



# Iterative Optimization


- In general, generalized linear discriminants with nonlinear activation and/or basis functions can no longer be optimized in closed form.
- Instead, we use iterative optimization schemes.
- Here: **Gradient Descent**.
  - Start with initial guess for parameter values.
  - Move towards a minimum of the error function by following the direction of steepest descent.
  - Iterate until convergence

$$y_k(\mathbf{x}) = g \left( \sum_{j=0}^M w_{kj} \phi_j(\mathbf{x}) \right) = g(\mathbf{w}^\top \phi(\mathbf{x}))$$



## Idea: Gradient Descent

- Start with an initial guess of parameter values  $w_{kj}^{(0)}$ .




# Idea: Gradient Descent

- Start with an initial guess of parameter values  $w_{kj}^{(0)}$ .
- Follow the gradient to move to a (local) minimum:

$$w_{kj}^{(\tau+1)} = w_{kj}^{(\tau)} - \eta \frac{\partial E(\mathbf{w})}{\partial w_{kj}} \Big|_{\mathbf{w}^{(\tau)}}$$

- $\eta$  is called the **learning rate**.
- This corresponds to a 1<sup>st</sup>-order Taylor expansion.
  - I.e., we approximate the error function by its tangent plane around the current point  $\mathbf{w}^{(\tau)}$ .
- Repeat this procedure for a number of steps.




# Idea: Gradient Descent

- Start with an initial guess of parameter values  $w_{kj}^{(0)}$ .
- Follow the gradient to move to a (local) minimum:

$$w_{kj}^{(\tau+1)} = w_{kj}^{(\tau)} - \eta \frac{\partial E(\mathbf{w})}{\partial w_{kj}} \Big|_{\mathbf{w}^{(\tau)}}$$

- $\eta$  is called the **learning rate**.
- This corresponds to a 1<sup>st</sup>-order Taylor expansion.
  - I.e., we approximate the error function by its tangent plane around the current point  $\mathbf{w}^{(\tau)}$ .
- Repeat this procedure for a number of steps.




## Idea: Gradient Descent

- Start with an initial guess of parameter values  $w_{kj}^{(0)}$ .
- Follow the gradient to move to a (local) minimum:

$$w_{kj}^{(\tau+1)} = w_{kj}^{(\tau)} - \eta \frac{\partial E(\mathbf{w})}{\partial w_{kj}} \Big|_{\mathbf{w}^{(\tau)}}$$

- $\eta$  is called the **learning rate**.
- This corresponds to a 1<sup>st</sup>-order Taylor expansion.
  - I.e., we approximate the error function by its tangent plane around the current point  $\mathbf{w}^{(\tau)}$ .
- Repeat this procedure for a number of steps.




## Idea: Gradient Descent

- Start with an initial guess of parameter values  $w_{kj}^{(0)}$ .
- Follow the gradient to move to a (local) minimum:

$$w_{kj}^{(\tau+1)} = w_{kj}^{(\tau)} - \eta \frac{\partial E(\mathbf{w})}{\partial w_{kj}} \Big|_{\mathbf{w}^{(\tau)}}$$

- $\eta$  is called the **learning rate**.
- This corresponds to a 1<sup>st</sup>-order Taylor expansion.
  - I.e., we approximate the error function by its tangent plane around the current point  $\mathbf{w}^{(\tau)}$ .
- Repeat this procedure for a number of steps.



# Discussion: Gradient Descent

## Advantages


- Simple approach for iterative optimization.
- Approximates the error function by its tangent plane around the current point in order to find the direction of steepest descent.

## Limitations

- Local optimization. Unless the error function is convex, will only converge to a local optimum.
- Relatively slow convergence (can be improved by second-order approaches).
- In practice, finding a good step size ([learning rate](#)) is important for fast convergence.

# Logistic Regression

1. Logistic Regression Formulation
2. Motivation and Background
3. Iterative Optimization
- 4. First-Order Gradient Descent**
5. Second-Order Gradient Descent
6. Error Function Analysis



# First-order Optimization

- Logistic regression uses the **binary cross-entropy error**:

$$E(\mathbf{w}) = - \sum_{n=1}^N (t_n \ln y(\mathbf{x}_n; \mathbf{w}) + (1 - t_n) \ln(1 - y(\mathbf{x}_n; \mathbf{w})))$$

- Properties
  - Convex function, so it has a unique minimum
  - But no closed-form solution
- We need to use iterative methods for optimization
  - Let's try (first-order) gradient descent:

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E(\mathbf{w})$$

# Gradient of the Cross-Entropy Error

$$E(\mathbf{w}) = - \sum_{n=1}^N (t_n \ln y_n + (1 - t_n) \ln(1 - y_n))$$

$$\begin{aligned}\nabla E(\mathbf{w}) &= - \sum_{n=1}^N \left( t_n \frac{\partial}{\partial \mathbf{w}} y_n + (1 - t_n) \frac{\partial}{\partial \mathbf{w}} (1 - y_n) \right) \\ &= - \sum_{n=1}^N \left( t_n \frac{y_n(1 - y_n)}{y_n} \phi_n + (1 - t_n) \frac{y_n(1 - y_n)}{(1 - y_n)} \phi_n \right) \\ &= - \sum_{n=1}^N ((t_n - t_n y_n - y_n + t_n y_n) \phi_n) \\ &= \sum_{n=1}^N (y_n - t_n) \phi_n\end{aligned}$$

$$y_n = y(\mathbf{x}_n; \mathbf{w})$$

$$\begin{aligned}\sigma'(a) &= \sigma(a)(1 - \sigma(a)) \\ \frac{\partial y_n}{\partial \mathbf{w}} &= y_n(1 - y_n) \phi_n\end{aligned}$$

$$\phi_n = \phi(\mathbf{x}_n)$$

- The gradient for logistic regression is

$$\nabla E(\mathbf{w}) = \sum_{n=1}^N (y_n - t_n) \phi_n$$

- We can plug this into gradient descent:


$$\begin{aligned}\mathbf{w}^{(\tau+1)} &= \mathbf{w}^{(\tau)} - \eta \nabla E(\mathbf{w}) \\ &= \mathbf{w}^{(\tau)} - \eta \sum_{n=1}^N (y_n - t_n) \phi_n\end{aligned}$$

*How should we choose  
the learning rate?*

- This update rule is known as the **Delta rule** (= LMS rule)
  - Simply feed back the input data points, weighted by the classification error.*


## Effects of the learning rate

$\eta$  too small




Convergence is slow

$\eta_{\text{opt}}$




Converges ideally in  
a single step

$\eta$  too large




Might not converge

## Example: Logistic Regression with Gradient Descent




Step 1

$$E(w) = 1.3682$$




Step 2

$$E(w) = 1.0947$$




Step 3

$$E(w) = 0.8822$$




Step 4


$$E(w) = 0.7210$$



Step 30


$$E(w) = 0.1514$$





Step 90

$$E(w) = 0.0777$$



# Discussion: Logistic Regression with Gradient Descent

## Advantages


- Simple iterative optimization scheme with a familiar update rule ([Delta rule](#)).

## Limitations

- Slow convergence
- Need to choose a suitable learning rate.

# Logistic Regression

1. Logistic Regression Formulation
2. Motivation and Background
3. Iterative Optimization
4. First-Order Gradient Descent
- 5. Second-Order Gradient Descent**
6. Error Function Analysis




## Second-Order Optimization

- So far, we have optimized the cross-entropy error with gradient descent:

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E(\mathbf{w})$$

- This is a first-order approximation, and it heavily depends on the learning rate  $\eta$ .

- Instead, we can apply a second-order optimization scheme that converges faster and is independent of the learning rate.



# Newton-Raphson Gradient Descent

- Second-order Newton-Raphson update scheme:


$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \mathbf{H}^{-1} \nabla E(\mathbf{w})$$

- Here,  $\mathbf{H} = \nabla \nabla E(\mathbf{w})$  is the Hessian matrix, i.e., the matrix of second derivatives:


$$\mathbf{H}_{ij} = \frac{\partial^2 E(\mathbf{w})}{\partial w_i \partial w_j}$$

- Properties
  - Local quadratic approximation
  - Much faster convergence by taking into account the curvature of the error function.

# Intuition




First-Order  
 $\eta = 0.005$




Second-Order

# Intuition




First-Order  
 $\eta = 0.005$




Second-Order

# Intuition




First-Order  
 $\eta = 0.005$




Second-Order

# Intuition



First-Order  
 $\eta = 0.005$



Second-Order


# Intuition


*First-Order needs  
another 16 steps...*


First-Order  
 $\eta = 0.005$

Second-Order

## Example: Logistic Regression with 2<sup>nd</sup>-order Optimization (IRLS)









Step 6

$$E(w) = 0.0544$$



Step 12

$$E(w) = 0.0001$$



# Discussion: Second-Order Optimization

## Advantages


- Faster convergence than first-order methods

## Limitations

- Second-order approach, relies on computing second derivatives.
- Computing (and inverting) the Hessian matrix is expensive for problems with many parameters.


# Logistic Regression

1. Logistic Regression Formulation
2. Motivation and Background
3. Iterative Estimation
4. First-Order Gradient Descent
5. Second-Order Gradient Descent
6. **Error Function Analysis**




# Error Function Analysis


- We have seen how to learn **generalized linear discriminant** models by optimizing an error function.
  - We observed problems with **least-squares classification** based on the squared error function.
  - We have seen that **logistic regression** behaves more robustly.
- *Let's analyze the cross-entropy error in more detail...*




# Error Contribution Plot




# Ideal Misclassification Error




# Squared Error



## Squared Error on Tanh



# Cross-Entropy Error



## Discussion: Cross-Entropy Error

### Advantages

- Minimizer of this error corresponds to class posteriors
- Convex function, unique minimum exists
- Robust to outliers

### Limitations

- No closed-form solution, requires iterative estimation

# References and Further Reading

- More information about [Logistic Regression](#) is available in Chapter 4.3 of Bishop's book.

Christopher M. Bishop  
Pattern Recognition and Machine Learning  
Springer, 2006

