

Elements of Machine Learning & Data Science

Winter semester 2025/26

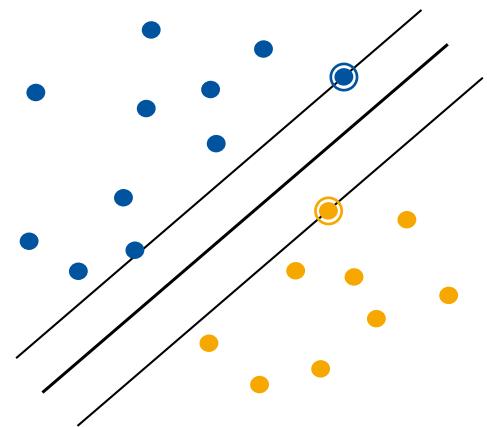
Lecture 14 – Support Vector Machines I

09.12.2025

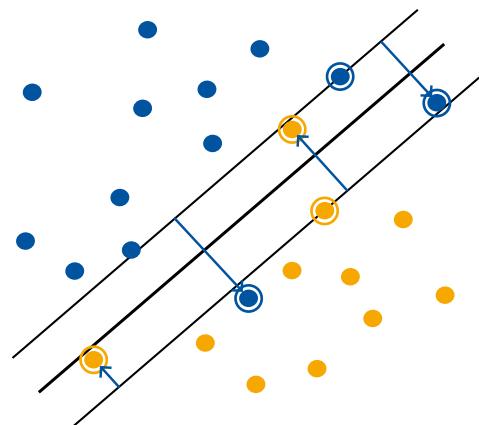
Prof. Bastian Leibe

Machine Learning Topics

- 8. Introduction to ML
- 9. Probability Density Estimation
- 10. Linear Discriminants
- 11. Linear Regression
- 12. Logistic Regression
- 13. Support Vector Machines**
- 14. Neural Network Basics



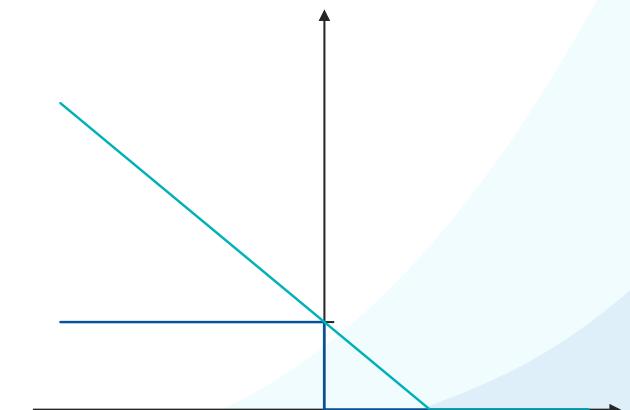
Maximum Margin Classification



Soft-Margin SVM

$$L_p(\mathbf{w}, b, \mathbf{a})$$
$$L_d(\mathbf{a})$$

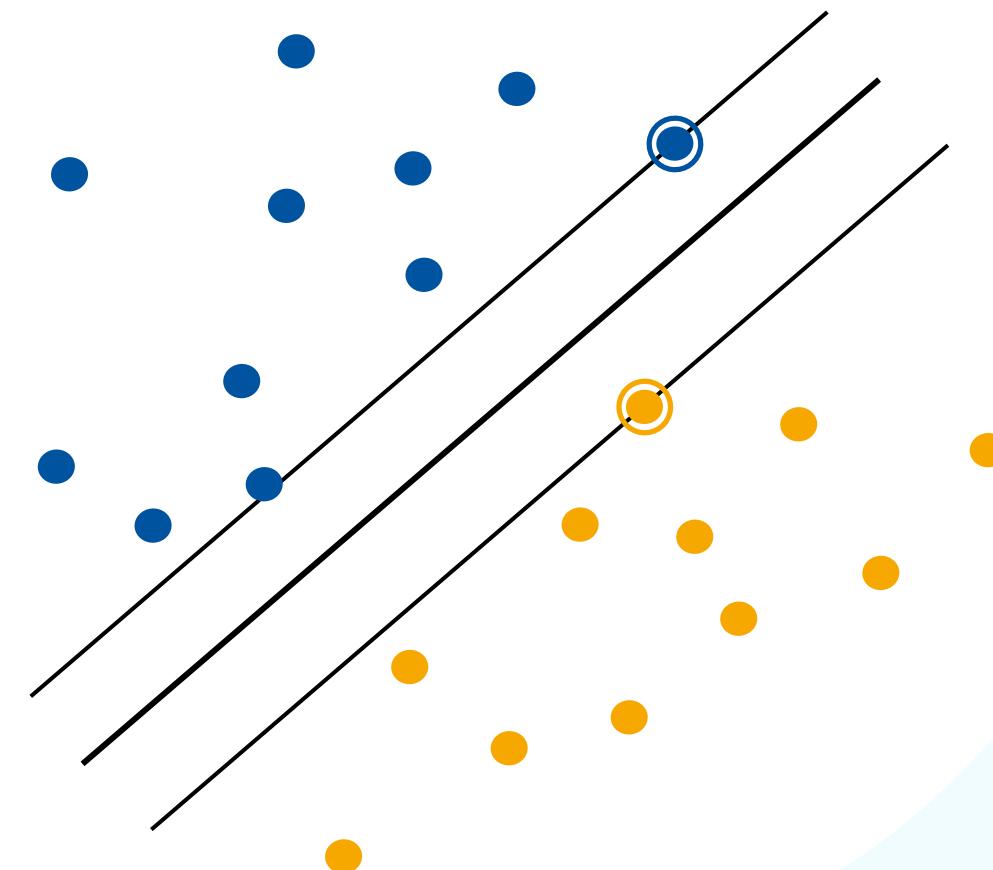
Primal & Dual Form



Hinge Loss

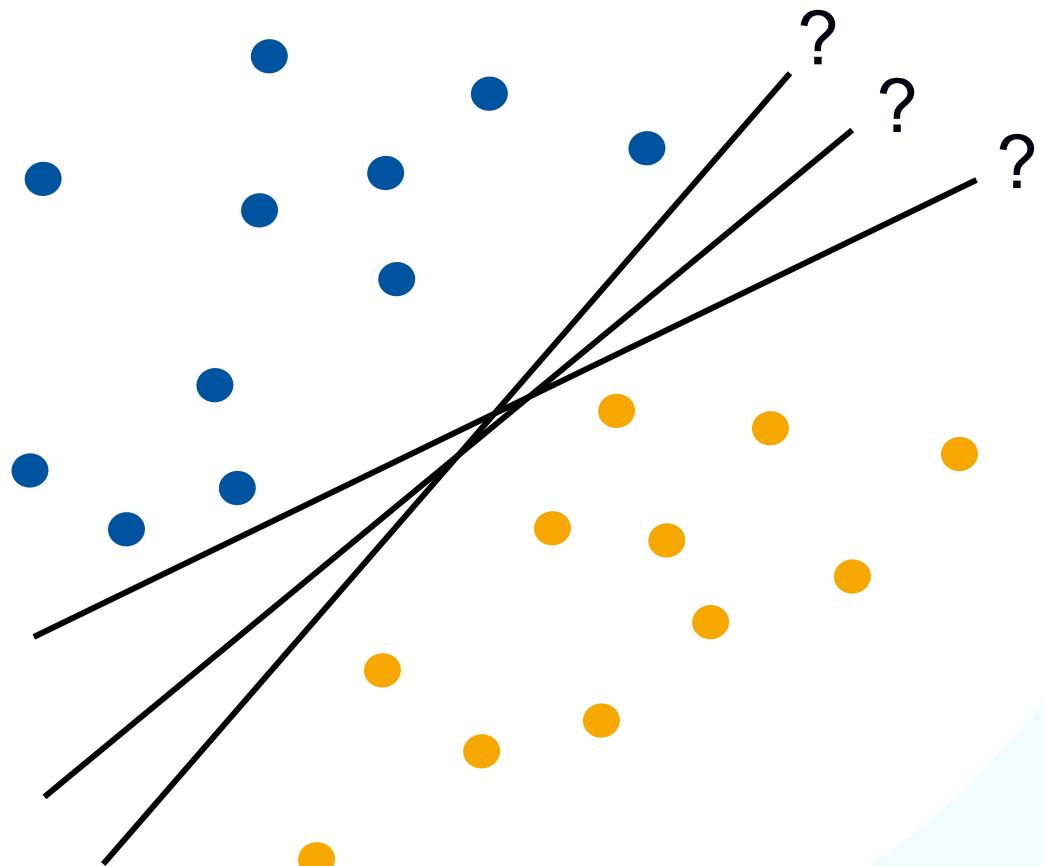
Support Vector Machines

1. Maximum Margin Classification
2. Primal Formulation
3. Dual Formulation
4. Soft-Margin SVMs
5. Non-linear SVMs
6. Error Function Analysis



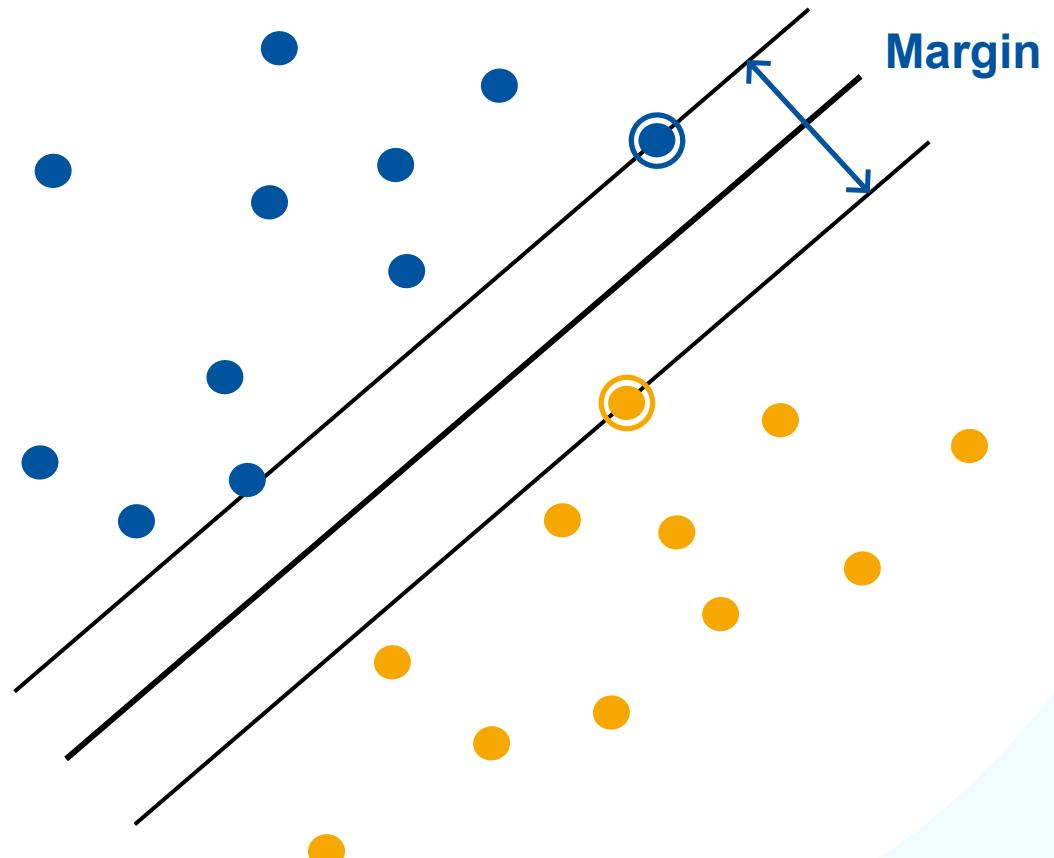
Maximum Margin Classification

- Overfitting is often a problem with linearly separable data
 - Which of the many possible decision boundaries is correct?
 - All of them have zero error on the training set...
 - However, they will perform differently on novel test data.
- *How can we select the classifier with the best generalization performance?*



Maximum Margin Classification

- Intuitively, we want to choose the classifier which leaves maximal “safety room” for future data points.
- This classifier has the largest **margin** between positive and negative points.
- It can be shown: The larger the margin, the lower the capacity for overfitting.

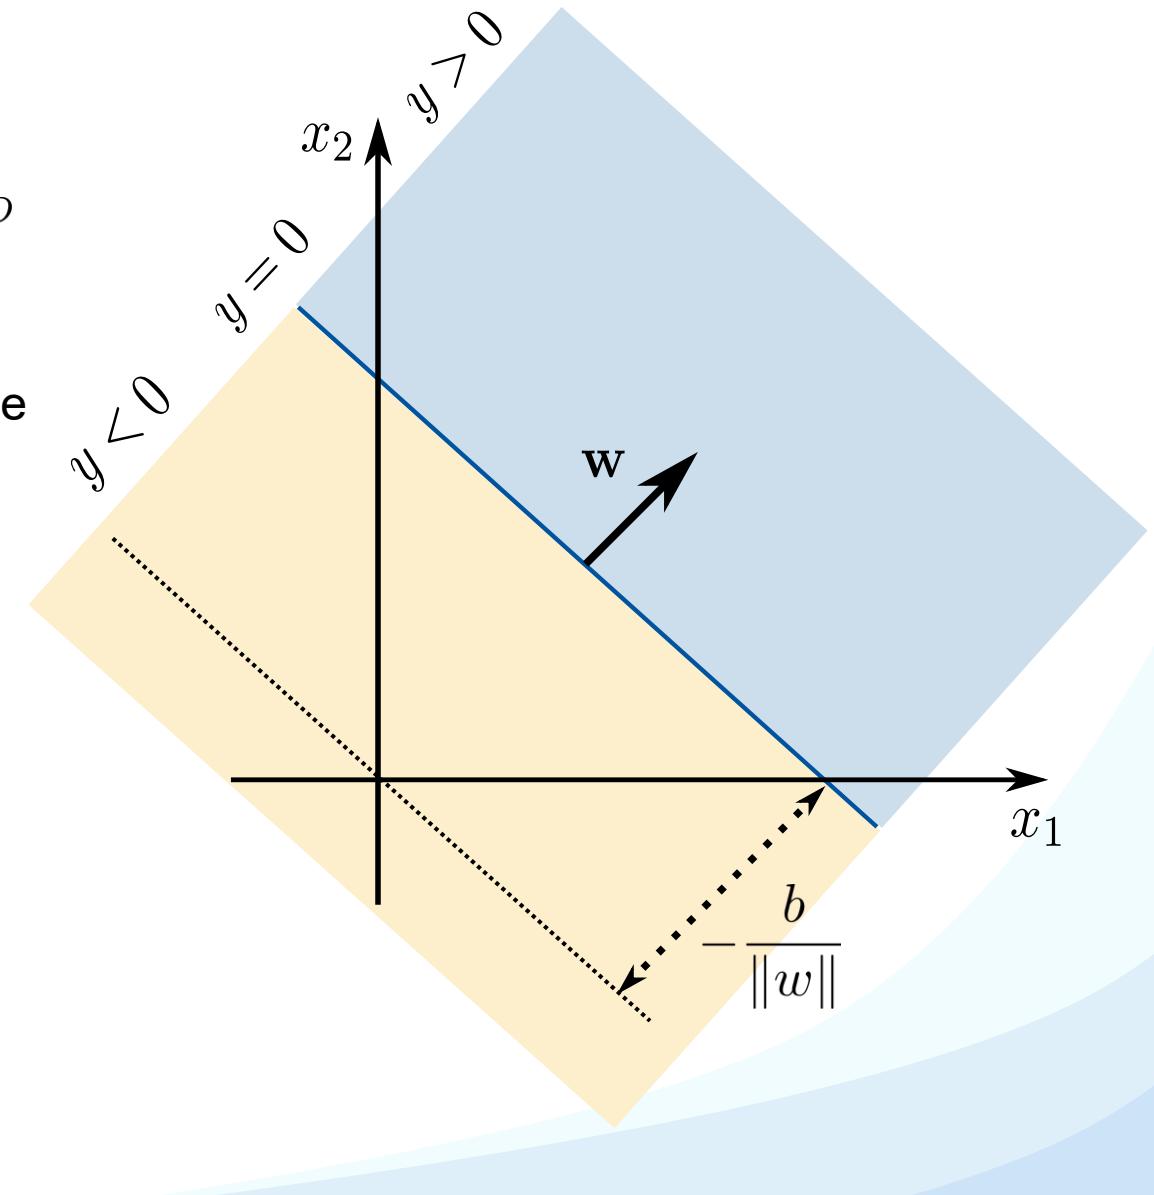


Intuition

- Let's first consider linearly separable data:
 - N training data points $\{(\mathbf{x}_i, t_i)\}_{i=1}^N$, $\mathbf{x}_i \in \mathbb{R}^D$
 - Binary labels $t_i \in \{-1, 1\}$
- A linear discriminant function models a hyperplane separating the data:

$$y(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + b$$

- Note that we denote the bias explicitly with b .*
- Decision rule
 - Decide for \mathcal{C}_1 if $y(\mathbf{x}) > 0$, else for \mathcal{C}_2 .



Support Vector Machines

- Assuming linearly separable data, we can always find a hyperplane with

$$\mathbf{w}^T \mathbf{x}_n + b \geq +1 \text{ for } t_n = +1$$

$$\mathbf{w}^T \mathbf{x}_n + b \leq -1 \text{ for } t_n = -1$$

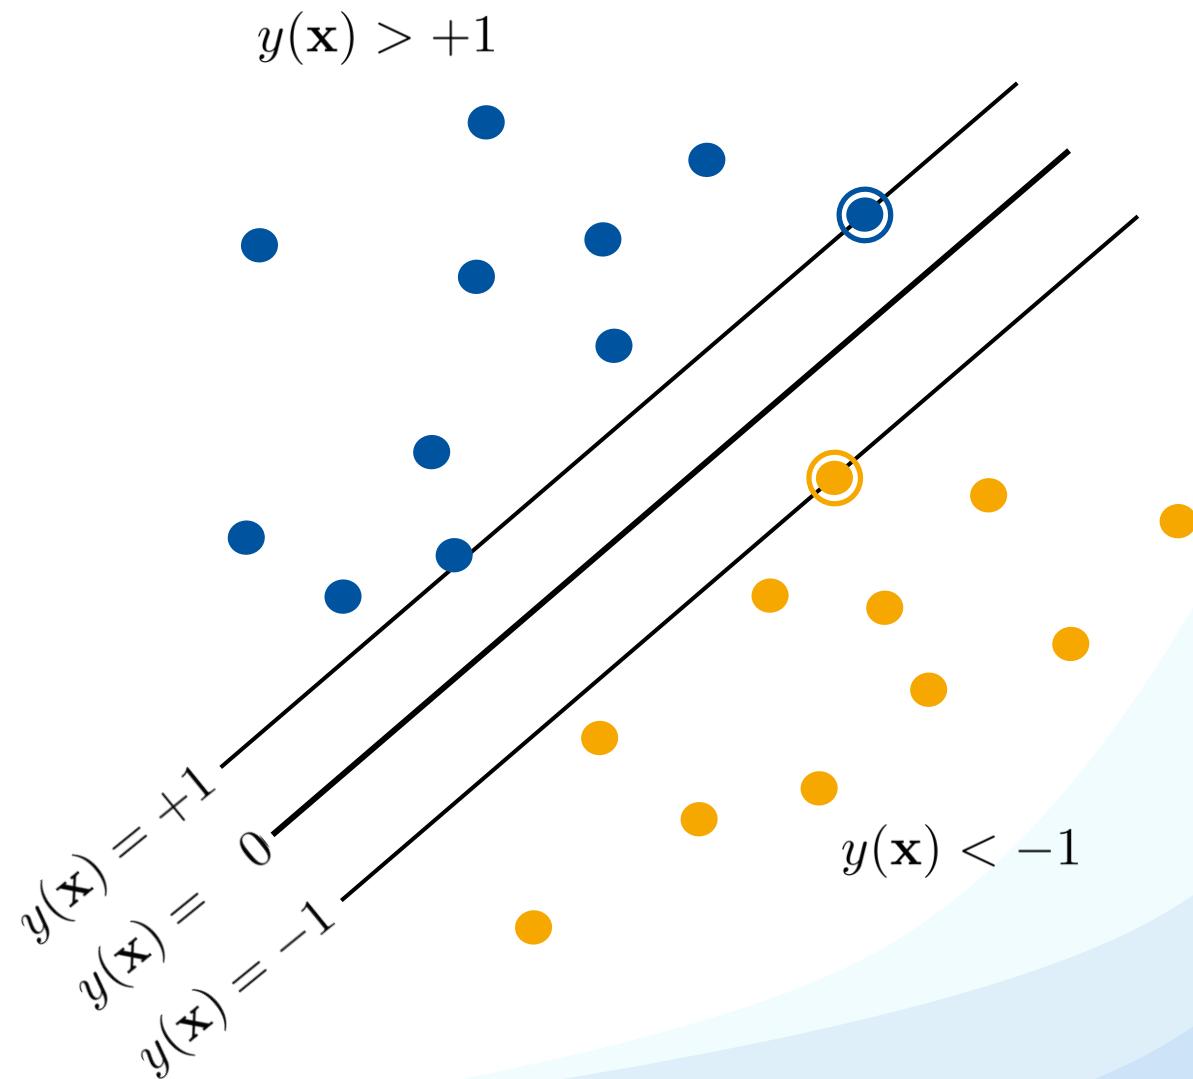
- In short:

$$t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 \quad \forall n$$

- We can rescale \mathbf{w} such that the equation holds exactly for the points on the margin:

$$t_n(\mathbf{w}^T \mathbf{x}_n + b) = 1$$

- There will be at least one such point.



- We can choose \mathbf{w} such that

$$\mathbf{w}^T \mathbf{x}_n + b = +1 \text{ for one } t_n = +1$$

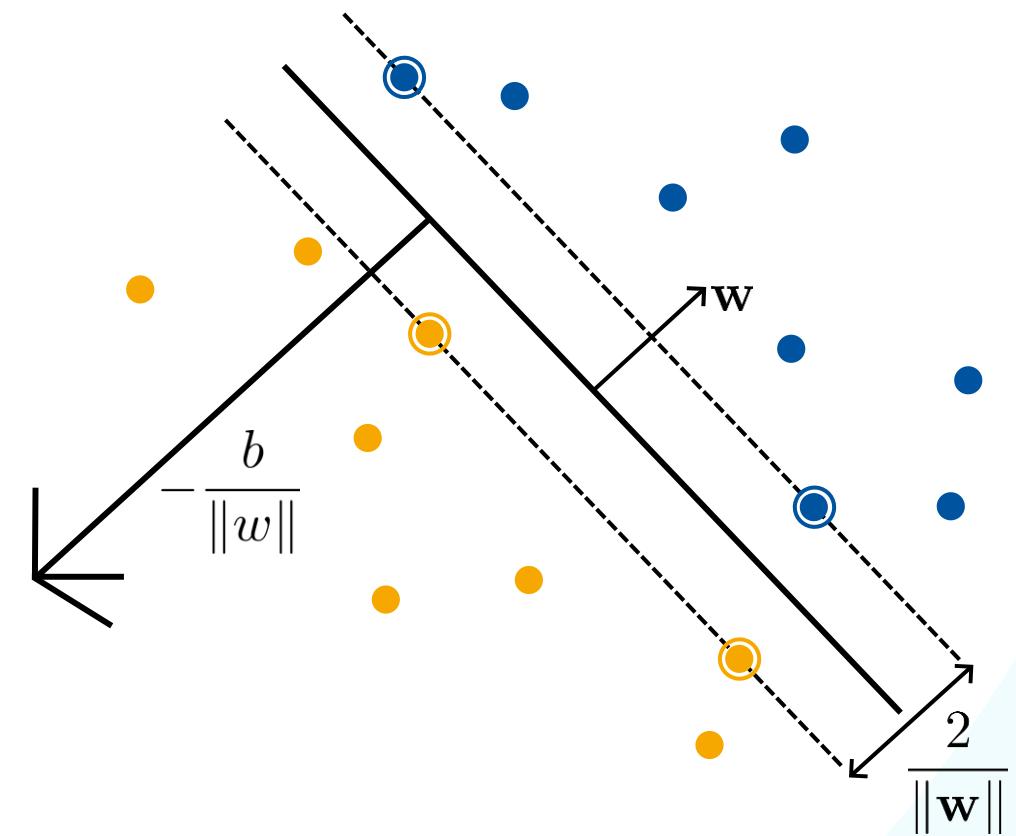
$$\mathbf{w}^T \mathbf{x}_n + b = -1 \text{ for one } t_n = -1$$

- The distance between those hyperplanes is then the margin:

$$d_- = d_+ = \frac{1}{\|\mathbf{w}\|}$$

$$d_- + d_+ = \frac{2}{\|\mathbf{w}\|}$$

\Rightarrow Maximize the margin by minimizing $\|\mathbf{w}\|^2$



- Optimization problem
 - Find the hyperplane with maximum margin by optimizing:

$$\arg \min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2$$

“Maximize the margin”

such that

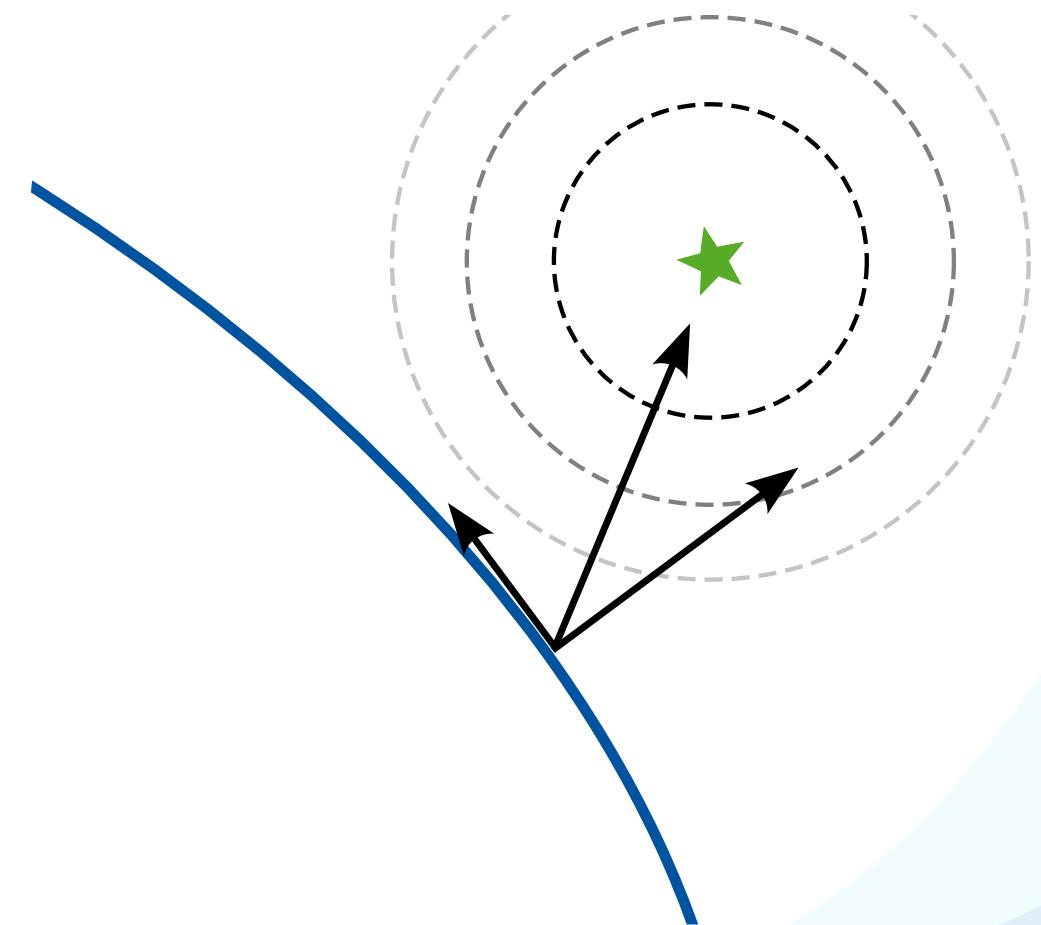
$$t_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1 \quad \forall n$$

*“such that each point
is on the correct side
of the margin”*

- This is a **quadratic programming problem** with linear constraints.

Support Vector Machines

1. Maximum Margin Classification
 - a) **Constrained Optimization**
2. Primal Formulation
3. Dual Formulation
4. Soft-Margin SVMs
5. Non-linear SVMs
6. Error Function Analysis



Constrained Optimization

- Recall the SVM objective:

$$\arg \min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2, \quad \text{such that} \quad t_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1 \quad \forall n$$

- This is a **constrained optimization problem**.
 - We want to optimize an objective $K(\mathbf{x})$ subject to constraints $f(\mathbf{x})$:

$\underset{\mathbf{x}}{\text{opt}} K(\mathbf{x})$	$\min \text{ or } \max$
such that	$f(\mathbf{x}) = 0$
	<i>equality constraints</i>
	$f(\mathbf{x}) \geq 0$
	<i>inequality constraints</i>

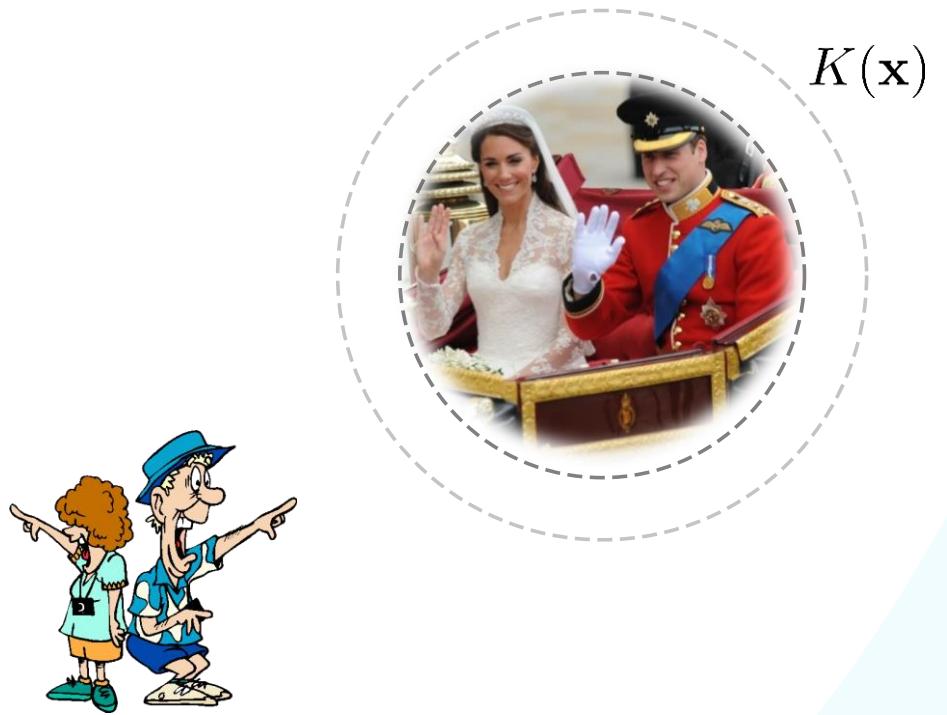
SVM

$$K(\mathbf{w}) = \frac{1}{2} \|\mathbf{w}\|^2$$
$$f_n(\mathbf{w}) = t_n(\mathbf{w}^\top \mathbf{x}_n + b) - 1 \geq 0 \quad \forall n$$

- We can solve such constrained optimization problems using the technique of **Lagrange multipliers**.

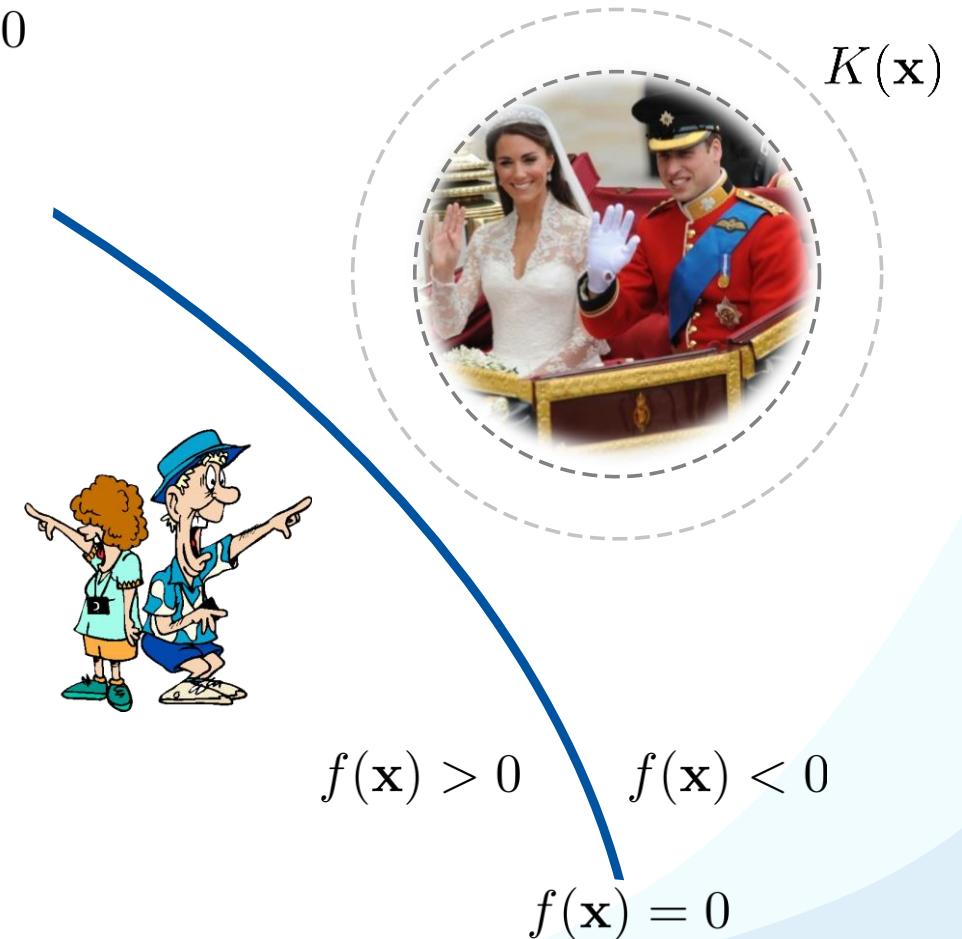
Lagrange Multipliers

- We want to maximize $K(\mathbf{x})$



Lagrange Multipliers

- We want to maximize $K(\mathbf{x})$ subject to constraints $f(\mathbf{x}) = 0$



Lagrange Multipliers

- We want to maximize $K(\mathbf{x})$ subject to constraints $f(\mathbf{x}) = 0$
- We can only move along $\nabla_{||} K = \nabla K + \lambda \nabla f$, with $\lambda \neq 0$.
- Add the constraints to the objective by introducing auxiliary variables λ :

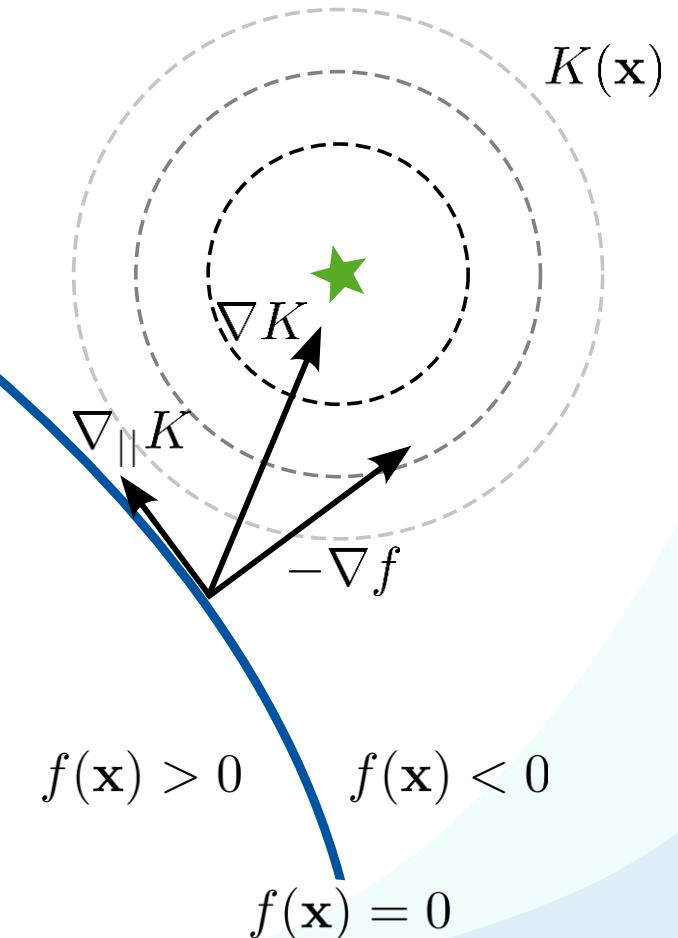
$$\mathcal{L}(\mathbf{x}, \lambda) = K(\mathbf{x}) + \lambda f(\mathbf{x})$$

- \mathcal{L} is called the **Lagrangian** form of the optimization problem, and λ is referred to as a **Lagrange multiplier**.
- Optimize \mathcal{L} :

$$\frac{\partial \mathcal{L}}{\partial \mathbf{x}} = \nabla_{||} K \stackrel{!}{=} 0$$

$$\frac{\partial \mathcal{L}}{\partial \lambda} = f(\mathbf{x}) \stackrel{!}{=} 0$$

The objective is maximized while satisfying the constraints.



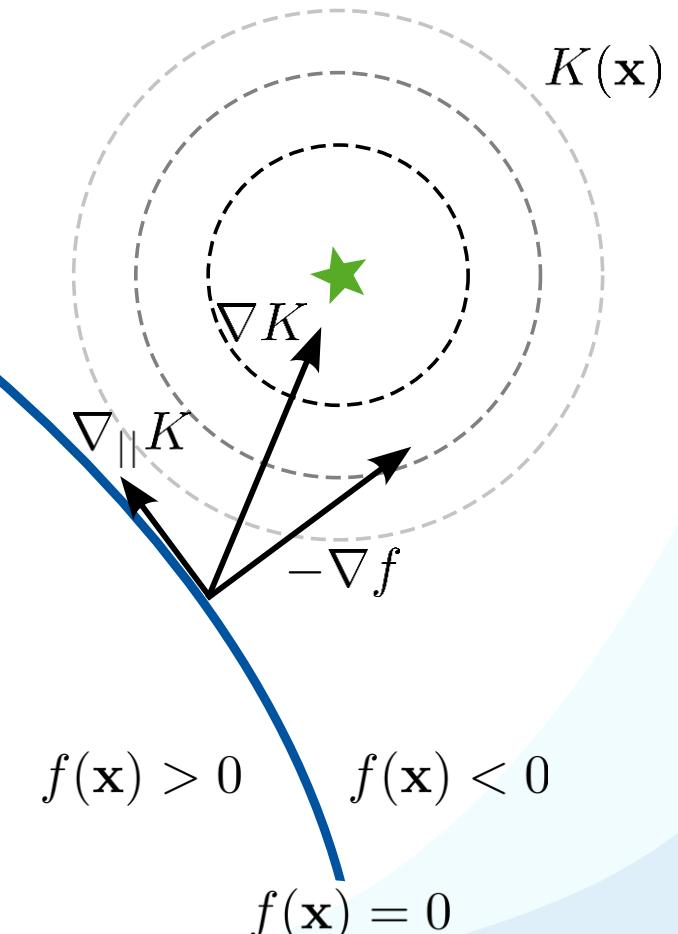
Inequality Constraints

- Now let's use inequality constraints $f(\mathbf{x}) \geq 0$.
- Optimize $\mathcal{L}(\mathbf{x}, \lambda) = K(\mathbf{x}) + \lambda f(\mathbf{x})$
 - Two cases
 - Solution lies on boundary:
 $\Rightarrow f(\mathbf{x}) = 0$ for some $\lambda > 0$
 - Solution lies inside $f(\mathbf{x}) > 0$:
 \Rightarrow Constraint inactive: $\lambda = 0$
- Karush-Kuhn-Tucker (KKT) conditions:

$$\begin{aligned}\lambda &\geq 0 \\ f(\mathbf{x}) &\geq 0 \\ \lambda f(\mathbf{x}) &= 0\end{aligned}$$

} In both cases:
 $\lambda f(\mathbf{x}) = 0$

*All valid solutions
need to fulfill the
KKT conditions.*



Maximization vs. Minimization

- Note: differences for maximization vs. minimization.
- If we want to **maximize** $K(\mathbf{x})$ subject to $f(\mathbf{x}) \geq 0$, we optimize the Lagrangian form

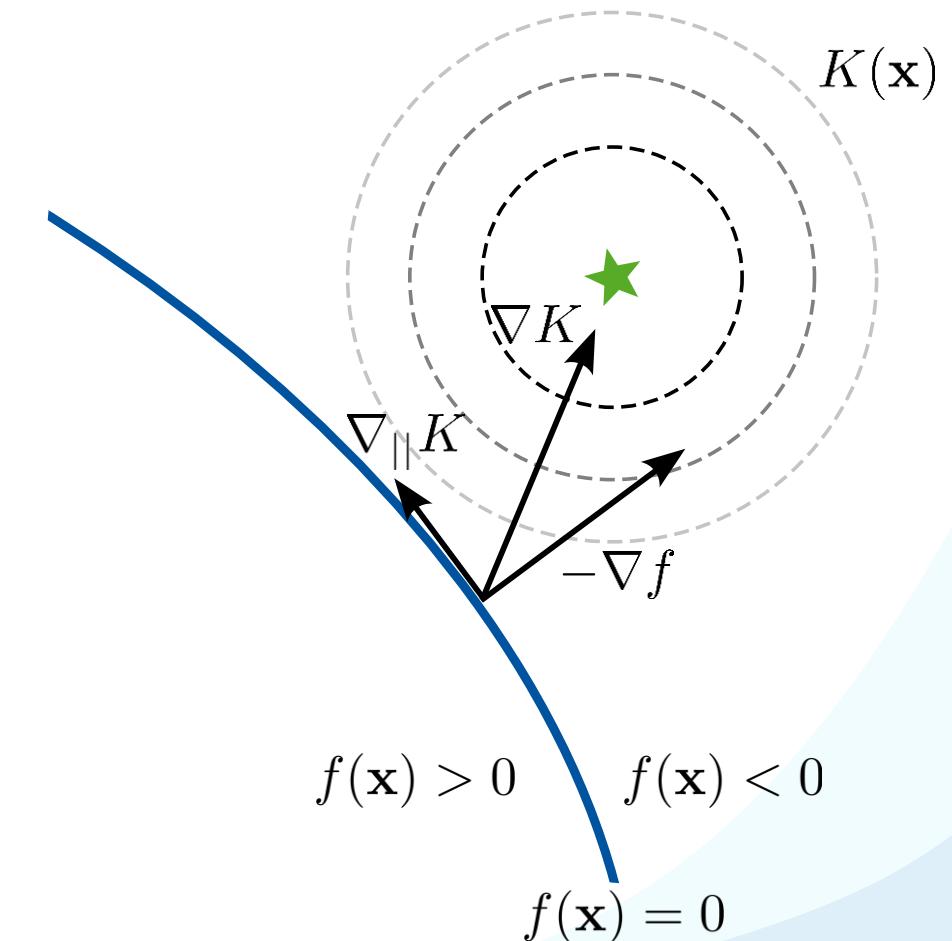
$$\mathcal{L}(\mathbf{x}, \lambda) = K(\mathbf{x}) + \lambda f(\mathbf{x})$$

- **maximize** w.r.t. \mathbf{x}
- **minimize** w.r.t. λ

- If we want to **minimize** $K(\mathbf{x})$ subject to $f(\mathbf{x}) \geq 0$, we optimize the Lagrangian form

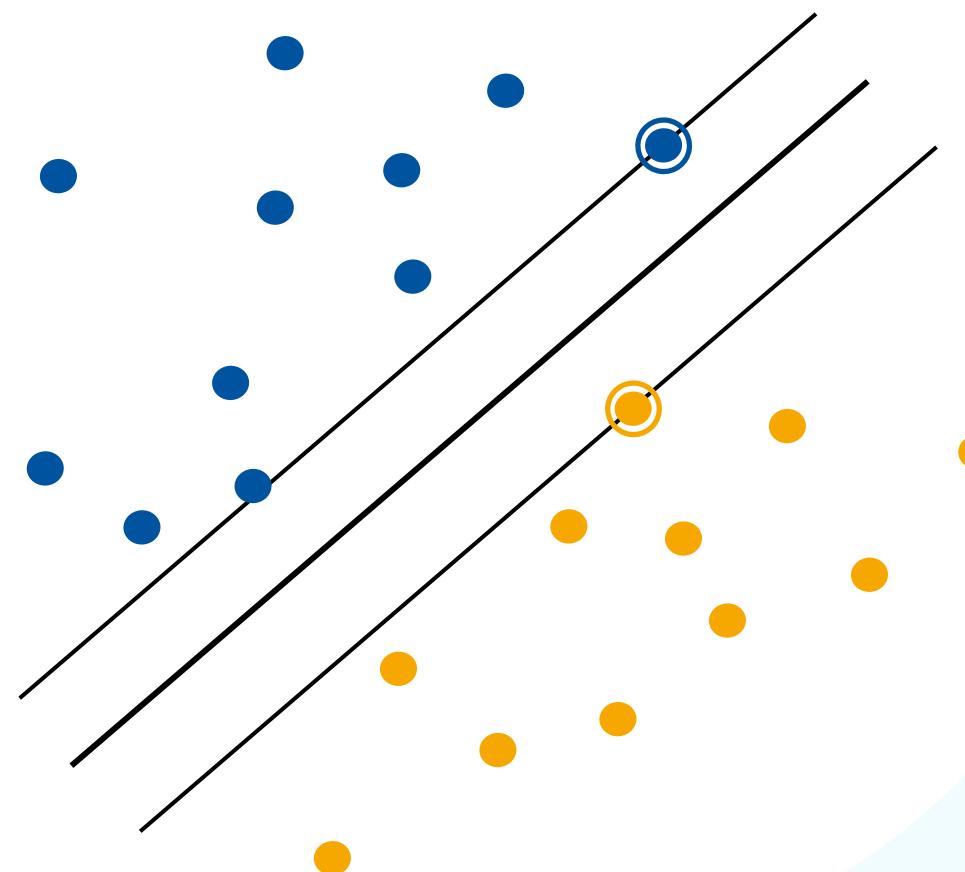
$$\mathcal{L}(\mathbf{x}, \lambda) = K(\mathbf{x}) - \lambda f(\mathbf{x})$$

- **minimize** w.r.t. \mathbf{x}
- **maximize** w.r.t. λ



Support Vector Machines

1. Maximum Margin Classification
2. **Primal Formulation**
3. Dual Formulation
4. Soft-Margin SVMs
5. Non-linear SVMs
6. Error Function Analysis



Primal SVM Formulation

- Recall the SVM objective:

$$\arg \min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2, \quad \text{such that} \quad t_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1 \quad \forall n$$

- We introduce positive Lagrange multipliers $a_n \geq 0$ and get the [primal form](#) of SVMs:

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{n=1}^N a_n [t_n(\mathbf{w}^\top \mathbf{x}_n + b) - 1]$$

Necessary and sufficient conditions:

$$\begin{aligned} a_n &\geq 0 \\ t_n(\mathbf{w}^\top \mathbf{x}_n + b) - 1 &\geq 0 \\ a_n[t_n(\mathbf{w}^\top \mathbf{x}_n + b) - 1] &= 0 \end{aligned}$$

KKT conditions:

$$\begin{aligned} \lambda &\geq 0 \\ f(\mathbf{x}) &\geq 0 \\ \lambda f(\mathbf{x}) &= 0 \end{aligned}$$

Lagrangian Formulation

- We want to minimize the primal form:

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{n=1}^N a_n [t_n (\mathbf{w}^\top \mathbf{x}_n + b) - 1]$$

$$\frac{\partial L(\mathbf{w}, b, \mathbf{a})}{\partial b} = \sum_{n=1}^N a_n t_n$$

$$\frac{\partial L(\mathbf{w}, b, \mathbf{a})}{\partial \mathbf{w}} = \mathbf{w} - \sum_{n=1}^N a_n t_n \mathbf{x}_n$$

- Setting the gradients for \mathbf{w}, b to zero, we get:

$$\frac{\partial L}{\partial b} = 0 \Rightarrow \sum_{n=1}^N a_n t_n = 0$$

$$\frac{\partial L}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_{n=1}^N a_n t_n \mathbf{x}_n$$

- The hyperplane is computed as a linear combination of training examples:

$$\mathbf{w} = \sum_{n=1}^N a_n t_n \mathbf{x}_n$$

$$\frac{\partial L}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_{n=1}^N a_n t_n \mathbf{x}_n$$

- Additionally, the solution needs to fulfill

$$a_n [t_n(\mathbf{w}^\top \mathbf{x}_n + b) - 1] = 0$$

- This implies $a_n > 0$ only for those points for which

$$[t_n(\mathbf{w}^\top \mathbf{x}_n + b) - 1] = 0$$

KKT conditions:

$$a_n \geq 0$$

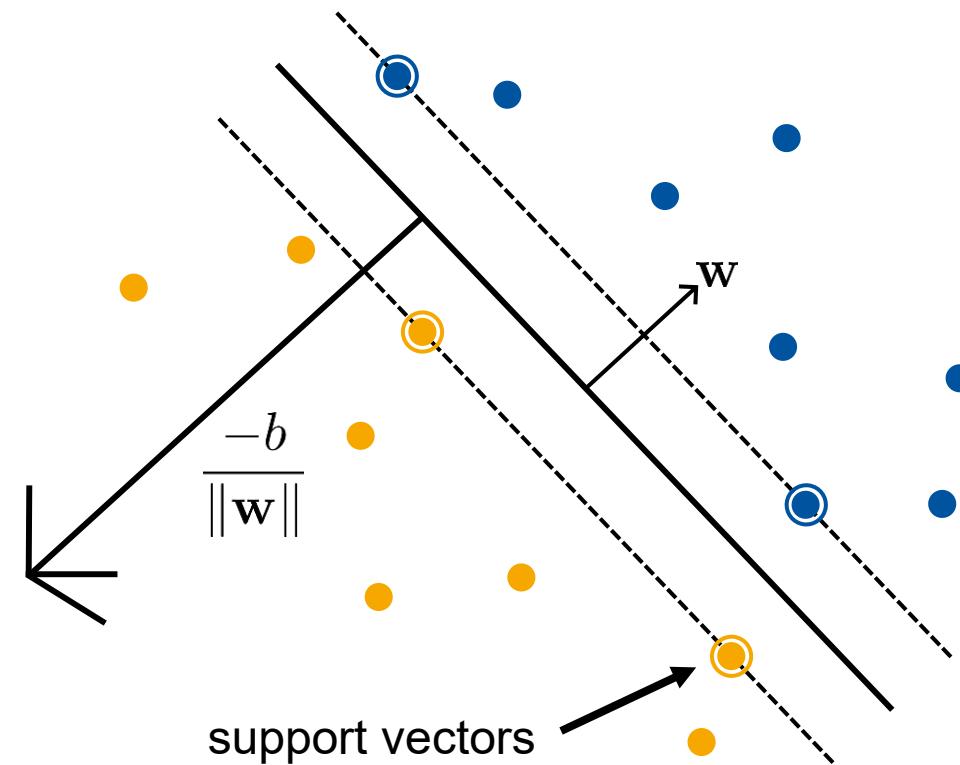
$$t_n(\mathbf{w}^\top \mathbf{x}_n + b) - 1 \geq 0$$

$$a_n [t_n(\mathbf{w}^\top \mathbf{x}_n + b) - 1] = 0$$

Only some data points influence the decision boundary!

Intuition

- The training points with $a_n > 0$ are called **support vectors**.
- They are the points on the margin.
- This makes the SVM robust to “too correct” points!



- We still need to find b .
- Observation: Any support vector \mathbf{x}_n satisfies

$$t_n y(\mathbf{x}_n) = t_n \left(\sum_{m \in \mathcal{S}} a_m t_m \mathbf{x}_m^\top \mathbf{x}_n + b \right) = 1$$

- Using $t_n^2 = 1$, we can derive

$$b = t_n - \sum_{m \in \mathcal{S}} a_m t_m \mathbf{x}_m^\top \mathbf{x}_n$$

- In practice, it is more robust to average over all support vectors:

$$b = \frac{1}{N_{\mathcal{S}}} \sum_{n \in \mathcal{S}} \left(t_n - \sum_{m \in \mathcal{S}} a_m t_m \mathbf{x}_m^\top \mathbf{x}_n \right)$$

Advantages

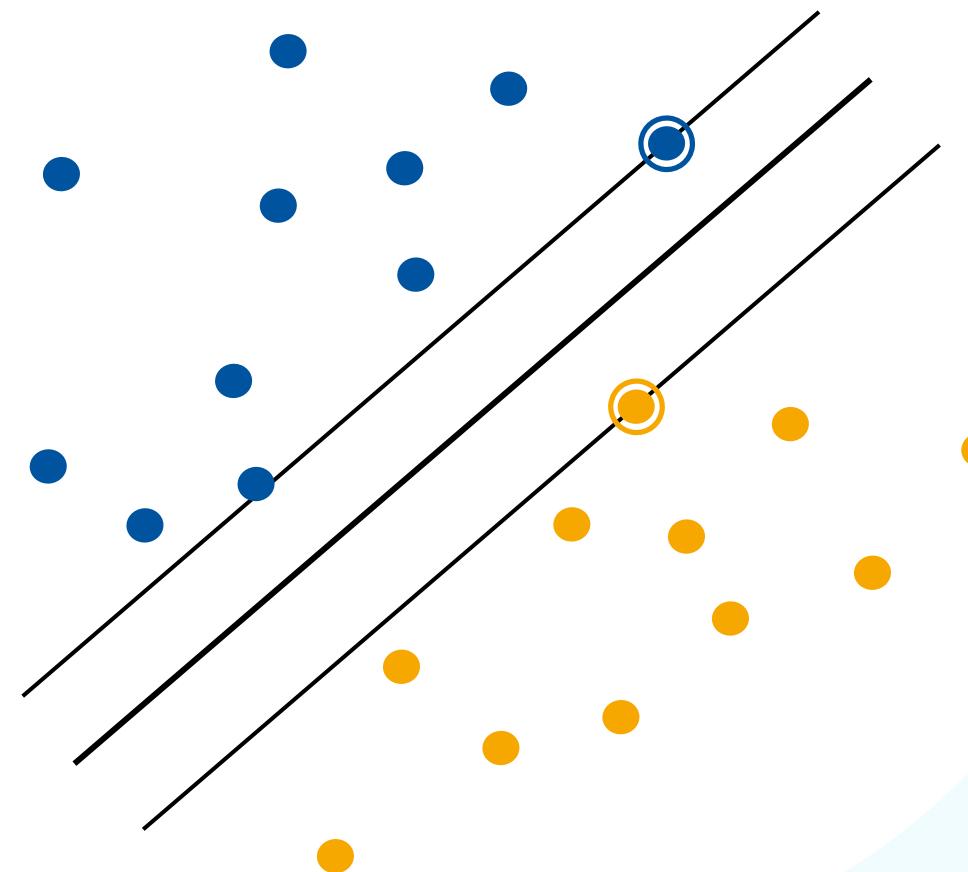
- SVMs yield a linear classifier with “guaranteed” generalization capability.
- Convex optimization, yields globally optimal solution.
- Solution depends only on a subset of the input data points, the **support vectors**.
- Automatic robustness against “too correct” data points.

Limitations

- Need to solve **quadratic programming** problem: time complexity for that is cubic in the number of variables.
- Here: Time complexity is in $\mathcal{O}(D^3)$.
- Scaling to high-dimensional data is difficult.

Support Vector Machines

1. Maximum Margin Classification
2. Primal Formulation
3. **Dual Formulation**
4. Soft-Margin SVMs
5. Non-linear SVMs
6. Error Function Analysis



Reminder: Primal SVM Formulation

- SVM objective:

$$\arg \min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2, \quad \text{such that} \quad t_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1 \quad \forall n$$

- This is a **Quadratic Programming (QP)** problem with **linear inequality constraints**.
 - In order to solve it, we have derived the **Lagrangian primal form**

$$L_p(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{n=1}^N a_n [t_n(\mathbf{w}^\top \mathbf{x}_n + b) - 1]$$

- We are *minimizing* this objective with respect to \mathbf{w} and b , and *maximizing* with respect to \mathbf{a} .

Solving a QP

- SVM objective:

$$\arg \min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2, \quad \text{such that} \quad t_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1 \quad \forall n$$

- Solving QPs is a well-understood problem
 - Typically done with the help of a [QP solver](#).
 - Solving a QP in K variables can be done in runtime $\mathcal{O}(K^3)$.
- In our case: $\mathbf{x}, \mathbf{w} \in \mathbb{R}^D$
 - #Variables: $D + 1$
 - ⇒ Complexity: $\mathcal{O}(D^3)$

Solving a QP

- SVM objective:

$$\arg \min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2, \quad \text{such that} \quad t_n(\mathbf{w}^\top \phi(\mathbf{x}_n) + b) \geq 1 \quad \forall n$$

- Solving QPs is a well-understood problem
 - Typically done with the help of a [QP solver](#).
 - Solving a QP in K variables can be done in runtime $\mathcal{O}(K^3)$.
 - In our case: $\mathbf{x}, \mathbf{w} \in \mathbb{R}^D$
 - #Variables: $D + 1$
 - ⇒ Complexity: $\mathcal{O}(D^3)$
 - With basis functions: $\phi(\mathbf{x}), \mathbf{w} \in \mathbb{R}^M$, $M \gg D$
 - #Variables: $M + 1$
 - ⇒ Complexity: $\mathcal{O}(M^3)$
- ⇒ [Curse of dimensionality](#), the SVM Primal Form does not scale well!

Dual SVM Formulation

- Idea
 - We will re-write the SVM primal form objective:

$$L_p(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{n=1}^N a_n [t_n(\mathbf{w}^\top \mathbf{x}_n + b) - 1]$$

in a dual form L_d that does not dependent on the dimensionality of the data.

- To do this, we will use the results we have derived in the previous section

$$\frac{\partial L}{\partial b} = 0 \Rightarrow \sum_{n=1}^N a_n t_n = 0$$

$$\frac{\partial L}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_{n=1}^N a_n t_n \mathbf{x}_n$$

Dual Form of the SVM Objective

- Maximize

$$L_d(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m (\mathbf{x}_m^\top \mathbf{x}_n)$$

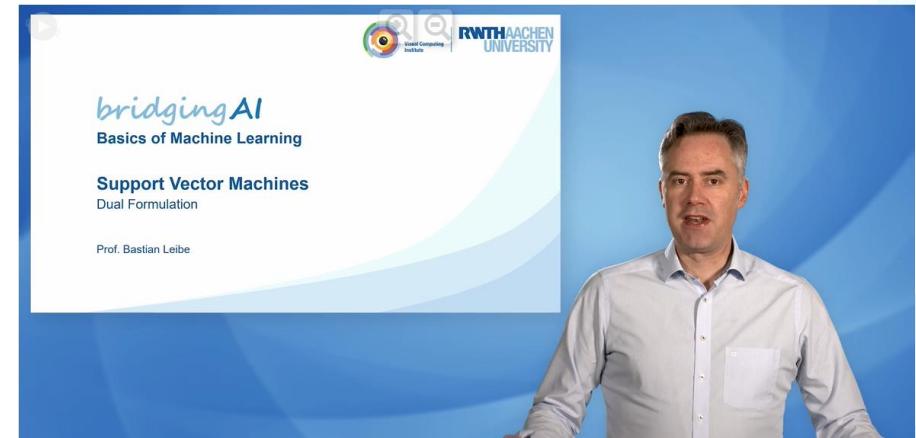
under the conditions

$$a_n \geq 0 \quad \forall n$$

$$\sum_{n=1}^N a_n t_n = 0$$

- We now have an optimization problem in N variables.
 \Rightarrow Complexity: $\mathcal{O}(N^3)$

For the derivation, please watch the video



Discussion

- What have we gained?
 - Previous complexity was $\mathcal{O}(D^3)$, now it is $\mathcal{O}(N^3)$.
 - *Isn't this much worse for large training sets???*
- However, the dual form has several advantages
 1. SVMs have sparse solutions: $a_n \neq 0$ only for support vectors.
 - This makes very efficient algorithms possible.
 - E.g., [Sequential Minimal Optimization \(SMO\)](#)
 - Effective runtime between $\mathcal{O}(N)$ and $\mathcal{O}(N^2)$.
 2. No dependency on the dimensionality anymore.
 - We can work with high-dimensional feature spaces!

Advantages

- Optimization problem only depends on the Lagrange multipliers a_n resulting in a worst-case runtime complexity of $\mathcal{O}(N^3)$.
- Since SVMs have sparse solutions and only few $a_n \neq 0$, specialized algorithms can solve the dual form very efficiently.
- The complexity of QP optimization no longer depends on the dimensionality of the feature space. This makes it possible to use very high-dimensional feature spaces.

Limitations

- Evaluating the SVM decision function

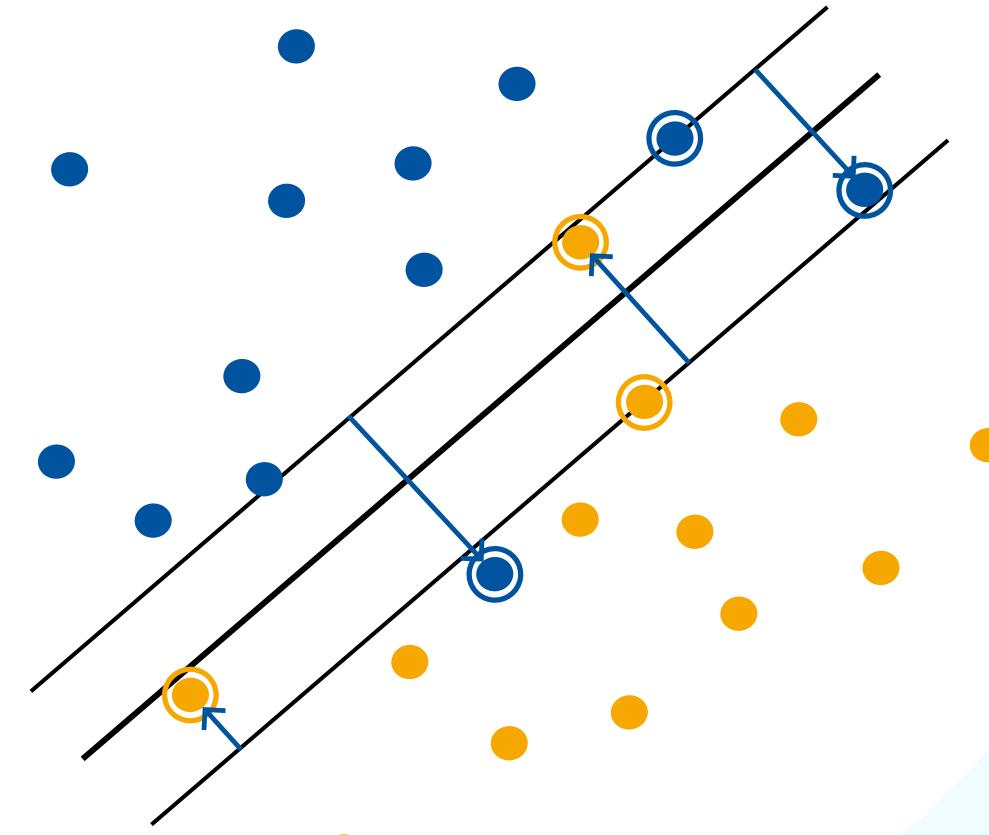
$$y(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x}) + b$$

$$\text{with } \mathbf{w} = \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n)$$

is still costly for high-dimensional feature spaces $\phi(\mathbf{x})$.

Support Vector Machines

1. Maximum Margin Classification
2. Primal Formulation
3. Dual Formulation
4. **Soft-Margin SVMs**
5. Non-linear SVMs
6. Error Function Analysis



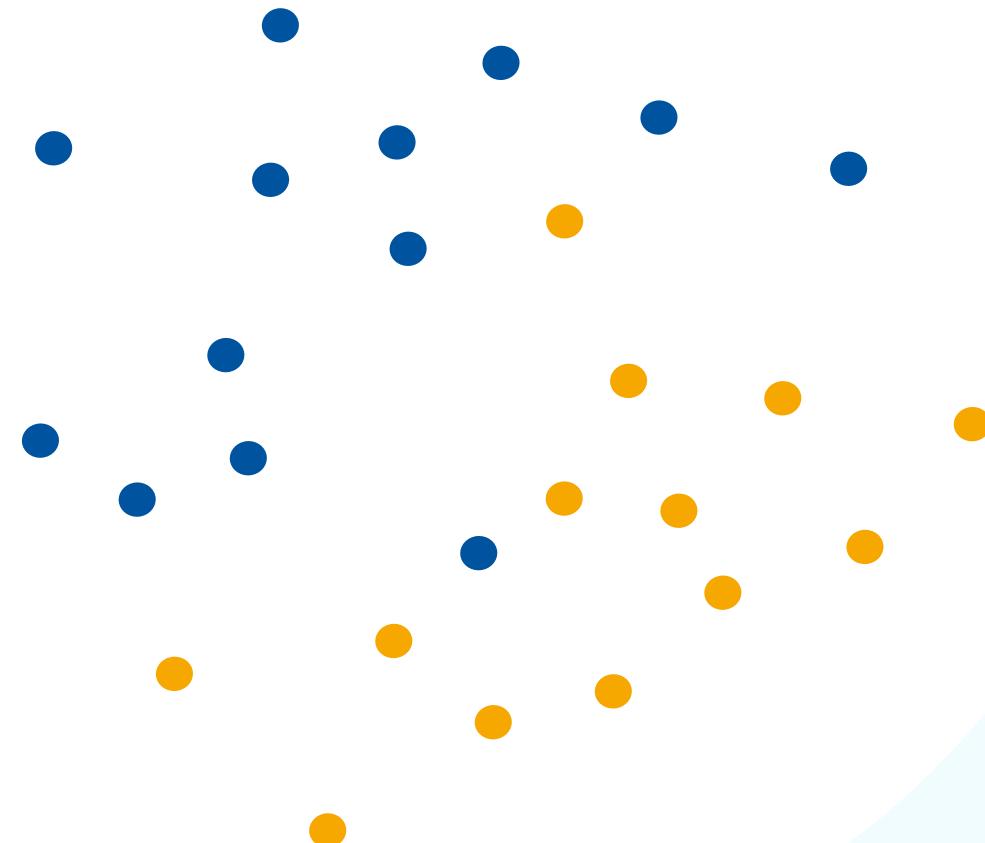
Soft-Margin SVM

- So far, we assumed linearly separable data.
 - Our current formulation has no solution if the data are not linearly separable!

$$\arg \min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2,$$

such that $t_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1 \quad \forall n$

- Need to introduce tolerance to outlier data points.
 - The resulting model is called **soft-margin SVM**.



Slack Variables

- For non-linearly separable data, not all constraints can be satisfied:

$$\mathbf{w}^T \mathbf{x}_n + b \geq +1 \quad \text{for } t_n = +1$$

$$\mathbf{w}^T \mathbf{x}_n + b \leq -1 \quad \text{for } t_n = -1$$

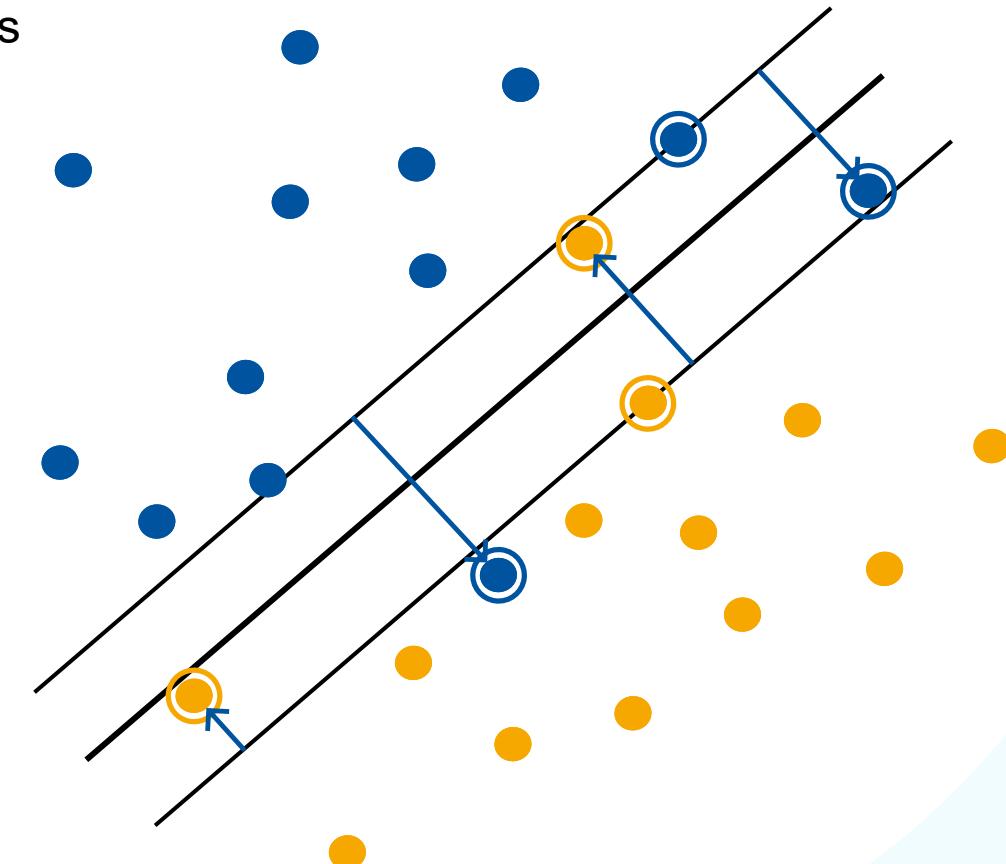
- Idea: Introduce **slack variables** $\xi_n \geq 0$:

$$\mathbf{w}^T \mathbf{x}_n + b \geq +1 - \xi_n \quad \text{for } t_n = +1$$

$$\mathbf{w}^T \mathbf{x}_n + b \leq -1 + \xi_n \quad \text{for } t_n = -1$$

⇒ We allow some datapoints to violate the constraint.

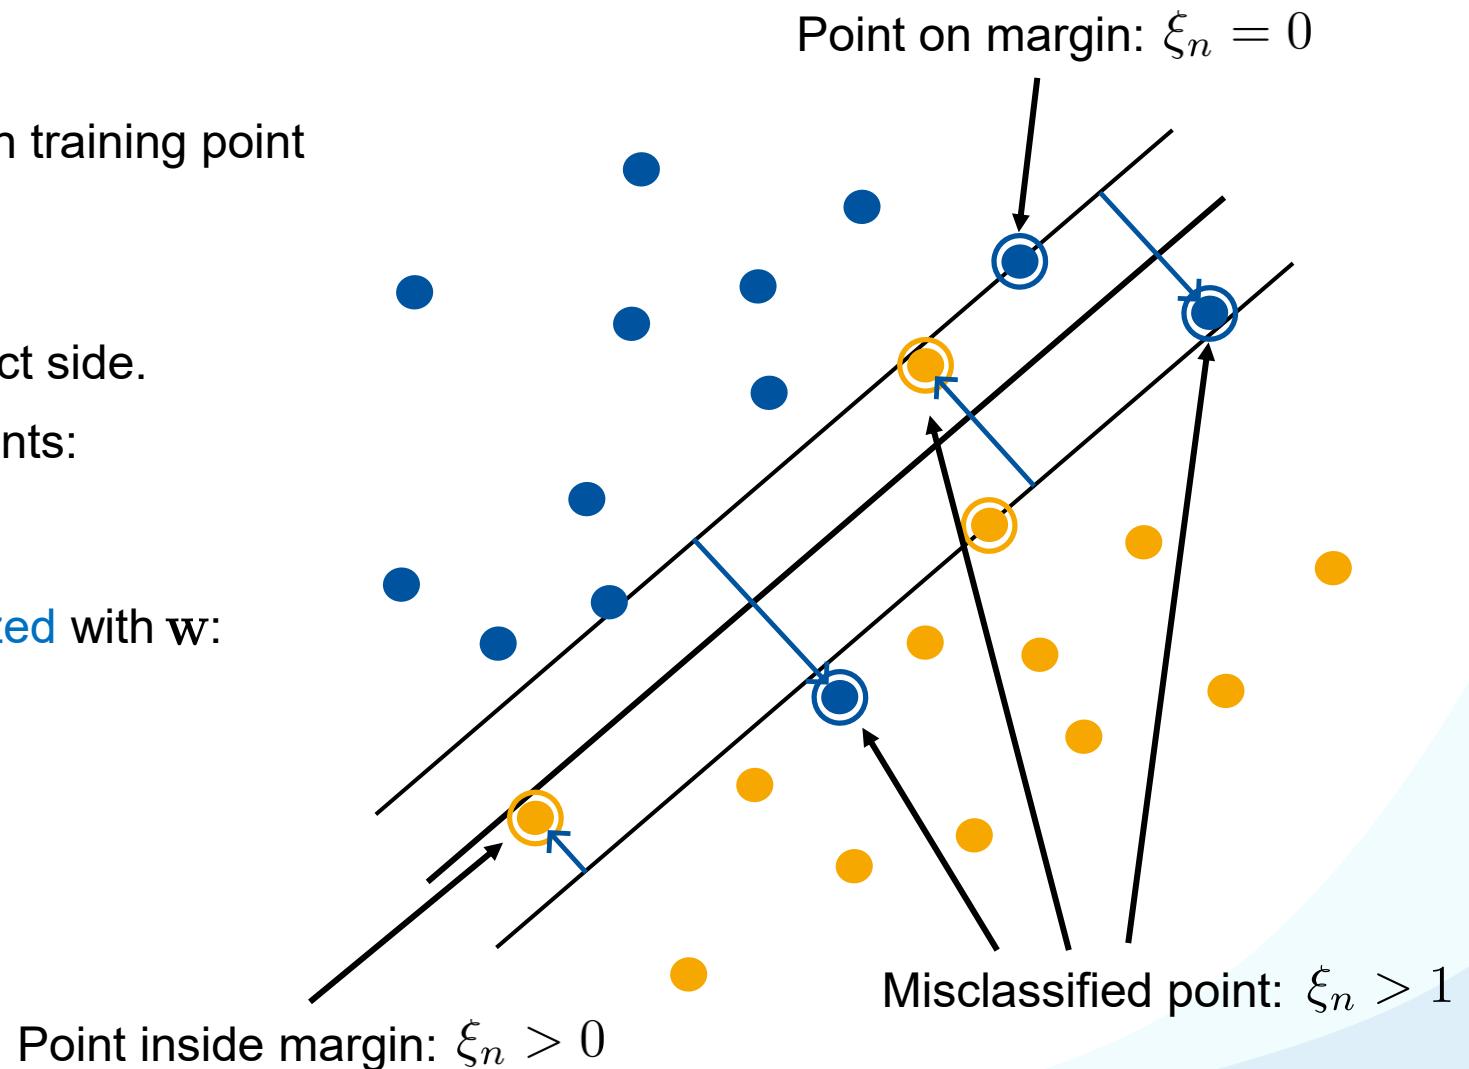
- For those points, the slack ξ_n makes up for the difference.



- Slack variables
 - One slack variable ξ_n for each training point
- Effect
 - $\xi_n = 0$ for points on the correct side.
 - **Linear penalty** for all other points:
$$\xi_n = |t_n - y(\mathbf{x}_n)|$$
- Slack variables are **jointly optimized** with \mathbf{w} :

$$\min \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n$$

where C is a tradeoff parameter.



New Primal Formulation

- Minimize

$$L_p = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n - \underbrace{\sum_{n=1}^N a_n [t_n(y(\mathbf{x}_n) - 1 + \xi_n)]}_{\text{Constraint}} - \underbrace{\sum_{n=1}^N \mu_n \xi_n}_{\text{Constraint}}$$
$$t_n y(\mathbf{x}_n) \geq 1 - \xi_n \quad \xi_n \geq 0$$

- KKT conditions

$$\begin{array}{ll} a_n \geq 0 & \mu_n \geq 0 \\ t_n y(\mathbf{x}_n) - 1 + \xi_n \geq 0 & \xi_n \geq 0 \\ a_n [t_n y(\mathbf{x}_n) - 1 + \xi_n] = 0 & \mu_n \xi_n = 0 \end{array}$$

New Dual Formulation

- Maximize

$$L_d(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m (\mathbf{x}_m^\top \mathbf{x}_n)$$

- Under the side conditions

$$0 \leq a_n \leq C \quad \forall n$$

$$\sum_{n=1}^N a_n t_n = 0$$

This is the only difference to before.

New Solution

- The decision hyperplane is again a linear combination of training samples:

$$\mathbf{w} = \sum_{n=1}^N a_n t_n \mathbf{x}_n$$

- This is still a sparse solution:
 - $a_n = 0$ for points on the correct side of the margin
 - Slack points with $\xi_n > 0$ are now also support vectors!
- Compute b by averaging over support vectors (points with $0 < a_n < C$):

$$b = \frac{1}{N_S} \sum_{n \in S} \left(t_n - \sum_{m \in S} a_m t_m \mathbf{x}_m^\top \mathbf{x}_n \right)$$

References and Further Reading

- More information about SVMs is available in Chapter 7.1 of Bishop's book.

Christopher M. Bishop
Pattern Recognition and Machine Learning
Springer, 2006

