



Visual Computing  
Institute

RWTH AACHEN  
UNIVERSITY

# Elements of Machine Learning & Data Science

Winter semester 2025/26

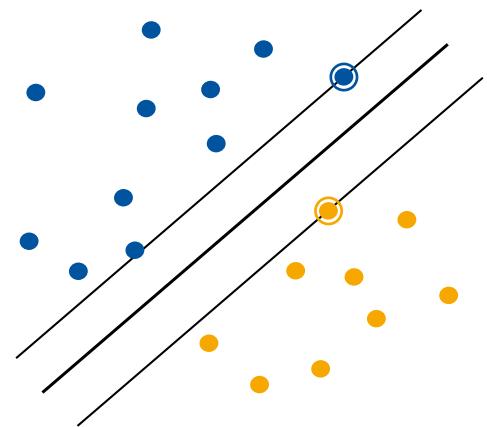
## Lecture 15 – Support Vector Machines II

15.12.2025

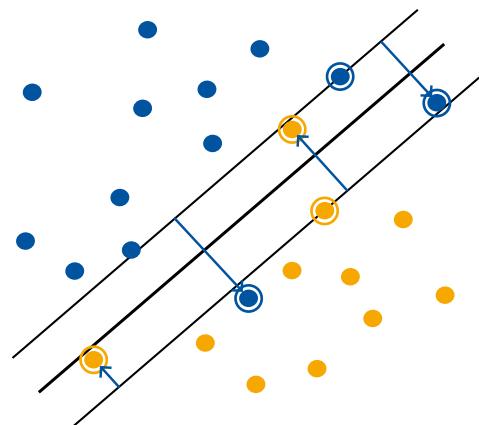
Prof. Bastian Leibe

# Machine Learning Topics

- 8. Introduction to ML
- 9. Probability Density Estimation
- 10. Linear Discriminants
- 11. Linear Regression
- 12. Logistic Regression
- 13. Support Vector Machines**
- 14. Neural Network Basics



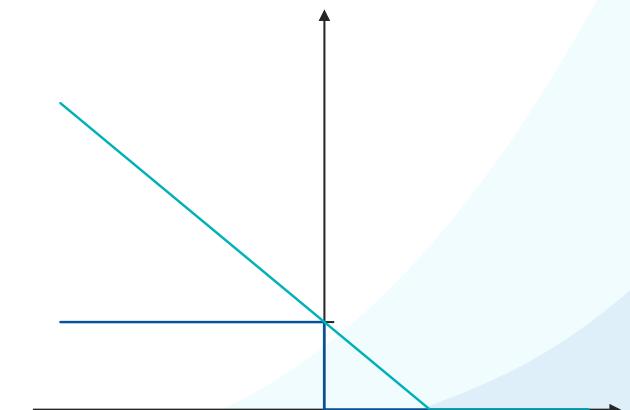
Maximum Margin Classification



Soft-Margin SVM

$$L_p(\mathbf{w}, b, \mathbf{a})$$
$$L_d(\mathbf{a})$$

Primal & Dual Form



Hinge Loss

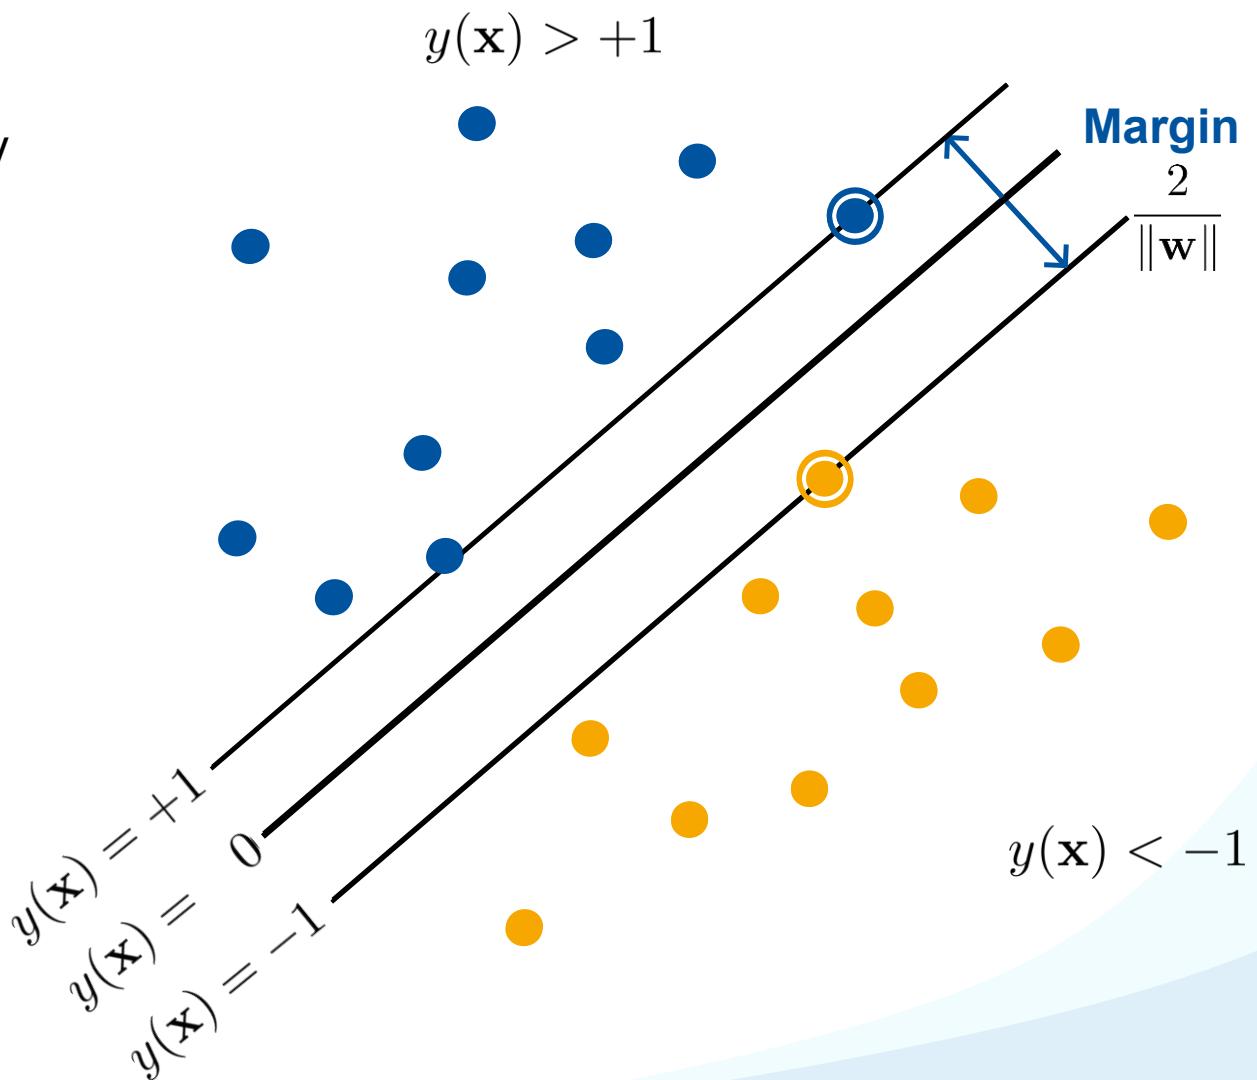
## Recap: Maximum Margin Classification

- Intuitively, we want to choose the classifier which leaves maximal “safety room” for future data points.
- This classifier has the largest **margin** between positive and negative points.
- We can rescale  $\mathbf{w}$  such that the distance of the points on the margin to the decision boundary is exactly 1.

$$t_n(\mathbf{w}^T \mathbf{x}_n + b) = 1$$

- If the data is linearly separable, then for all points, the following must hold:

$$t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 \quad \forall n$$



- Optimization problem
  - Find the hyperplane with maximum margin by optimizing:

$$\arg \min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2$$

*“Maximize the margin”*

such that

$$t_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1 \quad \forall n$$

*“such that each point  
is on the correct side  
of the margin”*

- This is a **quadratic programming problem** with linear constraints.

## Recap: Primal SVM Formulation

- Recall the SVM objective:

$$\arg \min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2, \quad \text{such that} \quad t_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1 \quad \forall n$$

- We introduce positive **Lagrange multipliers**  $a_n \geq 0$  and get the **primal form** of SVMs:

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{n=1}^N a_n [t_n(\mathbf{w}^\top \mathbf{x}_n + b) - 1]$$

Necessary and sufficient conditions:

$$\begin{aligned} a_n &\geq 0 \\ t_n(\mathbf{w}^\top \mathbf{x}_n + b) - 1 &\geq 0 \\ a_n[t_n(\mathbf{w}^\top \mathbf{x}_n + b) - 1] &= 0 \end{aligned}$$

KKT conditions:

$$\begin{aligned} \lambda &\geq 0 \\ f(\mathbf{x}) &\geq 0 \\ \lambda f(\mathbf{x}) &= 0 \end{aligned}$$

## Recap: Dual Form of the SVM Objective

- We can equivalently reformulate the SVM to maximize

$$L_d(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m (\mathbf{x}_m^\top \mathbf{x}_n)$$

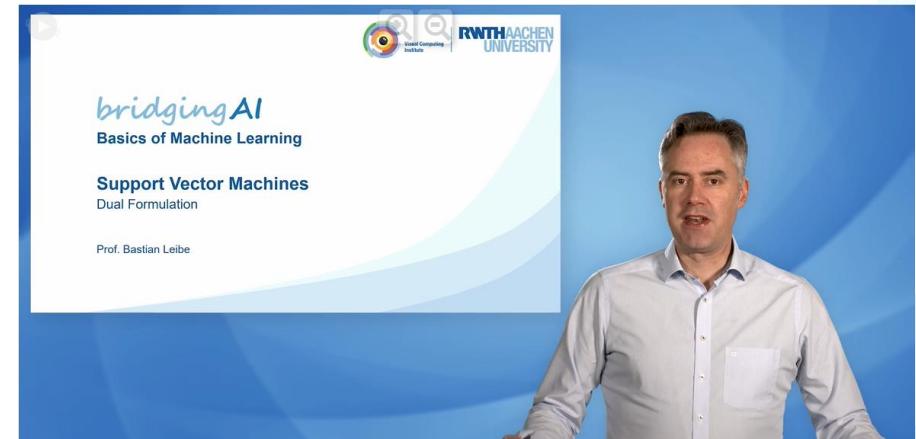
under the conditions

$$a_n \geq 0 \quad \forall n$$

$$\sum_{n=1}^N a_n t_n = 0$$

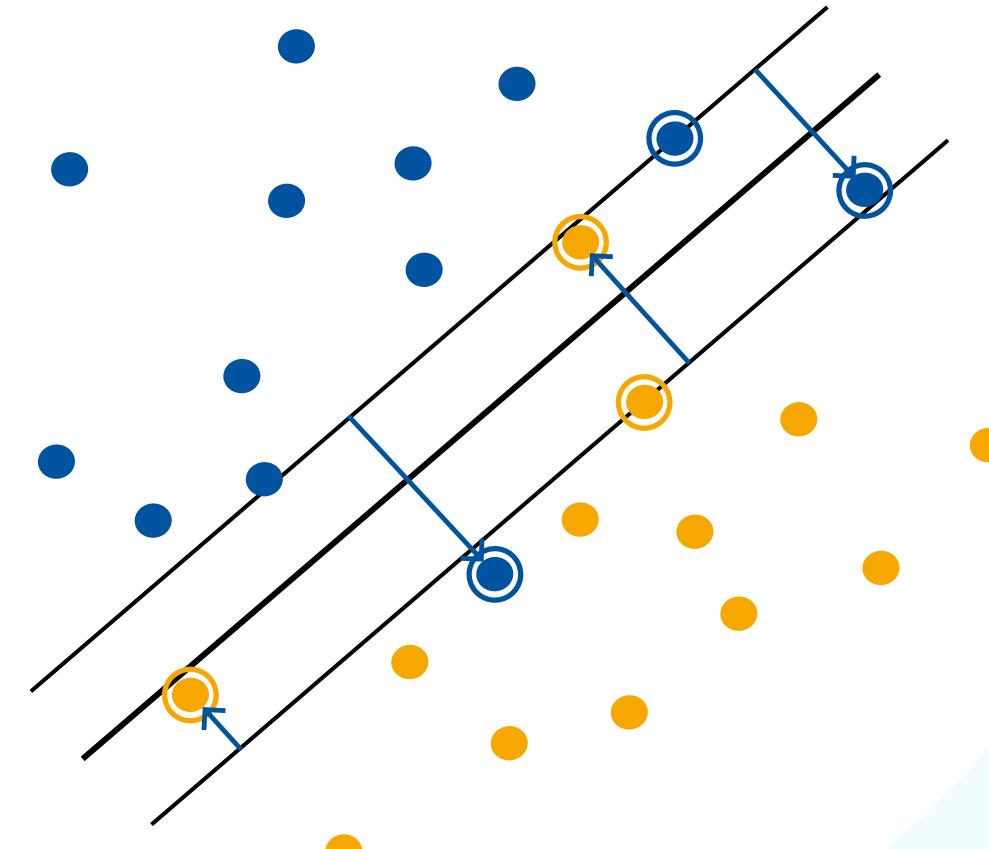
- We now have an optimization problem in  $N$  variables.  
⇒ Complexity:  $\mathcal{O}(N^3)$

*For the derivation, please watch the video*



# Support Vector Machines

1. Maximum Margin Classification
2. Primal Formulation
3. Dual Formulation
4. **Soft-Margin SVMs**
5. Non-linear SVMs
6. Error Function Analysis



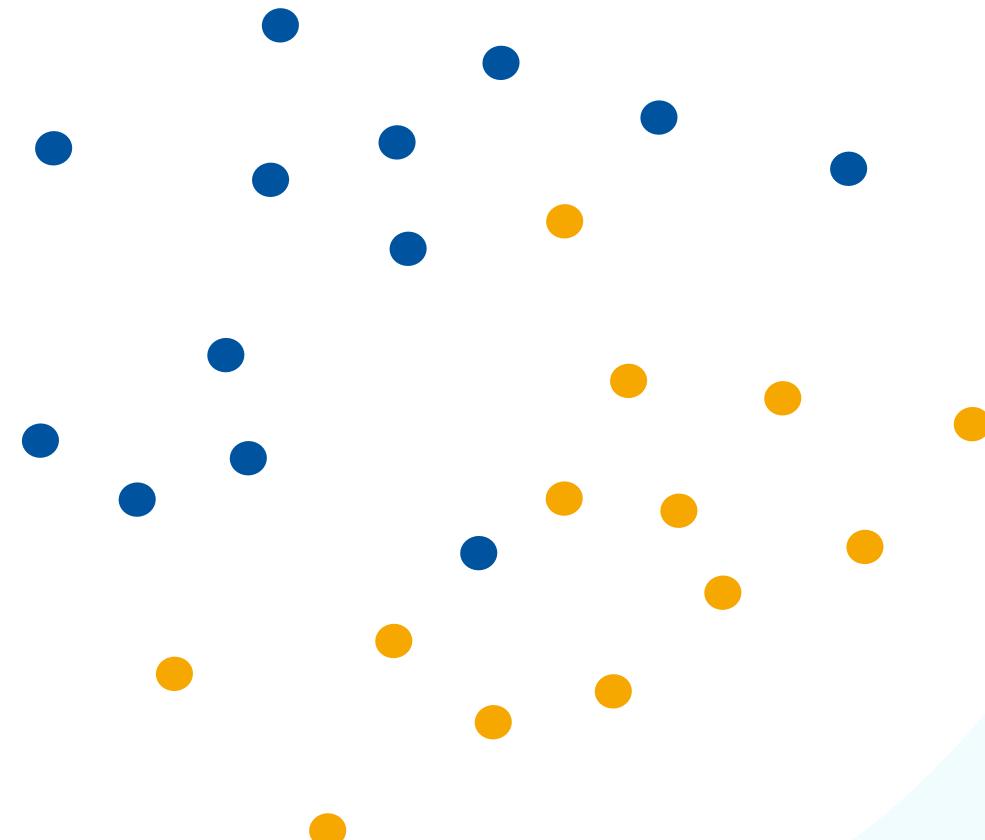
# Soft-Margin SVM

- So far, we assumed linearly separable data.
  - Our current formulation has no solution if the data are not linearly separable!

$$\arg \min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2,$$

such that  $t_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1 \quad \forall n$

- Need to introduce tolerance to outlier data points.
  - The resulting model is called **soft-margin SVM**.

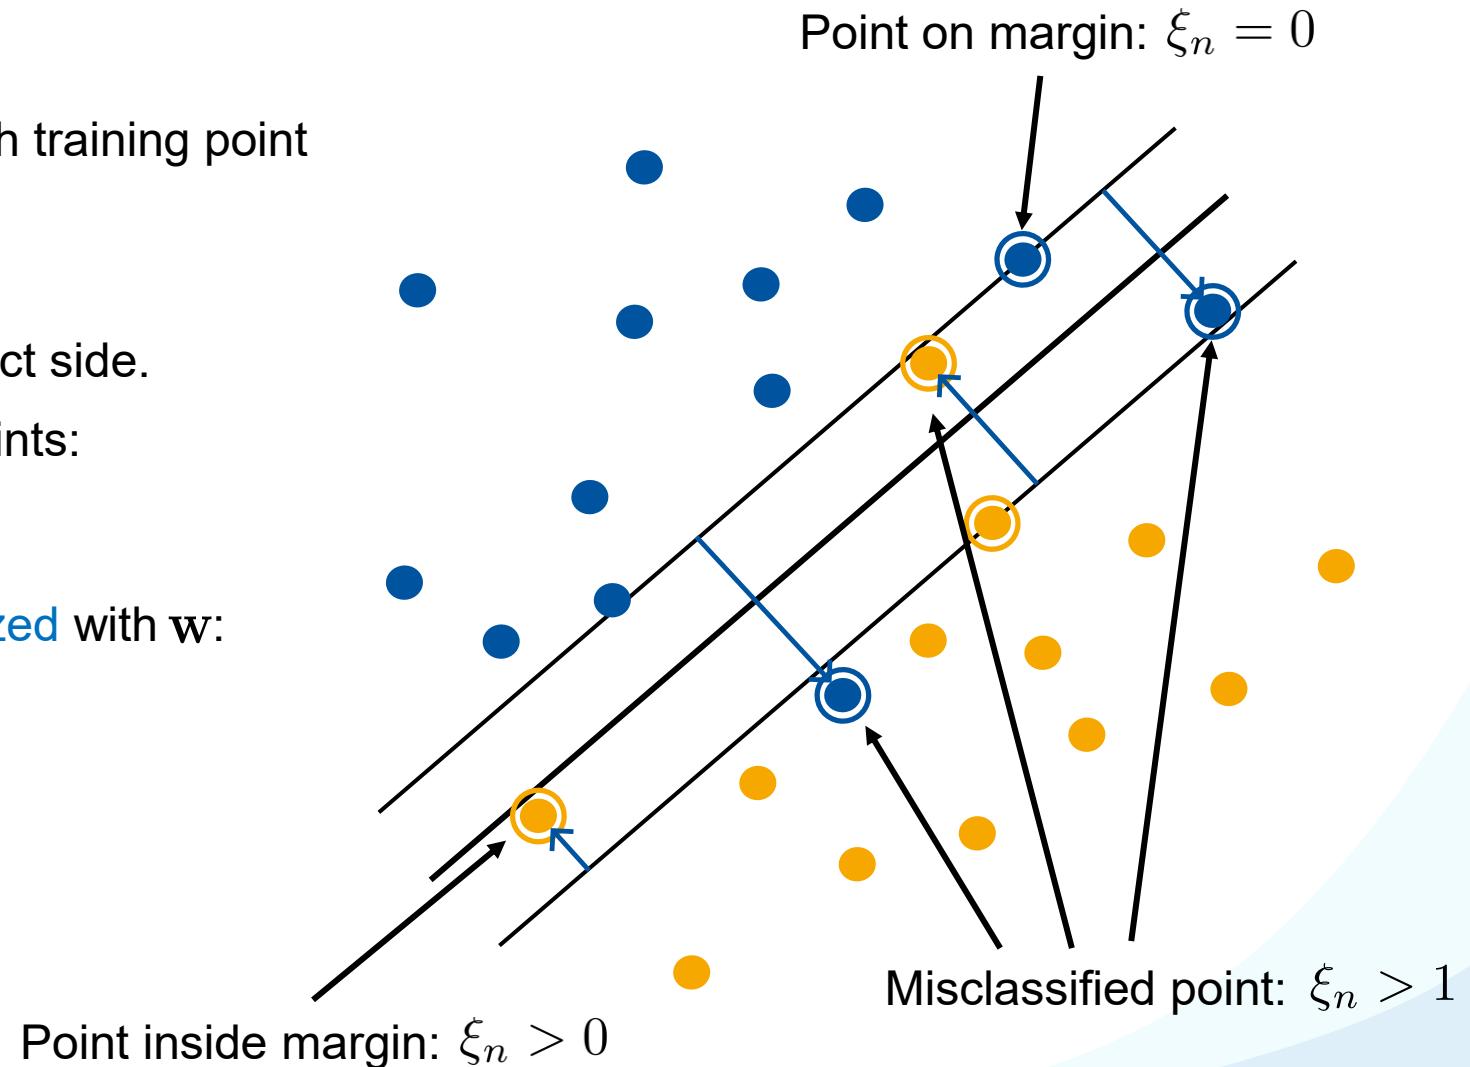


## Recap: Soft-Margin SVM with Slack Variables

- Slack variables
  - One slack variable  $\xi_n$  for each training point
- Effect
  - $\xi_n = 0$  for points on the correct side.
  - **Linear penalty** for all other points:
$$\xi_n = |t_n - y(\mathbf{x}_n)|$$
- Slack variables are **jointly optimized** with  $\mathbf{w}$ :

$$\min \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n$$

where  $C$  is a tradeoff parameter.



## Recap: New Primal Formulation

- Minimize

$$L_p = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n - \underbrace{\sum_{n=1}^N a_n [t_n(y(\mathbf{x}_n) - 1 + \xi_n)]}_{\text{Constraint}} - \underbrace{\sum_{n=1}^N \mu_n \xi_n}_{\text{Constraint}}$$
$$t_n y(\mathbf{x}_n) \geq 1 - \xi_n \quad \xi_n \geq 0$$

- KKT conditions

$$\begin{array}{ll} a_n \geq 0 & \mu_n \geq 0 \\ t_n y(\mathbf{x}_n) - 1 + \xi_n \geq 0 & \xi_n \geq 0 \\ a_n [t_n y(\mathbf{x}_n) - 1 + \xi_n] = 0 & \mu_n \xi_n = 0 \end{array}$$

## Recap: New Dual Formulation

- Maximize

$$L_d(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m (\mathbf{x}_m^\top \mathbf{x}_n)$$

- Under the side conditions

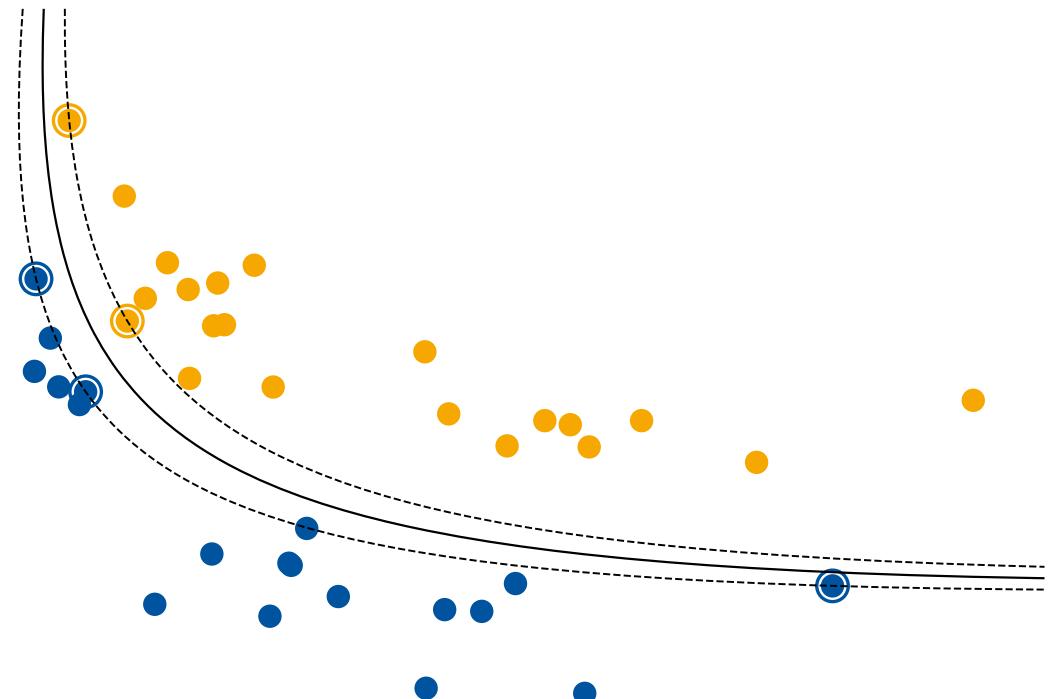
$$0 \leq a_n \leq C \quad \forall n$$

$$\sum_{n=1}^N a_n t_n = 0$$

*This is the only difference to before.*

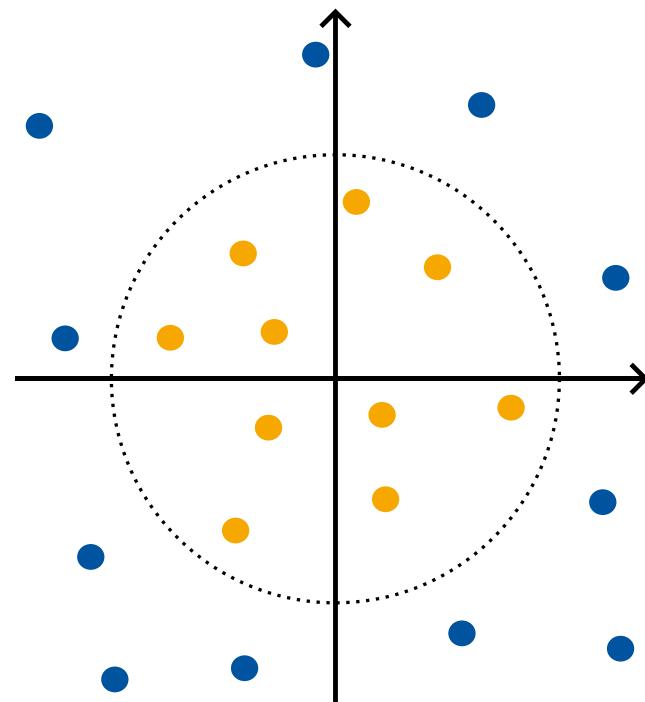
# Support Vector Machines

1. Maximum Margin Classification
2. Primal Formulation
3. Dual Formulation
4. Soft-Margin SVMs
5. **Non-Linear SVMs**
6. Error Function Analysis



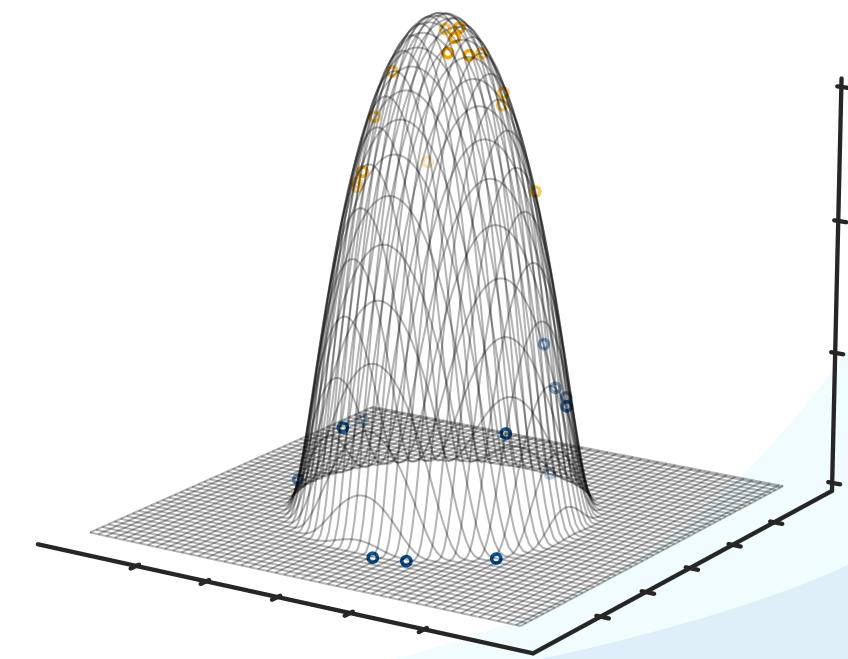
# Non-Linear SVMs

- So far, we have only considered linear decision boundaries.
- We now combine non-linear basis functions with SVMs.



$$\phi : \mathbb{R}^D \rightarrow \mathbb{R}^M$$

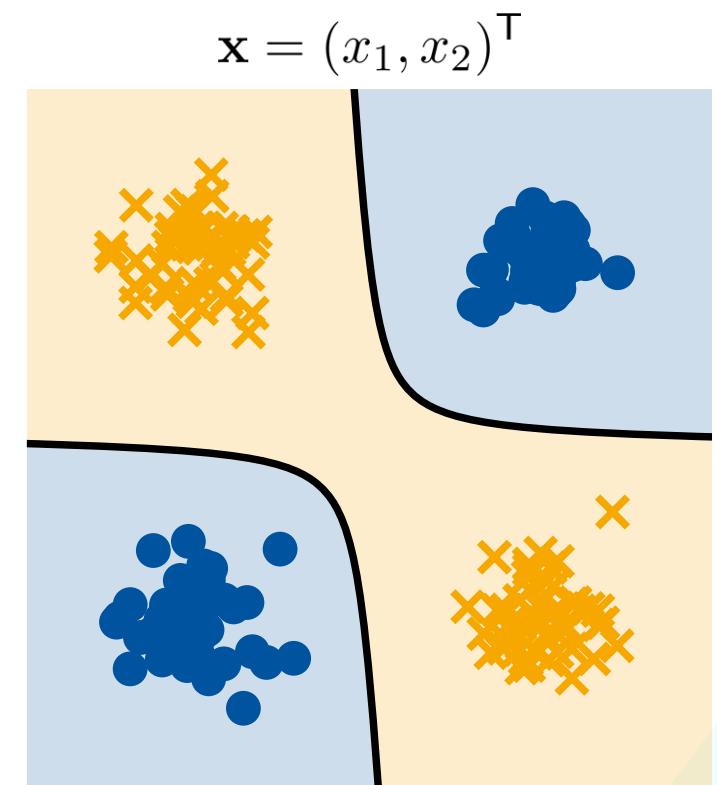
$$M \gg D$$



# Feature Spaces

- We have already seen non-linear basis functions:
  - Apply a nonlinear transformation  $\phi$  to the data points  $\mathbf{x}_n$ :
$$\mathbf{x} \in \mathbb{R}^D, \quad \phi : \mathbb{R}^D \rightarrow \mathbb{R}^M$$
  - Classify with a hyperplane in higher-dim. space  $\mathbb{R}^M$ :
$$\mathbf{w}^\top \phi(\mathbf{x}) + b = 0$$

$\Rightarrow$  Linear classifier in  $\mathbb{R}^M$ , nonlinear classifier in  $\mathbb{R}^D$ .
- Let us now apply this to SVMs...
  - We can train our SVM on the transformed features  $\phi(\mathbf{x})$  to get non-linear decision boundaries.
  - Usually,  $M \gg D$ : evaluating  $\mathbf{w}^\top \phi(\mathbf{x})$  can be quite expensive!



# The Kernel Trick

- On a closer look,  $\phi(\mathbf{x})$  only appears in the form of dot products:

$$L_d(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m (\phi(\mathbf{x}_m)^\top \phi(\mathbf{x}_n))$$

$$\begin{aligned} y(\mathbf{x}) &= \mathbf{w}^\top \phi(\mathbf{x}) + b \\ &= \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n)^\top \phi(\mathbf{x}) + b \end{aligned}$$

$$\boxed{\mathbf{w} = \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n)}$$

# The Kernel Trick

- On a closer look,  $\phi(\mathbf{x})$  only appears in the form of dot products:

$$L_d(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m k(\phi(\mathbf{x}_m)^\top \phi(\mathbf{x}_n))$$

$$y(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x}) + b$$

$$= \sum_{n=1}^N a_n t_n k(\phi(\mathbf{x}_n)^\top \phi(\mathbf{x})) + b$$

A blue oval encloses the equation  $k(\mathbf{x}, \mathbf{y}) = \phi(\mathbf{x})^\top \phi(\mathbf{y})$ . Three blue arrows point from the terms  $\phi(\mathbf{x}_m)^\top \phi(\mathbf{x}_n)$ ,  $\phi(\mathbf{x}_n)^\top \phi(\mathbf{x})$ , and  $\phi(\mathbf{x})^\top \phi(\mathbf{x})$  in the first equation to the corresponding terms in the enclosed equation.

$$k(\mathbf{x}, \mathbf{y}) = \phi(\mathbf{x})^\top \phi(\mathbf{y})$$

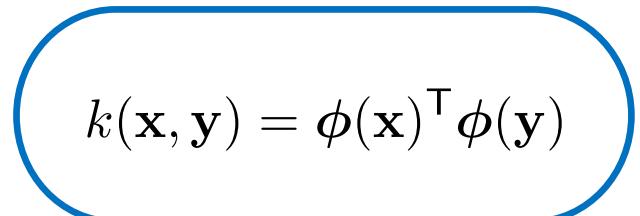
Define a **kernel function**  $k(\mathbf{x}, \mathbf{y}) = \phi(\mathbf{x})^\top \phi(\mathbf{y})$   
 $\Rightarrow$  Use the kernel instead of the dot product.

# The Kernel Trick

- On a closer look,  $\phi(\mathbf{x})$  only appears in the form of dot products:

$$L_d(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m k(\mathbf{x}_m, \mathbf{x}_n)$$

$$\begin{aligned} y(\mathbf{x}) &= \mathbf{w}^\top \phi(\mathbf{x}) + b \\ &= \sum_{n=1}^N a_n t_n k(\mathbf{x}_n, \mathbf{x}) + b \end{aligned}$$



$$k(\mathbf{x}, \mathbf{y}) = \phi(\mathbf{x})^\top \phi(\mathbf{y})$$

Define a kernel function  $k(\mathbf{x}, \mathbf{y}) = \phi(\mathbf{x})^\top \phi(\mathbf{y})$   
 $\Rightarrow$  Use the kernel instead of the dot product.

- $k(\cdot, \cdot)$  implicitly maps the data to some higher-dimensional space, without having to compute  $\phi(\mathbf{x})$ .

# When Can We Apply the Kernel Trick?

- In order for this to work,  $k(\cdot, \cdot)$  needs to define an implicit mapping.
- Formally
  - A function  $k(\mathbf{x}_1, \mathbf{x}_2) : \mathbb{R}^D \times \mathbb{R}^D \rightarrow \mathbb{R}$  is a **kernel function**, iff
    - There is a mapping  $\phi(\mathbf{x}) : \mathbb{R}^D \rightarrow \mathcal{H}$  such that
- *When will this be the case?*

- When is a function  $k(\mathbf{x}_1, \mathbf{x}_2)$  a valid kernel function? Two ways to check:

1. Every **Gram matrix**  $K$  of  $k$  is symmetric positive definite:

$$K = \begin{bmatrix} k(\mathbf{x}_1, \mathbf{x}_1) & \cdots & k(\mathbf{x}_1, \mathbf{x}_N) \\ \vdots & \ddots & \vdots \\ k(\mathbf{x}_N, \mathbf{x}_1) & \cdots & k(\mathbf{x}_N, \mathbf{x}_N) \end{bmatrix}$$

A matrix  $M$  is **positive definite** if all eigenvalues of  $K$  are positive.

- This is easy to verify for a given training set  $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$ .
- Unfortunately, it has to hold for **every possible** such set.

$\Rightarrow$  Very hard to prove in practice.

- When is a function  $k(\mathbf{x}_1, \mathbf{x}_2)$  a valid kernel function? Two ways to check:

2. We can construct valid kernels from other valid kernels:

- Given valid kernels  $k_1(\mathbf{x}, \mathbf{x}')$  and  $k_2(\mathbf{x}, \mathbf{x}')$ , the following combinations will also be valid

$$k(\mathbf{x}, \mathbf{x}') = c \cdot k_1(\mathbf{x}, \mathbf{x}')$$

$$k(\mathbf{x}, \mathbf{x}') = f(\mathbf{x})k_1(\mathbf{x}, \mathbf{x}')f(\mathbf{x}')$$

$$k(\mathbf{x}, \mathbf{x}') = \text{polynomial}(k_1(\mathbf{x}, \mathbf{x}')) \quad (\text{with nonnegative coefficients})$$

$$k(\mathbf{x}, \mathbf{x}') = \exp(k_1(\mathbf{x}, \mathbf{x}'))$$

$$k(\mathbf{x}, \mathbf{x}') = k_1(\mathbf{x}, \mathbf{x}') + k_2(\mathbf{x}, \mathbf{x}')$$

$$k(\mathbf{x}, \mathbf{x}') = k_1(\mathbf{x}, \mathbf{x}')k_2(\mathbf{x}, \mathbf{x}')$$

$$k(\mathbf{x}, \mathbf{x}') = k_3(\phi(\mathbf{x}), \phi(\mathbf{x}'))$$

$$k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^\top \mathbf{A} \mathbf{x}'$$

$\Rightarrow$  Much easier to apply in practice.

# New SVM Formulation

- Maximize

$$L_d(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m k(\mathbf{x}_m, \mathbf{x}_n)$$

under the constraints  $0 \leq a_n \leq C \quad \forall n$

$$\sum_{n=1}^N a_n t_n = 0$$

- Classify new data points using

$$y(\mathbf{x}) = \sum_{n=1}^N a_n t_n k(\mathbf{x}_n, \mathbf{x}) + b$$

## Example: Polynomial Kernel

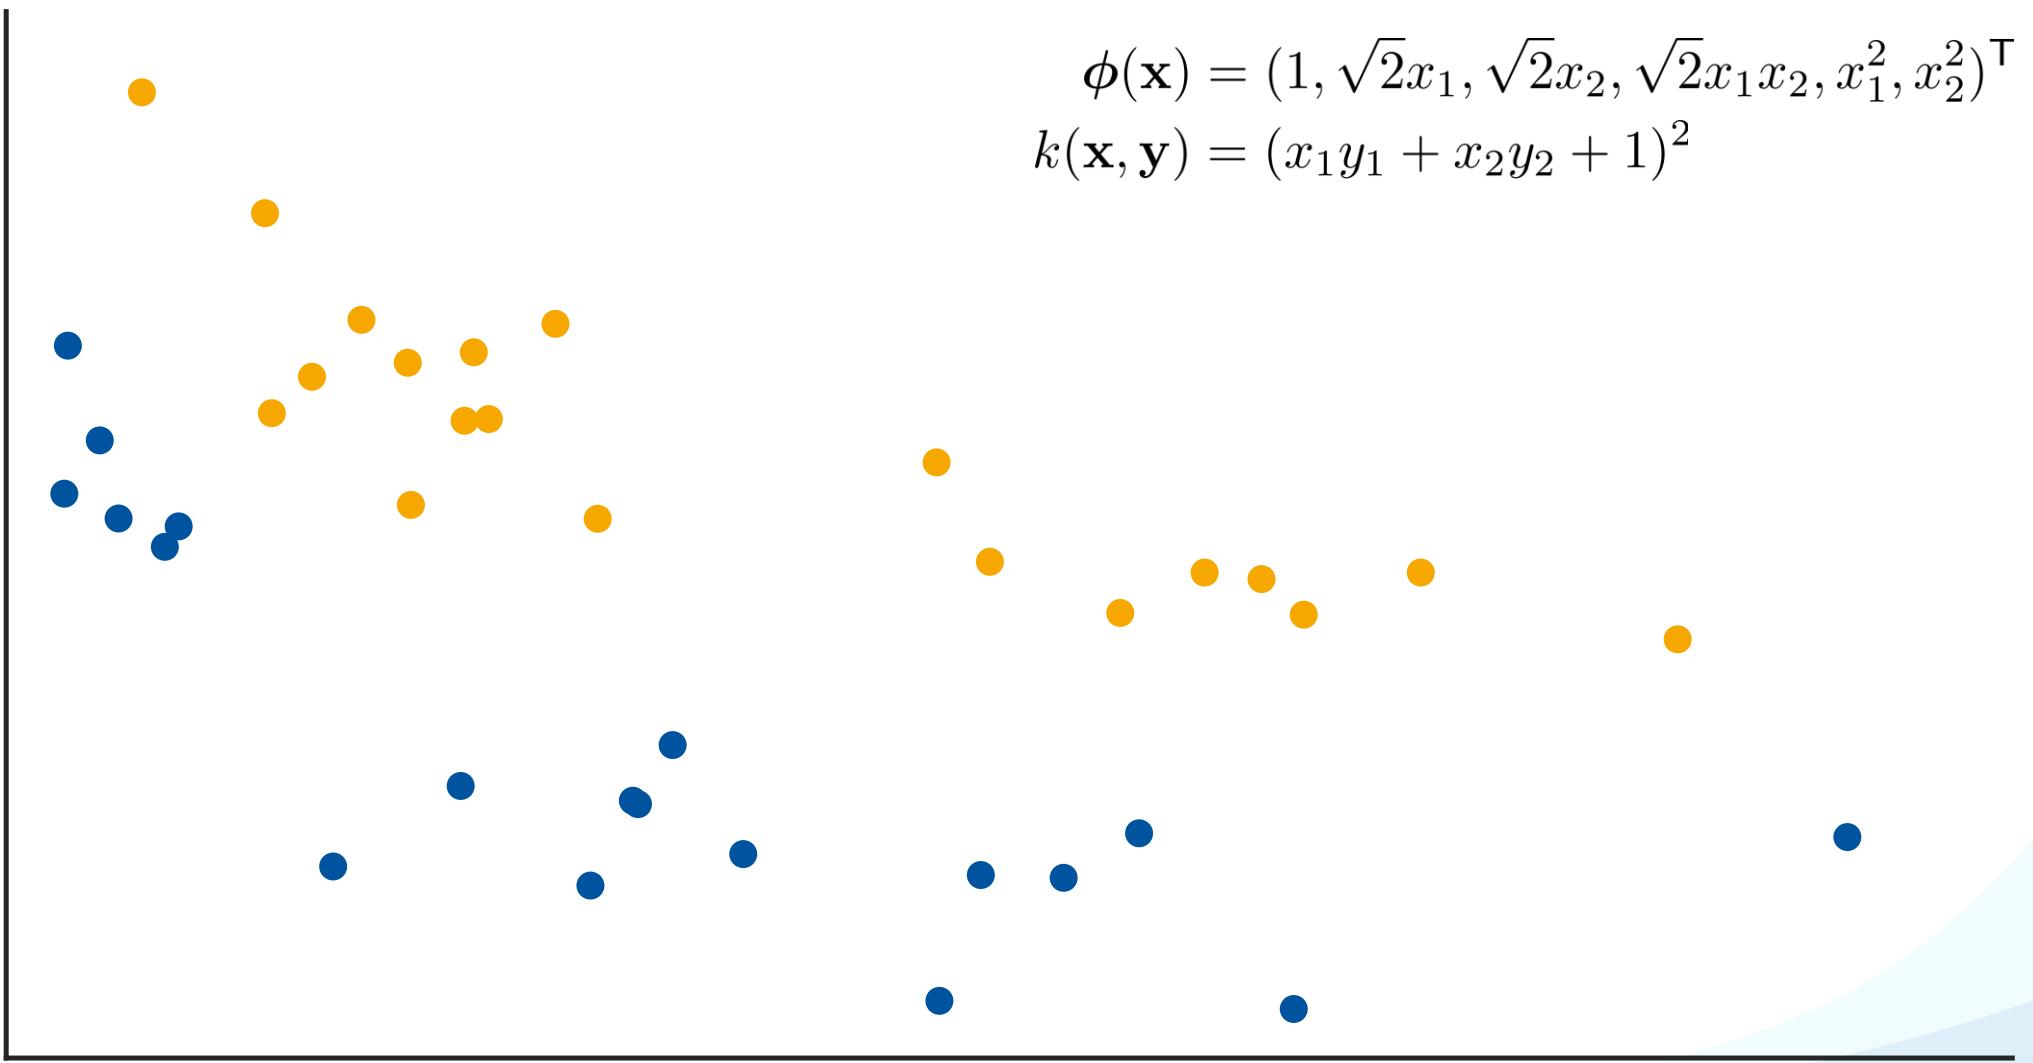
- We slightly adjust the polynomial basis function that we know:

$$\phi(\mathbf{x}) = (1, \sqrt{2}x_1, \sqrt{2}x_2, \sqrt{2}x_1x_2, x_1^2, x_2^2)^T$$

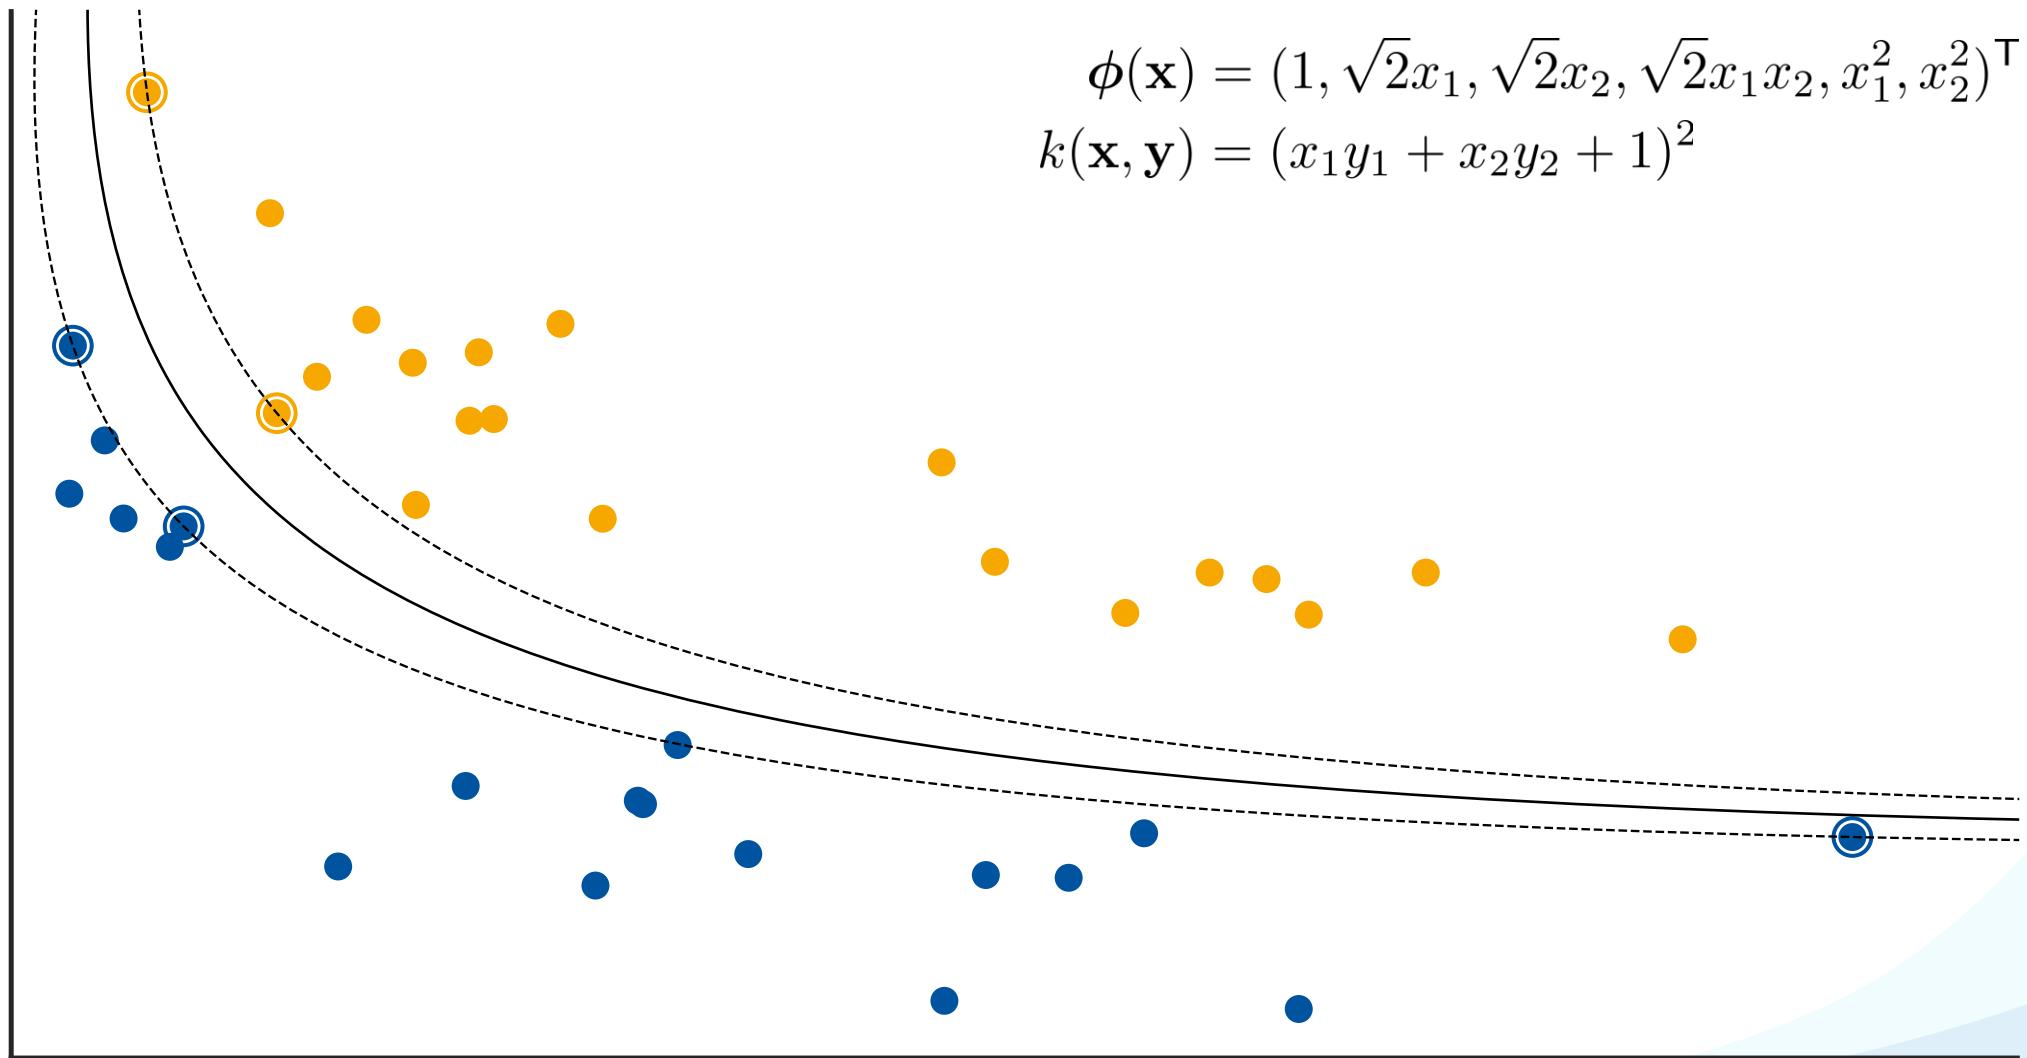
$$k(\mathbf{x}, \mathbf{y}) = (\mathbf{x}^T \mathbf{y} + 1)^2 = \phi(\mathbf{x})^T \phi(\mathbf{y})$$

- In fact,  $(\mathbf{x}^T \mathbf{y} + 1)^p$  is the kernel function for a polynomial of degree  $p$ .

## Example



# Example



## Advantages

- We can use high-dimensional or even infinite dimensional feature spaces
  - Since  $\phi(\mathbf{x})$  is never computed explicitly.
- We can work with non-vector space data
  - We can define kernel functions for arbitrary data types!
  - Graphs, Sets, Sequences, Histograms,
  - ...
- Simple to use and work very well in most cases

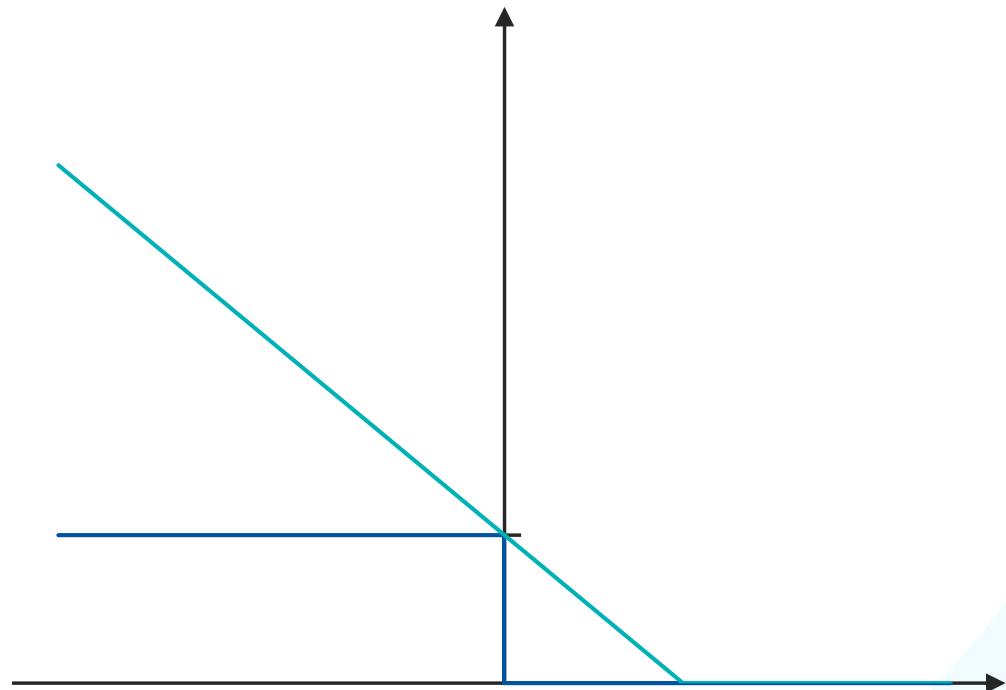
## Limitations

- Which kernel to choose?
  - **Model selection** problem
- How to choose kernel parameters?
  - **Hyperparameter optimization** problem, usually solved by performing a **grid search** over the validation set
- Evaluation speed scales linearly with number of support vectors

$$y(\mathbf{x}) = \sum_{n=1}^N a_n t_n k(\mathbf{x}_n, \mathbf{x}) + b$$

# Support Vector Machines

1. Maximum Margin Classification
2. Primal Formulation
3. Dual Formulation
4. Soft-Margin SVMs
5. Non-linear SVMs
6. **Error Function Analysis**



## Error Function Analysis

- We know how to formulate and optimize an SVM as a convex optimization problem:

$$\arg \min_{\mathbf{w}, b, \xi_n \in \mathbb{R}^+} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n$$

subject to the constraints

$$t_n y(\mathbf{x}_n) \geq 1 - \xi_n$$

# Error Function Analysis

- We know how to formulate and optimize an SVM as a convex optimization problem:

$$\arg \min_{\mathbf{w}, b, \xi_n \in \mathbb{R}^+} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n$$

*But what error function does this correspond to?*

subject to the constraints

$$t_n y(\mathbf{x}_n) \geq 1 - \xi_n$$

- Integrate the constraints into the objective function:
  - Rewrite as  $\xi_n \geq 1 - t_n y(\mathbf{x}_n)$
  - Thus, we obtain

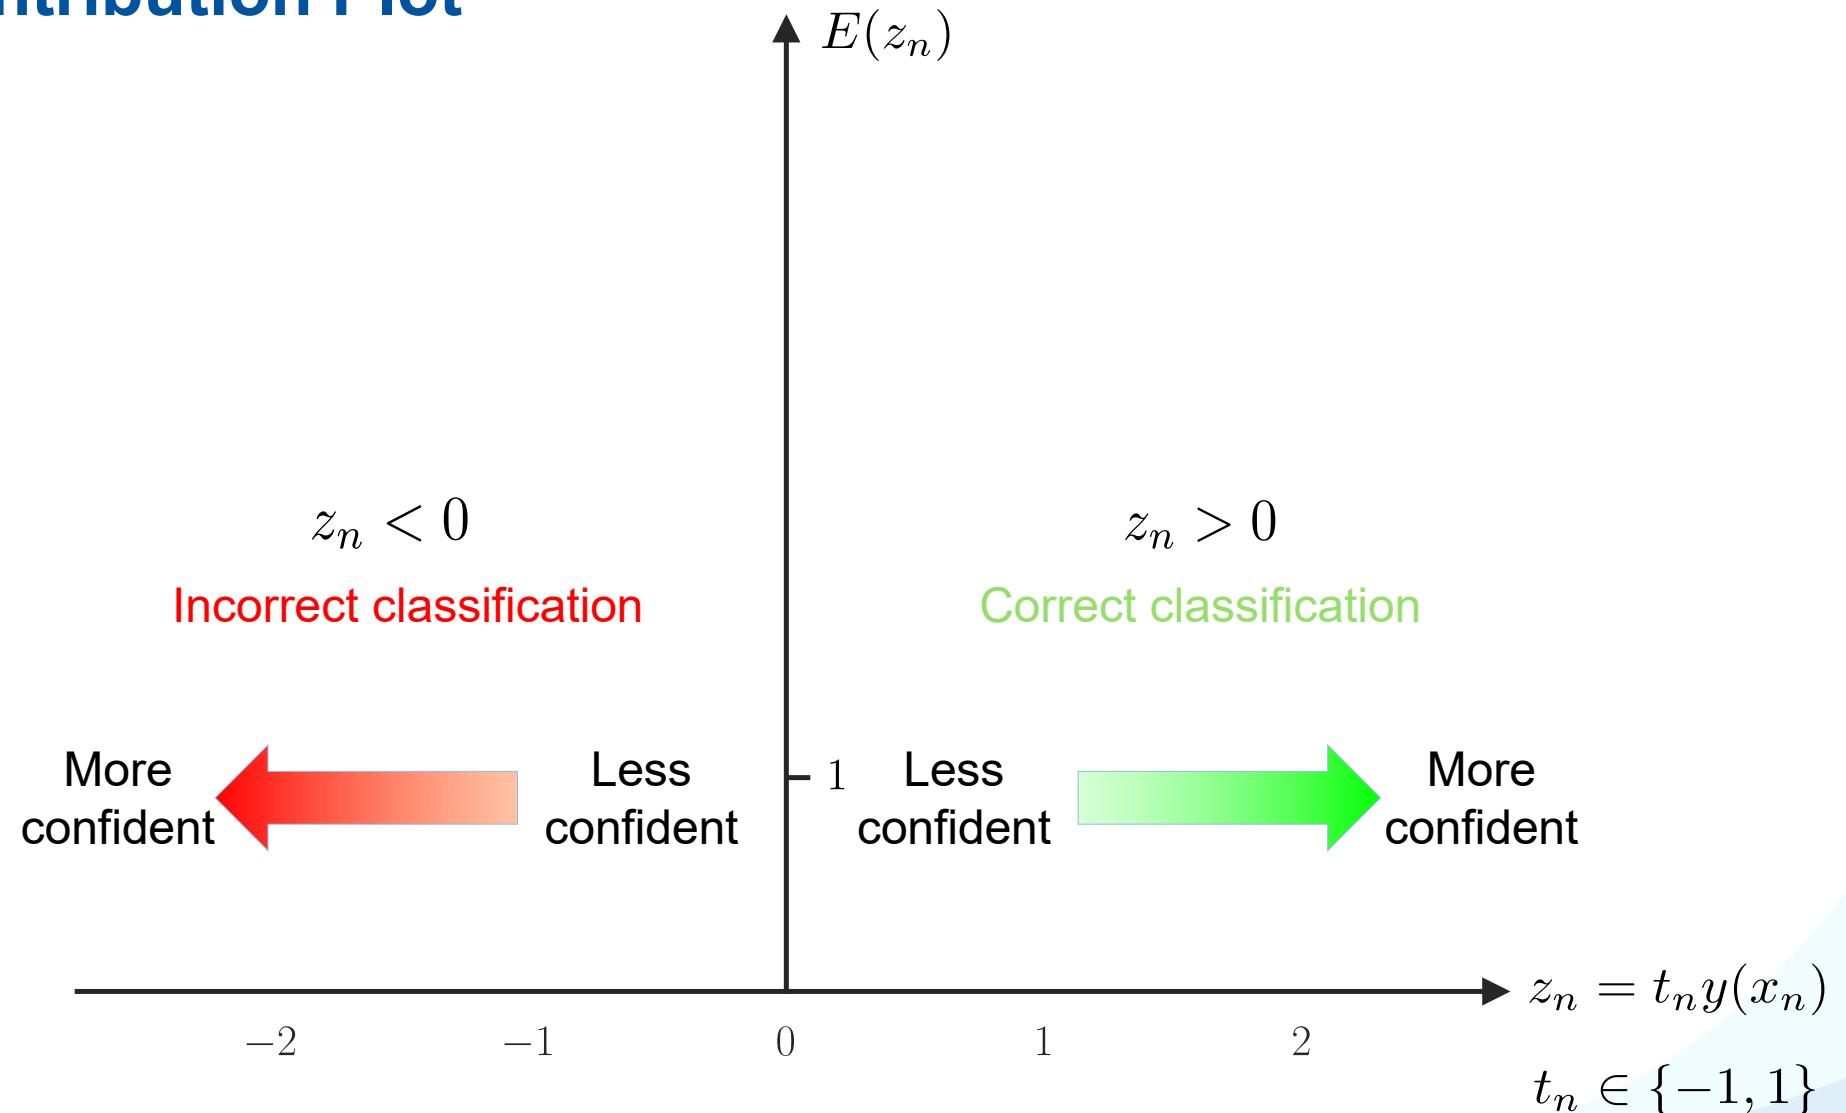
$$\min_{\mathbf{w}, b} E(\mathbf{w}) = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N [1 - t_n y(\mathbf{x}_n)]_+$$
$$[x]_+ \equiv \max\{x, 0\}$$

# The Hinge Loss

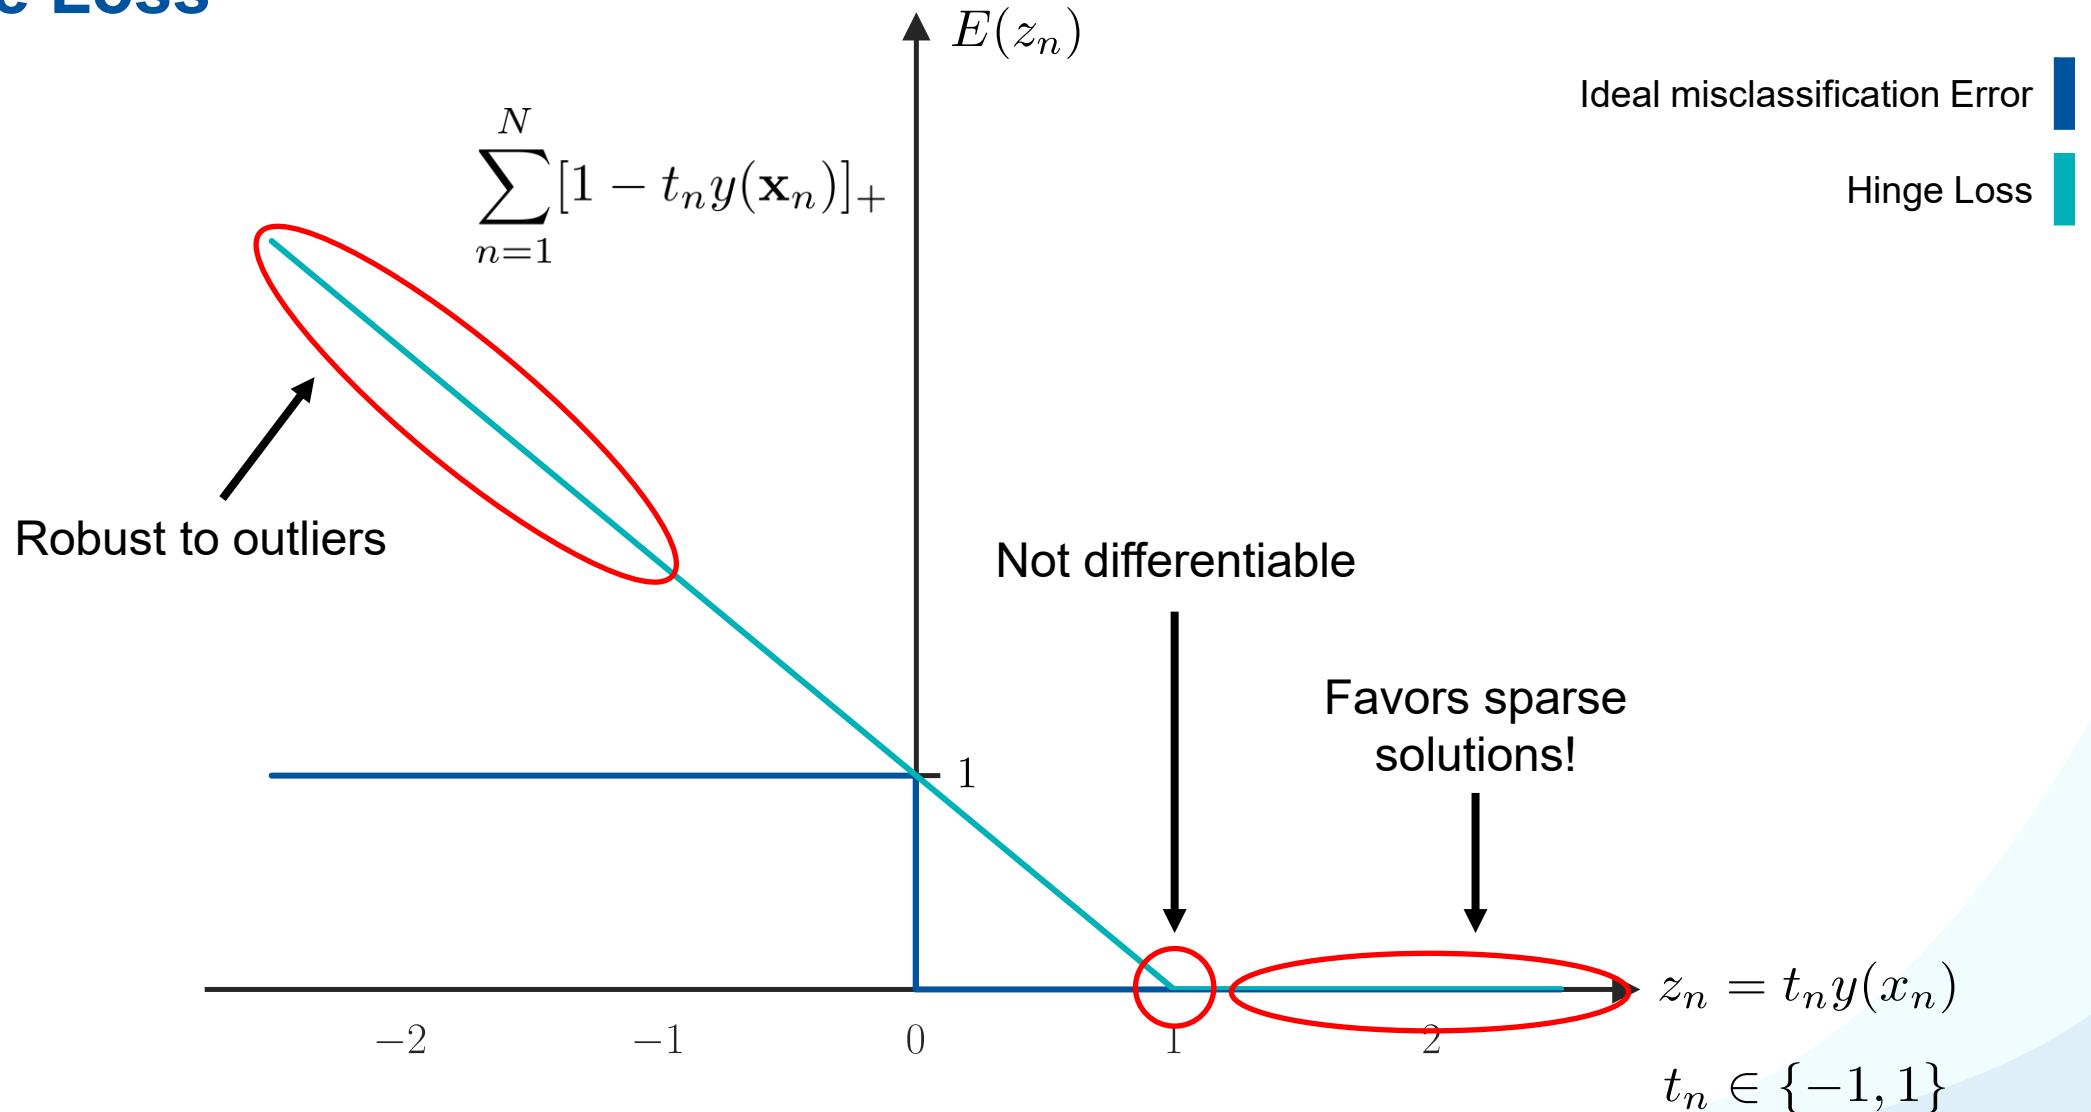
$$E(\mathbf{w}) = \underbrace{\frac{1}{2} \|\mathbf{w}\|^2}_{L_2 \text{ regularization}} + C \underbrace{\sum_{n=1}^N [1 - t_n y(\mathbf{x}_n)]_+}_{\text{Hinge loss}}$$

- Regularization bounds parameter size.
- Hinge Loss enforces sparsity:
  - Only a **subset of training data points** actually influences the decision boundary.
  - Still, all input dimensions are used.
- This formulation corresponds to an unconstrained optimization of a non-differentiable function.
  - Very efficient: stochastic (sub-)gradient descent.

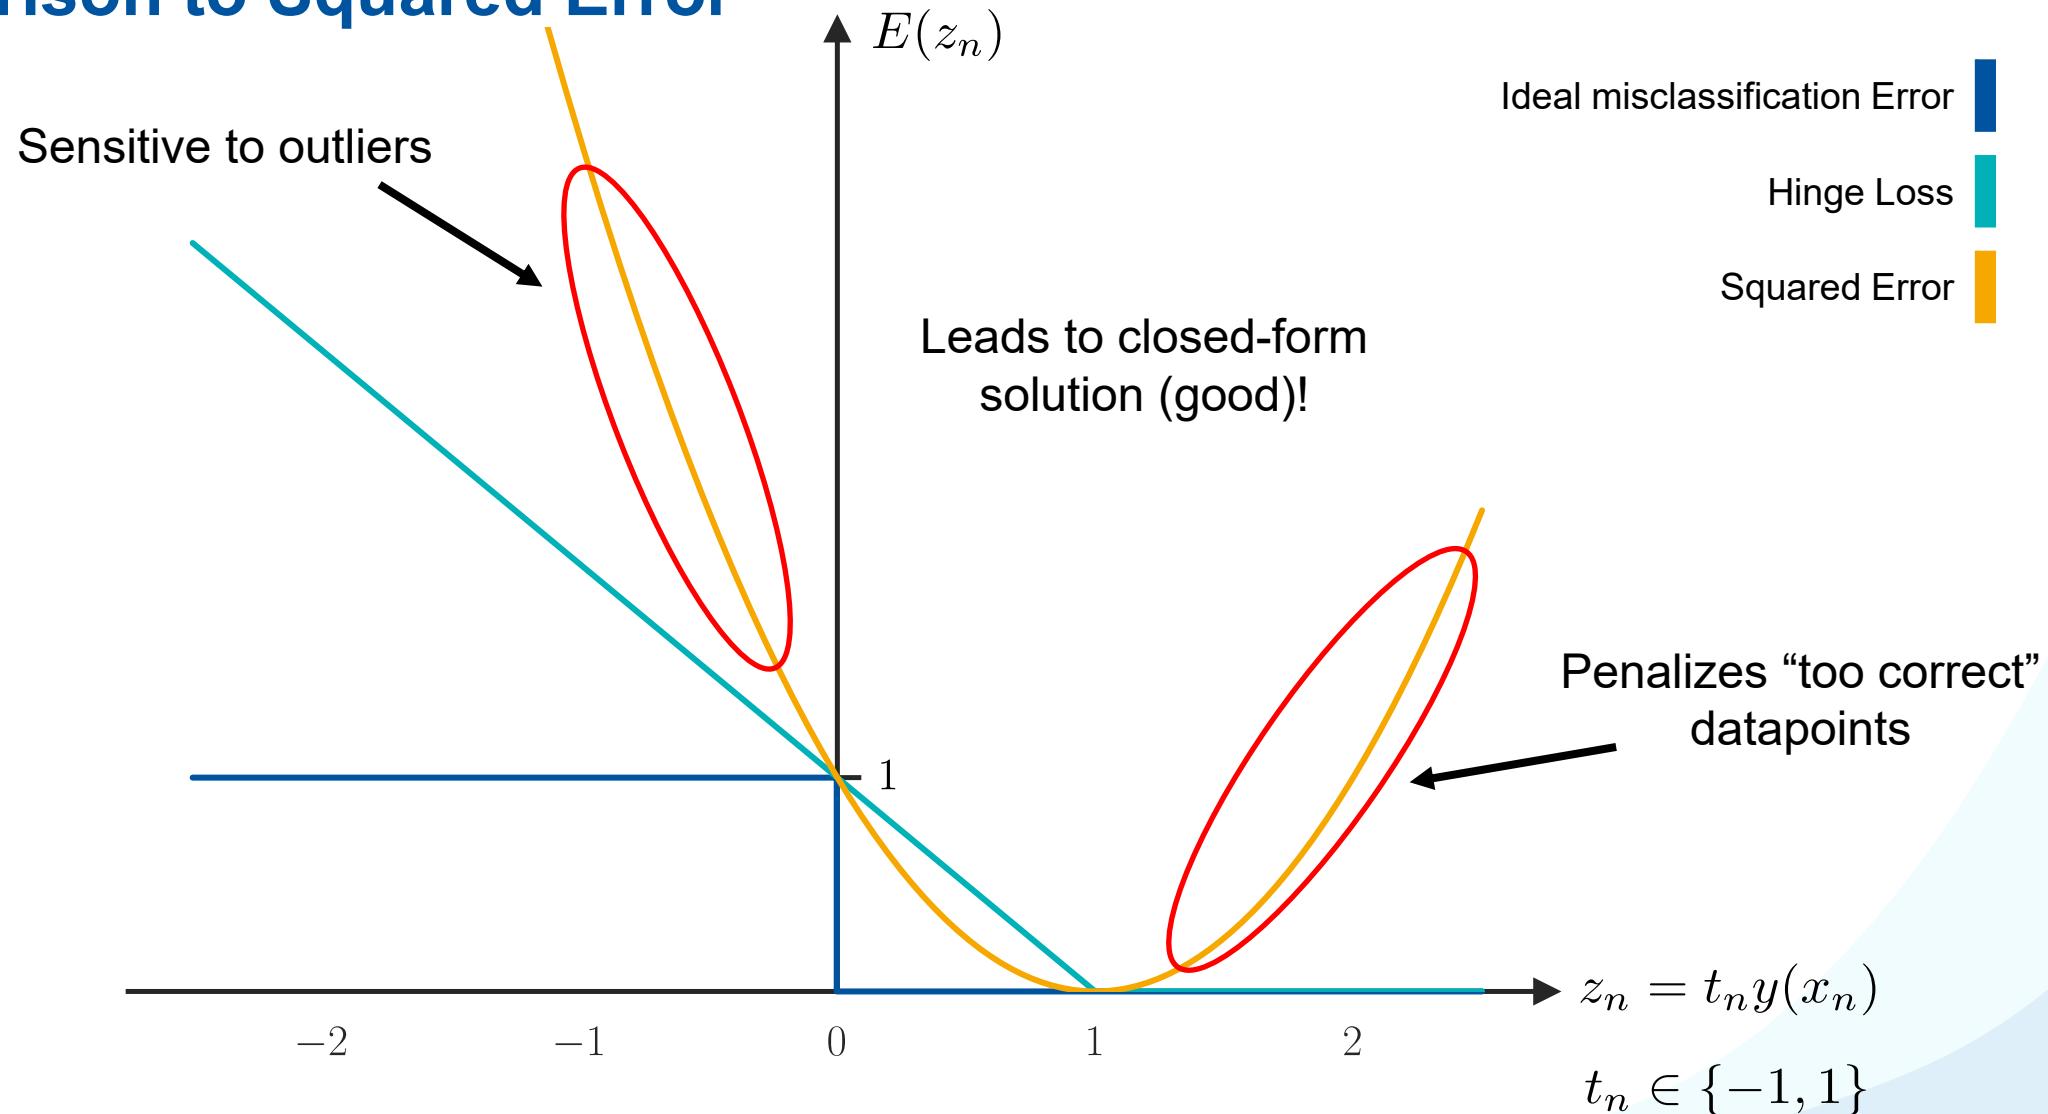
# Error Contribution Plot



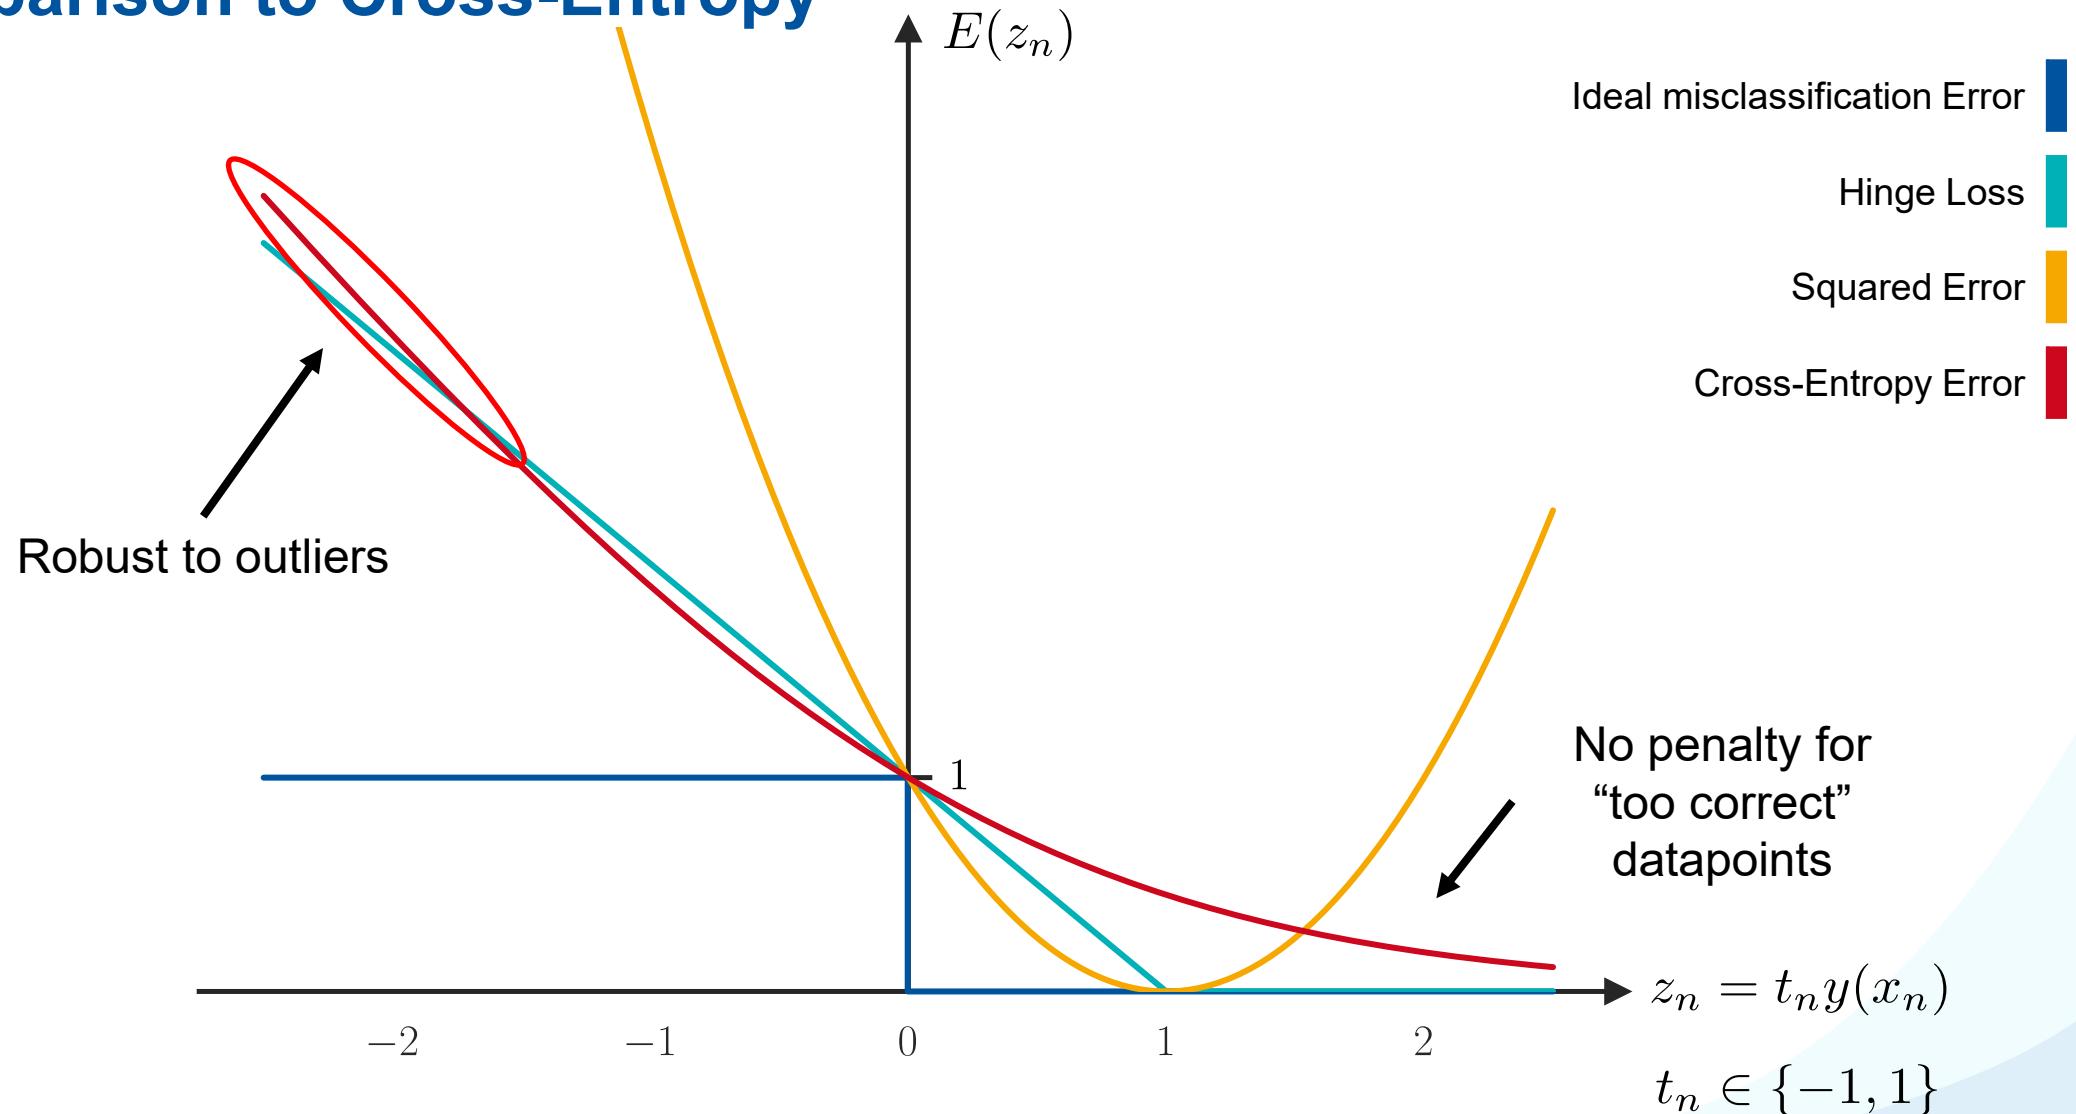
# Hinge Loss



## Comparison to Squared Error



## Comparison to Cross-Entropy



# Discussion: Hinge Loss

## Advantages

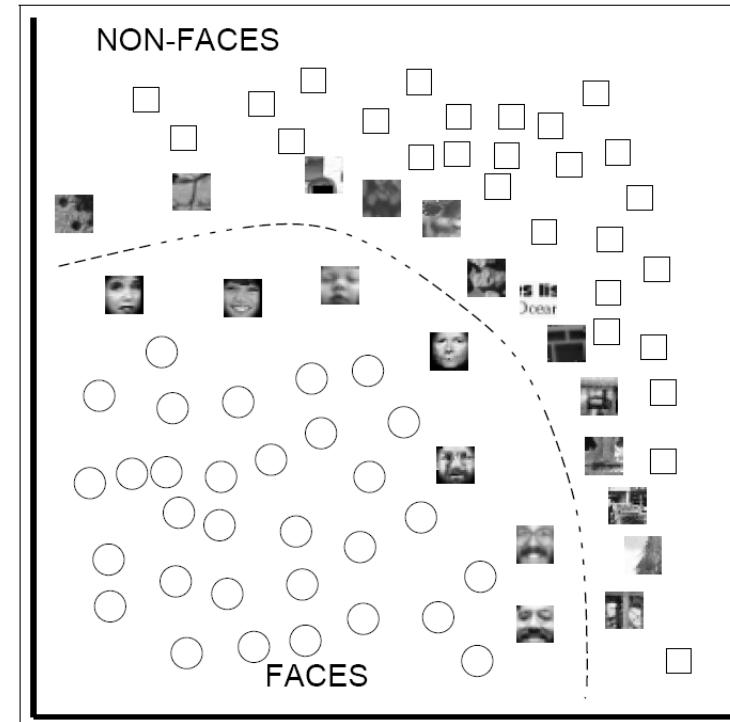
- Favors sparse solutions that only depend on a subset of training data points.
- Robust to outliers (only a linear penalty for misclassified points).
- Convex function, unique minimum exists.

## Limitations

- Not differentiable (cannot minimize this loss using standard gradient descent, but need to use [subgradient descent](#)).

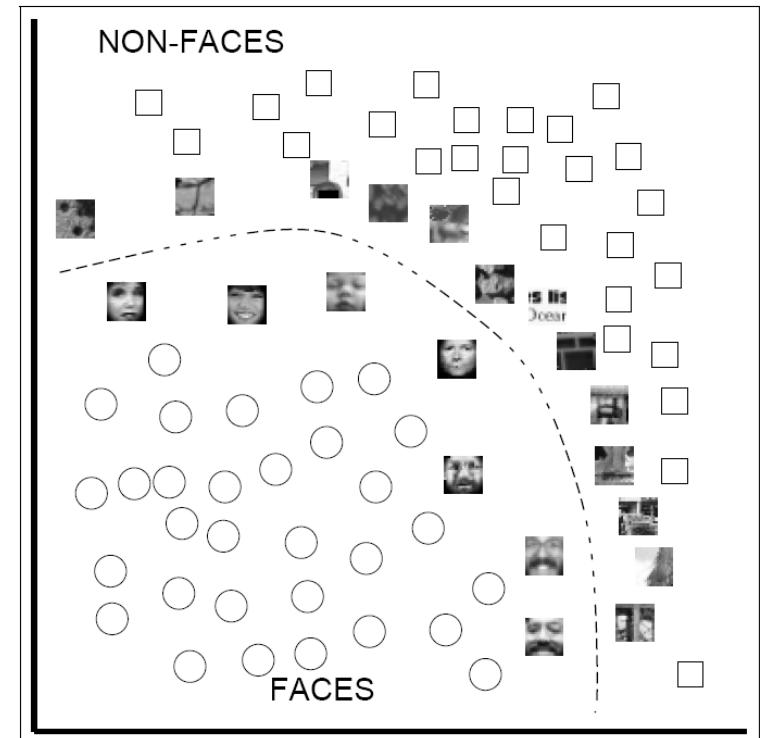
# Support Vector Machines

1. Maximum Margin Classification
2. Primal Formulation
3. Dual Formulation
4. Soft-Margin SVMs
5. Non-linear SVMs
6. Error Function Analysis
7. **Applications**



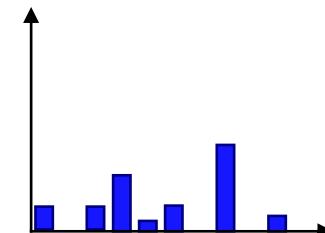
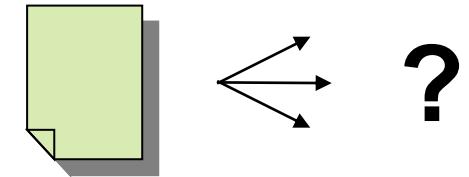
# Interpretation of Support Vectors

- Let's consider an image classification problem:
  - E.g., *faces* vs. *non-faces*
- The support vectors chosen by the SVM correspond to the hardest examples
  - The real faces closest to non-faces
  - The non-faces closest to real faces
- *We can visualize those support vectors again as images to get some insights into what the SVM reacts to...*



# Example Application: Text Classification

- Problem:
  - Classify a document in a number of categories
- Representation:
  - “Bag-of-words” approach
  - Histogram of word counts (on learned dictionary)
    - Very high-dimensional feature space ( $\sim 10,000$  dimensions)
    - Few irrelevant features
- This was one of the first applications of SVMs
  - T. Joachims (1997)

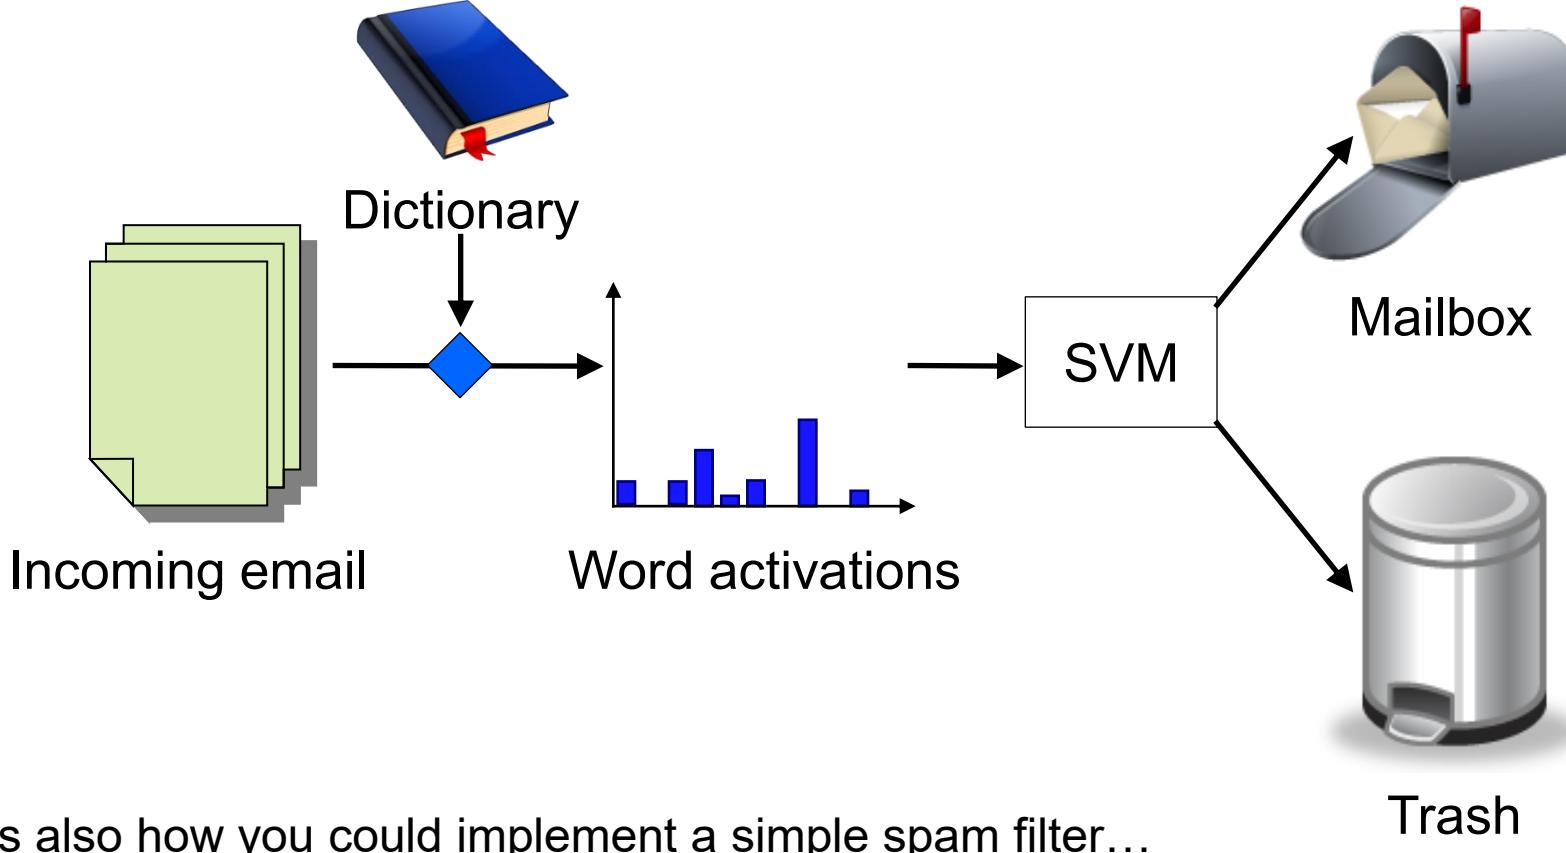


# Example Application: Text Classification

	Bayes	Rocchio	C4.5	k-NN	SVM (poly) degree $d =$					SVM (rbf) width $\gamma =$			
					1	2	3	4	5	0.6	0.8	1.0	1.2
earn	95.9	96.1	96.1	97.3	98.2	98.4	<b>98.5</b>	98.4	98.3	<b>98.5</b>	98.5	98.4	98.3
acq	91.5	92.1	85.3	92.0	92.6	94.6	<b>95.2</b>	95.2	95.3	95.0	95.3	95.3	<b>95.4</b>
money-fx	62.9	67.6	69.4	78.2	66.9	72.5	75.4	74.9	<b>76.2</b>	74.0	75.4	<b>76.3</b>	75.9
grain	72.5	79.5	89.1	82.2	91.3	93.1	<b>92.4</b>	91.3	89.9	<b>93.1</b>	91.9	91.9	90.6
crude	81.0	81.5	75.5	85.7	86.0	87.3	88.6	<b>88.9</b>	87.8	<b>88.9</b>	89.0	88.9	88.2
trade	50.0	77.4	59.2	77.4	69.2	75.5	76.6	77.3	<b>77.1</b>	76.9	78.0	<b>77.8</b>	76.8
interest	58.0	72.5	49.1	74.0	69.8	63.3	67.9	73.1	<b>76.2</b>	74.4	75.0	<b>76.2</b>	76.1
ship	78.7	83.1	80.9	79.2	82.0	85.4	86.0	<b>86.5</b>	86.0	<b>85.4</b>	86.5	87.6	87.1
wheat	60.6	79.4	85.5	76.6	83.1	84.5	85.2	<b>85.9</b>	83.8	<b>85.2</b>	85.9	85.9	85.9
corn	47.3	62.2	87.7	77.9	86.0	86.5	85.3	<b>85.7</b>	83.9	<b>85.1</b>	85.7	85.7	84.5
microavg.	<b>72.0</b>	<b>79.9</b>	<b>79.4</b>	<b>82.3</b>	84.2	85.1	85.9	86.2	85.9	86.4	86.5	86.3	86.2
					combined: <b>86.0</b>					combined: <b>86.4</b>			

- Results by T. Joachims (1997)

## Example Application: Text Classification

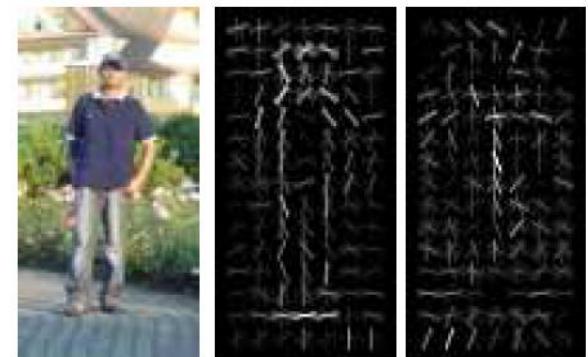


# Example Application: Visual Object Detection

## Sliding-Window Approach



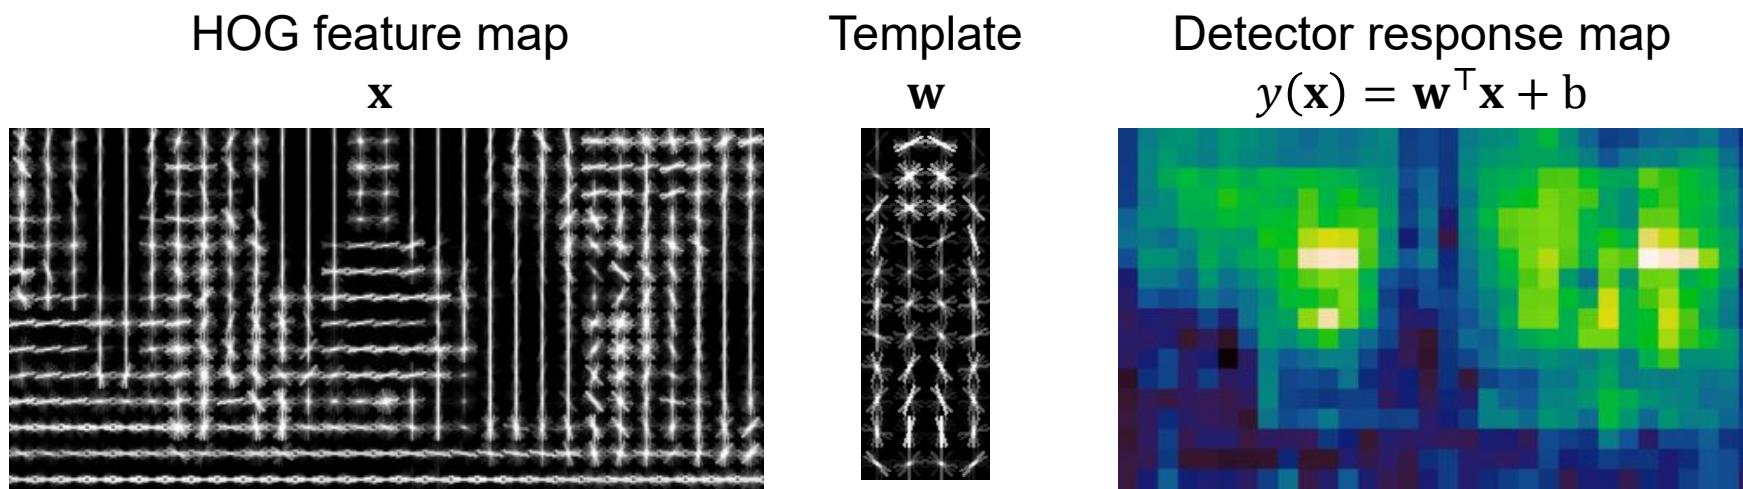
- Densely scan an image with a sliding window
- Object detection by per-window classification
  - E.g., use a histogram of gradients (HOG) feature representation
  - Train a linear SVM for *object* vs. *non-object* classification.



# Example Application: Pedestrian Detection with SVMs

- Train a pedestrian template using a linear SVM
- At test time, convolve feature map with learned template  $w$ 
  - Linear SVM classification function  $\Leftrightarrow$  convolution with “template”  $w$

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$



N. Dalal, B. Triggs, Histograms of Oriented Gradients for Human Detection, CVPR 2005

## Example Application: Pedestrian Detection with SVMs



N. Dalal, B. Triggs, Histograms of Oriented Gradients for Human Detection, CVPR 2005

## References and Further Reading

- More information about SVMs is available in Chapter 7.1 of Bishop's book.

Christopher M. Bishop  
Pattern Recognition and Machine Learning  
Springer, 2006

