

Elements of Machine Learning & Data Science

Winter semester 2025/26


Lecture 16 – Neural Networks Basics

16.12.2025


Prof. Bastian Leibe

Machine Learning Topics

- 8. Introduction to ML
- 9. Probability Density Estimation
- 10. Linear Discriminants
- 11. Linear Regression
- 12. Logistic Regression
- 13. Support Vector Machines
- 14. Neural Network Basics**



Multi-Layer Perceptrons




Backpropagation

$$\frac{1}{2}(y(\mathbf{x}) - t)^2$$

$$[1 - ty(\mathbf{x})]_+ - \sum_k \left(\mathbb{I}(t = k) \ln \frac{\exp(y_k(\mathbf{x}))}{\sum_j \exp(y_j(\mathbf{x}))} \right)$$
$$\frac{1}{2} \|\mathbf{w}\|^2$$


Losses & Regularizers



Stochastic Gradient Descent


Neural Network Basics

1. Perceptrons
2. Multi-Layer Perceptrons
3. Loss Functions
4. Backpropagation
5. Computational Graphs
6. Stochastic Gradient Descent




Perceptrons

- Inspired by biological neurons.
- The output is determined by the activation of the input nodes and a set of weights connecting input and output layers.



Basic Perceptron

- Input Layer:
 - Hand-designed features
- Outputs:
 - Linear outputs
$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$
 - Logistic outputs
$$y(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + w_0)$$
- Learning: finding optimal weights \mathbf{w} .



Multi-Class Networks


- One output node per class:
 - Linear outputs

$$y_k(\mathbf{x}) = \sum_{i=0}^D w_{ki} \mathbf{x}_i$$

- Logistic outputs

$$y_k(\mathbf{x}) = \sigma \left(\sum_{i=0}^D w_{ki} \mathbf{x}_i \right)$$

- We can do **multidimensional linear regression** or **multiclass classification** this way.




Non-Linear Basis Functions

- Apply a (fixed) mapping $\phi(\mathbf{x})$ to inputs:
 - Linear outputs

$$y_k(\mathbf{x}) = \sum_{j=0}^M w_{kj} \phi_j(\mathbf{x})$$

- Logistic outputs

$$y_k(\mathbf{x}) = \sigma \left(\sum_{j=0}^M w_{kj} \phi_j(\mathbf{x}) \right)$$




Connections to linear discriminants

- All of this should feel very familiar.
- Perceptrons are generalized linear discriminants!
- What does that mean?
 - We have the same limitations as before.
 - Can model any separable function perfectly, given the right input features.
 - For some tasks, this may require an exponential number of input features.

⇒ *It is the feature design that solves the task!*


Limitations so far

- Generalized linear discriminants (including perceptrons) are very limited.
 - A linear classifier cannot solve certain problems (e.g., XOR).
 - However, with a non-linear classifier based on suitable features, the problem becomes solvable.
 - So far, we have designed the features and kernels by hand.
- ⇒ How can we *learn* good feature representations?




Neural Network Basics

1. Perceptrons
2. **Multi-Layer Perceptrons**
3. Loss Functions
4. Backpropagation
5. Computational Graphs
6. Stochastic Gradient Descent



Multi-Layer Perceptrons

- Perceptrons are limited by having a fixed input mapping.




Multi-Layer Perceptrons

- Perceptrons are limited by having a fixed input mapping.
- Replace it with a **hidden layer** that learns suitable features.
- Output of hidden layer is input to next layer.
- Each layer computes a matrix multiplication and applies an elementwise **activation function** $g(\cdot)$:

$$\mathbf{z}(\mathbf{x}) = g^{(1)} \left(\mathbf{W}^{(1)} \mathbf{x} \right)$$

$$\mathbf{y}(\mathbf{x}) = g^{(2)} \left(\mathbf{W}^{(2)} \mathbf{z}(\mathbf{x}) \right)$$

- Key step: *Now we also make the earlier layer learnable!*
- This is known as a **Multi-Layer Perceptron (MLP)**.




General Network Structure

- Multi-Layer Perceptron model:

$$y_k(\mathbf{x}) = g^{(2)} \left(\sum_{i=0}^M w_{ki}^{(2)} g^{(1)} \left(\sum_{j=0}^D w_{ij}^{(1)} x_j \right) \right)$$


- Usually, each layer adds a bias term.
- Activation functions between layers should be non-linear.
 - For example: $g^{(2)}(a) = \sigma(a)$, $g^{(1)}(a) = \max\{a, 0\}$
 - With linear activations, successive layers would still compute a linear function.
- The hidden layer can have an arbitrary number of nodes.
- There can also be multiple hidden layers.



MLPs are Universal Approximators

$$y_k(\mathbf{x}) = g^{(2)} \left(\sum_{i=0}^M w_{ki}^{(2)} g^{(1)} \left(\sum_{j=0}^D w_{ij}^{(1)} x_j \right) \right)$$


- **Universal Approximator Theorem:**
 - A network with one hidden layer can approximate any continuous function of a compact domain arbitrarily well (assuming sufficient hidden nodes).
- ⇒ *Way more powerful than linear models!*



Learning with Hidden Units


- We now have a model that contains multiple layers of adaptive non-linear hidden units.
- How can we train such models?
 - Need to train *all* weights, not just last layer.
 - Learning the weights to the hidden units = learning features.
 - We don't know what the hidden units should do.
- Basic Idea: [Gradient Descent](#).

This is the main challenge of deep learning!



Neural Network Basics


1. Perceptrons
2. Multi-Layer Perceptrons
3. **Loss Functions**
4. Backpropagation
5. Computational Graphs
6. Stochastic Gradient Descent



Loss Functions

- We train Neural Networks by minimizing an **error function**
 - In principle, any differentiable objective function can be used here.
 - Typically, we use a combination of a **loss function** $L(t, y(\mathbf{x}))$ and a **regularizer** $\Omega(\mathbf{w})$:

$$E(\mathbf{w}) = \sum_{n=1}^N L(t_n, y(\mathbf{x}_n; \mathbf{w})) + \lambda\Omega(\mathbf{w})$$



Examples of Loss Functions


- We can use any of the loss functions we have seen so far to achieve different effects:

- L₂ loss (Squared Error)**

$$L(t, y(\mathbf{x})) = \frac{1}{2}(y(\mathbf{x}) - t)^2$$

- Binary Cross-Entropy loss**

$$L(t, y(\mathbf{x})) = -(t \ln \sigma(y(\mathbf{x})) + (1 - t) \ln(1 - \sigma(y(\mathbf{x}))))$$



⇒ Least-squares regression / classif.

- Hinge loss**

$$L(t, y(\mathbf{x})) = [1 - ty(\mathbf{x})]_+$$

⇒ Logistic regression

- Multi-Class Cross-Entropy loss**

$$L(t, \mathbf{y}(\mathbf{x})) = - \sum_k \left(\mathbb{I}(t = k) \ln \frac{\exp(y_k(\mathbf{x}))}{\sum_j \exp(y_j(\mathbf{x}))} \right)$$

⇒ SVM classification

⇒ Multi-class probabilistic classification

Examples of Regularization Terms

- Similarly, we can use any of the regularization terms we have seen so far:


- L_2 regularizer (“Weight Decay”)

$$\Omega(\mathbf{w}) = \frac{1}{2} \|\mathbf{w}\|_2^2$$

- L_1 regularizer

$$\Omega(\mathbf{w}) = \frac{1}{2} \|\mathbf{w}\|_1$$

- Since Neural Networks have many parameters, regularization becomes an important consideration.
 - Many of the more advanced NN “training tricks” can also be understood as a form of regularization




\Rightarrow Prevents overfitting

\Rightarrow Enforces sparsity (feature selection)


Neural Network Basics

1. Perceptrons
2. Multi-Layer Perceptrons
3. Loss Functions
4. **Backpropagation**
5. Computational Graphs
6. Stochastic Gradient Descent



Backpropagation

- We know a flexible model that is able to learn features.
- We also know how to compute an error estimate.
- Now we need to compute the gradients with respect to our parameters.




Approach 1: Naïve Analytical Differentiation

- Compute the gradients of each variable analytically.
- Scalar case is straightforward:

$$\Delta z = \frac{\partial z}{\partial y} \Delta y \quad \Delta y = \frac{\partial y}{\partial x} \Delta x$$

$$\Delta z = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} \Delta x$$

$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x}$$




Approach 1: Naïve Analytical Differentiation

- Compute the gradients of each variable analytically.
- Scalar case is straightforward.
- Multi-dimensional case: **Total derivative**
 - Need to sum over all paths to target variable:

$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y_1} \frac{\partial y_1}{\partial x} + \frac{\partial z}{\partial y_2} \frac{\partial y_2}{\partial x} + \dots = \sum_{i=1}^k \frac{\partial z}{\partial y_i} \frac{\partial y_i}{\partial x}$$

- With increasing depth, there will be exponentially many paths!




Approach 2: Incremental Analytical Differentiation

- Idea: compute the gradients layer by layer.
- Each layer below builds upon the results of the layer above.
- The gradient is propagated backwards through the layers.
- This is the [backpropagation](#) algorithm.


$$\frac{\partial E(\mathbf{w})}{\partial y_i} \quad \begin{array}{l} \downarrow \\ \frac{\partial E(\mathbf{w})}{\partial z_i} \end{array} \quad \begin{array}{l} \searrow \\ \frac{\partial E(\mathbf{w})}{\partial w_{ij}^{(2)}} \end{array}$$
$$\frac{\partial E(\mathbf{w})}{\partial x_i} \quad \begin{array}{l} \searrow \\ \frac{\partial E(\mathbf{w})}{\partial w_{ij}^{(1)}} \end{array}$$

Backpropagation



Automatic Differentiation

- Convert the network into a computational graph.
- Each layer/module just needs to specify how it affects the forward and backward passes.
- Apply **backpropagation** on the computational graph.
- Very general algorithm, used in today's Deep Learning packages.



Intuition

Let's assume we have a chain of functions

$$y = f_1(x)$$


$$z = f_2(y) = f_2(f_1(x))$$

How will the value of z change if we change x (i.e., $\frac{\partial z}{\partial x}$)? Use the chain rule!

Find out how y will affect z , then find out how x will affect y

Formula:


$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x}$$






Backpropagation algorithm

Backpropagation calculates the **gradient** of the **loss function** with respect to each weight by propagating the error gradient backwards through the network.

$$\frac{\partial E}{\partial x} = \frac{\partial E}{\partial y} \frac{\partial y}{\partial x} = \text{output gradient} \times \text{local gradient}$$



Example: Some Common Functions



A Larger Example

$$b = xy + z$$

Goal: find $\frac{\partial b}{\partial x}$, $\frac{\partial b}{\partial y}$ and $\frac{\partial b}{\partial z}$ when

$$x = 3, y = -5, z = 2$$

Step 1: Forward pass

Define intermediate values

$$a = xy = -15$$

$$b = a + z = -13$$

Step 2: Backward pass


$$\frac{\partial b}{\partial a} = 1$$

$$\frac{\partial b}{\partial z} = 1$$

$$\frac{\partial b}{\partial x} = \frac{\partial b}{\partial a} \frac{\partial a}{\partial x} = 1 \times (-5) = -5$$

$$\frac{\partial b}{\partial y} = \frac{\partial b}{\partial a} \frac{\partial a}{\partial y} = 1 \times 3 = 3$$


Done!



Multiple outputs

If there are more than one out-going edges at a node, we can just **sum the gradients!**

$$\frac{dL}{dx} = \frac{dL}{df_1} \cdot \frac{df_1}{dx} + \frac{dL}{df_2} \cdot \frac{df_2}{dx}$$



Discussion Backpropagation

Advantages


- Very general algorithm, widely used.
- Efficiently computes all gradients in the network using dynamic programming.
- The same concept can be applied to any differentiable function.
 - This makes it possible to define other types of layers.

Limitations

- Efficient evaluation of backpropagation requires storing all unit activations from forward pass.
 - The amount of memory necessary for this imposes a practical limit on the size of the network.
- Successful learning relies on the gradients to be propagated to the early network layers.
 - Numerical challenges may arise here.


Neural Network Basics

1. Perceptrons
2. Multi-Layer Perceptrons
3. Loss Functions
4. Backpropagation
5. Computational Graphs
6. Stochastic Gradient Descent



Backpropagation on Computational Graphs


- Neural networks define a computational graph computing some error L .
- Compute the weight's gradients via **backpropagation**.
- Only need to define two functions per node:
`forward(...)` to compute the output.
`backward(...)` to compute the gradient.



Example: Logistic Regression

Model output: $p(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + w_0) = 1/(1 + e^{-(w_0 + w_1 x_1 + w_2 x_2)})$


Loss function: $L(p, t) = \text{Cross-Entropy}(p, t) = -t \ln(p) - (1 - t) \ln(1 - p)$



Example: Logistic Regression

Model output: $p(\mathbf{x}) = \sigma(\mathbf{w}^\top \mathbf{x} + w_0)$

Loss function: $L(p, t) = \text{Cross-Entropy}(p, t)$



Grouping may simplify derivatives:

$$\frac{\partial \sigma(a)}{\partial a} = \sigma(a)(1 - \sigma(a))$$

$$\frac{\partial L(p, t)}{\partial p} = -\frac{t}{p} + \frac{1-t}{1-p}$$

Example: Neural Network

With computational graphs for matrix and tensor operations, we can apply backpropagation to neural networks.

Recall the MLP example


$$\mathbf{z} = g^{(1)}(\mathbf{W}^{(1)} \mathbf{x} + \mathbf{b}^{(1)})$$

$$\mathbf{y} = g^{(2)}(\mathbf{W}^{(2)} \mathbf{z} + \mathbf{b}^{(2)})$$

where


$$\mathbf{x} = \begin{bmatrix} x_1 \\ \vdots \\ x_d \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_k \end{bmatrix} \quad \mathbf{b}^{(1)} = \begin{bmatrix} W_{10}^{(1)} \\ \vdots \\ W_{h0}^{(1)} \end{bmatrix} \quad \mathbf{b}^{(2)} = \begin{bmatrix} W_{10}^{(2)} \\ \vdots \\ W_{k0}^{(2)} \end{bmatrix}$$

$$\mathbf{W}^{(1)} = \begin{bmatrix} W_{11}^{(1)} & \dots & W_{1d}^{(1)} \\ \vdots & & \vdots \\ W_{h1}^{(1)} & \dots & W_{hd}^{(1)} \end{bmatrix} \quad \mathbf{W}^{(2)} = \begin{bmatrix} W_{11}^{(2)} & \dots & W_{1h}^{(2)} \\ \vdots & & \vdots \\ W_{k1}^{(2)} & \dots & W_{kh}^{(2)} \end{bmatrix}$$




Example: Neural Network

The following computational graph represents such an MLP:



Neural Network Basics

1. Perceptrons
2. Multi-Layer Perceptrons
3. Loss Functions
4. Backpropagation
5. Computational Graphs
6. **Stochastic Gradient Descent**





Stochastic Gradient Descent

- Now that we have the gradients, we need to update the weights.
- We already know the basic equation for this
 - (1st-order) Gradient Descent

$$w_{kj}^{(\tau+1)} = w_{kj}^{(\tau)} - \eta \frac{\partial E(\mathbf{w})}{\partial w_{kj}} \Big|_{\mathbf{w}^{(\tau)}}$$

- Remaining Questions:
 - On what data do we want to apply this?
 - How should we choose the learning rate η ?



Stochastic vs. Batch Learning

- **Batch Learning**
 - Process the full dataset in one batch.
 - Compute the gradient based on all training examples.
- **Stochastic Learning**
 - Choose a single example from the training set.
 - Compute the gradient only based on this example.
 - This estimate will generally be noisy, which has some advantages.

$$w_{kj}^{(\tau+1)} = w_{kj}^{(\tau)} - \eta \frac{\partial E(\mathbf{w})}{\partial w_{kj}} \Big|_{\mathbf{w}^{(\tau)}}$$

$$E(\mathbf{w}) = \sum_{n=1}^N E_n(\mathbf{w})$$

$$w_{kj}^{(\tau+1)} = w_{kj}^{(\tau)} - \eta \frac{\partial E_n(\mathbf{w})}{\partial w_{kj}} \Big|_{\mathbf{w}^{(\tau)}}$$

Batch Learning

- Conditions of convergence are well understood.
- Many acceleration techniques only work in batch learning.
- Theoretical analysis of the weight dynamics and convergence rates are simpler.

Stochastic Learning

- Usually much faster than batch learning.
- Often results in better solutions.
- Can be used for tracking changes when the target distribution shifts.


Middle ground: Minibatches

Minibatches

- Idea
 - Process only a small batch of training examples together. $\mathcal{B} = \{(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_B, t_B)\} \subset \mathcal{D}$
 - Start with a small batch size & increase it as training proceeds.
- Advantages
 - Gradients will be more stable than for stochastic gradient descent, but still faster to compute than with batch learning.
 - Take advantage of redundancies in the training set.
 - Matrix operations are more efficient than vector operations.
- Caveat
 - Need to normalize error function by the minibatch size to use the same learning rate between minibatches

$$E(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N L(t_n, y(\mathbf{x}_n; \mathbf{w})) + \frac{\lambda}{N} \Omega(\mathbf{w})$$

Example



Choosing the Right Learning Rate


- Consider a simple 1D example:

$$w^{(\tau+1)} = w^{(\tau)} - \eta \frac{\partial E(w)}{\partial w}$$


- What is the optimal learning rate η_{opt} ?
- If E is quadratic, the optimal learning rate is given by the inverse of the Hessian:

$$\eta_{\text{opt}} = \left(\frac{\partial^2 E(w^{(\tau)})}{\partial w^2} \right)^{-1}$$

- For neural networks, the Hessian is usually infeasible to compute.




η too small




Convergence is slow

η_{opt}



Converges ideally in
a single step

η too large



Might not converge

Discussion: Stochastic Gradient Descent

Advantages

- Very simple, but still quite robust method.
- Minibatches offer a compromise between stability and faster computation.
- Stochasticity in minibatches is often beneficial for learning

Limitations

- Finding a good setting for the learning rate is very important for fast convergence.
 - Choosing the right learning rate is a challenge and requires experience.
 - A different learning rate may be optimal for different parts of the network
- Following the direction of steepest descent is not always the fastest way to the optimum
 - E.g., in highly correlated data

$$\frac{\partial E(\mathbf{w})}{\partial \mathbf{w}}$$


References and Further Reading

- More information about [Neural Networks and Deep Learning](#) is available in the following book.

I. Goodfellow, Y. Bengio, A. Courville
Deep Learning
MIT Press, 2016

<https://www.deeplearningbook.org/>

