

# Elements of Machine Learning & Data Science

Winter semester 2025/26

## Lecture 18 – Data Quality and Preprocessing

12.01.2026

Prof. Bastian Leibe

slides by Prof. Wil van der Aalst



# Elements of Machine Learning & Data Science

## Winter semester 2025/26

### **Part 3: Empirical Analysis and Performance Optimization**

Content by Prof. Holger Hoos  
Chair for AI Methodology (AIM)

Content by Prof. Wil van der Aalst  
Chair of Process and Data Science (PADS)

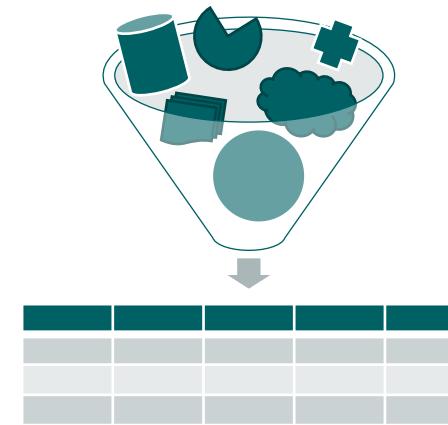
# Empirical Analysis and Performance Evaluation Topics

## 15. Data Quality and Preprocessing

16. Responsible Data Science

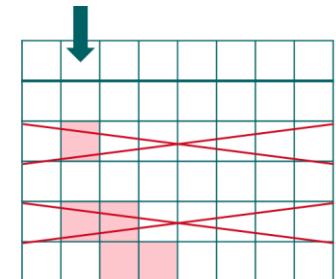
17. Evaluation

18. Performance Optimization

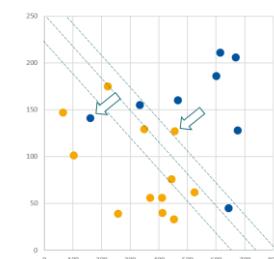


*Data Extraction*

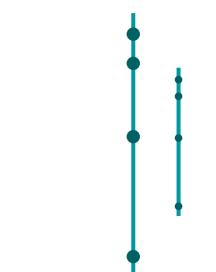
*Preprocessing*



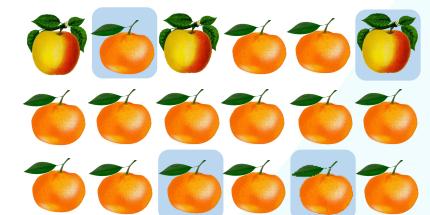
*Missing Values*



*Outliers*



*Normalization*



*Subset Selection*

# Let's Take A Step Back: How to Get the Data?



# 80/20

It is not uncommon that 80% of the effort/time in a data science project is devoted to finding, extracting, cleaning, and transforming the data. Only 20% is concerned with analysis.

# The Two Biggest Hurdles in Practice: Getting the Data and Implementing Changes

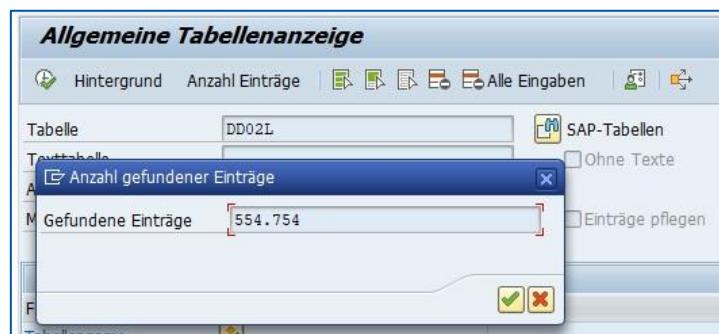


Generated using DALL E3

# Real-Life Examples

CIO of a US bank: “We reduced the number of applications from 12.000 to 8.000” : -)

An SAP installation has hundreds of thousands of tables.



DD02L is the SAP table for SAP Tables.

Tables may have hundreds of columns (e.g. EKPO has > 300 fields).

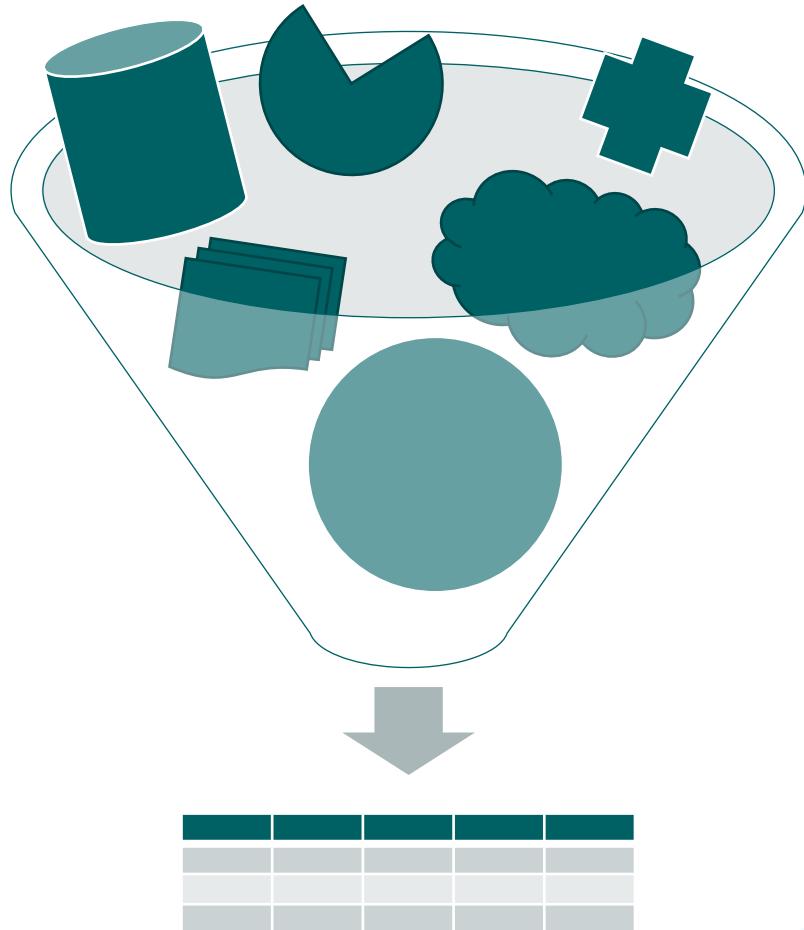
EKPO – Purchasing Document Item  
#1 MANDT – Client  
#2 EBELN – Purchasing Document Number  
#3 EBELP – Item Number of Purchasing Document  
#4 LOEKZ – Deletion indicator in purchasing document  
#5 STATU – RFQ status  
...  
#299 POL\_ID – Order List Item Number  
#300 CONS\_ORDER – Purchase Order for Consignment

Organizations such as Siemens have 70 SAP installations.



# Data Quality & Preprocessing

1. Introduction
2. Missing Values
3. Outliers
4. Transformation & Normalization
5. Reduction



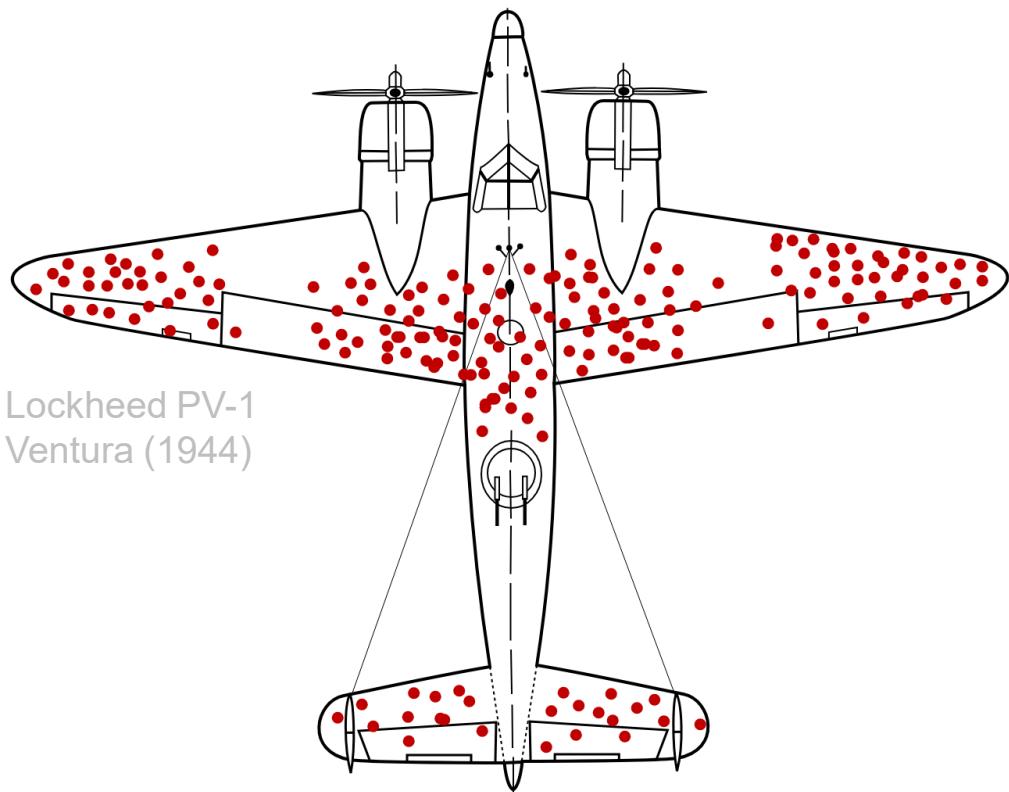
# Data Science Pipeline

- Garbage in, garbage out
- Possible **problems** (big data, security), **errors** (data quality), **biases** (e.g., survivorship bias) everywhere
- Problems, errors and biases **propagate**

**Goal:** increase data quality and modify the data to suit the analysis question and applied techniques



## Example: Survivorship Bias



[https://en.wikipedia.org/wiki/Survivorship\\_bias](https://en.wikipedia.org/wiki/Survivorship_bias)



A Canadian study in 2011 revealed that 97.4 percent of Porsches from the last 25 years are still on the road.



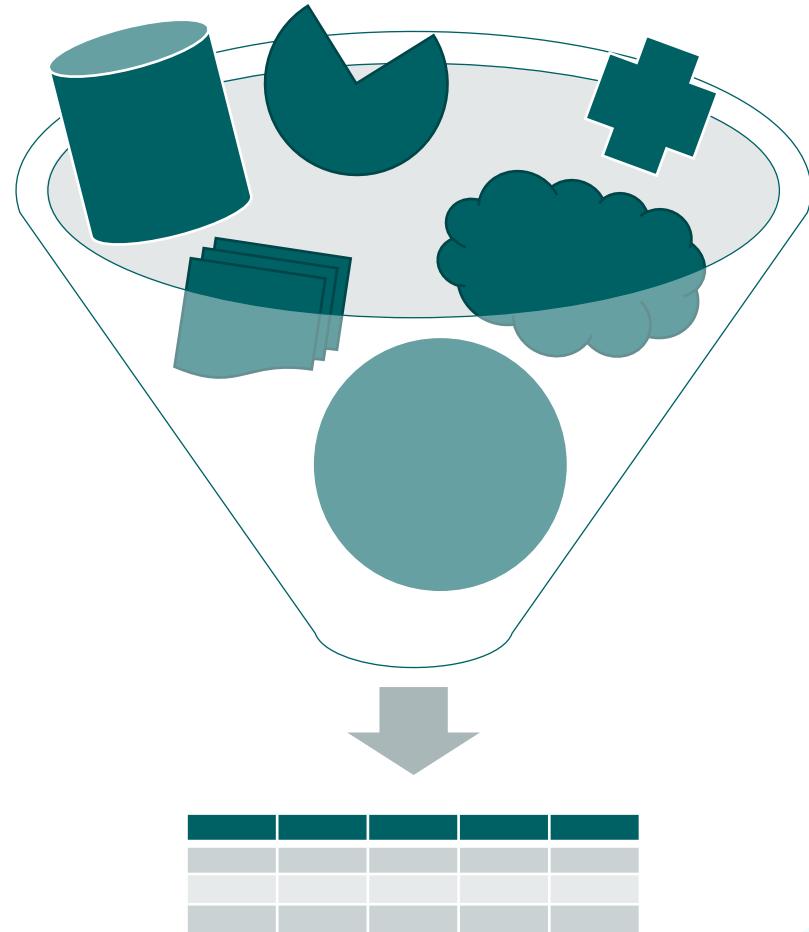
# Data Quality Aspects

- **Accuracy** (wrong value)
- **Completeness** (missing value)
- **Consistency** (different conventions/formats)
- **Timeliness** (outdated values)

Name	Age	Siblings	Date of Admission
Sara Johnson	55	0	30.09.2022
NAME	17		23-11-22
Smith	28	2	8/24/22
Emma Miller	2	56	May 10 <sup>th</sup> , 22
Jones	187	3	220701
...	...	...	

# Data Quality & Preprocessing

1. Introduction
2. **Missing Values**
3. Outliers
4. Transformation & Normalization
5. Reduction



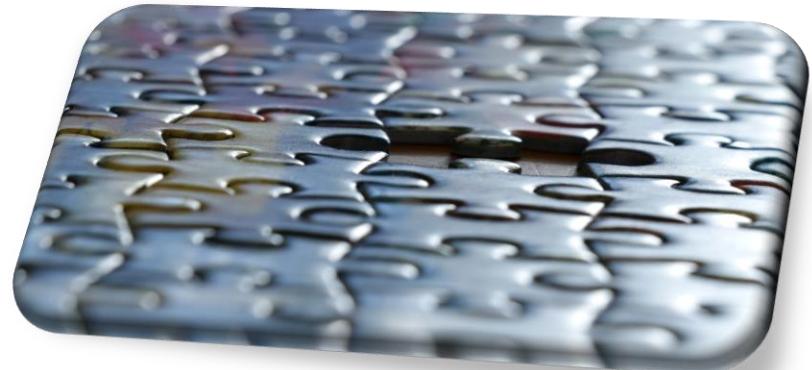
# Detecting Missing Values

**Missing values may be obvious...**

- Empty value
- NaN / NA

**... or may be disguised!**

- Default value
- Invalid value



# Handling Missing Values

- 1) Fill in manually
- 2) Ignore
- 3) Fill in a derived value



# Handling Missing Values: Ignore



## Discard the feature

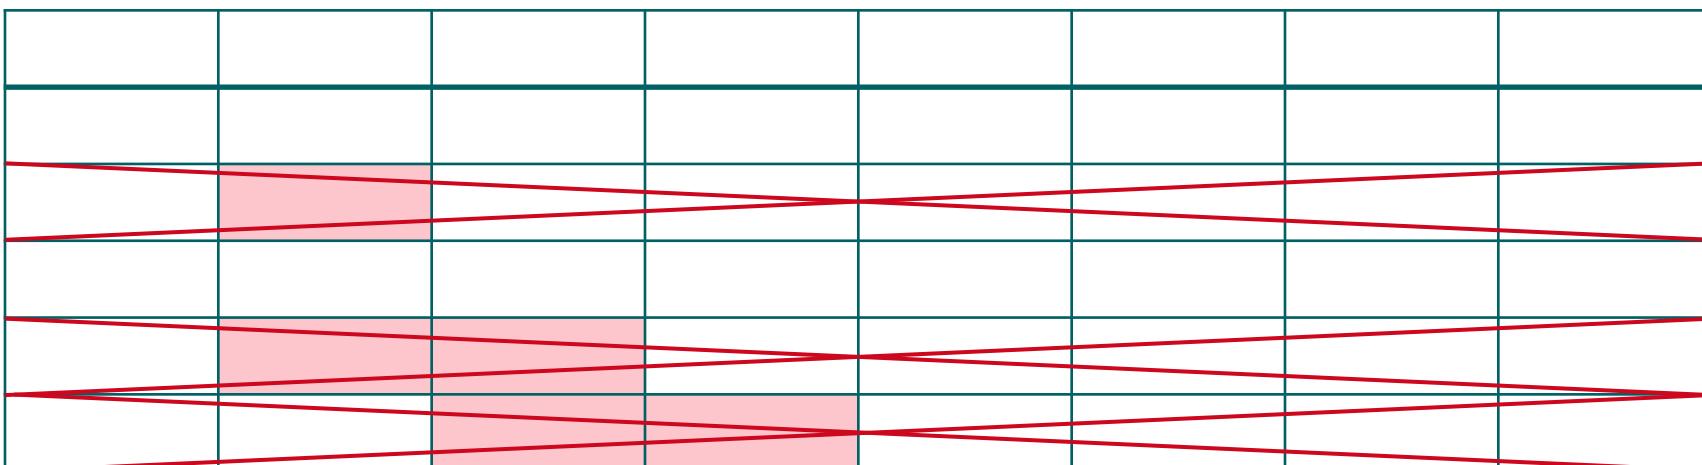
- The whole feature is removed from the data
- Usually done if the number of missing values is too large to allow meaningful analysis  
(as a rule of thumb, if more than 60 % of the feature values are missing)


# Handling Missing Values: Ignore



## Discard the instance

- The entire instance is simply **discarded**
- Usually done when the whole instance becomes unusable (e.g., labeling attribute for classification is missing)
- If the data set misses a lot of values, this technique may make the whole data set **unusable** or introduce a **bias**

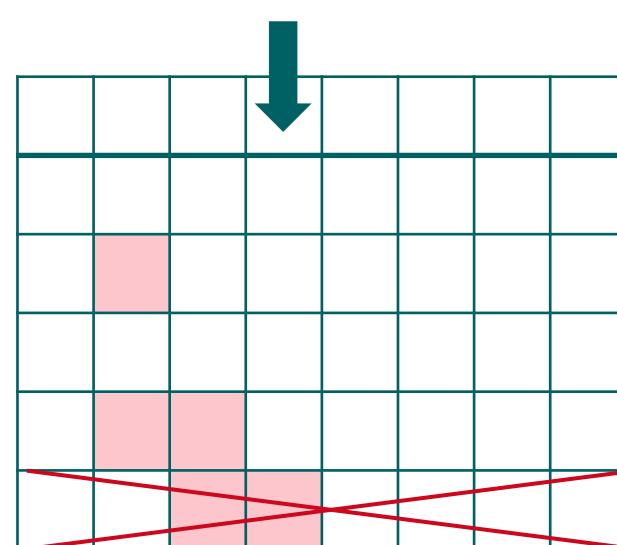
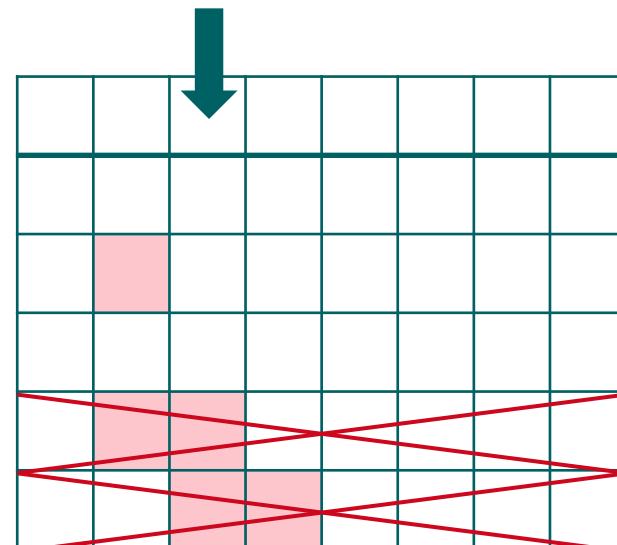
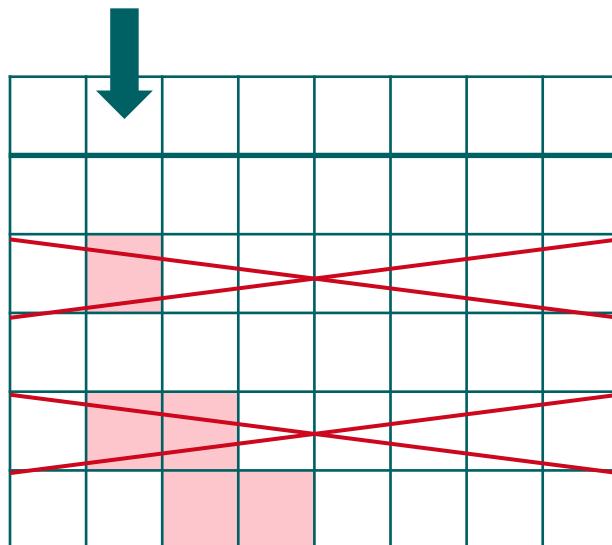


# Handling Missing Values: Ignore



**Ignore the instance only for features where the value is missing**

- The instance is ignored when analyzing features where it misses a value
- Information for other features remains usable



# Handling Missing Values: Create



## Mean/median/mode of the whole feature

- Compute mean/median/mode and fill the gaps accordingly
- Example: compute yearly income

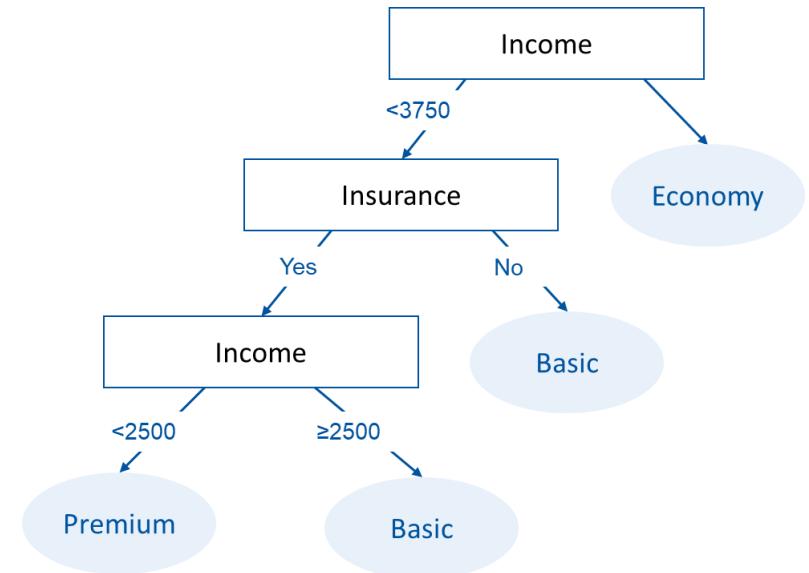
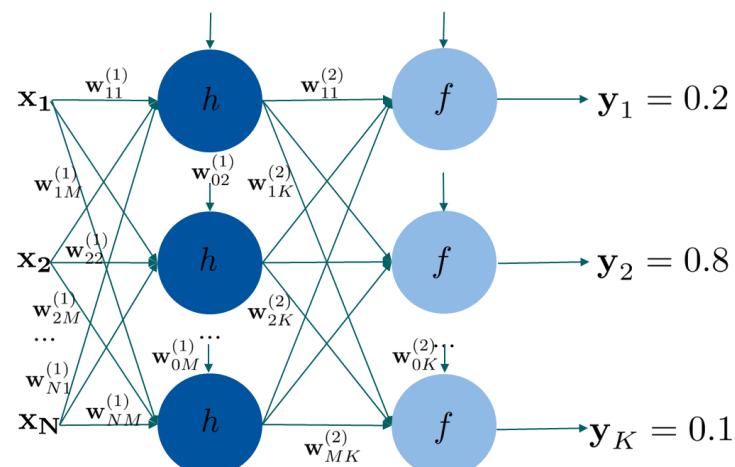
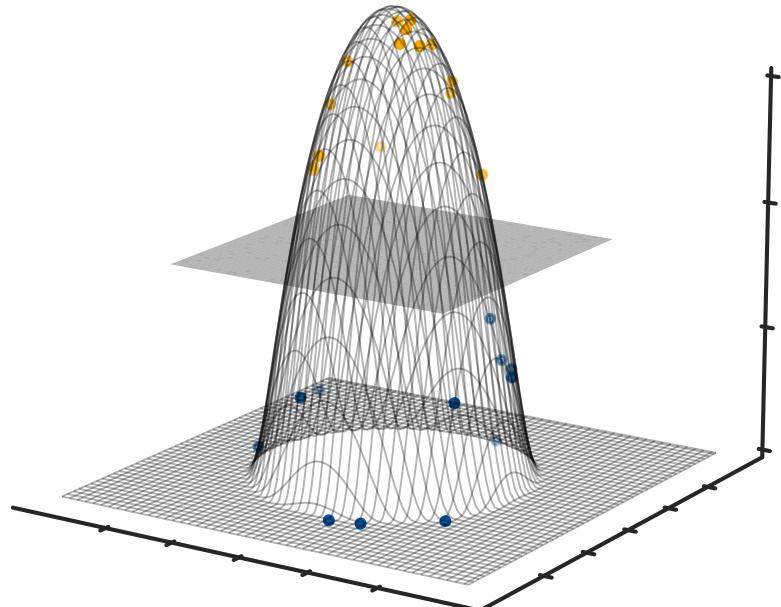
## Mean/median/mode of all instances belonging to the same class

- Compute mean/median/mode only based on instances with the same class label
- Higher chances to be accurate compared to the overall mean/median/mode
- Example: compute income for a 20-year-old Student living in Aachen, Germany

# Handling Missing Values: Create

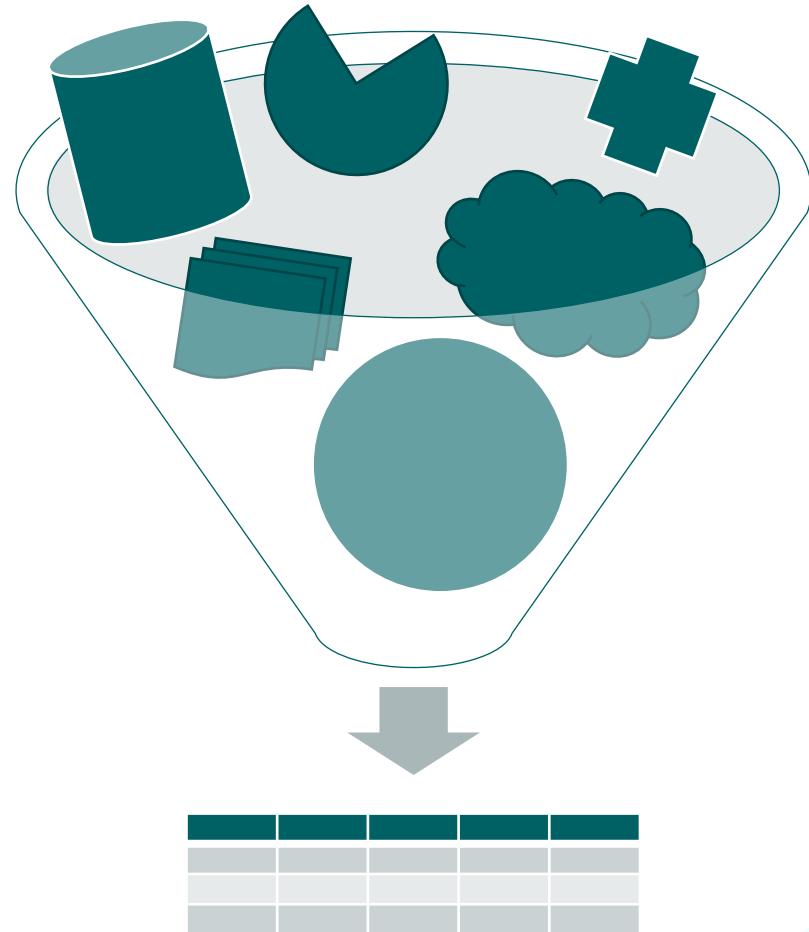
## Complex derived value (use a predictor model)

- Fill in the value given by a suitable prediction model
- E.g., decision trees, regression, NNs, SVMs...

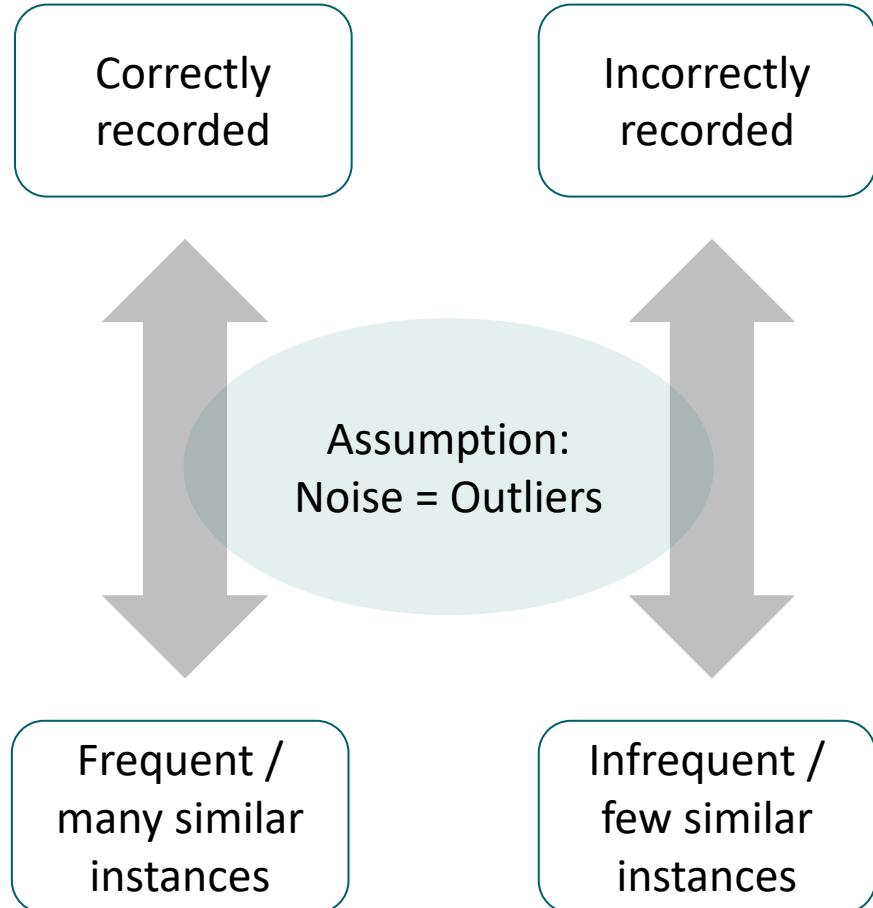


# Data Quality & Preprocessing

1. Introduction
2. Missing Values
- 3. Outliers**
4. Transformation & Normalization
5. Reduction



# Introduction

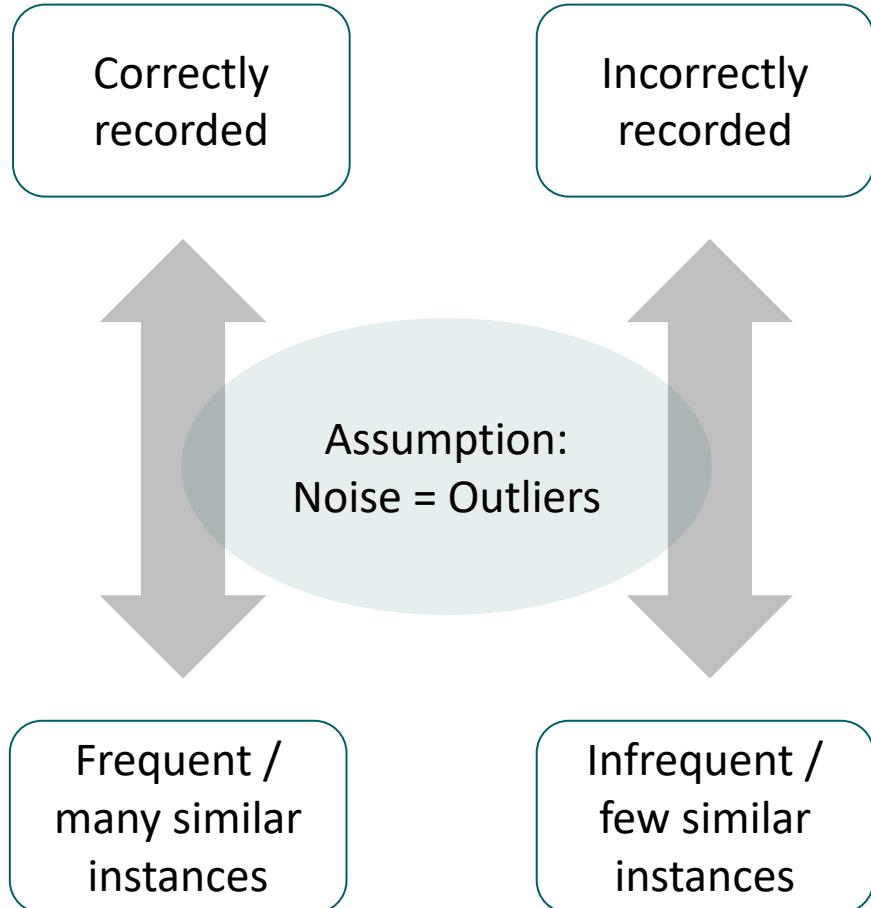


## What is noise?

- We assume that noise causes outliers
- Thus, **outliers** indicate noise

e.g., 25.5 centimeters of snow in Rome

# Outlier Detection

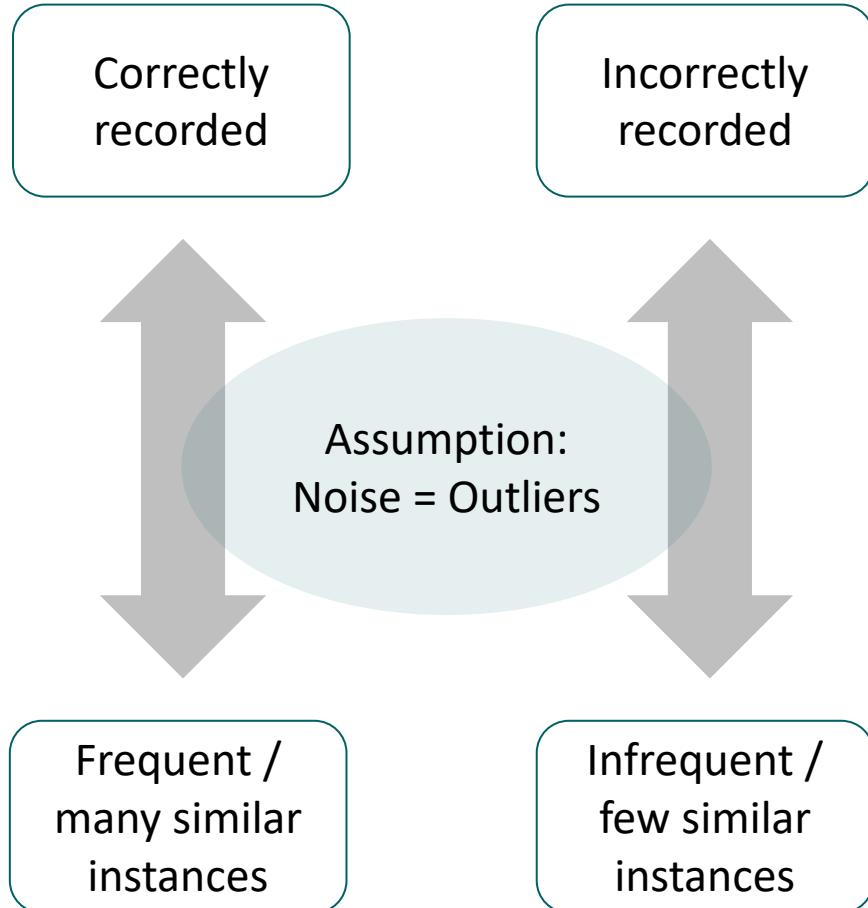


## How to detect outliers?

- Boxplots
- Decision trees
- Regression
- SVMs
- Clustering
- ...

→ Predictor models can be used to **define** outliers

# Outlier Handling



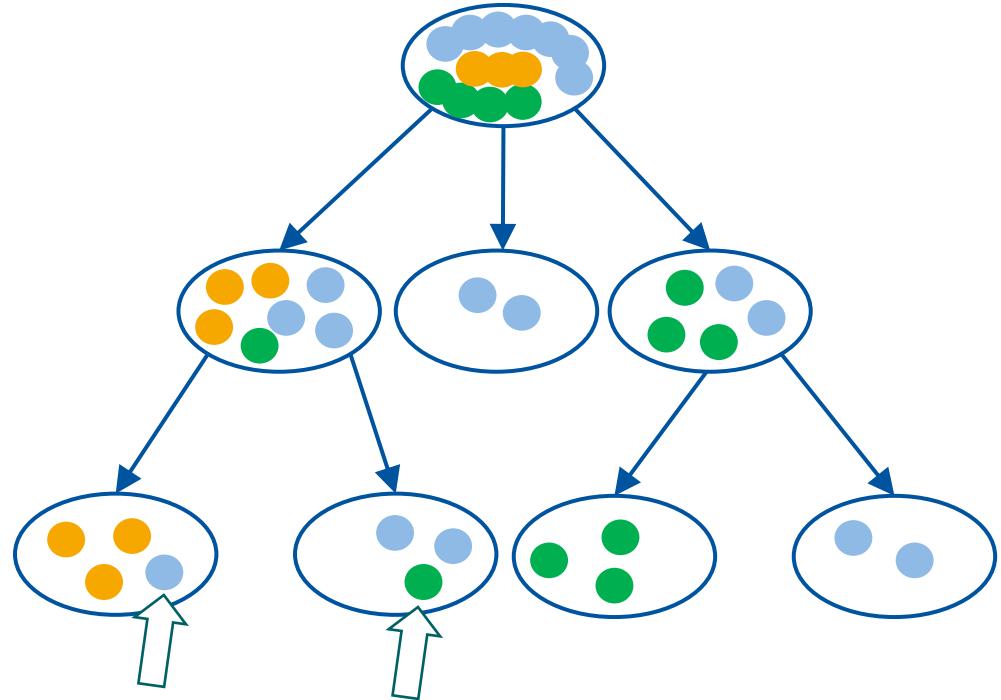
## How to handle outliers?

Outliers can be **handled as missing values**:

- Fill in a correct value manually
- Ignore the feature/instance
- Replace with a derived value

→ Predictor models can be used to **replace** outliers

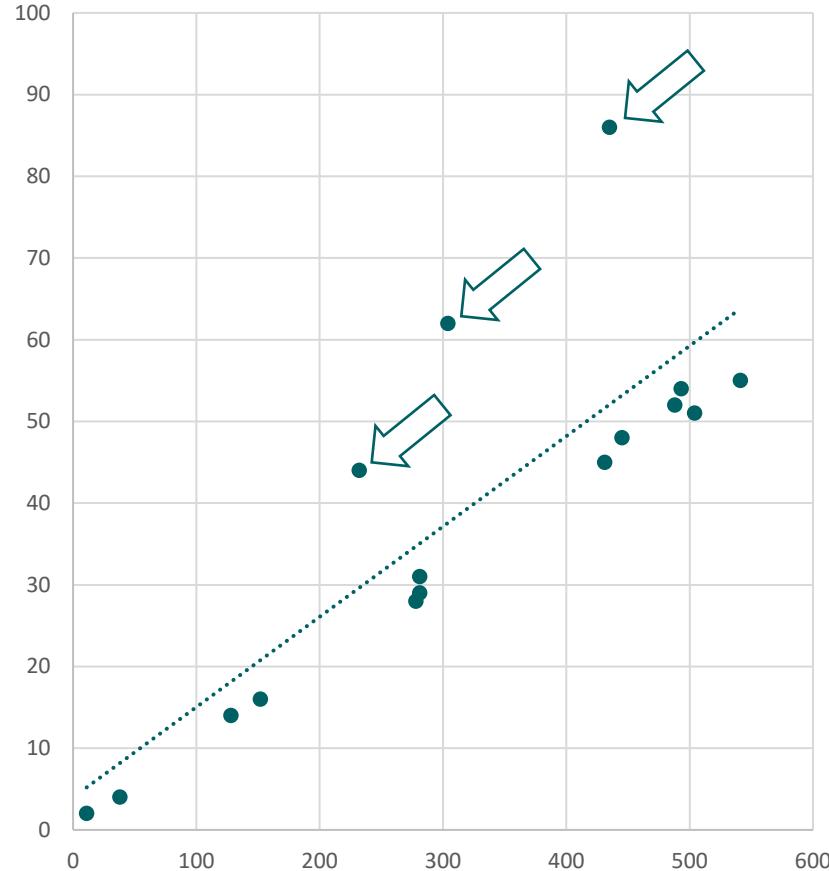
# Outlier Detection - Decision Trees



## How to detect outliers?

- Every leaf node is assigned a class label
- Instances in that leaf node with a **non-matching class label** can be considered outliers

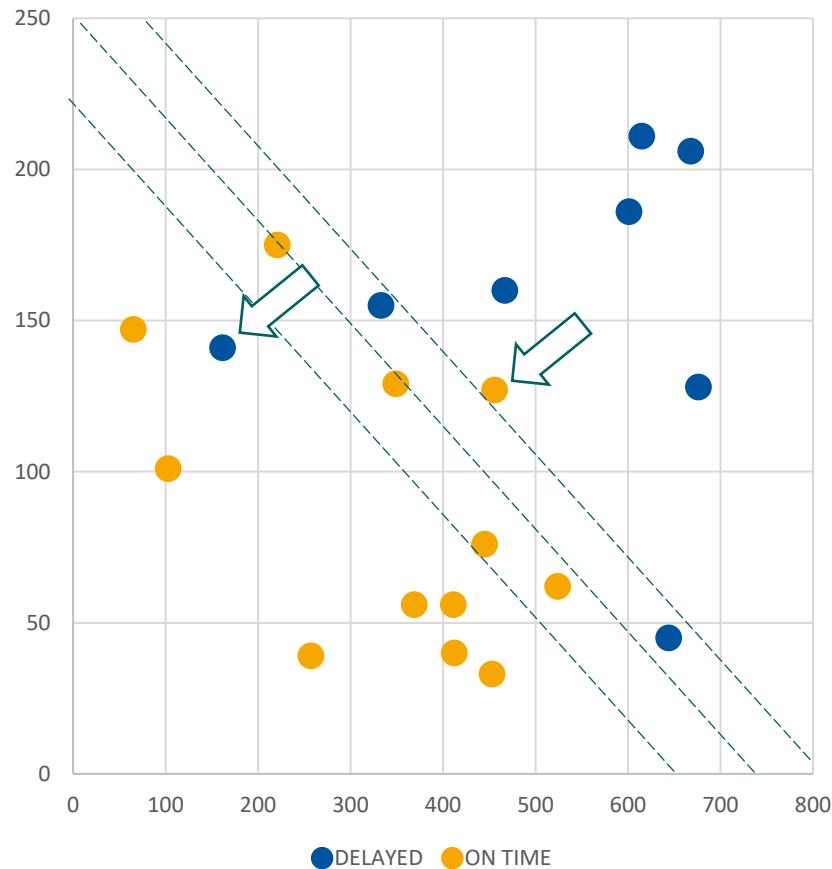
# Outlier Detection - Regression



## How to detect outliers?

- Instances which are **far away** from the predicted value are considered outliers
- The definition of 'far away' depends on an **error function** and **threshold**

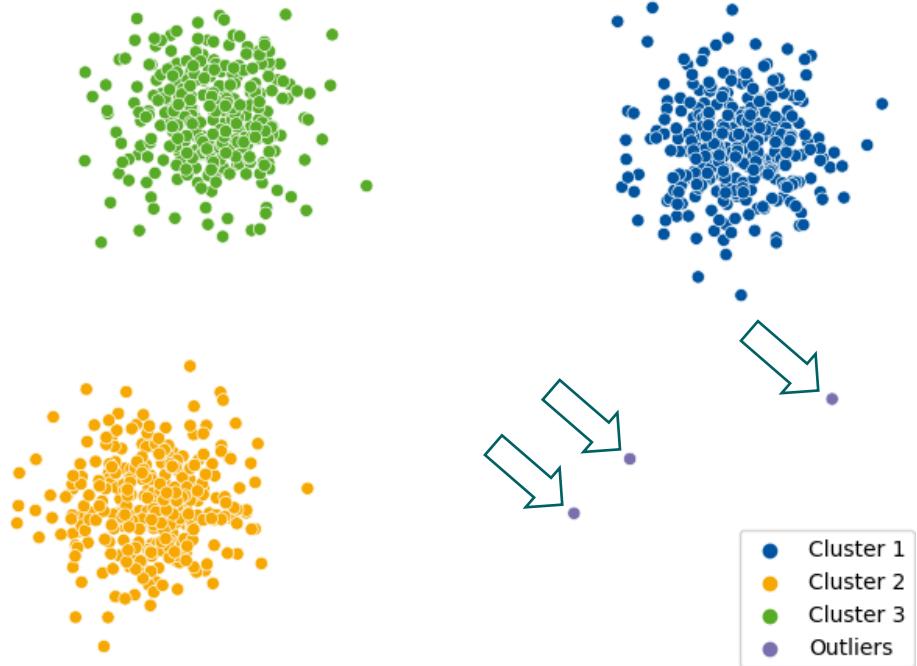
# Outlier Detection – SVM



## How to detect outliers?

- Instances which are (too far) on the [wrong side](#) of the hyperplane are considered outliers
- Soft margin may be used to define how far

# Outlier Detection - Clustering

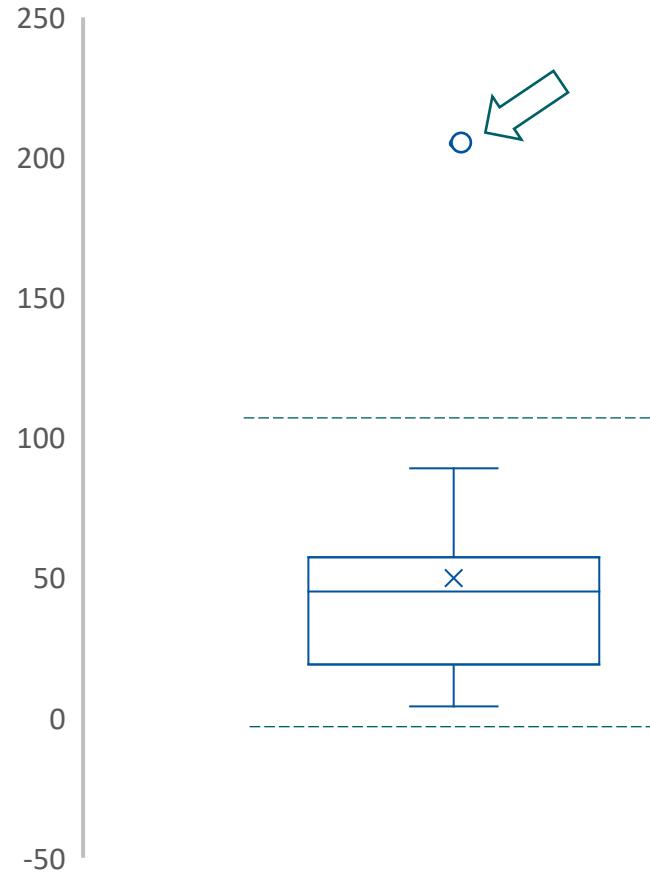


## How to detect outliers?

- Instances **outside of any cluster** can be considered outliers

(Of course, a prerequisite for this is that the clustering algorithm itself can handle outliers...)

# Outlier Detection - Boxplots



## How to detect outliers?

- Instances **above the upper fence**
- Instances **below the lower fence**

→ Outlier handling option:

Clamp values to the nearest fence or take median value

# Outlier Handling



[1]

## How to handle outliers?

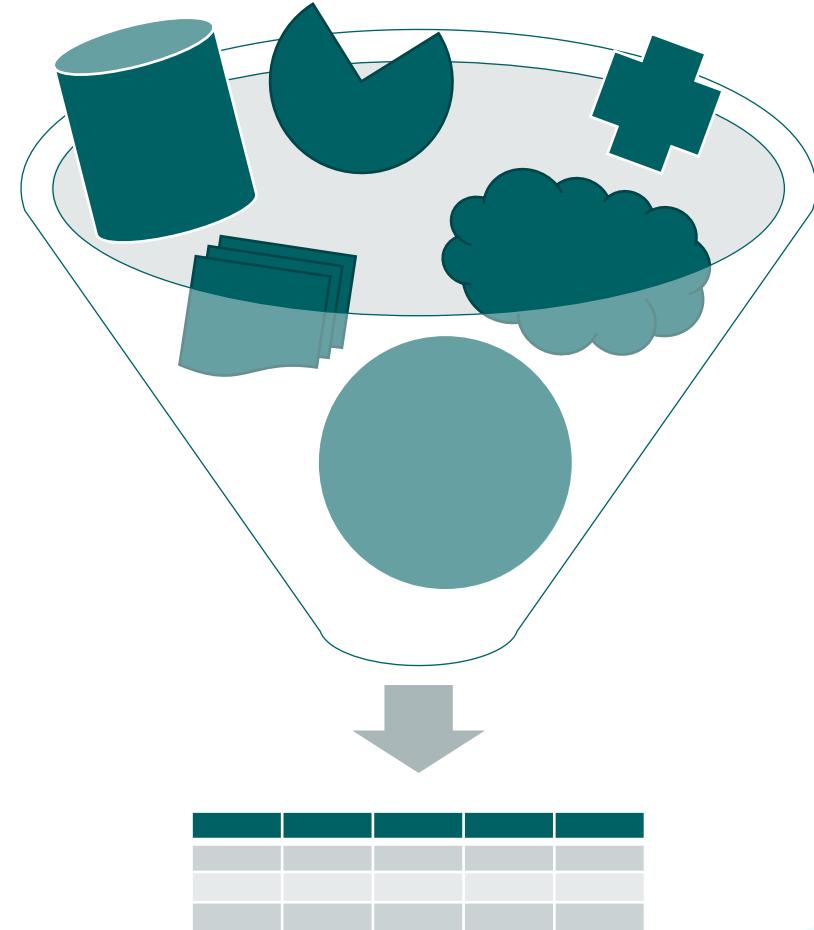
Outliers can be [handled as missing values](#):

- 1) Fill in a correct value manually
- 2) Ignore the feature/instance
- 3) Replace with a derived value

Again: the appropriate method depends on the data and purpose

# Data Quality & Preprocessing

1. Introduction
2. Missing Values
3. Outliers
- 4. Transformation & Normalization**
5. Reduction



# Preprocessing – Preparing the Data for Analysis

- Transformation: change the data to the right data type
- Normalization: adjust the influence of features
- Reduction: make the data smaller for analysis



[1]

## Preprocessing – Transformation

- One-hot encoding: categorical to numerical
- Binning: numerical to categorical

The diagram illustrates the transformation of a dataset from its original form to a one-hot encoded representation. A blue arrow points from the left table to the right table, indicating the mapping.

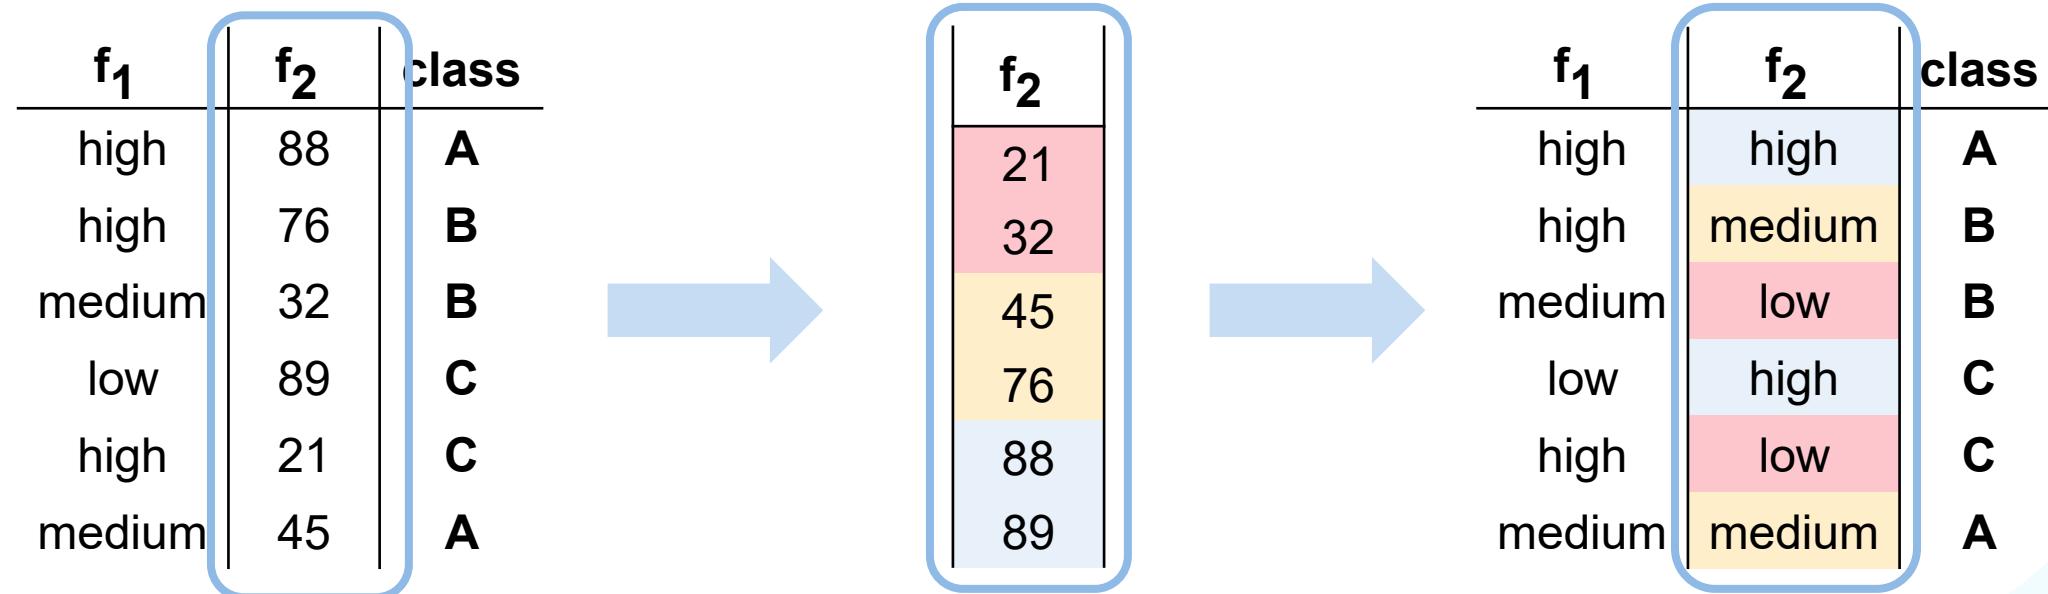
$f_1$	$f_2$	class
high	88	A
high	76	B
medium	32	B
low	89	C
high	21	C
medium	45	A

→

$f_1$ - high	$f_1$ - medium	$f_1$ - low	$f_2$	class
1	0	0	88	A
1	0	0	76	B
0	1	0	32	B
0	0	1	89	C
1	0	0	21	C
0	1	0	45	A

# Preprocessing – Transformation

- One-hot encoding: categorical to numerical
- Binning: numerical to categorical



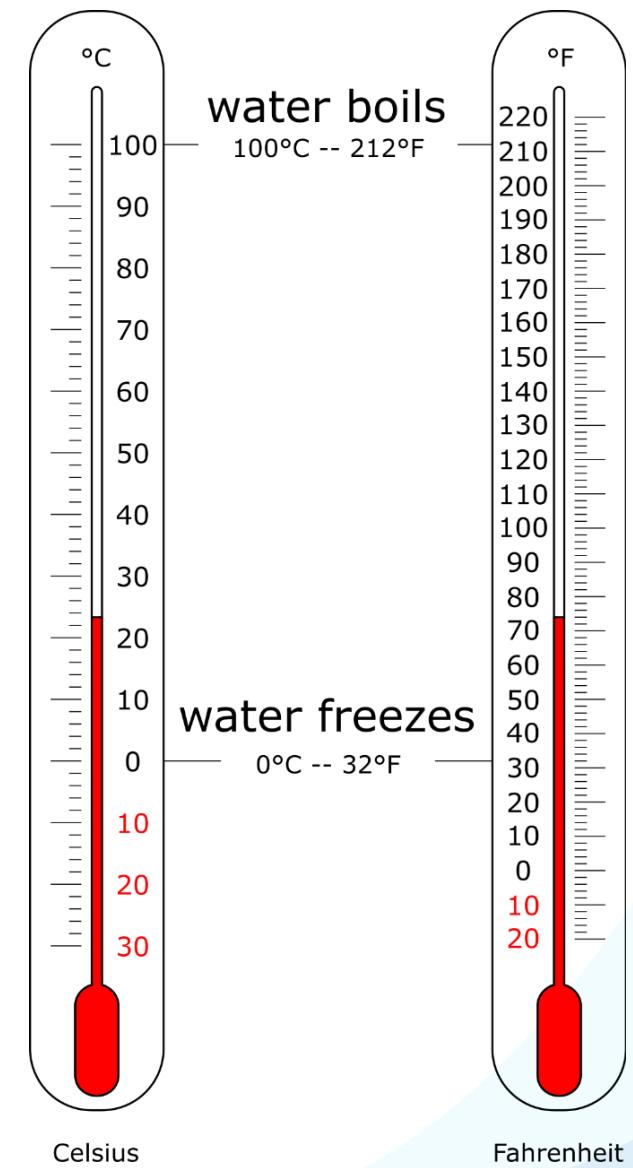
# Preprocessing – Normalization

## Adjusting the influence of features

- Feature weight and range often depends on the chosen unit (km, mm, miles, ...)
  - Algorithms tend to give more weight to features with a large range
- May introduce an unwanted bias
- May hinder interpretability
- Scales may be non-linear (e.g. logarithmic)

Sum of squared errors:

$$\frac{1}{2} \sum_{i=1}^N (t_i - \mathbb{M}(\mathbf{x}_i))^2$$



# Preprocessing – Normalization

## Min-max normalization

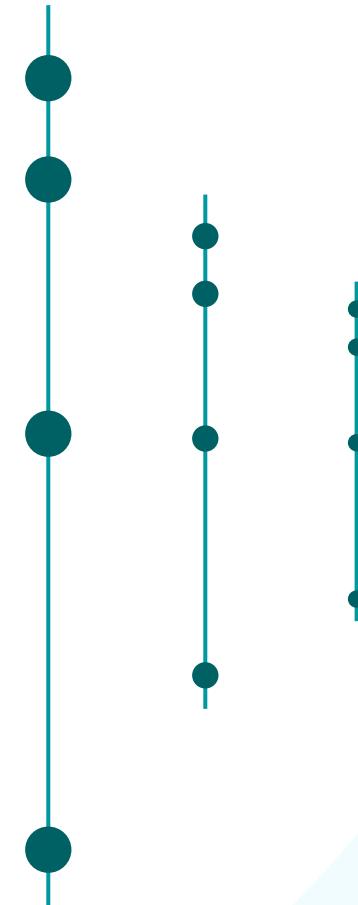
- Maps the values onto a **predefined range** [low, high]
- Preserves **relative differences**, i.e., relations between the data values

We normalize feature  $d$  by replacing its value for each instance  $i$  as follows:

$$\text{norm}(\mathbf{x}_i[d]) = \frac{\mathbf{x}_i[d] - d_{\min}}{d_{\max} - d_{\min}} \cdot (\text{high} - \text{low}) + \text{low}$$

maximal value  
of feature  $d$

minimal value  
of feature  $d$



# Preprocessing – Normalization

## Min-max normalization

d
11
82
33
12
76

$$\begin{aligned}d_{min} &= 11 \\d_{max} &= 82\end{aligned}$$



Consider  
high = 100  
low = 5

$$\text{norm}(\mathbf{x_i}[d]) = \frac{\mathbf{x_i}[d] - d_{\min}}{d_{\max} - d_{\min}} \cdot (\text{high} - \text{low}) + \text{low}$$

# Preprocessing – Normalization

## Min-max normalization

<b>d</b>	<b>norm(d)</b>	<b>norm(d)</b>
11	$(11 - 11)/(82 - 11) \cdot (100 - 5) + 5$	5
82	$(82 - 11)/(82 - 11) \cdot (100 - 5) + 5$	100
33	$(33 - 11)/(82 - 11) \cdot (100 - 5) + 5$	34.44
12	$(12 - 11)/(82 - 11) \cdot (100 - 5) + 5$	6.34
76	$(76 - 11)/(82 - 11) \cdot (100 - 5) + 5$	91.97

≈

$$\text{norm}(\mathbf{x_i}[d]) = \frac{\mathbf{x_i}[d] - d_{\min}}{d_{\max} - d_{\min}} \cdot (\text{high} - \text{low}) + \text{low}$$

# Preprocessing – Normalization

## Standard score (Z-score) normalization

- Uses the standard deviation to quantify the significance of the difference between a value and the overall mean
- Range is  $[-\infty, \infty]$ , but 0 has a clear meaning
- Useful when actual minimum and maximum of the attribute are unknown
- Useful when outliers may impact min-max normalization

For each  $i$ :

$$\text{norm}(\mathbf{x}_i[d]) = \frac{\mathbf{x}_i[d] - \bar{d}}{\text{sd}(d)}$$

$\bar{d}$  is the mean of all values of feature  $d$ :

$$\frac{1}{N} \sum_{i=1}^N \mathbf{x}_i[d]$$

$\text{sd}(d)$  is the standard deviation of feature  $d$ :

$$\sqrt{\left( \frac{\sum_{i=1}^N (\mathbf{x}_i[d] - \bar{d})^2}{N-1} \right)}$$

# Preprocessing – Normalization

## Standard score (Z-score) normalization

d
11
82
33
12
76



$$\bar{d} = 42.8$$
$$\text{sd}(d) = 34.259$$

$$\text{norm}(\mathbf{x_i}[d]) = \frac{\mathbf{x_i}[d] - \bar{d}}{\text{sd}(d)}$$

# Preprocessing – Normalization

## Standard score (Z-score) normalization

d		norm(d)	
11		$(11 - 42.8)/34.259$	
82	$\bar{d} = 42.8$	$(82 - 42.8)/34.259$	
33	$sd(d) = 34.259$	$(33 - 42.8)/34.259$	
12		$(12 - 42.8)/34.259$	
76		$(76 - 42.8)/34.259$	

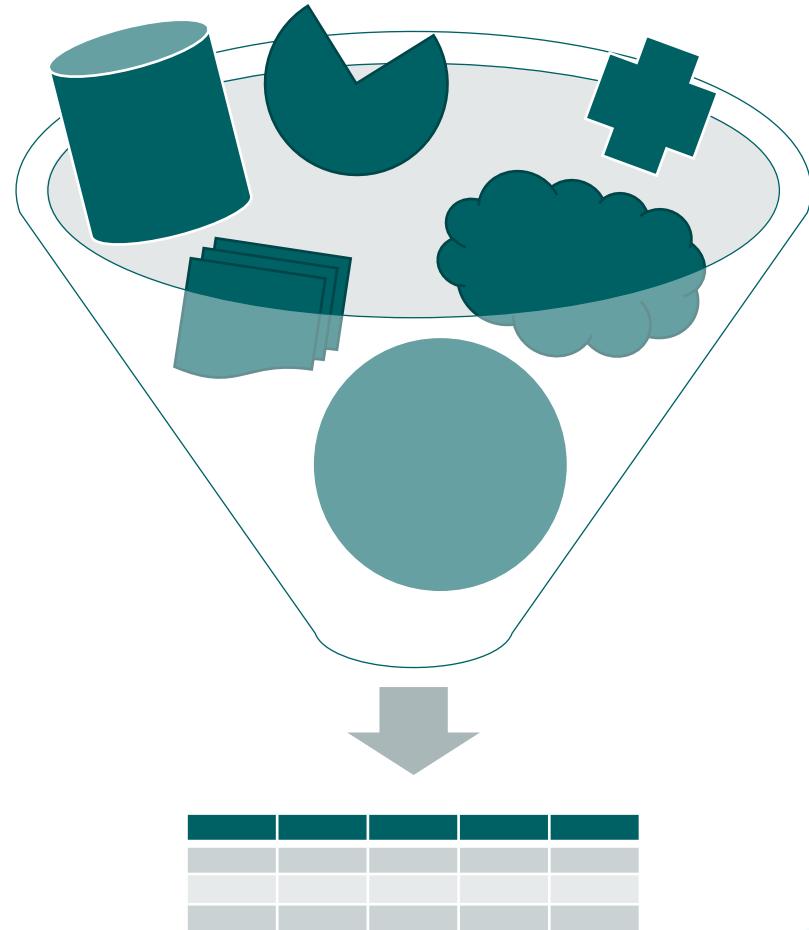
$\approx$

norm(d)
-0.93
1.14
-0.29
-0.90
0.97

$$\text{norm}(\mathbf{x_i}[d]) = \frac{\mathbf{x_i}[d] - \bar{d}}{sd(d)}$$

# Data Quality & Preprocessing

1. Introduction
2. Missing Values
3. Outliers
4. Transformation & Normalization
5. **Reduction**



# Preprocessing – Data Reduction

- Analysis may become unfeasible due to size of data
- Goal: reduce the data size but maintain same (or similar) analysis results
- Feature reduction: remove or replace some features
- Instance reduction: remove, replace or aggregate some instances

Feature reduction

ID	$f_1$	$f_2$	...	$f_D$
1				
2				
...				
N				

Instance reduction



# Preprocessing – Feature Reduction

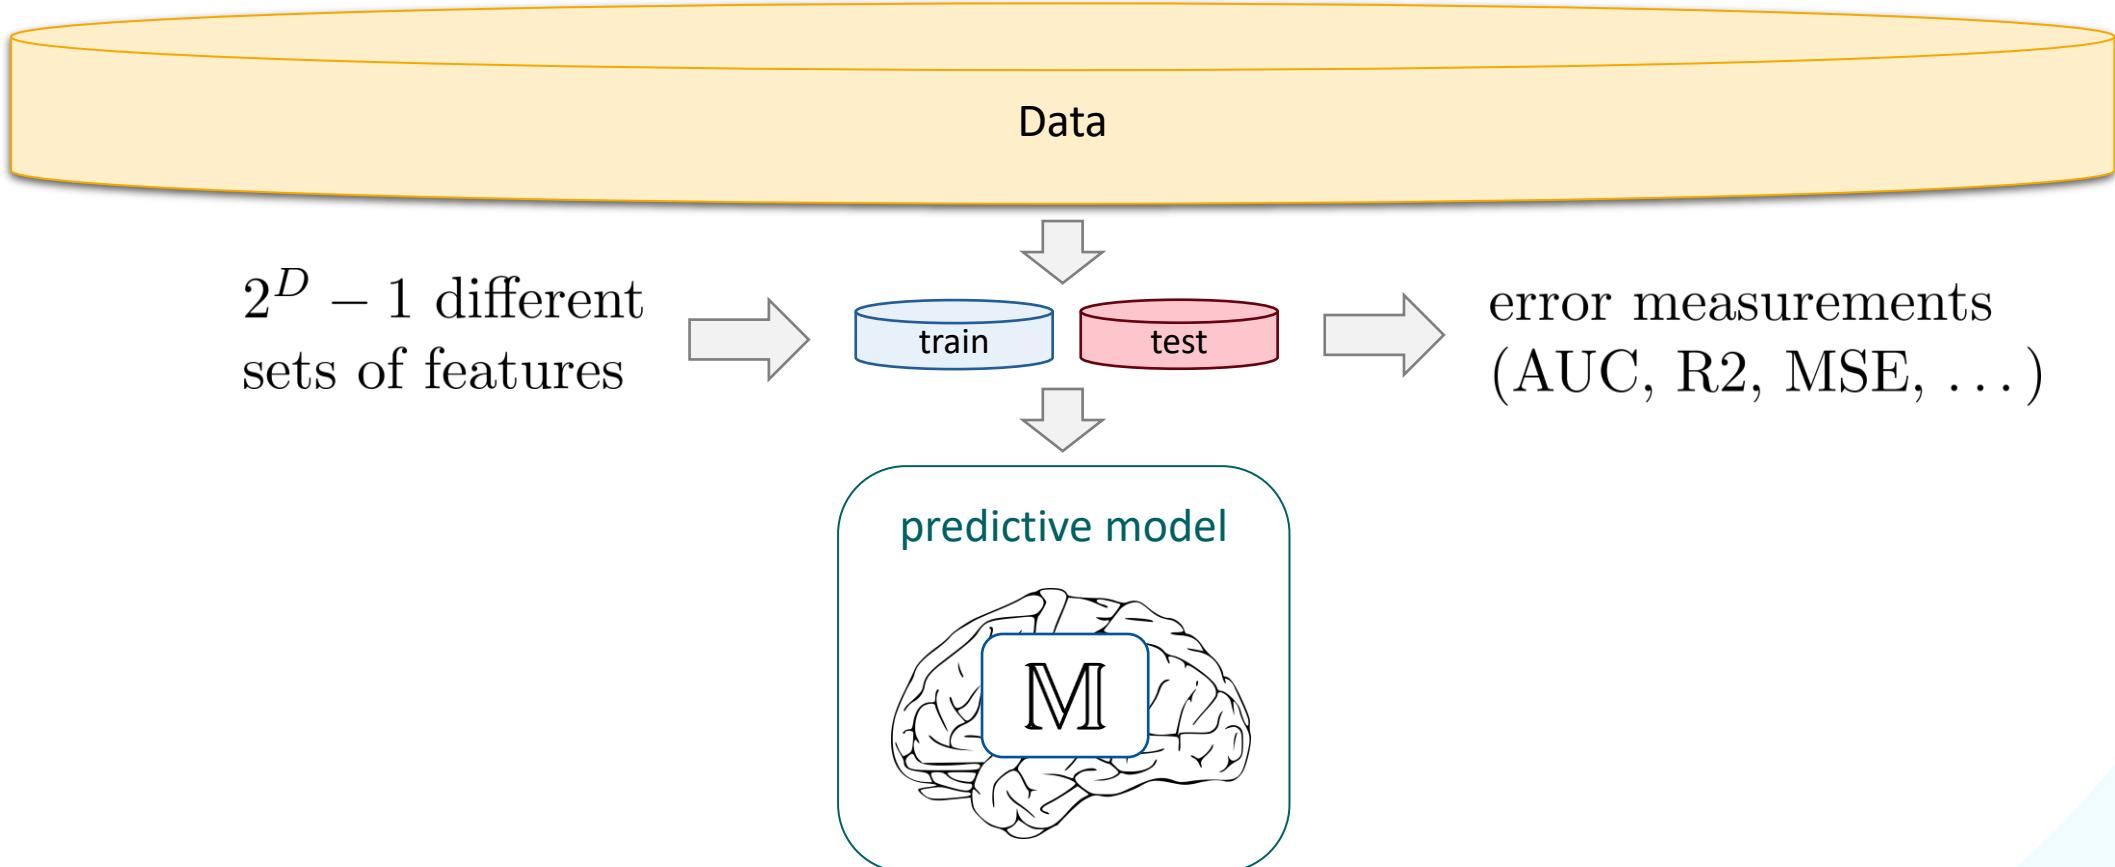
## Projecting data on fewer dimensions

- [Autoencoders](#) (compare text mining): a special type of NN that transforms the input data into a representation with fewer dimensions (encoding). [Learning a good representation is key!](#)
- [Principal Component Analysis \(PCA\)](#): represent the original features by a few orthogonal (uncorrelated) variables that capture most of the variability

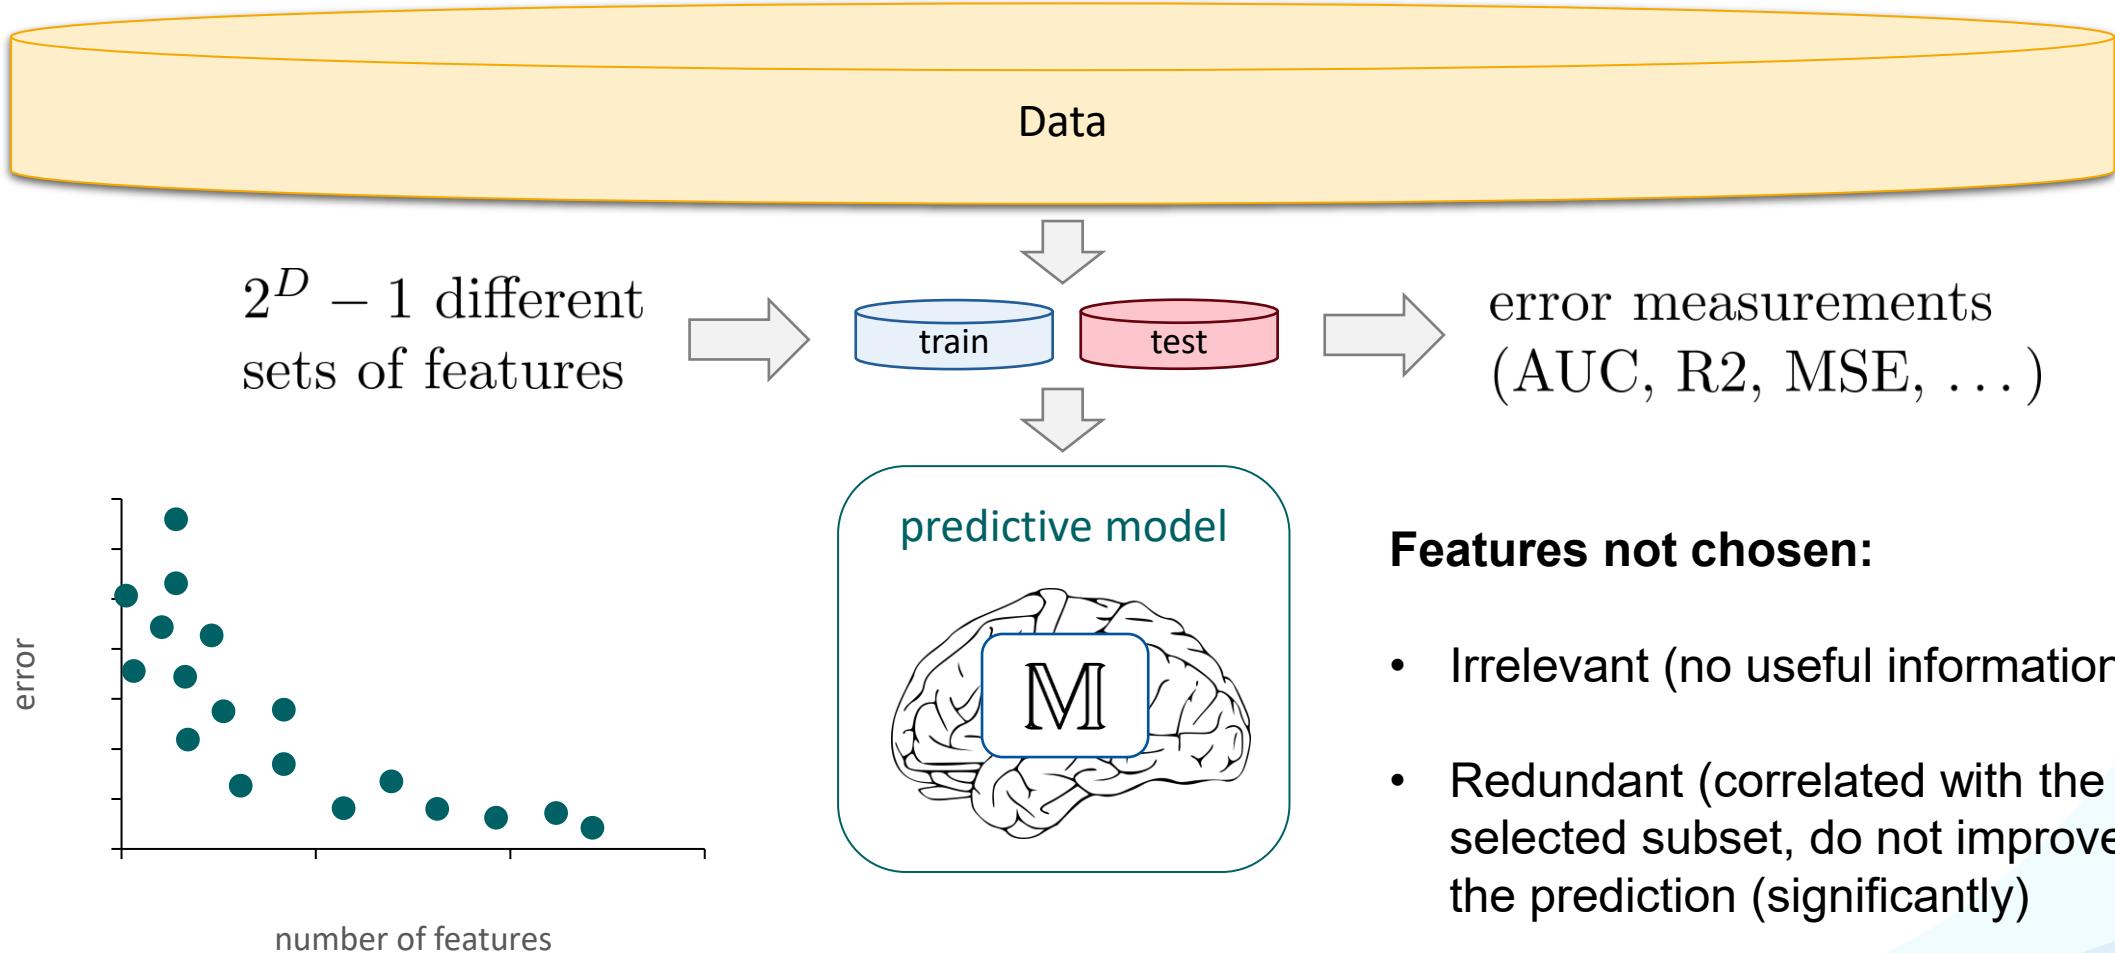
## Feature subset selection: detect and remove [irrelevant/redundant features](#)

- Use domain knowledge (e.g., remove identifiers)
- Exploit dependencies (e.g., delete features that can be estimated from others using regression)
- Model-driven (e.g. delete features that are not used in a constructed decision tree or, more general, features that can be left out without reducing the quality of the model much)

# Preprocessing – Feature Subset Selection

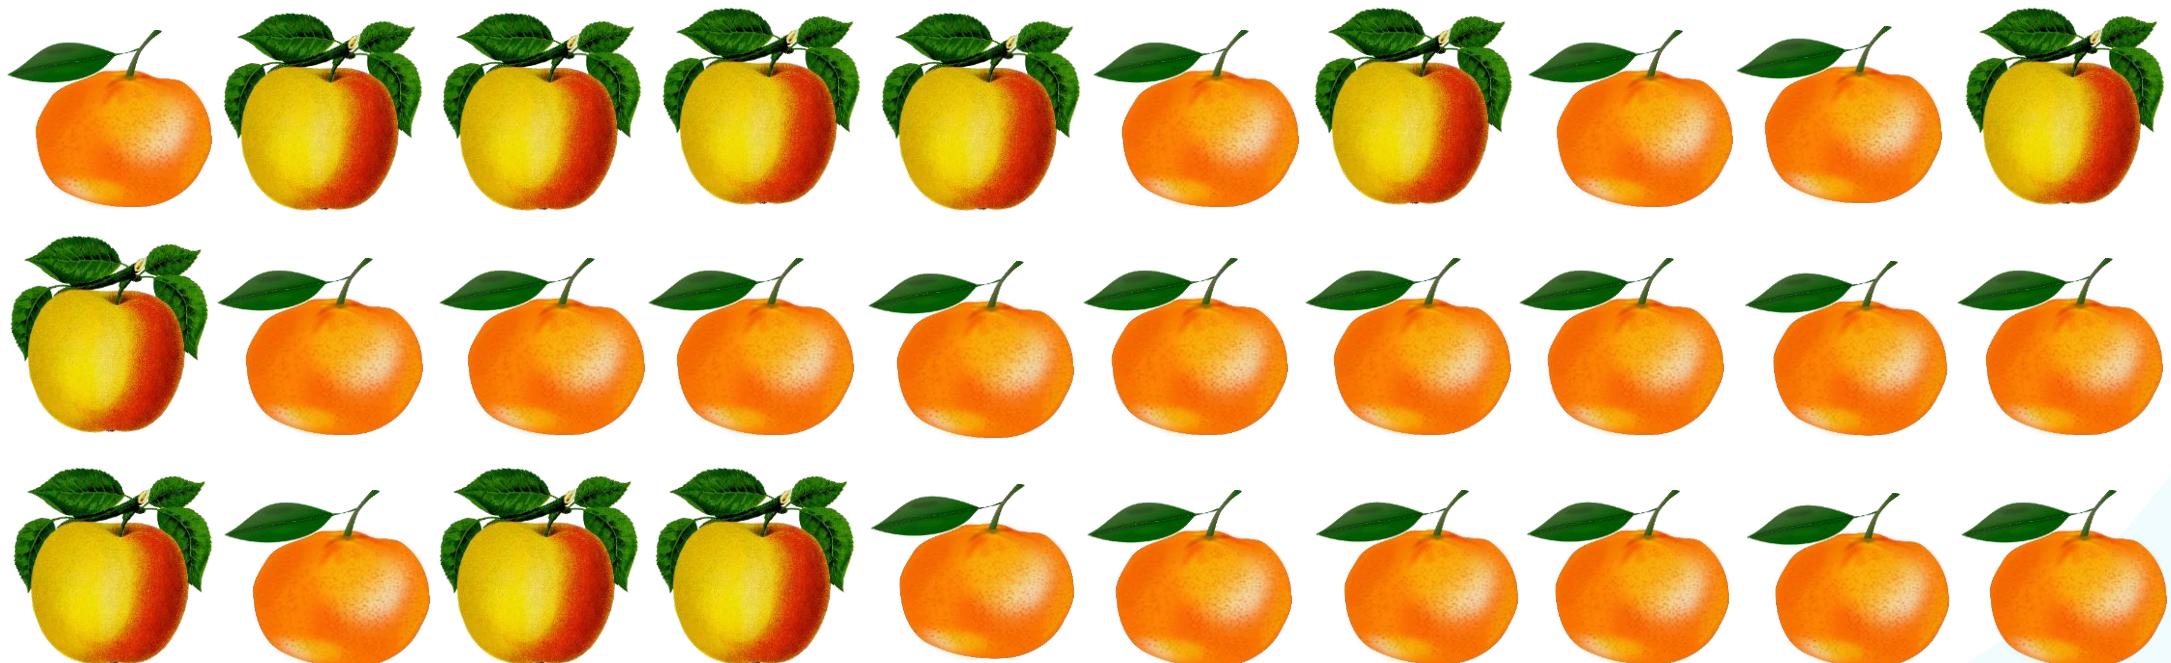


# Preprocessing – Feature Subset Selection

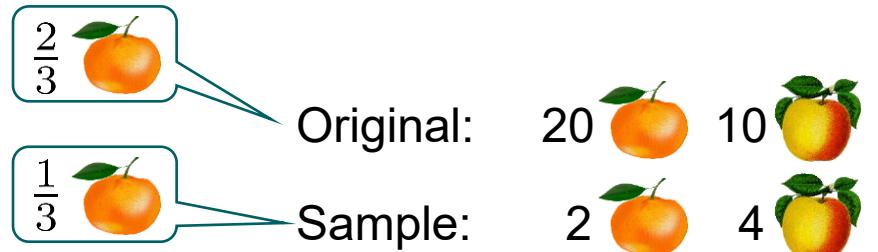


## Preprocessing – Sampling

Goals: make the data smaller, remove or introduce biases

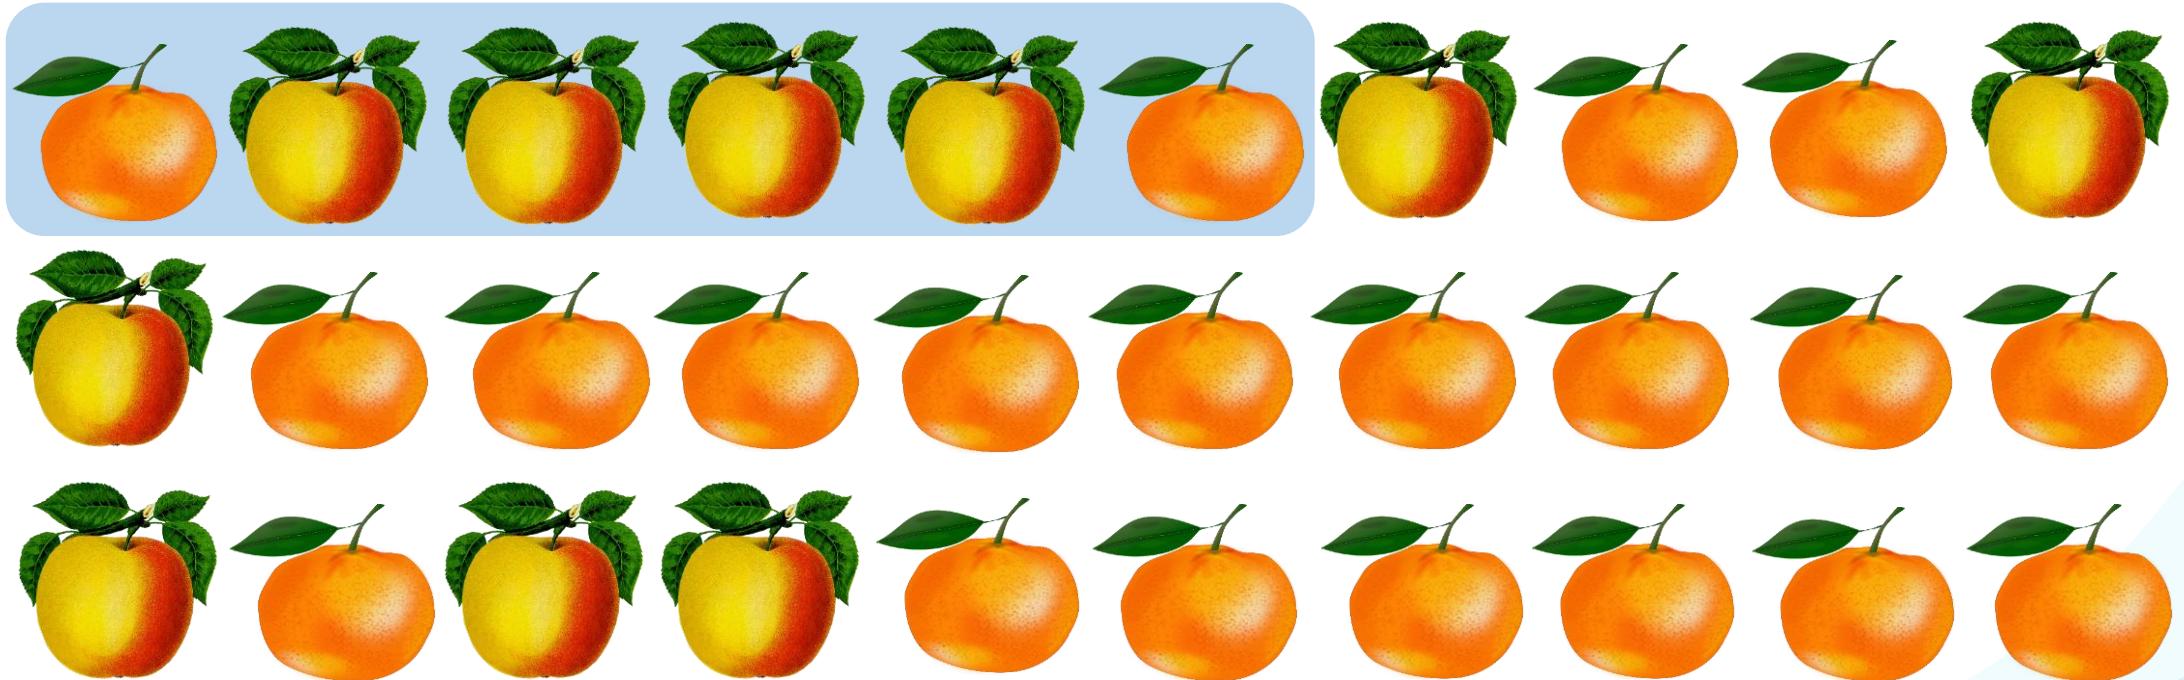


## Preprocessing – Sampling

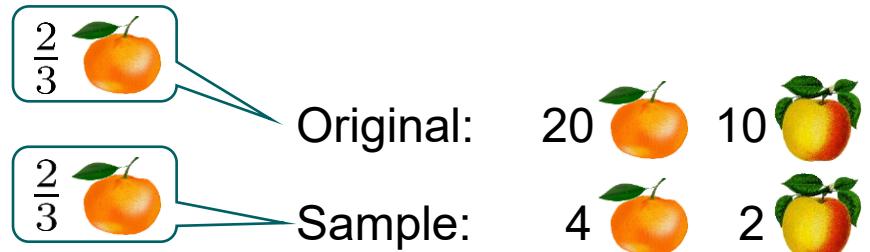


Top sampling:

take the first  $N$  instances

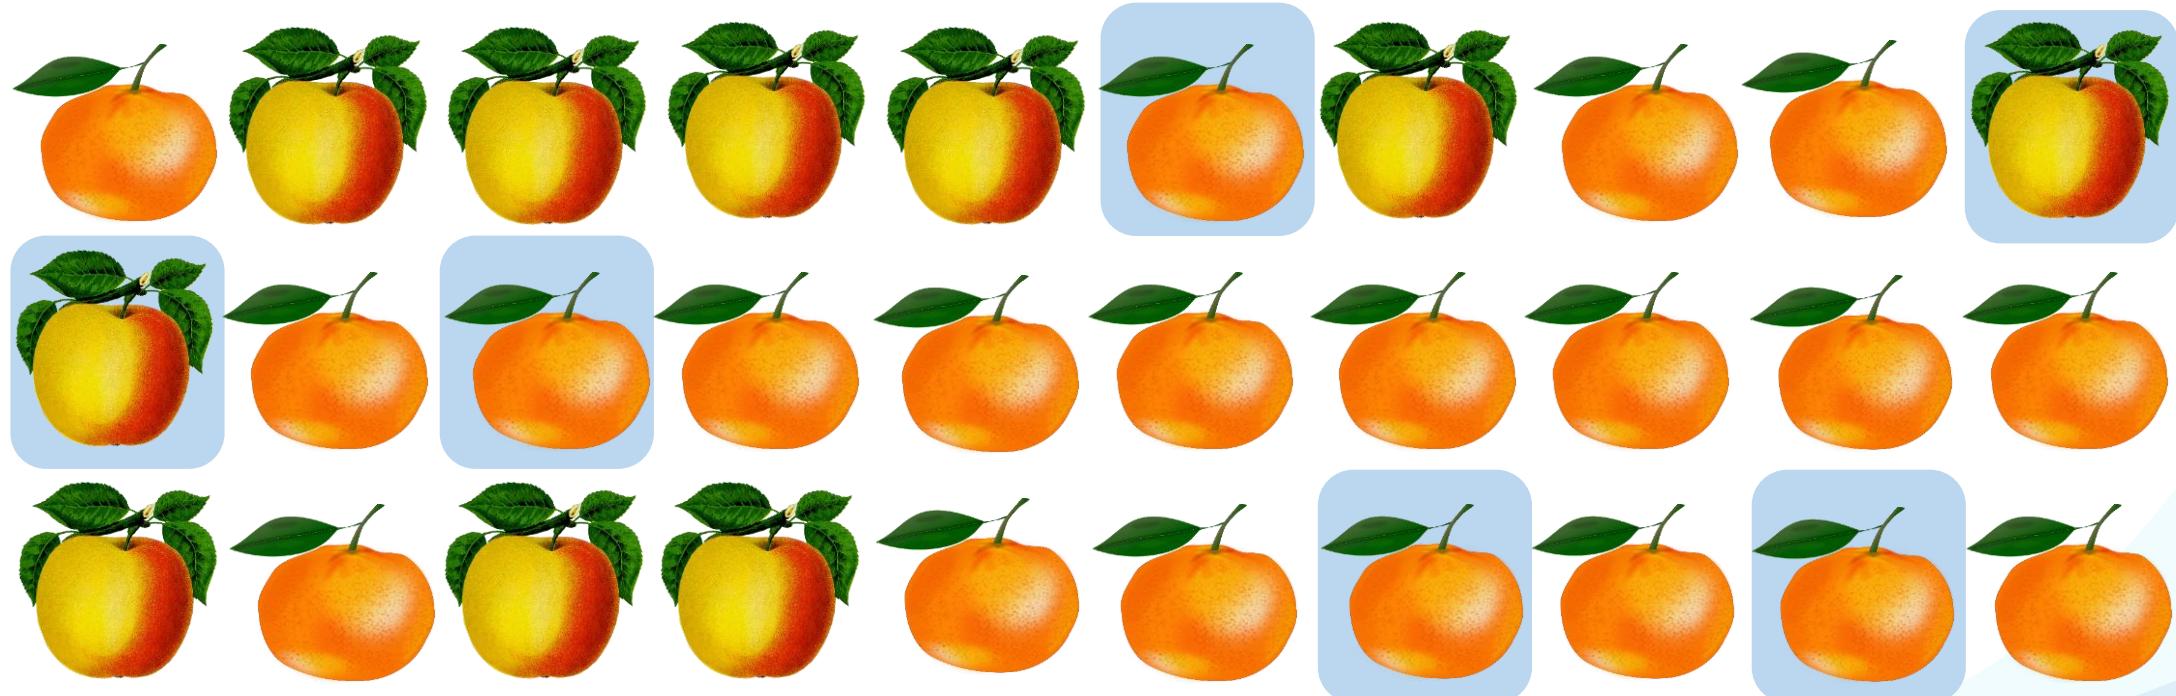


## Preprocessing – Sampling

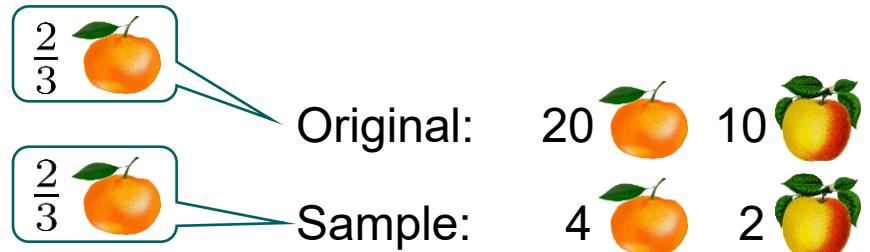


Random sampling:

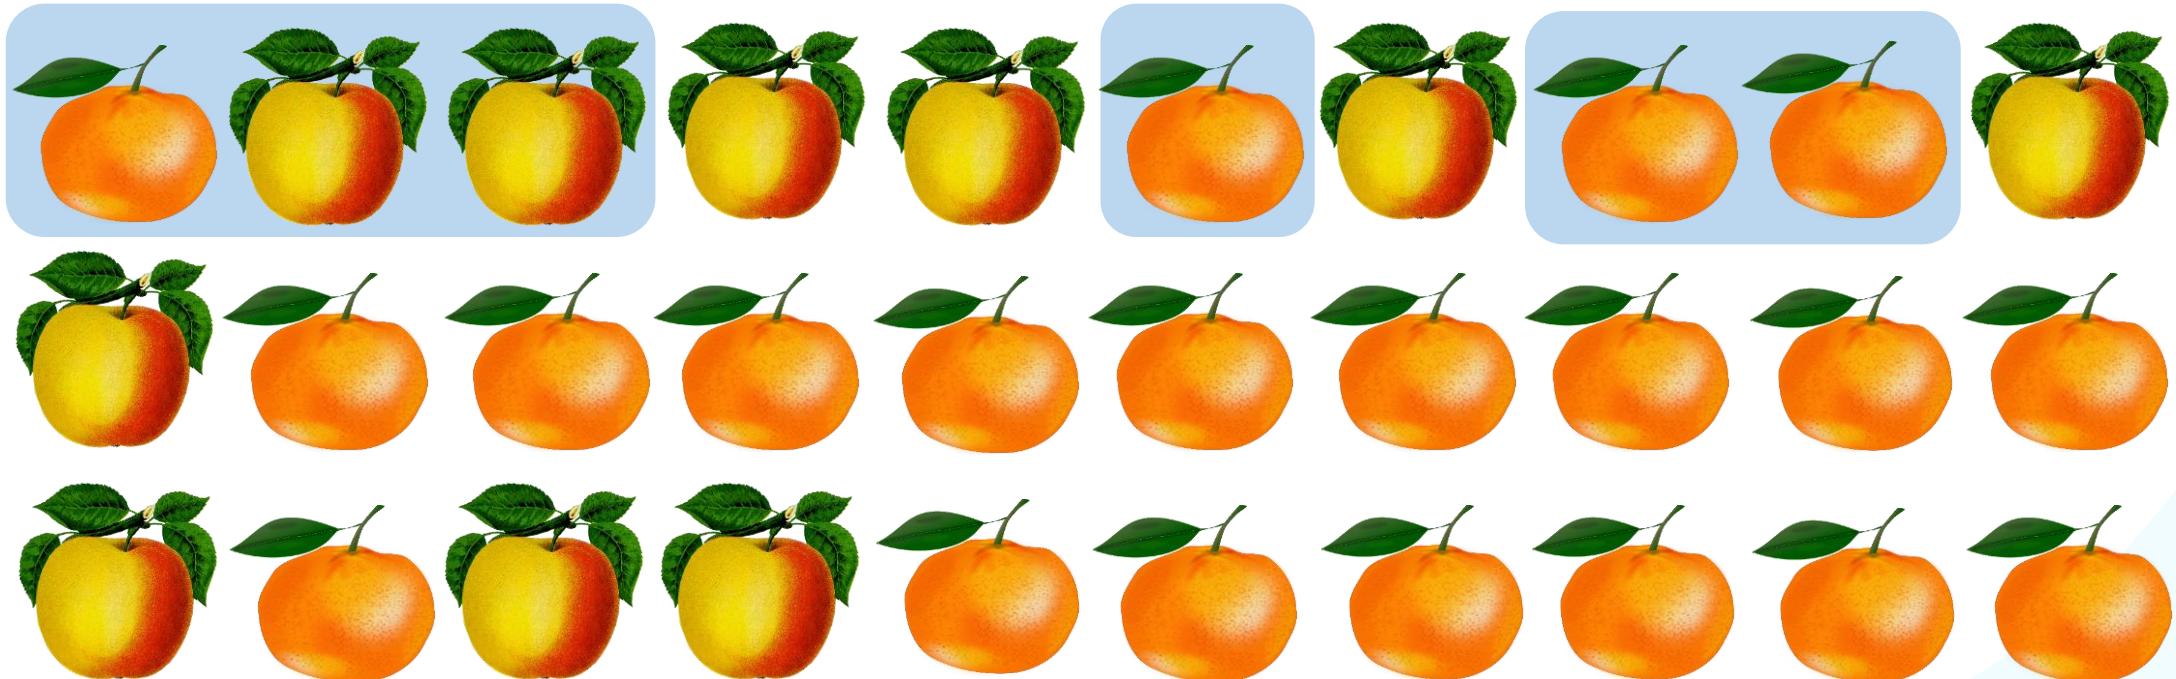
take  $N$  arbitrary instances (based on random generator)



## Preprocessing – Sampling



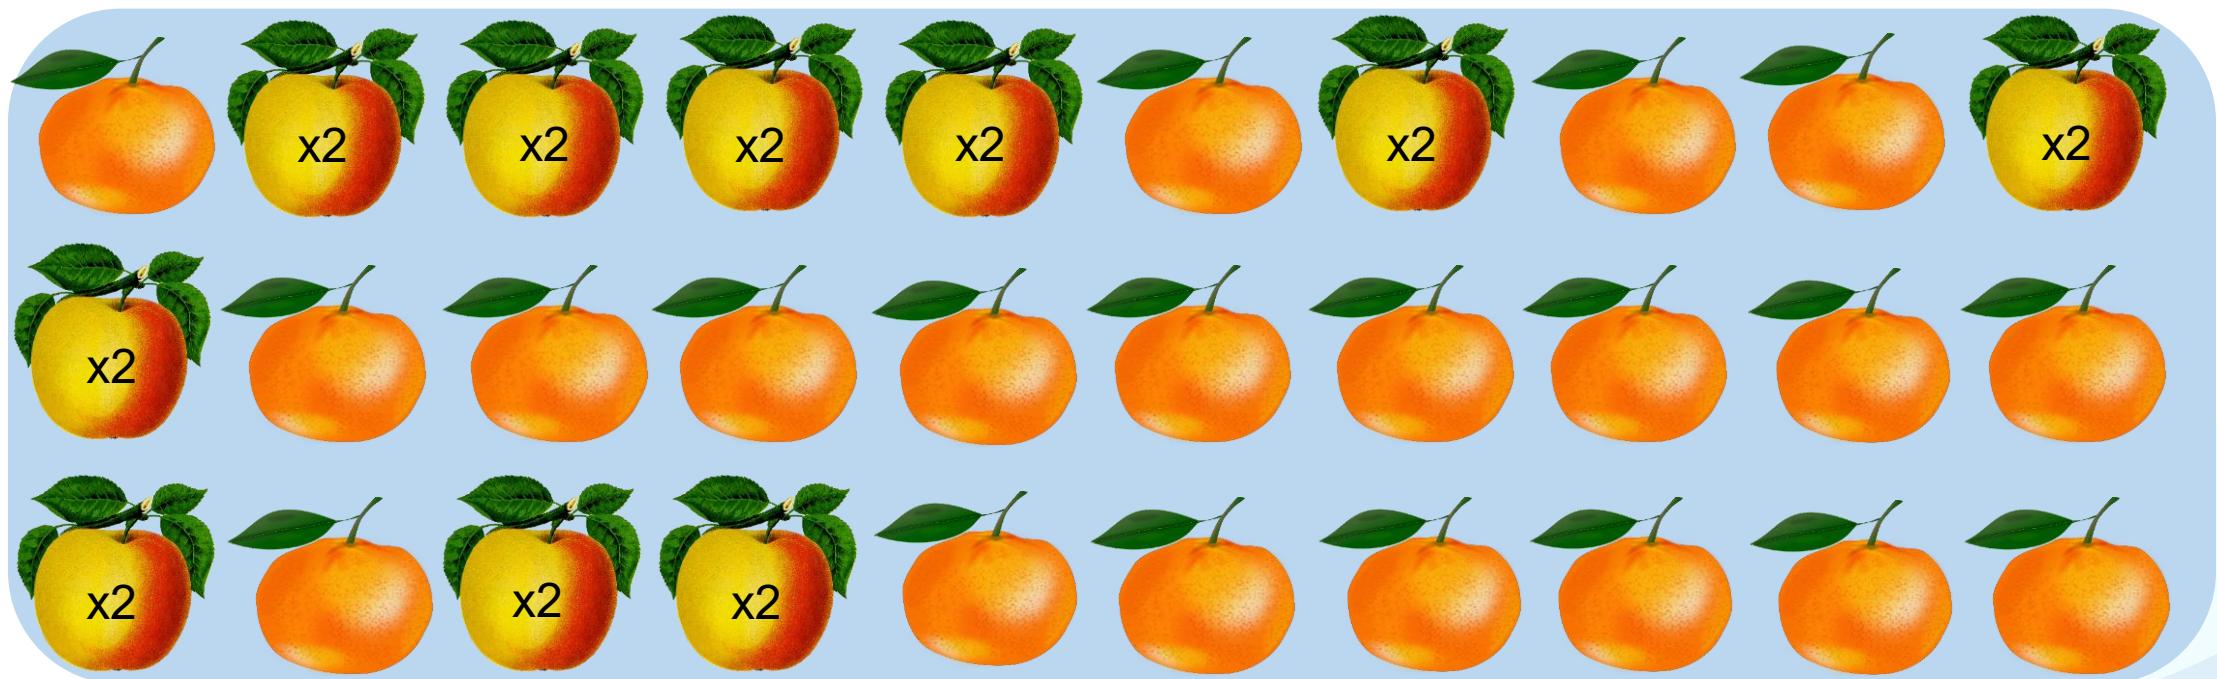
**Stratified sampling:** ensure that relative frequencies are maintained (e.g., take the same percentage from every group)



## Preprocessing – Sampling



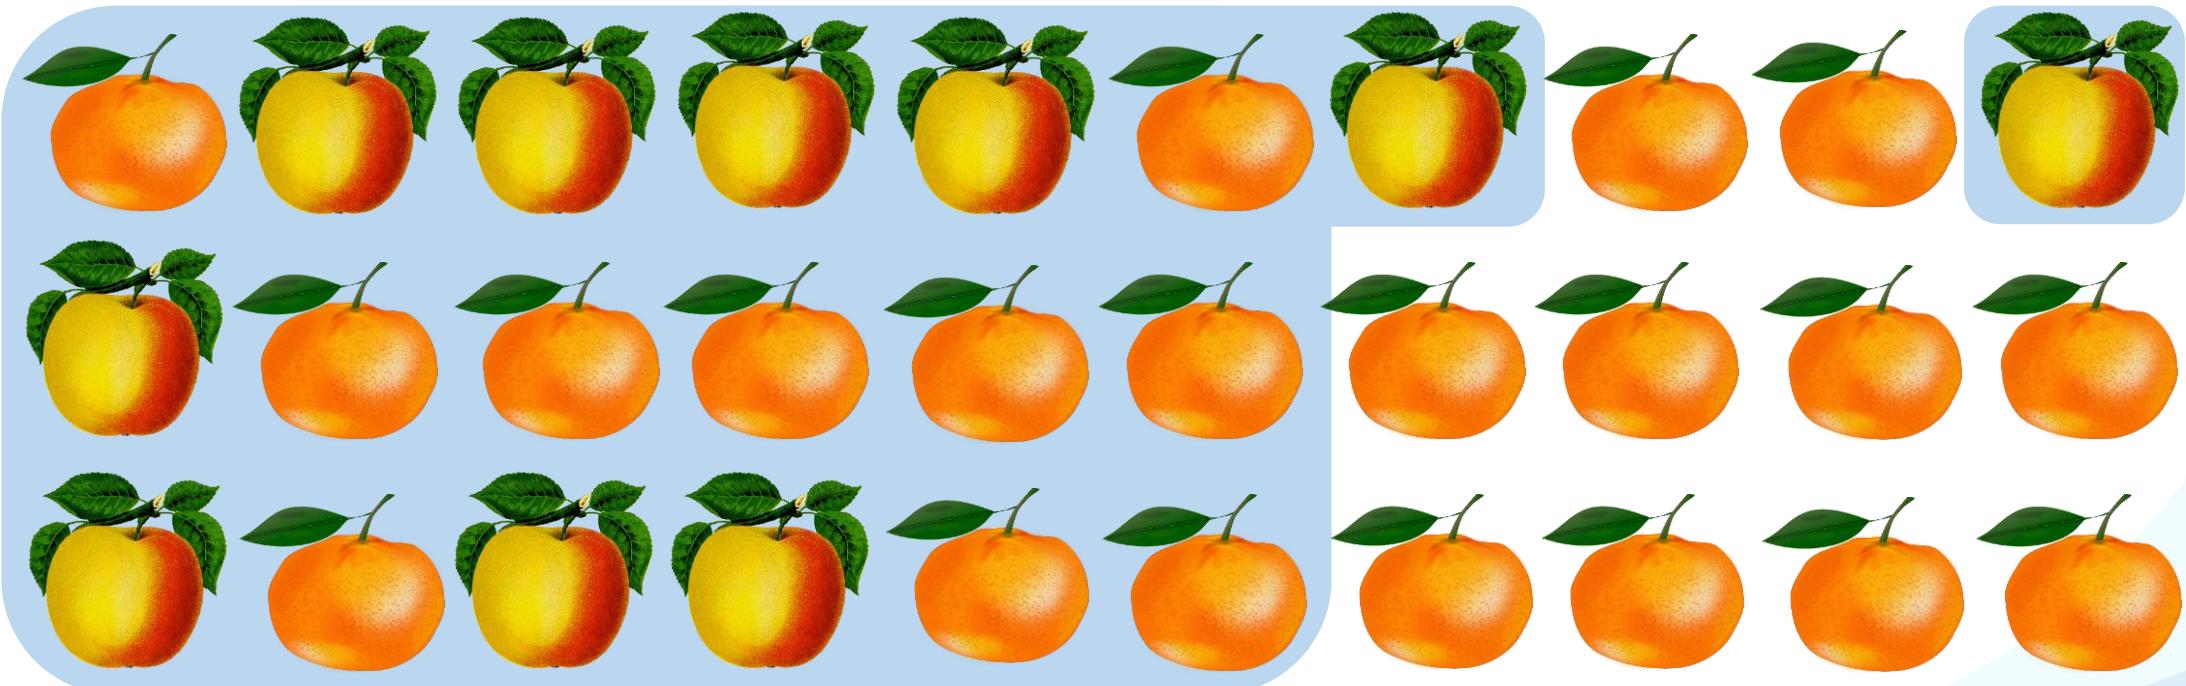
**Over-sampling:** ensure a certain distribution  
(e.g. equal frequency for each group) by duplicating under-represented instances



## Preprocessing – Sampling



**Under-sampling:** ensure a certain distribution  
(e.g. equal frequency for each group) by leaving out over-represented instances



## To Conclude



**Goal:** increase data quality and modify the data to suit the analysis question and applied techniques

**Best strategy/solution:** depends on the data, context and goal of the analysis

### Data quality aspects

- Missing data
- Noise/outliers
- Semantic problems

Garbage in, Garbage out (GIGO)

### Data preprocessing

- Transformation
- Normalization
- Data reduction



80/20